
„Systemnahe Programmierung“ (BA-INF 034) Wintersemester 2023/2024


Dr. Matthias Frank, Dr. Matthias Wübbeling

Institut für Informatik 4
Universität Bonn

E-Mail: *{matthew, matthias.wuebbeling}@cs.uni-bonn.de*
Sprechstunde: nach Vereinbarung

2. Fortgeschrittene Konzepte der Systemprogrammierung

2. Betriebssysteme/Threads
„Kommunikation innerhalb von
Prozessen bzw. zwischen
Prozessen eines Rechners“



Teil 2



[2.1. Processes](#)

[2.2. Threads](#)

[2.3. Interprocess Communication IPC](#)

[...](#)

Literature

- Mark Mitchell, Jeffrey Oldham, and Alex Samuel. Advanced Linux Programming. New Riders Publishing. First edition, 2001. Chapter 4. <http://www.advancedlinuxprogramming.com/>
- Abraham Silberschatz, Peter Baer Galvin, und Greg Gagne. Operating System Concepts. John Wiley & Sons. Eighth edition, 2008. Chapters 4 and 7.
- Andrew S. Tanenbaum. Modern Operating Systems. Prentice Hall. Third edition, 2007. Chapters 2 and 3.
- W. Richard Stevens, Stephen A. Rago. Advanced Programming in the UNIX Environment. Addison-Wesley. Second Edition, 2005. Chapter 11.
- Blaise Barney. POSIX Threads Programming Tutorial. Lawrence Livermore National Laboratory. <https://computing.llnl.gov/tutorials/pthreads/>



2.2. Threads – Outline

■ 2.2.1. Fundamentals

- Threads in Linux
- Thread Synchronization
- Deadlocks
- Important Threading Mechanisms



2.2.1. Motivation

- A process – as discussed so far – is a program executed sequentially with a single thread of control
- Processes allow parallelism...
 - Different applications are executed in parallel (separate processes)
 - A single application can consist of several (cooperating) processes
- ... but: Processes are **independent** execution units
 - Communication among processes is relatively complex
 - Switching context from process to process is expensive
- Idea: Support multiple threads of control within a single process

more on communication between
processes in subsection 2.3. IPC

Threads (1)

- Threads provide **multiple execution flows** within one process
 - Threads operate within the environment of a process
 - All threads run independently, executing in parallel
- Threads are **less independent** than processes
 - Share the same address space
 - Can access the same global variables
- **No built-in protection** between threads
 - Data is shared among threads
 - Threads can access (and manipulate!) each others stack!
 - A thread might interfere with the execution of other threads

Threads (2)


- All **threads share**
 - Address space
 - Set of open files
 - Set of child processes
 - Sets of alarms and signals
- **Each thread needs a separate**
 - Stack – Stores local variables of the functions in execution
 - Program counter – “Pointer” to the next command in program code
 - Registers – Data currently used by the processor when executing the thread
 - State – Stores whether the thread is executing, blocked, ...



Benefits of Using Threads

- **Parallel execution** of multiple activities
 - Sharing of processor time
 - Utilization of the processor while some activities are blocked
- Utilization of **multiprocessor / multicore systems**
 - Multiple threads of control can run truly in parallel on multiprocessor or multicore architectures
- **Responsiveness**
 - Long-running computations in a background thread do not block the main application thread
- **Resource sharing**
 - Threads operate in the same address space and can therefore easily share resources

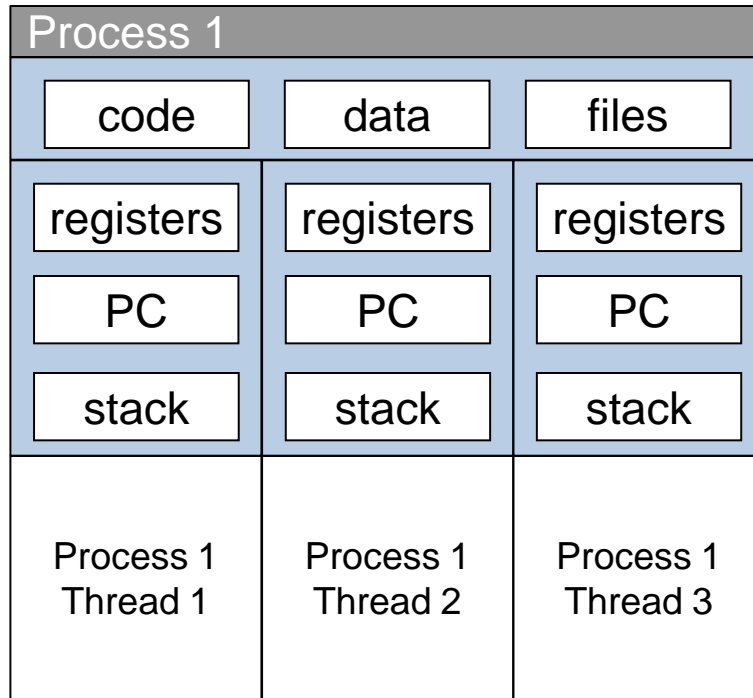
Threads – Example Scenario



- User interface thread:
 - **Should remain responsive** even while the application processes data
 - Might be idle over longer periods of time
- Processing thread
 - Responsible for **time-consuming computations**
 - Requires a lot of CPU time
- Communication thread:
 - Sending, receiving and processing messages
 - **Blocked** while accessing external resources


Threads vs. Processes (1)

Thread model



- All threads within a process share a **common address space**
- **Fast** context switches
- **Efficient** communication and cooperation

Process Model



- Processes are **isolated from each other**
- Communication and cooperation is **complex**

Threads vs. Processes (2)

- **Advantages** over Processes
 - Low context switching overhead
 - Shared state
 - Efficient communication and cooperation among threads
 - Easier to program?
- **Disadvantages**
 - Shared state: Threads can interfere with each others operation
 - Errors in one thread (e.g., a crash) can affect all others
 - More difficult to synchronize?
- Due to the similarities of threads and processes, threads are often called **lightweight processes**

Performance comparison – processes vs. threads

- „ ... the following table compares timing results for the *fork()* subroutine and the *pthread_create()* subroutine.“

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
AMD 2.4 GHz Opteron (8cpus/node)	41.07	60.08	9.01	0.66	0.19	0.43
IBM 1.9 GHz POWER5 p5-575 (8cpus/node)	64.24	30.78	27.68	1.75	0.69	1.10
IBM 1.5 GHz POWER4 (8cpus/node)	104.05	48.64	47.21	2.01	1.00	1.52
INTEL 2.4 GHz Xeon (2 cpus/node)	54.95	1.54	20.78	1.64	0.67	0.90
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.54	1.07	22.22	2.03	1.26	0.67

- „Timings reflect 50,000 process/thread creations, were performed with the time utility, and units are in seconds, no optimization flags.“

Table and *quotations* from:

Blaise Barney. POSIX Threads Programming Tutorial. Lawrence Livermore National Laboratory.

<https://computing.llnl.gov/tutorials/threads/>

Operating System Support for Threads

- **Threading support** can be provided on the
 - User level → User threads
 - Kernel level → Kernel threads
- **Multithreading** on the user level can be implemented following different models
 - Many-to-One (cf. slides 15 ff.)
 - One-to-One
 - Many-to-Many

User Threads


- Implemented by a thread library at the user level
 - Provides support for creating, managing and scheduling without support from the kernel
- Kernel is unaware of threads
 - Sees only a single thread of execution – the process
- **Advantages**
 - Fast to create and manage (No kernel involvement required)
- **Disadvantages**
 - Cannot take advantage of multiprocessors or multicores
 - Blocking system call in one thread blocks all threads

Kernel Threads

- Implemented as part of the core operating system
- Kernel responsible for creation, scheduling and management
- **Advantages**
 - Able to take advantage of multiprocessors or multicores
 - Kernel can schedule another thread upon blocking on a system call in one thread
- **Disadvantages**
 - Slower to create and manage than user threads

Multithreading Models (1)


- Many-to-One
 - Many user-level threads are mapped to a single kernel-level thread
 - Has all the advantages and disadvantages of user threads
 - Example: GNU Portable Threads



```
int main(void) {  
    int thread = 0;  
    while(1) {  
        if thread == 0 {  
            thread0();  
            thread = 1;  
        }  
        else {  
            thread1();  
            thread = 0;  
        }  
    }  
}
```



Multithreading Models (2)

- One-to-One
 - Maps each user-level thread to a kernel-level thread
 - Has all the advantages and disadvantages of kernel threads
 - Example: **Linux, Windows, Solaris**




Multithreading Models (3)

- Many-to-Many
 - Multiplexing many user-level threads to many kernel-level threads
 - Number of kernel-level threads \leq number of user-level threads
 - Mapping decision based on various criteria:
 - Application properties
 - Hardware properties (e.g., number of cores)
→ Most flexible model
 - Example: Older versions of Solaris
 - Golang revives N:M mapping



Overview

- Fundamentals
- **2.2.2. Threads in Linux**
- Thread Synchronization
- Deadlocks
- Important Threading Mechanisms



2.2.2. Implementation of Threads in Linux (1)

- On the **operating system level**, Linux does **not really distinguish between processes and threads**
 - Uses the term task to cover both
 - Difference: Threads share their address space with each other
- A new thread is created using the `clone()` system call
 - Generates a new child process just like the `fork()` system call
 - Allows the child process to share parts of its execution context with the calling process
 - Flags allow to configure what is to be shared, e.g., the memory space, the file descriptor table, the signal handlers, ...

Implementation of Threads in Linux (2)

- Effects
 - Threads in Linux are usually based on the **one-to-one multithreading model**
 - Originally, each thread not only had a different thread ID but also a different process ID – this has changed in newer kernel versions...
- **What does this mean for the programmer?**
 - `clone()` system call is usually hidden behind a **user-level thread implementation like Pthreads**
 - Usually: **No need to worry about the implementation** – but **always good to be aware of how it is mapped internally**.

Implementation of Threads in Linux (3)

- (Recently) new possibility: C11 (this is NOT C++11 !)
 - Features the new `<threads.h>`
 - APIs very similar to pthread
 - Support for: Threading, Mutual Exclusion, Condition variables, Thread-local storage
 - Link with API: <https://en.cppreference.com/w/c/thread>
 - **Support in glibc** since Version 2.28 but...
 - Glibc 2.3 supports “Native POSIX Thread Library” providing pthread_* syscalls in Linux.

POSIX Threads

- **POSIX:** “Portable Operating System Interface” for software compatible with variants of the UNIX operating system
 - Family of standards from the area of operating systems
 - Administered by the IEEE (Institute of Electrical and Electronics Engineers)
- POSIX Threads (or **PThreads**)
 - Standard programming interface for threads
 - Defines procedures for creating, manipulating and destroying threads
 - Implementations exist for various operating systems (albeit not all fully standards compliant)

POSIX Threads in Linux

- Implementations provided by the GNU C library on Linux
 - Originally: [LinuxThreads](#) (Now obsolete but still available)
 - Since the Linux 2.6 kernel: [NPTL](#) (Native POSIX Threads Library)
 - Note: The description in the “Advanced Linux Programming” book is still based on Linux Threads!
- Functions and data types are declared in `<pthread.h>`
- Included in the library `libpthread`
 - Add `-lpthread` as a parameter to your linker call

PThreads Overview

- Tasks
 - Creating threads
 - Managing thread attributes
 - Passing parameters to threads
 - Exiting threads
 - Joining threads
 - Cancelling threads

→ Explained in detail on the following slides

PThread Thread Creation

- **Create a new thread**, specify which function it should execute and initiate the execution of the new thread:

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start_routine)(void*), void *arg);
```

- **thread** – Data structure representing the created thread (= thread identifier)
 - **attr** – The attributes of the new thread. If set to NULL, then the default attributes will be used.
 - **start_routine** – Pointer to the thread function executed by the newly created thread
 - **arg** – Arguments passed to the start routine
- Return value: 0 if OK, error code otherwise
 - PThread library executes clone for each create internally

PThread Thread Function

- Initial function executed by a new thread
 - Might call any other function
 - Returning from this function is one way of terminating the thread.
- Function of type: `void* threadFunctionName (void*)`
 - Any function of this type qualifies as a thread function
 - Allows any type for both the argument and the return value → Use casting for accessing thread-specific data
- A function pointer of type `void* (*) (void*)` must be provided to create a new thread
 - New thread executes the thread function specified by the function pointer

PThread Thread Creation – Example

```
#include <pthread.h>
#include <stdio.h>

pthread_t workThreadID;

void *workThreadFunction(void *thread_arg) {
    printf("Your work thread says: Hello world!\n");
    return ((void *) 0);
}

int main(void) {
    pthread_create(&workThreadID, NULL, workThreadFunction, NULL);
    printf("Your main thread says: Hello world!\n");
    sleep(2);
    return 0;
}
```

declaration of thread function



creation of thread
(no attributes, no args)



PThread Thread Attributes (1)

- Can be used to **configure the fundamental behavior** of a new thread
- Initializing
 - `int pthread_attr_init(pthread_attr_t *attr);`
 - Initializes attr with the default value for all attributes
- Destroying
 - `int pthread_attr_destroy(pthread_attr_t *attr);`
- The available set of thread attributes is implementation dependent.
- Accessing attributes (set & get)
 - `int pthread_attr_setAttrName (pthread_attr_t *attr, AttrType t);`
 - `int pthread_attr_getAttrName (pthread_attr_t *attr, AttrType *t);`

PThread Thread Attributes (2)

- **Detach state** (exact semantics will be explained later – cf. slide 34 ff.)
 - `int pthread_attr_setdetachstate (pthread_attr_t *attr, int detachstate);`
 - `int pthread_attr_getdetachstate(const pthread_attr_t *attr, int *detachstate);`
 - `detachstate` – either `PTHREAD_CREATE_DETACHED` or `PTHREAD_CREATE_JOINABLE`
- **Thread Scheduling**
 - `pthread_attr_setschedpolicy / pthread_attr_getschedpolicy`
 - `pthread_attr_setscope / pthread_attr_getscope`
 - `pthread_attr_setinheritsched / pthread_attr_getinheritsched`
- **Stack size and stack address**
 - `pthread_attr_setstacksize / pthread_attr_getstacksize`
 - `pthread_attr_setstackaddr / pthread_attr_getstackaddr`

PThread Thread Attributes – Example

- Creating a thread in the detached state:

```
int main(void) {  
    pthread_attr_t wtAttr;  
  
    pthread_attr_init(&wtAttr);  
    pthread_attr_setdetachstate(&wtAttr, PTHREAD_CREATE_DETACHED);  
    pthread_create(&workThreadID, &wtAttr, workThreadFunction, NULL);  
    pthread_attr_destroy(&wtAttr);  
  
    [...]  
}
```

- Remark: Changes to wtAttr after calling pthread_create have no effect on the thread

PThread Thread Parameter Passing

Function of type: `void* threadFunctionName (void*)`

- Useful to **pass data to a newly created thread**
- Each thread function expects exactly one parameter of type `void*`
 - Flexible type of the argument
 - Limitation: **Only a single parameter possible**
- Solution:
 - Define a **structure** for each thread that **contains the expected parameters**
 - Pass a **pointer** to the structure **as the single argument** to the thread
- Be careful: **Consider the lifetime of variables** passed to a thread using pointers

PThread Thread Parameter Passing – Example

```
struct wtParams {  
    int param1;  
    char param2;  
};  
  
void* workThreadFunction(void* thread_arg) {  
    struct wtParams* wtp = (struct wtParams*) thread_arg;  
    printf("Param 1: %d\n", wtp->param1);  
    printf("Param 2: %c\n", wtp->param2);  
    return ((void*) 0);  
}  
  
int main(void) {  
    struct wtParams args;  
    args.param1 = 4;  
    args.param2 = 'x';  
    pthread_create(&workThreadID, NULL, workThreadFunction, &args);  
    [...]  
}
```

argument passed and used
as a pointer!

in this case: pointer
to global struct

- Be careful: **Consider the lifetime of variables** passed to a thread using pointers (do not do this at home)!



PThread Ending / Exiting a Thread ...

... three ways to do this!

1. Implicitly

- Returning from the thread function (see example on prev. slide)
- Return value of the thread function is used as return value of the thread

2. Explicitly using `pthread_exit`

- `void pthread_exit(void *value_ptr);`
 - `value_ptr` – The return value of the thread
- Terminates the calling thread and makes `value_ptr` available as the return value of the thread

3. Thread can be **cancelled by another thread** in the same process

- See below

PThread Joining a Thread

- Problem: How to determine when a thread is finished?
 - Example: Main thread should wait until all other threads are done before ending the program (Note: ugly workaround using sleep in the first example – slide 28).
 - Second example: Thread needs return value of a child thread it created previously.
- We need a function that blocks until a thread is finished and then collects the return value of the thread: `pthread_join`
- `int pthread_join(pthread_t thread, void **value_ptr);`
 - `thread` – Identifier of the target thread
 - `value_ptr` – The result value passed to `pthread_exit` by the terminating thread

PThread Joining a Thread – Example

```
void* workThreadFunction(void* thread_arg) {  
    int answer = 42;  
    return ((void*) answer);  
}  
  
int main(void) {  
    int threadResult;  
    pthread_create(&workThreadID, NULL, workThreadFunction, NULL);  
    pthread_join(workThreadID, (void*) &threadResult);  
    printf("Thread result: %d\n", threadResult);  
}
```

return value is integer,
but casted to pointer type!

result available in
integer variable

- Be careful again: **Consider lifetime of variables** returned as result!
 - Do not return pointer to data placed on the stack of the thread!
 - Make clear who is responsible for freeing dynamic memory.
 - Do not try this at home!?

PThread Thread Detach State (1)

- A thread can be in one of two states: **joinable** (default) or **detached**
 - Using the detachstate attribute at thread creation time
 - Calling pthread_detach to detach a running thread
 - `int pthread_detach(pthread_t thread);`
 - **Inverse operation (reattaching a detached thread) is not possible!**
- **Joinable threads**
 - Not automatically cleaned up when the thread terminates
 - Requires another thread to call `pthread_join` to clean up the exit state of the thread
 - Allows **collecting result data from the thread**

PThread Thread Detach State (2)

- Detached threads
 - Immediately cleaned up when the thread terminates
 - Impossible to synchronize on their completion
 - No retrieval of result data possible
 - Only useful if
 - Thread is **completely independent** or
 - Thread **communicates** its (result) data to its peers **differently** (e.g., using shared state)

PThread C Cancelling Threads (1)

- Up to now: Every thread runs to completion
 - Returning from its thread function
 - Calling `pthread_exit`
- There can be a need to stop a thread prematurely, for example:
 - All threads must be stopped prior to exiting the program.
 - The user cancelled the operation.
 - The results of the computations of a thread are not required anymore.
- Solution: Cancelling threads
 - `int pthread_cancel(pthread_t thread);`
- Note: Later we discuss whether this is a good solution...

PThread Cancelling Threads (2)

- Each **thread manages** its own **cancelability**
 - `int pthread_setcancelstate(int state, int *oldstate);`
 - **Enable cancelling** by setting state to `PTHREAD_CANCEL_ENABLE`
 - **Disable cancelling** by setting state to `PTHREAD_CANCEL_DISABLE`
 - `oldstate` returns the previous cancel state of the thread
 - Attempts to cancel a thread are ignored if cancelling is disabled
- **Two cancelling types** are possible
 - `int pthread_setcanceltype(int type, int *oldtype);`
 - Set to **deferred (synchronous)** by setting type to `PTHREAD_CANCEL_DEFERRED`
 - Set to **asynchronous cancelable** by setting type to `PTHREAD_CANCEL_ASYNCHRONOUS`
 - `oldtype` returns the previous cancel type of the thread

PThread Cancelling Threads (3)

- **Asynchronously** cancelable
 - Thread can be cancelled at any point of its execution – irrespective of the operations it currently performs or the state it is in
- **Synchronously** cancelable (**DEFAULT**)
 - Thread can only be cancelled at specified cancellation points
 - Explicit cancellation points
 - Specified in the code by the programmer
 - Calling: `void pthread_testcancel(void);`
- Implicit cancellation points
 - Cancellation points are automatically added with other commands, for example `pthread_join`, `pthread_cond_wait`, `sem_wait`, ...

PThread Cancelling Threads – Example

```
void* workThreadFunction(void* thread_arg) {  
    pthread_setcancelstate(PTHREAD_CANCEL_ENABLE, NULL);  
    pthread_setcanceltype(PTHREAD_CANCEL_DEFERRED, NULL);  
    // Do fancy stuff...  
    pthread_testcancel();  
    // Do even more fancy stuff...  
}  
  
int main(void) {  
    pthread_create(&workThreadID, NULL, workThreadFunction, NULL);  
    pthread_cancel(workThreadID);  
    pthread_join(workThreadID, NULL);  
}
```

- Where is the thread cancelled?
 - Which parts have been executed when `pthread_cancel` is called?
 - What if “Do fancy stuff” or one of the functions called in there include an implicit cancellation point?

Cancelling Threads – Disadvantages

- **Asynchronous** canceling is **not applicable** in most scenarios
 - We have to deal with aborts at any point of the execution flow.
 - Might leave resources accessed by the thread in an inconsistent state
- **Synchronous** canceling is **difficult** to get it right
 - Are all implicit and explicit cancellation points correctly covered?
 - What do the sub routines do that might affect the cancelling behavior?
- **Recommendation:**
 - **Avoid canceling** threads in most application scenarios.
 - **Explicitly communicate** the request to exit to the thread **using a shared state variable** and react in the thread accordingly.

USE MULTITHREADING, THEY SAID



IT'S EASY, THEY SAID

Recap

- Threads as “lightweight processes”
- **Advantages**
 - Low context-switching overhead
 - Shared state and memory
 - Efficient communication among threads
- **Disadvantages**
 - Shared state and memory
 - If threads crash, the entire process is crashed
 - Synchronization
- Various types of threads
 - Kernel vs. user threads
 - 1:1, n:1, n:m mapping

Overview

- Fundamentals
- Threads in Linux
- **2.2.3. Thread Synchronization**
- Deadlocks
- Important Threading Mechanisms

2.2.4. Thread Synchronization Problem

- Threads execute *independently* of each other
 - No particular ordering of events is guaranteed
 - Each program execution can result in a different execution sequence
 - Possible on multiprocessor / multicore systems: Truly parallel execution of commands
- Threads are *not independent* of each other
 - Threads accessing the same data
 - One thread building upon the results of another thread
 - Threads requiring the same resources (e.g., a file handler)
 - ...

Race Conditions

- The **result** of a computation critically **depends on the sequence or the timing of events**
 - Multiple threads access a shared resource
 - Operations are performed in parallel / quasi parallel
 - Some proper order of operations is implicitly required but the system fails to enforce it



Race Conditions – Example (1)

- Adding integer values to a queue


```
struct listNode {  
    struct listNode* succ;  
    int value;  
};  
  
struct listNode* firstNode = NULL;  
  
void addValue(int value) {  
    struct listNode* newNode = (struct listNode*) malloc(sizeof(struct listNode));  
    newNode->value = value;  
    newNode->succ = firstNode;  
    firstNode = newNode;  
}
```

Thread 1: addValue(42);

Thread 2: addValue(13);


Race Conditions – Example (2)

Start situation:



Thread 1: Command 1
Thread 1: Command 2
Thread 1: Command 3
Thread 1: Command 4
Thread 2: Command 1
Thread 2: Command 2
Thread 2: Command 3
Thread 2: Command 4


Schedule 1



Sequence of commands

Thread 1: Command 1
Thread 1: Command 2
Thread 1: Command 3
Thread 2: Command 1
Thread 2: Command 2
Thread 2: Command 3
Thread 2: Command 4
Thread 1: Command 4

Schedule 2



Result

```

void addValue(int value) {
    struct listNode* newNode = (struct listNode*)
    malloc(sizeof(struct listNode));
    newNode->value = value;
    newNode->succ = firstNode;
    firstNode = newNode;
}
  
```

Critical Sections (also called: Critical Regions) (1)

- Definition: A **critical section** is a part of the code that accesses shared resources that must not be concurrently accessed by more than one thread of execution.
 - Example: Command 3+4 of the function `addValue` together form a critical section
- Insufficient definition: A critical section is a piece of code that can only be executed by one thread at a time
 - Other parts of the code might access the same resource → All these code parts together belong to the critical section.
 - We have to protect the resources not the code sections!
- We need **synchronization primitives** to prevent the concurrent access to critical sections

Critical Sections (also called: Critical Regions) (2)

- Tanenbaum's **four conditions for a good solution** of the critical section problem:
 1. **No two threads** may be inside their **critical section simultaneously**.
 2. **No assumptions** are to be made about **speeds or the number of CPUs**.
 3. **No thread** running outside its critical region **may block other threads**.
 4. **No thread** should have to **wait forever** to enter its critical region.

Thread Synchronization

- Goals
 - Prevent concurrent access to critical sections
 - Provide control over the ordering of relevant system events
 - Approaches
 - 2.2.3.1. **Mutex** variables
 - 2.2.3.2. **Condition** variables
 - 2.2.3.3. **Semaphores**
 - 2.2.3.4. **Barriers**
 - 2.2.3.5. **Read-write locks**
- Explained in detail on the following slides

} **Primary means of implementing
thread synchronization**

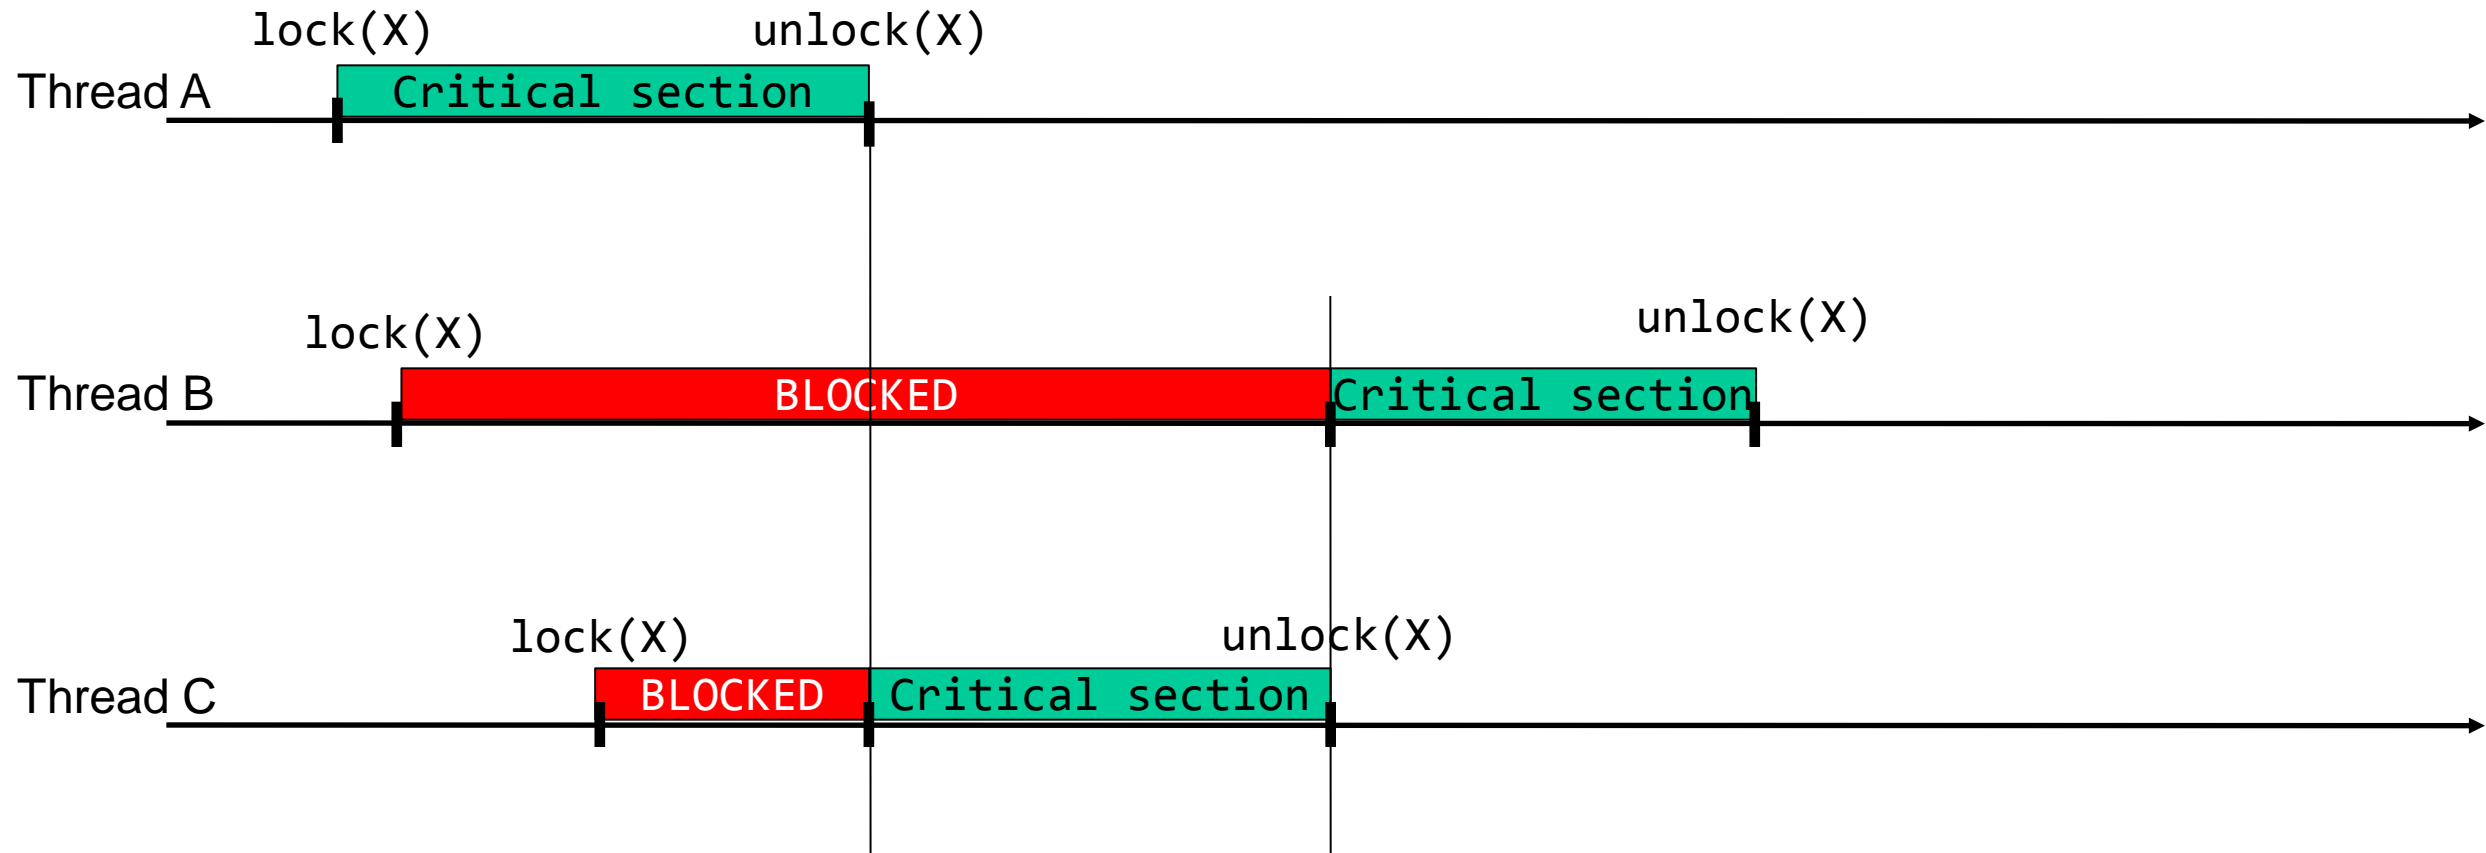
2.2.3.1. Mutex Variables – Motivation

- **Mutual Exclusion**
- **Goal:** Protect a critical section against simultaneous access by two or more threads
 - Thread is only allowed to enter if no other thread is in the critical section
 - Once a thread leaves a critical section, it should be accessible to other threads again

Mutex Variables – Basic Concept

- Critical regions are **protected with** the help of **locks**
 - Thread must acquire lock before entering the critical region
 - Thread releases lock upon leaving the critical region
- **Invariant: At any given time, only one thread can have a lock on a mutex variable**
- Thread A **requests a lock on mutex X**. Two cases:
 - No thread holds a lock on X: A acquires the lock and is able to proceed
 - Another thread already holds a lock on X: A is blocked
- Thread A **releases its lock on mutex X**
 - Unblock one of the threads waiting for X (if such a thread exists)

Mutex Variables – Example



Note: No FIFO guarantees provided!

PThread Mutexes

- Data structures
 - `pthread_mutex_t` – Mutex variable data type
 - `pthread_mutexattr_t` – Mutex attributes data type
- Initializing a mutex variable
 - Dynamically
 - `int pthread_mutex_init (pthread_mutex_t *mutex, pthread_mutexattr_t *attr);`
 - Statically at declaration time
 - `pthread_mutex_t mutexVar = PTHREAD_MUTEX_INITIALIZER;`
- Destroying a mutex variable
 - `int pthread_mutex_destroy (pthread_mutex_t *mutex);`

PThread Mutex Attributes

- Initializing
 - `int pthread_mutexattr_init(pthread_mutexattr_t *attr);`
 - Initializes attr with the default value for all attributes
- Destroying
 - `int pthread_mutexattr_destroy(pthread_mutexattr_t *attr);`
- Important mutex attribute: Setting / getting the type of a mutex
 - `int pthread_mutexattr_settype (pthread_mutexattr_t *attr, int type);`
 - `int pthread_mutexattr_gettype (const pthread_mutexattr_t *attr, int *type);`
 - Example:
 - `pthread_mutexattr_settype(&attr, PTHREAD_MUTEX_RECURSIVE);`

PThread Mutex Types

- `PTHREAD_MUTEX_NORMAL`
 - Attempting to lock a mutex twice from the same thread results in a deadlock; behavior when unlocking unlocked mutexes undefined
 - Fast mutex – no checks or counters required
- `PTHREAD_MUTEX_ERRORCHECK`
 - Relocking a locked mutex or unlocking an unlocked mutex results in an error
- `PTHREAD_MUTEX_RECURSIVE`
 - Multiple locks from the same thread possible – requires the same number of calls to `pthread_mutex_unlock`
- `PTHREAD_MUTEX_DEFAULT`
 - Recursive locking and unlocking error behavior implementation dependent

PThread Locking and Unlocking Mutexes

- Acquiring a lock on a mutex
 - `int pthread_mutex_lock(pthread_mutex_t *mutex);`
 - Blocks the thread until the lock can be acquired
 - `int pthread_mutex_trylock(pthread_mutex_t *mutex);`
 - Never blocks the thread but indicates the success of the locking attempt in its return value
 - 0 – Mutex has been locked successfully.
 - EBUSY – Mutex has not been locked, because it was already locked.
- Releasing a lock on a mutex
 - `int pthread_mutex_unlock(pthread_mutex_t *mutex);`

PThread Mutexes – Example

- Revisiting the race condition example:
- Only one thread can access the marked code (*) at a time – others are blocked
- Prevents race condition in adding values to the queue
- No guarantees which thread can add its node first!

```
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;

void addValue(int value) {
    struct listNode* newNode = (struct listNode*) malloc(sizeof(struct listNode));
    newNode->value = value;
    pthread_mutex_lock(&mutex);
    (*) newNode->succ = firstNode;
    (*) firstNode = newNode;
    pthread_mutex_unlock(&mutex);
}
```

2.2.3.2. Condition Variables – Motivation

- Thread might need to **wait for an event or a certain condition** before proceeding
 - Thread can be put to sleep while waiting (with others continuing their work)
 - Thread should be notified once the condition is fulfilled
- Examples:
 - Waiting for a computation being started by the user
 - Waiting for data being added to a linked list
 - Consumer thread waiting for data input prepared by a producer thread



First Idea: Busy Waiting (Spinning)

- Goal: Waiting for an event using only mutexes
- Idea: Check repeatedly to see if a condition is true

```
void* waitingThread(void* thread_arg) {  
    int localFlag = 0;  
    while(!localFlag) {  
        pthread_mutex_lock(&threadFlagMutex);  
        localFlag = threadFlag;  
        pthread_mutex_unlock(&threadFlagMutex);  
    }  
    // Condition fulfilled - Do work!  
}
```

```
void setFlag(int newFlagValue) {  
    pthread_mutex_lock(&threadFlagMutex);  
    threadFlag = newFlagValue;  
    pthread_mutex_unlock(&threadFlagMutex);  
}
```

Main disadvantage:

Consumes processor time
without performing any work!



Condition Variables – Basic Concept

- A thread can **wait on a condition variable**
- A waiting thread is **blocked until another thread signals the same condition variable** either
 - Waking a single thread or
 - Waking all waiting threads by broadcasting the signal
- **Condition variables** must always be **used together with a mutex**
 - Lock mutex before waiting on condition variable
 - Waiting on the condition variable implicitly unlocks the mutex
 - Lock on the mutex is reacquired automatically when the thread is unblocked
 - Acquire lock on mutex before signalling the condition variable

Condition Variables – Example

Consumer thread 1

lock(X) wait(Y)



Producer thread

lock(X) signal(Y) unlock(X)



Consumer thread 2

lock(X) wait(Y)



unlock(X)



X ... Mutex variable
Y ... Condition variable



PThread Condition Variables

- Data structures
 - `pthread_cond_t` – Condition variable data type
 - `pthread_condattr_t` – Condition attributes data type
- Initializing a condition variable
 - Dynamically
 - `int pthread_cond_init(pthread_cond_t *cond, const pthread_condattr_t *attr);`
 - Statically at declaration time
 - `pthread_cond_t condVar = PTHREAD_COND_INITIALIZER;`
- Destroying a condition variable
 - `int pthread_cond_destroy(pthread_cond_t *cond);`

PThread Condition Variable Attributes

- Allows to configure the properties of a new condition variable
- Initializing
 - `int pthread_condattr_init(pthread_condattr_t *attr);`
 - Initializes attr with the default value for all attributes
- Destroying
 - `int pthread_condattr_destroy(pthread_condattr_t *attr);`
- There is **usually no need to set any condition variable attributes to special** values

PThread Waiting On a Condition Variable

- `int pthread_cond_wait (pthread_cond_t *cond, pthread_mutex_t *mutex);`
 - `cond` – Condition variable to wait on
 - `mutex` – Mutex variable associated with `cond`
 - Precondition: `mutex` must be locked by the calling thread
 - Effect: Atomically releases `mutex` and blocks the calling thread on the condition variable `cond` until `cond` “is signalled”
 - Postcondition: `mutex` is again locked by the calling thread

```
pthread_mutex_lock(&mutex);  
pthread_cond_wait(&cond, &mutex);  
// Signal has been received  
// Do actual work here...  
pthread_mutex_unlock(&mutex);  
// ... or here
```

PThread Condition Variable Timed Wait (1)

- `int pthread_cond_timedwait (pthread_cond_t *cond, pthread_mutex_t *mutex, const struct timespec *abstime);`
 - `abstime` – Absolute time value limiting how long the calling thread can remain blocked waiting for the signal
- Behavior identical to `pthread_cond_wait`, except
 - Wait time for the signal is limited
 - Returns `ETIMEDOUT` if `abstime` is passed without cond being signaled
- Note: `pthread_cond_timedwait` expects an **absolute time value** not the length of a time interval!



PThread Condition Variable Timed Wait (2)

- Example:

```
int result;
struct timespec tspec;
struct timeval tval;

pthread_mutex_lock(&mutex);
result = gettimeofday(&tval, NULL);
tspec.tv_sec = tval.tv_sec;
tspec.tv_nsec = tval.tv_usec * 1000;
tspec.tv_sec += 5;
result = pthread_cond_timedwait(&cond, &mutex, &tspec);
if (result == ETIMEDOUT) {
    printf("Timeout!\n");
} else {
    printf("Signal received!\n");
}
pthread_mutex_unlock(&mutex);
```

- Waits for a maximum of 5 seconds
- But: Can remain blocked longer as it needs to reacquire the lock on the mutex.

PThread Signaling on a Condition Variable

- `int pthread_cond_signal(pthread_cond_t *cond);`
 - Effect: Unblocks **at least one** of the threads waiting for cond (if there are threads waiting)
- `int pthread_cond_broadcast(pthread_cond_t *cond);`
 - Effect: Unblocks all threads waiting for cond
- For predictable scheduling behavior:
 - Lock mutex variable mutex before signalling on a condition variable and release lock on mutex afterwards.
 - Signaled threads remain blocked until lock on mutex is released again.
 - Only one of the signaled threads at a time gets the lock on mutex.

PThread Condition Variables – Waking Threads (1)

- Problem
 - pthread_cond_signal unblocks **AT LEAST** one thread
 - Specification allows “spurious wakeups” from pthread_cond_wait and pthread_cond_timedwait
 - Returning from pthread_cond_wait or pthread_cond_timedwait **DOES NOT IMPLY ANYTHING** about the value of the predicate
- Solution
 - Always reevaluate the predicate upon returning from a wait or a timed wait

PThread Condition Variables – Waking Threads (2)

- Example

```
pthread_mutex_lock(&threadFlagMutex);  
while(!threadFlag) {  
    pthread_cond_wait(&threadFlagCondition, &threadFlagMutex);  
}  
pthread_mutex_unlock(&threadFlagMutex);  
// Condition fulfilled - Do work!
```

- Remark: Also covers the case where an event is signaled before the thread waits on the condition variable
- Works similarly for `pthread_cond_timedwait`
 - Using absolute time values is actually helpful here!
 - Idea: Only repeat if “result != ETIMEDOUT”

2.2.3.3. Semaphores – Motivation

- Some applications need to **limit access** to sections **based on a counter**, for example:
 - Limited number of N resources available → Only proceed if not all of them are in use.
 - Count the number of entries in a queue → Only proceed if the queue is not empty.
- Goal: **Implement a shared counter** that controls the access to critical sections
- Generalization of the **mutex** approach which can be seen as a **binary semaphore**

Semaphores – Basic Concept

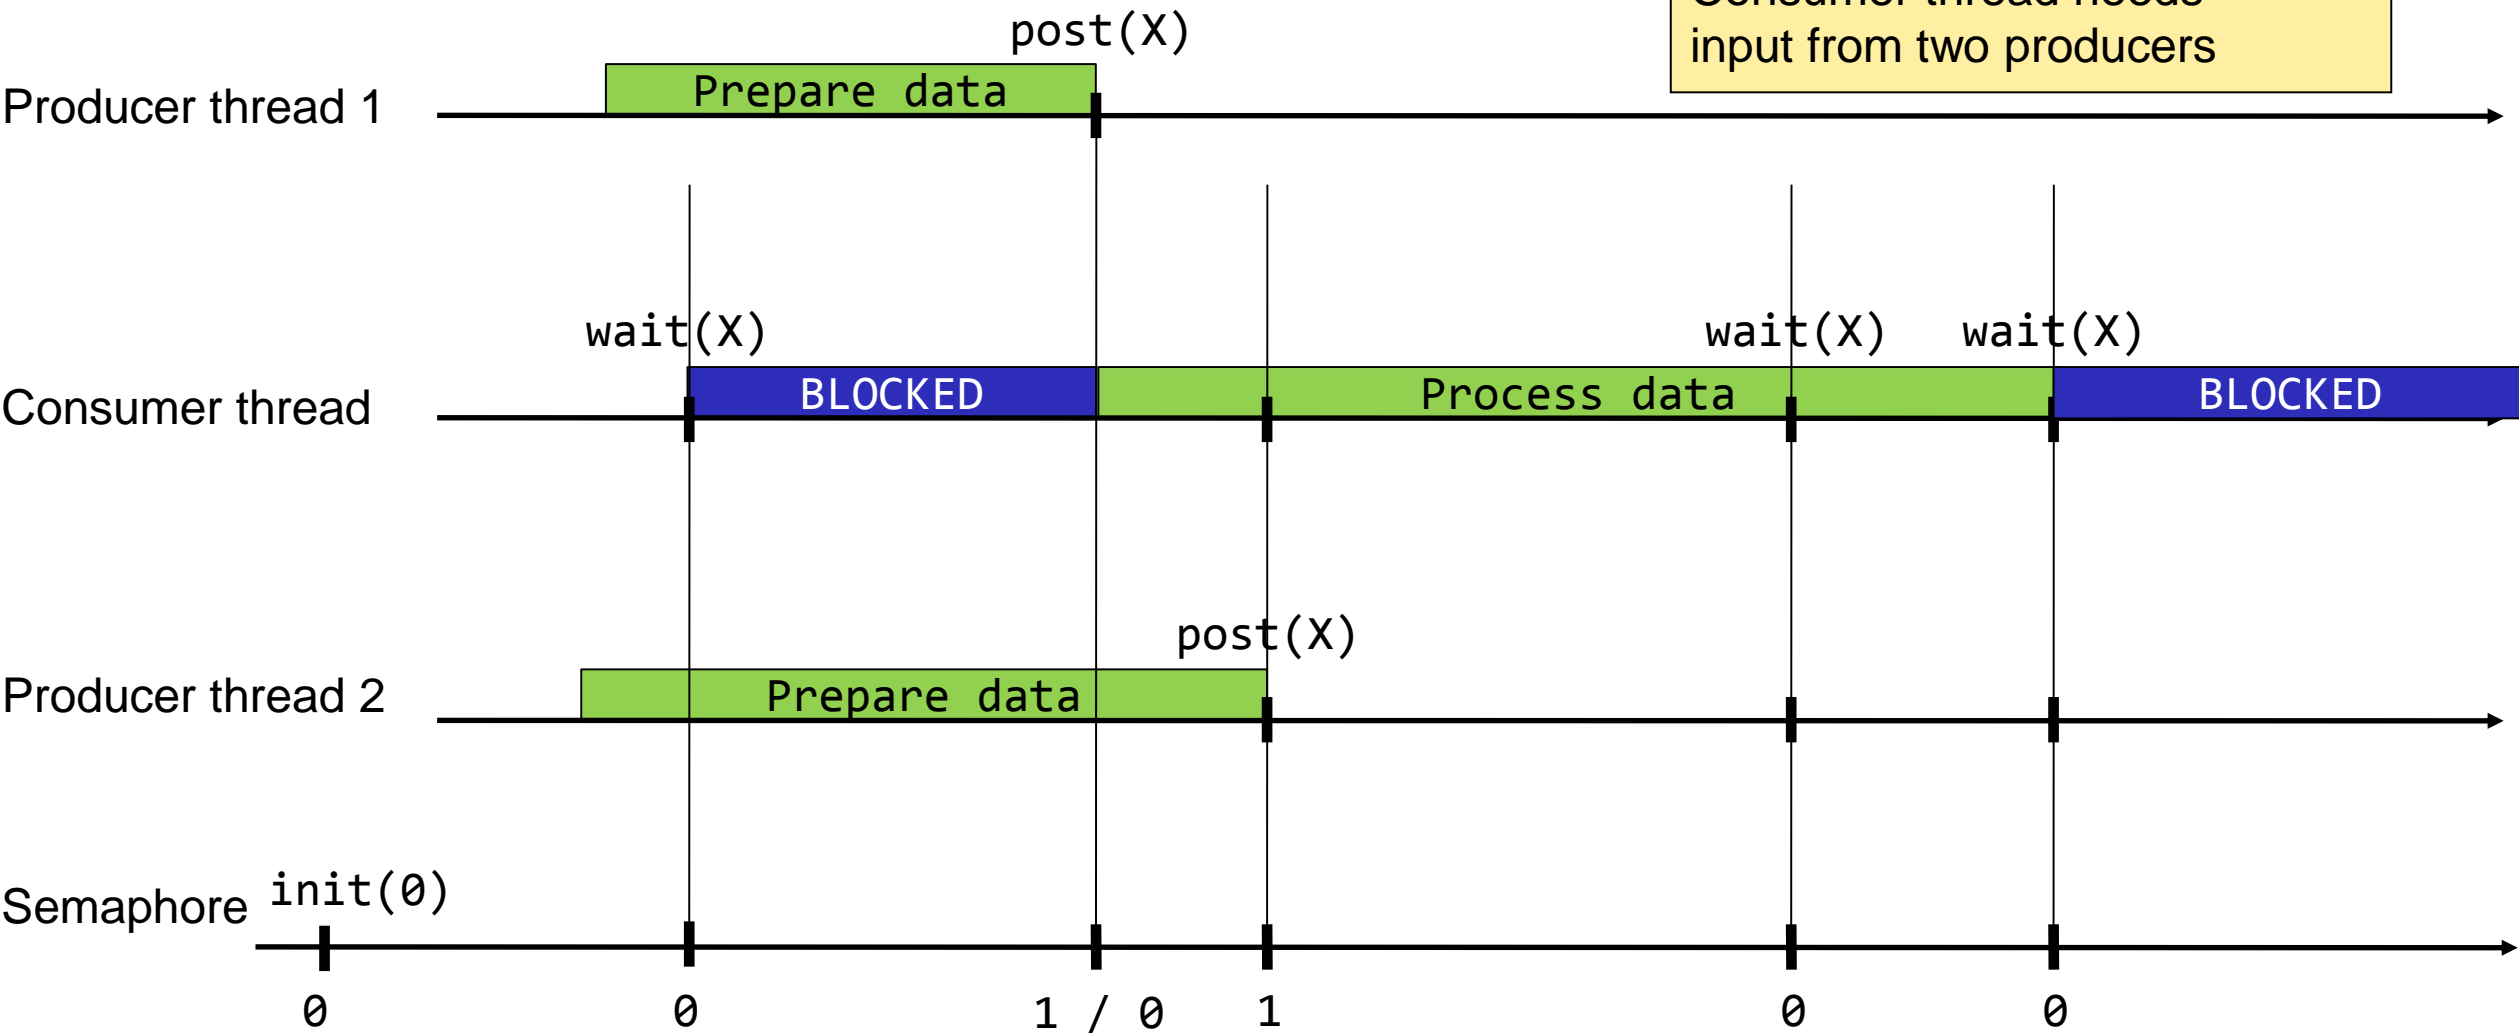
- Use a **counter to synchronize threads**
 - Value > 0 : Thread is allowed to proceed
 - Value $= 0$: Thread is blocked
- **Wait operation**
 - If counter is zero, then block the calling thread
(the block will be released when the semaphore value is > 0)
 - Counter value is decremented by 1
 - wait returns
- **Post operation**
 - Increment counter value by 1
 - Unblock one of the waiting threads (if any) if the counter was zero before

Semaphores – Application Examples

- **Mutual exclusion** = **Mutex cf. 2.2.3.1.**
 - Initialize counter with 1
 - Process entering critical section calls wait operation
 - Process leaving critical section calls post operation
- Variants of the mutual exclusion problem
 - E.g., no more than X threads are allowed to enter a critical section at a time.
- Consumer waiting for input (see next slide)
- ...

Semaphores – Example

Example application:
Consumer thread needs
input from two producers



PThread Semaphores

- Requires: `#include <semaphore.h>`
- Data structures
 - `sem_t` – Semaphore data type
- **Initializing** a semaphore
 - `int sem_init(sem_t *sem, int pshared, unsigned int value);`
 - `sem` – The semaphore
 - `pshared` – Indicates whether the semaphore is shared among processes
 - `value` – The initial value of the semaphore
- **Destroying** a semaphore
 - `int sem_destroy(sem_t *sem);`

PThread Operating on Semaphores

- **Waiting** on a semaphore
 - `int sem_wait (sem_t *sem);`
 - Blocks the thread until the lock on the semaphore can be acquired
 - `int sem_trywait (sem_t *sem);`
 - Never blocks the thread but indicates the success of the locking attempt in its return value
 - `0` – Semaphore has been locked successfully.
 - `EAGAIN` – Semaphore has not been locked, because it was already locked.
- **Releasing** a lock on a semaphore
 - `int sem_post(sem_t *sem);`

PThread Semaphores Example

- Using a semaphore to **implement mutex** functionality

```
void* workThread(void* thread_arg) {  
    sem_wait(&workThreadSem);  
    // Critical section  
    sem_post(&workThreadSem);  
}  
  
void init() {  
    sem_init(&workThreadSem, 0, 1);  
    [...]  
}
```

- Counter of the semaphore is always zero while a thread is within the critical section.
→ No other thread is able to enter until the thread calls `sem_post` which sets the semaphore counter back to one.

Semaphores

Brinch Hansen (1973): “The semaphore is an elegant synchronizing tool for an ideal programmer who never makes mistakes. But unfortunately of using semaphores incorrectly can be quite serious“

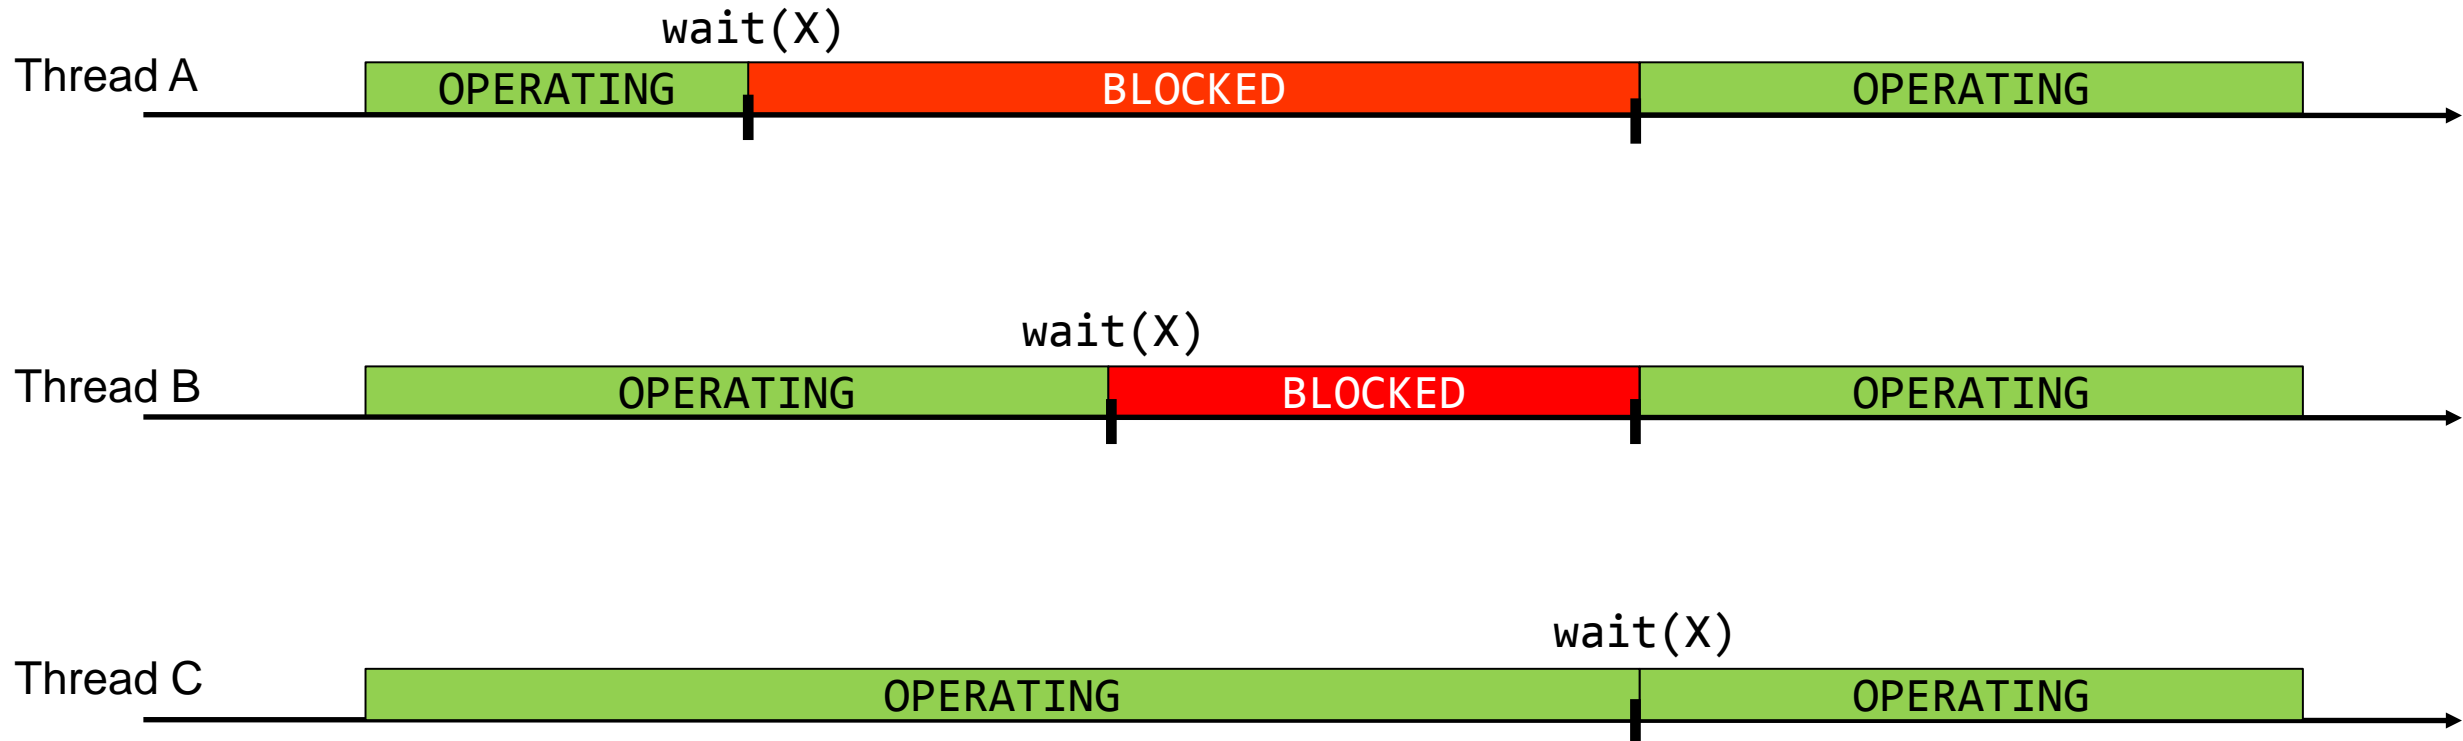
2.2.3.4. Barriers – Motivation

- Threads need to “meet up” at certain points of the program execution
 - Only when all threads have arrived at this point, the execution of the individual threads can continue.
 - Threads “wait for each other”
- Example applications:
 - Threads operating on different parts of a large problem synchronizing their (intermediate) results

Barriers – Basic Concept

- Barrier is initialized with the number of threads participating n_{Barr}
- Barrier manages a counter to record the number of threads already waiting at the barrier
- Wait operation
 - Increment counter: $\text{counter}++$
 - If ($\text{counter} < n_{\text{Barr}}$) then block calling thread
 - If ($\text{counter} == n_{\text{Barr}}$) then unblock all threads and re-initialise counter with 0

Barriers – Example



`init nbarr = 3`



PThread Barriers (1)

- Data structures
 - `pthread_barrier_t` – Barrier data type
 - `pthread_barrierattr_t` – Barrier attributes data type
- **Initializing** a barrier
 - `int pthread_barrier_init(pthread_barrier_t *restrict barrier, const pthread_barrierattr_t *restrict attr, unsigned count);`
 - **count** – The **number of threads** that must call `pthread_barrier_wait` for the threads to become unlocked
- **Destroying** a barrier
 - `int pthread_barrier_destroy(pthread_barrier_t *barrier);`

PThread Barriers (2)

- Waiting on the barrier
 - `int pthread_barrier_wait(pthread_barrier_t *);`
- Example

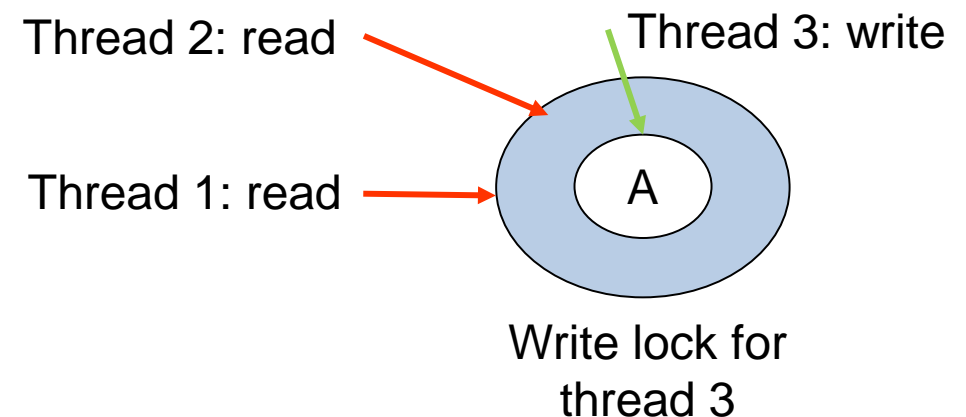

thread(s) waiting here

```
void* workThread(void* thread_arg) {  
    // Do a lot of work here...  
    pthread_barrier_wait(&allThreadsFinishedBarrier);  
    printf("All threads finished!\n");  
}  
  
void init() {  
    int i;  
    pthread_barrier_init(&allThreadsFinishedBarrier, NULL, 5)  
    for (i=0; i<5; i++) {  
        pthread_create(&threadID[i], NULL, workThreadFunction, NULL);  
    }  
    [...]  
}
```

init with 5

2.2.3.5. Read-Write Locks

- Motivation:
 - Reading data concurrently from multiple threads is not harmful.
 - Only **write operations can interfere** with concurrent read or write operations
- **Idea:** Treat locking requests for reading shared data differently than locking requests for writing data



Read-Write Locks – Basic Concept

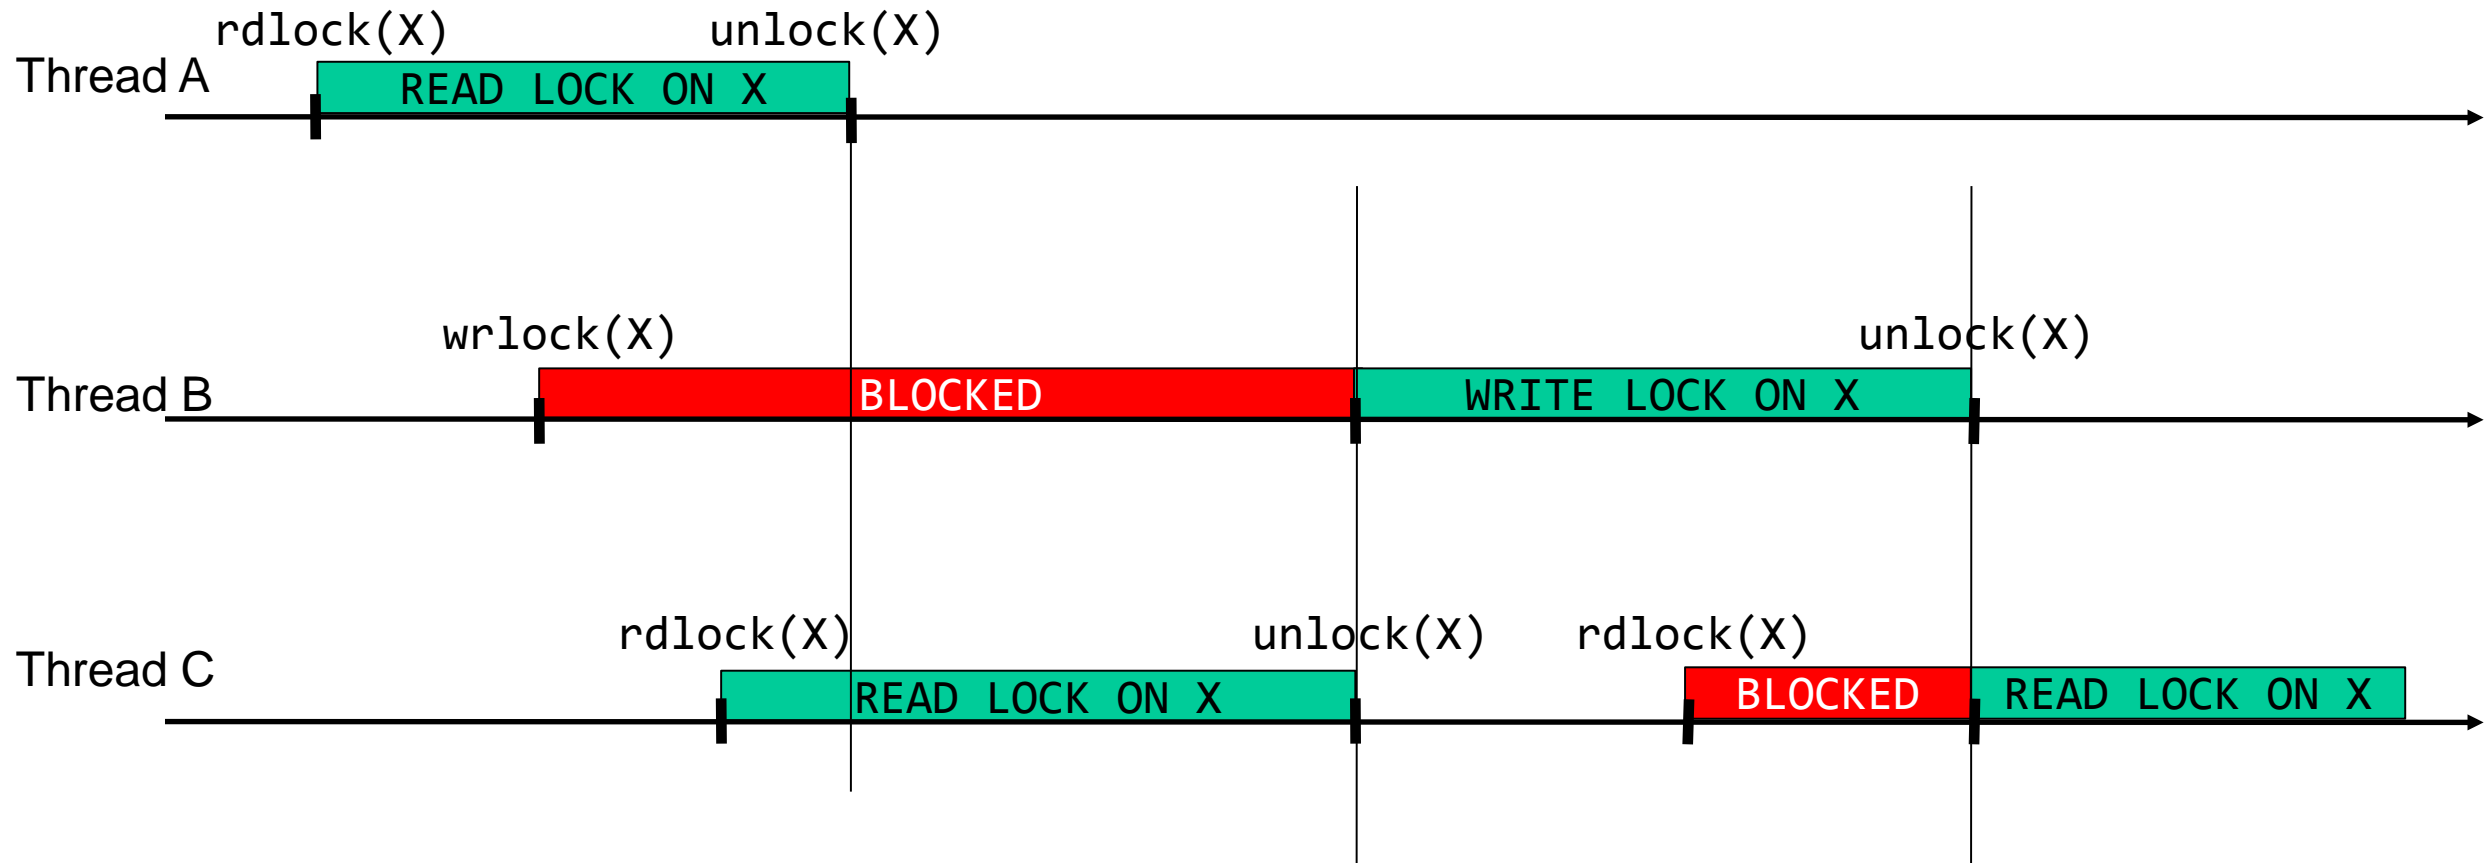
- Servicing a **request for a read** lock on X
 - if no other thread holds a lock on X
 - if other threads only hold read locks on X
 - Otherwise: Requesting thread is blocked

- Servicing a **request for a write** lock on X
 - only if no other thread holds a lock (of any kind) on X
 - Otherwise: Requesting thread is blocked

		Lock set		
Lock requested		Null	Read	Write
	Read	+	+	-
	Write	+	-	-

+ request will acquire lock
 - request will block thread

Read-Write Lock – Example



PThread Read-Write Lock

- Data structures
 - `pthread_rwlock_t` – Read-write lock data type
 - `pthread_rwlockattr_t` – Read-write lock attributes data type
- **Initializing** a read-write lock
 - Dynamically
 - `int pthread_rwlock_init (pthread_rwlock_t *rwlock, const pthread_rwlockattr_t *attr);`
 - Statically at declaration time
 - `pthread_rwlock_t rwlock = PTHREAD_RWLOCK_INITIALIZER;`
- **Destroying** a read-write lock
 - `int pthread_rwlock_destroy(pthread_rwlock_t *rwlock);`

PThread Read-Write Lock – Locking and Unlocking

- Acquiring a read lock
 - `int pthread_rwlock_rdlock (pthread_rwlock_t *rwlock);`
 - `int pthread_rwlock_tryrdlock(pthread_rwlock_t *rwlock);`
- Acquiring a write lock
 - `int pthread_rwlock_wrlock(pthread_rwlock_t *rwlock);`
 - `int pthread_rwlock_trywrlock(pthread_rwlock_t *rwlock);`
- Releasing a lock
 - `int pthread_rwlock_unlock(pthread_rwlock_t *rwlock);`
 - No different functions required for releasing read and write locks!

Waiting for Threads

- PThreads:
 - pthread_join for joinable threads
 - Missing join produces zombies
 - Manual management for detached threads
 - Global variables
 - Counters
 - ...

Thread communication

- PThread
 - Return codes
 - Pass by reference
 - Global variables
- Problem:
 - Requires synchronization

Share by communication

- Communicate results instead of using shared memory
 - Used in functional programming
 - Unless it's easier to solve with Mutexes
 - In most cases channels are easier than sync methods
 - Very useful in concurrent network development
- Shared data-structures still require traditional synchronization methods
 - Dynamic Lists
 - Arrays
 - Hash-maps

Channels

- Channels for communication between routines
 - Read on one end
 - Write on an other end
 - Block on full buffer
 - Thread-safe




src: <http://golangtutorials.blogspot.de/>

Overview

- Fundamentals
- Threads in Linux
- Thread Synchronization
- **2.2.4. Deadlocks**
- Important Threading Mechanisms

2.2.4. Deadlocks

- A **deadlock** is a situation when **two or more threads** (or processes) are **waiting for the other to finish some action** (e.g., accessing a resource, signaling an event) **thereby blocking each others progress**.



Deadlock Conditions


1. **Mutual exclusion:** Resources are exclusive – they cannot be accessed by more than one thread at a time.
2. **Non preemption:** Resources can only be released voluntarily by the holding thread and cannot be forcibly removed.
3. **Hold & wait:** Threads can hold resources while waiting for other resources.
4. **Circular wait:** There exists a closed loop of threads in which each thread waits for a resource held by its successor in the loop.

All four conditions **must be fulfilled for a deadlock to occur!**

Wait-For Graphs (1)

Resource-Allocation Graphs

- Two types of **nodes**:
 - Processes or threads are represented by **circles**
 - Resources are represented by **rectangles**
- Edges** only possible between different types of nodes
 - Process → Resource**: Process is waiting for resource
 - Resource → Process**: Resource is held by process
- Example:



Cycle in the graph → Deadlock

Wait-For Graphs (2)


Process Wait-For Graphs

- Nodes are processes
- Edge $P1 \rightarrow P2$: Process P1 waits for Process P2
- Example:



Cycle in the graph → Deadlock

- Resource-allocation graphs are **easily converted** to process wait-for graphs:



Handling Deadlocks (1)

- **Deadlock prevention** (more also on slide 105)
 - Idea: Ensure that one of the necessary deadlock conditions cannot hold
 - Achieve this by constraining how processes / threads can request resources
- **Deadlock avoidance**
 - Supervisor (e.g., the OS) needs information on which resources processes / threads are going to request
 - Decides which resource requests can be satisfied and which have to be delayed

Handling Deadlocks (2)

- **Deadlock detection**

- **Assumption:** Deadlock situations can occur
- **Strategy:** (Periodically) examine the state of the system to detect deadlock situations
- Recover from deadlocks (e.g., by terminating some processes / threads)

- **Suspecting deadlocks**

- Like deadlock detection but already act when a deadlock is suspected
- Allows for less intricate detection mechanisms

Deadlock Prevention

- Some approaches:
 - Allow threads to lock only one resource at a time → Resource must be released before another resource is requested (**Hold & wait condition**)
 - Force threads to release all held resources if they request another resource that cannot be served immediately (**No preemption condition**)
 - Impose a total ordering of resources → Resources can only be requested in this order (**Circular wait condition**)
- Applicability depends on the system properties
- **Deadlock prevention can be expensive!**

Deadlocks – Relevance for the Programmer

- Thread synchronization can provide for the first three deadlock conditions
 - Danger of deadlocks is a real problem!
 - Not using thread synchronization is usually not an option.
- Multithreaded applications are prime candidates for the deadlock problem as multiple threads compete for resources
- Programmer needs to be aware of the deadlock problem
 - Careful programming can avoid many potential deadlock situations
 - Deadlock prevention is most effectively and efficiently implemented with application knowledge

Overview

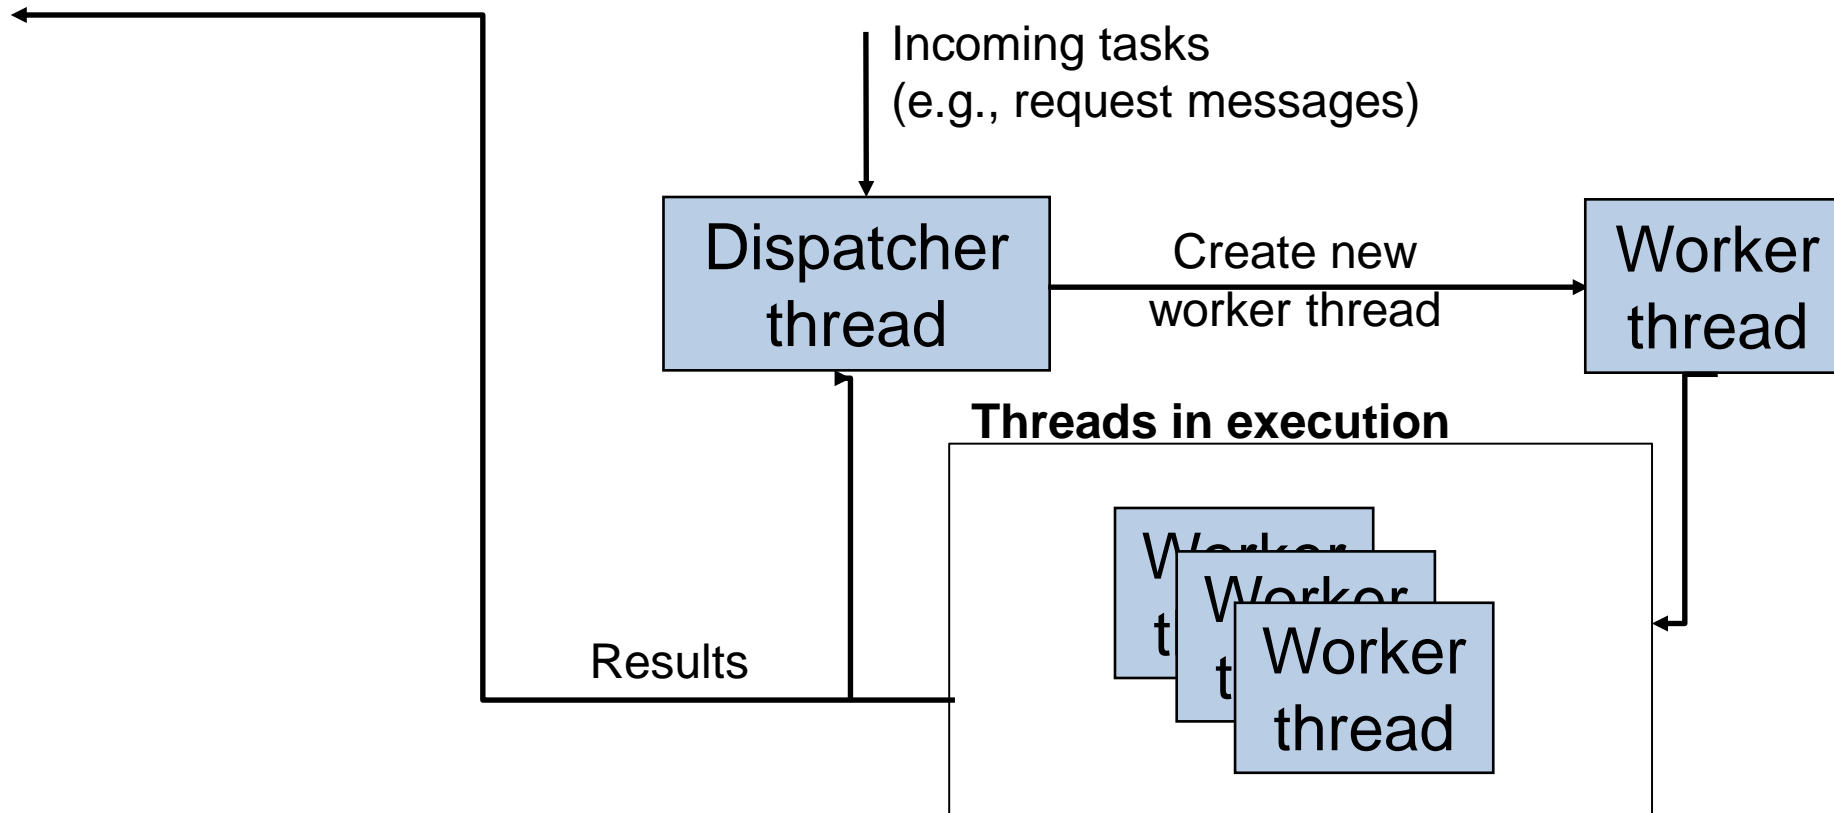
- Fundamentals
- Threads in Linux
- Thread Synchronization
- Deadlocks
- **2.2.5. Important Threading Mechanisms**

2.2.5. ... Worker Threads (1)

- Motivation
 - Programs need to **execute long running tasks**
 - The **application should not block** while tasks are executed
 - The GUI of interactive applications should remain responsive
 - More tasks might come in while one task is executed
- Approach
 - **Dispatcher** thread receives all incoming tasks
 - Dispatcher creates a **separate worker thread** for each long running task
 - If required: Finished worker threads **report their results** back to the dispatcher (which might forward the result)
 - Worker threads **terminate** after completing their task

Worker Threads (2)

- General Model:



Worker Threads (3)

- Advantages
 - Simple mechanism
 - Provides for parallelism in the execution of tasks
- Thread synchronization requirements
 - Worker threads accessing shared data
 - Worker threads returning result data to the dispatcher or the original requester
 - Expensive if many tasks are used
 - Increased complexity

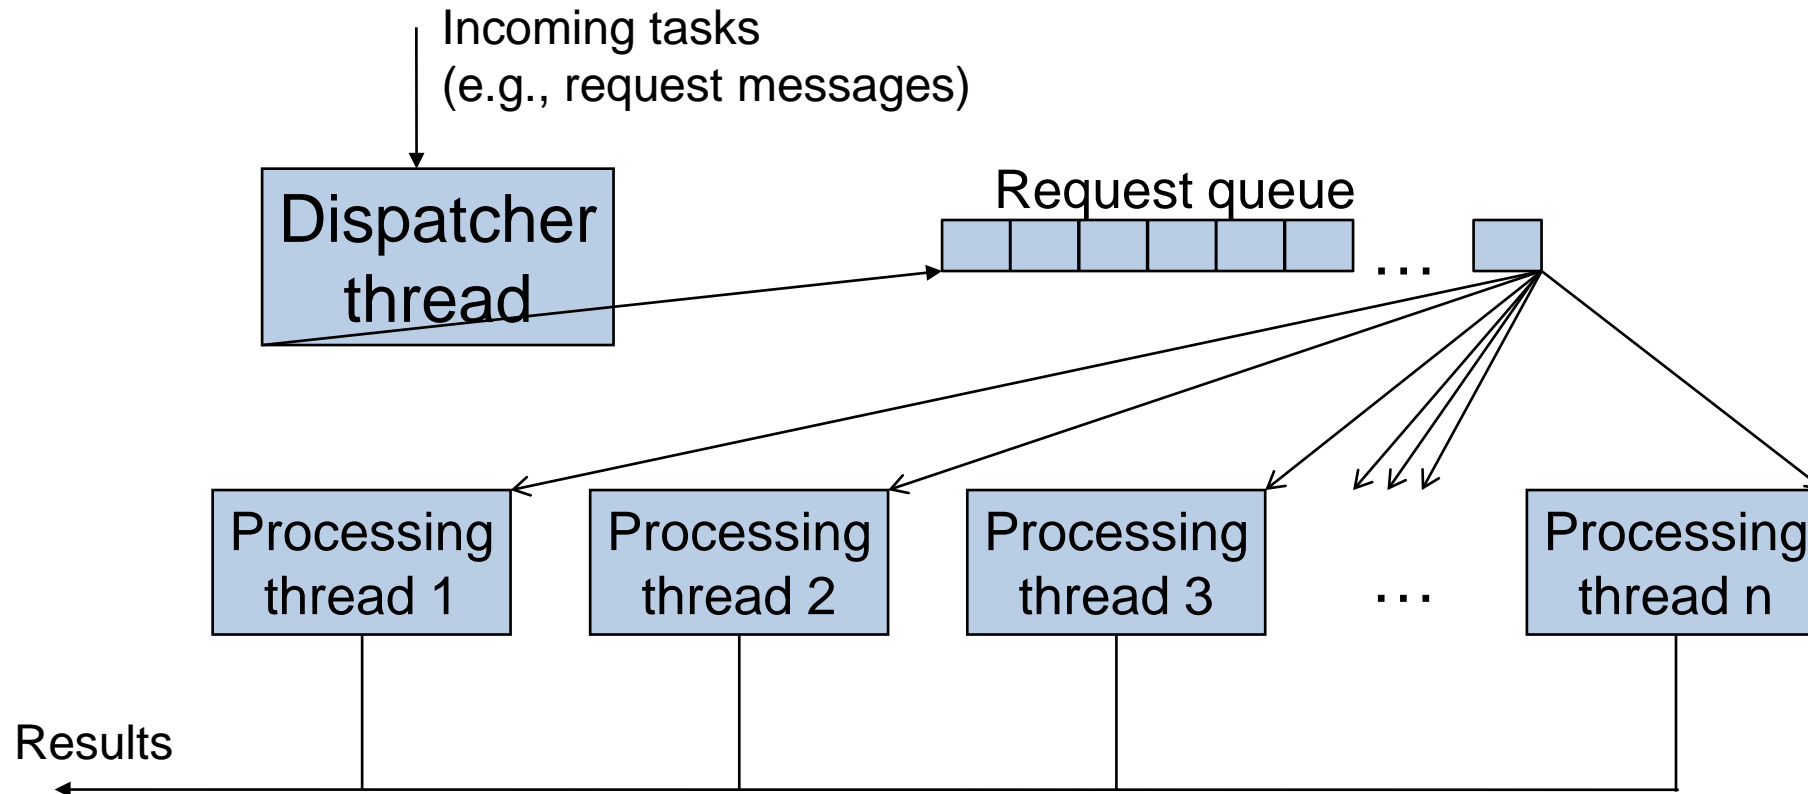


Thread Pools (1)

- Motivation
 - Many applications use threads to perform recurring tasks (e.g., to process incoming messages). Most of them are **short-lived**.
 - **Tasks are independent** of each other.
 - Creating threads on-the-fly for each task can be costly.
- Approach
 - Create a **fixed number** of processing threads (the **thread pool**) at **startup time**.
 - Add new tasks to a **task queue** and notify threads waiting for work.
 - Processing threads becoming available **check the queue** for tasks and wait if currently no task is available.

Thread Pools (2)

- General model



Thread Pools (3)

- **Advantages**
 - Efficient use of multithreading
- Thread synchronization requirements
 - Accessing the request queue for adding or removing tasks
 - Processing threads accessing shared data
 - Processing threads returning result data to the dispatcher or the original requester
- Variants
 - **Dynamically adjust the number** of threads in the pool based on the workload in the system.
 - Use a **priority queue** to handle tasks waiting for a thread.

Thread-Safe Libraries

- Code or code libraries are **thread safe** if they function correctly when executed in parallel by multiple threads
 - Multiple threads might call the same functions
 - Multiple threads might call functions accessing shared data
- Example: Function for printing data to a screen (e.g., `printf` !!!)
 - Data might be buffered before output
 - What if two threads are accessing the buffer in parallel?
- **Documentation** of functions or code libraries should provide information on **whether thread safety is provided or not!**
- In some cases, special reentrant versions of functions or libraries are available

Achieving Thread Safety

- Ensure **re-entrance of functions**
 - Only use **purely local state** – no access of shared data
 - Partial execution of the function only affects the stack of the thread
- Use **mutual exclusion**
 - Serialize access to critical sections using thread synchronization primitives
- Use **thread-local data**
 - Create copies of data for each thread and operate only on this thread-local data
- Use **atomic operations for accessing shared data**
- **Share memory by communication**

Thread-Safe Libraries – Pros and Cons

- **Importance** of thread safety
 - Library function calls are used as black boxes within threads
 - It might not be obvious, which functions of libraries a thread uses
 - Interrelation of library functions often unclear
- **Why** is **not** every library made **thread-safe**?
 - Programming overhead
 - Performance overhead of using thread synchronization mechanisms
 - Thread-safety not relevant in many scenarios
 - Programmer using a library has better overview on the concurrency issues in the application (but does not necessarily know the inner workings of the affected library!).

Summary

- **Threads** are a powerful mechanism for parallel programming
 - More lightweight than processes
 - Allow for efficient communication and cooperation
- Standardized support for thread programming in C is available using the **PThreads library**
- Synchronizing the execution of threads is a big challenge for parallel programming
 - Many **thread synchronization** mechanisms
 - Care must be taken to **avoid deadlocks**
 - Share memory by communicating (if possible)



Summary (2)

- What did we learn in this section?
 - how to **create and handle threads** (2.2.1. / 2.2.2.)
 - how **threads communicate** with each other: (2.2.2.)
 - parameters with creating a new thread
 - return of results from a terminating thread
 - shared variables, data structures, ...
- how to **synchronize threads** in critical sections (2.2.3.x)
 - mutex, semaphore (2.2.3.1. + 2.2.3.3.)
 - condition variable (2.2.3.2.)
 - read/write locks (2.2.3.5.)
 - barriers (2.2.3.4.)
 - channels (2.2.3.5.)
- problem of **deadlocks** (2.2.4.)
- prominent **thread programming models** (2.2.5.)