


# Neural Networks

Lecture

## #1 The Basics



**Cyrill Stachniss**

**Summer term 2024 – Cyrill Stachniss**

# 5 Minute Preparation for Today



<https://www.youtube.com/watch?v=8mwFVKKRePQ>

# **Photogrammetry & Robotics Lab**

## **Intro to Neural Networks Part 1: Network Basics**

**Cyrill Stachniss**

---

The slides have been created by Cyrill Stachniss.

# Image Classification

input

classifier

output



"cat"



"5"

# Semantic Segmentation



**“a label for  
each pixel”**



# Neural Networks

- Machine learning technique
- Often used for classification, semantic segmentation, and related tasks
- First ideas discussed in the 1950/60ies
- Theory work on NNs in the 1990ies
- Increase in attention from 2000 on
- Deep learning took off around 2010
- CNNs for image tasks from 2012 on

# **Part 1**

## **Neural Networks Basics**

# Neural Network



What is a **neuron**?

fundamental unit  
(of the brain)



What is a **network**?

connected elements

**neural networks are connected  
elementary (computing) units**



# Biological Neurons


Biological neurons are the **fundamental units** of the brain that

- Receive sensory input from the external world or from other neurons
- Transform and relay signals
- Send signals to other neurons and also motor commands to the muscles

# Artificial Neurons


Artificial neurons are the fundamental units of artificial neural networks that

- Receive **inputs**
- **Transform** information
- Create an **output**



# Neurons


- Receive **inputs / activations** from sensors or other neurons
- **Combine / transform** information
- Create an **output / activation**



# Neurons as Functions


We can see a neuron as a function

- Input given by  $x \in \mathbb{R}^N$
- Transformation of the input data can be described by a function  $f$
- Output  $f(x) = \hat{y} \in \mathbb{R}$



# Neural Network

- NN is a network/graph of neurons
- Nodes are neurons
- Edges represent input-output connections of the data flow




# Neural Network as a Function

- The whole network is again a function
- Multi-layer perceptron or MLP is often seen as the “vanilla” neural network

input layer


hidden layers

output layer

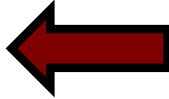


# Neural Networks are Functions

- Neural networks are functions
- Consist of connected artificial neurons
- Input layer takes (sensor) data
- Output layer provides the function result (information or command)
- Hidden layers do some computations



# Different Types of NNs

- Perceptron
- MLP – Multilayer perceptron 
- Autoencoder
- CNN – Convolutional NN
- RNN – Recurrent NN
- LSTM – Long/short term memory NN
- GANs – Generative adversarial network
- Graph NN
- Transformer
- ...



# A mostly complete chart of Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org

Backfed Input Cell

Input Cell

Noisy Input Cell

Hidden Cell

Probabilistic Hidden Cell

Spiking Hidden Cell

Output Cell

Match Input Output Cell

Recurrent Cell


Memory Cell

Different Memory Cell


Kernel

Convolution or Pool


Perceptron (P)




Feed Forward (FF)




Radial Basis Network (RBF)




Deep Feed Forward (DFF)




Recurrent Neural Network (RNN)




Long / Short Term Memory (LSTM)




Gated Recurrent Unit (GRU)




Auto Encoder (AE)




Variational AE (VAE)




Denoising AE (DAE)




Sparse AE (SAE)




Markov Chain (MC)



Hopfield Network (HN)




Boltzmann Machine (BM)














Restricted BM (RBM)




Deep Belief Network (DBN)







# Multi-layer Perceptron (MLP)



# Multi-layer Perceptron Seen as a Function



# Image Classification Example



$x$



input  
image

$f_{NN}(\cdot)$



function that maps  
images to labels

$\hat{y}$



label

# What is the Network's Input?

**An image consists of individual pixels.**



image

# What is the Network's Input?

**An image consists of individual pixels.**

**Each pixel stores an intensity value.**


pixel intensities



image

# What is the Network's Input?

pixel intensities



**An image consists of individual pixels.**


**Each pixel stores an intensity value.**



image



# What is the Network's Input?




**An image consists of individual pixels.**

**Each pixel stores an intensity value.**

**We have  $N+1$  such intensity values.**

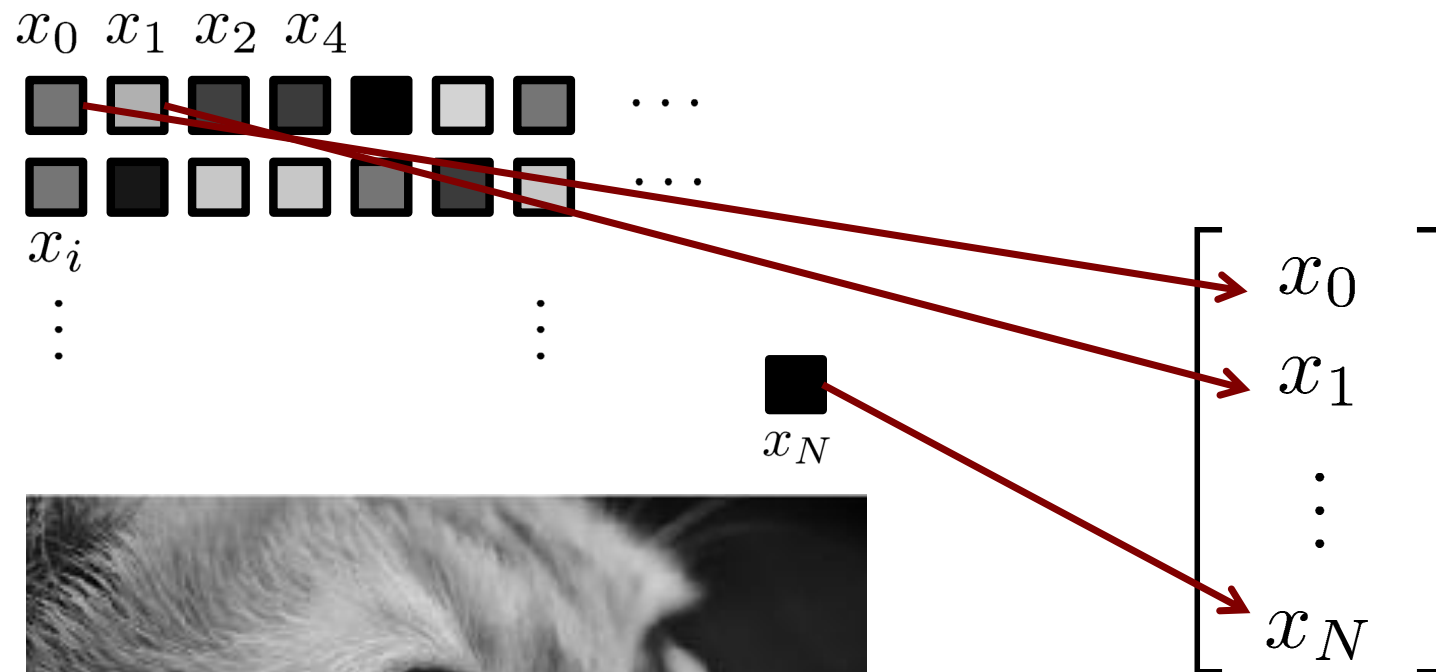
# What is the Network's Input?



$$\begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_N \end{bmatrix}$$

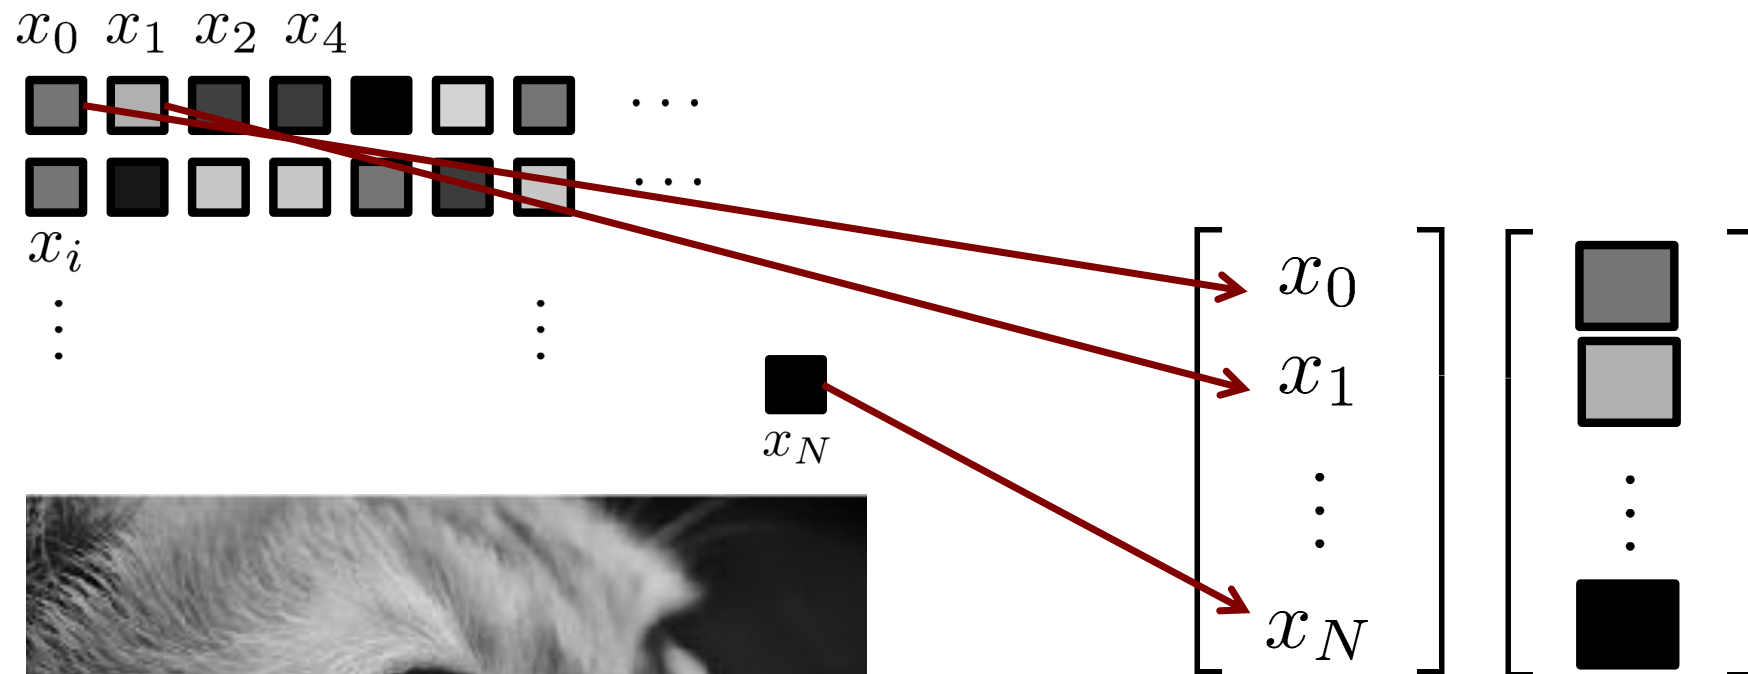
**Arrange all the intensity values in a  $N+1$  dim vector.**

# What is the Network's Input?




**Arrange all the intensity values in a  $N+1$  dim vector.**

# What is the Network's Input?




**Arrange all the intensity values in a  $N+1$  dim vector.**

# Input Layer of the Network




**This vector is  
the input layer  
of our network!**

# What is the Network's Output?




# What is the Network's Output?



Is it a...  
cat or a  
dog or a  
human or a  
...?

# What is the Network's Output?



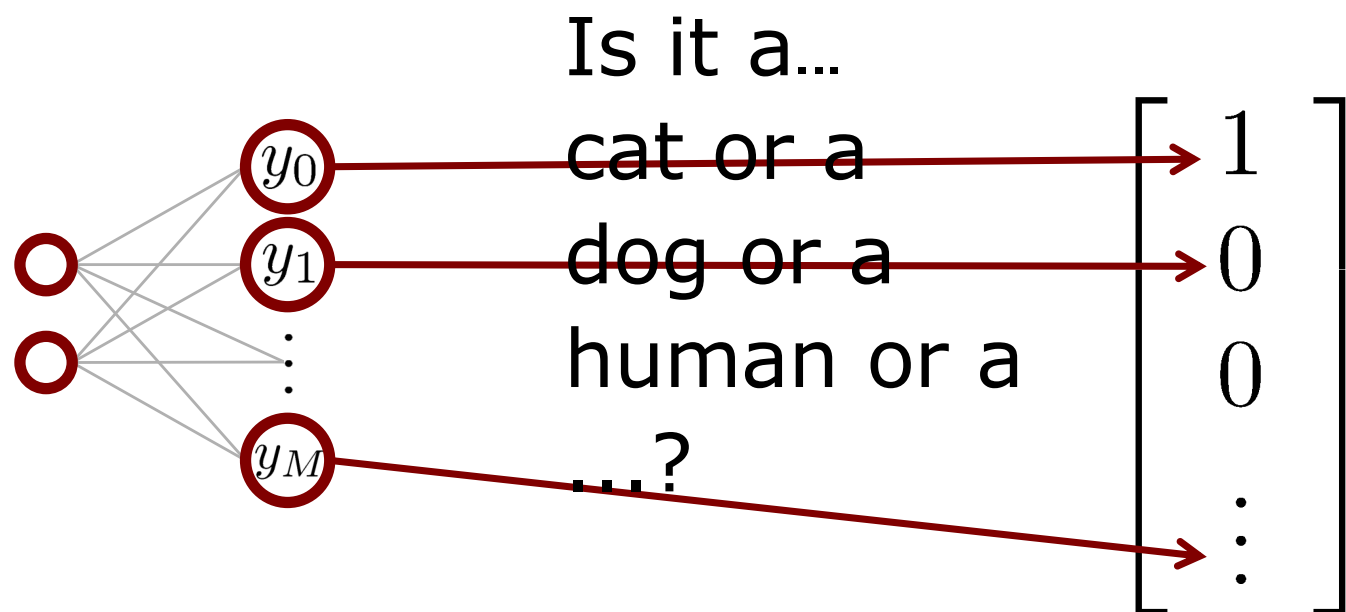
Is it a...  
cat or a  
dog or a  
human or a  
...?

$$\begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \end{bmatrix}$$

**indicator  
vector**




# What is the Network's Output?



**indicator  
vector**

# What is the Network's Output?




Is it a...  
cat or a  
dog or a  
human or a  
...?

$$\begin{bmatrix} 98\% \\ 1\% \\ 0.1\% \\ \vdots \end{bmatrix}$$


**we are  
never  
certain...**

# Output of the Network



**the output layer is vector  
indicating an  
activation/likelihood for  
each label**

# Image Classification




pixels intensities  
are the values of  
the input layer

output layer is a  
vector of likelihoods  
for the possible labels


# Multi-layer Perceptron

## Let's Look at a Single Neuron




# Multi-layer Perceptron


## Let's Look at a Single Neuron



# Perceptron (Single Neuron)




# Perceptron (Single Neuron)





# Function Behind a Neuron



(input) activations  $a_i$

weights  $w_i$

bias  $b$

activation function  $\sigma(\cdot)$


output activation  $a$

# Function Behind a Neuron

A neuron gets activated ( $a$ ) through

- A weighted sum of input activations  $w_i, a_i$
- A bias activation  $b$
- An activation function  $\sigma(\cdot)$

$$a = \sigma(w_0 a_0 + w_1 a_1 + \dots + w_n a_n + b)$$



# Similarity to Convolutions?

- A neuron is similar to a convolution
- Remember linear shift-invariant kernels used as local operators


$$a = \sigma(w_0a_0 + w_1a_1 + \dots + w_na_n + b)$$

**This part looks like the convolutions  
used for defining local operators**

**Additionally: activation function and bias**


# Activation Function

- Biological neurons are either active or not active
- We can see this as a step function:



- Bias tells us where the activation happens

# Activation Function



- We can model this behavior through


$$a = \begin{cases} 0 & \sum_i w_i a_i \leq -b \\ 1 & \text{otherwise} \end{cases}$$

- Non-smooth functions (eg, steps) have disadvantages later down the line...

# Sigmoid Activation Function


- Common activation function is a sigmoid (also called logistic function)
- Smooth function
- Squeezes values to  $[0,1]$

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$




# ReLU Activation Function

- Most commonly used one is the so-called “rectified linear unit” or ReLU
- $\sigma(x) = \text{ReLU}(x) = \max(0, x)$
- Often advantages for deep networks



# Neuron Activation

- A neuron is only activated if  $x > 0$



- If  $a = \text{ReLU}(w_0a_0 + w_1a_1 + \dots + w_na_n + b) > 0$
- the weighted activations are larger than the negative bias  $-b$



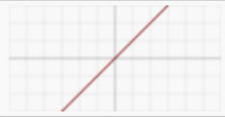


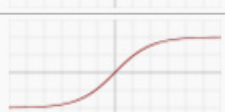
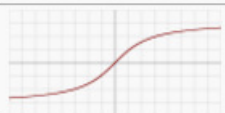
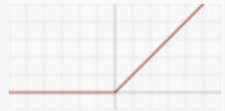


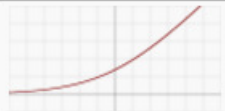
# Common Activation Functions

There are different activation functions

- `sigmoid()`
- `ReLU()`
- `tanh()`
- `atan()`
- `softplus()`
- `identity()`
- `step-function()`
- ...

**ReLU is often used**


# Illustration

Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
SoftPlus		$f(x) = \log_e(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$

[Courtesy of S. Sharma]

# Function Behind a Neuron


- Neuron gets activated if the weighted sum of input activations is large enough (larger than the negative bias)



$$a = \sigma(w_0a_0 + w_1a_1 + \dots + w_na_n + b)$$


- This is the case for **all neurons** in the neural network

# For All Neurons...



**These are a lot of values!**

# Let's Use a Matrix Notation



$$a_0^{(1)} = \sigma \left( w_{00}a_0^{(0)} + w_{01}a_1^{(0)} + \dots + w_{0n}a_n^{(0)} + b_0^{(1)} \right)$$

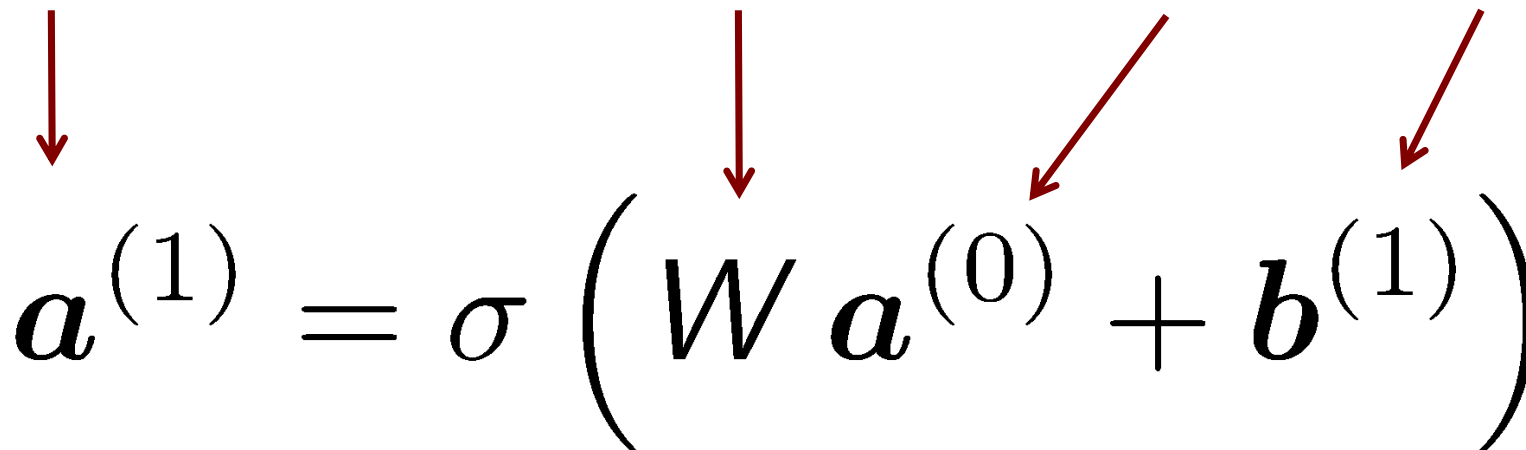
$$\begin{bmatrix} a_0^{(1)} \\ a_1^{(1)} \\ \vdots \\ a_n^{(1)} \end{bmatrix} = \sigma \left( \begin{bmatrix} w_{00} & w_{01} & \dots & w_{0n} \\ w_{10} & w_{11} & \dots & w_{1n} \\ \vdots & \vdots & \vdots & \vdots \\ w_{k0} & w_{k1} & \dots & w_{kn} \end{bmatrix} \begin{bmatrix} a_0^{(0)} \\ a_1^{(0)} \\ \vdots \\ a_n^{(0)} \end{bmatrix} + \begin{bmatrix} b_0^{(1)} \\ b_1^{(1)} \\ \vdots \\ b_n^{(1)} \end{bmatrix} \right)$$

# Each Layer Can Be Expressed Through Matrix Multiplications


**layer 1**

**layer 0**

$$\begin{bmatrix} a_0^{(1)} \\ a_1^{(1)} \\ \vdots \\ a_n^{(1)} \end{bmatrix} = \sigma \left( \begin{bmatrix} w_{00} & w_{01} & \dots & w_{0n} \\ w_{10} & w_{11} & \dots & w_{1n} \\ \vdots & \vdots & \ddots & \vdots \\ w_{k0} & w_{k1} & \dots & w_{kn} \end{bmatrix} \begin{bmatrix} a_0^{(0)} \\ a_1^{(0)} \\ \vdots \\ a_n^{(0)} \end{bmatrix} + \begin{bmatrix} b_0^{(1)} \\ b_1^{(1)} \\ \vdots \\ b_n^{(1)} \end{bmatrix} \right)$$


$$\mathbf{a}^{(1)} = \sigma \left( \mathbf{W} \mathbf{a}^{(0)} + \mathbf{b}^{(1)} \right)$$

# Do It Layer by Layer...



# Do It Layer by Layer...

**input** = layer 0  $x = a^{(0)}$

layer 1

$$\sigma \left( W^{(1)} a^{(0)} + b^{(1)} \right) = a^{(1)}$$

layer 2

$$\sigma \left( W^{(2)} a^{(1)} + b^{(2)} \right) = a^{(2)}$$

⋮

layer k = **output**


$$\sigma \left( W^{(k)} a^{(k-1)} + b^{(k)} \right) = a^{(k)} = \hat{y}$$

**That not much more than linear algebra...**



# Feedforward Networks

- MLPs are feedforward networks
- The information flows from left to right
- There are no loops



- Such networks are called feedforward networks
- There exist other variants (eg, RNNs)

# **Example: Handwritten Digit Recognition**

# Handwritten Digit Recognition


504192




28x28 pixel image

= 5

# Handwritten Digit Recognition




# A Basic MLP Recognizing Digits



# Images to Digits - A Mapping from 784 to 10 Dimensions



**28x28  
pixel  
input  
images  
(784 dim)**



$$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ \vdots \end{bmatrix}$$

**output  
vector**

**(10 dim)**

# **What Happens in the Layers?**

# What Happens in the 1<sup>st</sup> Layer?



**784 input activations = pixel intensities**

**784 weights = weights for pixel intensities**




# What Happens in the 1<sup>st</sup> Layer?

784 input activations = pixel intensities


784 weights = weights for pixel intensities

**treat activations and weights as images**




# What Happens in the 1<sup>st</sup> Layer?

$a_i$




$w_i$



**weights tell us  
what matters for  
activating the neuron!**

**individual “weight images” for a neuron  
support individual patterns in the image**

# Link to Local Operators Defines Through Convolutions



**weights tell us  
what matters**


- Direct link to defining image operators through convolutions

**Here:**

- Global (not local) operators
- Weight matrix does not (yet) “slide over image”


# Weights & Bias = Patterns

- **Weights define** the **patterns** to look for in the image
- **Bias** tells us how well the image must **match** the pattern
- Activation functions "**switches the neuron on**" if it matches the pattern




# What Happens in the 2<sup>nd</sup> Layer?

- The weights in layer 2 tell us which 1<sup>st</sup> layer patterns should be combined
- The deeper we go, the more patterns get arranged and combined



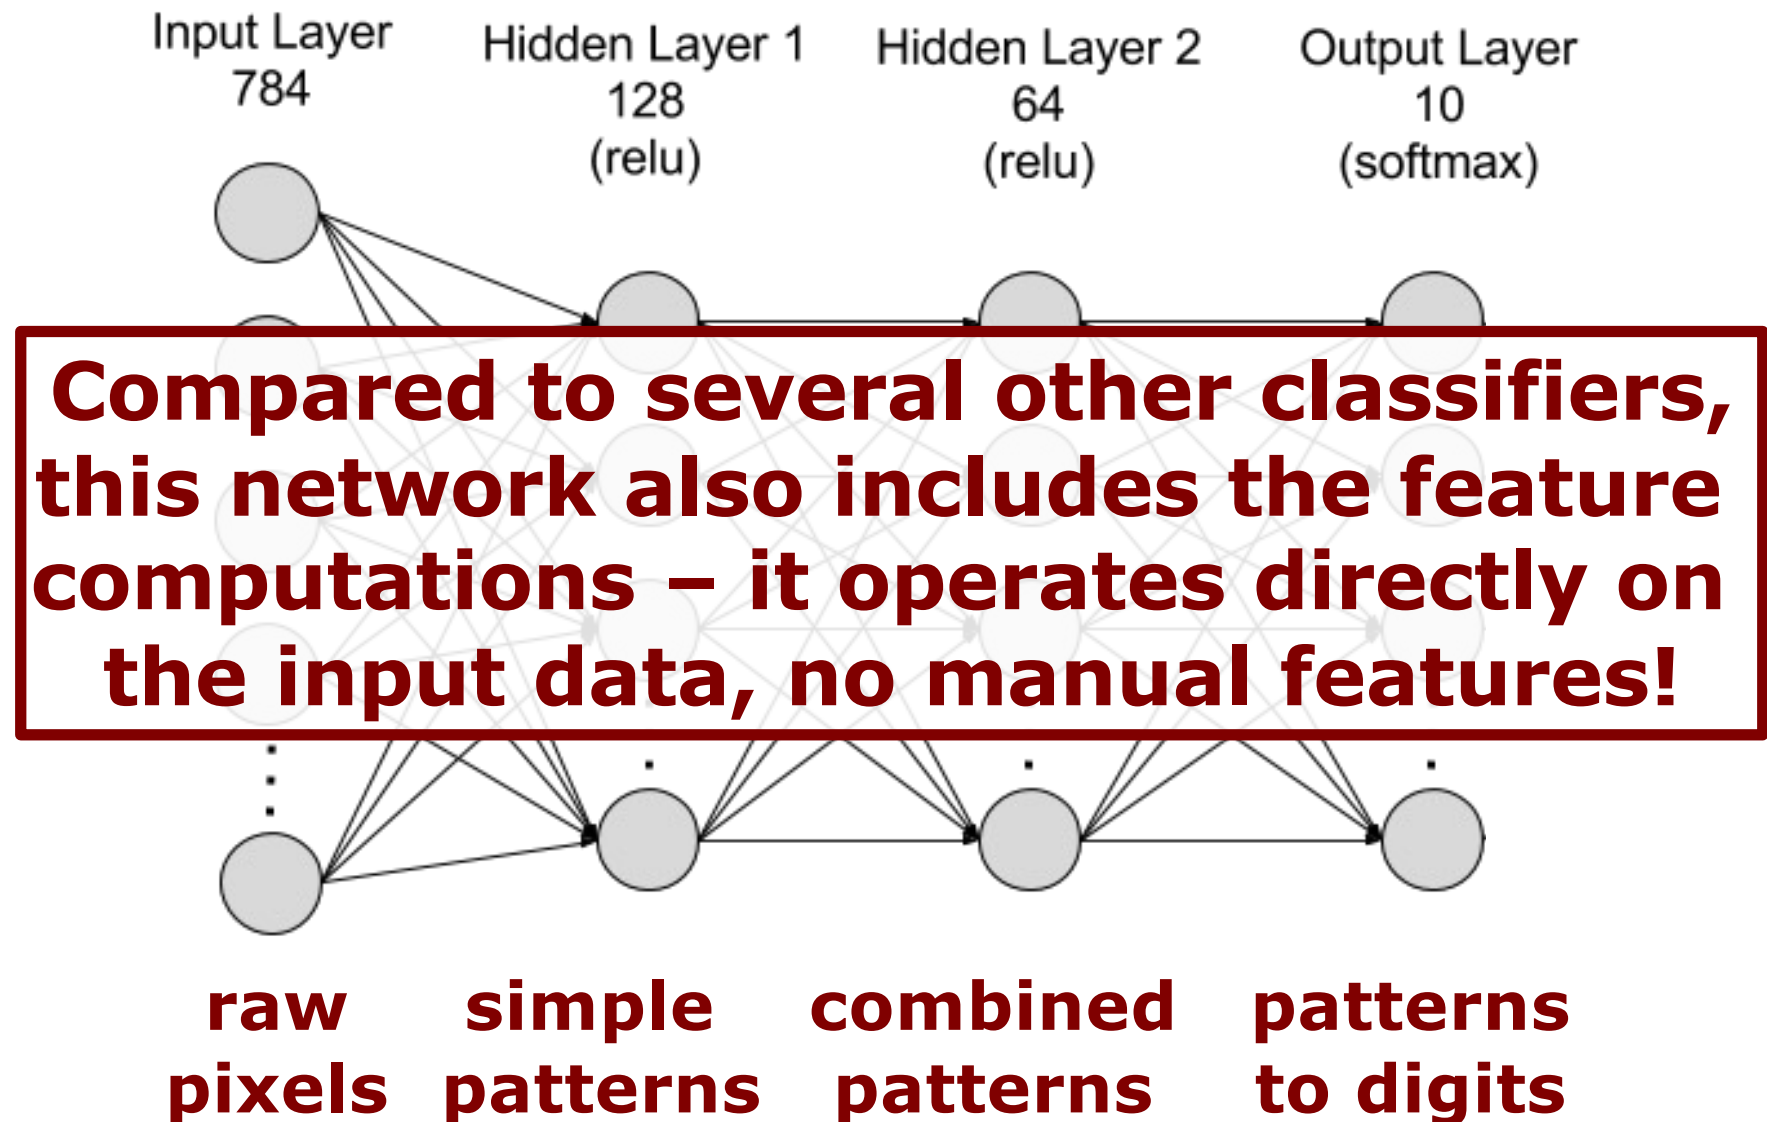
- The last layer decides, which final patterns make up a digit

# What Happens in the Layers?




**raw pixels    simple patterns    combined patterns    patterns to digits**

# No Manual Features



# Classification Performance




Such a simple MLP achieves a correct classification for  $\sim 96\%$  of the examples



# Classification Performance

- A simple MLP achieves a classification accuracy of  $\sim 96\%$
- Note that there are tricky cases




- That is a good performance for a simple model!
- Improved networks achieve  $\sim 99\%$

# **How to Design a Neural Network?**

# How to Make the Network Compute What We Want?

- So far, the network is a recipe for sequentially performing computations
- **Structure** and **parameters** are the **design choices**
- How to set them?

**Learning!**



# Summary – Part 1

- What are neurons and neural networks
- Lots of different networks exists
- Focus: multi-layer perceptrons (MLP)
- Activations, weights, bias
- Networks have many parameters
- “It’s just a bunch of matrices and vectors”
- MLP for simple image classification
- Part 2: Learning the parameters

# Literature & Resources

- Online Book by Michael Nielsen, Chapter 1:  
<http://neuralnetworksanddeeplearning.com/chap1.html>
- Nielsen, Chapter 1, Python3 code:  
<https://github.com/MichalDanielDobrzanski/DeepLearningPython>
- MNIST database:
  - <http://yann.lecun.com/exdb/mnist/>
- Grant Sanderson, Neural Networks  
<https://www.3blue1brown.com/>
- Alpaydin, Introduction to Machine Learning