

The logo of the University of Bonn, featuring a blue square with a white curved line and a grey square.

UNIVERSITÄT **BONN**

Juergen Gall

Linear Filtering
MA-INF 2201 - Computer Vision
WS24/25

Organization

- Lecturer:
 - Prof. Dr. Juergen Gall
 - <http://gall.cv-uni-bonn.de/>
- Teaching:

Prof. Dr. Juergen Gall

Teaching

Research Interests

Job Offers/Theses

Publications

Software

Data

Projects

Conferences

Teaching

Talks

Awards

PhD Seminar - Graphics, Vision, Audio for Intelligent Systems, SS13, WS13/14, SS14, WS14/15, SS15, WS15/16, SS16, WS16/17, SS17, WS17/18, SS18, WS18/19, SS19, WS19/20

PhD Machine Learning Seminar, Dates

BA-INF 062 - Begleitseminar zur Bachelorarbeit: Computer Vision, SS14, WS14/15, SS15, WS15/16, SS16, WS16/17, SS17, WS17/18, SS18, WS18/19, SS19, WS19/20, SS20, WS20/21, SS21, WS21/22, SS22, WS22/23, SS23, WS23/24

MA-INF 2201 - Computer Vision, WS13/14, WS14/15, WS15/16, WS16/17, WS17/18, WS18/19, WS19/20, WS21/22, WS22/23, **WS23/24**

MA-INF 2213 - Advanced Computer Vision (Computer Vision II), SS14, SS15, SS16, SS17, SS18, SS19, SS20, SS21, SS22, SS23

MA-INF 2218 - Video Analytics, SS17, SS18, SS19, SS21, SS22, SS23

Organization

- Slides:

<http://gall.cv-uni-bonn.de/teaching/Lectures/cv23.html>

MA-INF 2201 - Computer Vision

4L + 2E, WS23/24

Lecturer

Juergen Gall

Content

Tentative: Linear filters, Edges, Derivatives, Hough Transform, Segmentation, Graph Cuts, Mean Shift
Background Subtraction, Temporal Filtering, Active Appearance Models, Shapes, Optical Flow, 2D Trajectory
Estimation, Articulated Pose Estimation, Deformable Meshes, RGBD Vision

Material

The slides and recordings are available at [Slides/Recordings](#).

Prerequisites

Basic knowledge of linear algebra, analysis, probability theory, Python programming

Organization

- Slides via sciebo (<https://hochschulcloud.nrw/en/>):



- Password: MA-INF2201

Organization

- Structure: 4+2 SWS
- Lecture:
 - Tuesday, 10:15-11:45, HSZ / HS 3
 - Friday, 10:15-11:45, HSZ / HS 3
 - Lectures will be recorded
- Exercise:
 - Start: 16.10.
 - Wednesday, 12:15-13:45, 2.025, Informatikzentrum
 - Wednesday, 14:15-15:45, 2.025, Informatikzentrum

Overview – Part 1

- Image
- Image processing
 - Filtering
 - Edges






Overview – Part 1

- Detecting objects with single template



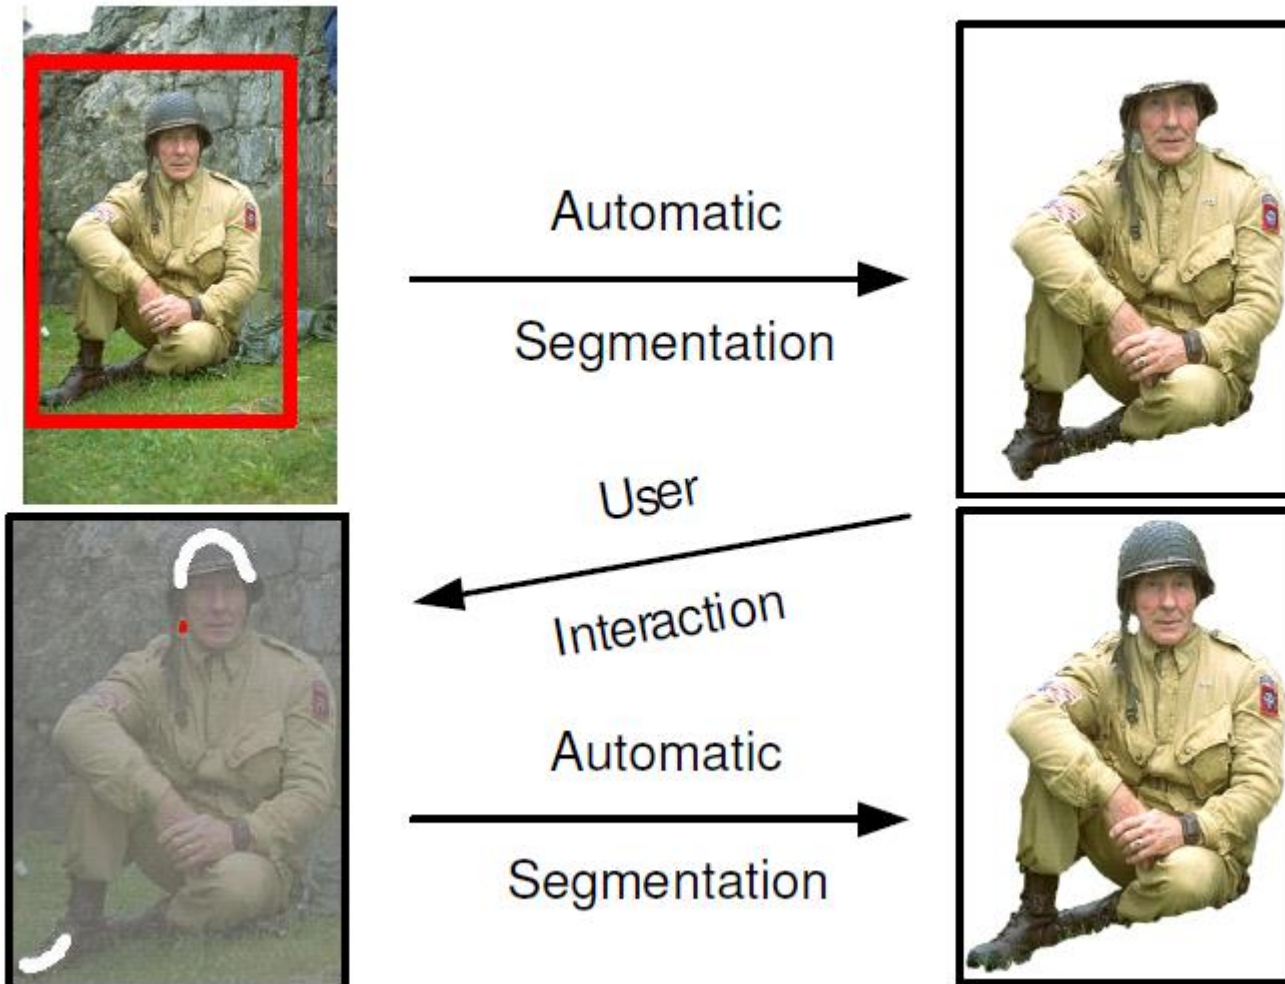
Overview – Part 1

- Track objects




Overview – Part 1

- Segment objects



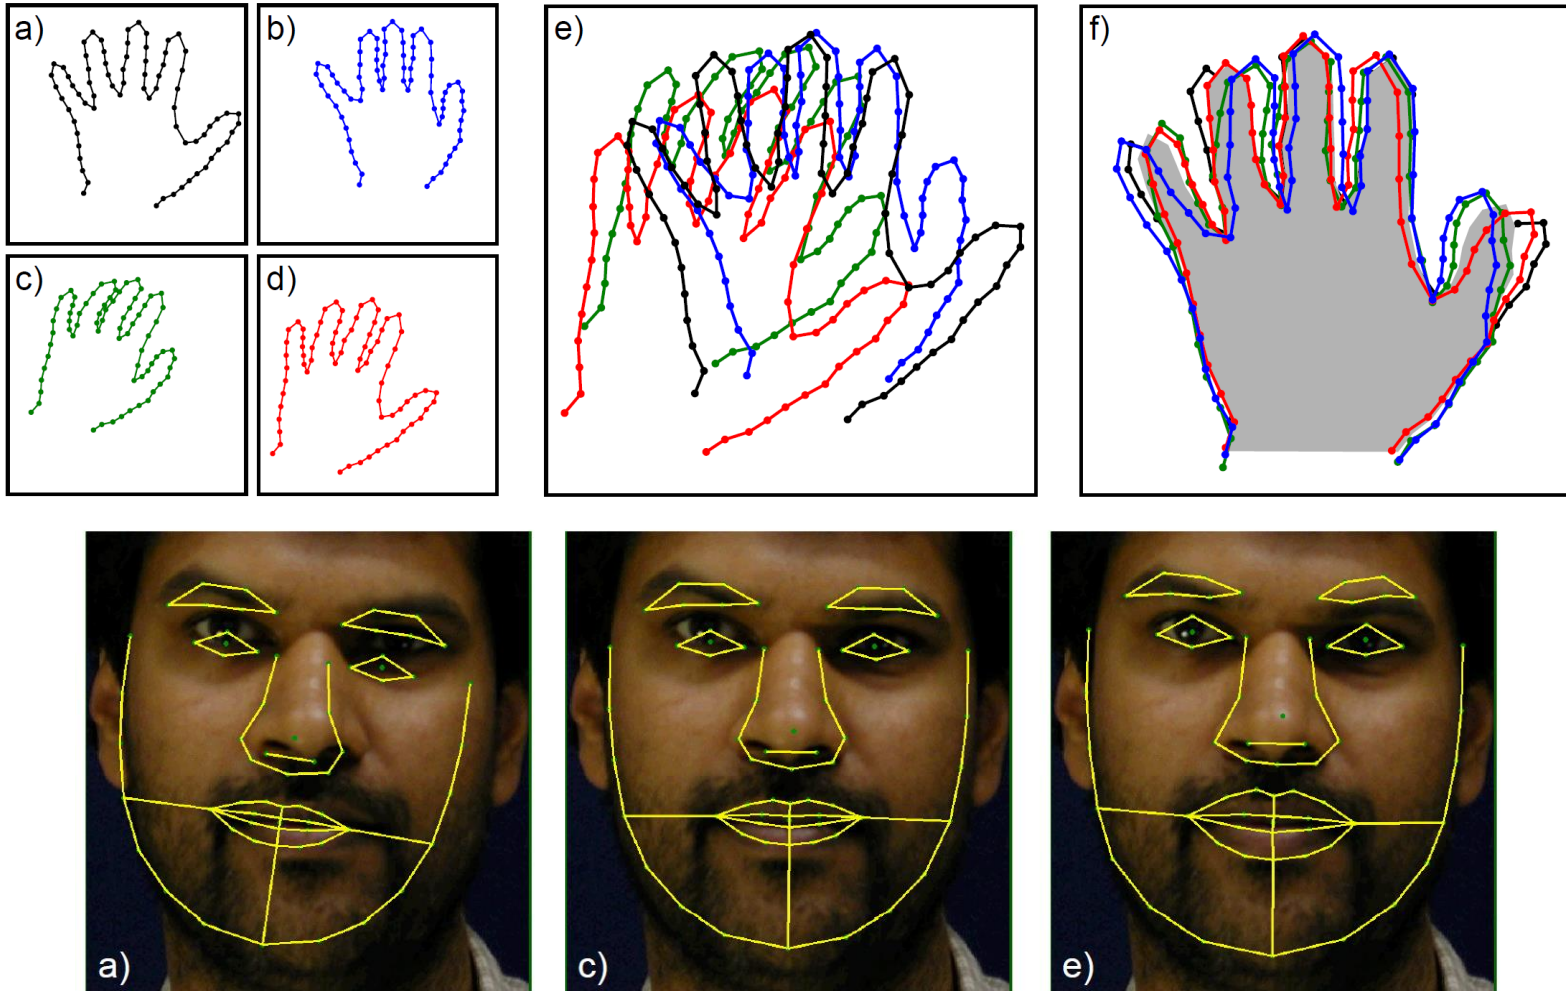
Overview – Part 1

- Motion segmentation
- Optical Flow




Overview – Part 1

- Statistical shape models




Overview – Part 1

- Align images / Panorama views



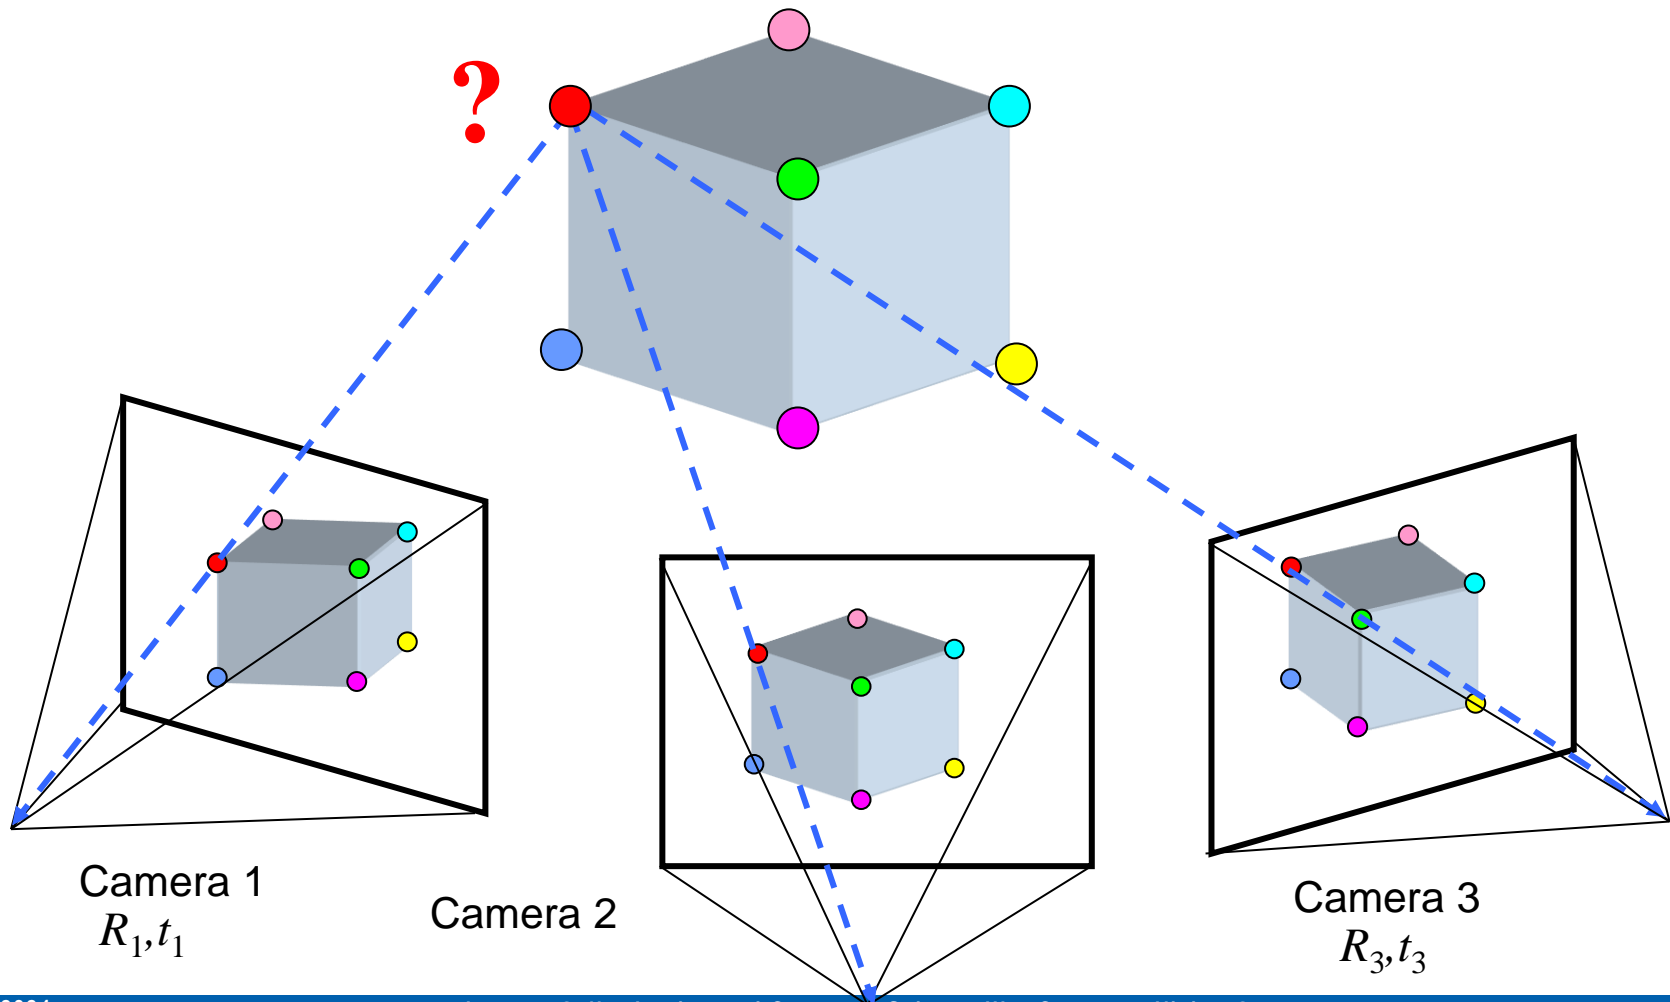
Overview – Part 2

- Cameras
- Camera calibration






Overview – Part 2

- 3D Reconstruction / Stereo / Structure from motion




Overview – Part 2

- Pose estimation




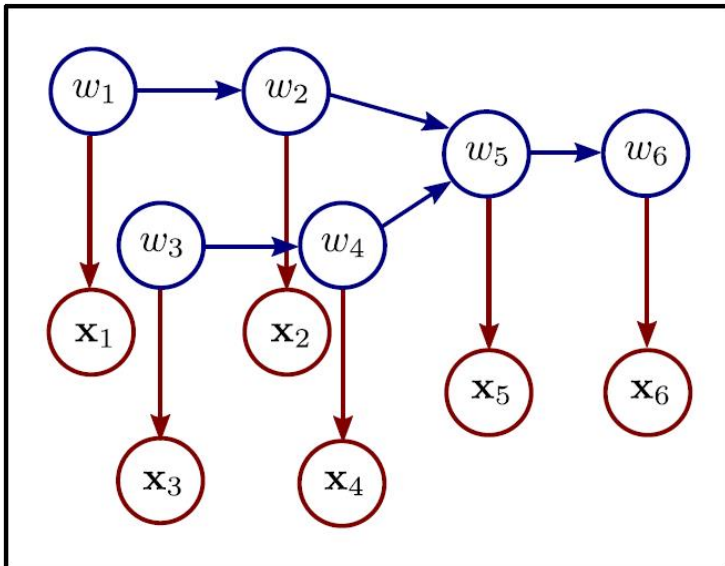
Overview – Part 2


- Depth sensors



Overview – Bonus


- Graphical models





Chapter 2 and 3. Computer Vision - Algorithms and Applications, Szeliski, Richard, Springer, 2011

Image as function




45	60	98	127	132	133	137	133
46	65	98	123	126	128	131	133
47	65	96	115	119	123	135	137
47	63	91	107	113	122	138	134
50	59	80	97	110	123	133	134
49	53	68	83	97	113	128	133
50	50	58	70	84	102	116	126
50	50	52	58	69	86	101	120

Mapping from image domain (pixels)
to values

Gray image: $f : \Omega \rightarrow \mathbb{R}$

Color images: $f : \Omega \rightarrow \mathbb{R}^3$

Image represented as matrix or
tensor



Color Spaces

There are many
color spaces

Popular: Intensity,
RGB, Lab, HSV



(a) RGB



(b) R



(c) G



(d) B



(e) rgb



(f) r



(g) g




(h) b




(i) L^*



(j) a^*



(k) b^*



(l) H



(m) S



(n) V

Pixel transforms

Image operator h :

$$g(\mathbf{x}) = h(f(\mathbf{x}))$$

$$g(i, j) = h(f(i, j)) \quad \mathbf{x} = (i, j)$$

Blend two images linearly

$$g(\mathbf{x}) = (1 - \alpha)f_0(\mathbf{x}) + \alpha f_1(\mathbf{x})$$

Gamma correction:

$$g(\mathbf{x}) = [f(\mathbf{x})]^{1/\gamma}$$

Histogram equalization

Pixels as distribution:

$$p_f(y) = \frac{|\{x \in \Omega : f(x) = y\}|}{|\Omega|}$$



$$F_f(y) = \int_0^y p_f(y') dy' = \sum_{y'=0}^y p_f(y')$$

Cumulative distribution function

Make CDF linear:

$$h(f(x)) = F(f(x)) \cdot 255$$

$$p_h(z) = \frac{|\{x \in \Omega : h(f(x)) = z\}|}{|\Omega|} = \frac{1}{255}$$



Proof

Probability theory:

$$F_h(z) = \int_0^z p_h(z') dz' = \int_0^{h^{-1}(z)} p_f(y') dy' = F_f(h^{-1}(z))$$

We get:

$$y = h^{-1}(z)$$


$$p_h(z) = \frac{d}{dz} F_h(z) = \frac{d}{dz} F_f(h^{-1}(z)) = p_f(h^{-1}(z)) \frac{d}{dz} h^{-1}(z)$$

Since $h(f(x)) = F(f(x)) \cdot 255$ and $\frac{d}{dy} h(y) = 255 \cdot p_f(y)$:

$$p_f(h^{-1}(z)) \frac{d}{dz} h^{-1}(z) = \frac{1}{255} \cdot \underbrace{\frac{d}{dy} h(y) \Big|_{y=h^{-1}(z)} \cdot \frac{d}{dz} h^{-1}(z)}_{\frac{d}{dz} h(h^{-1}(z)) = 1}$$

$$p_h(z) = \frac{1}{255}$$

Linear filtering



Source: S. Lazebnik

Motivation: Image denoising

How can we reduce noise in a photograph?



Source: S. Lazebnik

Moving average

Let's replace each pixel with a weighted average of its neighborhood

The weights are called the filter kernel

What are the weights for the average of a 3x3 neighborhood?

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

“box filter”

Defining convolution

Let f be the image and h be the kernel. The output of convolving f with h is denoted: $g = f * h$

$$g(i, j) = \sum_{k, l} f(i - k, j - l) h(k, l) = \sum_{k, l} f(k, l) h(i - k, j - l)$$

45	60	98	127	132	133	137	133
46	65	98	123	126	128	131	133
47	65	96	115	119	123	135	137
47	63	91	107	113	122	138	134
50	59	80	97	110	123	133	134
49	53	68	83	97	113	128	133
50	50	58	70	84	102	116	126
50	50	52	58	69	86	101	120

 $f(x, y)$

*

0.1	0.1	0.1
0.1	0.2	0.1
0.1	0.1	0.1

=

69	95	116	125	129	132
68	92	110	120	126	132
66	86	104	114	124	132
62	78	94	108	120	129
57	69	83	98	112	124
53	60	71	85	100	114

 $g(x, y)$

Convolution

Continuous convolution:

$$g(\boldsymbol{x}) = \int f(\boldsymbol{x} - \boldsymbol{u})h(\boldsymbol{u})d\boldsymbol{u}$$

h is called impulse response function:

$$h * \delta = h \qquad \delta(x) = \begin{cases} \infty & \text{if } x=0 \\ 0 & \text{otherwise} \end{cases}$$

Proof:

$$g(x) = \int \delta(x - u)h(u)du = h(x)$$

Key properties

Linearity: $\text{filter}(f1 + f2) = \text{filter}(f1) + \text{filter}(f2)$

Shift invariance: same behavior regardless of pixel location:
 $\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$

Theoretical result: any linear shift-invariant operator can be represented as a convolution

Proof

Linear shift-invariant operator T as a convolution

Using $f(i, j) = \sum_{p=-\infty}^{\infty} \sum_{q=-\infty}^{\infty} f(p, q) \delta(i - p, j - q)$ with $\delta(p, q) = \begin{cases} 1 & \text{if } p=0 \text{ and } q=0 \\ 0 & \text{otherwise} \end{cases}$ we get:

$$\begin{aligned}
 (T \circ f)(i, j) &= T \left\{ \sum_{p=-\infty}^{\infty} \sum_{q=-\infty}^{\infty} f(p, q) \delta(i - p, j - q) \right\} \\
 &= \sum_{p=-\infty}^{\infty} \sum_{q=-\infty}^{\infty} f(p, q) (T \circ \delta)(i - p, j - q) && \text{linear} \\
 &= \sum_{p=-\infty}^{\infty} \sum_{q=-\infty}^{\infty} f(i - p, j - q) (T \circ \delta)(p, q) && \text{shift-invariant} \\
 &= (f * (T \circ \delta))(i, j)
 \end{aligned}$$

Properties in more detail

Commutative: $a * b = b * a$

- Conceptually no difference between filter and signal

Associative: $a * (b * c) = (a * b) * c$



- Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
- This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$

Distributes over addition: $a * (b + c) = (a * b) + (a * c)$


Scalars factor out: $ka * b = a * kb = k(a * b)$

Identity: unit impulse $e = [\dots, 0, 0, 1, 0, 0, \dots]$,
 $a * e = a$

Border padding and output size




zero



wrap



clamp




mirror

Rectangular filter



Rectangular filter




Rectangular filter



Integral image

Precompute average



Get average values of any size by reading only 4 values!

Integral image

3	2	7	2	3
1	5	1	3	4
5	1	3	5	1
4	3	2	1	6
2	4	1	4	8

Original image

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

Integral image

$$s(i, j) = \sum_{k=0}^i \sum_{l=0}^j f(k, l)$$

Recursive:

$$s(i, j) = s(i - 1, j) + s(i, j - 1) - s(i - 1, j - 1) + f(i, j)$$

$$17 + 19 - 11 + 3 = 28$$

Integral image

3	2	7	2	3
1	5	1	3	4
5	1	3	5	1
4	3	2	1	6
2	4	1	4	8

Original image

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

Integral image

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

Integral image

4 values independent of size:

$$s(i_1, j_1) - s(i_1, j_0 - 1) - s(i_0 - 1, j_1) + s(i_0 - 1, j_0 - 1)$$

$$48 - 13 - 14 + 3 = 24 = 5 + 1 + 3 + 1 + 3 + 5 + 3 + 2 + 1$$

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

?

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters



Original

0	0	0
1	0	0
0	0	0

?

Practice with linear filters



Original

0	0	0
1	0	0
0	0	0



Shifted *left*
By 1 pixel

Practice with linear filters



Original

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

?


Practice with linear filters



Original

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1



Blur (with a
box filter)

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$


1	1	1
1	1	1
1	1	1




Sharpening filter

- Accentuates differences with local average

Sharpening



before



after

Sharpening

What does blurring take away?



−



=



Let's add it back:




+




=



Image rotation





?



It is linear, but not a spatially invariant operation. There is no convolution.

Smoothing with box filter revisited



What's wrong with this picture?



Gaussian Kernel

Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should renormalize weights to sum to 1 in any case)

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

5 x 5, $\sigma = 1$

Choosing kernel width


Rule of thumb: set filter half-width to about 3σ

$$P(\mu - \sigma < X \leq \mu + \sigma) \approx 68\%$$


$$P(\mu - 2\sigma < X \leq \mu + 2\sigma) \approx 95.5\%$$

$$P(\mu - 3\sigma < X \leq \mu + 3\sigma) \approx 99.7\%$$

Effect of σ



Gaussian vs. box filtering



Source: S. Lazebnik

Gaussian filters

Remove “high-frequency” components from the image (low-pass filter)

Convolution with itself is another Gaussian

- So can smooth with small- σ kernel, repeat, and get same result as larger- σ kernel would have
- Convoluting two times with Gaussian kernel with std. σ is same as convoluting once with kernel with std. dev. $\sigma\sqrt{2}$

Separable kernel

- Factors into product of two 1D Gaussians

Separability of the Gaussian filter

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Separability example

2D convolution
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors
into a product of 1D
filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution
along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution
along the remaining column:

Why is separability useful?

What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?

- $O(n^2 m^2)$

What if the kernel is separable?

- $O(n^2 m)$

Is a kernel separable?

- Kernel as matrix:

$$K = \sum_i \sigma_i \mathbf{u}_i \mathbf{v}_i^T$$

- Singular Value Decomposition (SVD)
- If only the first singular value σ_0 is non-zero, kernel is separable
- Vertical and horizontal kernels:

$$\sqrt{\sigma_0} \mathbf{u}_0 \text{ and } \sqrt{\sigma_0} \mathbf{v}_0^T$$

- Approximation: $K \approx \sigma_0 \mathbf{u}_0 (\mathbf{v}_0)^T$

SVD - Definition

$$\mathbf{A}[m \times n] = \mathbf{U}[m \times r] \mathbf{\Sigma}[r \times r] (\mathbf{V}[n \times r])^T$$

A: Input data matrix

- $m \times n$ matrix

U: Left singular vectors

- $m \times r$ matrix


Σ : Singular values

- $r \times r$ diagonal matrix (r : rank of the matrix A)


V: Right singular vectors

- $n \times r$ matrix

$$\mathbf{A} \approx \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_i \sigma_i \mathbf{u}_i \circ \mathbf{v}_i^T$$



$$\mathbf{A} \approx \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_i \sigma_i \mathbf{u}_i \circ \mathbf{v}_i^T$$



σ_i ... scalar
 \mathbf{u}_i ... vector
 \mathbf{v}_i ... vector

SVD - Properties

It is always possible to decompose a real matrix A into $A = U \Sigma V^T$, where

U, Σ, V : unique

U, V : column orthonormal:

- $U^T U = I; V^T V = I$ (I : identity matrix)
- (Cols. are orthogonal unit vectors)

Σ : diagonal

- Entries (singular values) are positive, and sorted in decreasing order ($\sigma_1 \geq \sigma_2 \geq \dots \geq 0$)


SVD – Example

$A = U \Sigma V^T$ - example:

$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.18 & 0 \\ 0.36 & 0 \\ 0.18 & 0 \\ 0.90 & 0 \\ 0 & 0.53 \\ 0 & 0.80 \\ 0 & 0.27 \end{bmatrix} \times \begin{bmatrix} 9.64 & 0 \\ 0 & 5.29 \end{bmatrix} \times \begin{bmatrix} 0.58 & 0.58 & 0.58 & 0 & 0 \\ 0 & 0 & 0 & 0.71 & 0.71 \end{bmatrix}$$

Geometric interpretation

Change of basis vectors:



$$\mathbf{A}_{M \times N} = \mathbf{U}_{M \times P} \mathbf{\Sigma}_{P \times P} \mathbf{V}_{P \times N}^T$$

$$\mathbf{A}\mathbf{V} = \mathbf{U}\mathbf{\Sigma} \quad \text{or} \quad \mathbf{A}\mathbf{v}_j = \sigma_j \mathbf{u}_j$$

Noise

Salt and pepper noise:
contains random
occurrences of black and
white pixels

Impulse noise: contains
random occurrences of white
pixels

Gaussian noise: variations in
intensity drawn from a
Gaussian normal distribution



Original



Salt and pepper noise

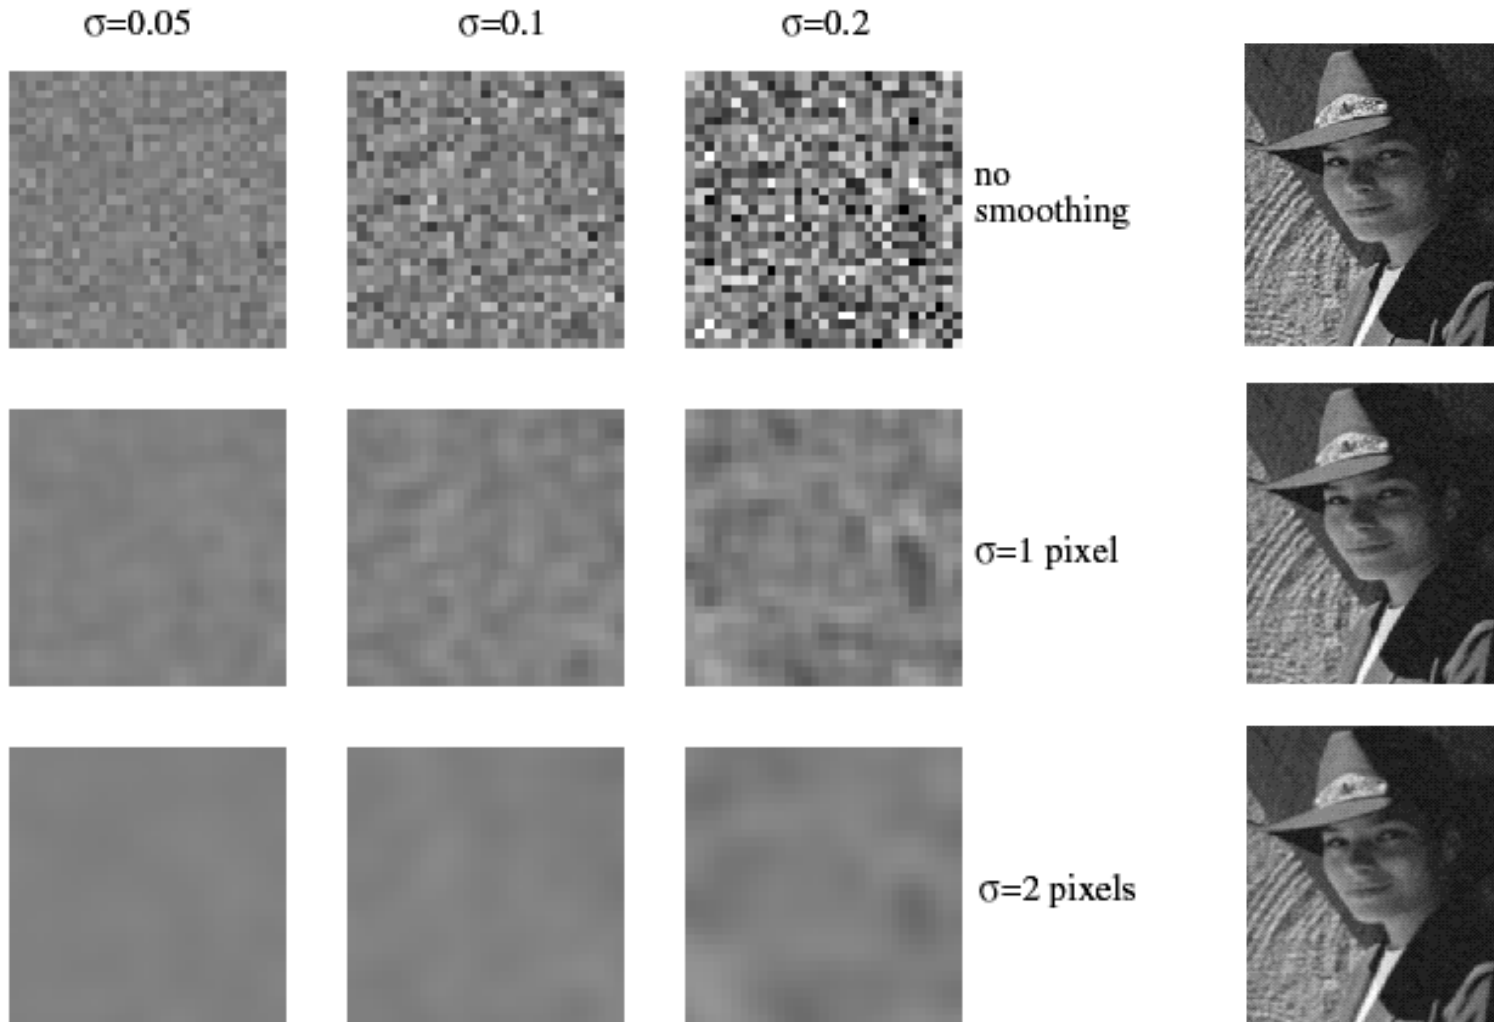


Impulse noise



Gaussian noise

Reducing Gaussian noise



Smoothing with larger standard deviations suppresses noise, but also blurs the image

Reducing salt-and-pepper noise

What's wrong with the results?

3x3



5x5




7x7



Alternative idea: Median filtering

A median filter operates over a window by selecting the median intensity in the window




Is median filtering linear?

Median filter

What advantage does median filtering have over Gaussian filtering?

Robustness to outliers

filters have width 5 :



Gaussian vs. median filtering

3x3

5x5

7x7

Gaussian




Median




Bilateral filter

Filter is data dependent


$$g(i, j) = \frac{\sum_{k,l} f(k, l) w(i, j, k, l)}{\sum_{k,l} w(i, j, k, l)}$$



Input



Domain kernel



Range kernel


$$d(i, j, k, l) = \exp \left(-\frac{(i - k)^2 + (j - l)^2}{2\sigma_d^2} \right)$$

$$r(i, j, k, l) = \exp \left(-\frac{\|f(i, j) - f(k, l)\|^2}{2\sigma_r^2} \right)$$


Bilateral filter

Filter is data dependent


$$g(i, j) = \frac{\sum_{k,l} f(k, l) w(i, j, k, l)}{\sum_{k,l} w(i, j, k, l)}$$



Input



Bilateral filter




Output

$$w(i, j, k, l) = \exp \left(-\frac{(i - k)^2 + (j - l)^2}{2\sigma_d^2} - \frac{\|f(i, j) - f(k, l)\|^2}{2\sigma_r^2} \right)$$

Slow, but methods for approximation exist

Binary images

Neighborhoods


 \mathcal{N}_4


4 Neighborhood


 \mathcal{N}_8


8 Neighborhood


Erosion



(a)




(b)




(c)

Change a foreground pixel to background if it has a background pixel as a 4-neighbor.


Dilation



(a)



(b)




(c)

Change a background pixel to foreground if it has a foreground pixel as a 4-neighbor.



Original image




Initial threshold

Opening = Dilate(Erode)



Original image




After opening

Closing = Erode(Dilate)



Original image



After closing

Thank you for your attention.

UNIVERSITÄT  **BONN**