



UNIVERSITÄT **BONN**

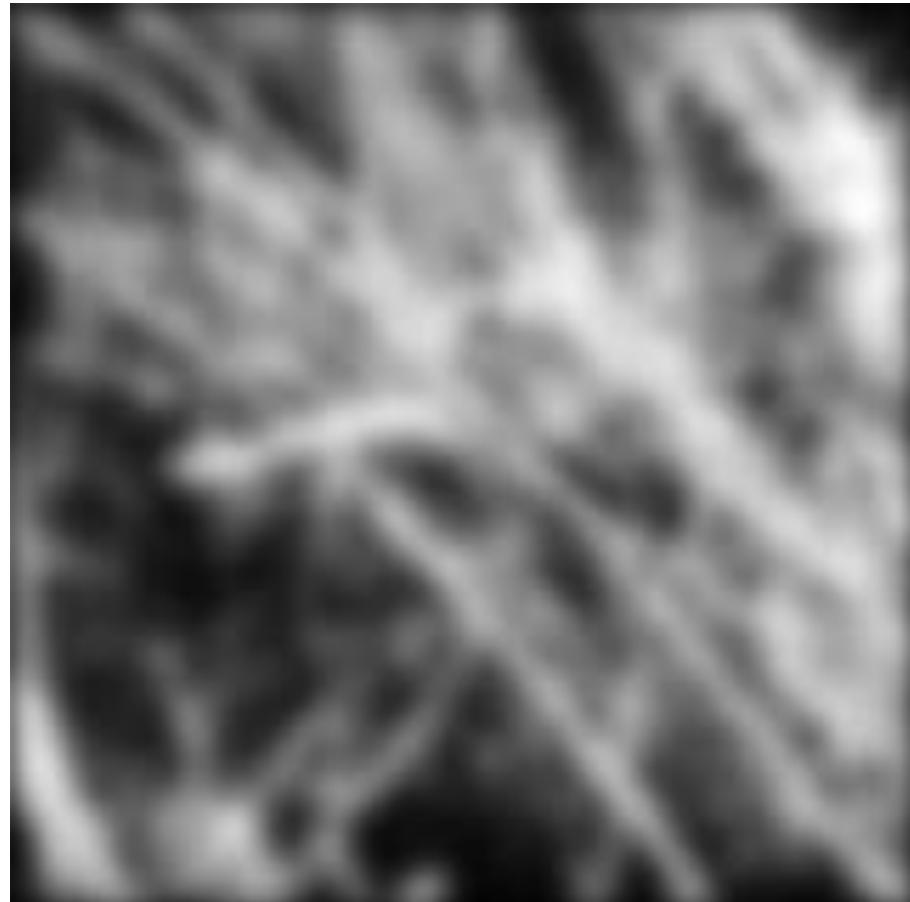
Juergen Gall

Fourier Transform and Image Pyramids
MA-INF 2201 - Computer Vision
WS24/25

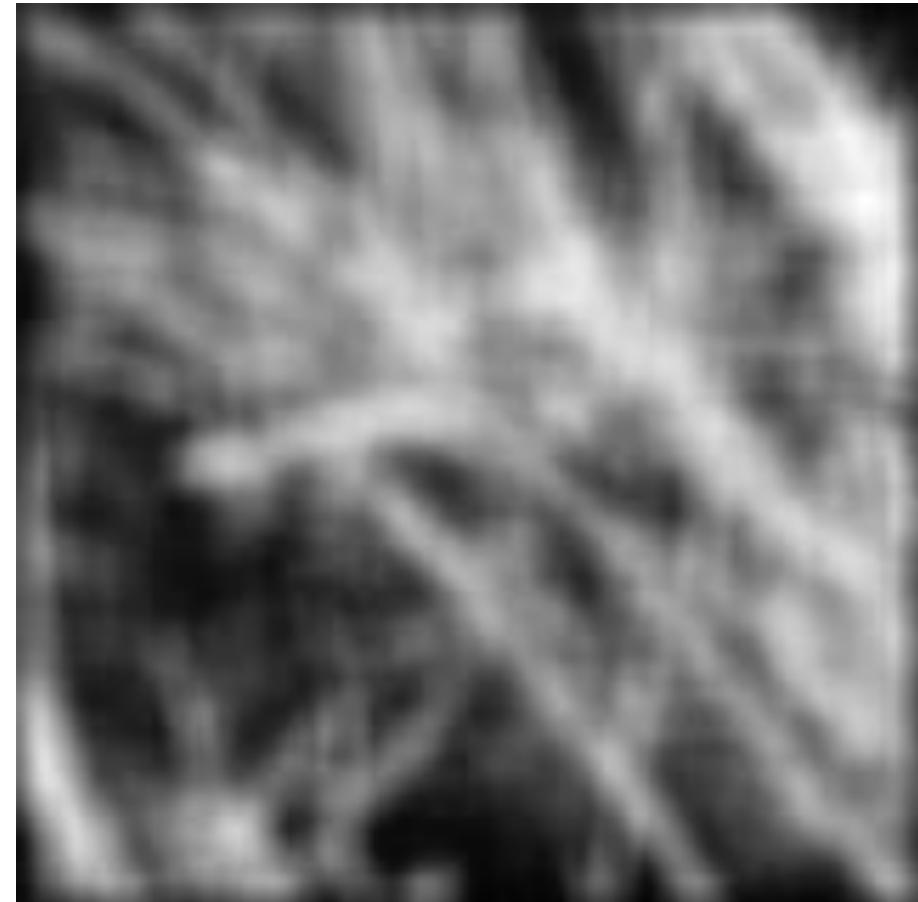
Linear filtering

Gauss vs. box filter

Gaussian



Box filter



A sum of sines

Our building block:

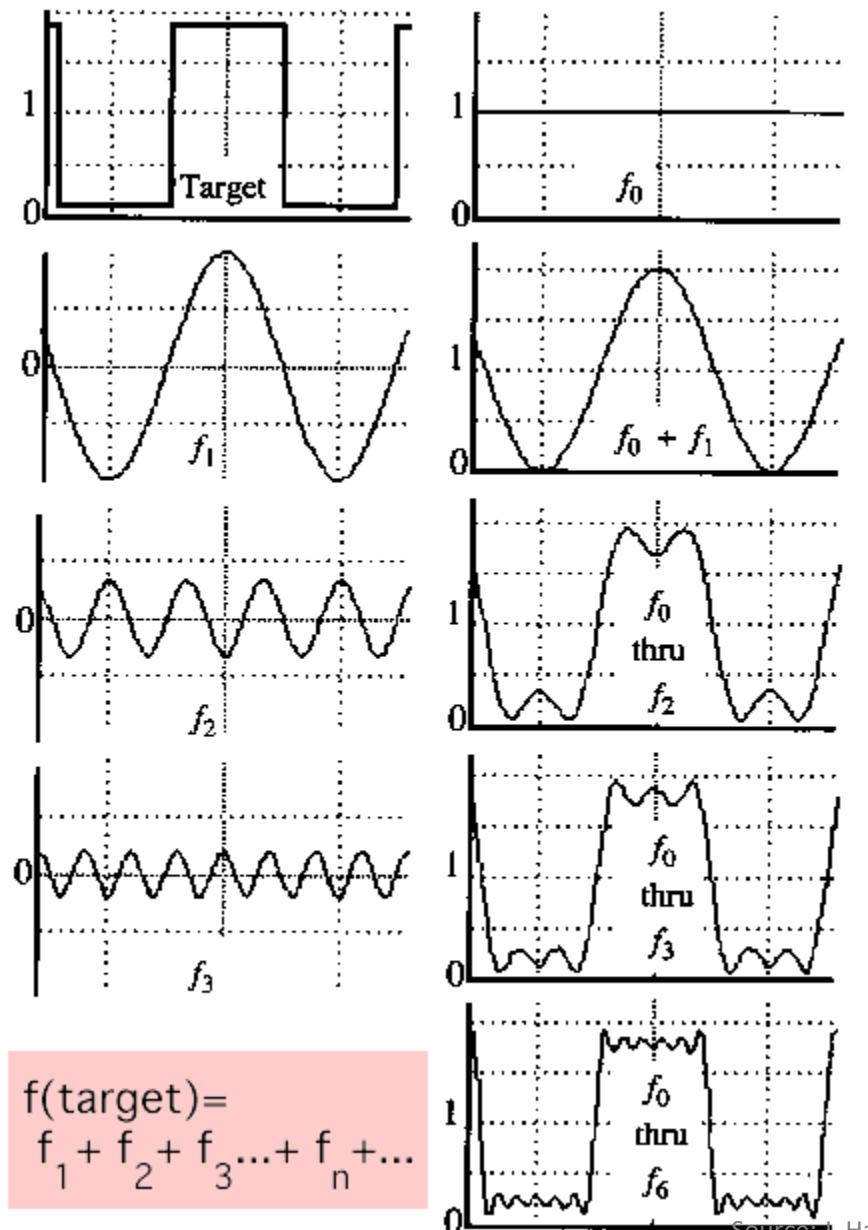
$$A \sin(\omega x + \phi)$$

A: Amplitude

ω : angular frequency

ϕ : phase

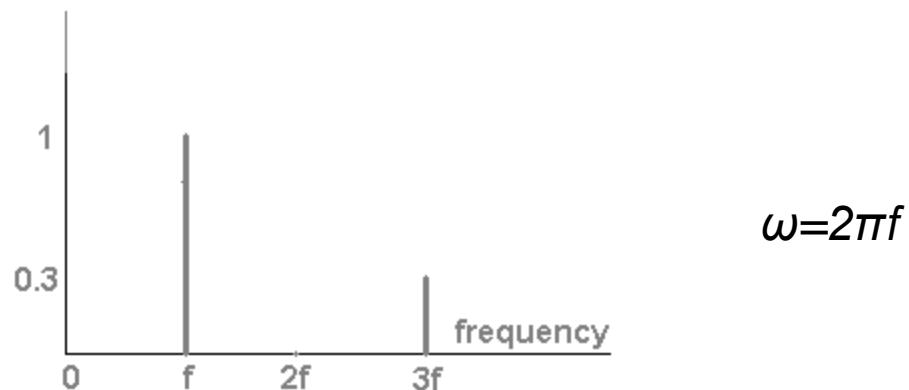
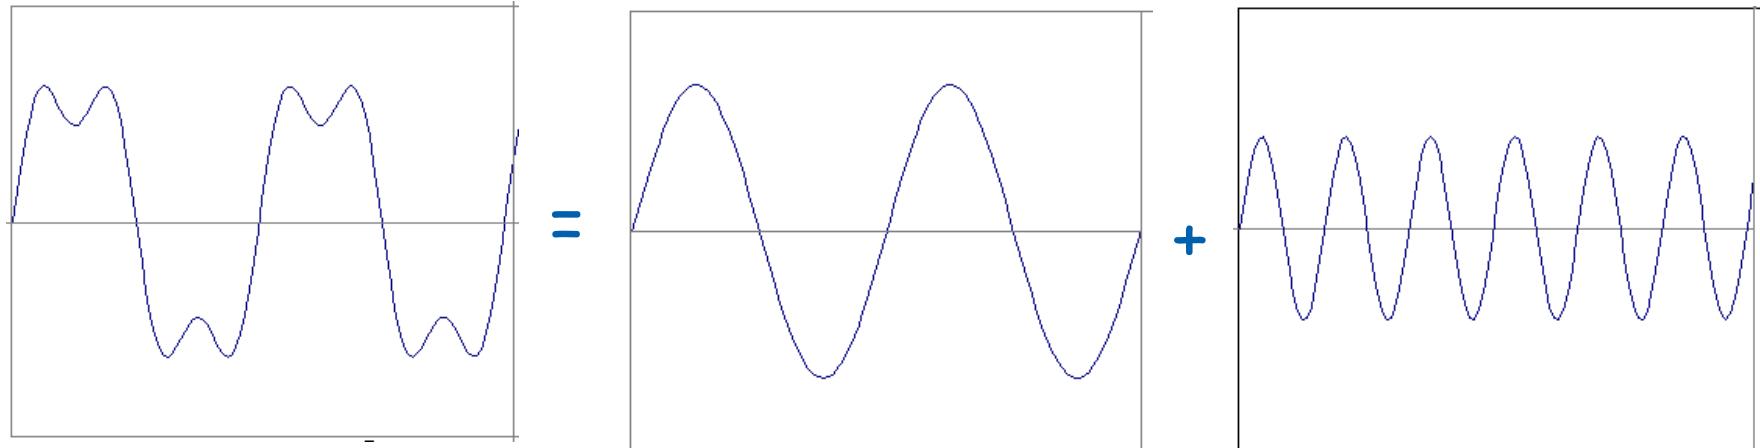
Add enough of them to get any signal $f(x)$ you want!



Source: J. Hays

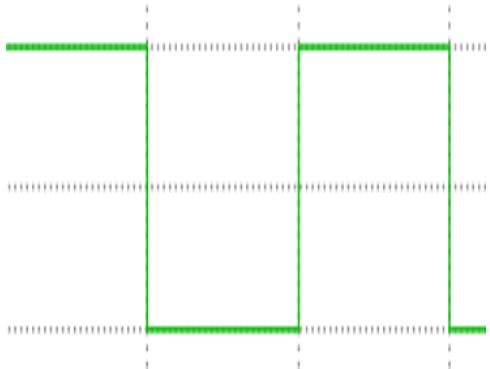
Frequency Spectra

Example: $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f) t)$



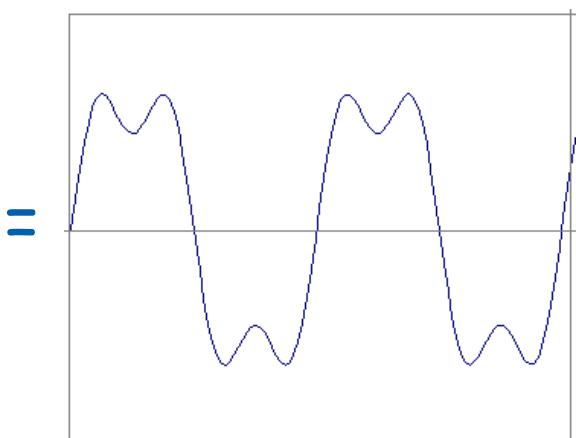
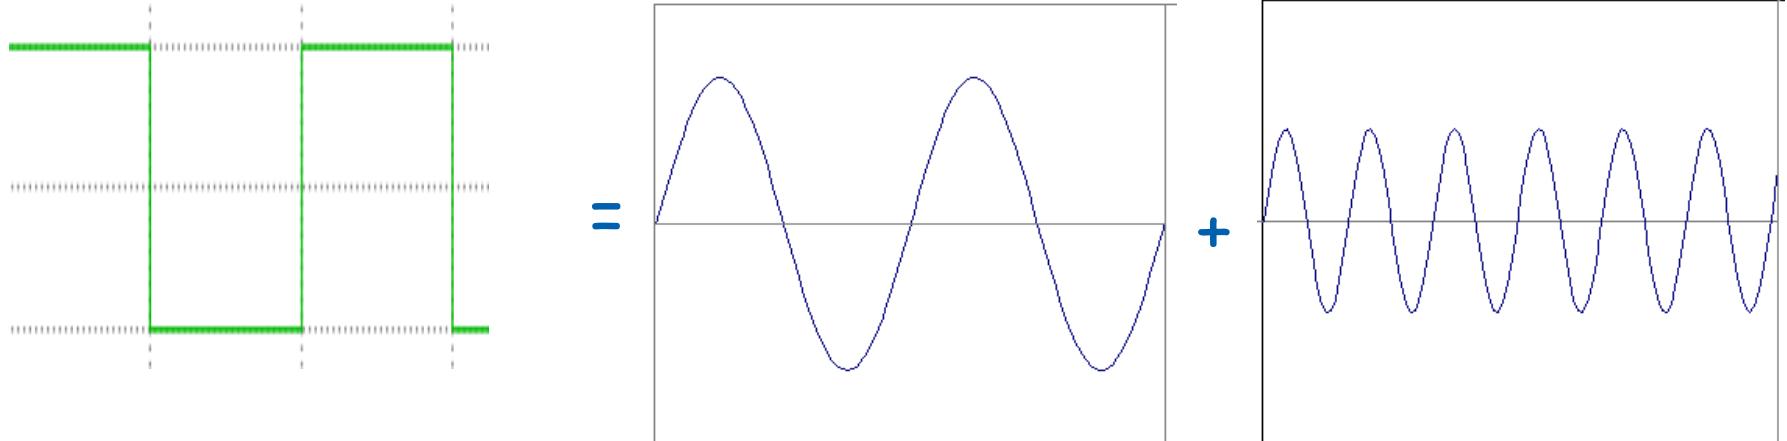
Source: A. Efros

Frequency Spectra



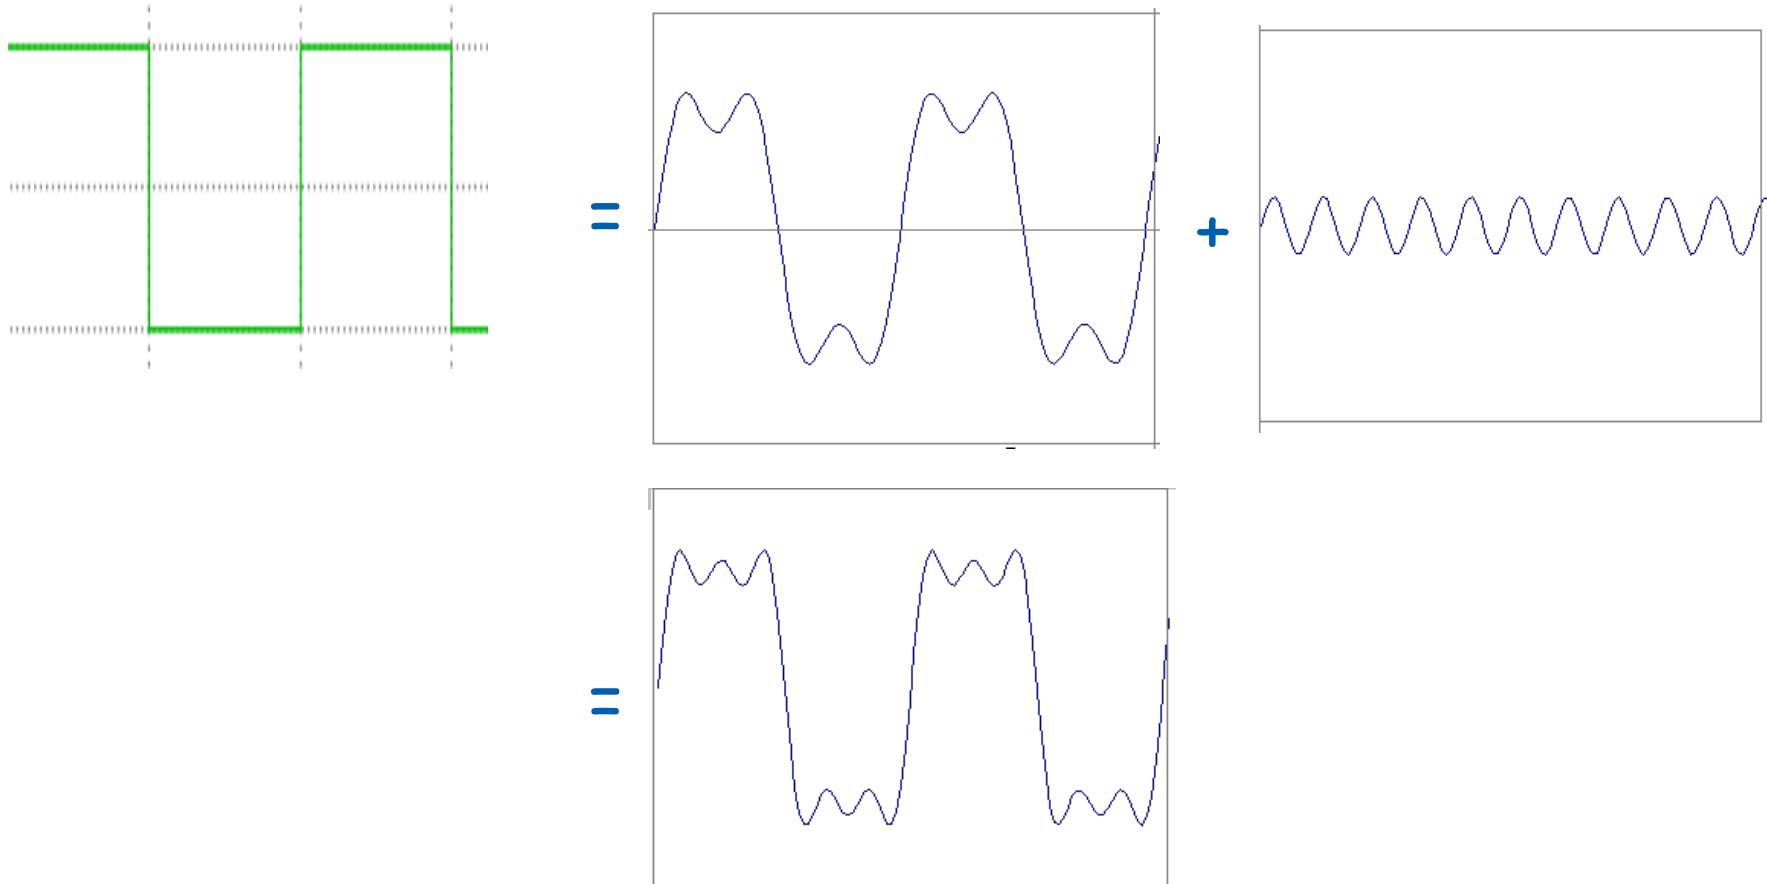
Source: J. Hays

Frequency Spectra



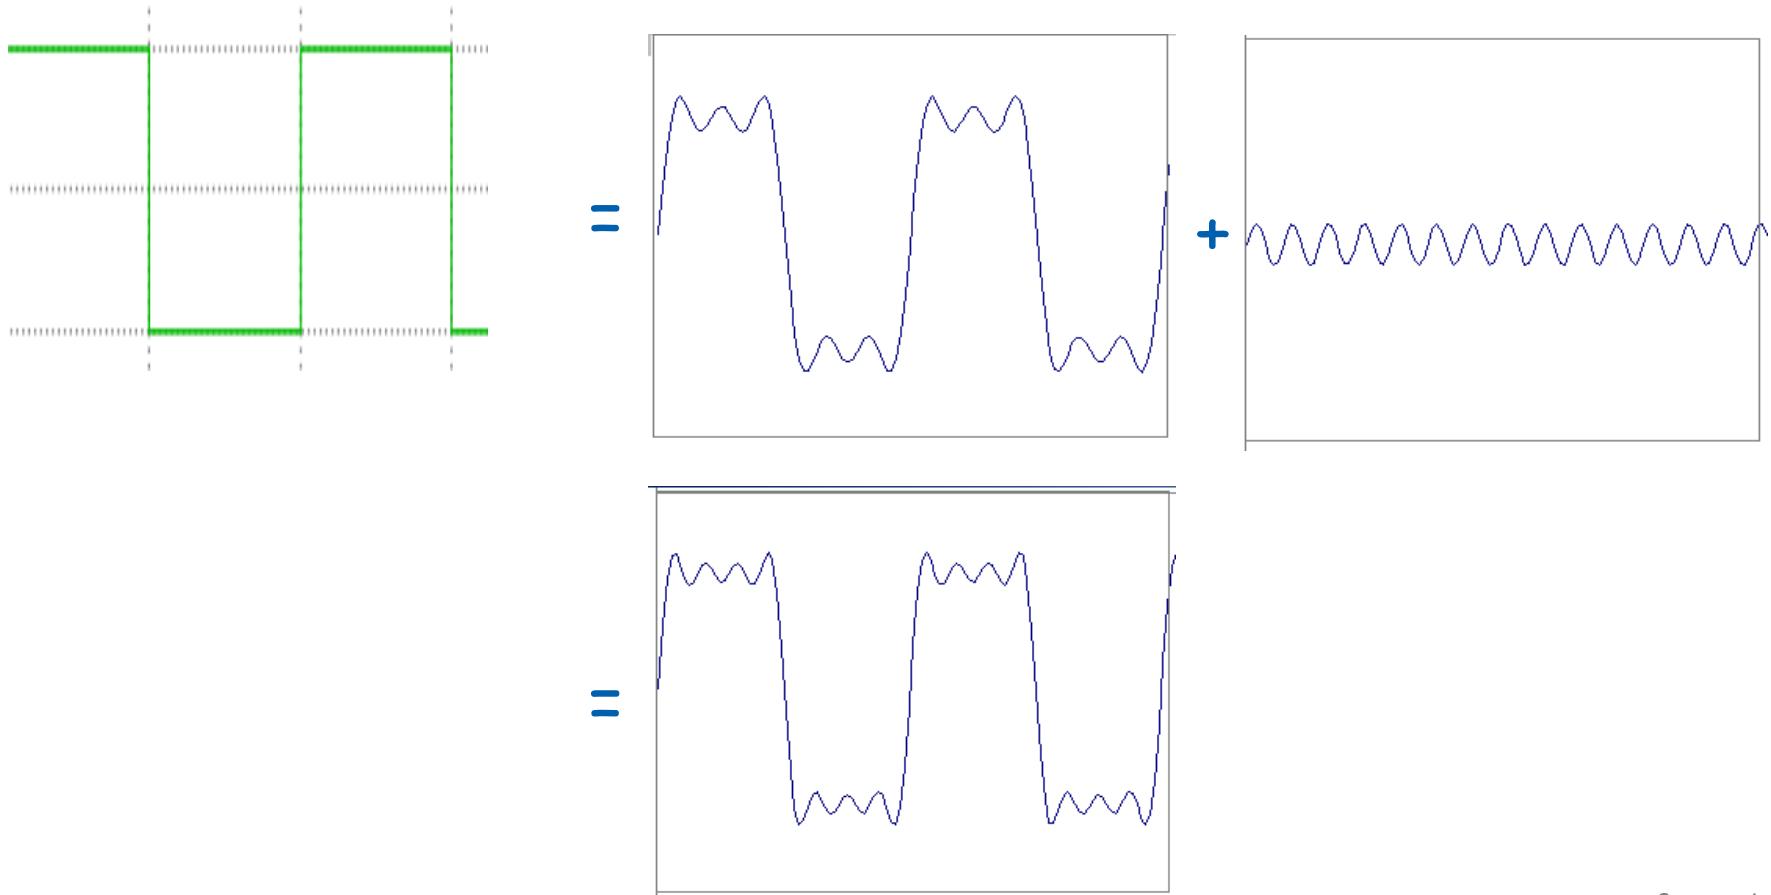
Source: J. Hays

Frequency Spectra



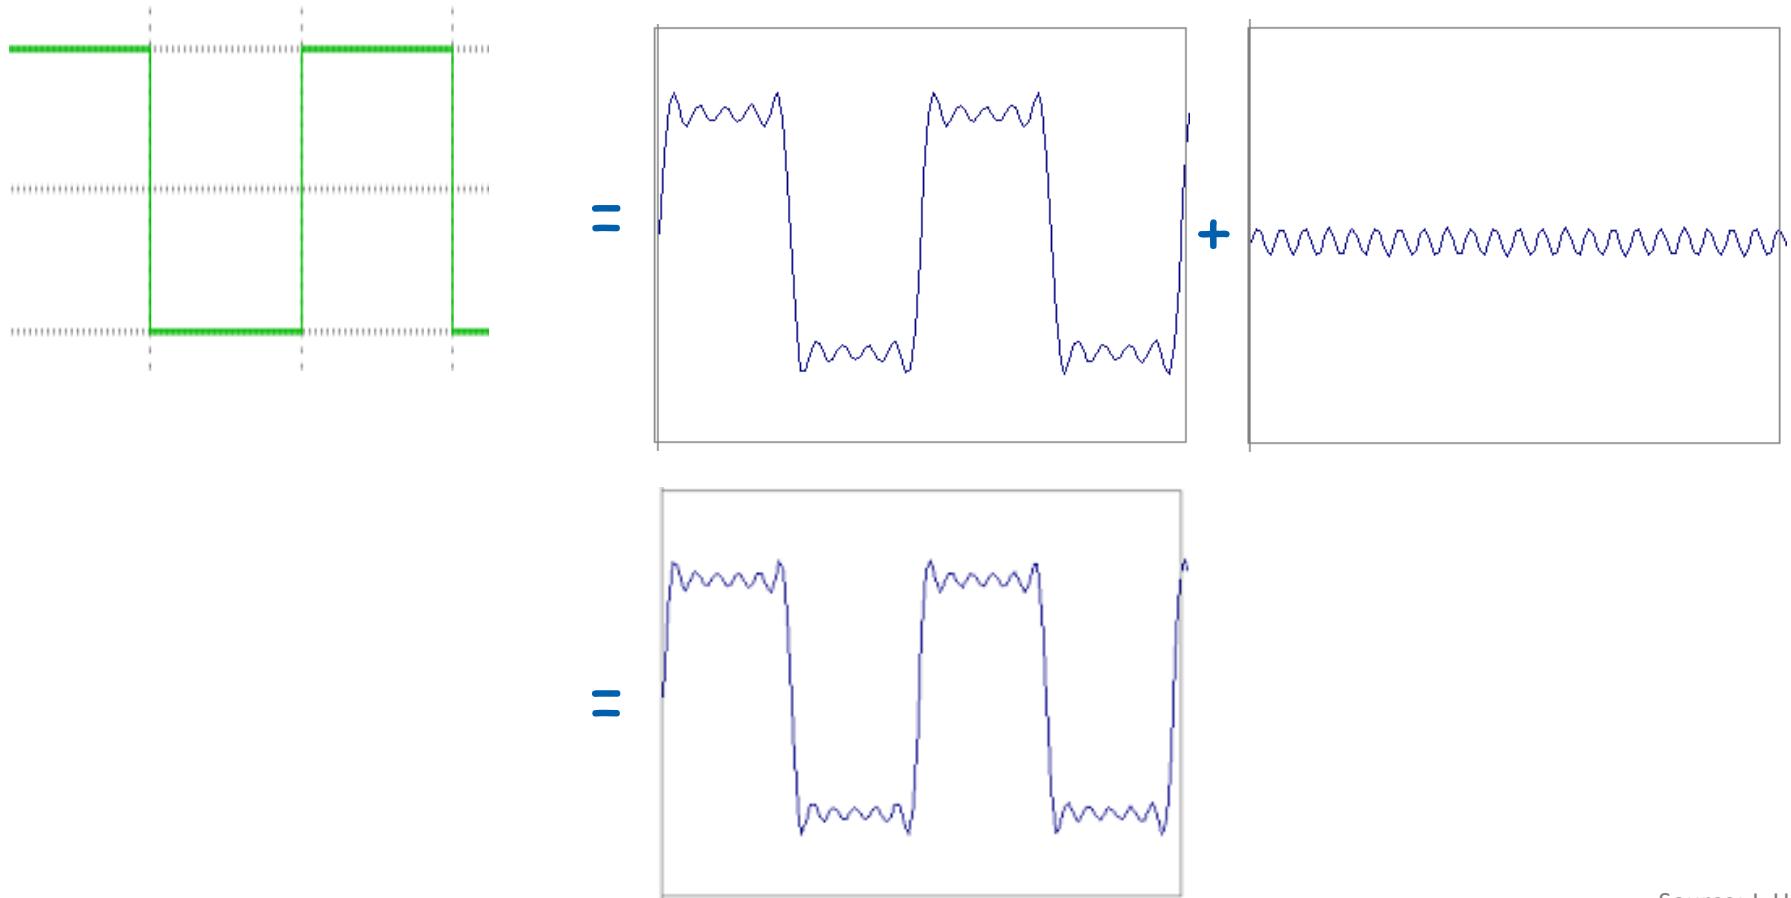
Source: J. Hays

Frequency Spectra



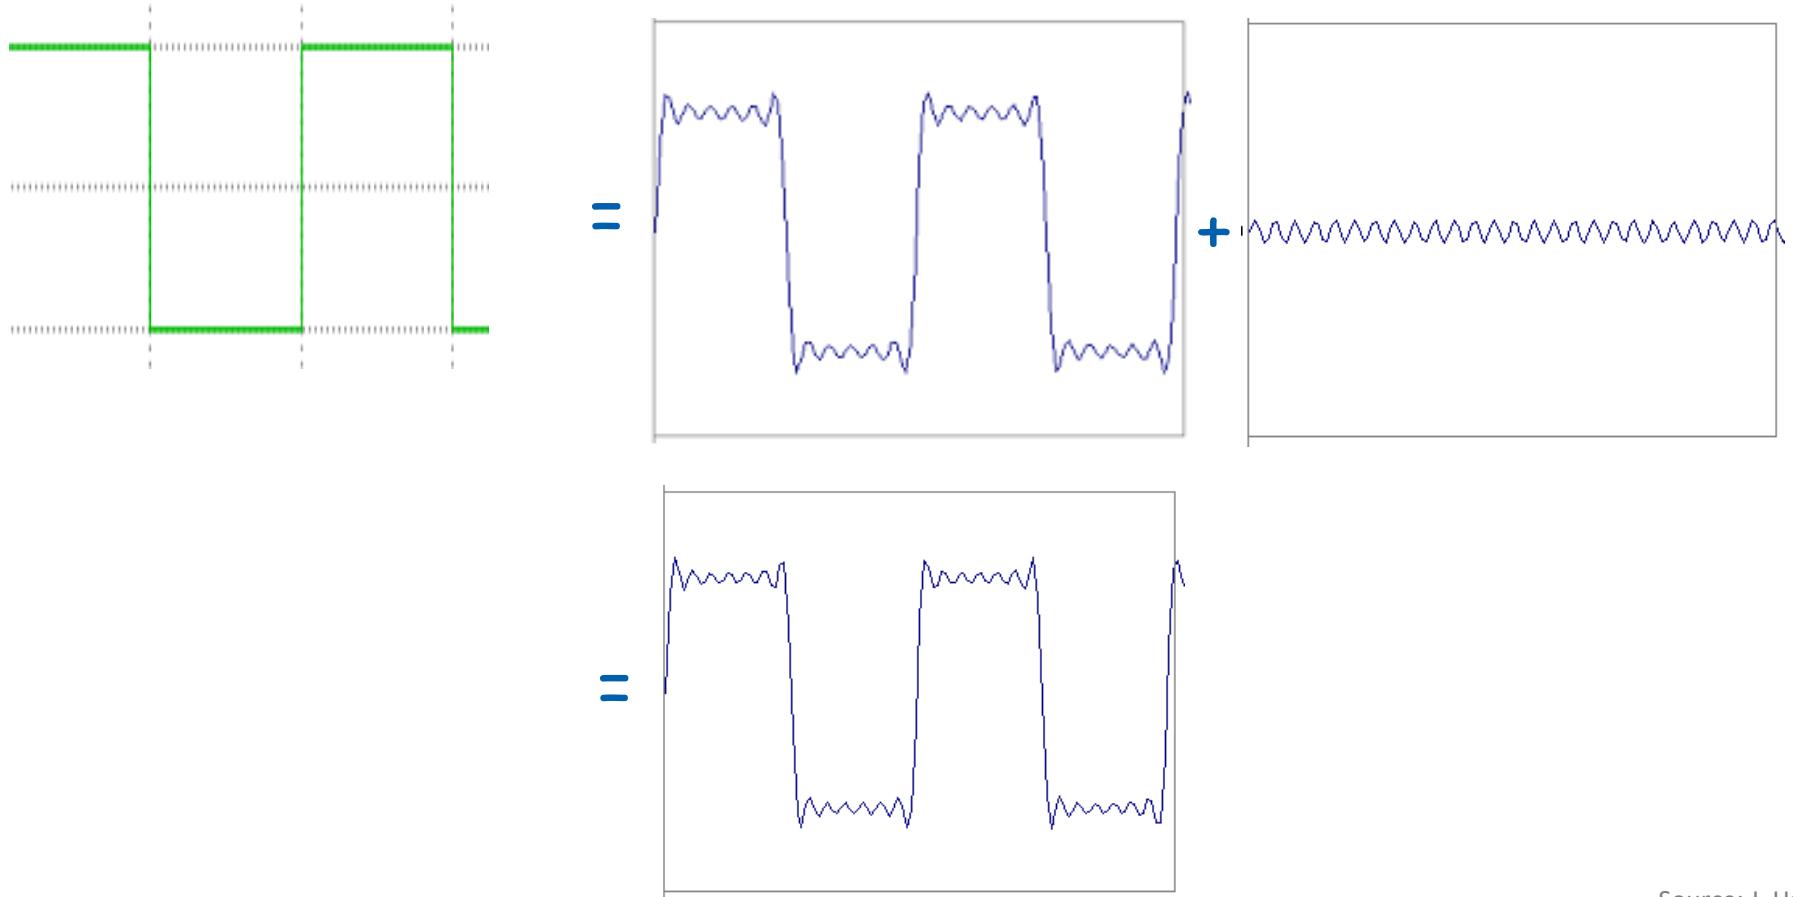
Source: J. Hays

Frequency Spectra



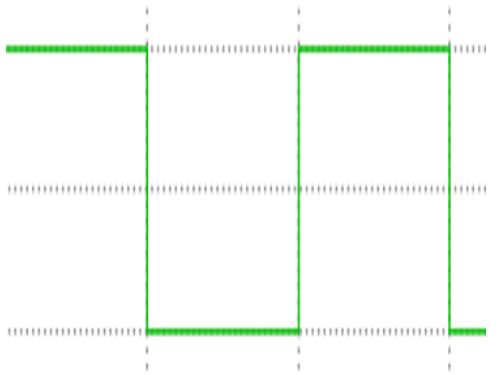
Source: J. Hays

Frequency Spectra



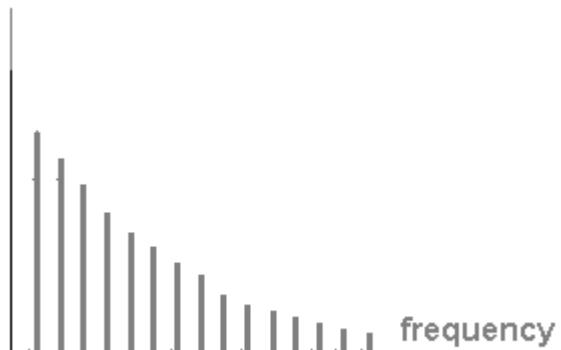
Source: J. Hays

Frequency Spectra



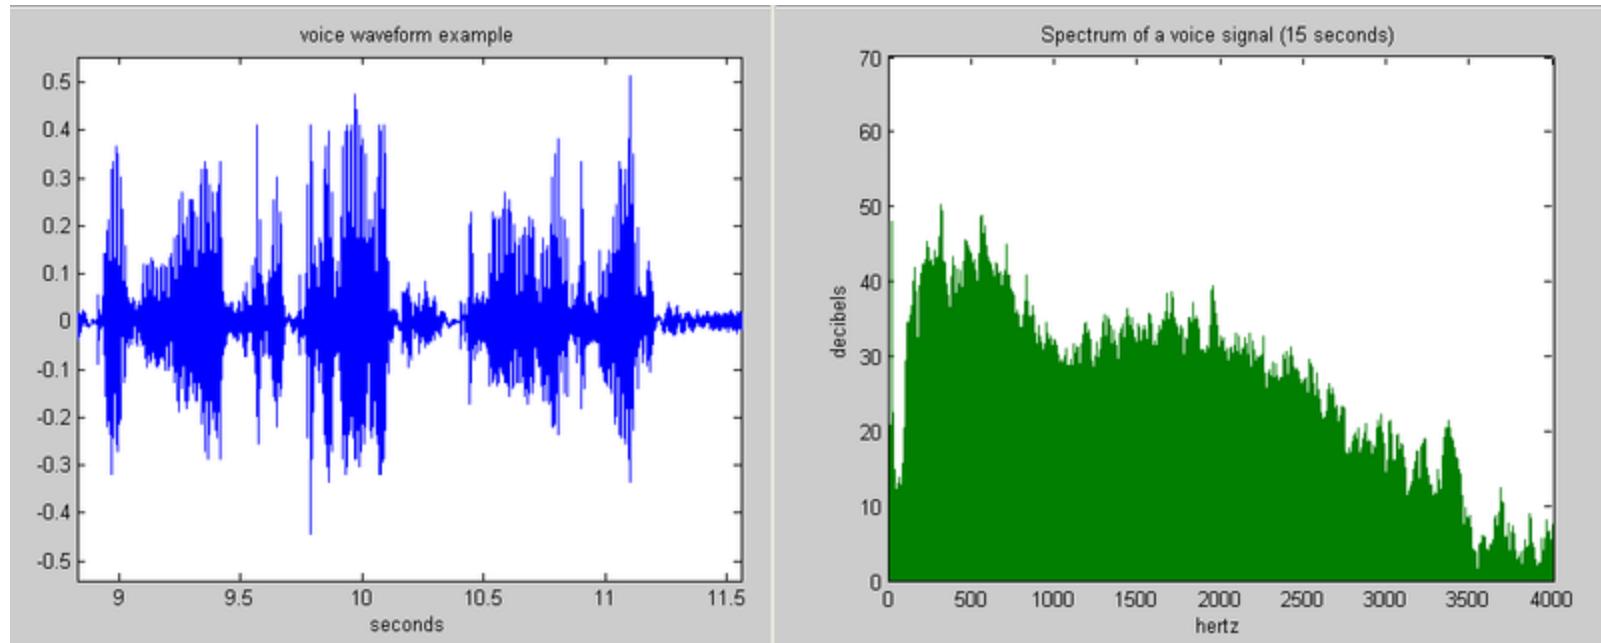
=

$$A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kt)$$



Example: Music

We think of music in terms of frequencies at different magnitudes



Source: D. Hoiem

Continuous Fourier transform

Euler's formula:

$$e^{\pm j\theta} = \cos(\theta) \pm j\sin(\theta)$$

Complex numbers:

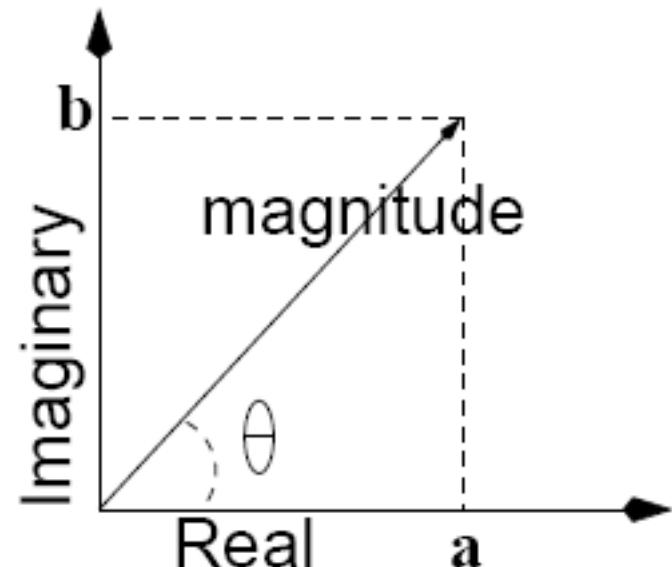
$$x = a + jb, \text{ where } j = \sqrt{-1}$$

$$\phi(x) = \tan^{-1}(b/a) \quad \text{phase}$$

$$|x| = \sqrt{a^2 + b^2} \quad \text{magnitude}$$

$$x^* = a - jb \quad \text{conjugate}$$

$$x = |x| e^{j\phi(x)}$$



Continuous Fourier transform

1D Continuous Fourier Transform:

$$f(x) = \int_{u=-\infty}^{\infty} F(u)e^{+2\pi jux} du$$

The inverse Fourier transform

$$F(u) = \int_{x=-\infty}^{\infty} f(x)e^{-2\pi jux} dx$$

The Fourier transform

2D Continuous Fourier Transform:

$$f(x, y) = \int_{u=-\infty}^{\infty} \int_{v=-\infty}^{\infty} F(u,v)e^{+2\pi j(ux+vy)} dudv$$

The inverse transform

$$F(u, v) = \int_{x=-\infty}^{\infty} \int_{y=-\infty}^{\infty} f(x,y)e^{-2\pi j(ux+vy)} dxdy$$

The transform

Discrete Fourier transform

1D Discrete Fourier Transform:

$$f(x) = \sum_{u=0}^{N-1} F(u) e^{\frac{+2\pi j u x}{N}}$$

$$F(u) = \frac{1}{N} \sum_{x=0}^{N-1} f(x) e^{\frac{-2\pi j u x}{N}}$$

2D Discrete Fourier Transform:

$$f(x, y) = \sum_{u=0}^{N-1} \sum_{v=0}^{M-1} F(u, v) e^{+2\pi j (\frac{ux}{N} + \frac{vy}{M})}$$

$$F(u, v) = \frac{1}{N} \frac{1}{M} \sum_{x=0}^{N-1} \sum_{y=0}^{M-1} f(x, y) e^{-2\pi j (\frac{ux}{N} + \frac{vy}{M})}$$

Fourier Transform

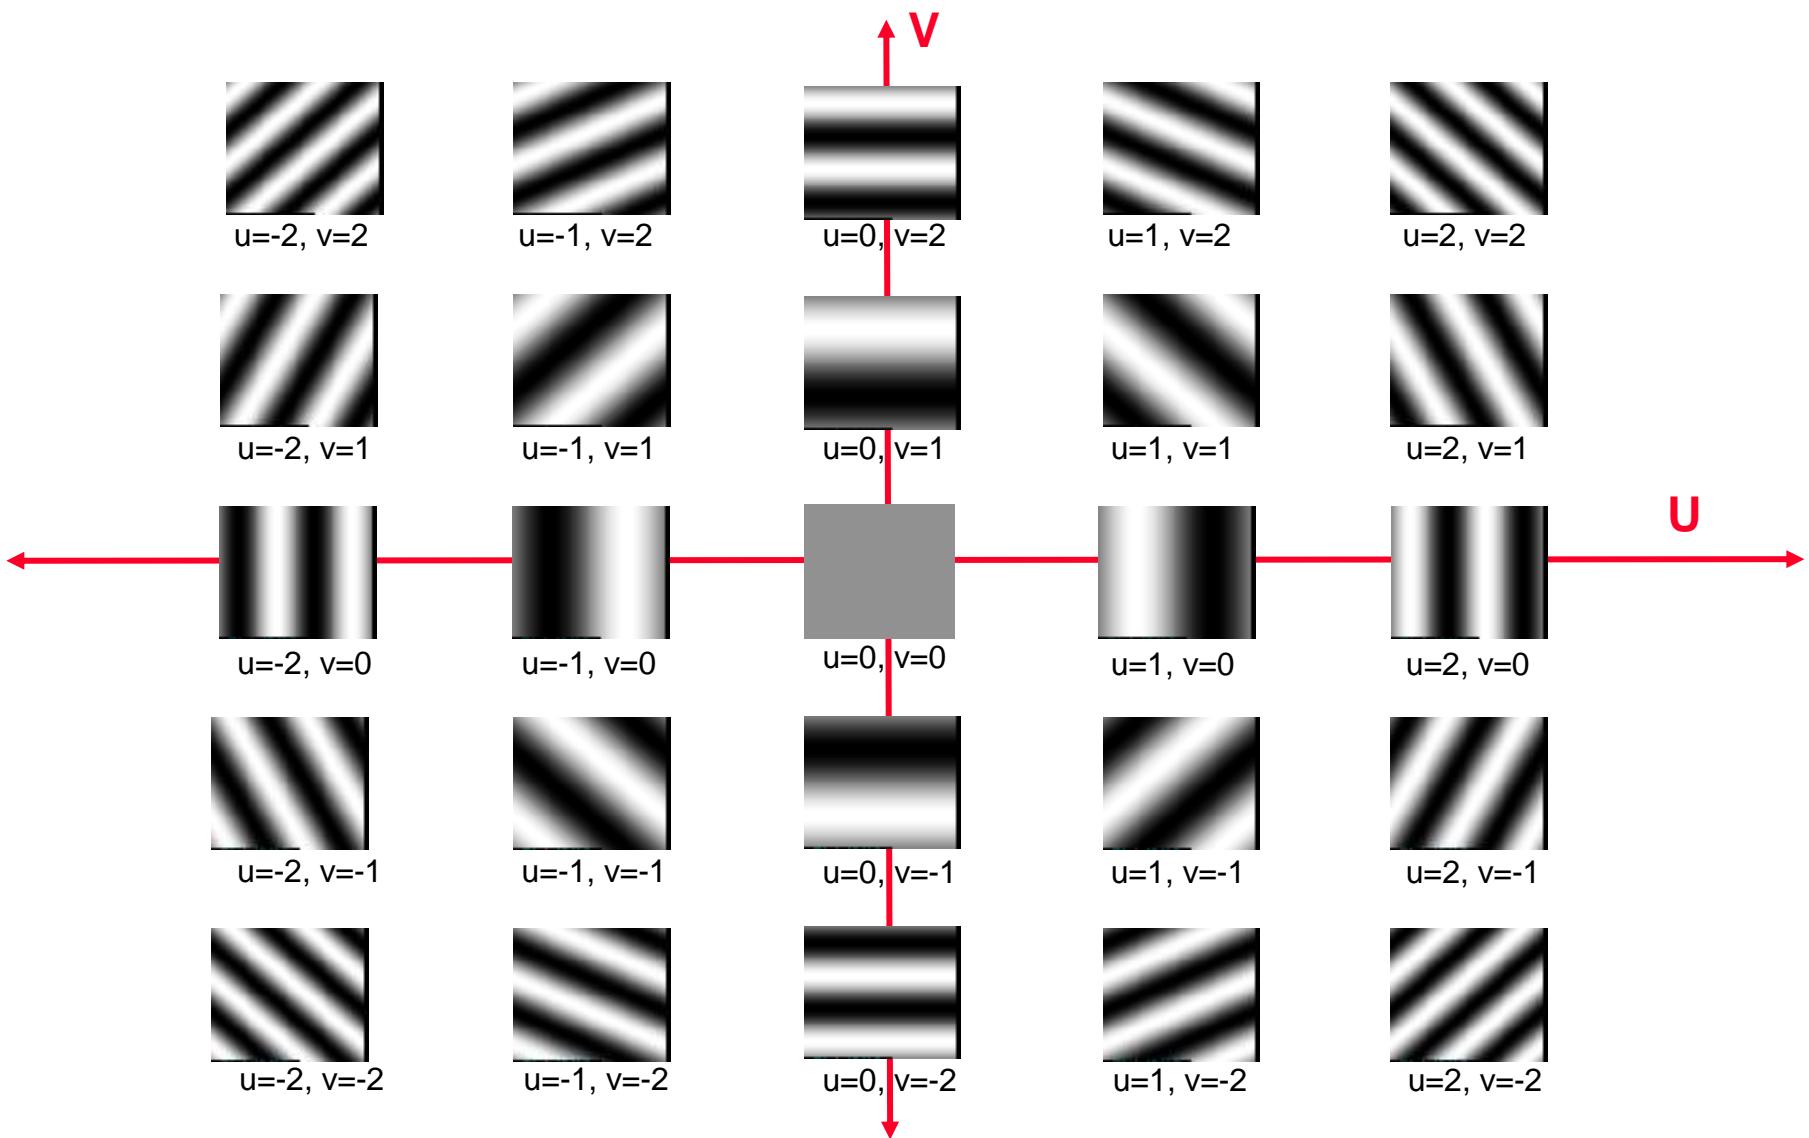
Fourier transform stores the magnitude and phase at each frequency

- Magnitude encodes how much signal there is at a particular frequency
- Phase encodes spatial information (indirectly)

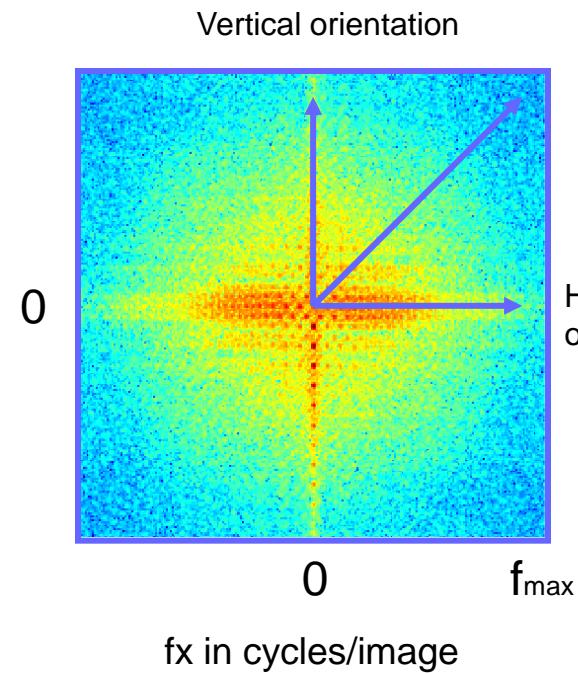
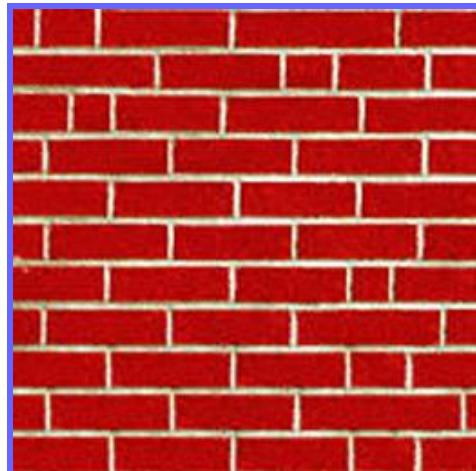
$$\text{Amplitude: } A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

$$\text{Phase: } \phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

The 2D Basis Functions

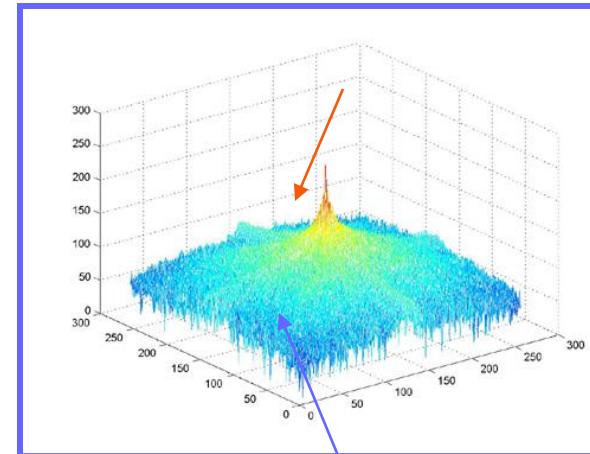


How to interpret a Fourier Spectrum



Log power spectrum

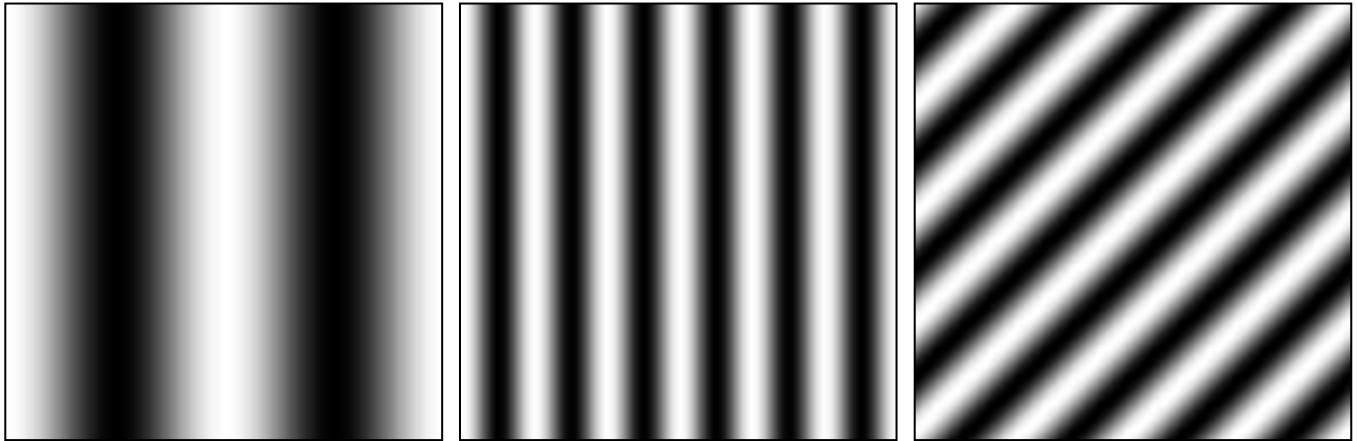
Low spatial frequencies



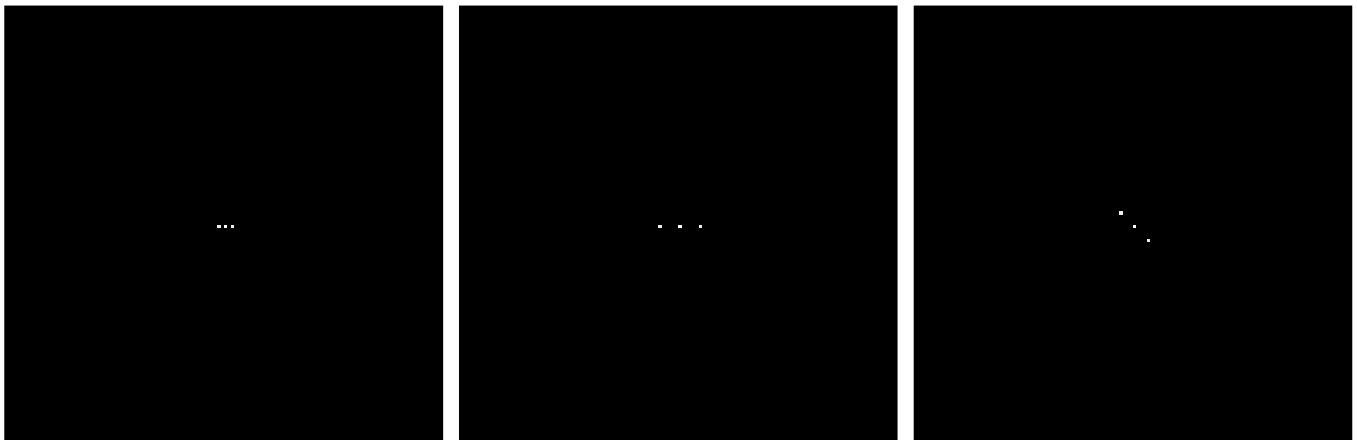
**High
spatial
frequencies**

Fourier analysis in images

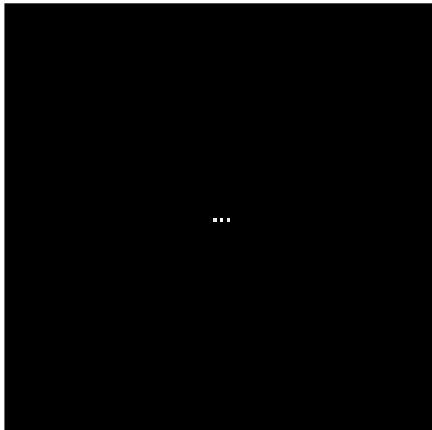
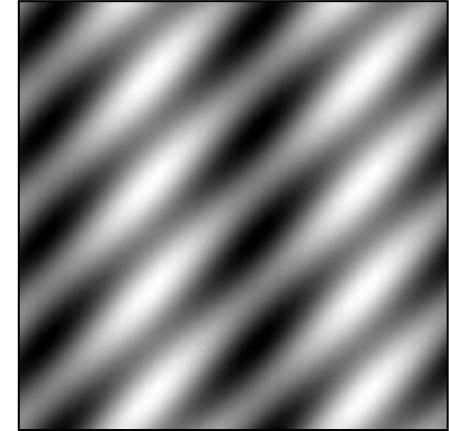
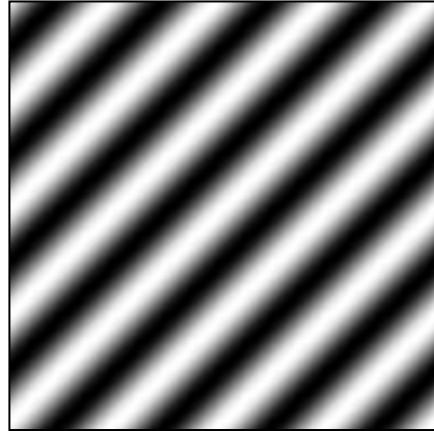
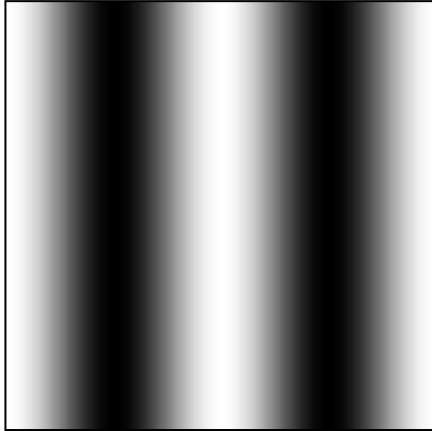
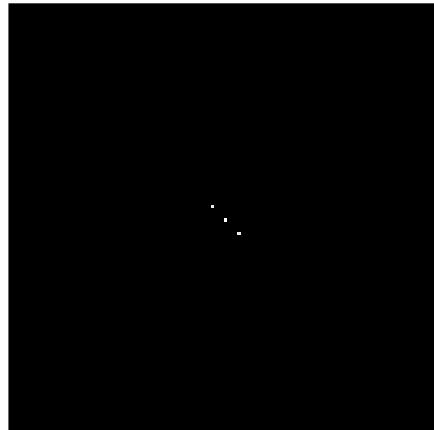
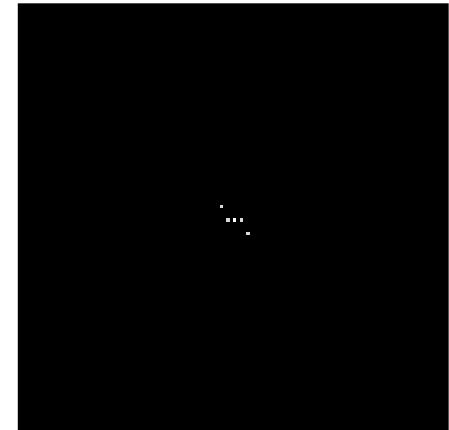
Intensity Image



Fourier Image

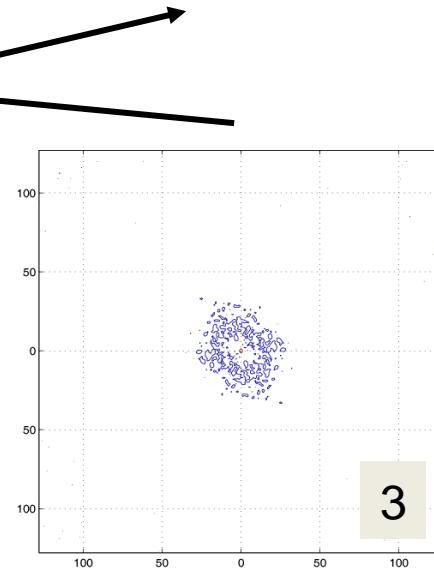
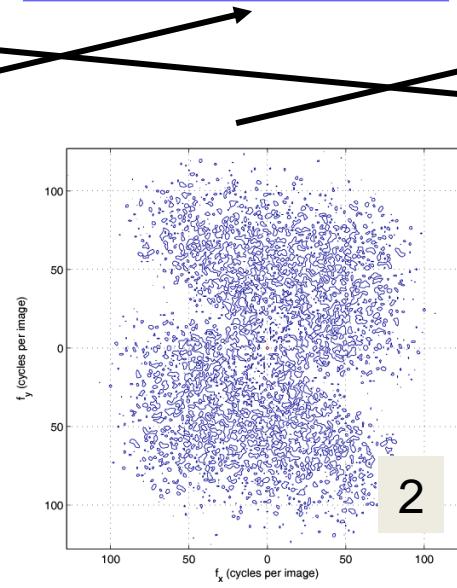
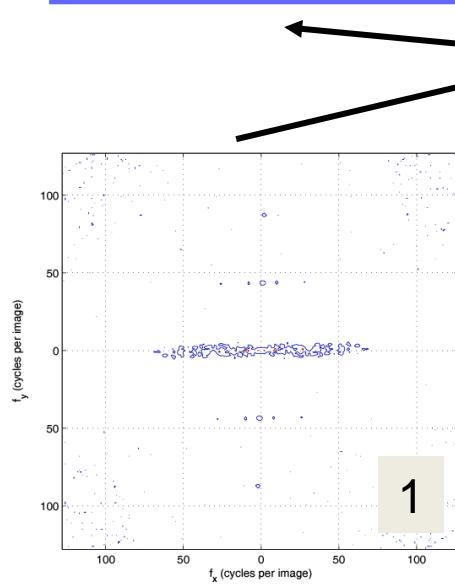
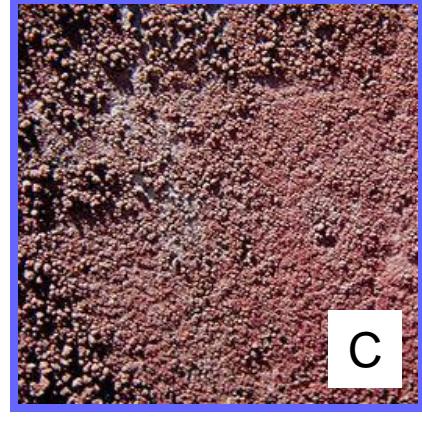
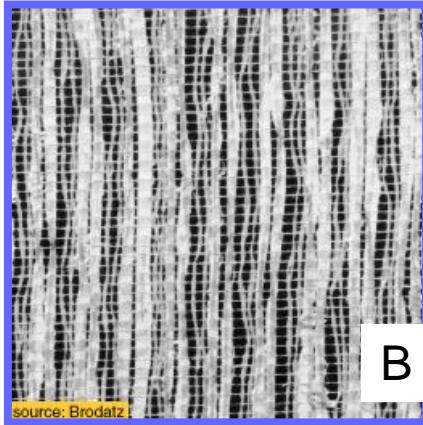


Signals can be composed

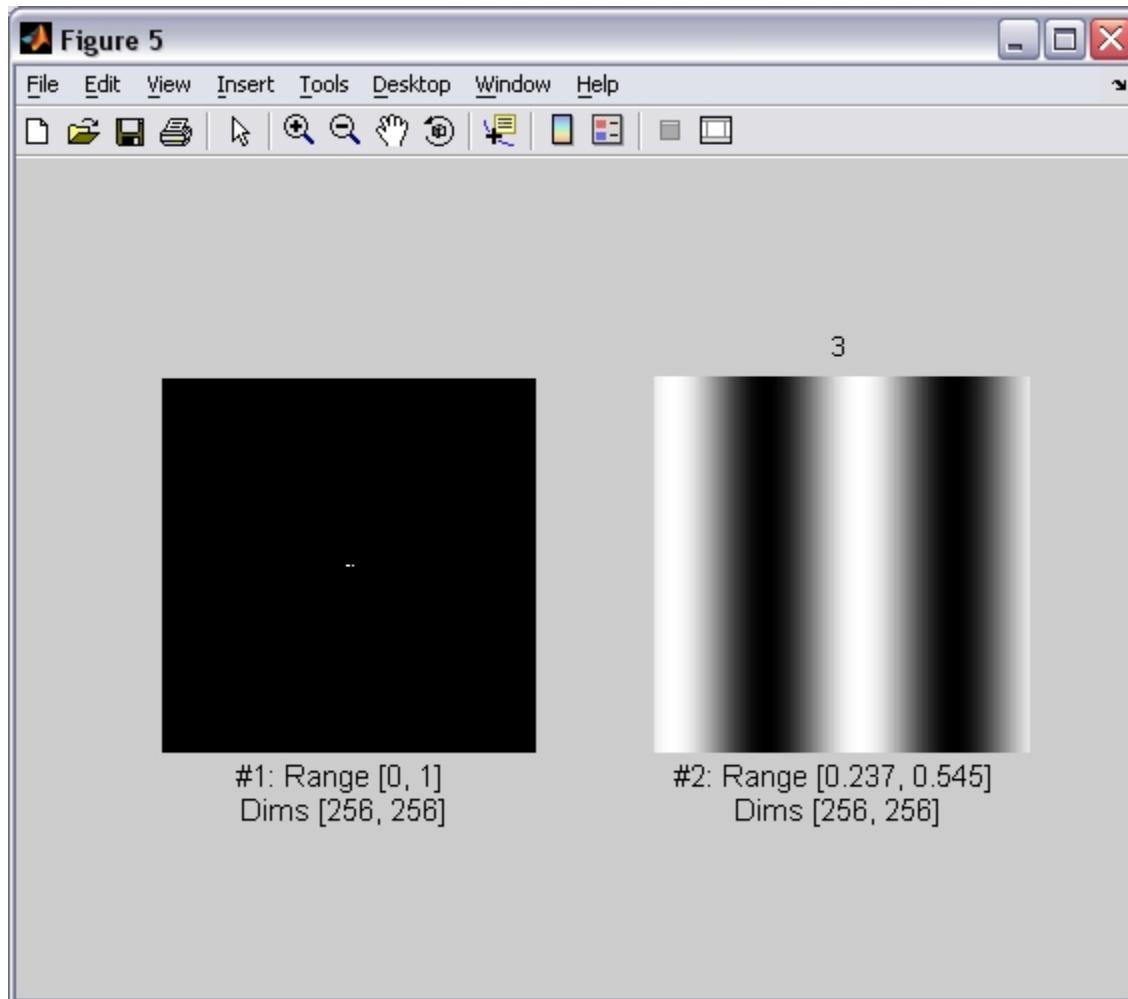
 $+$  $=$ 

Source: J. Hays

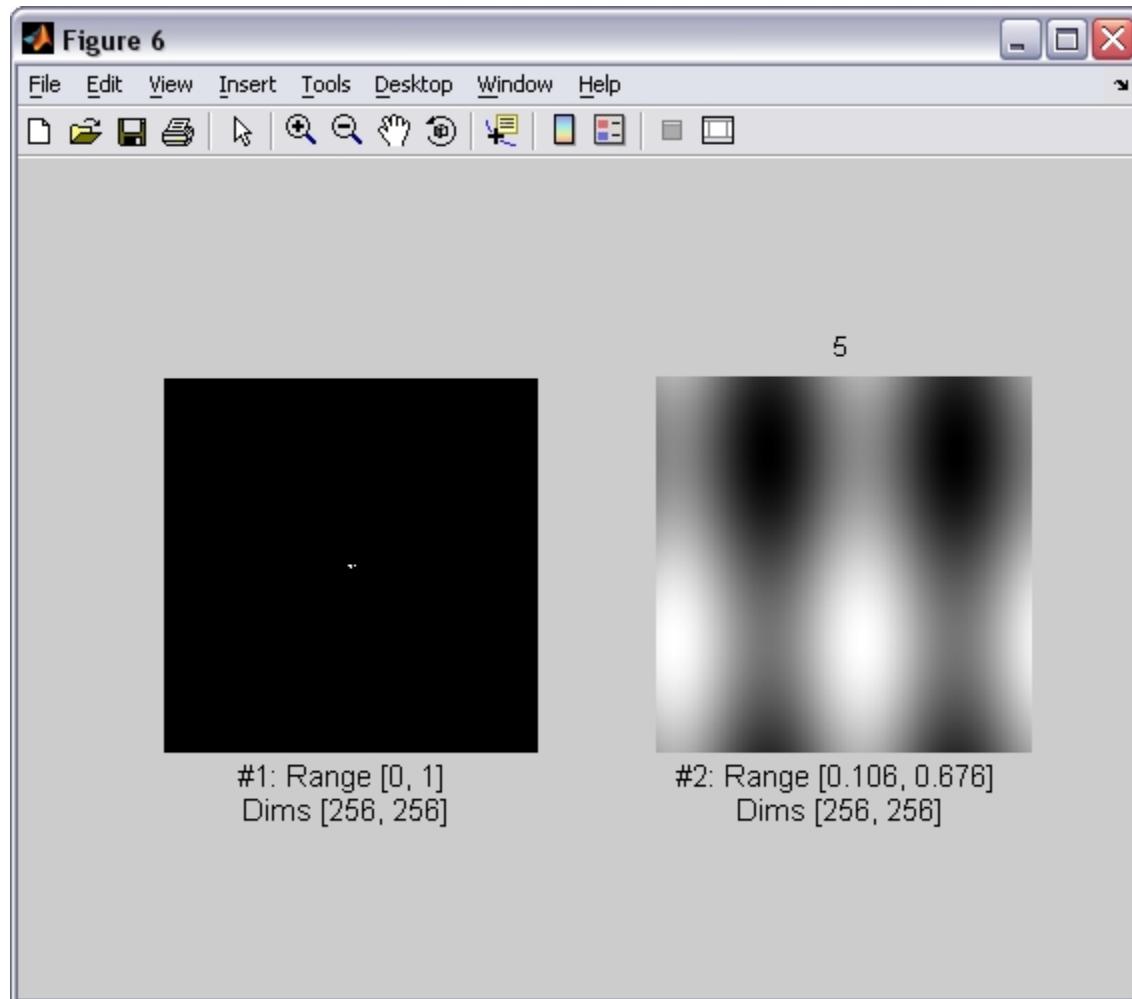
Fourier Amplitude Spectrum

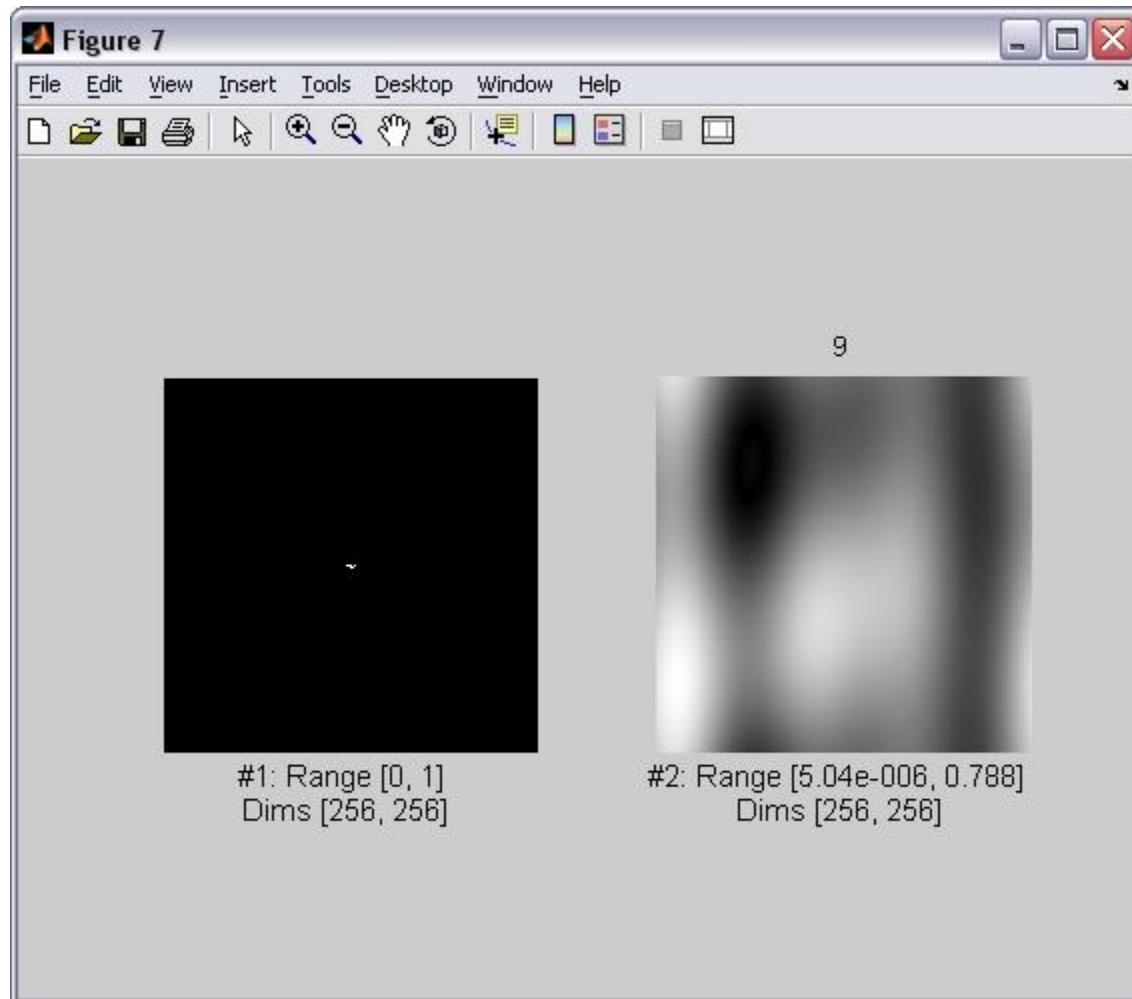


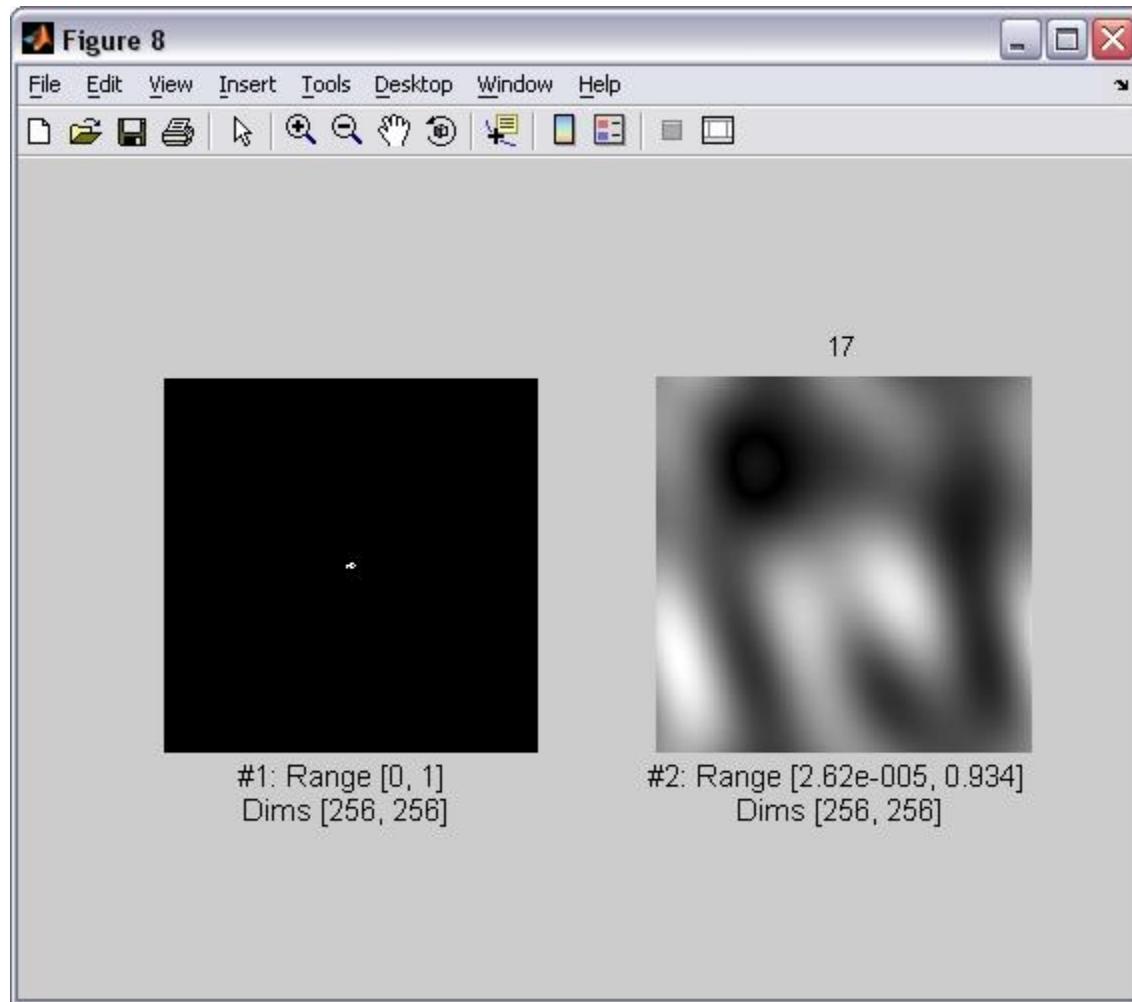
3 Fourier components

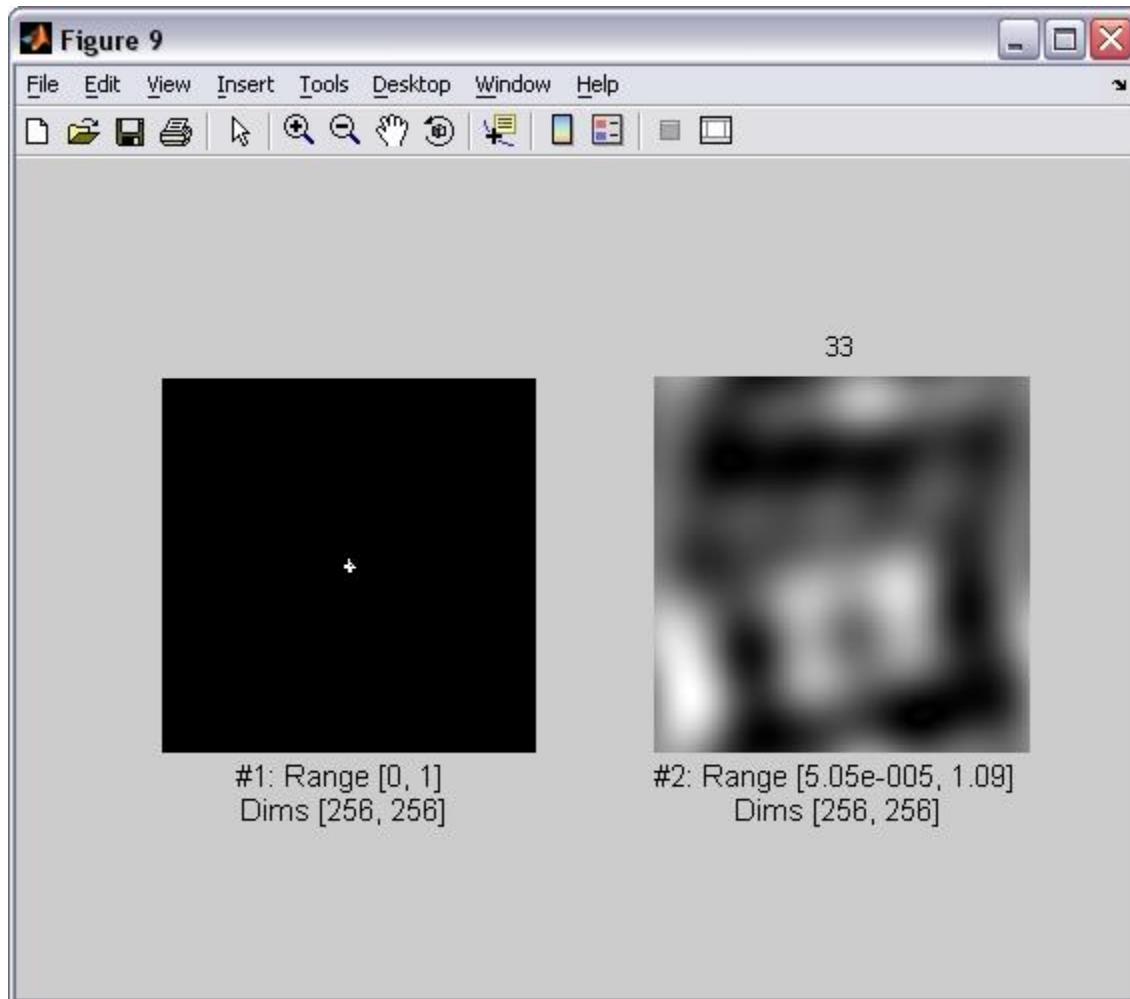


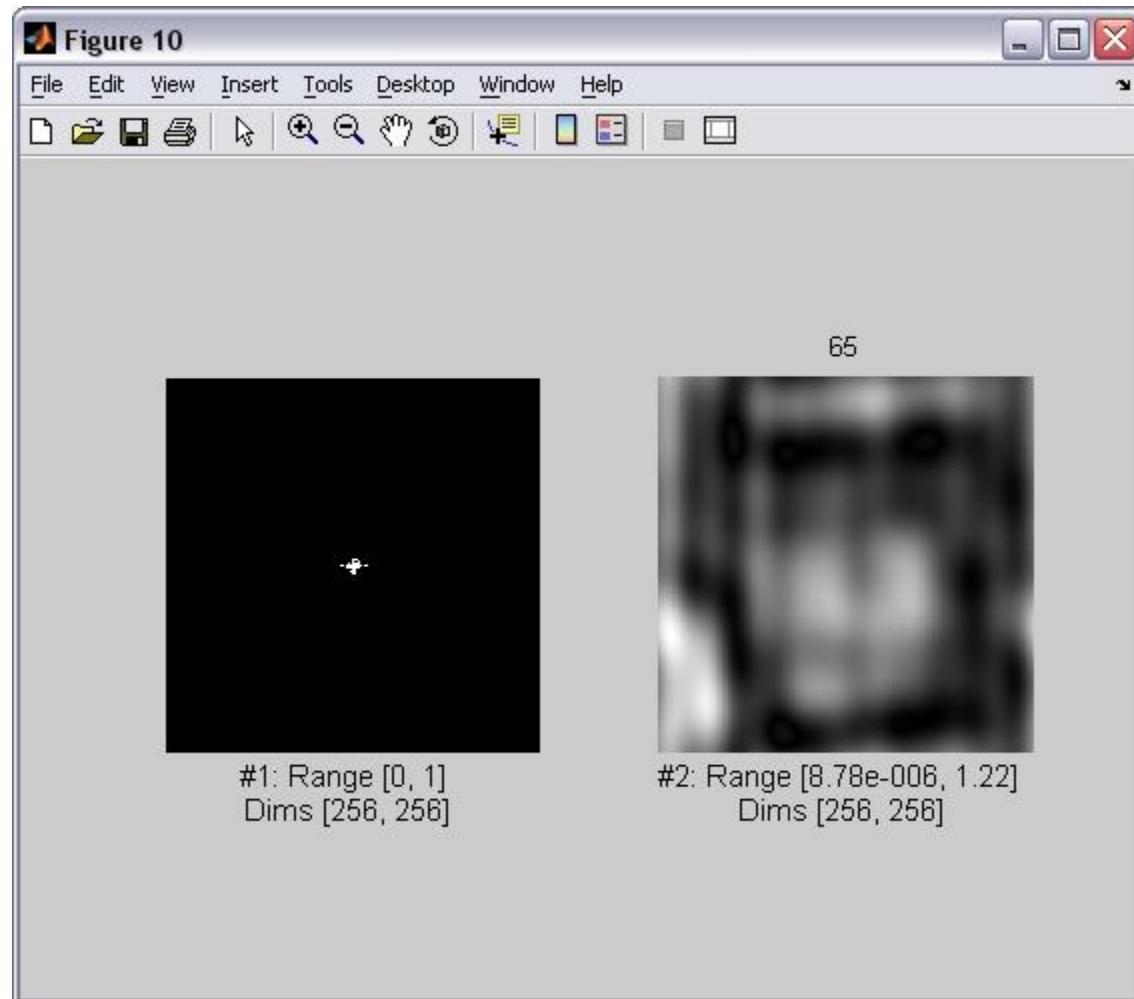
Descending order of magnitude.



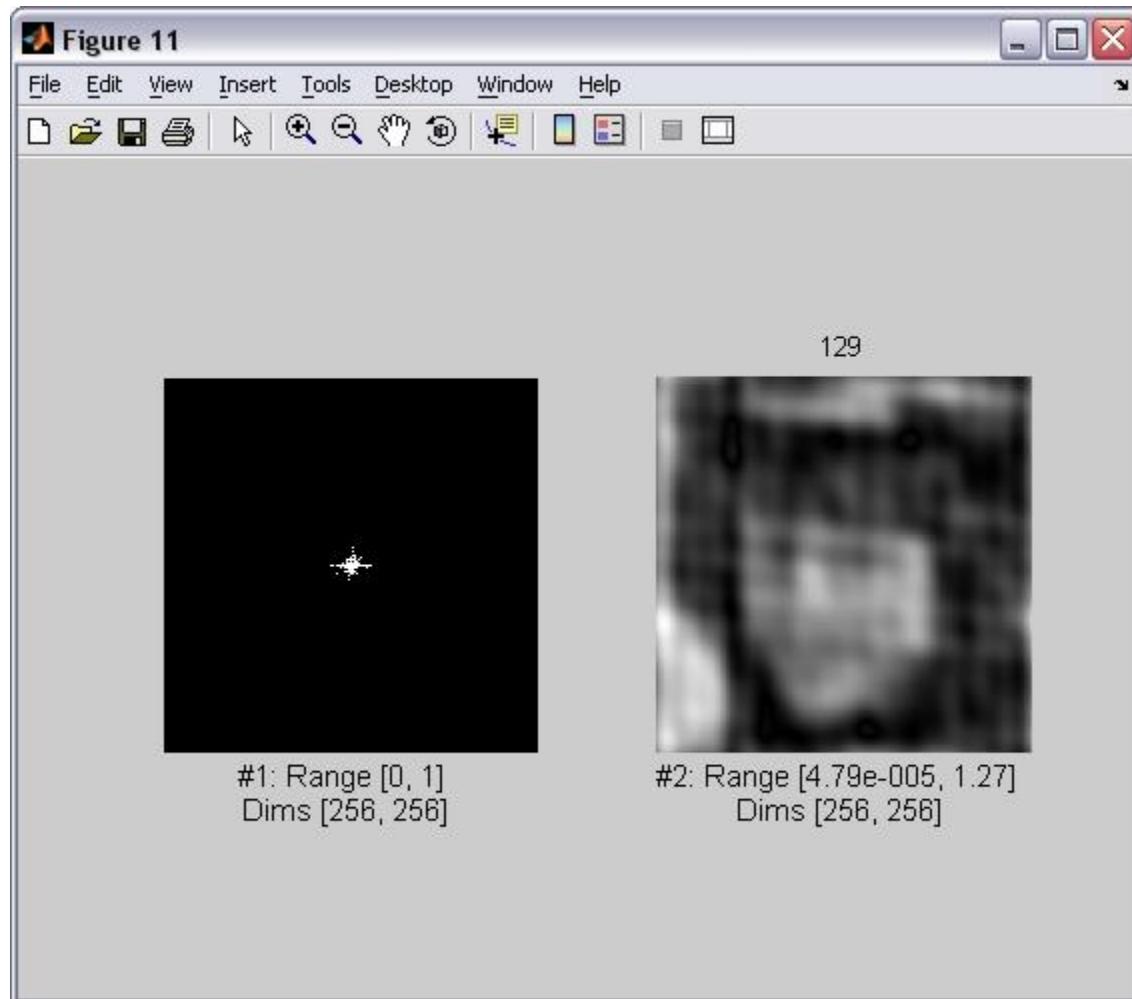




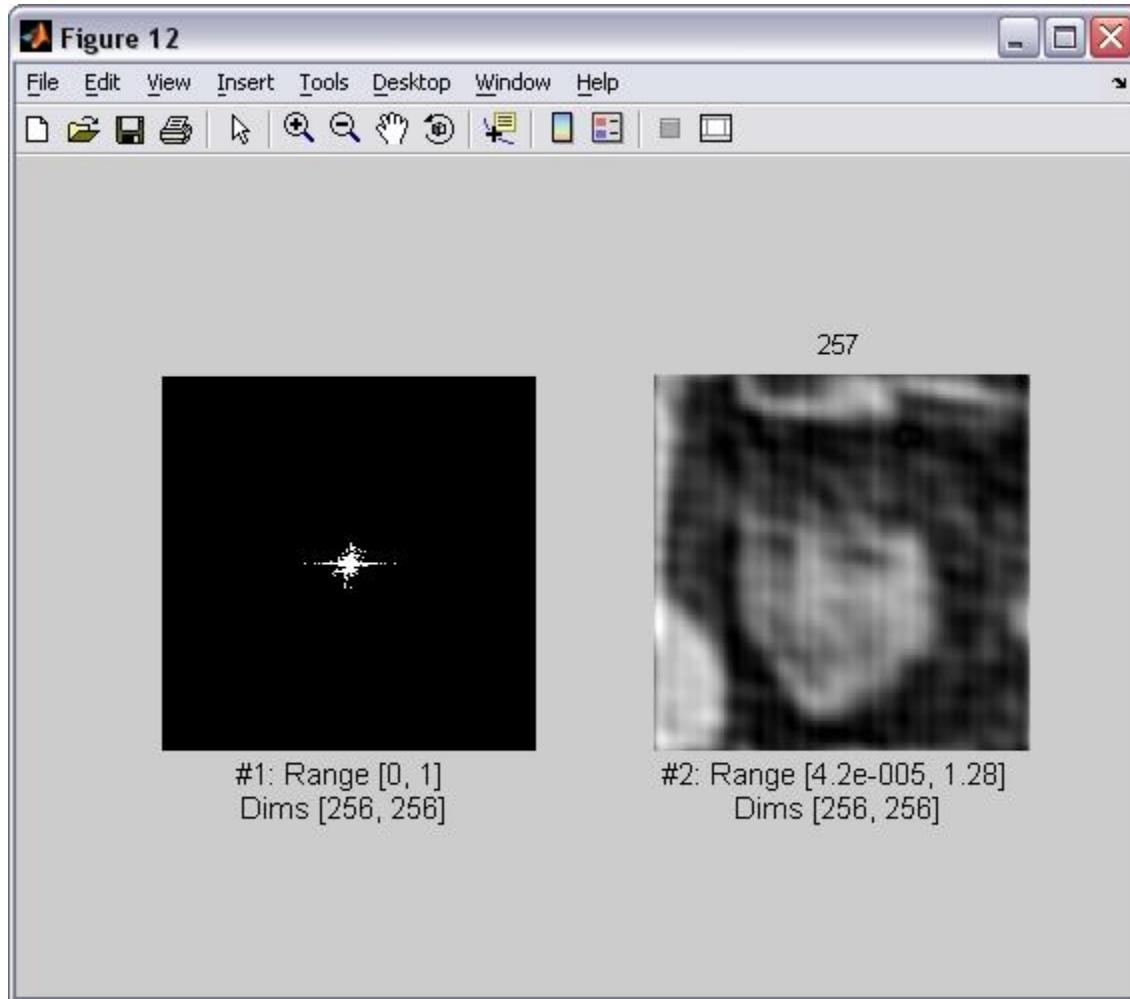




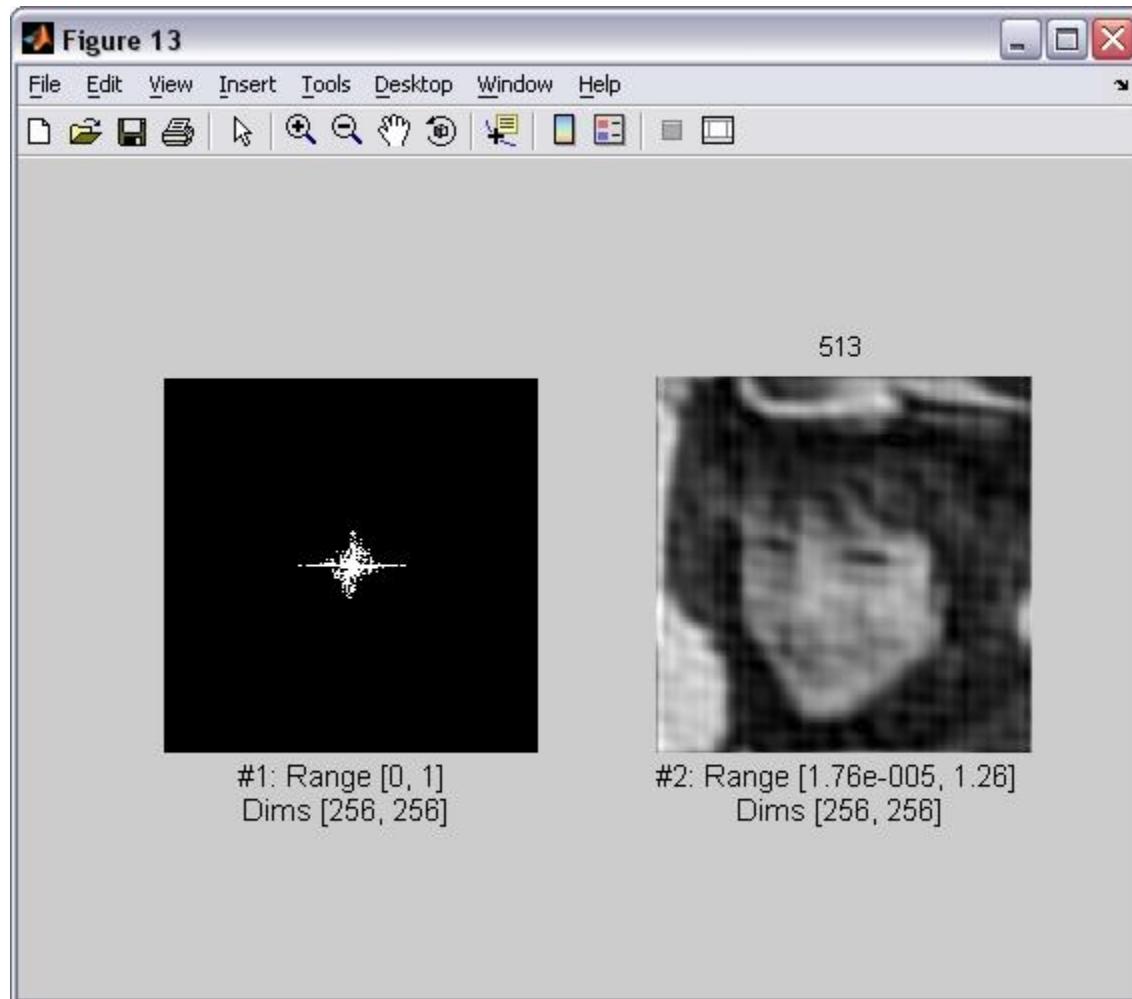
129



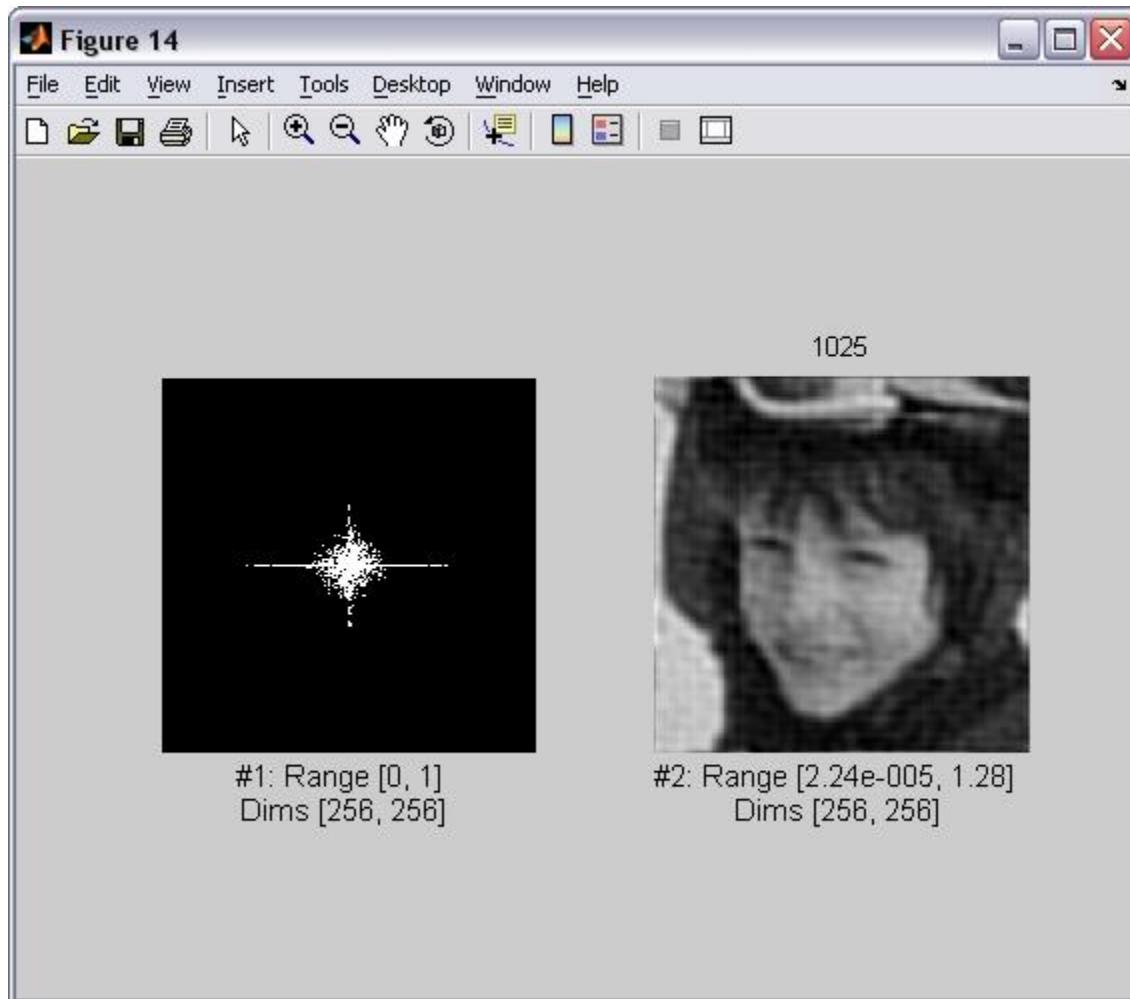
257



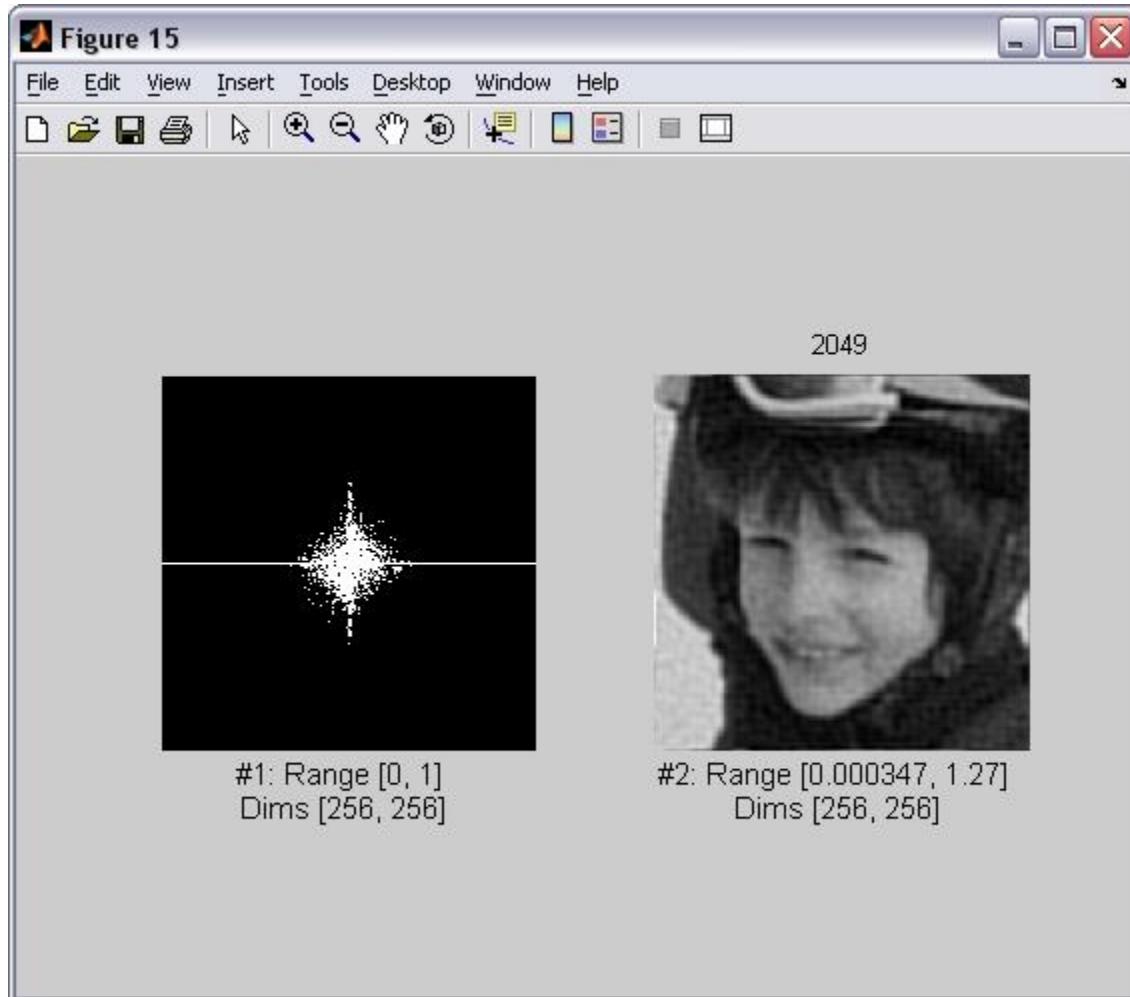
513



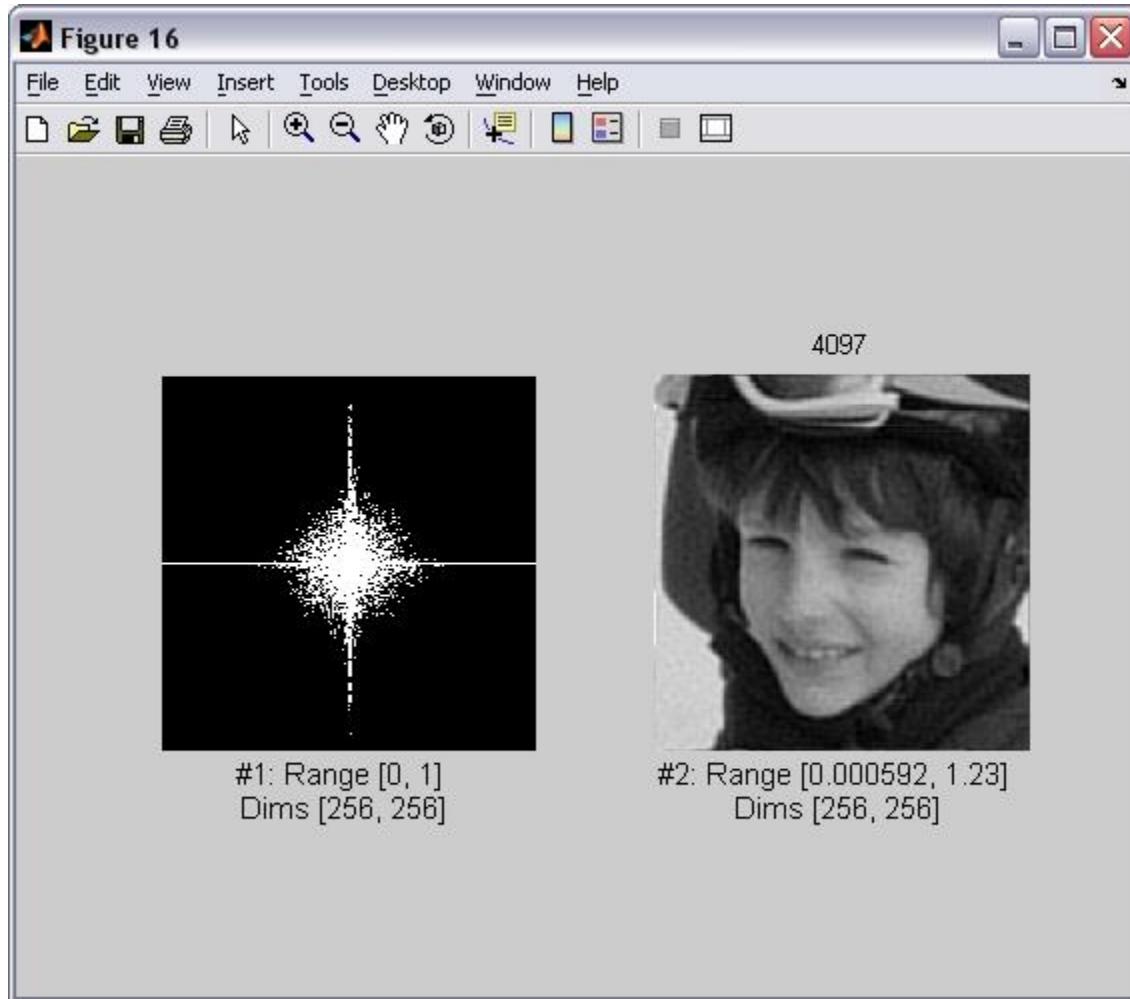
1025



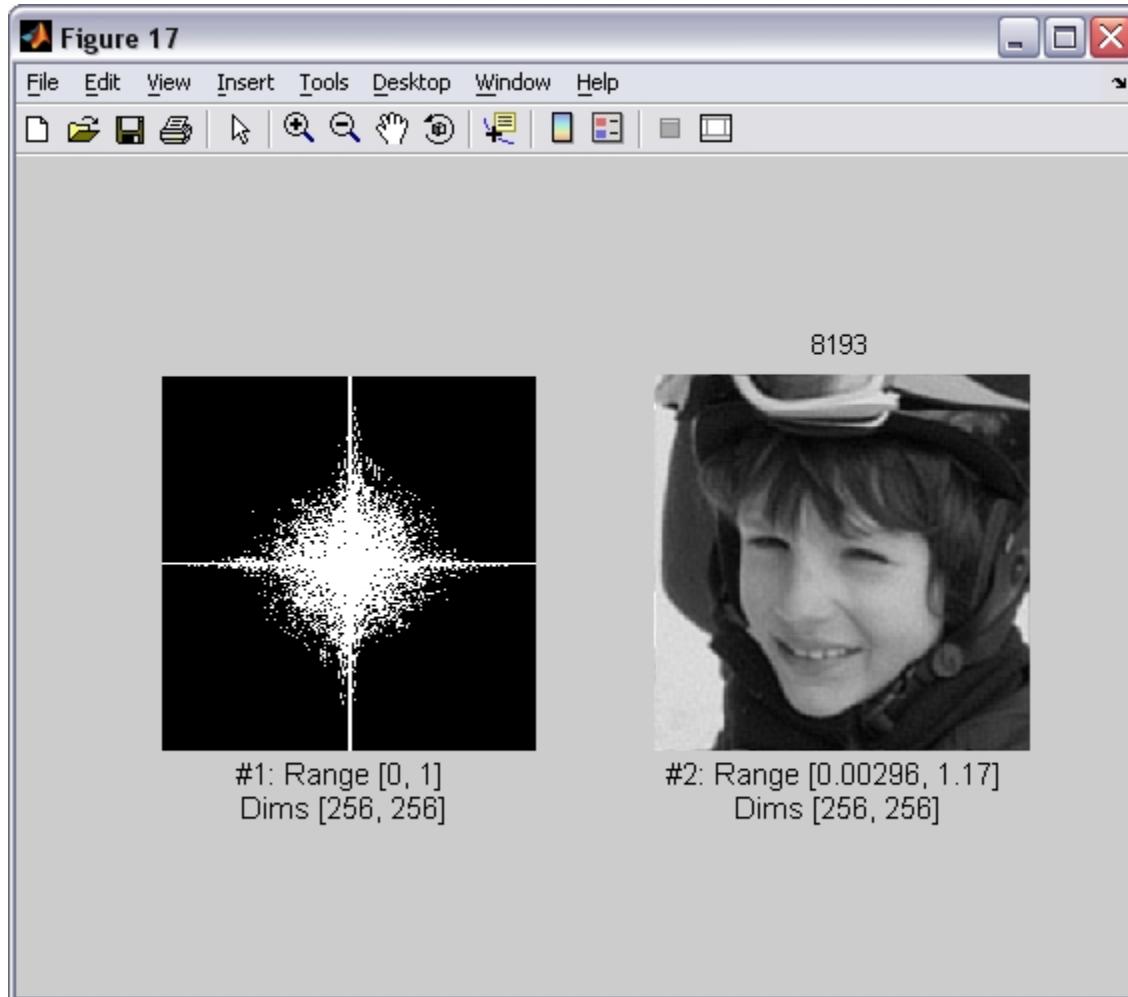
2049



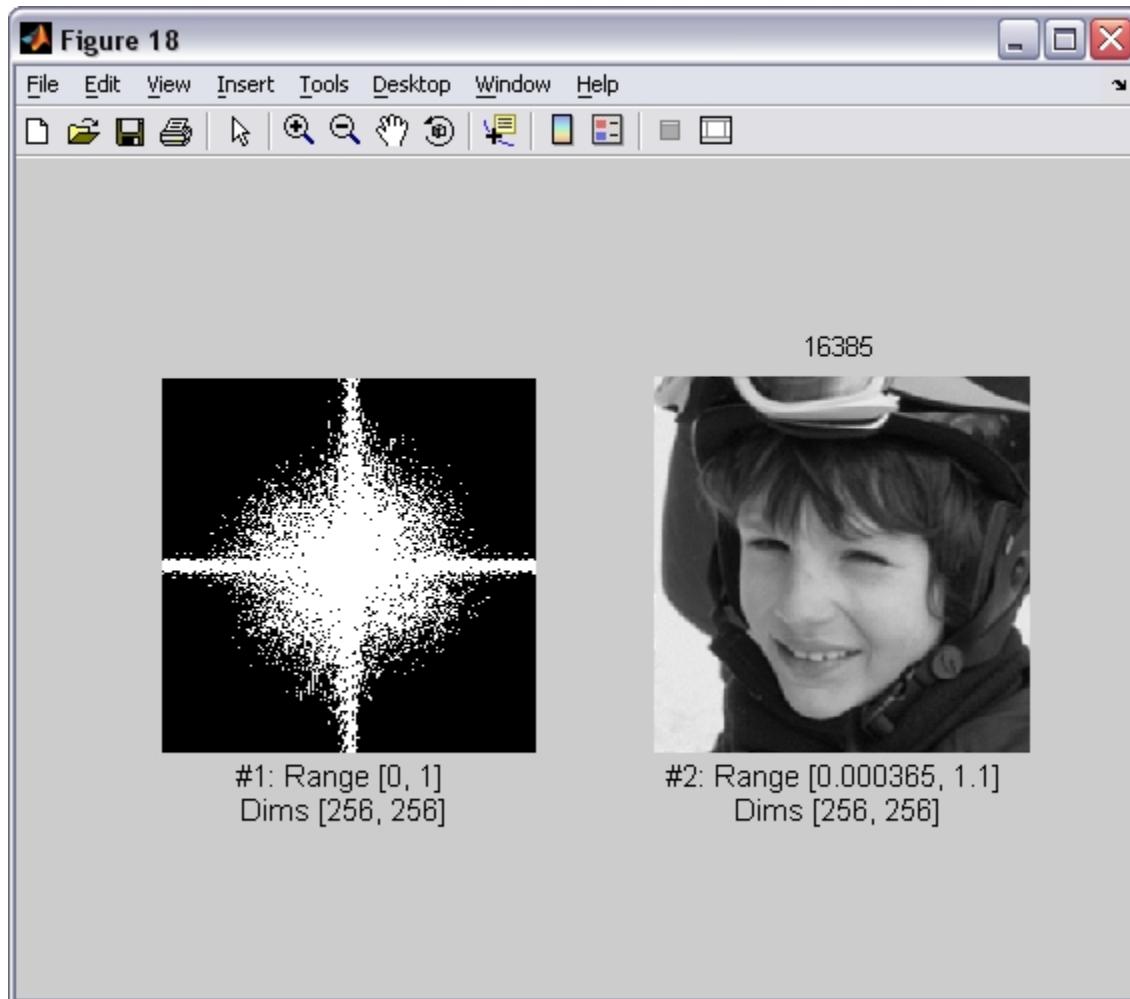
4097



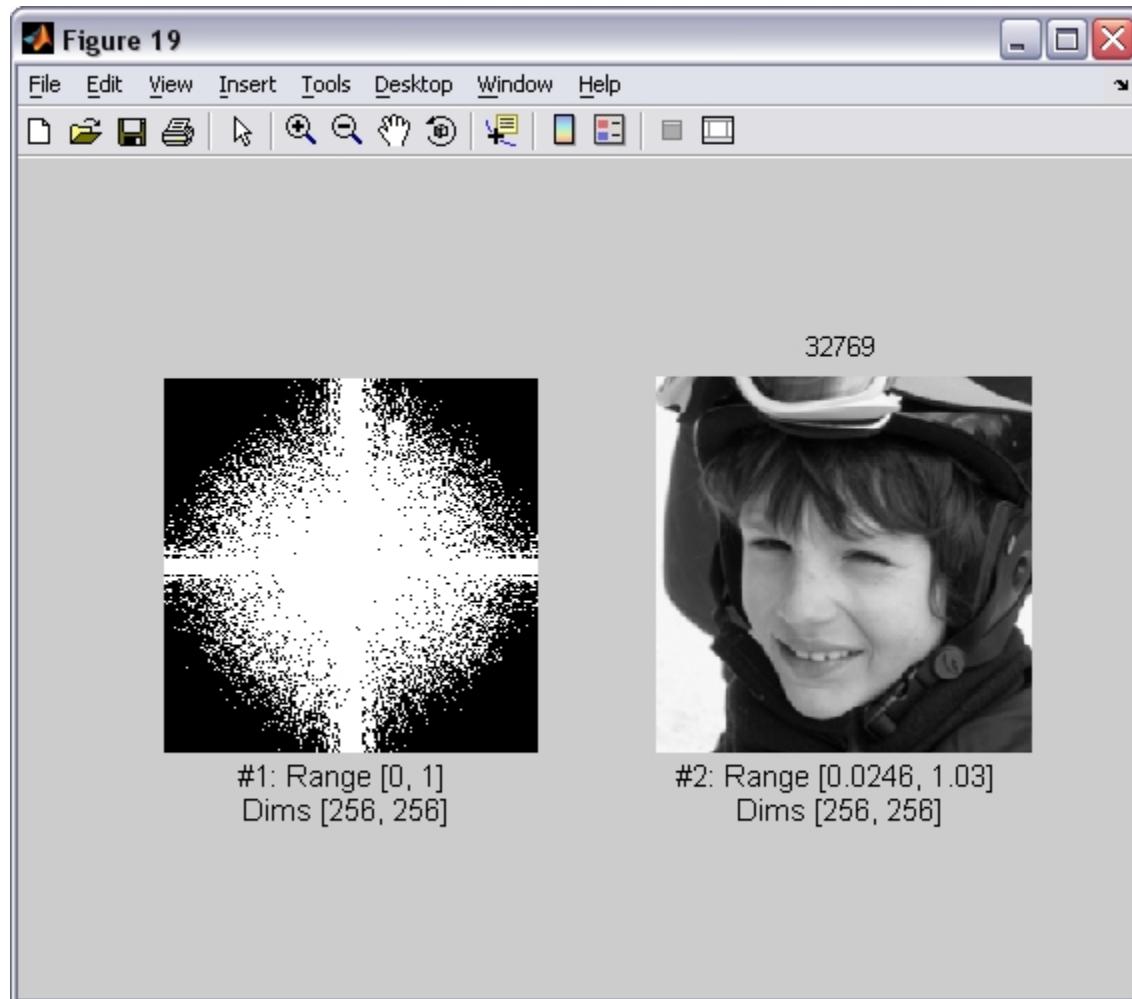
8193



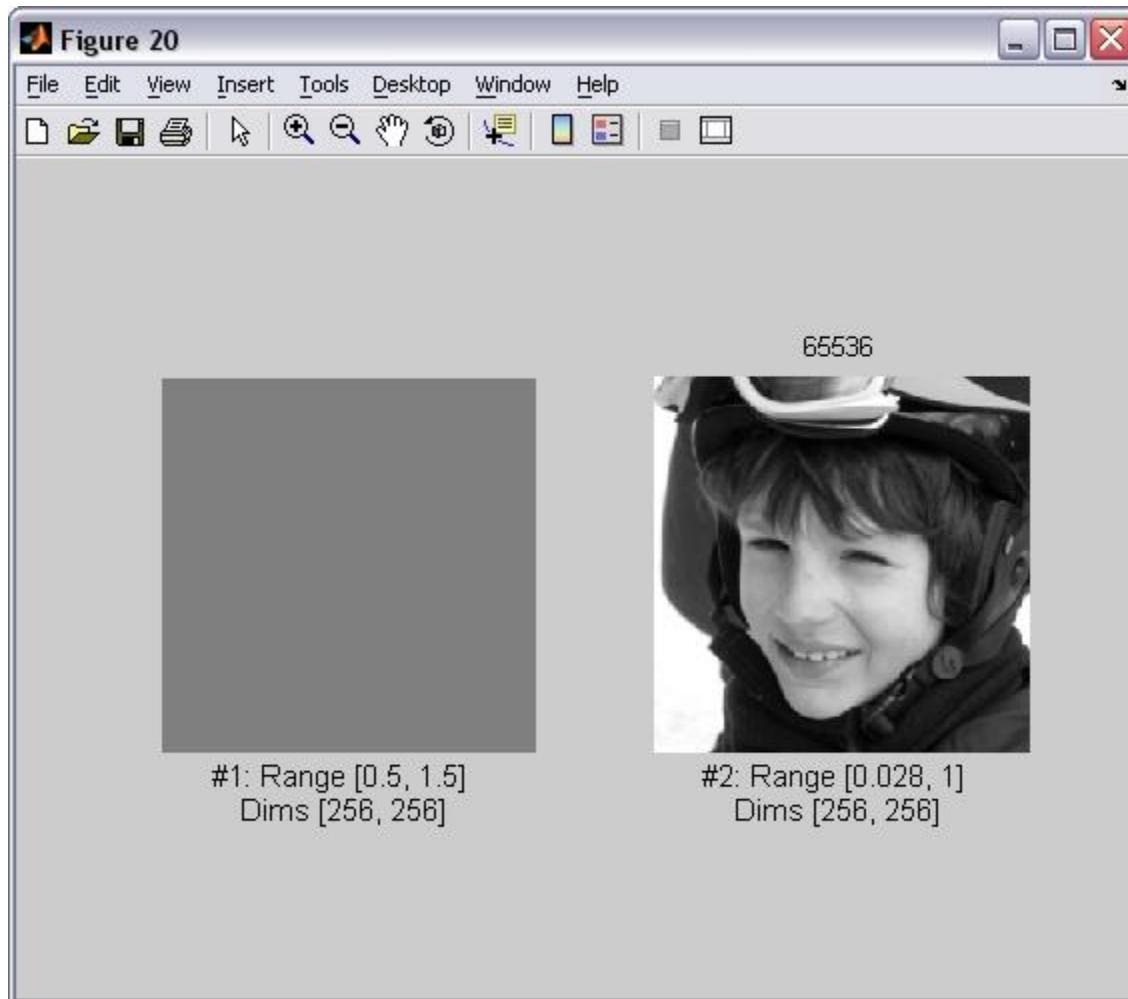
16385



32769

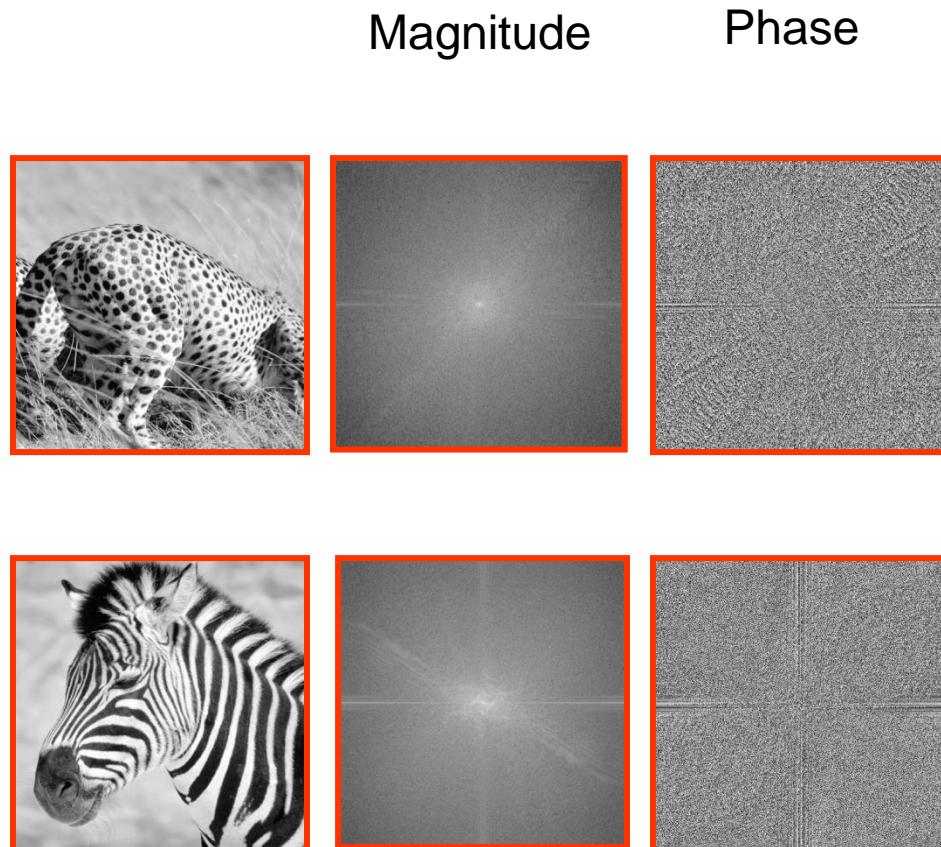


65536



Fourier Transform

- Fourier transform of a real function is complex
 - difficult to plot, visualize
 - instead, we can think of the phase and magnitude of the transform
- Phase is the phase of the complex transform
- Magnitude is the magnitude of the complex transform

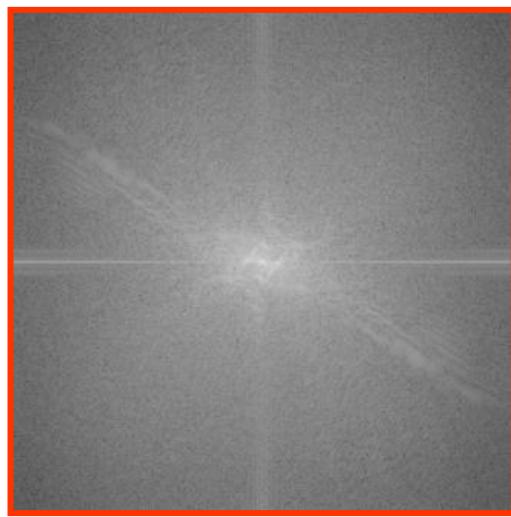
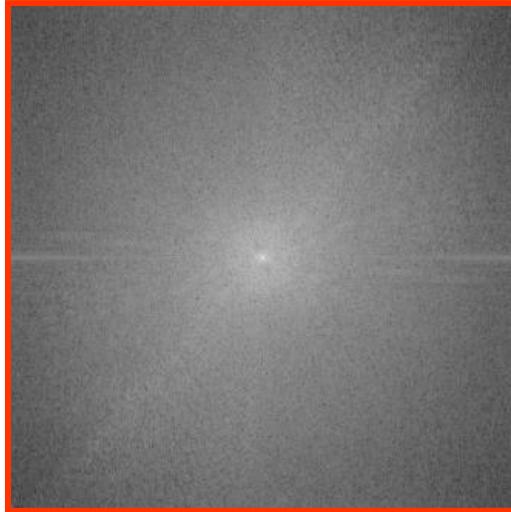


Source: D. Forsyth

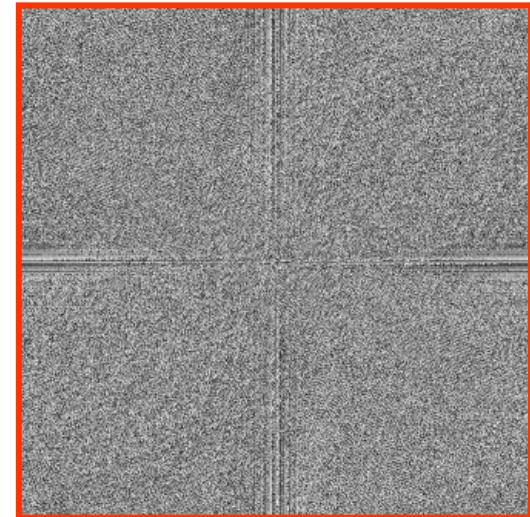
Fourier Transform



Magnitude

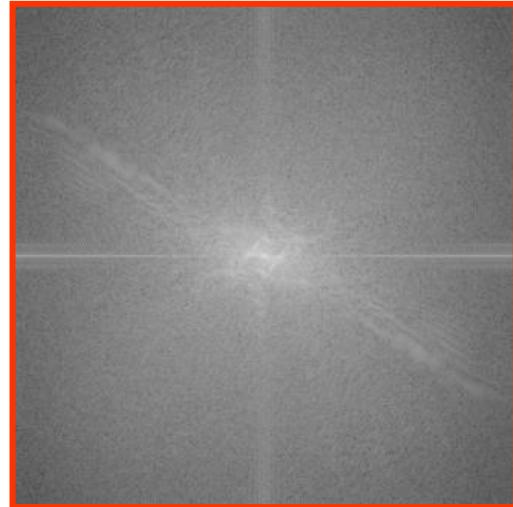
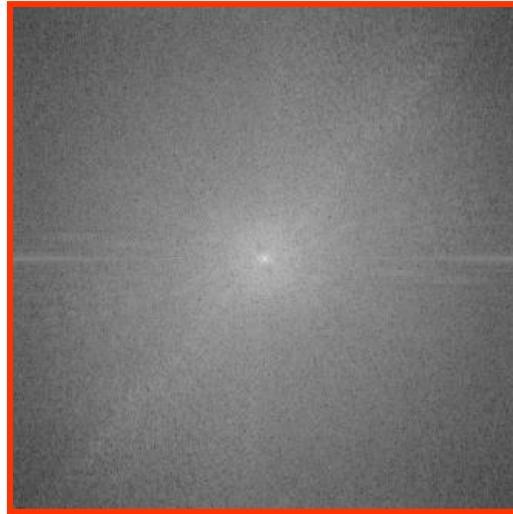


Phase



Fourier Transform

Magnitude



Phase

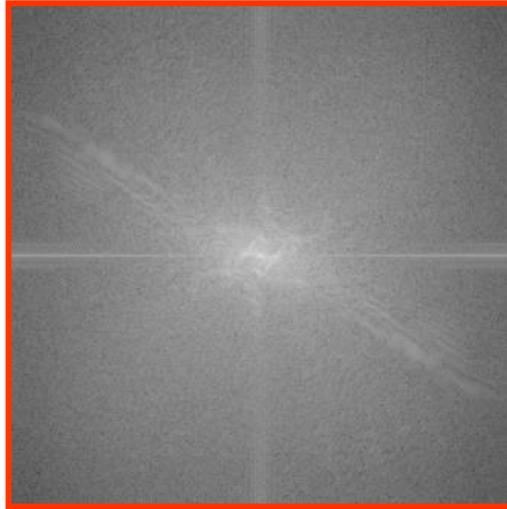


Fourier Transform

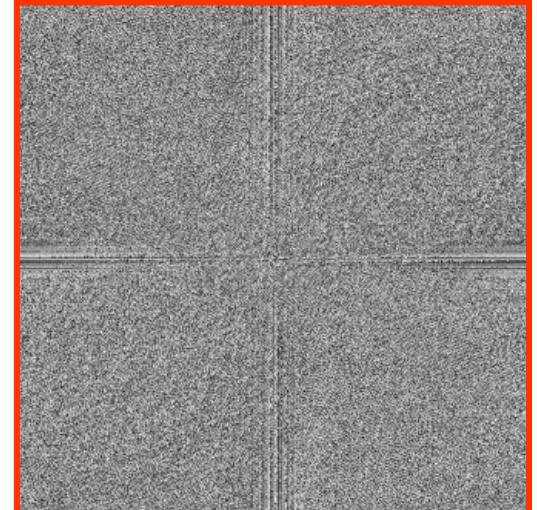
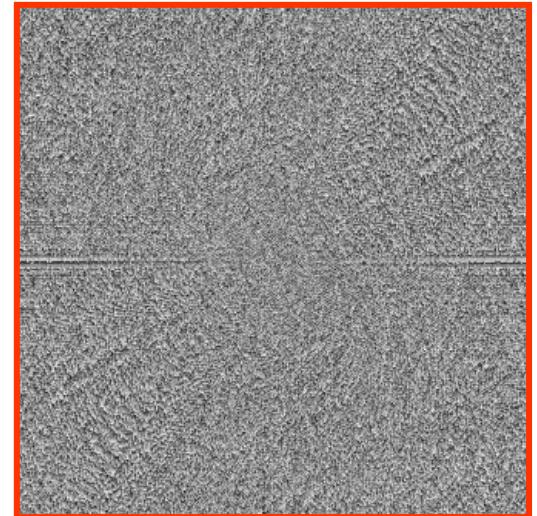
?

?

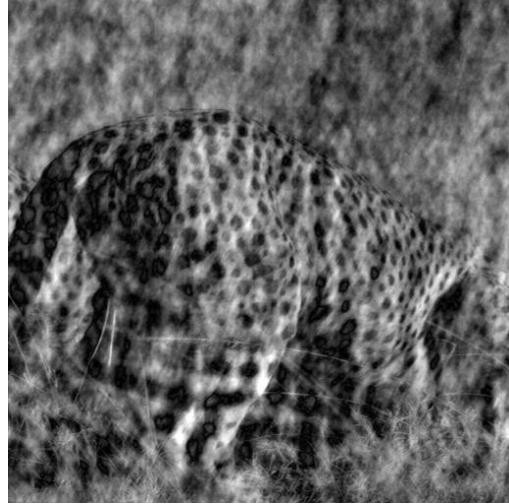
Magnitude



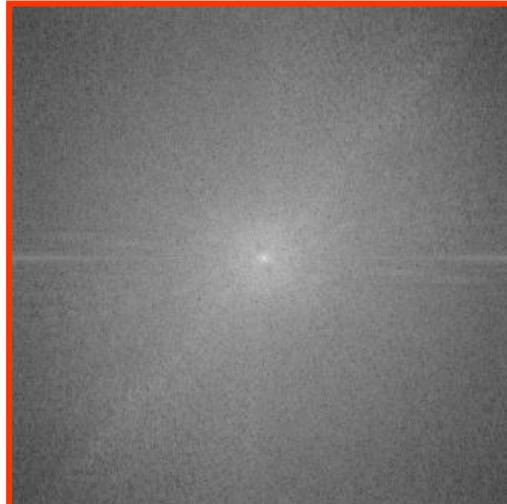
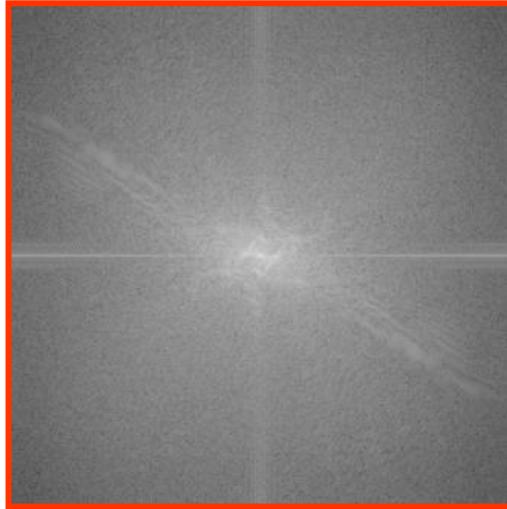
Phase



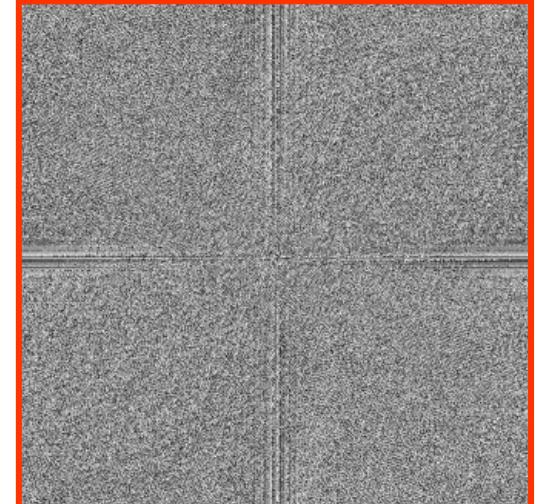
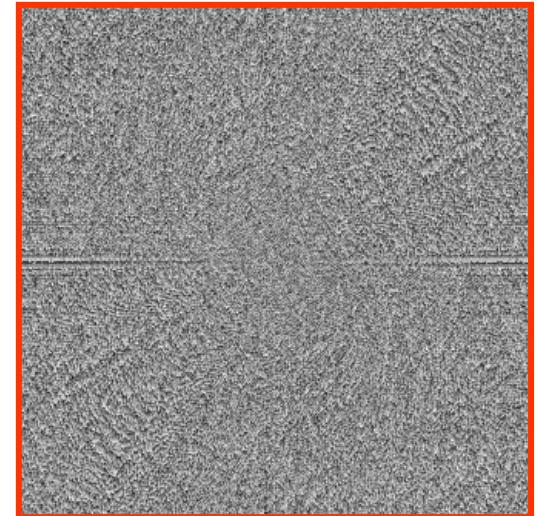
Fourier Transform



Magnitude



Phase



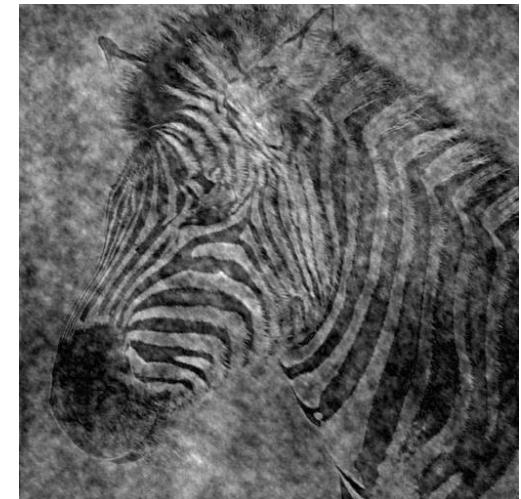
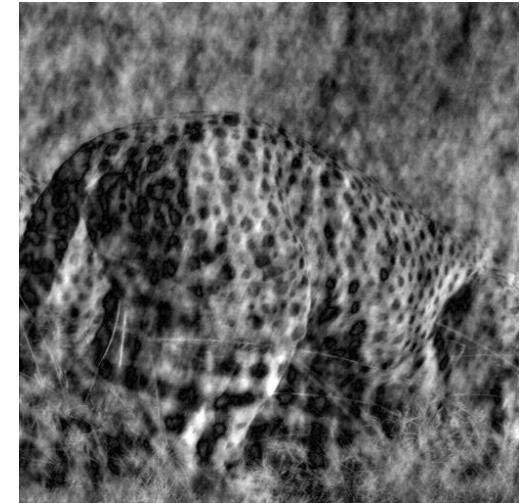
Phase and Magnitude



Image with cheetah phase
(and zebra magnitude)



Image with zebra phase
(and cheetah magnitude)



Source: D. Forsyth

The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

$$\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$

- The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms

$$f[gh] = f[g]*f[h]$$

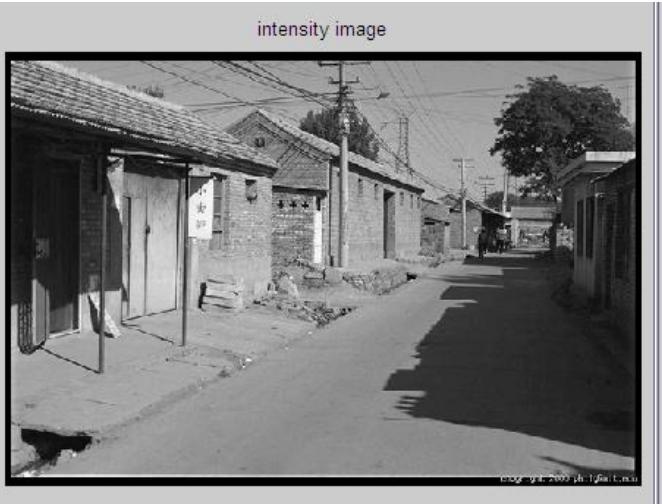
- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!

Properties of Fourier Transforms

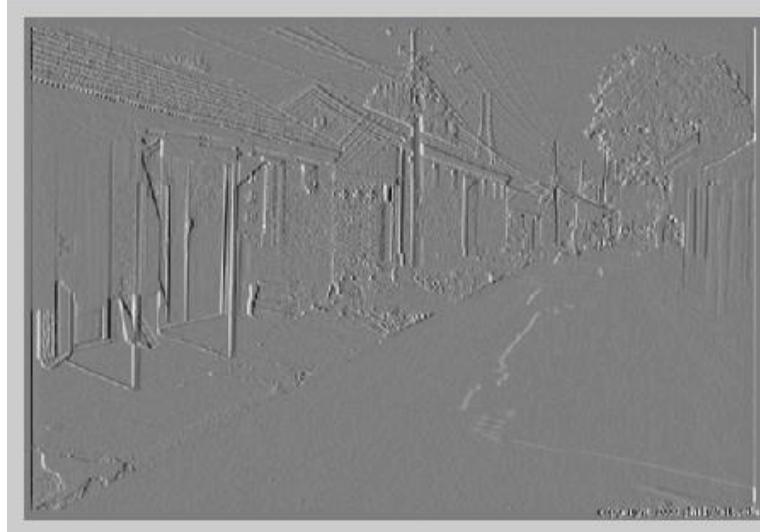
Property	Signal	Transform
superposition	$f_1(x) + f_2(x)$	$F_1(\omega) + F_2(\omega)$
shift	$f(x - x_0)$	$F(\omega)e^{-j\omega x_0}$
reversal	$f(-x)$	$F^*(\omega)$
convolution	$f(x) * h(x)$	$F(\omega)H(\omega)$
correlation	$f(x) \otimes h(x)$	$F(\omega)H^*(\omega)$
multiplication	$f(x)h(x)$	$F(\omega) * H(\omega)$
differentiation	$f'(x)$	$j\omega F(\omega)$
domain scaling	$f(ax)$	$1/aF(\omega/a)$
real images	$f(x) = f^*(x)$	$\Leftrightarrow F(\omega) = F(-\omega)$
Parseval's Theorem	$\sum_x [f(x)]^2$	$= \sum_\omega [F(\omega)]^2$

Filtering in spatial domain

1	0	-1
2	0	-2
1	0	-1

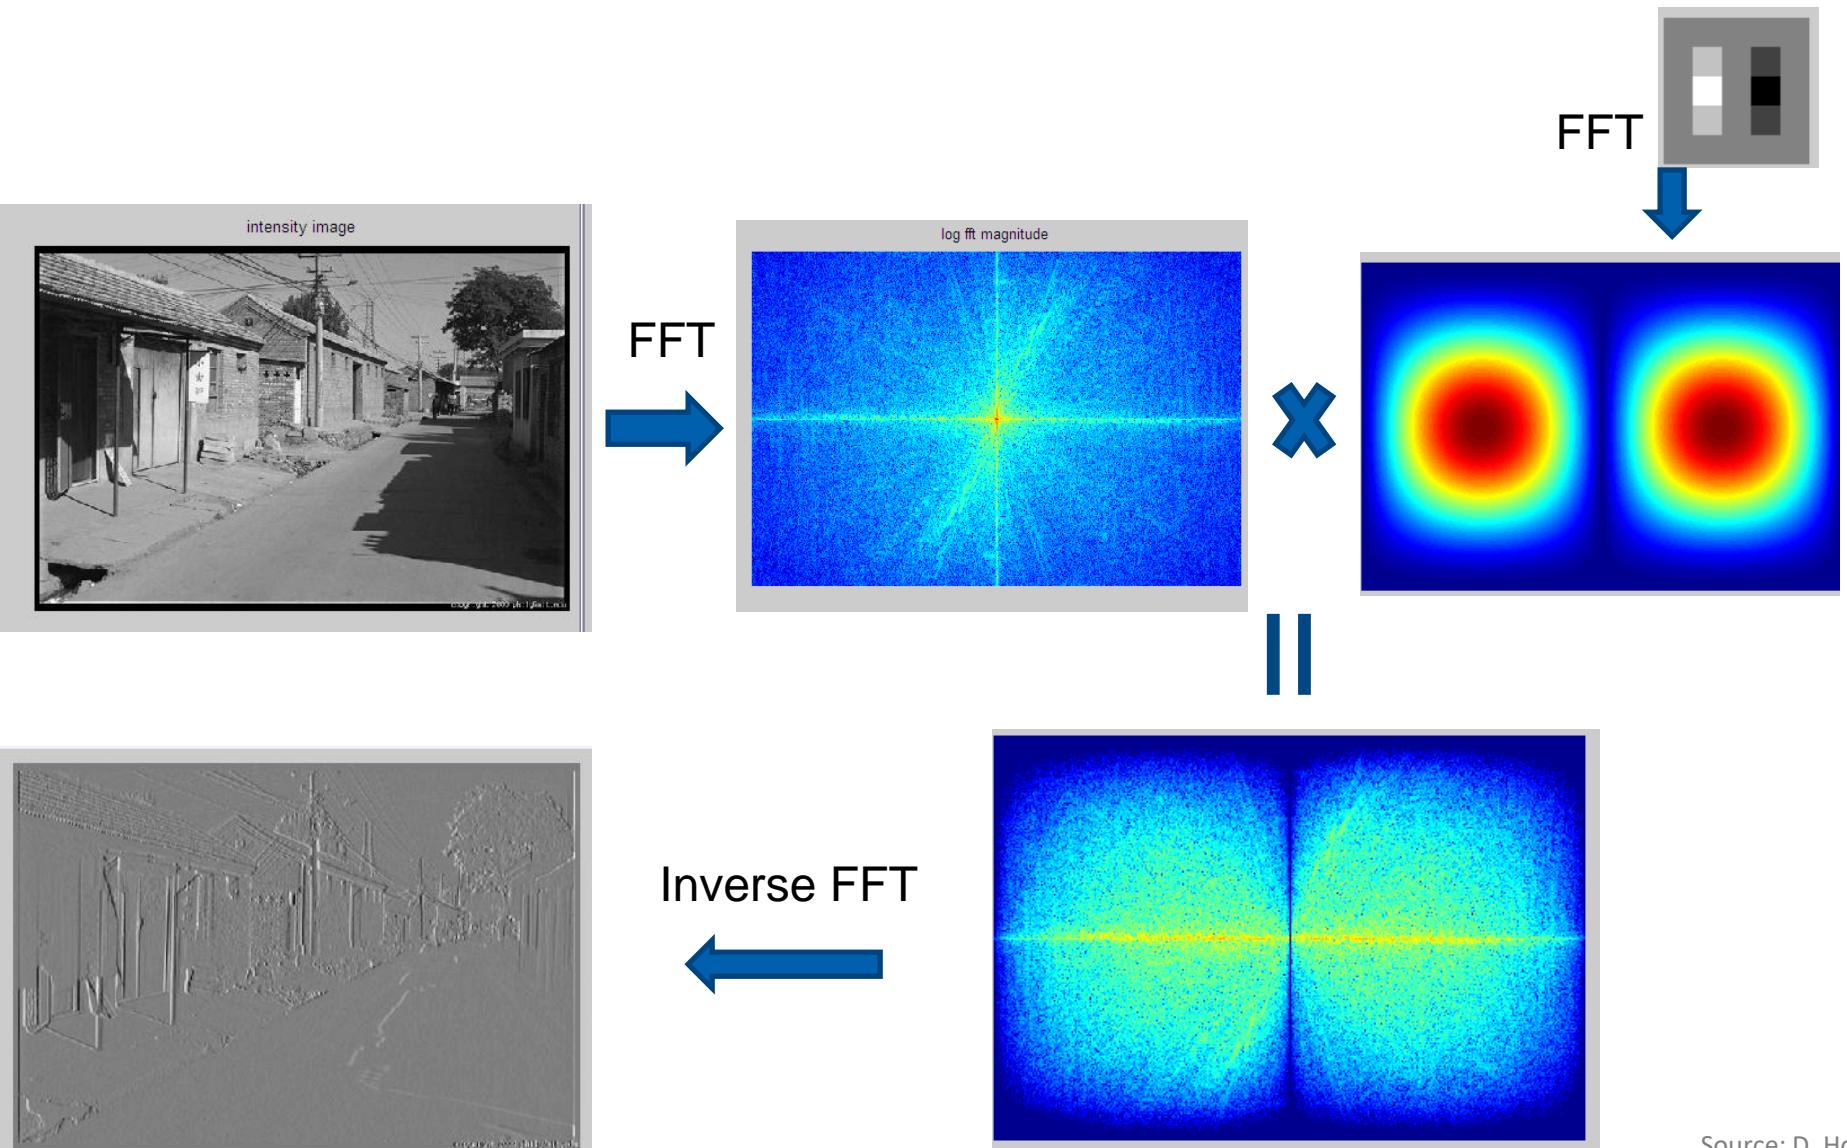


$$\begin{matrix} * & \end{matrix} = \begin{matrix} \text{blurred image} \end{matrix}$$



Source: J. Hays

Filtering in frequency domain

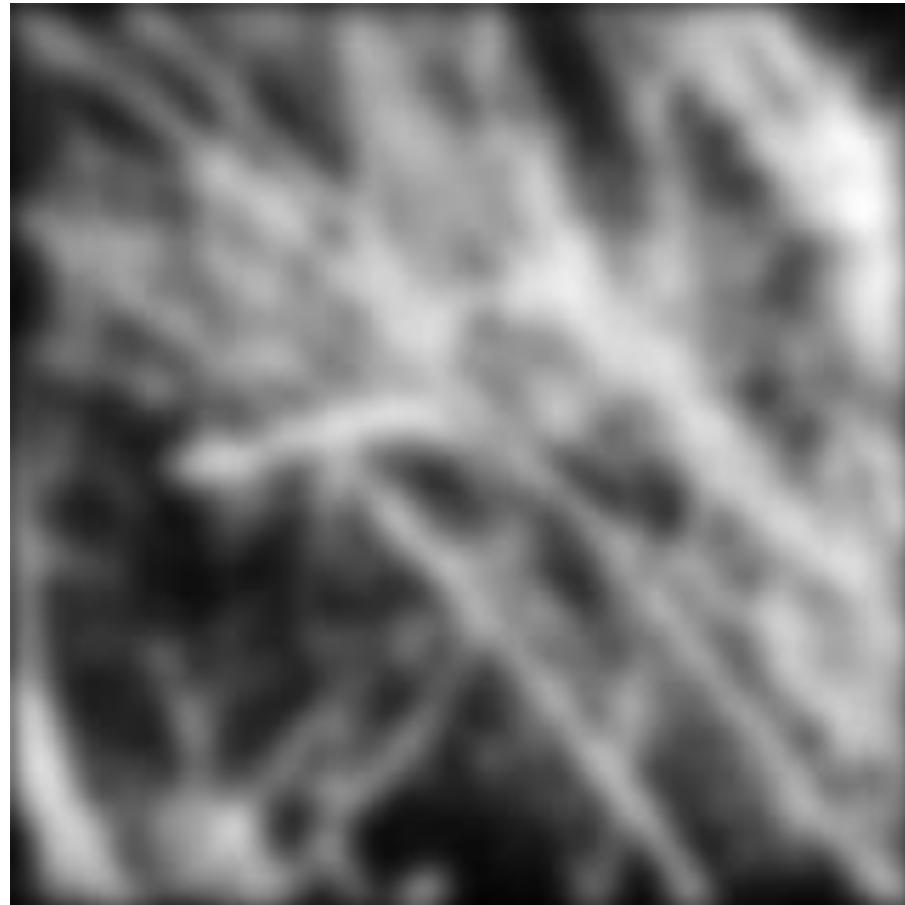


Source: D. Hoiem

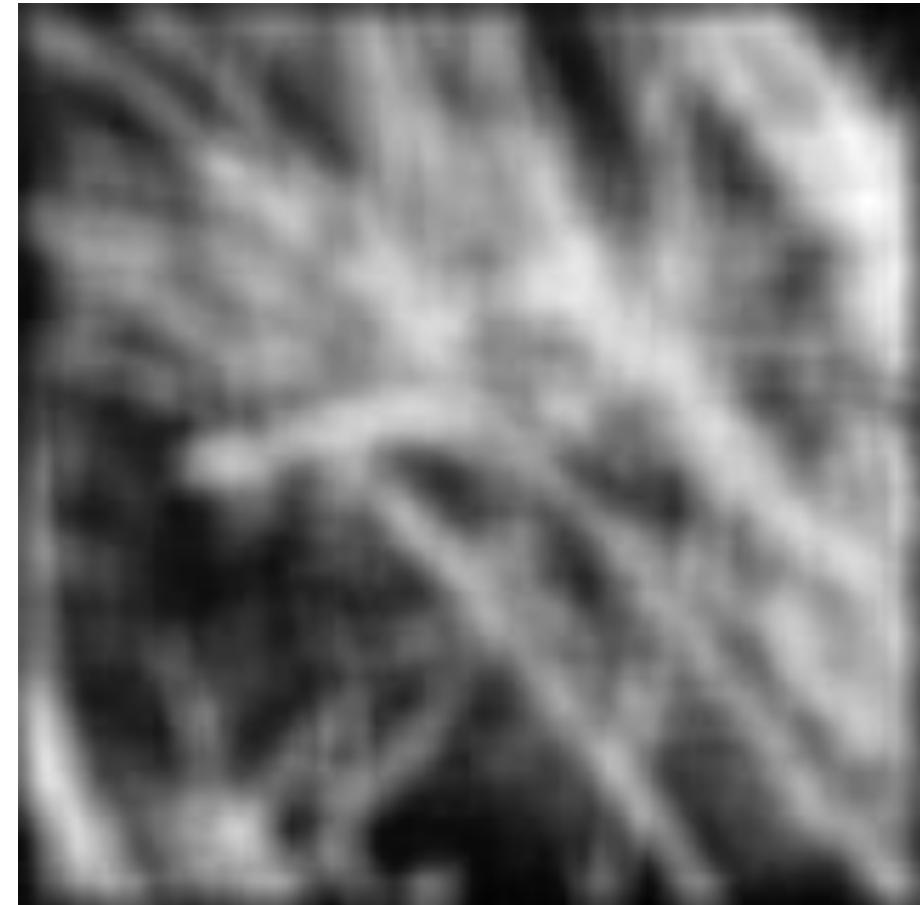
Linear filtering revisited

Gauss vs. box filter

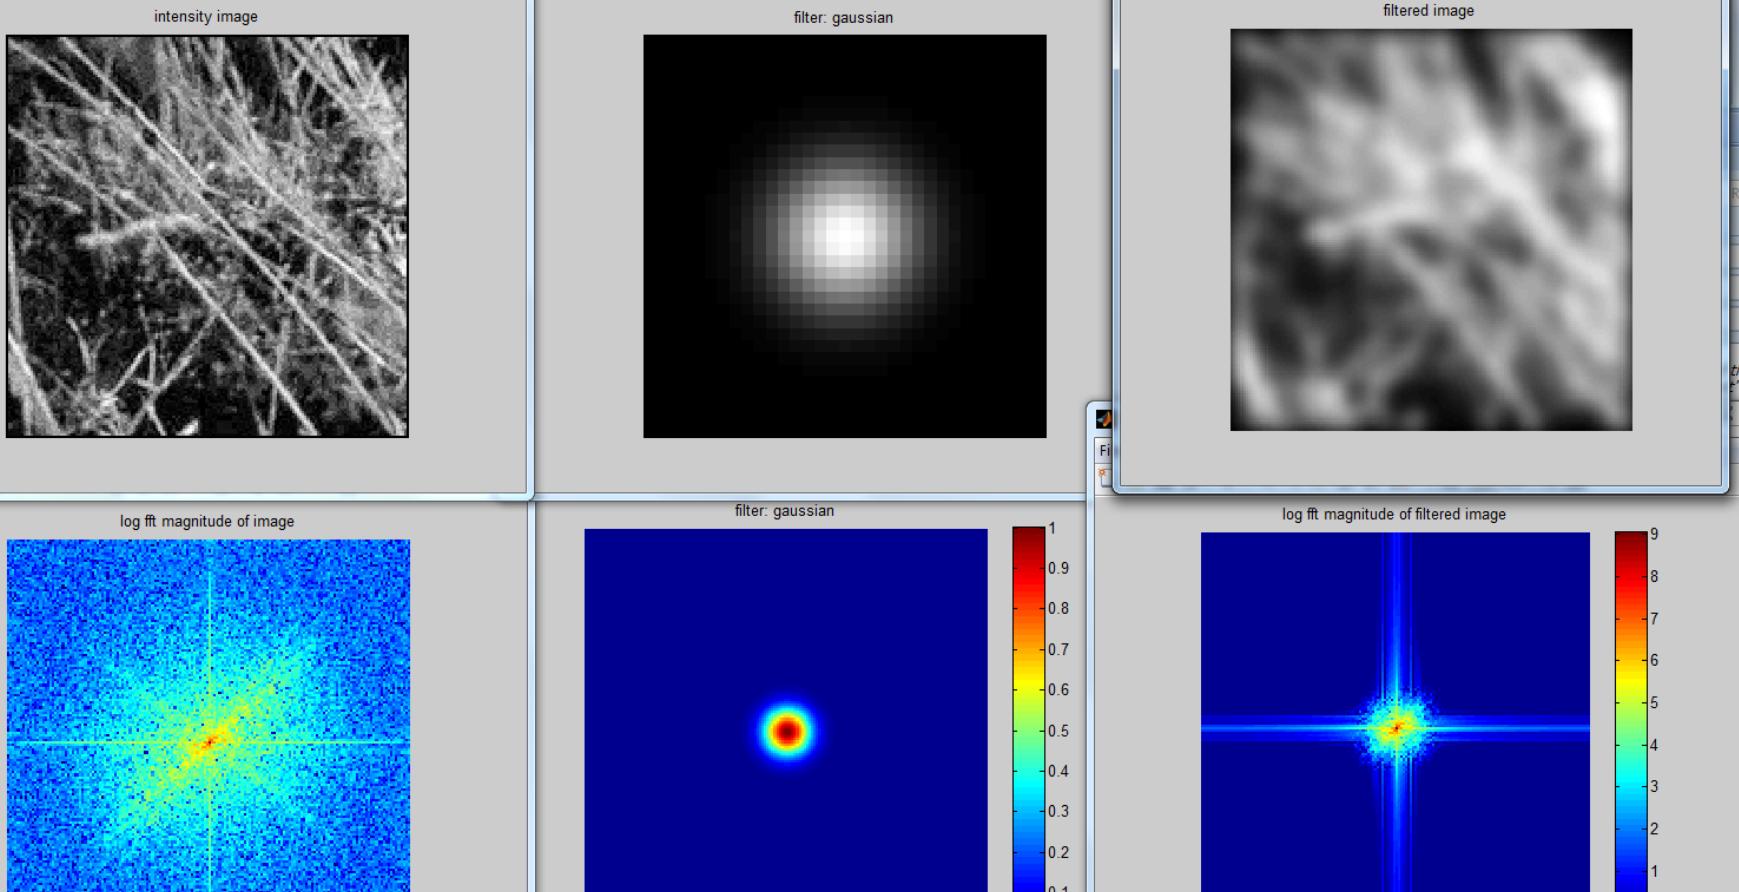
Gaussian



Box filter

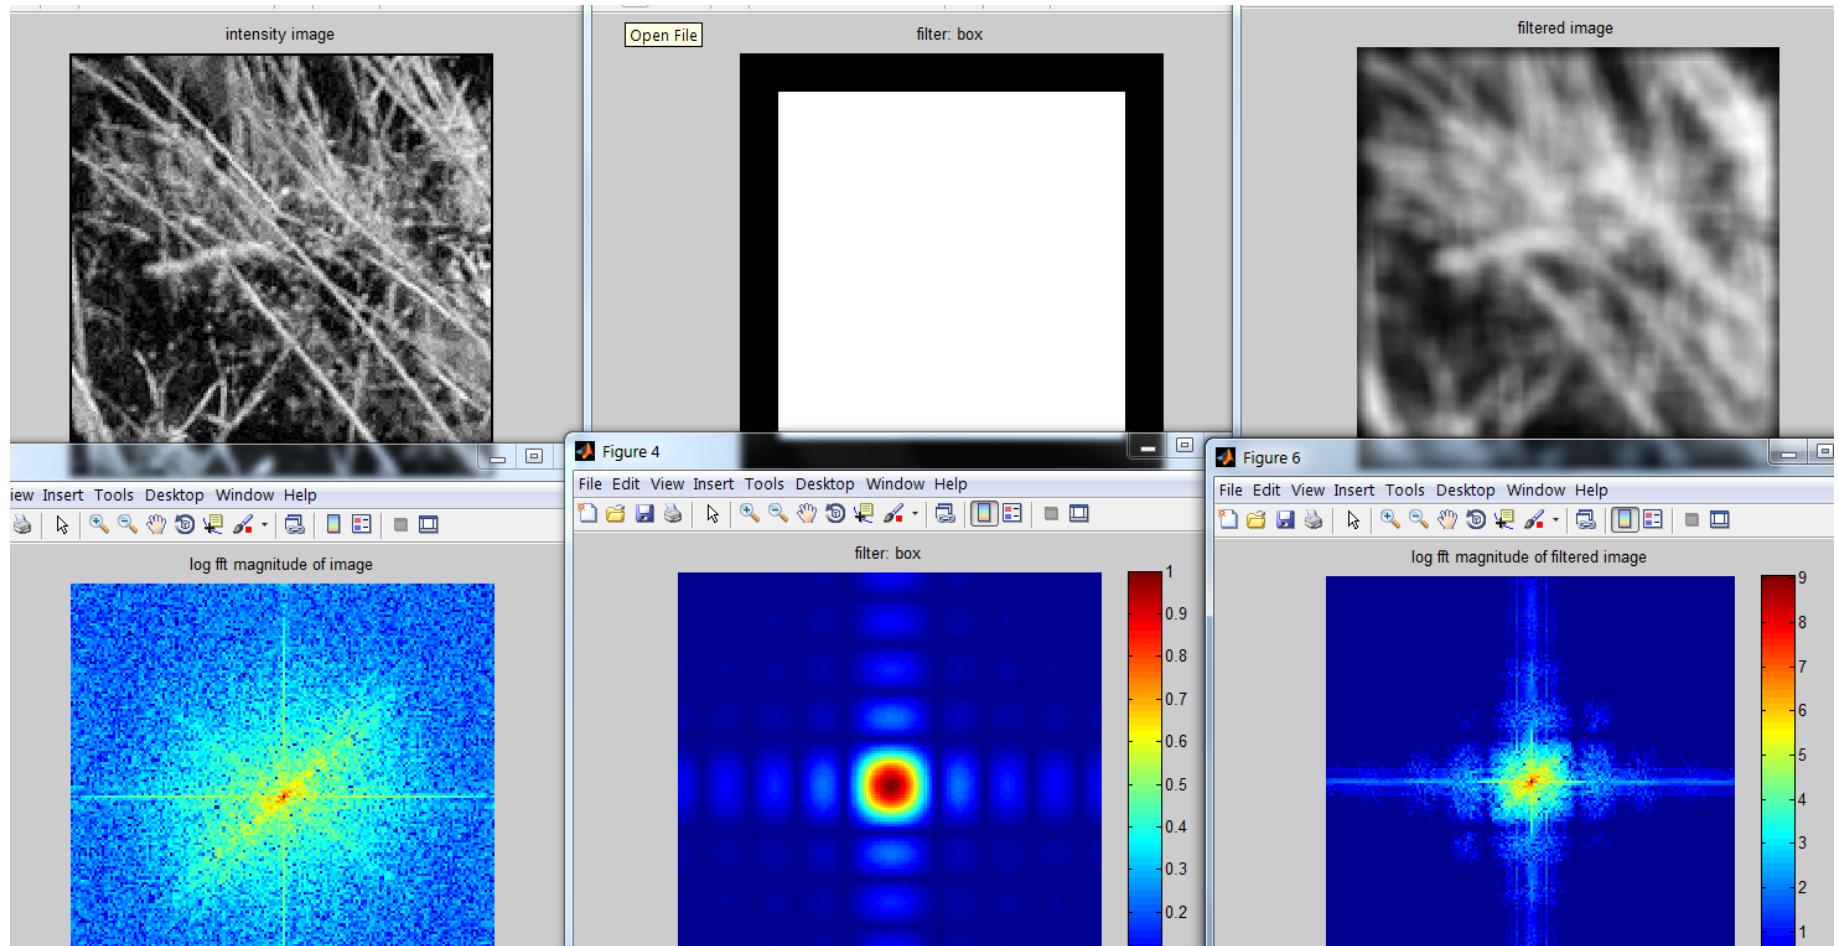


Gaussian



Source: J. Hays

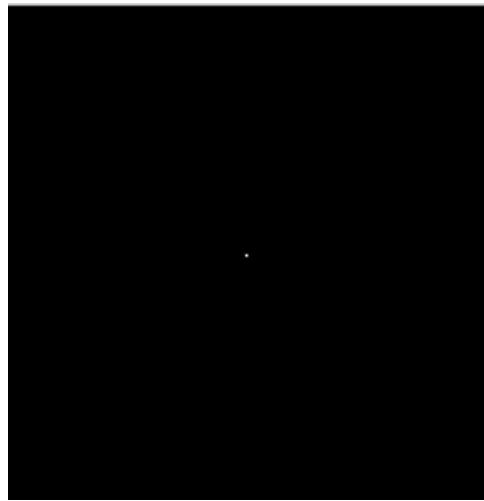
Box filter



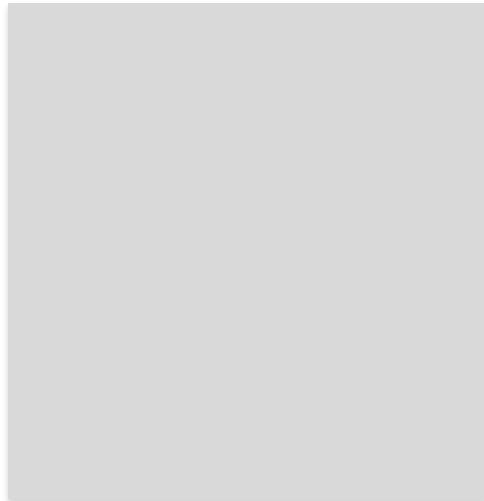
Source: J. Hays

Some important Fourier Transforms

Image



Magnitude FT



Source: A. Torralba

Continuous Fourier transform

1D Continuous Fourier Transform:

$$f(x) = \int_{u=-\infty}^{\infty} F(u)e^{+2\pi jux}du$$

The inverse Fourier transform

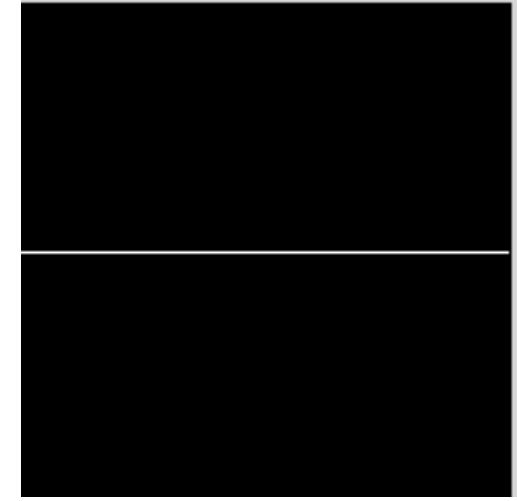
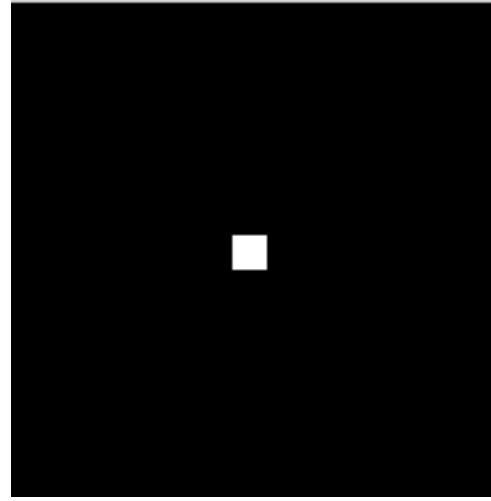
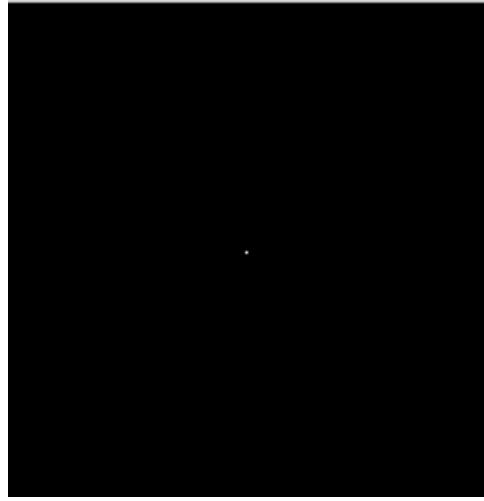
$$F(u) = \int_{x=-\infty}^{\infty} f(x)e^{-2\pi jux}dx$$

The Fourier transform

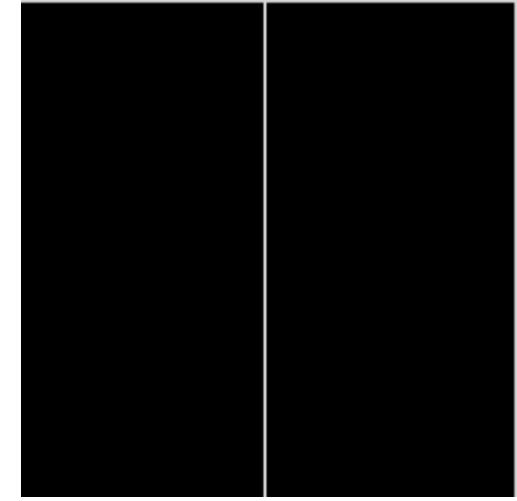
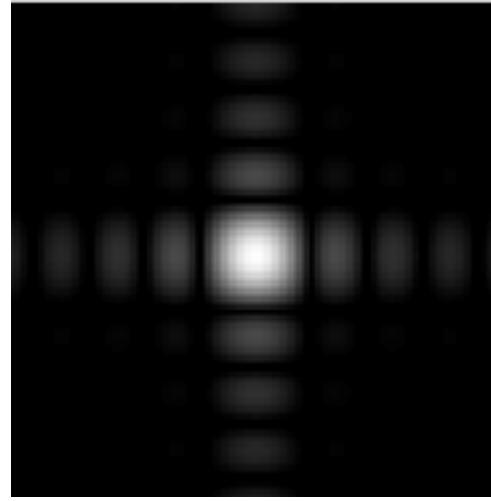
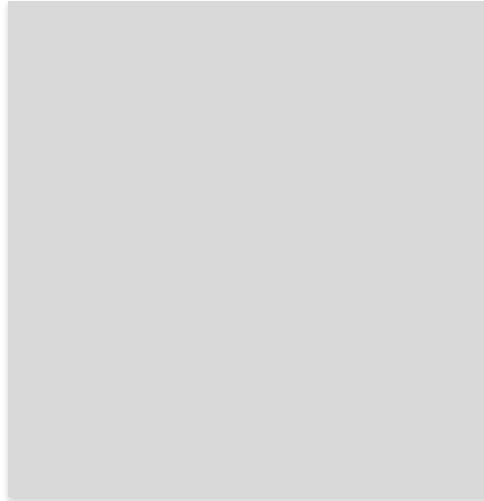
$$g(x) = \int \delta(x - u)h(u)du = h(x)$$

Some important Fourier Transforms

Image



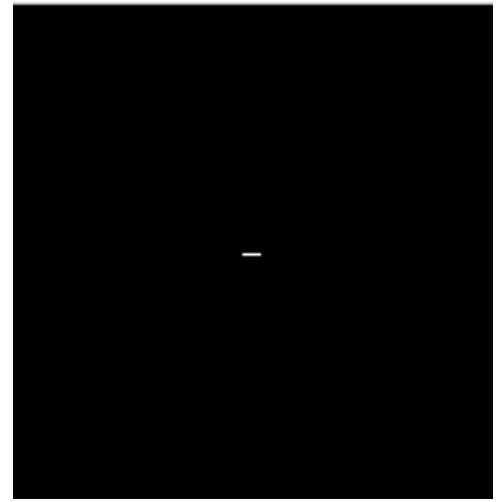
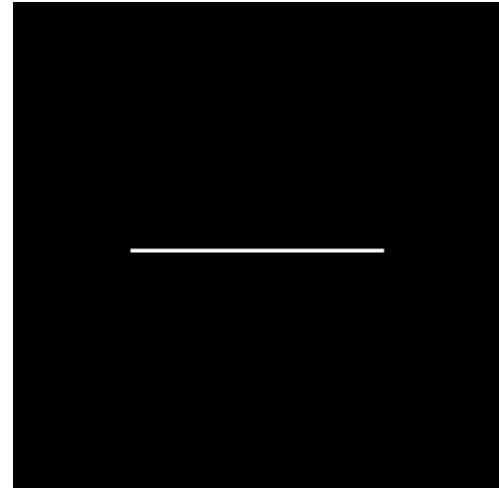
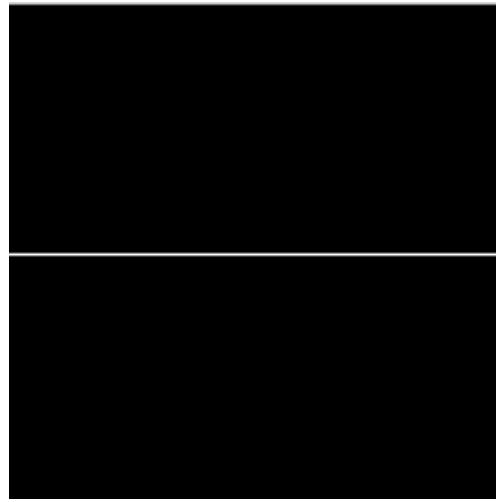
Magnitude FT



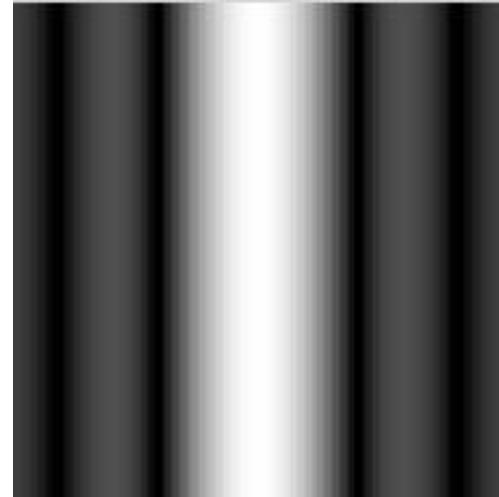
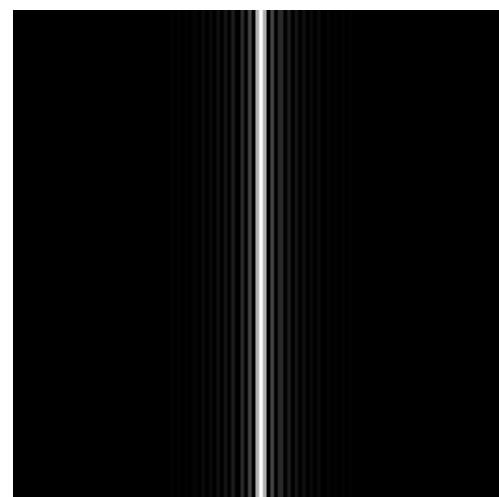
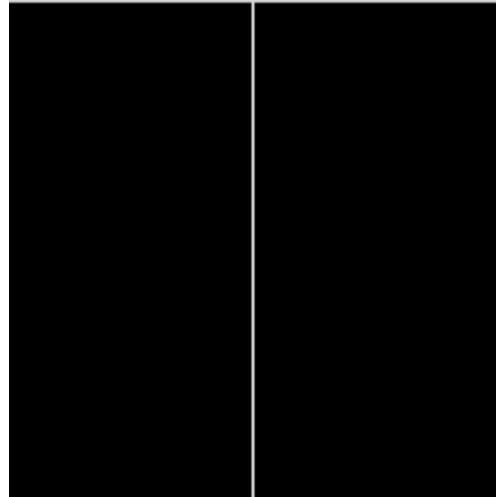
Source: A. Torralba

Some important Fourier Transforms

Image



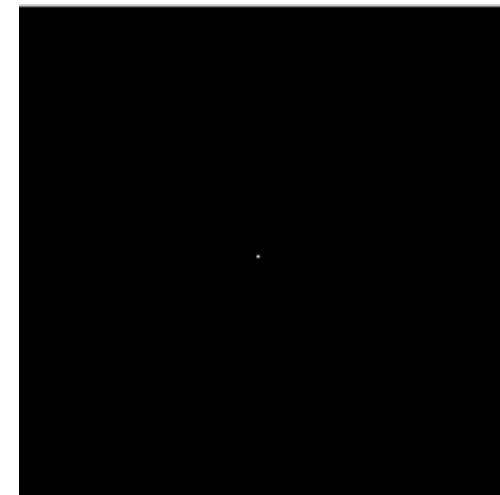
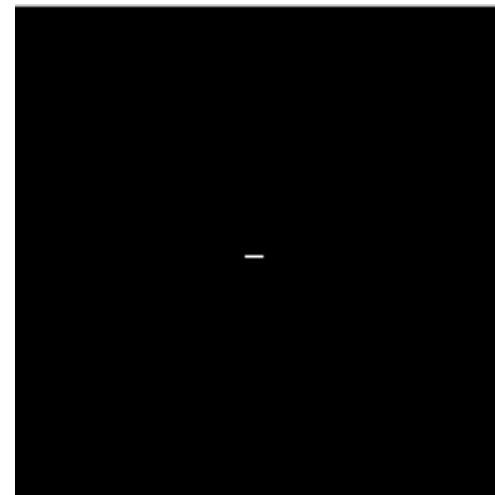
Magnitude FT



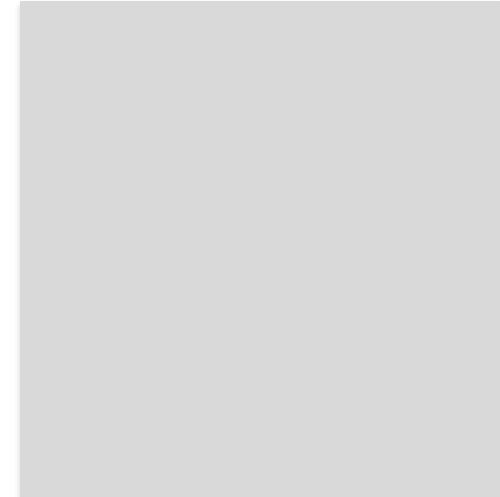
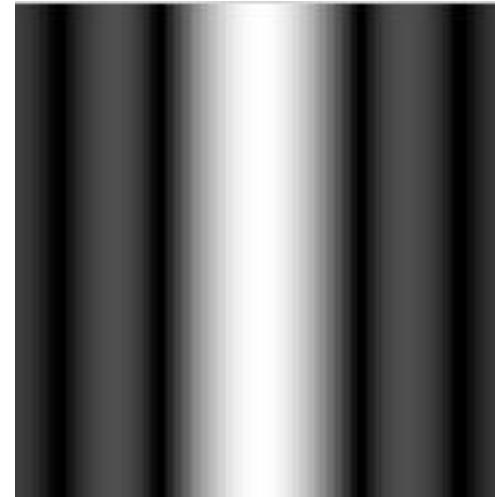
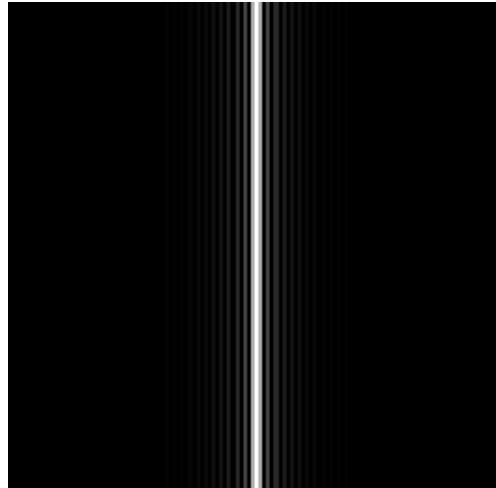
Source: A. Torralba

Some important Fourier Transforms

Image

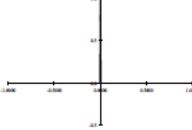
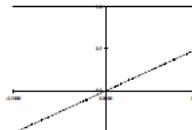
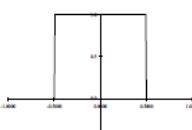
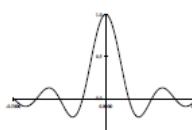
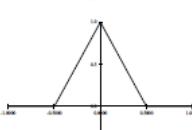
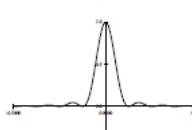
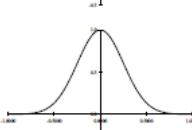
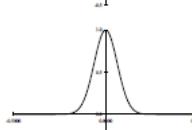
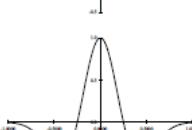
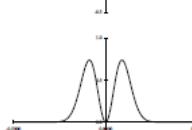
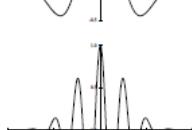
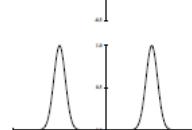


Magnitude FT



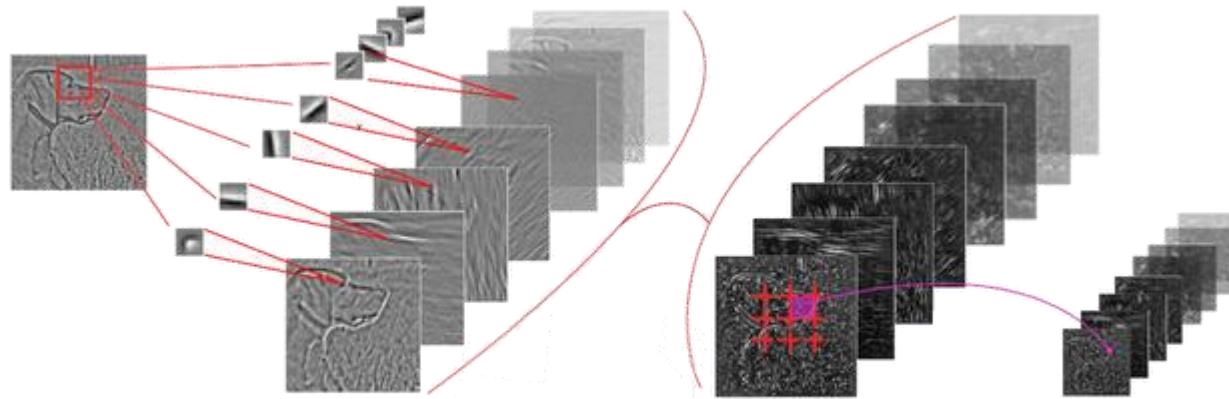
Source: A. Torralba

Pairs - Functions

Name	Signal	Transform
impulse		$\delta(x) \Leftrightarrow 1$ 
shifted impulse		$\delta(x - u) \Leftrightarrow e^{-j\omega u}$ 
box filter		$\text{box}(x/a) \Leftrightarrow a \text{sinc}(a\omega)$ 
tent		$\text{tent}(x/a) \Leftrightarrow a \text{sinc}^2(a\omega)$ 
Gaussian		$G(x; \sigma) \Leftrightarrow \frac{\sqrt{2\pi}}{\sigma} G(\omega; \sigma^{-1})$ 
Laplacian of Gaussian		$(\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2})G(x; \sigma) \Leftrightarrow -\frac{\sqrt{2\pi}}{\sigma} \omega^2 G(\omega; \sigma^{-1})$ 
Gabor		$\cos(\omega_0 x)G(x; \sigma) \Leftrightarrow \frac{\sqrt{2\pi}}{\sigma} G(\omega \pm \omega_0; \sigma^{-1})$ 

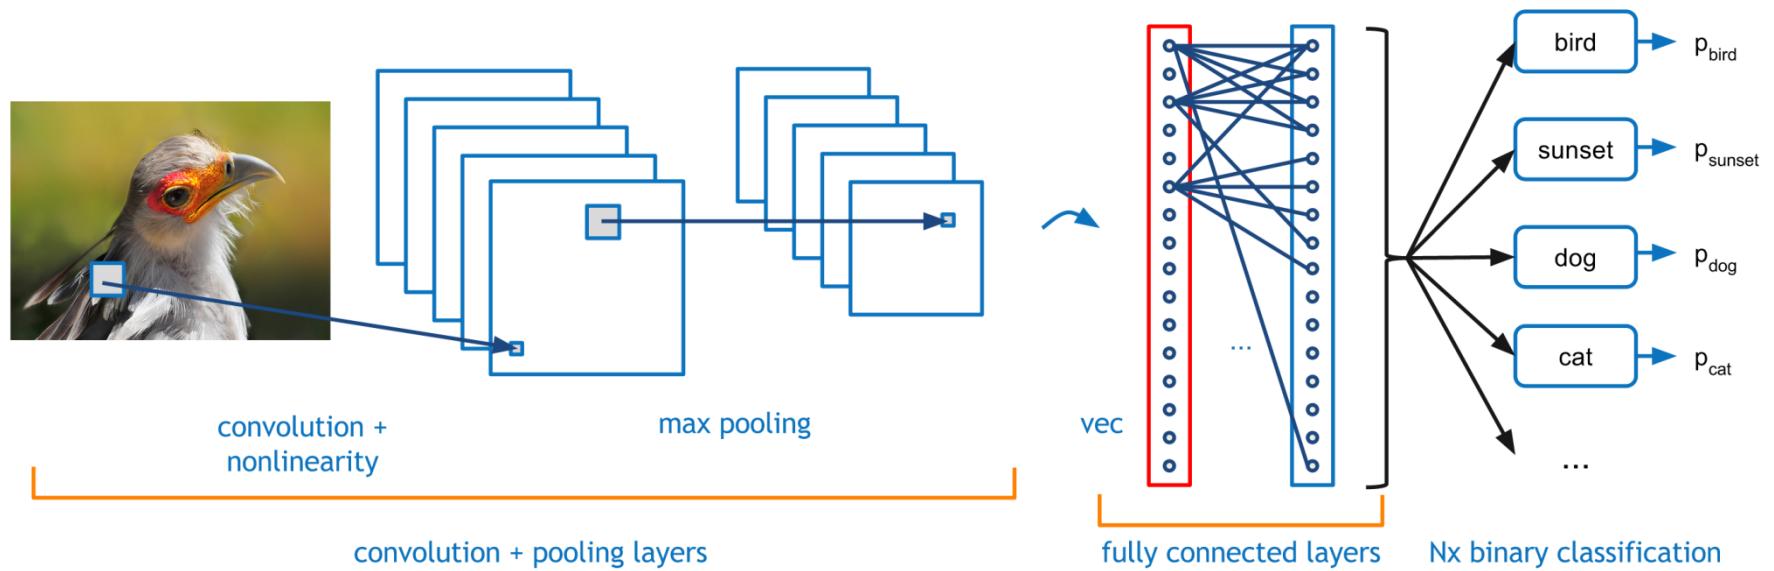
FFT and Convolutional Networks

- Convolutional networks learn convolution kernels from training data



FFT and Convolutional Networks

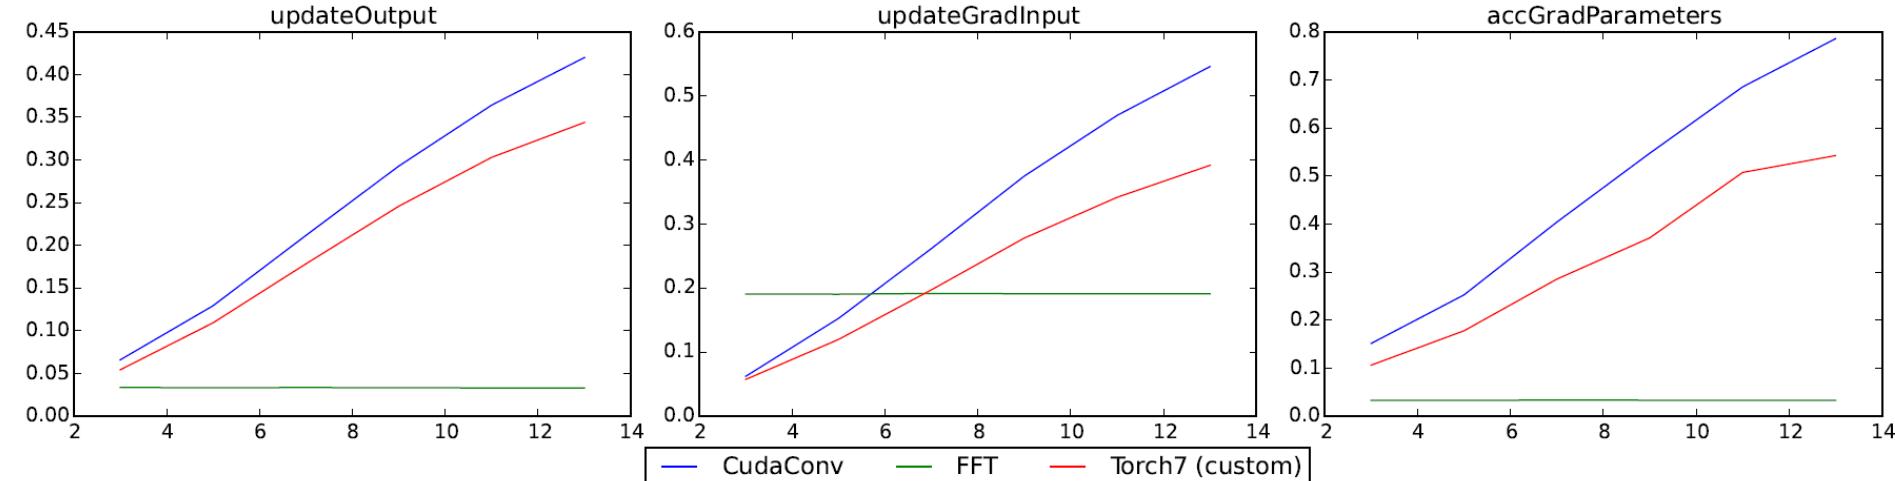
- Convolutional networks learn convolution kernels from training data
- Can be used for image classification



FFT and Convolutional Networks

- Convolutional networks are trained on GPUs
- Pixel-wise multiplications perfect for GPU
- NVIDIA CUDA Fast Fourier Transform library (cuFFT)
- Speed up training for larger kernels

Time (seconds) versus kernelSize
batchSize=128, inputSize=32, nInputPlanes=96, nOutputPlanes=256



[Mathieu et al. Fast Training of Convolutional Networks through FFTs, ICLR2014]

Image Scaling

This image is too big to fit on the screen. How can we reduce it?

How to generate a half-sized version?

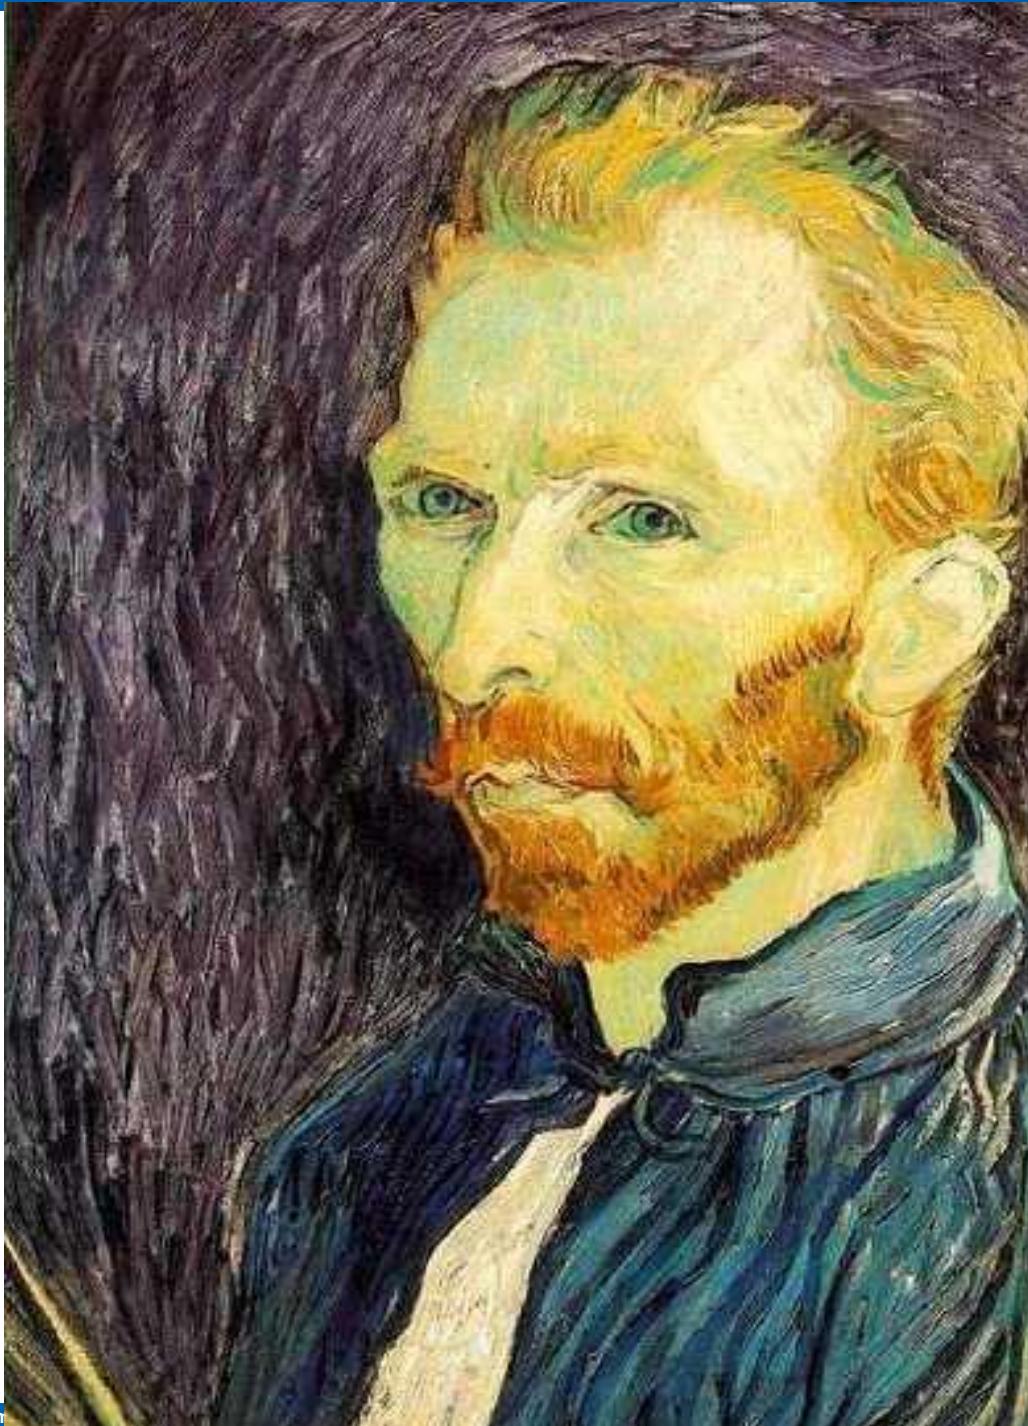
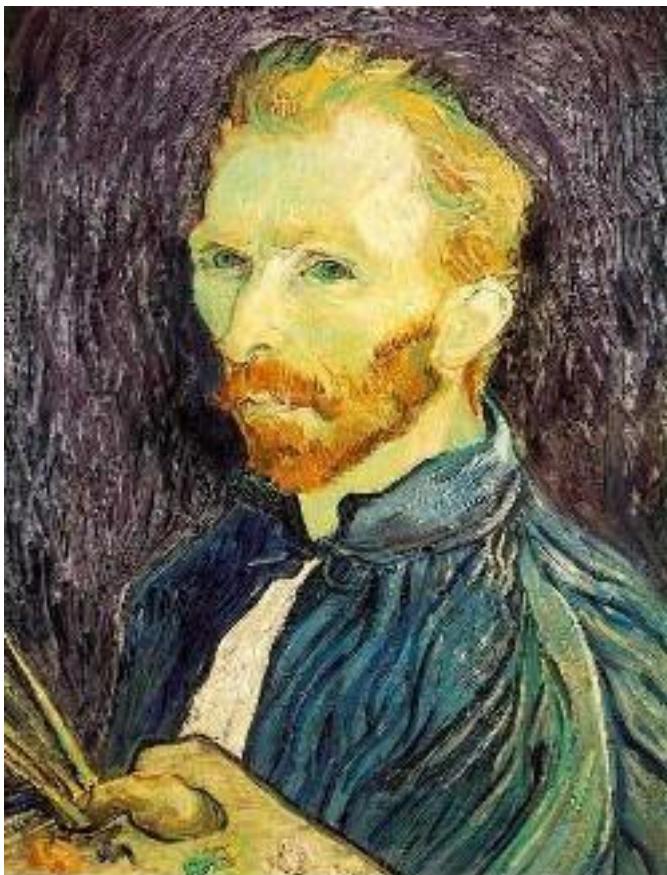


Image Sub-Sampling



1/4

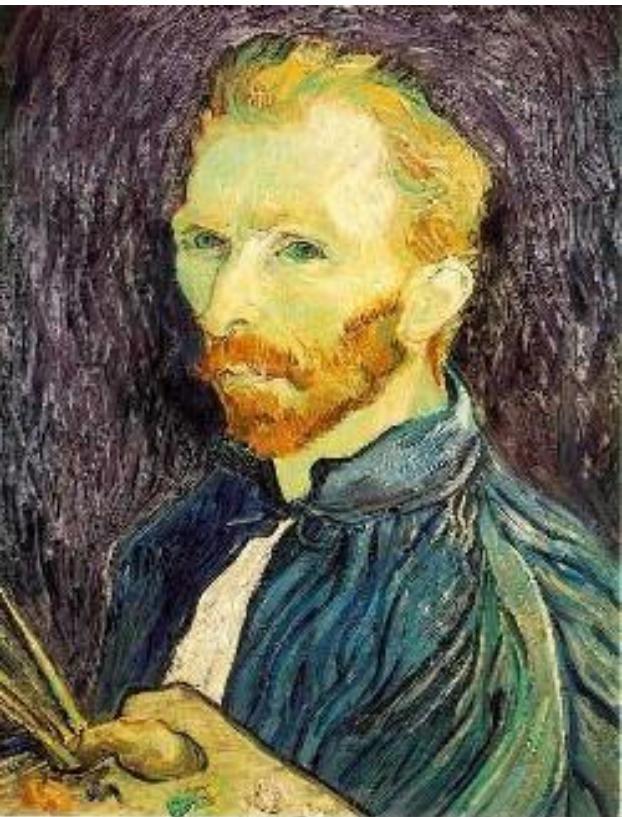


1/8

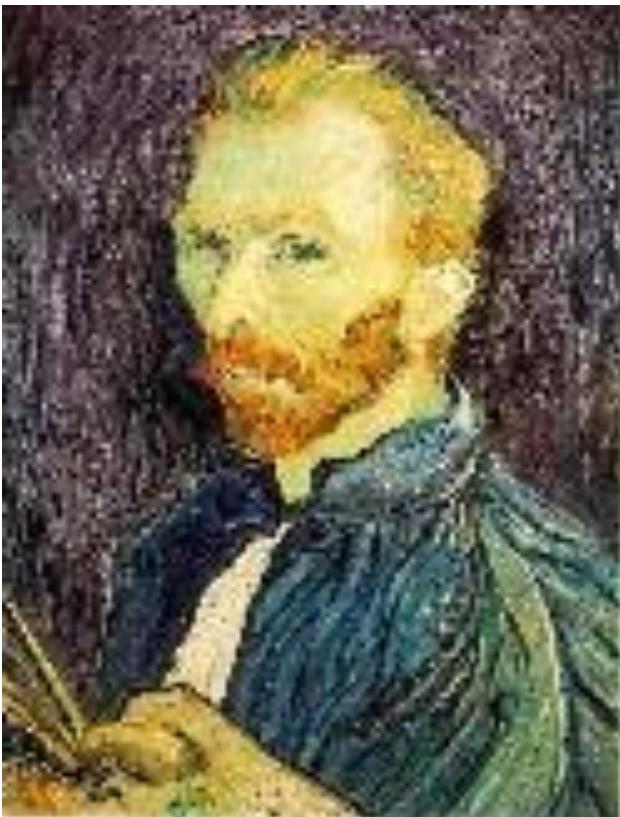
Throw away every other row and column to create a $1/2$ size image
- called *image sub-sampling*

Source: S. Narasimhan

Image Sub-Sampling



1/2



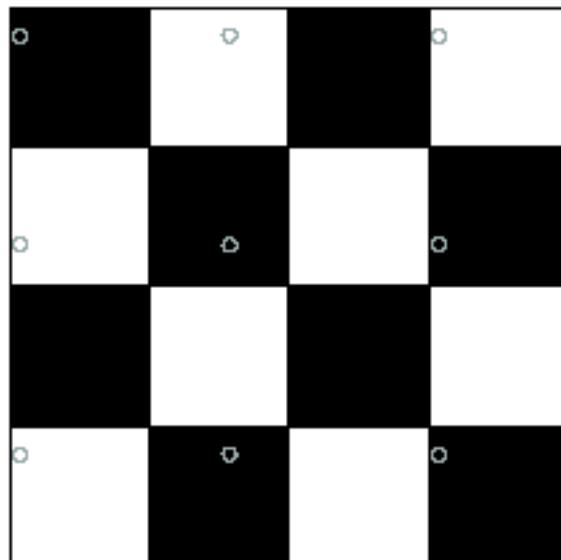
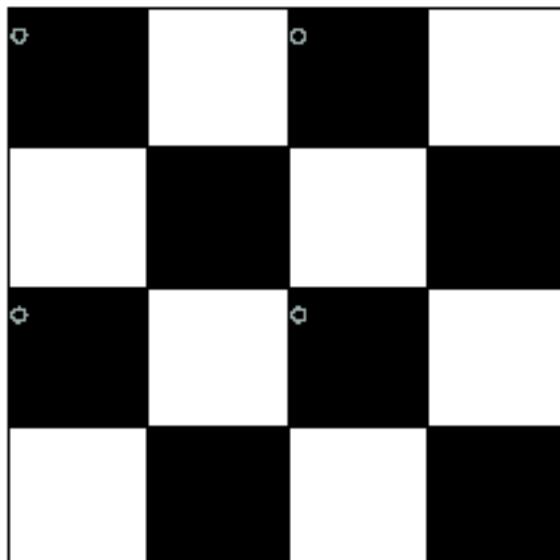
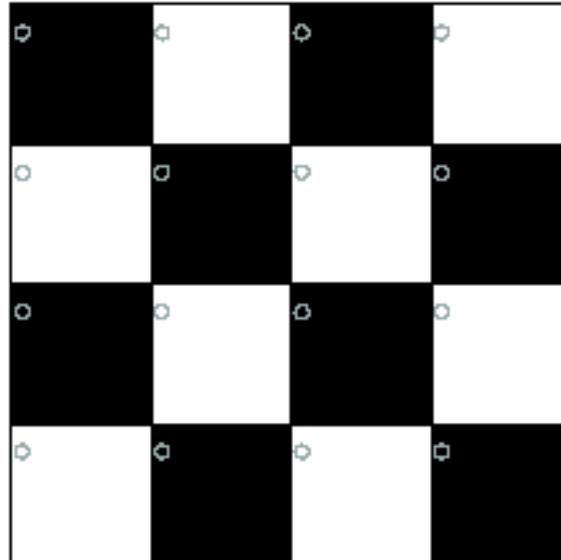
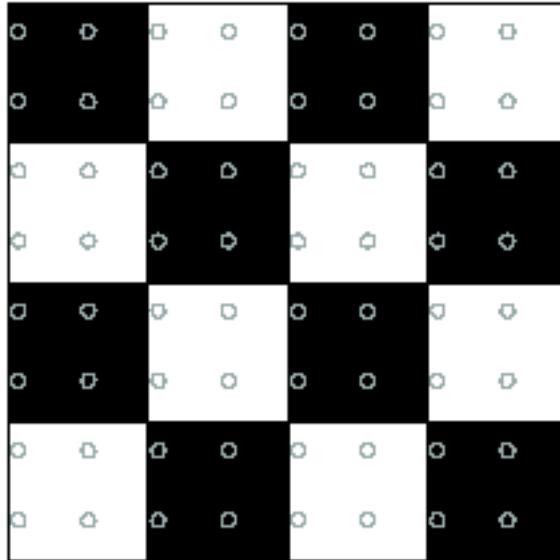
1/4 (2x zoom)



1/8 (4x zoom)

Source: S. Narasimhan

Sampling

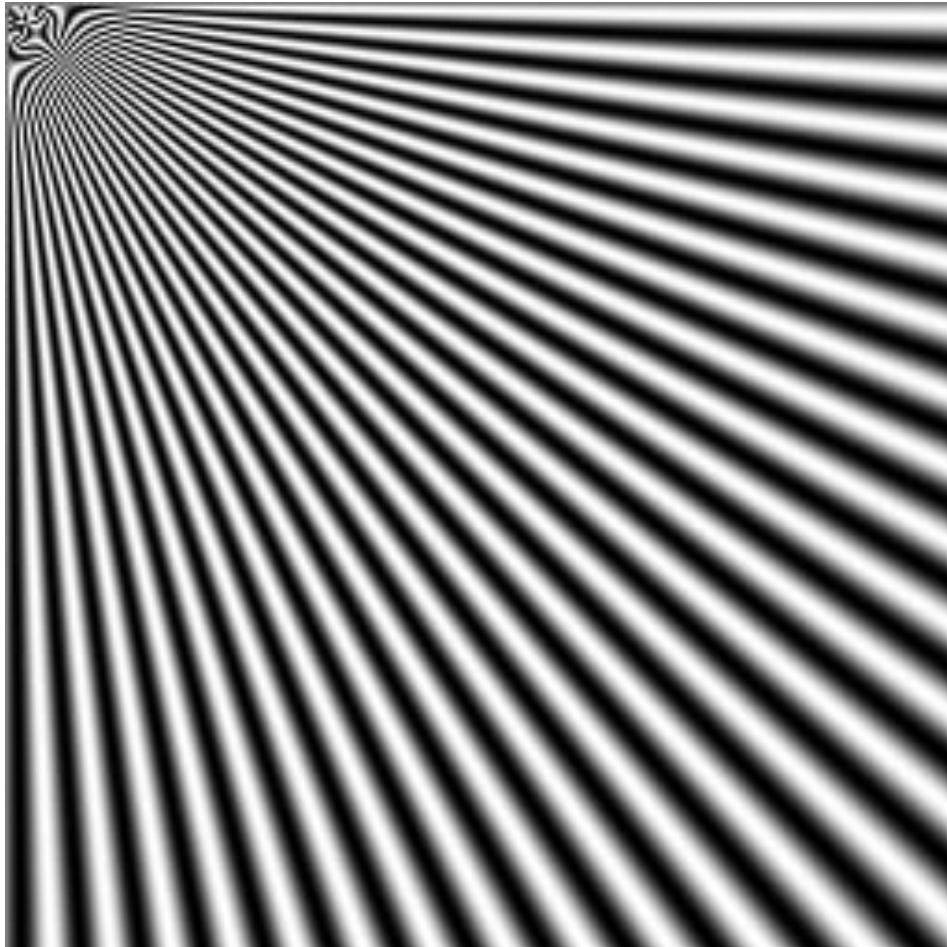


Good sampling:

- Sample often or,
- Sample wisely

Bad sampling:

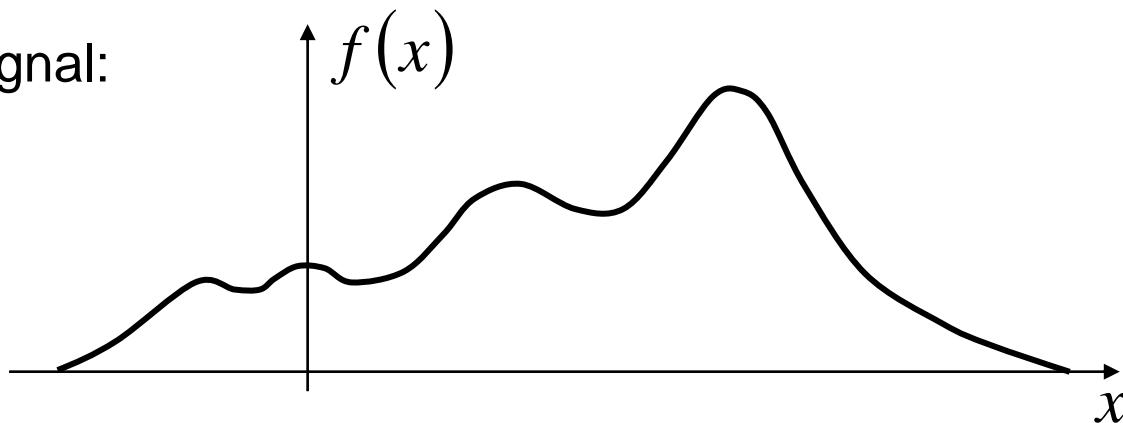
Aliasing



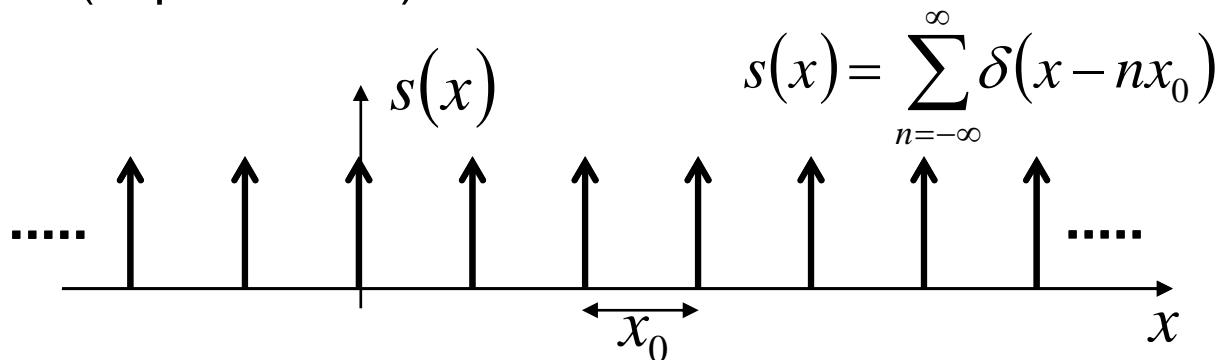
Source: S. Narasimhan

Sampling Theorem

Continuous signal:



Shah function (Impulse train):



Sampled function:

$$f_s(x) = f(x)s(x) = f(x) \sum_{n=-\infty}^{\infty} \delta(x - nx_0)$$

Source: S. Narasimhan

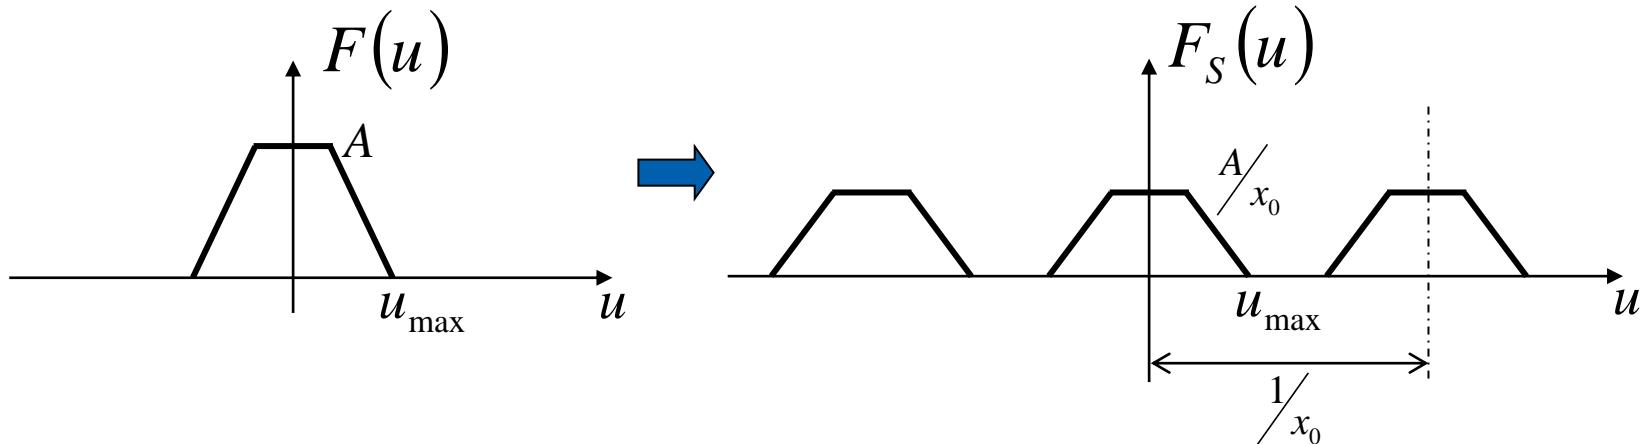
Sampling Theorem

Sampled function:

$$f_s(x) = f(x)s(x) = f(x) \sum_{n=-\infty}^{\infty} \delta(x - nx_0)$$

Sampling frequency $\frac{1}{x_0}$

$$F_s(u) = F(u) * S(u) = F(u) * \frac{1}{x_0} \sum_{n=-\infty}^{\infty} \delta\left(u - \frac{n}{x_0}\right)$$

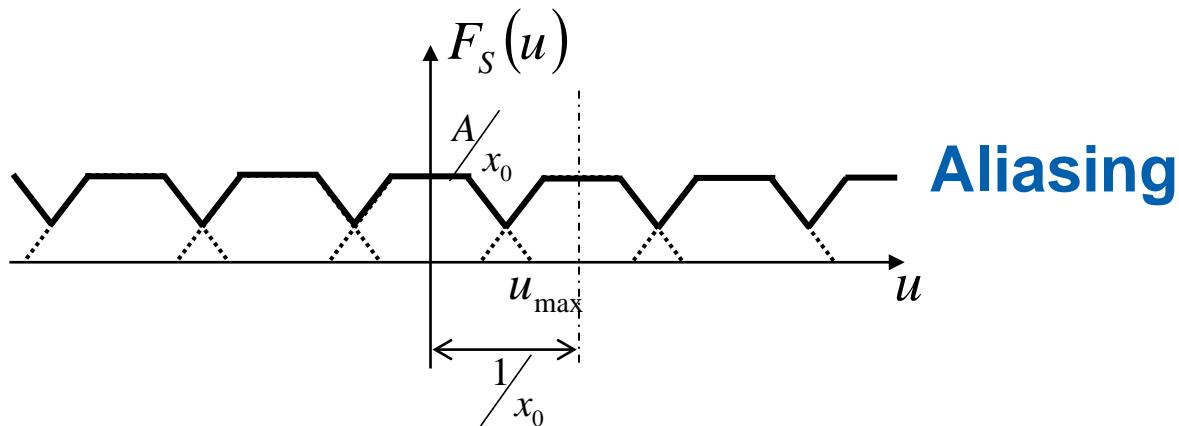


Only if $u_{\max} \leq \frac{1}{2x_0}$

Source: S. Narasimhan

Nyquist Theorem

If $u_{\max} > \frac{1}{2x_0}$



When can we recover $F(u)$ from $F_S(u)$?

Only if $u_{\max} \leq \frac{1}{2x_0}$ (Nyquist Frequency)

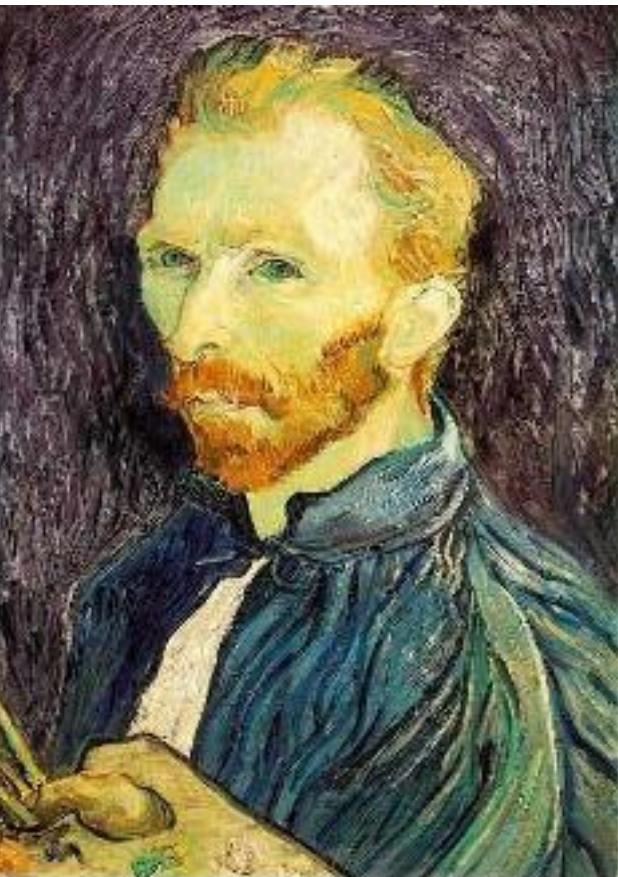
We can use

$$C(u) = \begin{cases} x_0 & |u| < \frac{1}{2x_0} \\ 0 & \text{otherwise} \end{cases}$$

Then $F(u) = F_S(u)C(u)$ and $f(x) = \text{IFT}[F(u)]$

Sampling frequency must be greater than $2u_{\max}$

Image Sub-Sampling



1/2



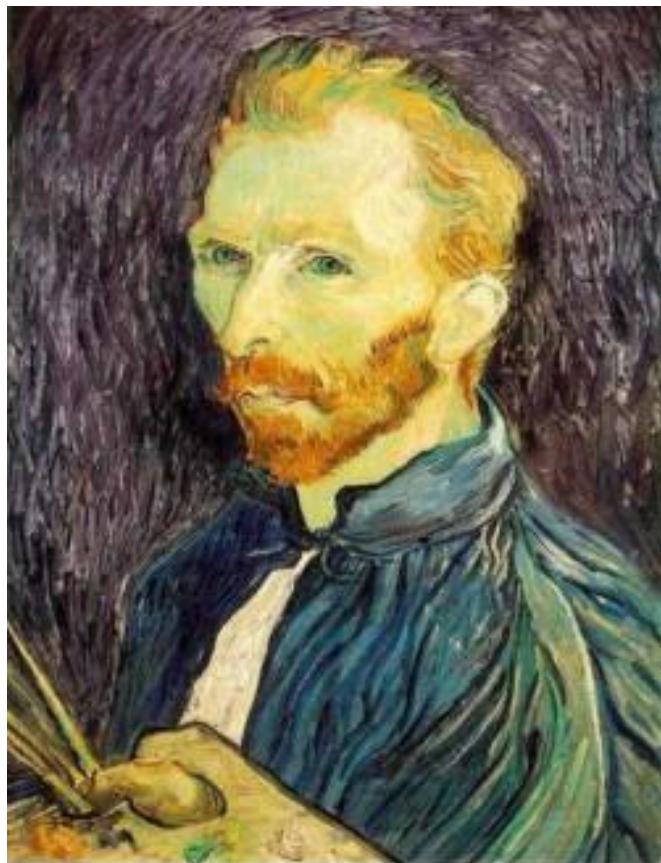
1/4 (2x zoom)



1/8 (4x zoom)

Source: S. Narasimhan

Sub-Sampling with Gaussian Pre-Filtering



Gaussian 1/2



G 1/4

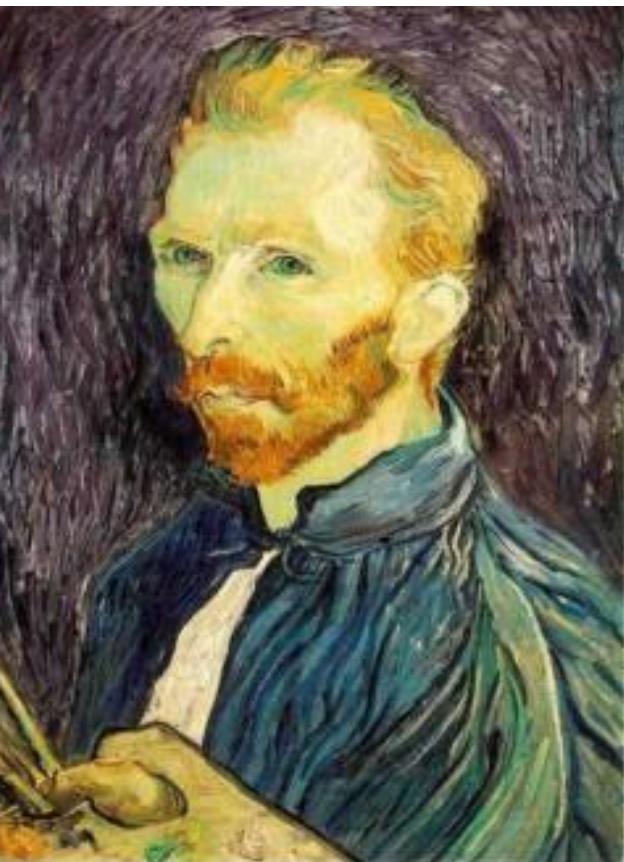


G 1/8

- Solution: filter the image, *then* subsample
 - Filter size should double for each $\frac{1}{2}$ size reduction.

Source: S. Narasimhan

Sub-Sampling with Gaussian Pre-Filtering



Gaussian 1/2



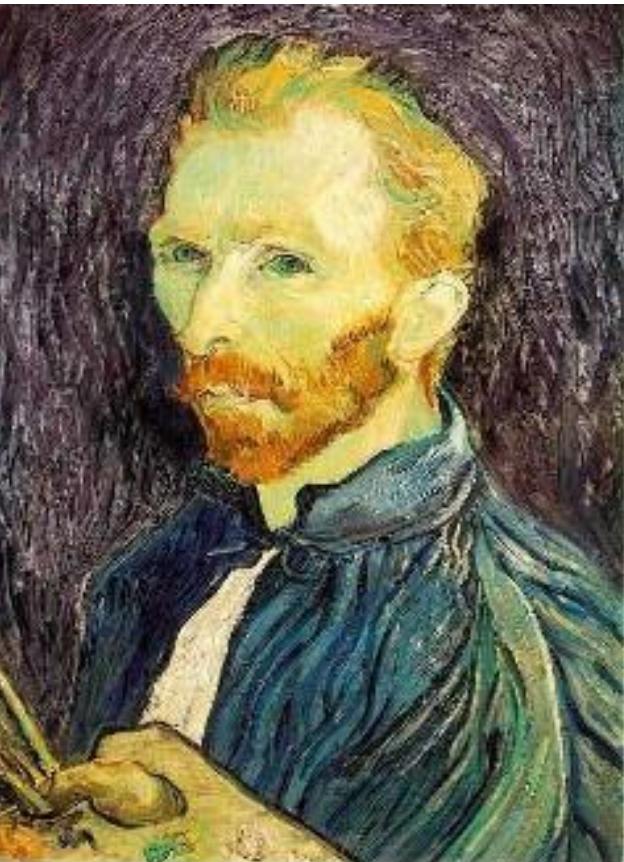
G 1/4



G 1/8

Source: S. Narasimhan

Compare with...



1/2



1/4 (2x zoom)



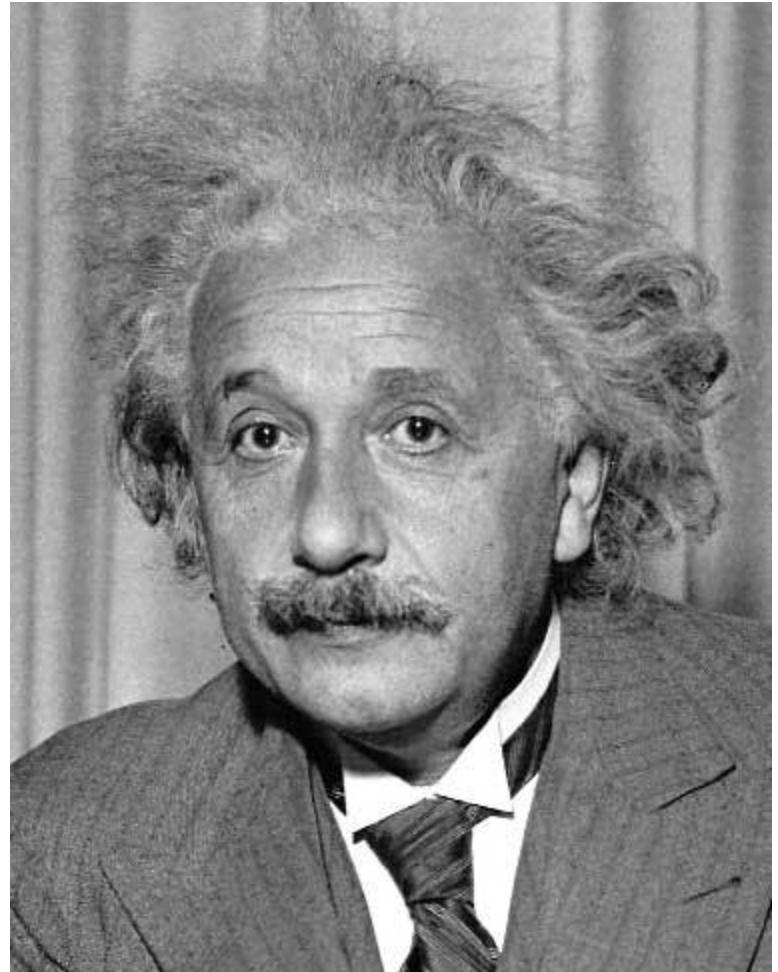
1/8 (4x zoom)

Source: S. Narasimhan

Template matching

Goal: find  in image

Main challenge: What is a good similarity or distance measure between two patches?



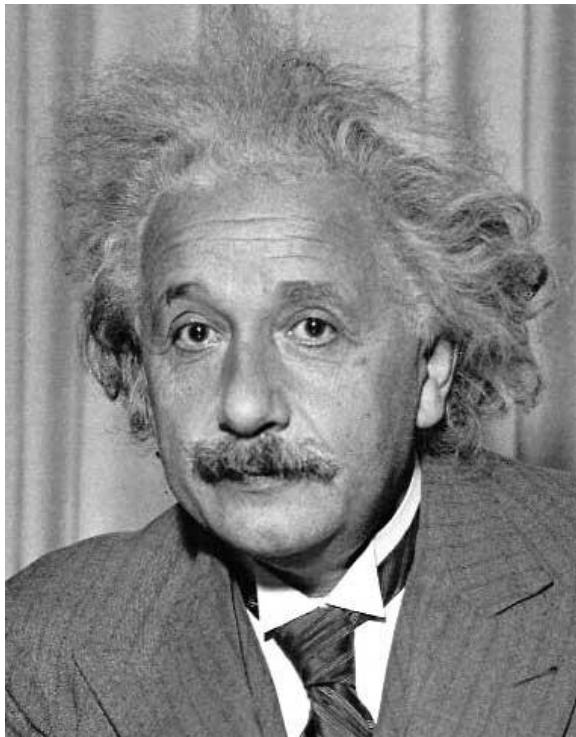
Source: D. Hoiem

Correlation

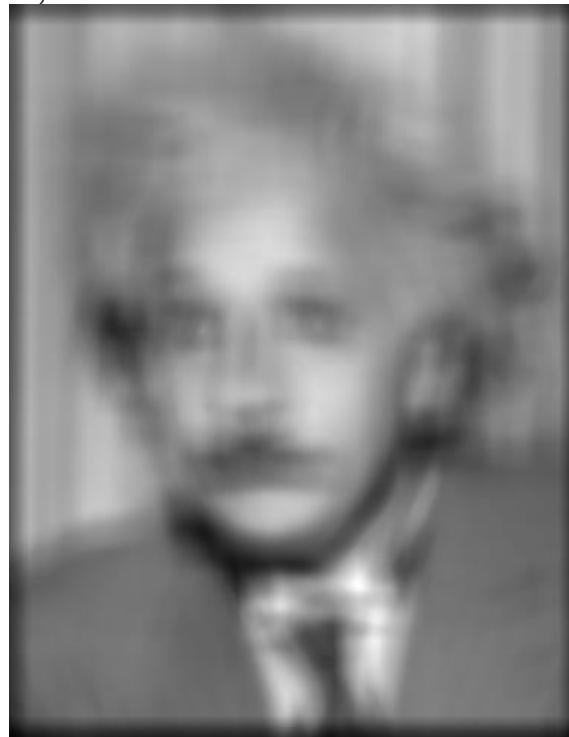
Goal: find  in image

- Filter the image with eye patch (correlation)

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k, n+l]$$



Input



Filtered Image

f = image
g = filter

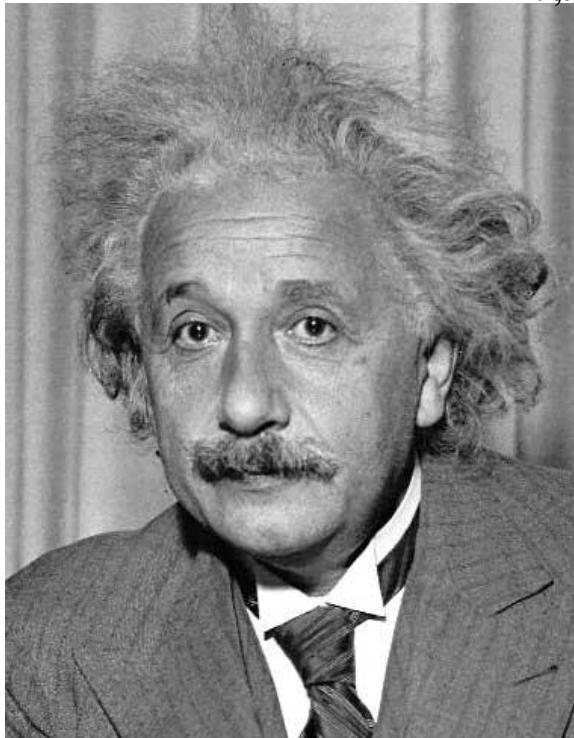
What went wrong?

Zero-mean correlation

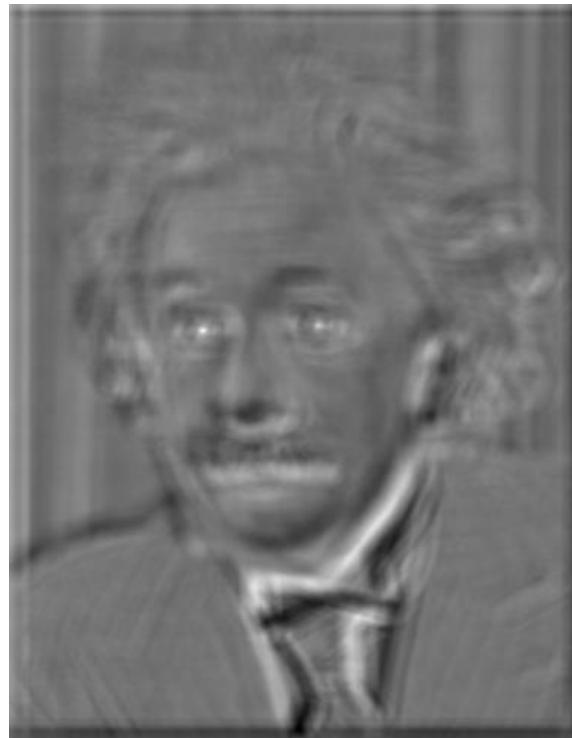
Goal: find  in image

Method 1: filter the image with zero-mean eye

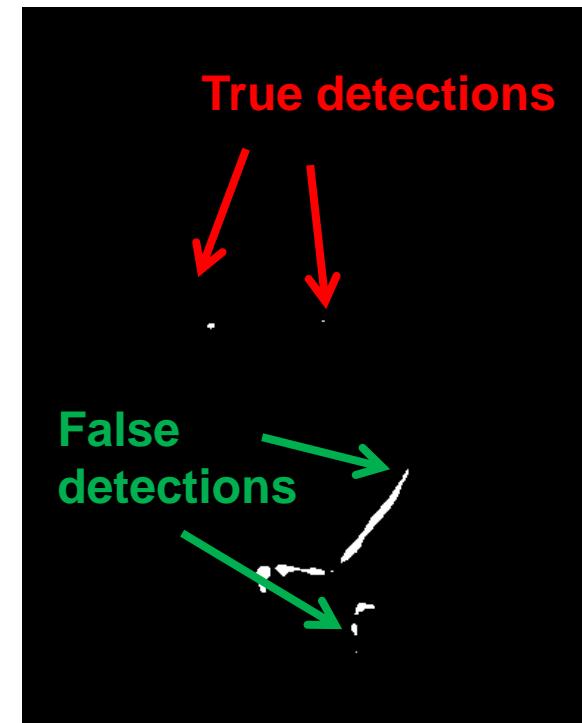
$$h[m, n] = \sum_{k, l} (g[k, l] - \bar{g})(f[m + k, n + l])$$



Input



Filtered Image (scaled)



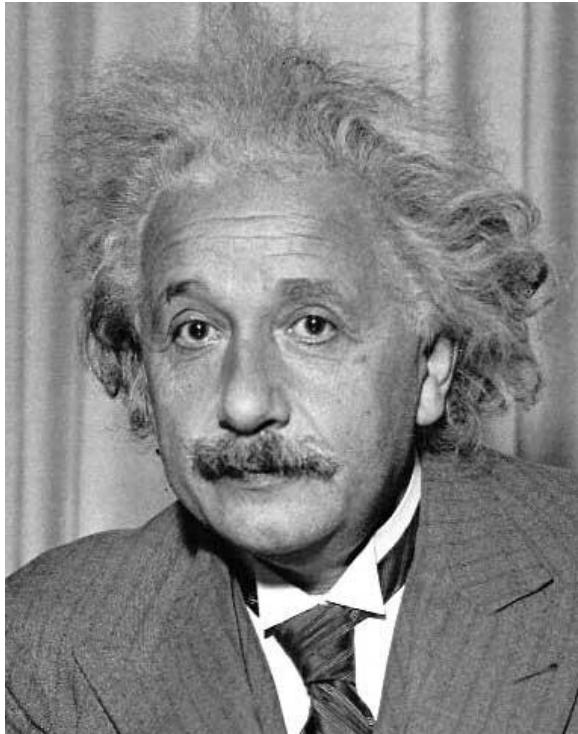
Thresholded Image

Source: D. Hoiem

Sum Square Difference

Goal: find  in image

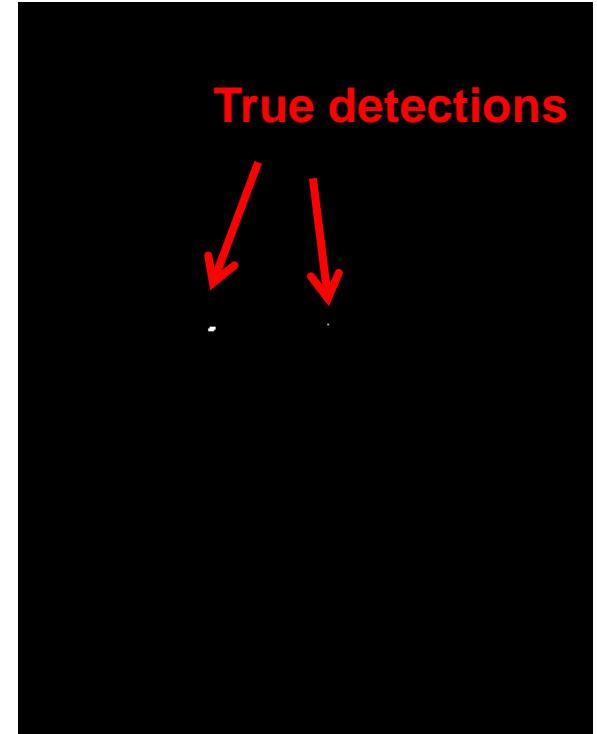
Method 2: SSD $h[m,n] = \sum_{k,l} (g[k,l] - f[m+k, n+l])^2$



Input



1- sqrt(SSD)

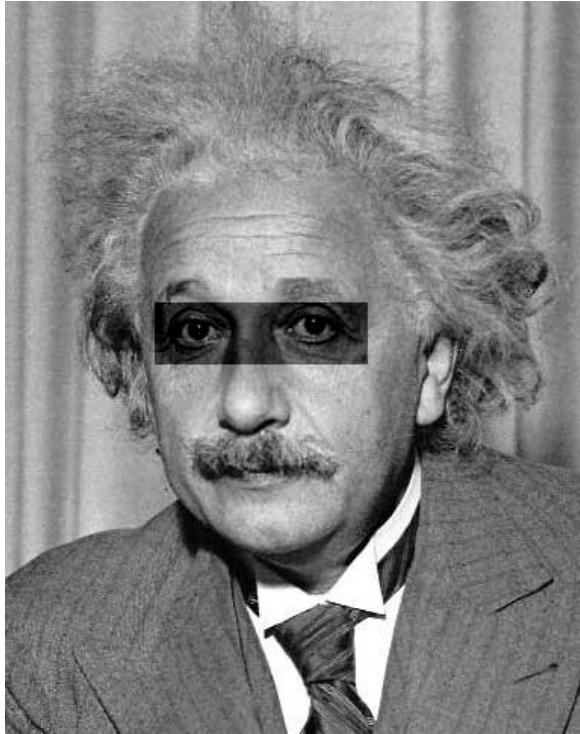


Thresholded Image

Sum Square Difference

Goal: find  in image

Method 2: SSD $h[m, n] = \sum_{k, l} (g[k, l] - f[m + k, n + l])^2$



Input



1- sqrt(SSD)

Source: D. Hoiem

Normalized Cross Correlation

Goal: find  in image

Method 3: Normalized cross-correlation

$$h[m, n] = \frac{\sum_{k,l} (g[k, l] - \bar{g})(f[m + k, n + l] - \bar{f}_{m,n})}{\left(\sum_{k,l} (g[k, l] - \bar{g})^2 \sum_{k,l} (f[m + k, n + l] - \bar{f}_{m,n})^2 \right)^{0.5}}$$

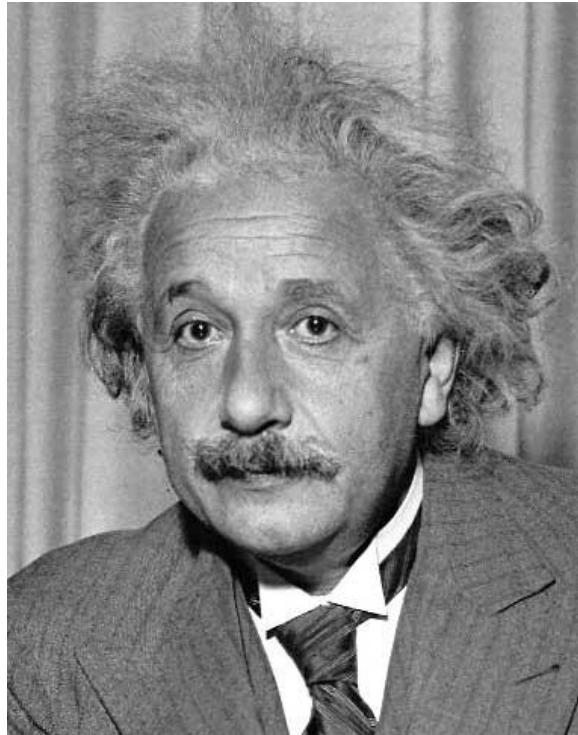
mean template
↓
 $\sum_{k,l} (g[k, l] - \bar{g})(f[m + k, n + l] - \bar{f}_{m,n})$

mean image patch
↓
 $\sum_{k,l} (g[k, l] - \bar{g})^2 \sum_{k,l} (f[m + k, n + l] - \bar{f}_{m,n})^2$

Normalized Cross Correlation

Goal: find  in image

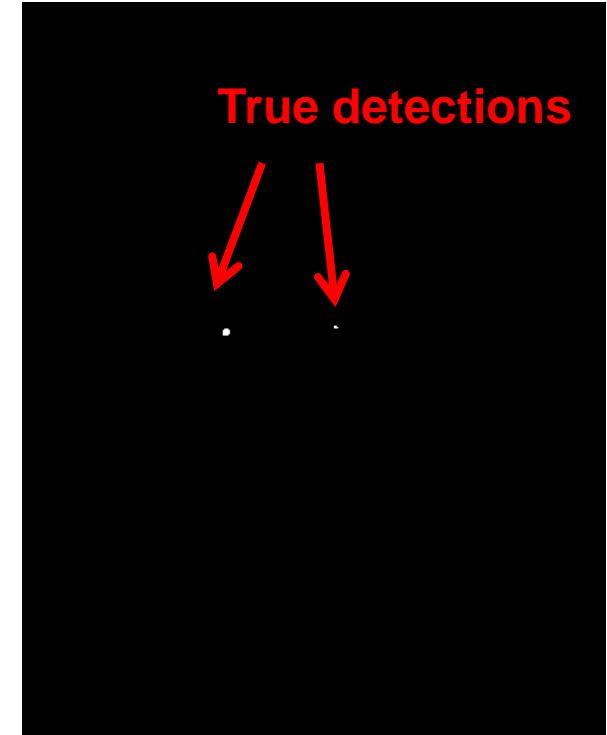
Method 3: Normalized cross-correlation



Input



Normalized X-Correlation

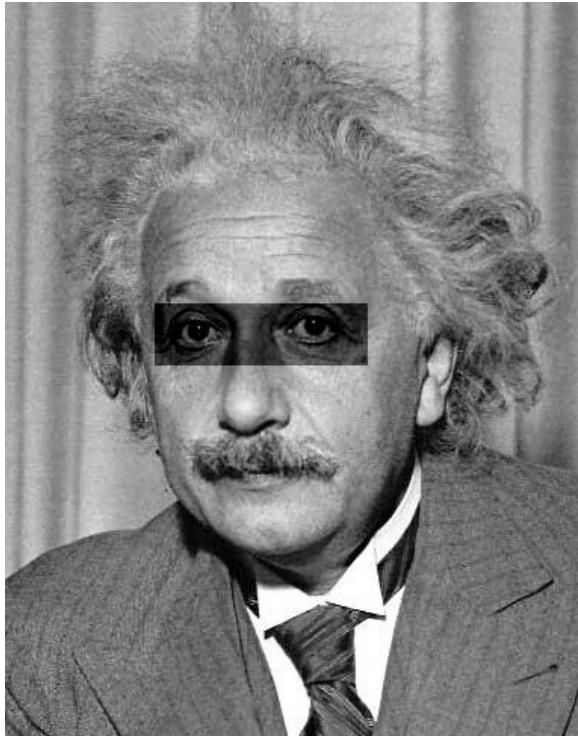


True detections
Thresholded Image

Normalized Cross Correlation

Goal: find  in image

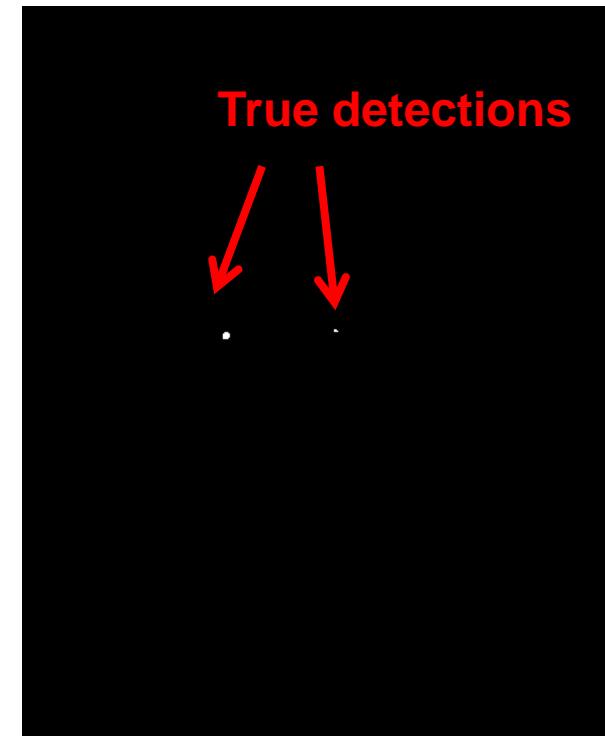
Method 3: Normalized cross-correlation



Input



Normalized X-Correlation



Thresholded Image

Fast template matching

Template



Search Region



How make it faster?

Source: S. Narasimhan

Multi-resolution

Template



Search Region

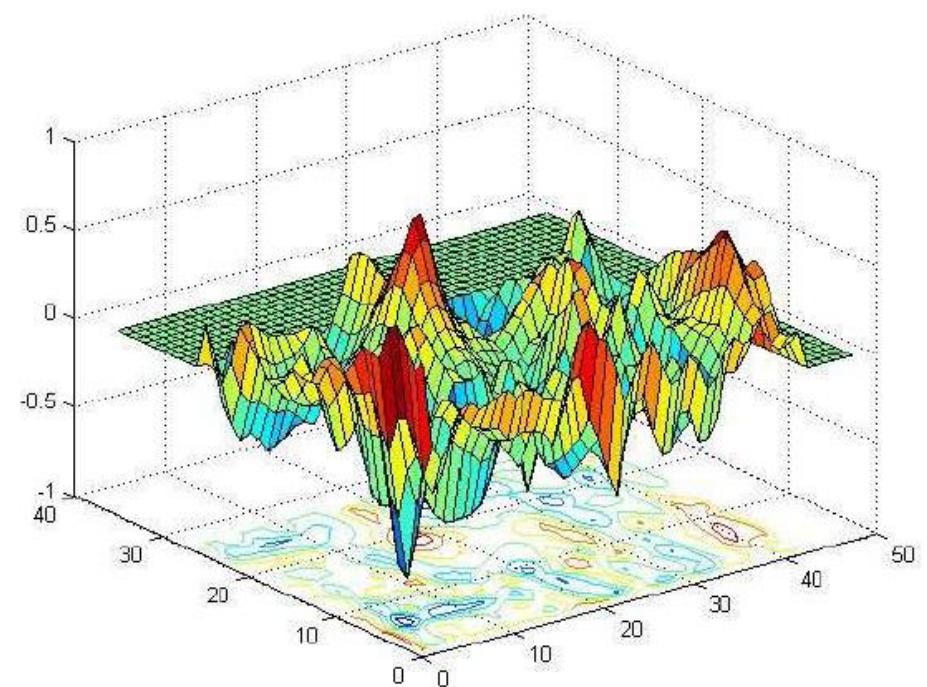
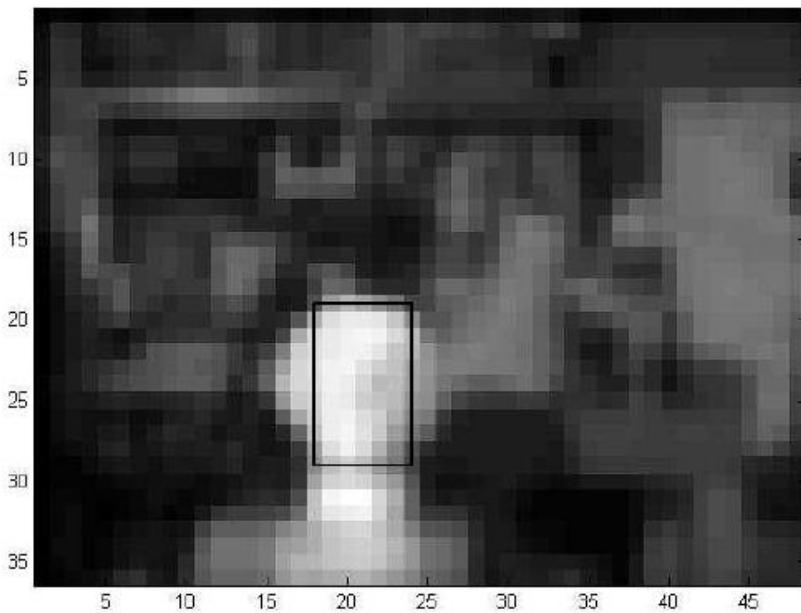
Original Image



Source: S. Narasimhan

Level 3 Search

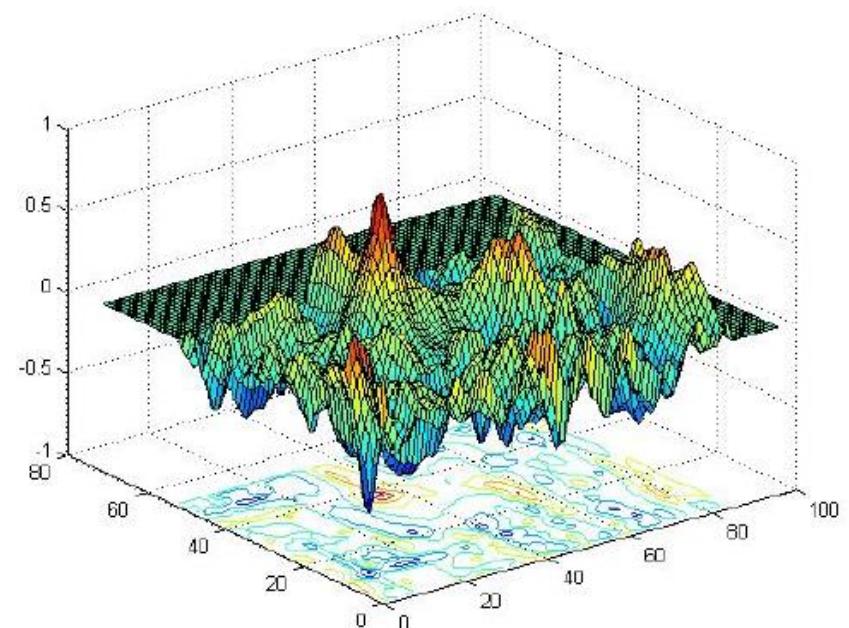
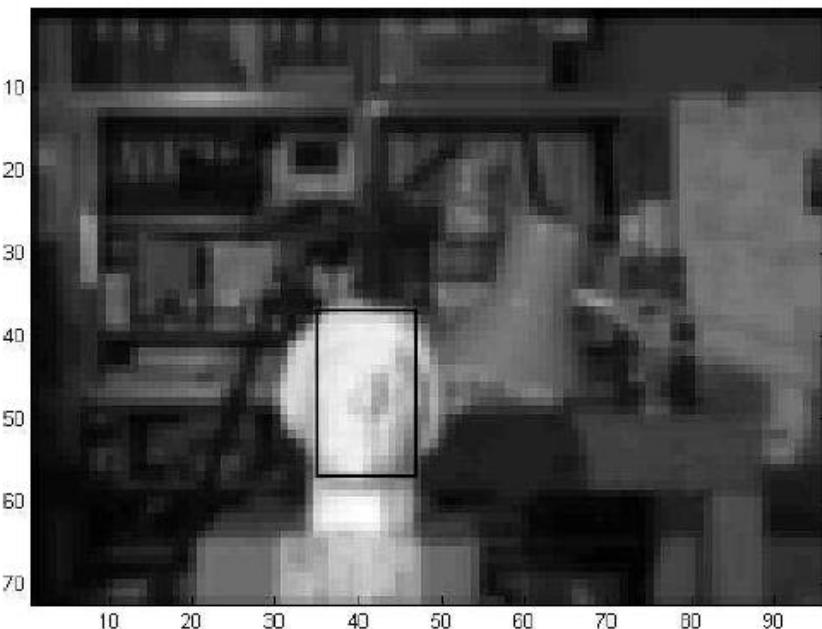
- At the lowest pyramid level, we search the entire image with the correlation template



Source: S. Narasimhan

Level 2 Search

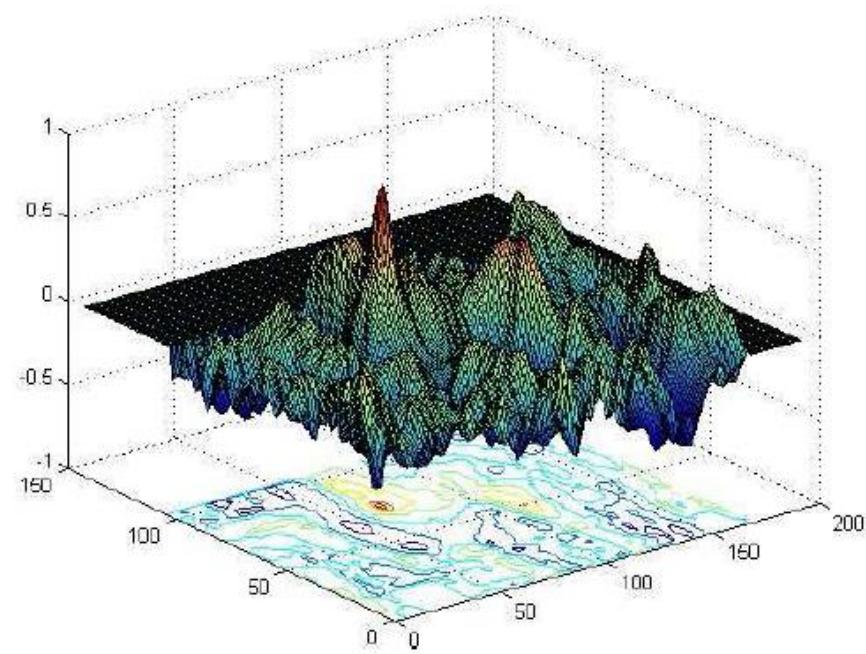
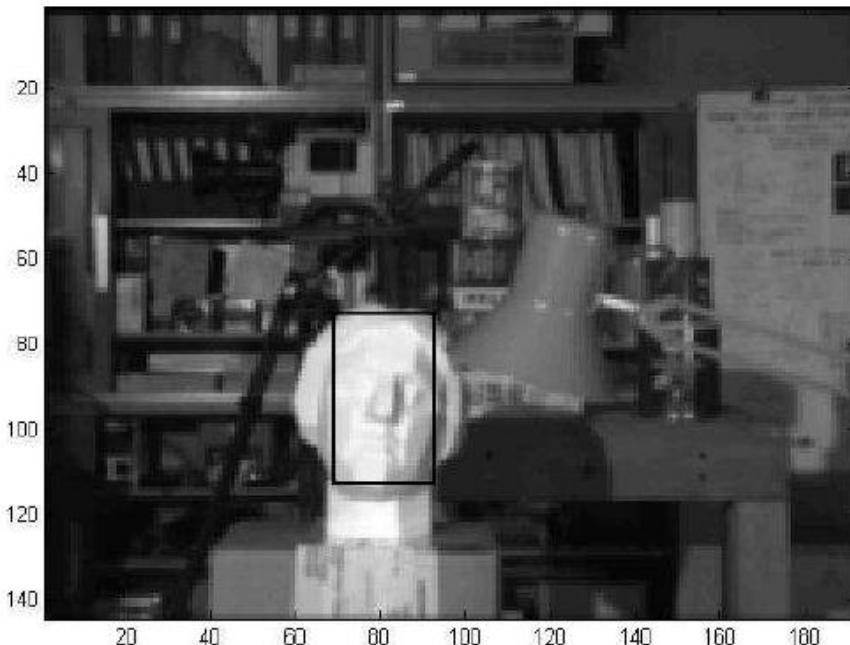
- Subsequent searches are constrained to a neighborhood of only several pixels in the x and y directions



Source: S. Narasimhan

Level 1 Search

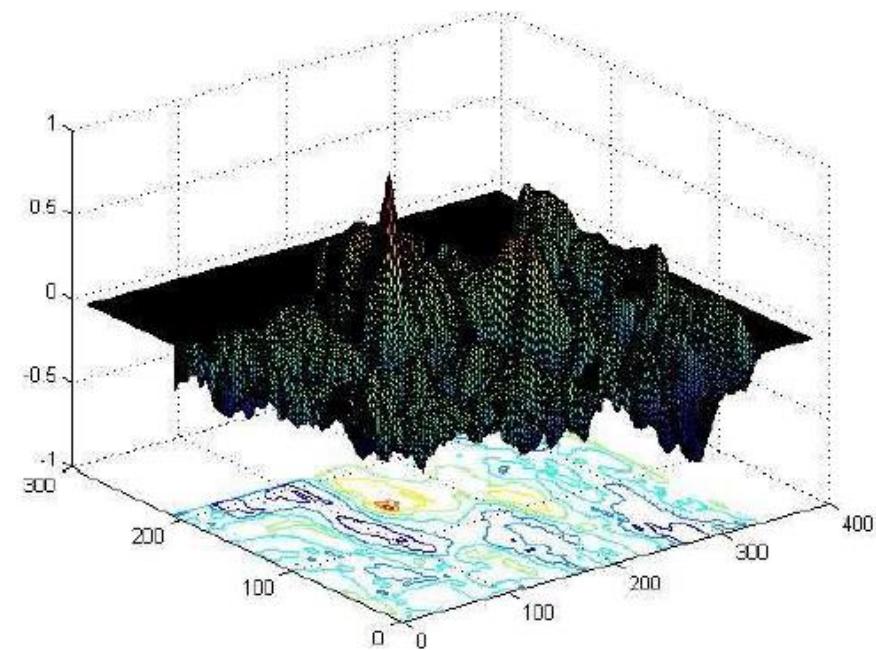
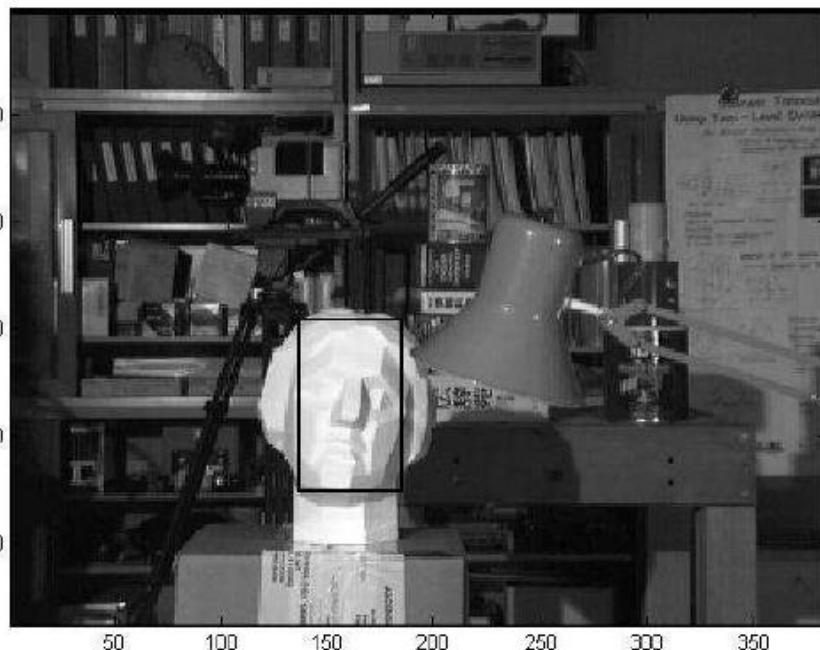
- Subsequent searches are constrained to a neighborhood of only several pixels in the x and y directions



Source: S. Narasimhan

Level 0 Search

- In the end, the total time (in Matlab) was reduced from ≈ 31 seconds to ≈ 0.5 seconds while obtaining the same template match



Source: S. Narasimhan

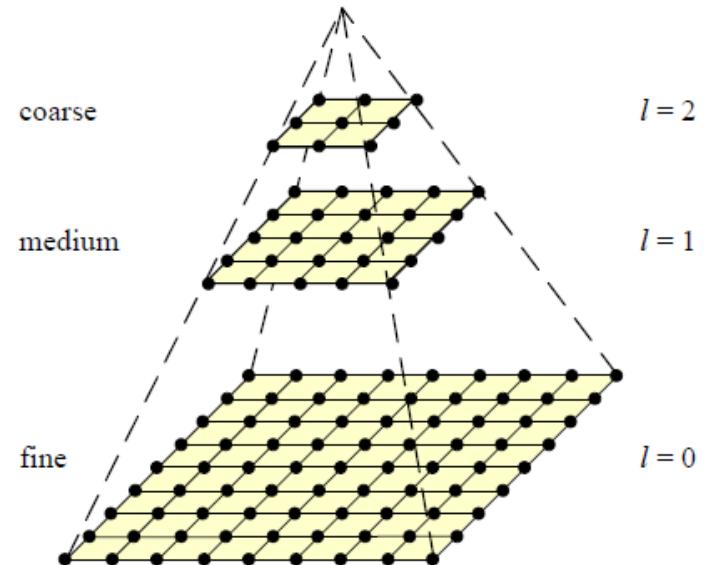
Coarse-to-fine Image Registration

Compute Gaussian pyramid

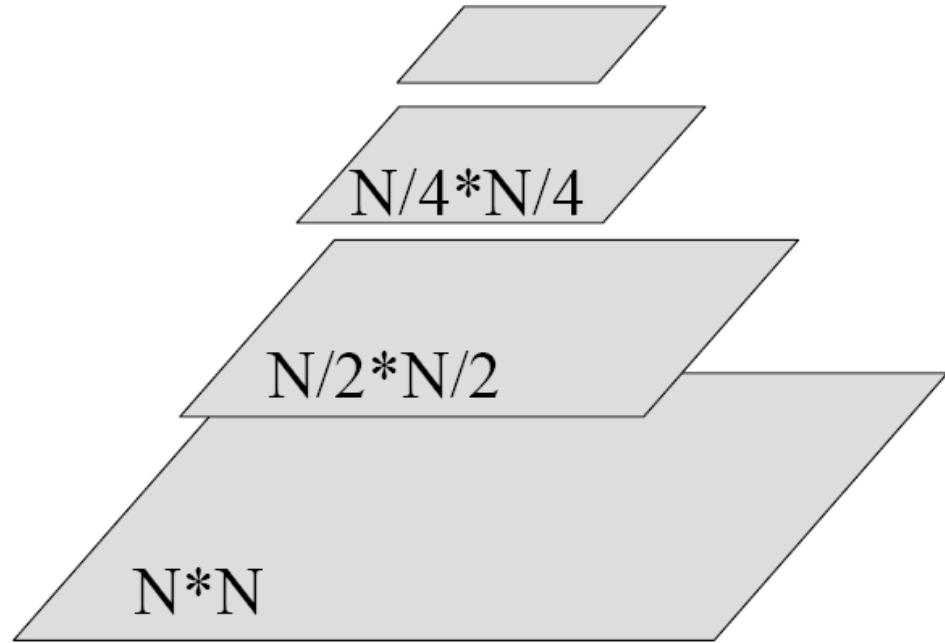
Align with coarse layer

Successively align with finer layers

Search smaller range



Space requirements

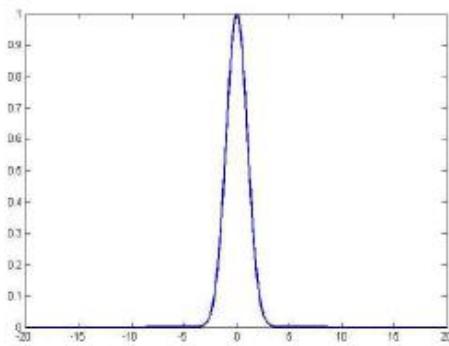


$$N^2 + \frac{1}{4}N^2 + \frac{1}{16}N^2 + \dots = 1\frac{1}{3}N^2$$

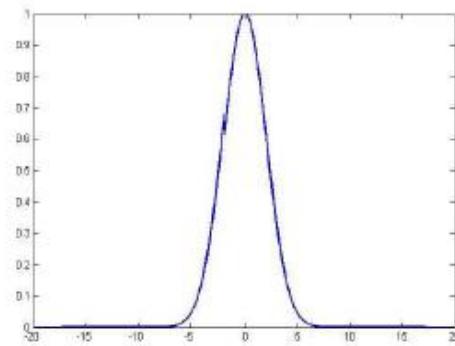
Source: S. Narasimhan

The Gaussian Pyramid

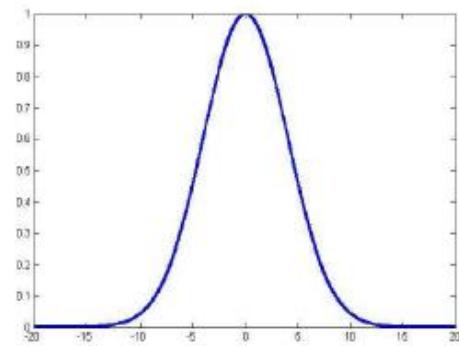
- Smooth with Gaussians because
 - a Gaussian*Gaussian=another Gaussian
- Synthesis
 - smooth and downsample
- Gaussians are low pass filters, so repetition is redundant
- Kernel width doubles with each level



Level 1



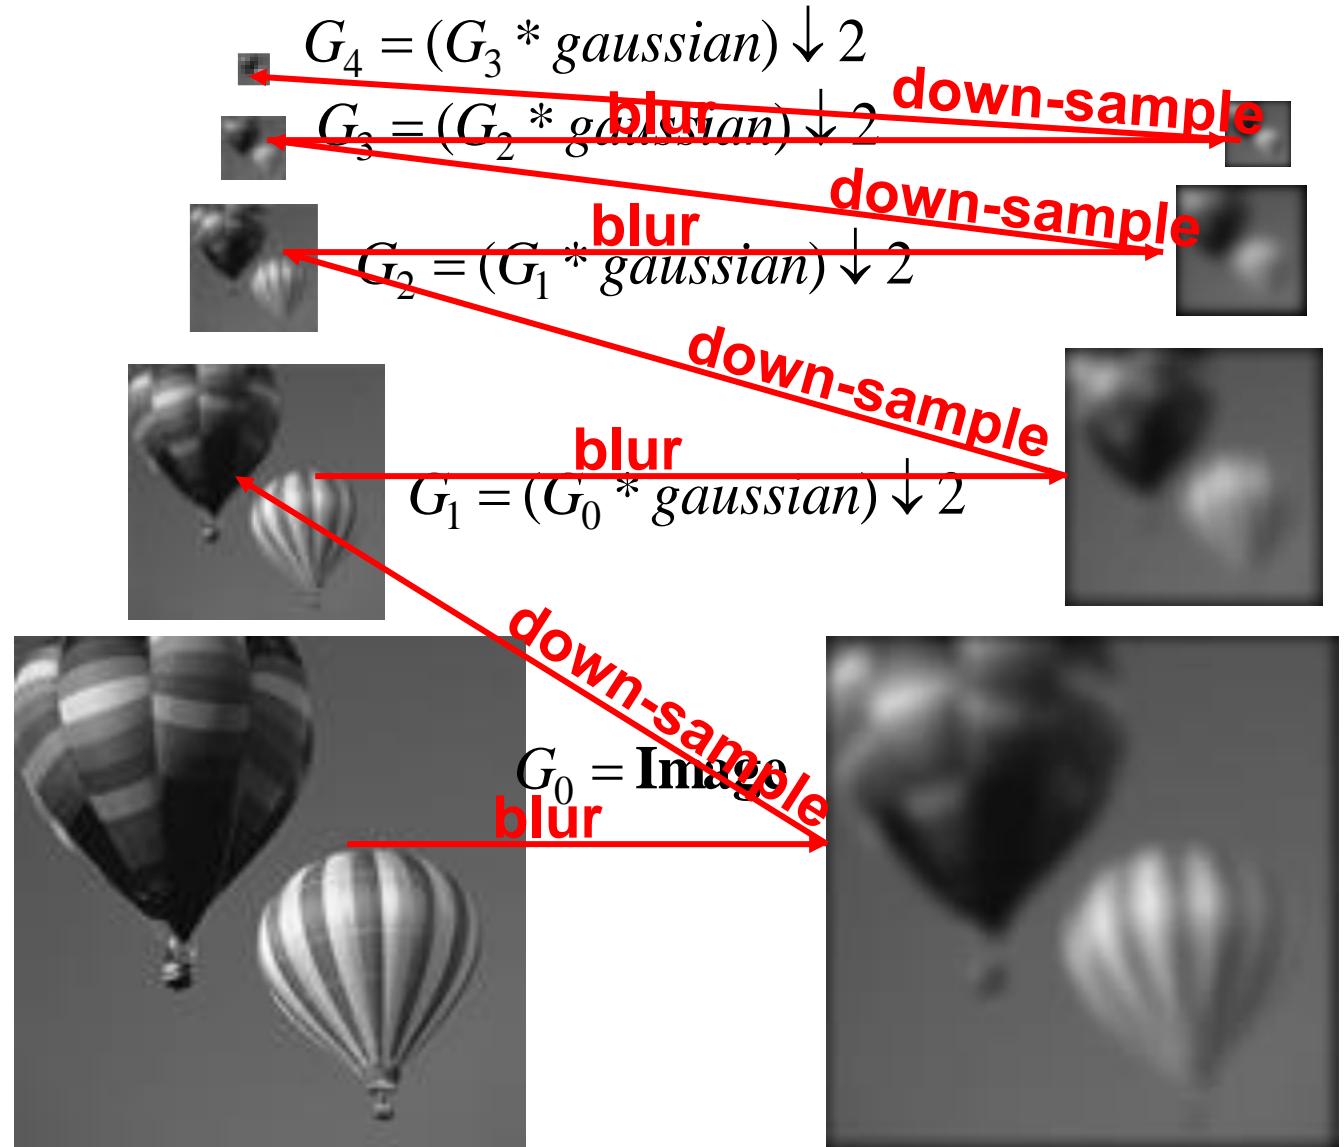
Level 2



Level 3

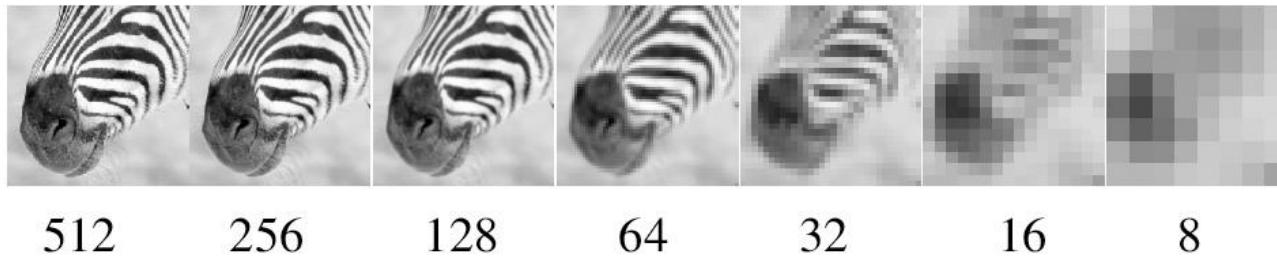
Gaussian pyramid

Low resolution



High resolution

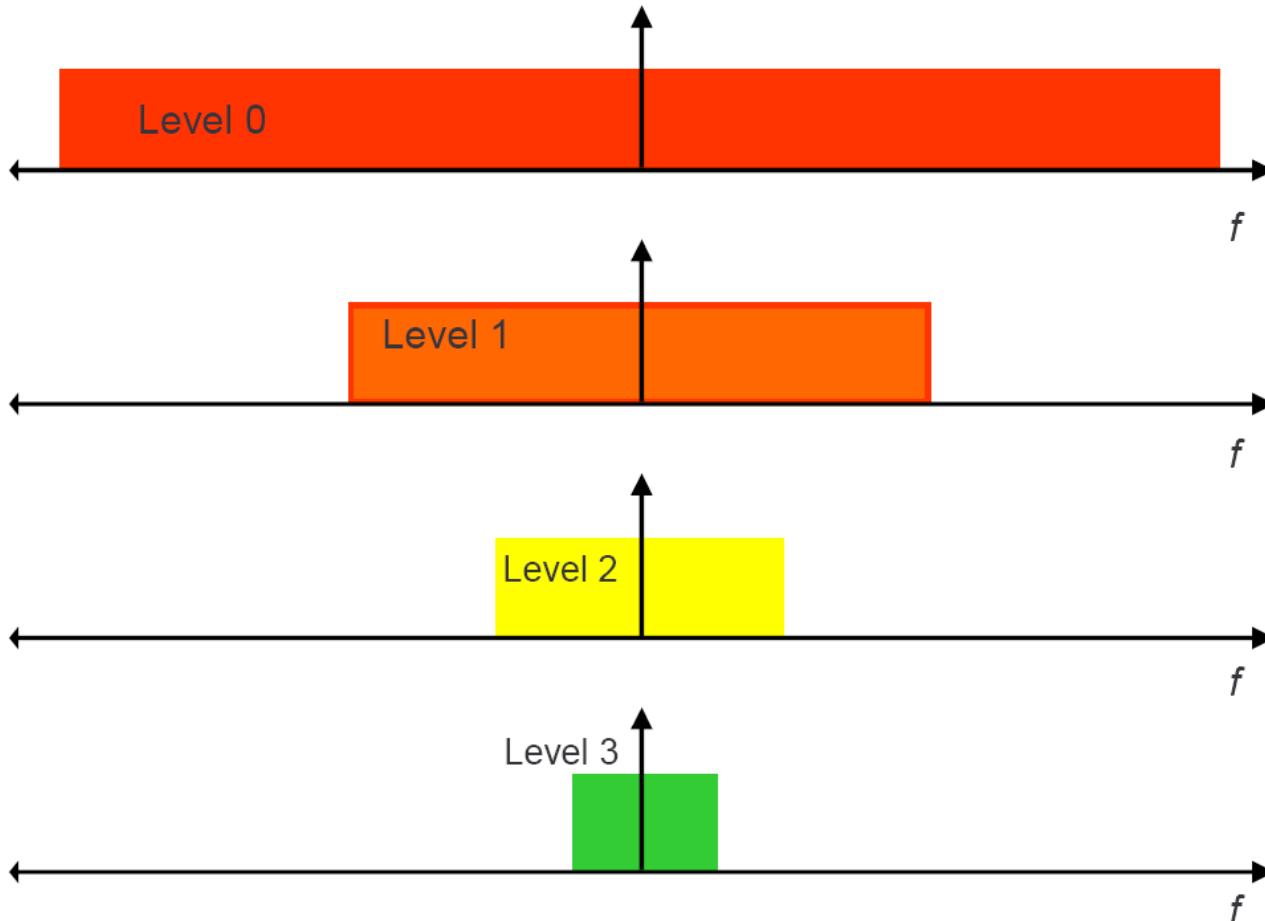
Gaussian pyramid



Source: D. Forsyth

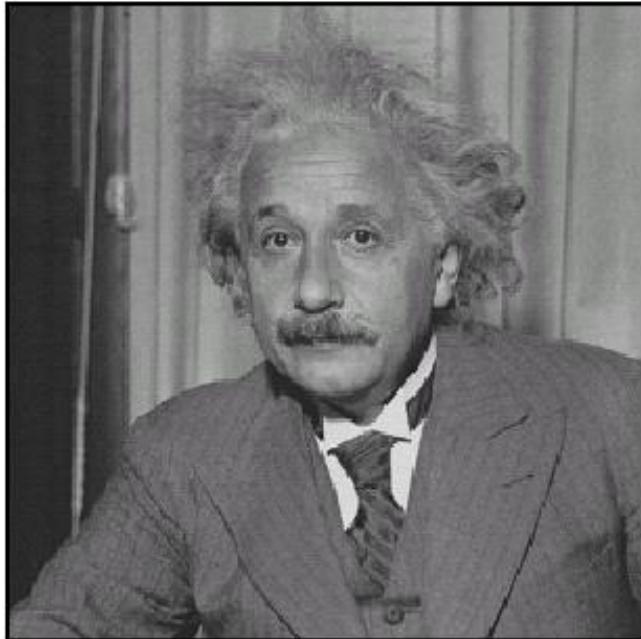
Frequencies

Gaussian Pyramid Frequency Composition

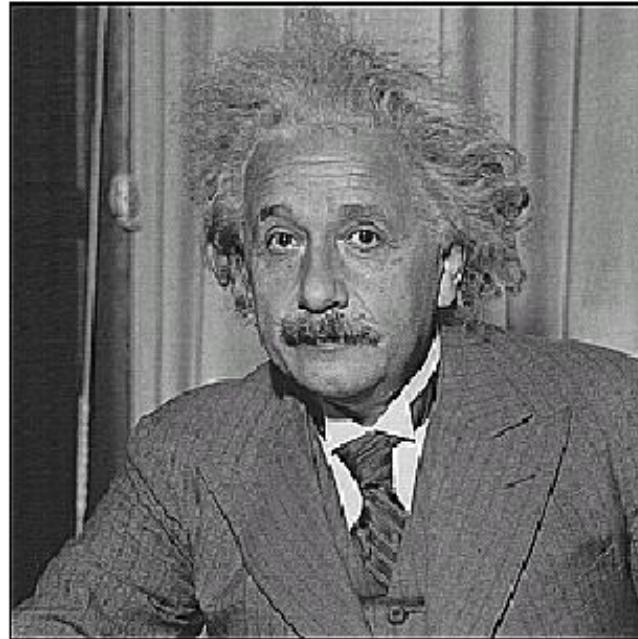


Source: S. Narasimhan

Sharpening revisited



before



after

Source: D. Lowe

Sharpening

What does blurring take away?



-



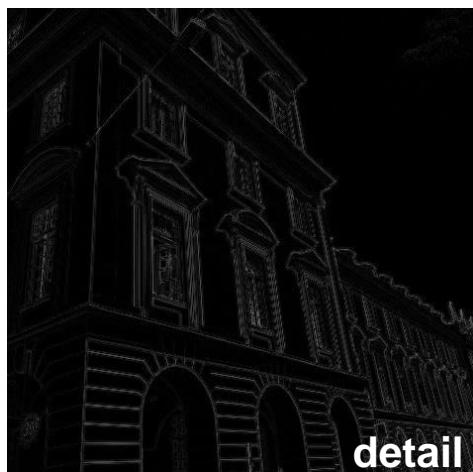
=



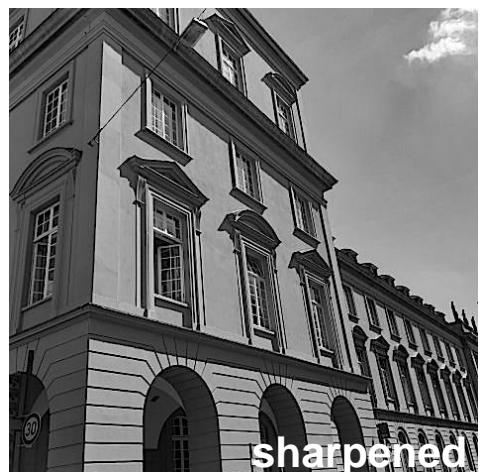
Let's add it back:



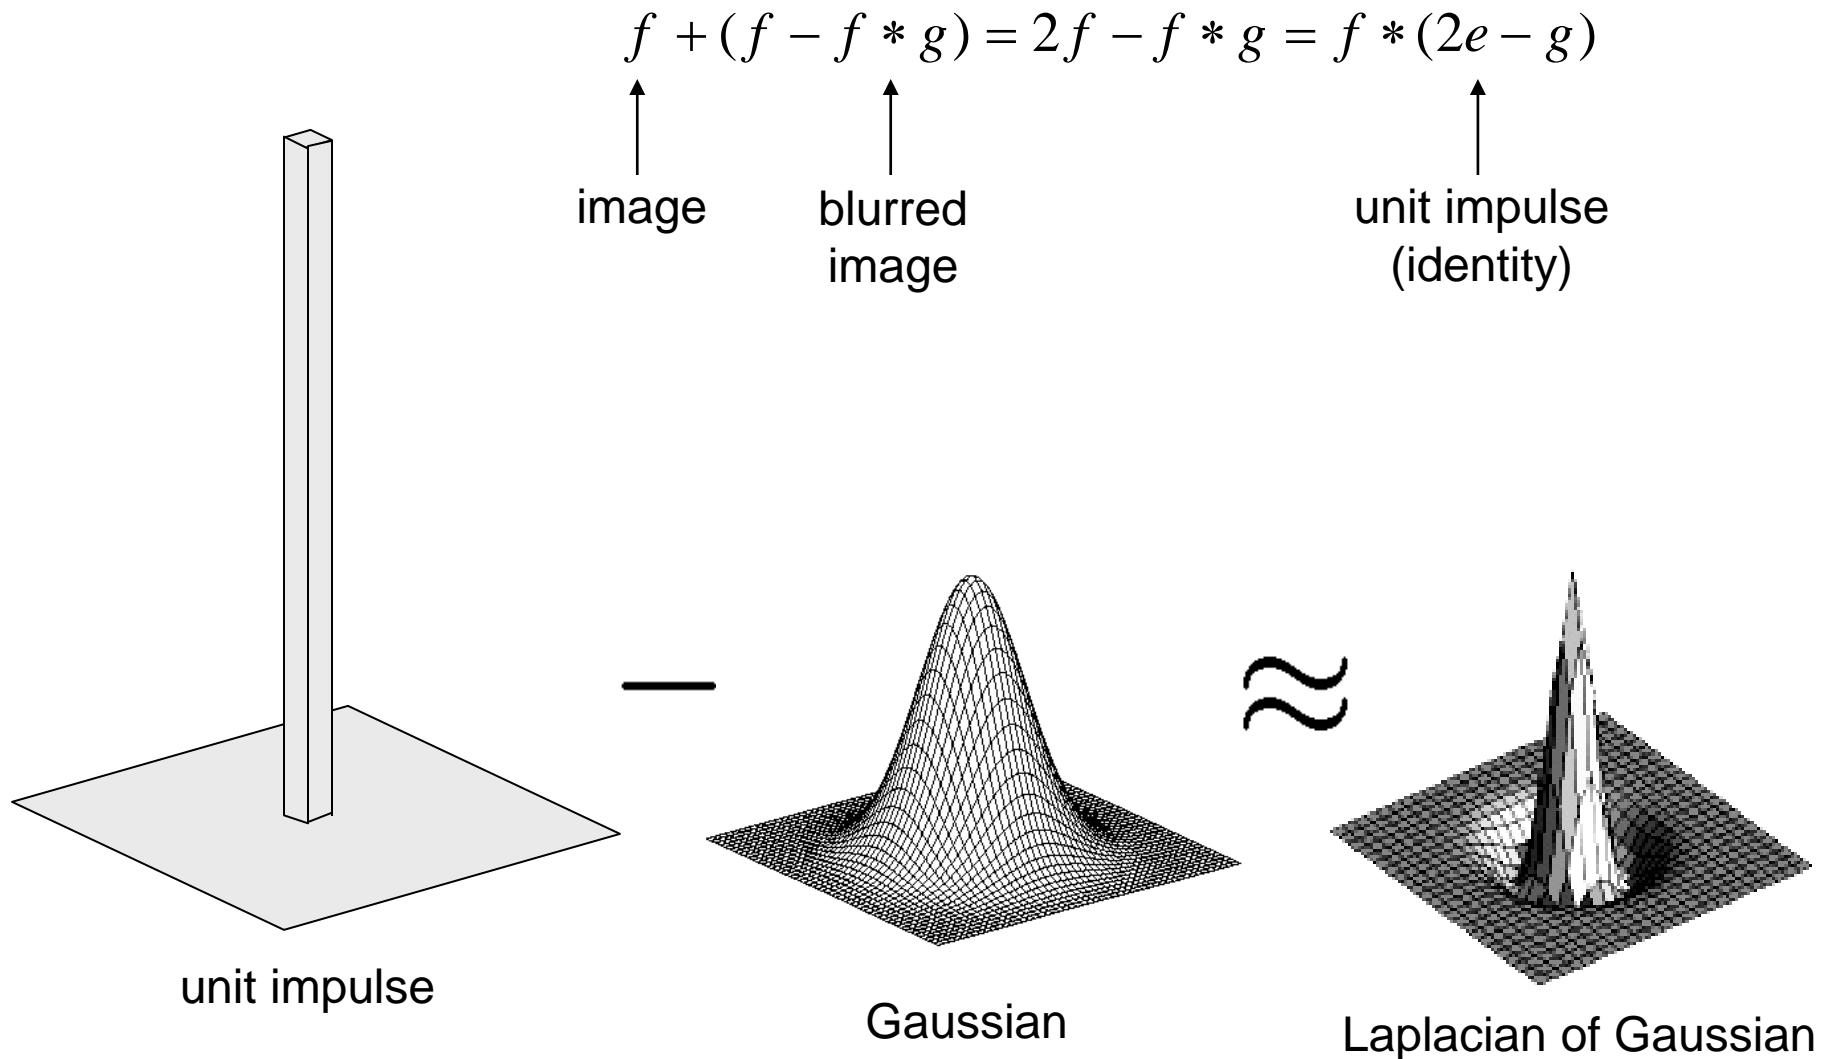
+



=



Laplacian of Gaussian

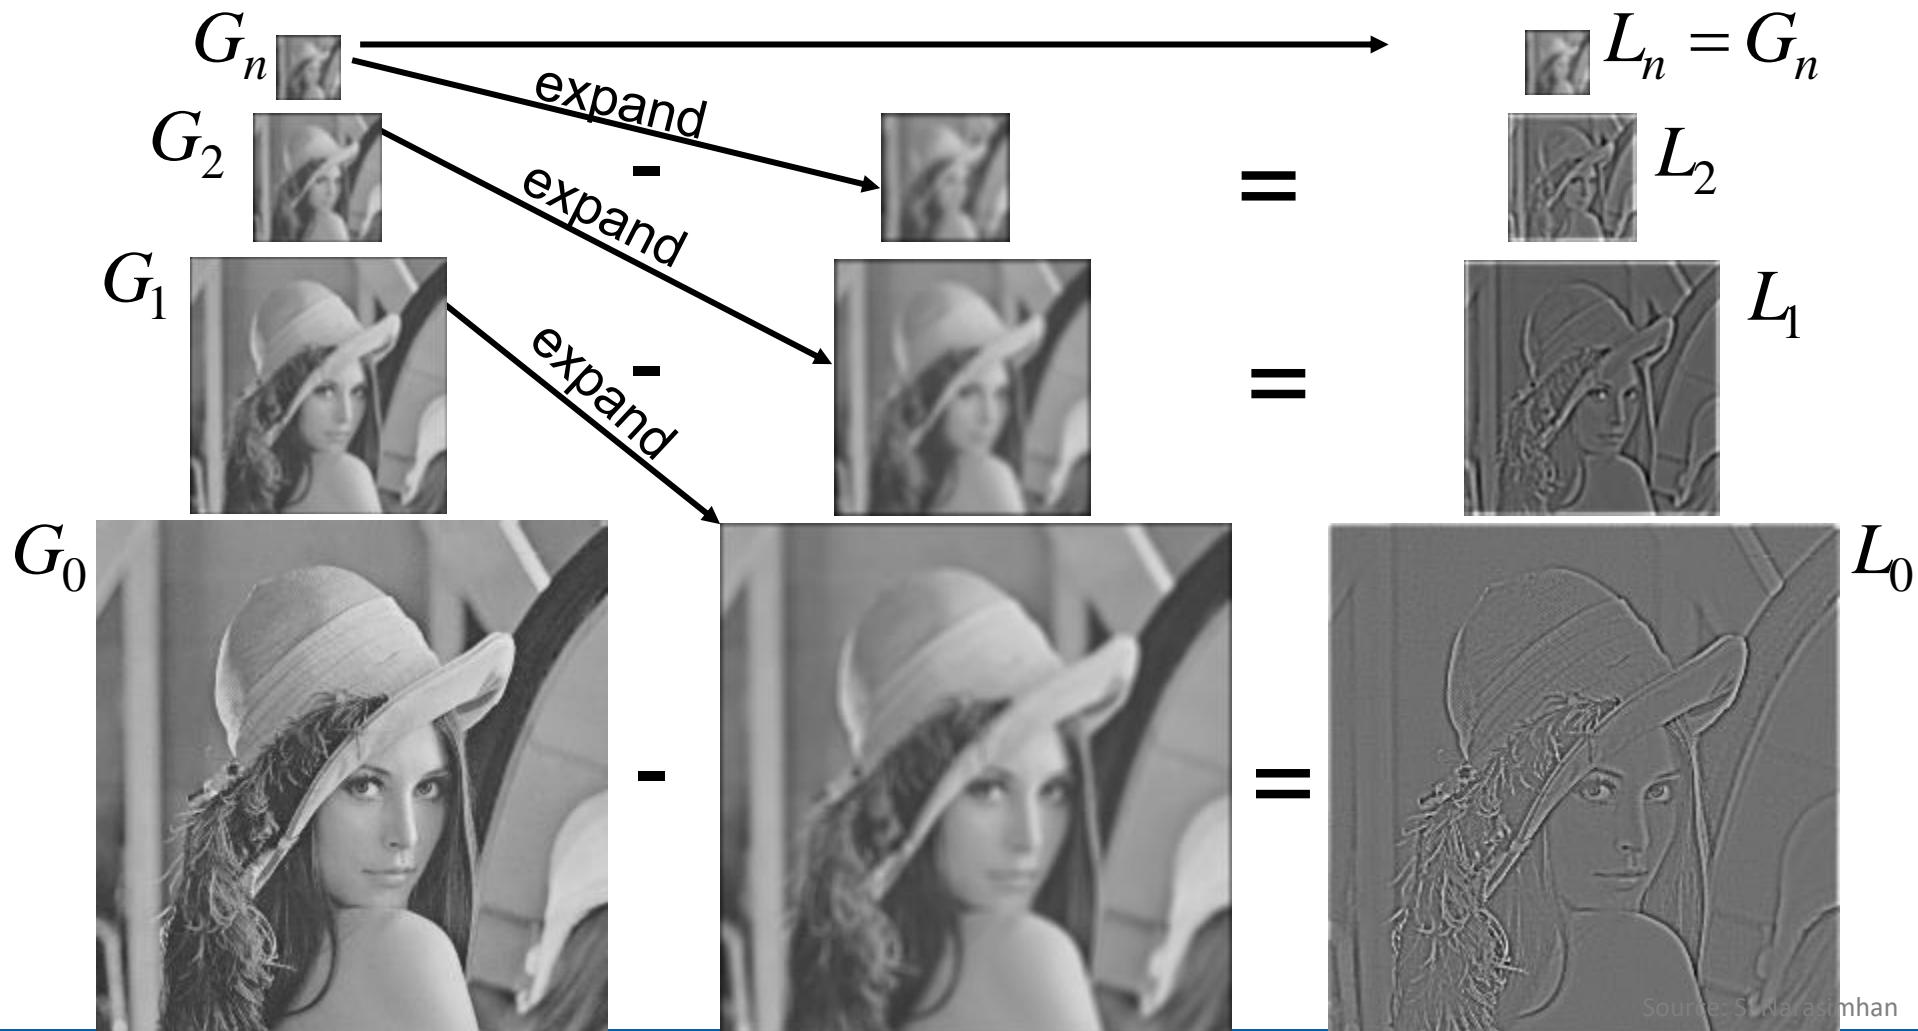


Laplacian pyramid

Gaussian Pyramid

$$L_i = G_i - \text{expand}(G_{i+1})$$

Laplacian Pyramid



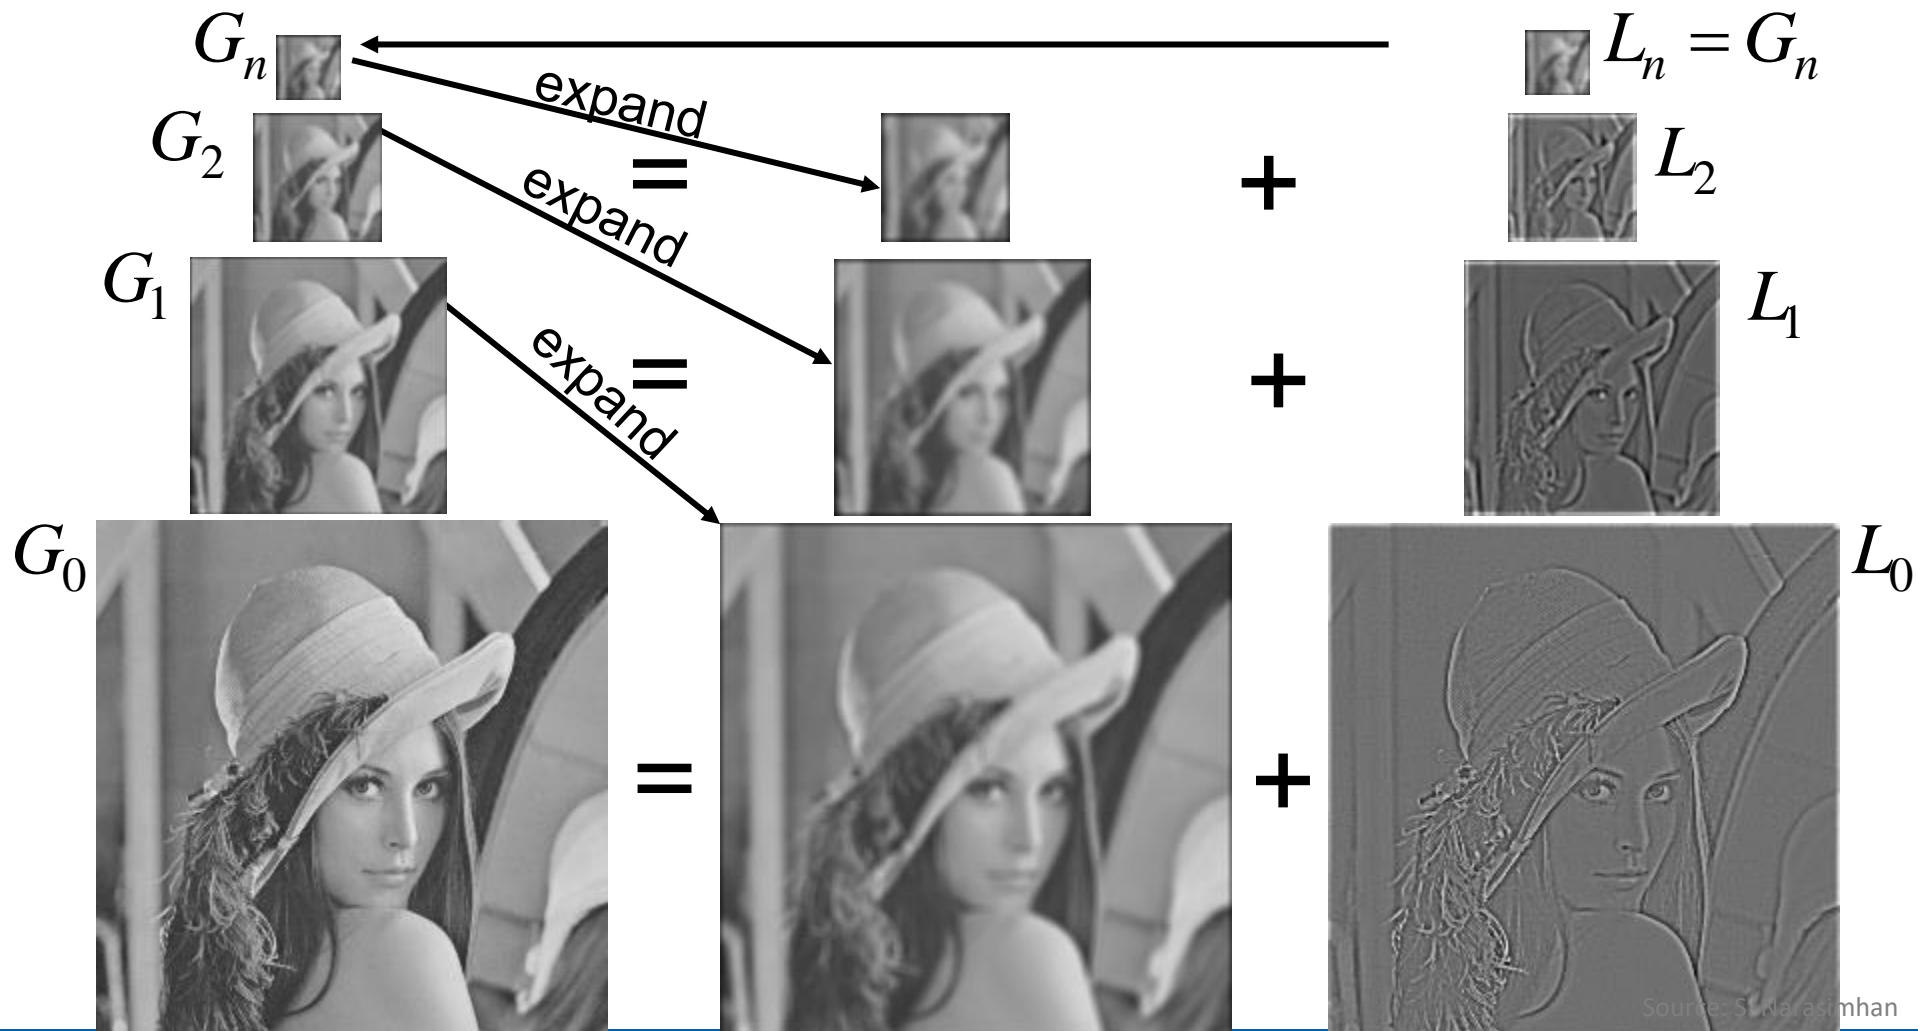
Source: S. Narasimhan

Laplacian pyramid

Gaussian Pyramid

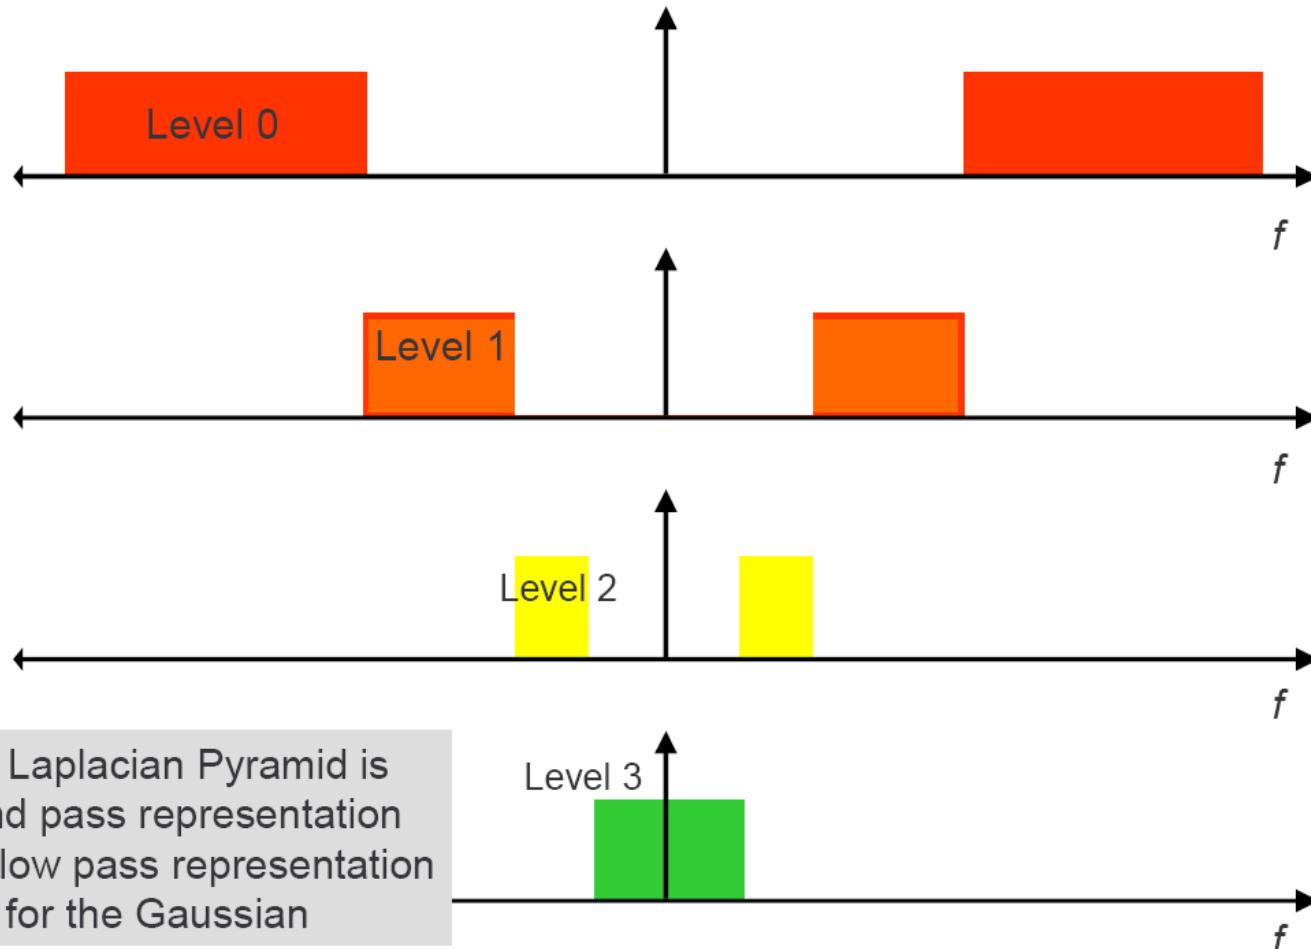
$$G_i = L_i + \text{expand}(G_{i+1})$$

Laplacian Pyramid

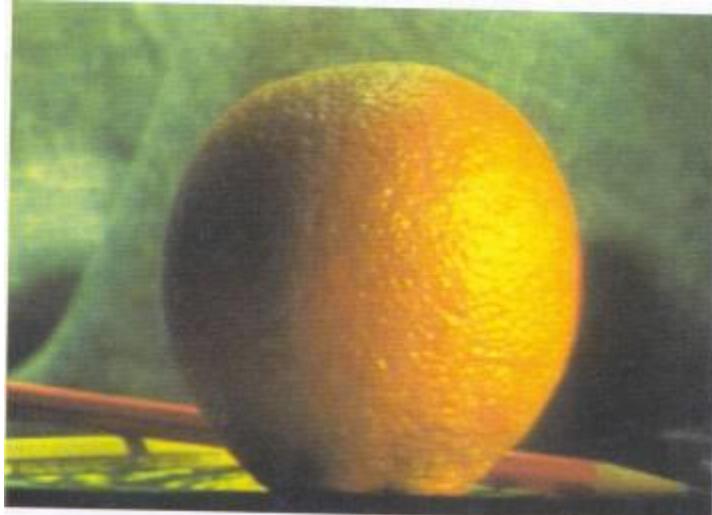
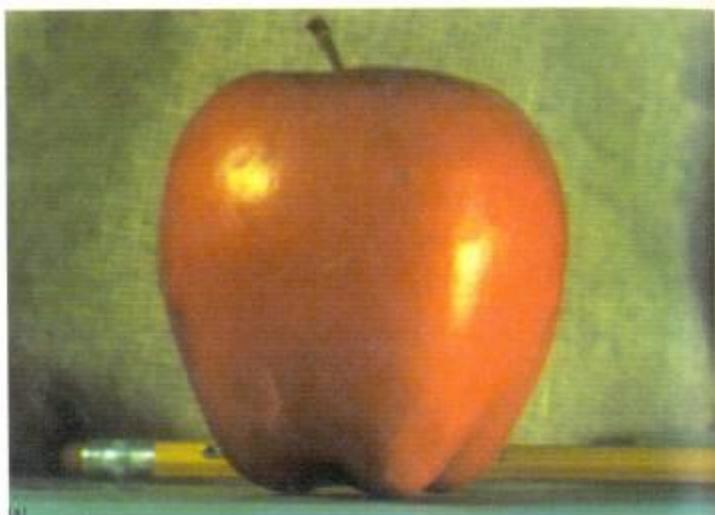


Frequencies

Laplacian Pyramid Frequency Composition



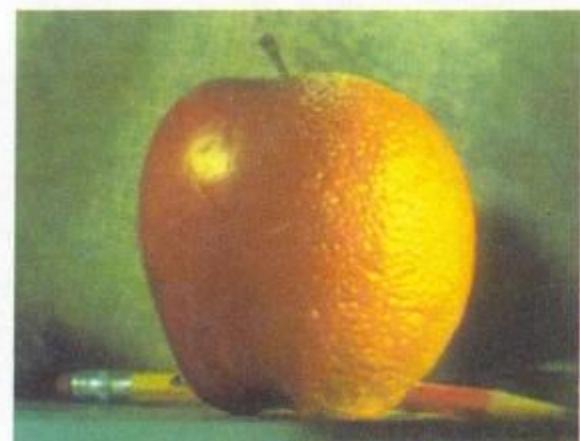
Pyramid Blending



(d)



(h)

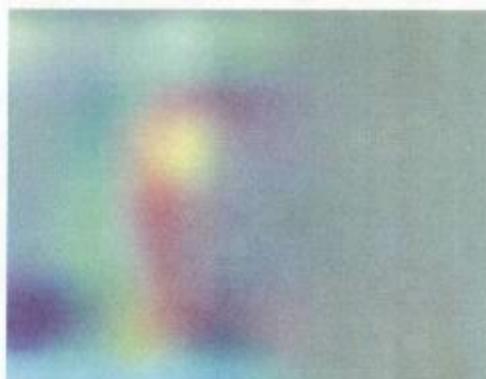


(l)

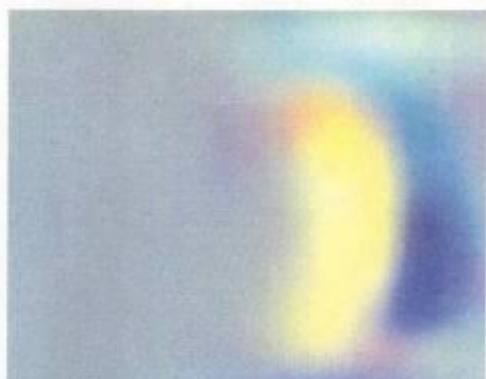
Source: A. Efros

laplacian
level

4



(c)



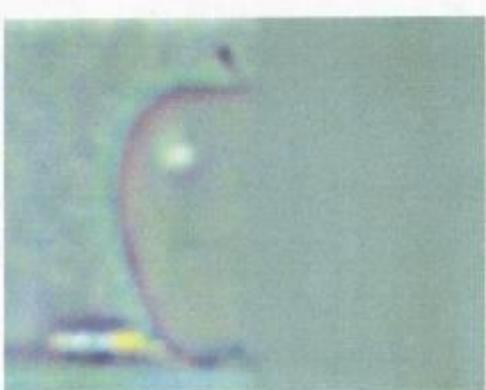
(g)



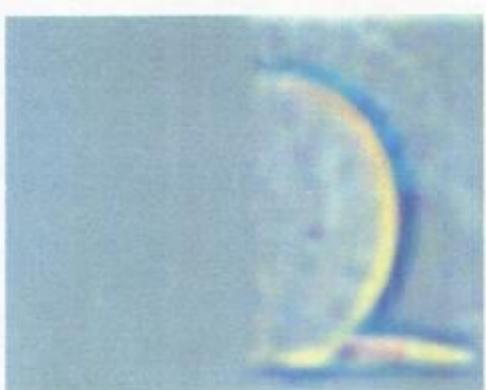
(k)

laplacian
level

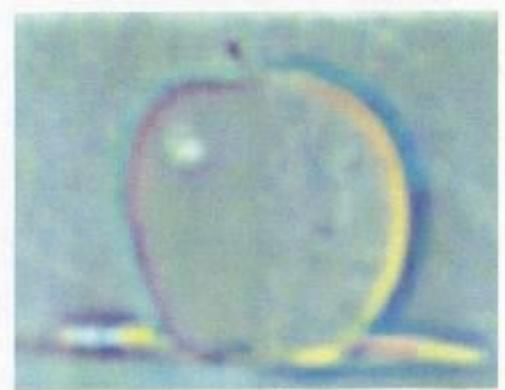
2



(b)



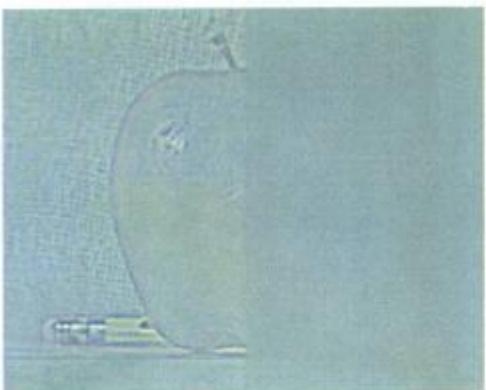
(f)



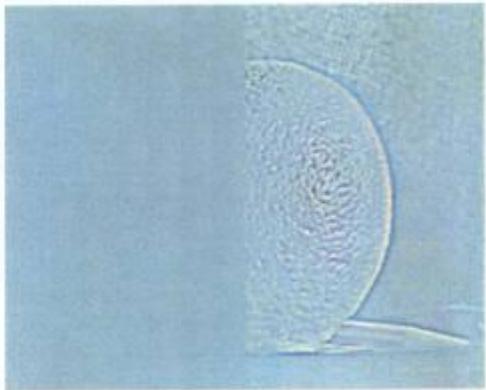
(j)

laplacian
level

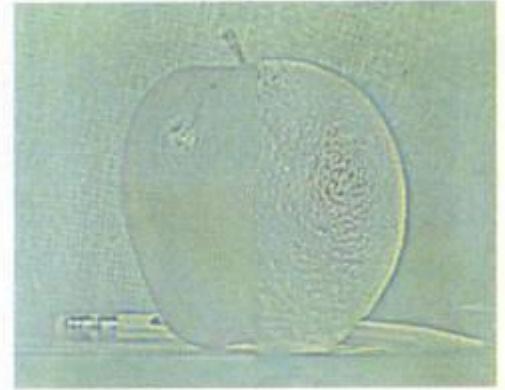
0



(a)



(e)



(i)

left pyramid

right pyramid

blended pyramid

source: A. Efros

Blending Regions



Source: A. Efros

Laplacian Pyramid: Blending

General Approach:

1. Build Laplacian pyramids LA and LB from images A and B
2. Build a Gaussian pyramid GR from selected region R
3. Form a combined pyramid LS from LA and LB using GR as weights:
 - $LS_l(i,j) = GR_l(i,j) * LA_l(i,j) + (1 - GR_l(i,j)) * LB_l(i,j)$
4. Collapse the LS pyramid to get the final blended image ($LS_0 = LS_0 + \text{expand}(LS_{l+1})$)

[P. Burt and E. Adelson. A multiresolution spline with application to image mosaics. ACM Trans. Graph. 1983]

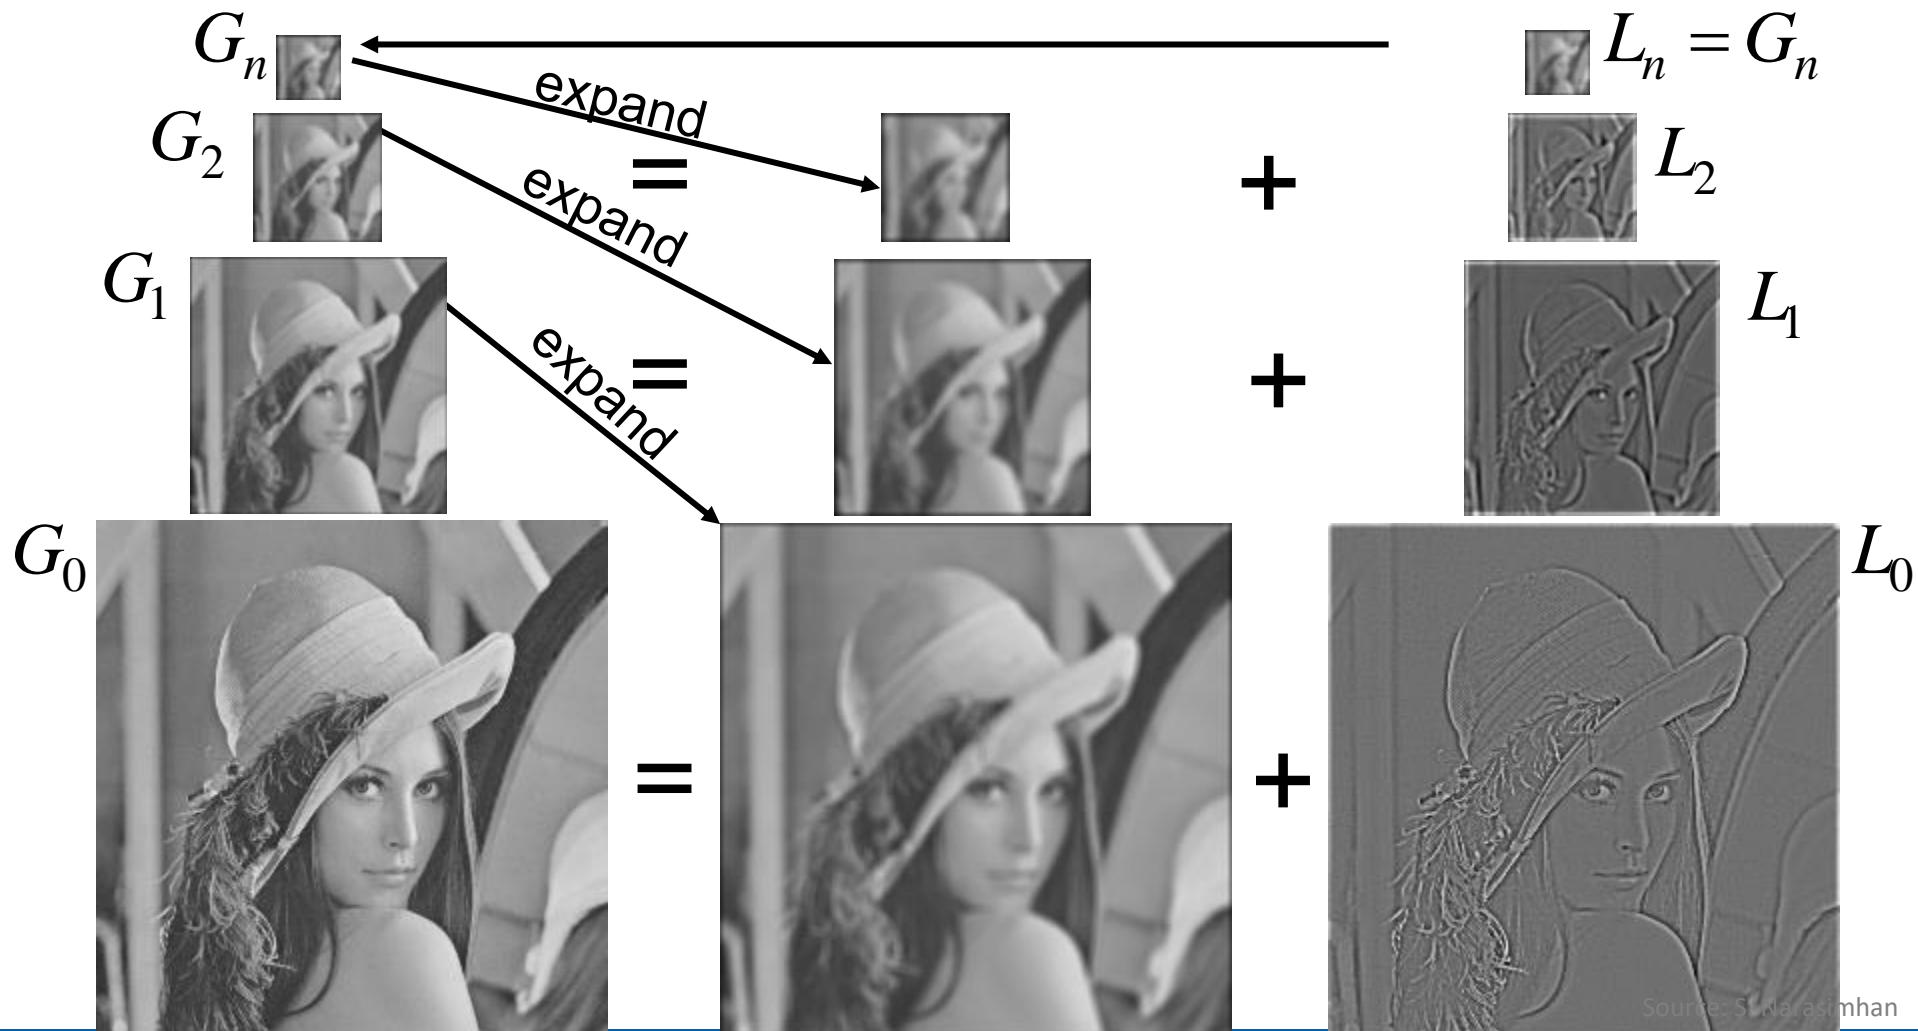
Source: A. Efros

Laplacian pyramid

Gaussian Pyramid

$$G_i = L_i + \text{expand}(G_{i+1})$$

Laplacian Pyramid



Source: S. Narasimhan

Blending Regions



Source: A. Efros

UNIVERSITÄT BONN

