

The logo of the University of Bonn, featuring a blue square with a white curved line and a grey square.


UNIVERSITÄT **BONN**

Juergen Gall

Edges and Corners  
MA-INF 2201 - Computer Vision  
WS24/25

# Origin of edges


Edges are caused by a variety of factors:




# Characterizing edges

- An edge is a place of rapid change in the image intensity function


image



intensity function  
(along horizontal scanline)



first derivative



edges correspond to  
extrema of derivative

# Derivatives with convolution

For 2D function  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$


To implement above as convolution, what would be the associated filter?

# Partial derivatives of an image




$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

-1
1



# Finite difference filters

Other approximations of derivative filters exist:

**Prewitt:**  $M_x =$ 

-1	0	1
-1	0	1
-1	0	1

 $;$   $M_y =$ 

1	1	1
0	0	0
-1	-1	-1

**Sobel:**  $M_x =$ 

-1	0	1
-2	0	2
-1	0	1

 $;$   $M_y =$ 

1	2	1
0	0	0
-1	-2	-1

**Roberts:**  $M_x =$ 




0	1
-1	0

 $;$   $M_y =$ 

1	0
0	-1

# Image gradient

The gradient of an image:  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$





- The gradient points in the direction of most rapid increase in intensity
- The gradient direction is given by
$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$
- The edge strength is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2}$$

# Effects of noise

Consider a single row or column of the image

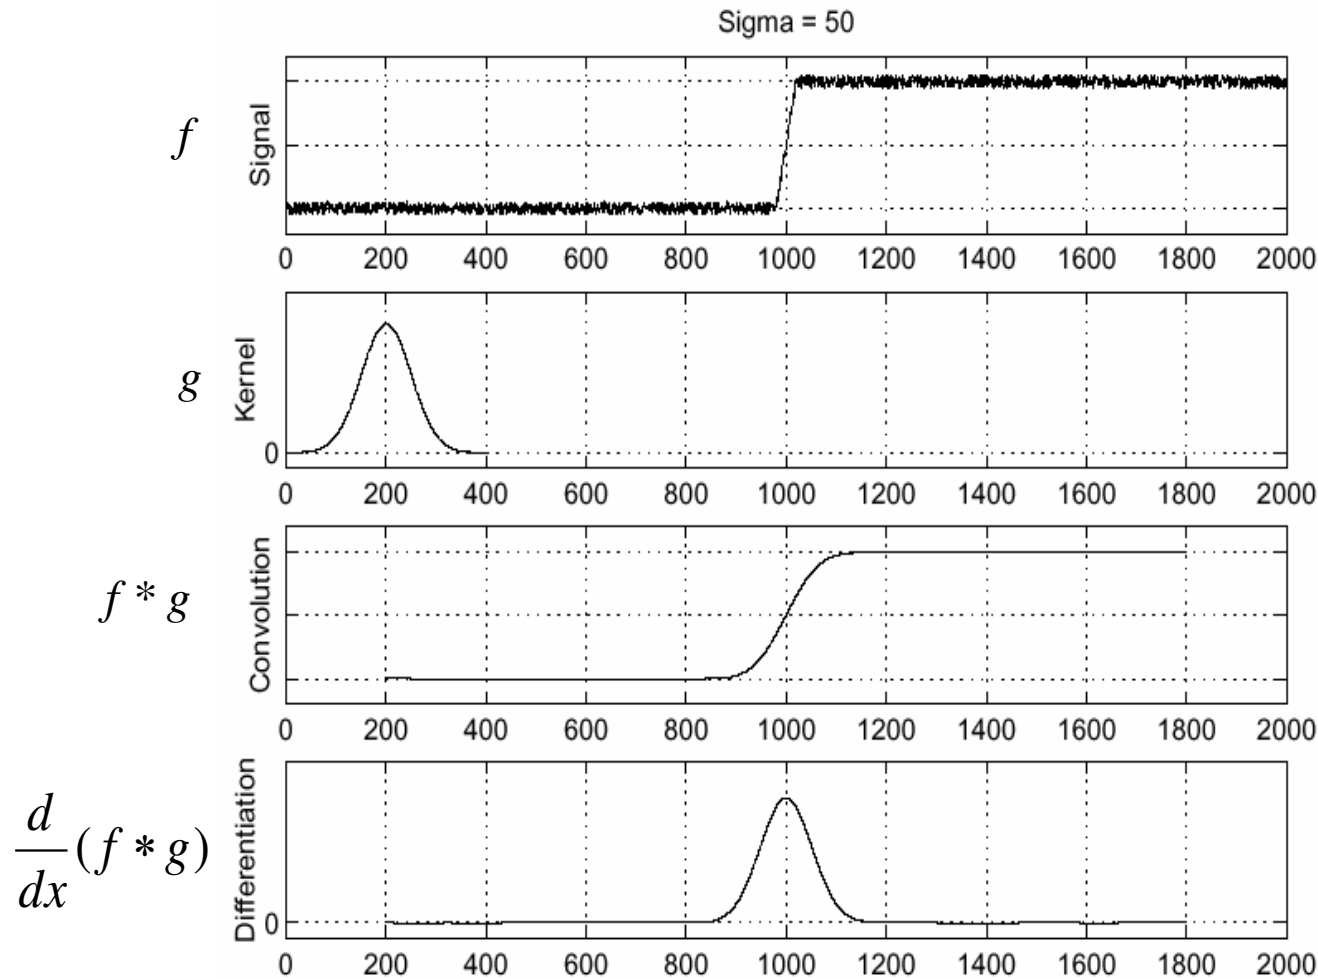
- Plotting intensity as a function of position gives a signal



Where is the edge?



# Solution: smooth first




- To find edges, look for peaks in  $\frac{d}{dx}(f * g)$

# Derivative theorem of convolution

Differentiation is convolution,  
and convolution is associative:


$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

This saves us one operation:





Source: S. Seitz


# Derivative of Gaussian filter



x-direction




y-direction



# Scale of Gaussian derivative filter

Smoothed derivative removes noise, but blurs edge.  
Also finds edges at different “scales”




1 pixel

3 pixels



7 pixels

# Review: Smoothing vs. derivative filters



## Smoothing filters

- Gaussian: remove “high-frequency” components; “low-pass” filter
- Can the values of a smoothing filter be negative?
- What should the values sum to?
  - **One:** constant regions are not affected by the filter



## Derivative filters

- Derivatives of Gaussian
- Can the values of a derivative filter be negative?
- What should the values sum to?
  - **Zero:** no response in constant regions
- High absolute value at points of high contrast

# The Canny edge detector




# The Canny edge detector



Magnitude of gradient

# The Canny edge detector





How to turn these thick regions of the gradient into curves?

Magnitude of gradient



# Non-maximum suppression

- Check if pixel is local maximum along gradient direction,  
select single max across width of the edge
- requires checking interpolated pixels  $p$  and  $r$



# The Canny edge detector



Magnitude of gradient

# The Canny edge detector



Non-maximum suppression

# The Canny edge detector



Problem:  
pixels along  
this edge  
didn't  
survive the  
thresholding

Threshold

# The Canny edge detector




Problem:  
Too much  
noise

Lower threshold



# Hysteresis thresholding

Use a high threshold to start edge curves, and a low threshold to continue them.



	$\ \nabla f(x, y)\  \geq t_1$	definitely an edge
$t_0 \geq$	$\ \nabla f(x, y)\  < t_1$	maybe an edge, depends on context
	$\ \nabla f(x, y)\  < t_0$	definitely not an edge

# Hysteresis thresholding



Low threshold



High threshold



Hysteresis threshold

# Canny edge detector


1. Filter image with derivative of Gaussian
2. Find magnitude and orientation of gradient
3. **Non-maximum suppression:**
  - Thin wide “ridges” down to single pixel width
4. **Linking and thresholding (hysteresis):**
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them

J. Canny, ***A Computational Approach To Edge Detection***, IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.



# Relevance of edges

## Application shape matching



# Not all edges are relevant...


Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/resources.html>




image



human segmentation



gradient magnitude



# Is boundary detection solved?


$$\text{Precision} = \frac{tp}{tp + fp}$$

$$\text{Recall} = \frac{tp}{tp + fn}$$


P. Arbelaez et al. **Contour Detection and Hierarchical Image Segmentation**. TPAMI 2011

Saining Xie and Zhuowen Tu **Holistically-Nested Edge Detection**. IJCV 2017

Machine Learning for boundary detection



Want to associate a model with observed features




[Fig from Marszalek & Schmid, 2007]

For example, the model could be a line, a circle, or an arbitrary shape.

# Chamfer distance

- Average distance to nearest feature

$$D_{chamfer}(T, I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$




- $T$ : template shape  $\rightarrow$  a set of points
- $I$ : image to search  $\rightarrow$  a set of points
- $d_i(t)$ : min distance for point  $t$  to some point in  $I$

# Chamfer distance

Average distance to nearest feature

$$D_{chamfer}(T, I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$



**Edge image**

*Can it be  
implemented as  
correlation or  
convolution?*



# Distance transform

*Image features (2D)*


*Distance Transform*

1	0	1	2	3	4	3	2
1	0	1	2	3	3	2	1
1	0	1	2	3	2	1	0
1	0	0	1	2	1	0	1
2	1	1	2	1	0	1	2
3	2	2	2	1	0	1	2
4	3	3	2	1	0	1	2
5	4	4	3	2	1	0	1


**Distance Transform** is a function  $D(\cdot)$  that for each image pixel  $p$  assigns a non-negative number  $D(p)$  corresponding to distance from  $p$  to the nearest feature in the image  $I$

Features could be edge points, foreground points,...


# Distance transform



original



edges



distance transform

Value at  $(x,y)$  tells how far that position is from the nearest edge point (or other binary image structure)



# Distance transform (1D)

Two pass  $O(n)$  algorithm for 1D  $L_1$  norm

1. Initialize: For all  $j$   
 $D[j] \leftarrow 1_{\mathbf{P}}[j]$  // 0 if  $j$  is in  $\mathbf{P}$ , infinity otherwise


P. Felzenszwalb and D. Huttenlocher. **Distance Transforms of Sampled Functions**.  
Cornell Computing and Information Science TR2004-1963

Source: D. Huttenlocher

# Distance transform (1D)

Two pass  $O(n)$  algorithm for 1D  $L_1$  norm

1. Initialize: For all  $j$   
 $D[j] \leftarrow 1_P[j]$  // 0 if  $j$  is in  $\mathbf{P}$ , infinity otherwise
2. Forward: For  $j$  from 1 up to  $n-1$   
 $D[j] \leftarrow \min(D[j], D[j-1]+1)$
3. Backward: For  $j$  from  $n-2$  down to 0  
 $D[j] \leftarrow \min(D[j], D[j+1]+1)$



$\infty$	0	$\infty$	0	$\infty$	$\infty$	$\infty$	0	$\infty$
$\infty$	0	1	0	1	2	3	0	1
1	0	1	0	1	2	1	0	1

P. Felzenszwalb and D. Huttenlocher. **Distance Transforms of Sampled Functions**.  
 Cornell Computing and Information Science TR2004-1963

Source: D. Huttenlocher


# Distance Transform (2D)

- 2D case analogous to 1D
  - Initialization
  - Forward and backward pass
    - Fwd pass finds closest above and to left
    - Bwd pass finds closest below and to right

-	1
1	0

0	1
1	-



$\infty$	$\infty$	$\infty$	$\infty$
$\infty$	0	$\infty$	$\infty$
$\infty$	0	$\infty$	$\infty$
$\infty$	$\infty$	$\infty$	$\infty$

$\infty$	$\infty$	$\infty$	$\infty$
$\infty$	0	1	$\infty$
$\infty$	0	$\infty$	$\infty$
$\infty$	$\infty$	$\infty$	$\infty$

$\infty$	$\infty$	$\infty$	$\infty$
$\infty$	0	1	2
$\infty$	0	1	2
$\infty$	1	2	3

2	1	2	3
1	0	1	2
1	0	1	2
2	1	2	3

# Distances

Minkowski distance:

$$d_p(x, y) = \left( \sum_i |x_i - y_i|^p \right)^{1/p}$$

$$d_\infty(x, y) = \lim_{p \rightarrow \infty} \left( \sum_i |x_i - y_i|^p \right)^{1/p} = \max_i (|x_i - y_i|)$$

Special cases: L1-norm, L2-norm (Euclidean)

# Approximate Euclidean

+c5	+c4	+c3	+c4	+c5
+c4	+c2	+c1	+c2	+c4
+c3	+c1	0	+c1	+c3
+c4	+c2	+c1	+c2	+c4
+c5	+c4	+c3	+c4	+c5

Parallel mask

Forward:

for  $i = (\text{size} + 1)/2, \dots, \text{lines}$  do

for  $j = (\text{size} + 1)/2, \dots, \text{columns}$  do

$$v_{i,j} = \underset{(k,l) \in \text{forward mask}}{\text{minimum}} (v_{i+k,j+l} + c(k,l))$$

Backward:

for  $i = \text{lines} - (\text{size} - 1)/2, \dots, 1$  do

for  $j = \text{columns} - (\text{size} - 1)/2, \dots, 1$  do

$$v_{i,j} = \underset{(k,l) \in \text{backward mask}}{\text{minimum}} (v_{i+k,j+l} + c(k,l))$$

+c5	+c4	+c3	+c4	+c5
+c4	+c2	+c1	+c2	+c4
+c3	+c1	0		



		0	+c1	+c3
+c4	+c2	+c1	+c2	+c4
+c5	+c4	+c3	+c4	+c5



Forward mask

Backward mask

Chamfer 3-4

$$3 \times \sqrt{2} \approx 4$$

$$3 \times 1 = 3$$

+4	+3	+4
+3	0	+3
+4	+3	+4

	+11		+11	
+11	+7	+5	+7	+11
	+5	0	+5	
+11	+7	+5	+7	+11
	+11		+11	

Chamfer 5-7-11


$$5 \times \sqrt{5} \approx 11$$

$$5 \times \sqrt{2} \approx 7$$


$$5 \times 1 = 5$$

G. Borgefors. **Distance transformations in digital images**. Computer Vision, Graphics, and Image Processing, 34:344–371, 1986.


# Approximate Euclidean




City block




Chessboard




City block




Chessboard




Octagonal




Chamfer 3-4




Octagonal




Chamfer 3-4




Chamfer 5-7-11



Euclidean



Chamfer 5-7-11




Euclidean

# Chamfer distance

Average distance to nearest feature

$$D_{chamfer}(T, I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$




Edge image



Distance transform image

# Chamfer distance



Edge image




Distance transform image



# Line fitting

- Many objects characterized by presence of straight lines



Lines can be detected by the Hough Transform

# Hough space





image space





Hough (parameter) space

Connection between image  $(x,y)$  and Hough  $(m,b)$  spaces:

- A line in the image corresponds to a point in Hough space



# Hough space



Connection between image  $(x,y)$  and Hough  $(m,b)$  spaces:

- A line in the image corresponds to a point in Hough space
- What does a point  $(x_0, y_0)$  in the image space map to?
- The solutions of  $b = -x_0m + y_0$  (line in Hough space)

# Hough space



- What are the line parameters for the line that contains both  $(x_0, y_0)$  and  $(x_1, y_1)$ ?
- It is the intersection of the lines  $b = -x_0m + y_0$  and  $b = -x_1m + y_1$

# Hough algorithm





image space




Hough (parameter) space

**How can we use this to find the most likely parameters (m,b) for the most prominent line in the image space?**

- Let each edge point in image space *vote* for a set of possible parameters in Hough space
- Accumulate votes in discrete set of bins; parameters with the most votes indicate line in image space.

# Polar representation for lines

Issues with usual  $(m,b)$  parameter space: can take on infinite values, undefined for vertical lines.



$d$  : perpendicular distance from line to origin

$\theta$  : angle the perpendicular makes with the x-axis

$$x \cos \theta - y \sin \theta = d$$

Point in image space  $\rightarrow$  sinusoid segment in Hough space

# Hough transform algorithm


Using the polar parameterization:


$$x \cos \theta - y \sin \theta = d$$

Basic Hough transform algorithm


1. Initialize  $H[d, \theta] = 0$
2. for each edge point  $I[x, y]$  in the image
  - for  $\theta = 0$  to  $180$  // some quantization
  - $d = x \cos \theta - y \sin \theta$
  - $H[d, \theta] += 1$
3. Find the value(s) of  $(d, \theta)$  where  $H[d, \theta]$  is maximum
4. The detected line in the image is given by  $d = x \cos \theta - y \sin \theta$

H: accumulator array (votes)





**Image space  
edge coordinates**




**Votes**



Bright value = high vote count  
Black = no votes



Square :



# Lines




# Hough transform for circles

- Circle: center  $(a,b)$  and radius  $r$

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- For a fixed radius  $r$ , unknown gradient direction




# Hough transform for circles

- Circle: center  $(a,b)$  and radius  $r$

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$


- For a fixed radius  $r$ , unknown gradient direction



# Hough transform for circles

- Circle: center  $(a,b)$  and radius  $r$   

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$
- For an unknown radius  $r$ , unknown gradient direction



# Hough transform for circles

- Circle: center  $(a,b)$  and radius  $r$   

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$
- For an unknown radius  $r$ , unknown gradient direction






Image space




Hough space

# Hough transform for circles

- Circle: center  $(a,b)$  and radius  $r$   

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$
- For an unknown radius  $r$ , **known** gradient direction



# Hough transform for circles

For every edge pixel  $(x,y)$  :

For each possible radius value  $r$ :

For each possible gradient direction  $\theta$ :  
*// or use estimated gradient*


$$a = x - r \cos(\theta)$$

$$b = y - r \sin(\theta)$$

$$H[a,b,r] += 1$$


end

end







Original



Edges




Votes: Penny




Note: a different Hough transform (with separate accumulators) was used for each circle radius (quarters vs. penny).


Original



Edges



Votes: Quarter




Combined detections


# Parameters for analytic curves

Analytic Form	Parameters	Equation
Line	$d, \theta$	$x\cos\theta + y\sin\theta = d$
Circle	$x_0, y_0, r$	$(x-x_0)^2 + (y-y_0)^2 = r^2$
Parabola	$x_0, y_0, \rho, \theta$	$(y-y_0)^2 = 4\rho(x-x_0)$
Ellipse	$x_0, y_0, a, b, \theta$	$(x-x_0)^2/a^2 + (y-y_0)^2/b^2 = 1$

# Impact of noise on Hough



**Image space  
edge coordinates**




**Votes**


What difficulty does this present for an implementation?

Source: K. Grauman

# Impact of noise on Hough



**Image space  
edge coordinates**



**Votes**

Here, everything appears to be “noise”, or random edge points, but we still see peaks in the vote space.

Source: K. Grauman

# Voting: practical tips

Minimize irrelevant tokens first (take edge points with significant gradient magnitude)

Choose a good grid / discretization


- Too coarse: large votes obtained when too many different lines correspond to a single bucket
- Too fine: miss lines because some points that are not exactly collinear cast votes for different buckets

Vote for neighbors also (smoothing in accumulator array) or use mean shift

Utilize direction of edge to reduce free parameters by 1


To read back which points voted for “winning” peaks, keep tags on the votes.

# Mean Shift



**Objective : Find the densest region**  
 Distribution of identical billiard balls

# Mean Shift



**Objective : Find the densest region**  
 Distribution of identical billiard balls



# Mean Shift

Region of  
interest

Center of  
mass

Mean Shift  
vector

**Objective : Find the densest region**  
Distribution of identical billiard balls

# Mean Shift

Region of  
interest

Center of  
mass

Mean Shift  
vector

**Objective : Find the densest region**  
Distribution of identical billiard balls

# Mean Shift

Region of  
interest

Center of  
mass

Mean Shift  
vector

**Objective : Find the densest region**  
Distribution of identical billiard balls

Source: Y. Ukrainitz and B. Sarel

# Mean Shift

Region of  
interest

Center of  
mass


Mean Shift  
vector

**Objective : Find the densest region**  
Distribution of identical billiard balls

# Mean Shift

Region of  
interest

Center of  
mass




**Objective : Find the densest region**  
Distribution of identical billiard balls

# Kernel density estimation

$$P(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n K(\mathbf{x} - \mathbf{x}_i)$$

A function of some finite number of data points  
 $\mathbf{x}_1 \dots \mathbf{x}_n$



Data

In practice one uses the forms:

$$K(\mathbf{x}) = c \prod_{i=1}^d k(x_i) \quad \text{or} \quad K(\mathbf{x}) = ck(\|\mathbf{x}\|)$$

Same function on each dimension

Function of vector length only

$$P(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n K(\mathbf{x} - \mathbf{x}_i)$$





A function of some finite number of data points  
 $\mathbf{x}_1 \dots \mathbf{x}_n$

## Examples:

- Epanechnikov Kernel  $K_E(\mathbf{x}) = \begin{cases} c(1 - \|\mathbf{x}\|^2) & \|\mathbf{x}\| \leq 1 \\ 0 & \text{otherwise} \end{cases}$

- Uniform Kernel  $K_U(\mathbf{x}) = \begin{cases} c & \|\mathbf{x}\| \leq 1 \\ 0 & \text{otherwise} \end{cases}$

- Normal Kernel  $K_N(\mathbf{x}) = c \cdot \exp\left(-\frac{1}{2}\|\mathbf{x}\|^2\right)$



Source: Y. Ukrainitz and B. Sarel

# Gradient ascent

$$\nabla P(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \nabla K(\mathbf{x} - \mathbf{x}_i)$$

Using the  
Kernel form:

$$K(\mathbf{x} - \mathbf{x}_i) = ck \left( \left\| \frac{\mathbf{x} - \mathbf{x}_i}{h} \right\|^2 \right)$$

Size of window



# Gradient ascent

$$\begin{aligned}\nabla P(x) &= \frac{c}{N} \sum_i \nabla k \left( \|x - x_i\|^2 \right) \\&= \frac{c}{N} \sum_i k' \left( \|x - x_i\|^2 \right) 2(x - x_i) \\&= \frac{2c}{N} \sum_i g \left( \|x - x_i\|^2 \right) (x_i - x) && \text{Use: } g(x) = -k'(x) \\&= \frac{2c}{N} \left( \left\{ \sum_i x_i g \left( \|x - x_i\|^2 \right) \right\} - \left\{ \sum_i g \left( \|x - x_i\|^2 \right) \right\} x \right) \\&= \frac{2c}{N} \sum_i g \left( \|x - x_i\|^2 \right) \left( \frac{\sum_i x_i g \left( \|x - x_i\|^2 \right)}{\sum_i g \left( \|x - x_i\|^2 \right)} - x \right)\end{aligned}$$

# Gradient ascent

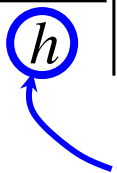
$$\nabla P(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \nabla K(\mathbf{x} - \mathbf{x}_i)$$

Using the  
Kernel form:

$$K(\mathbf{x} - \mathbf{x}_i) = ck \left( \left\| \frac{\mathbf{x} - \mathbf{x}_i}{h} \right\|^2 \right)$$

We get :

Size of window



$$\nabla P(\mathbf{x}) = \frac{2c}{N} \sum_i g(\|x - x_i\|^2) \left( \frac{\sum_i x_i g(\|x - x_i\|^2)}{\sum_i g(\|x - x_i\|^2)} - x \right)$$

$$g(\mathbf{x}) = -k'(\mathbf{x})$$

# Mean shift

$$\nabla P(\mathbf{x}) = \frac{c}{n} \sum_{i=1}^n \nabla k_i = \frac{2c}{N} \sum_i g\left(\|x - x_i\|^2\right) \left( \frac{\sum_i x_i g\left(\|x - x_i\|^2\right)}{\sum_i g\left(\|x - x_i\|^2\right)} - x \right)$$


Yet another Kernel density estimation !


Simple Mean Shift procedure:

- Compute mean shift vector

$$\mathbf{m}(\mathbf{x}) = \left[ \frac{\sum_{i=1}^n \mathbf{x}_i g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)}{\sum_{i=1}^n g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)} - \mathbf{x} \right]$$

- Translate the Kernel window by  $\mathbf{m}(\mathbf{x})$





## Updated Mean Shift Procedure:

- Find all modes using the Simple Mean Shift Procedure
- Prune modes by perturbing them (find saddle points and plateaus)
- Prune nearby – take highest mode in the window

# Hough transform: pros and cons

## Pros

All points are processed independently, so can cope with occlusion

Some robustness to noise: noise points unlikely to contribute consistently to any single bin

Can detect multiple instances of a model in a single pass

## Cons


Complexity of search time increases exponentially with the number of model parameters

Non-target shapes can produce spurious peaks in parameter space

Quantization: hard to pick a good grid size

# Generalized Hough transform

What if want to detect arbitrary shapes defined by boundary points and a reference point?



At each boundary point, compute displacement vector:  $\mathbf{r} = \mathbf{a} - \mathbf{p}_i$ .

For a given model shape: store these vectors in a table indexed by gradient orientation  $\theta$ .

[D. Ballard, Generalizing the Hough Transform to Detect Arbitrary Shapes, 1980]

# Generalized Hough transform

To *detect* the model shape in a new image:

For each edge point


- Index into table with its gradient orientation  $\theta$
- Use retrieved  $\mathbf{r}$  vectors to vote for position of reference point

Peak in this Hough space is reference point with most supporting edges

*Assuming translation is the only transformation here, i.e., orientation and scale are fixed.*

# Example

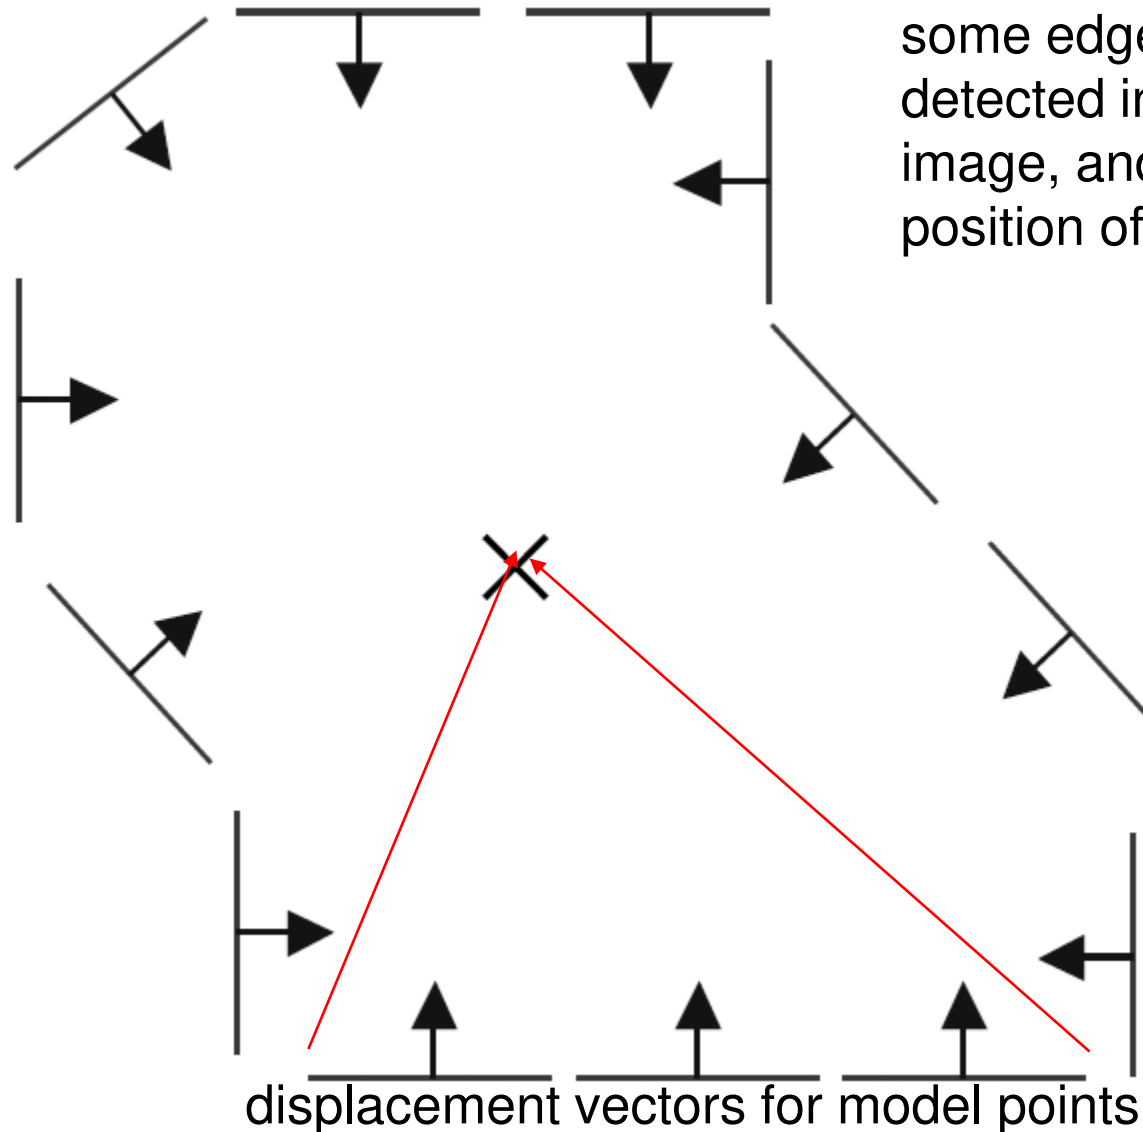
Say we've already stored a table of displacement vectors as a function of edge orientation for this model shape.






# Example


Now we want to look at some edge points detected in a *new* image, and vote on the position of that shape.




# Example




# Example




# Example




# Example



# Example



# Example



# Hough voting for object detection




- Parts of an object provide useful spatial information
- Machine Learning: learn how to vote

Gall et al. **Hough Forests for Object Detection, Tracking, and Action Recognition.** IEEE Transactions on Pattern Analysis and Machine Intelligence, Vol. 33, No. 11, 2188-2202, 2011.



# Hough voting for object detection



Gall et al. **Hough Forests for Object Detection, Tracking, and Action Recognition.** IEEE Transactions on Pattern Analysis and Machine Intelligence, Vol. 33, No. 11, 2188-2202, 2011.

The logo graphic of the University of Bonn, featuring a blue square in the top-left corner and a grey square in the bottom-right corner, separated by a white curved line.

UNIVERSITÄT **BONN**