

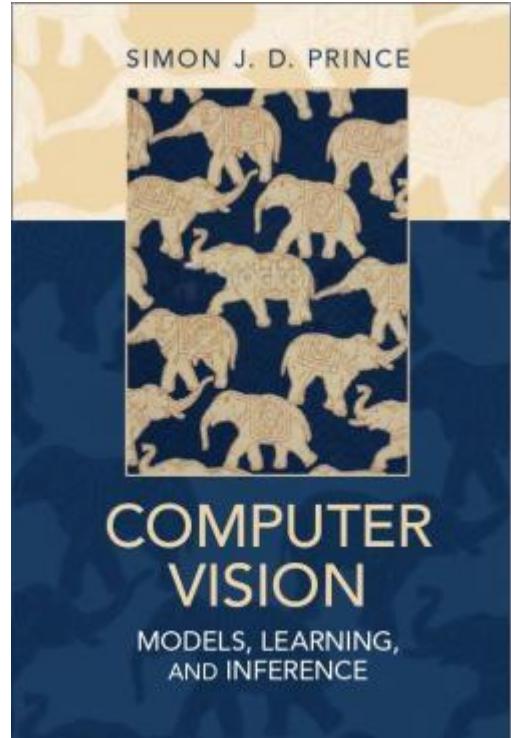


UNIVERSITÄT **BONN**

Juergen Gall

Transformations and Shapes
MA-INF 2201 - Computer Vision
WS24/25

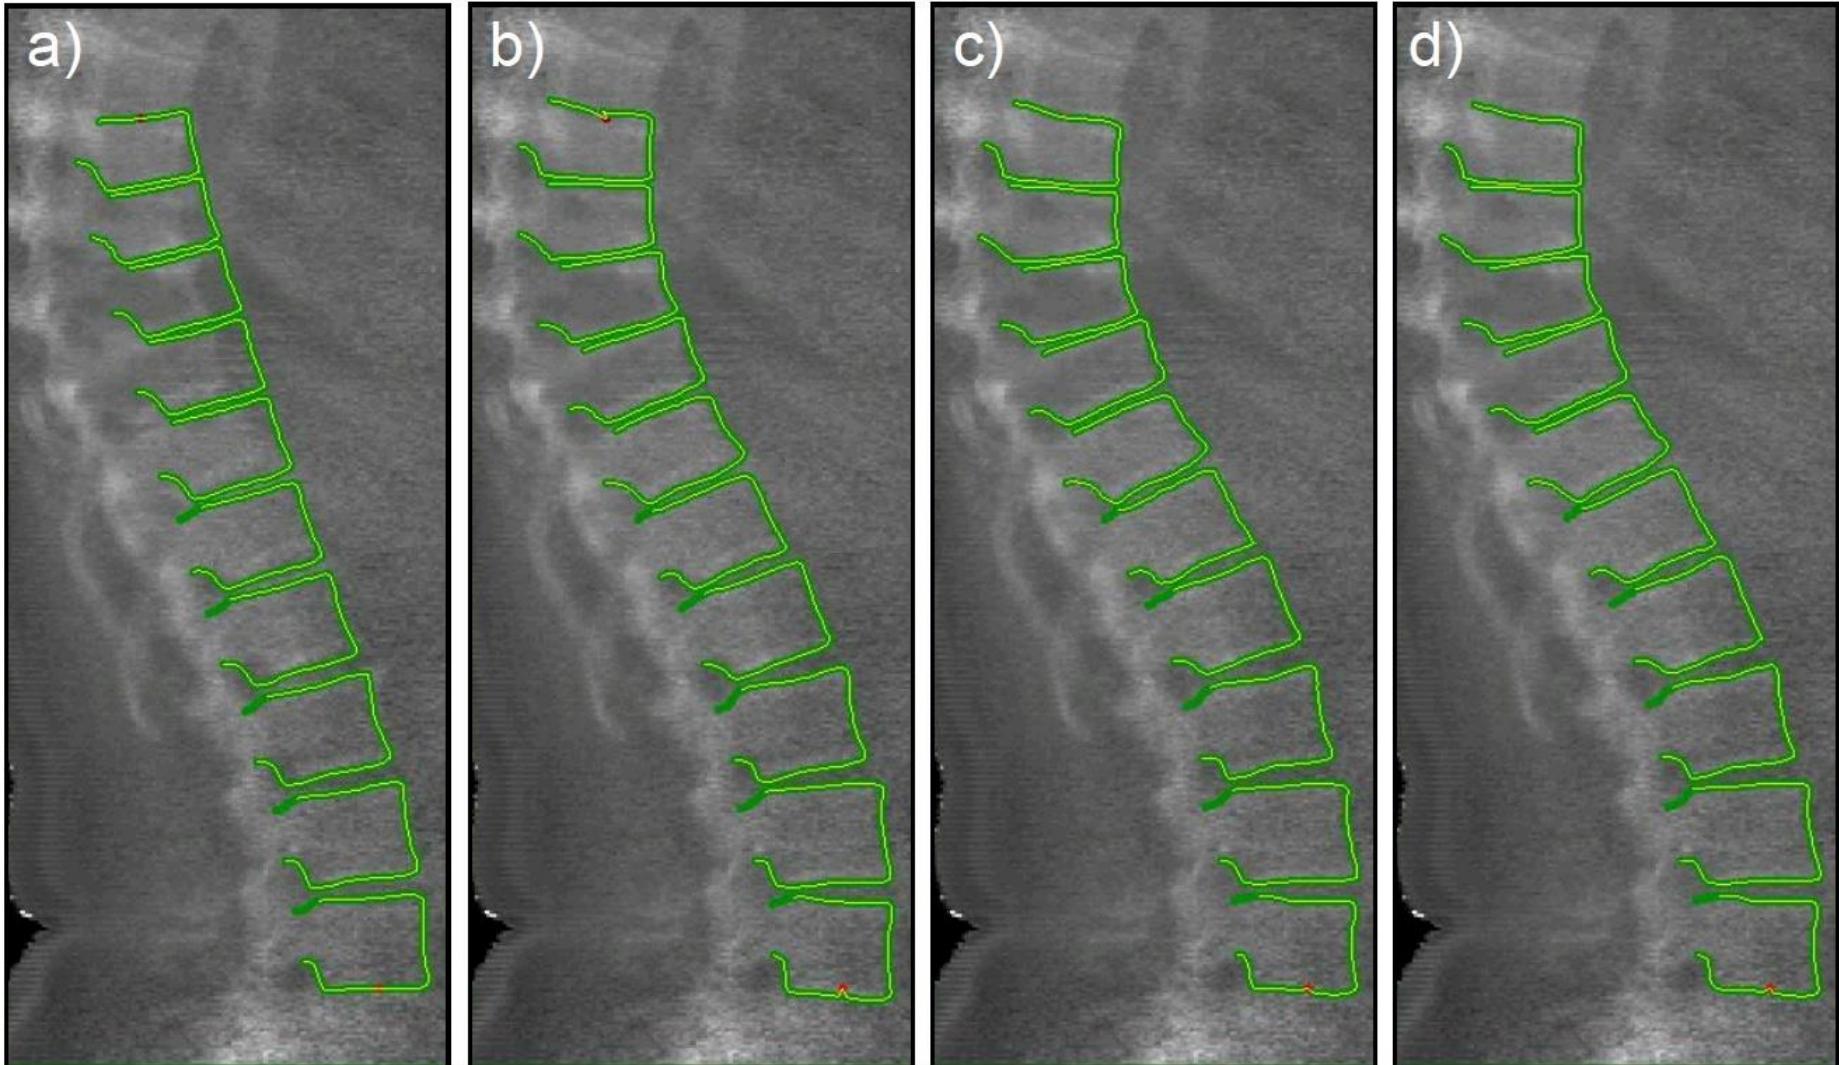
Literature



Chapter 17 Models for shape

S. Prince. **Computer Vision: Models, Learning, and Inference.** Cambridge University Press 2012

Motivation: fitting shape model

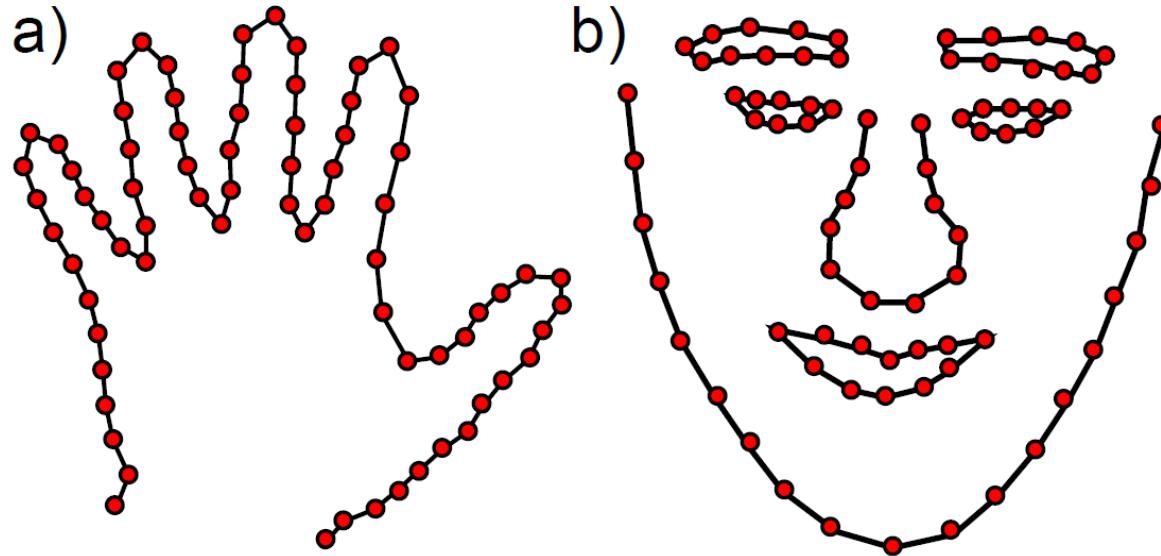


Source: T. Cootes

What is shape?

- Kendall (1984) – Shape “is all the geometrical information that remains when location, scale and rotational effects are filtered out from an object”
- In other words, it is whatever is invariant to a similarity transformation

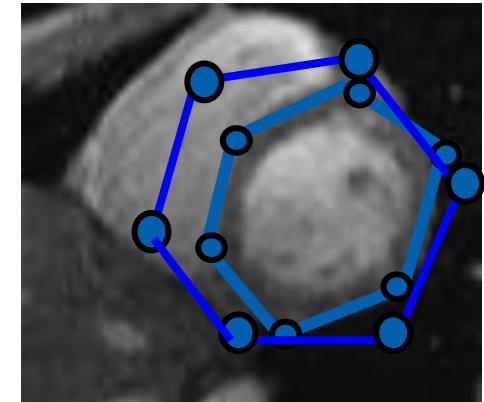
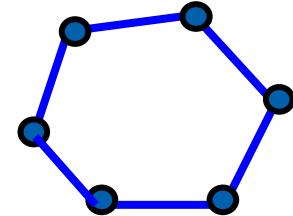
Landmark Points



- Landmark points can be thought of as discrete samples from underlying contour
 - Ordered (single continuous contour)
 - Ordered with wrapping (closed contour)
 - More complex organisation (collection of closed and open)

Recap: Snakes

- A simple elastic snake is defined by:
 - A set of n points,
 - An internal energy term (tension, bending, plus optional shape prior)
 - An external energy term (gradient-based)
- To use to segment an object:
 - Initialize in the vicinity of the object
 - Modify the points to minimize the total energy



Source: K. Grauman

Probability or not?

Snakes: minimize cost function

$$E_{total} = E_{internal} + E_{external}$$

Probabilistic view:

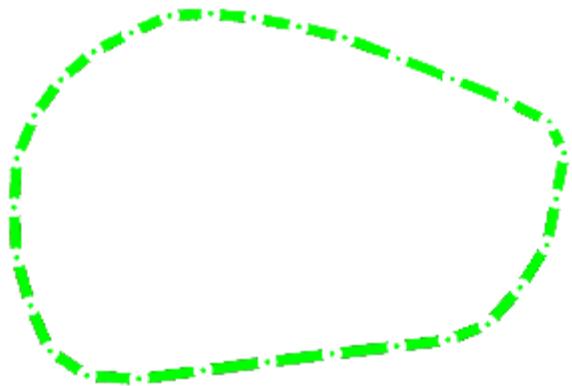
$$\begin{aligned}\hat{\mathbf{W}} = \operatorname{argmax}_{\mathbf{W}} [Pr(\mathbf{W}|\mathbf{x})] &= \operatorname{argmax}_{\mathbf{W}} [Pr(\mathbf{x}|\mathbf{W})Pr(\mathbf{W})] \\ &= \operatorname{argmax}_{\mathbf{W}} [\log[Pr(\mathbf{x}|\mathbf{W})] + \log[Pr(\mathbf{W})]]\end{aligned}\quad \text{likelihood} \quad \text{prior}$$

Weak priors: smoothness etc...

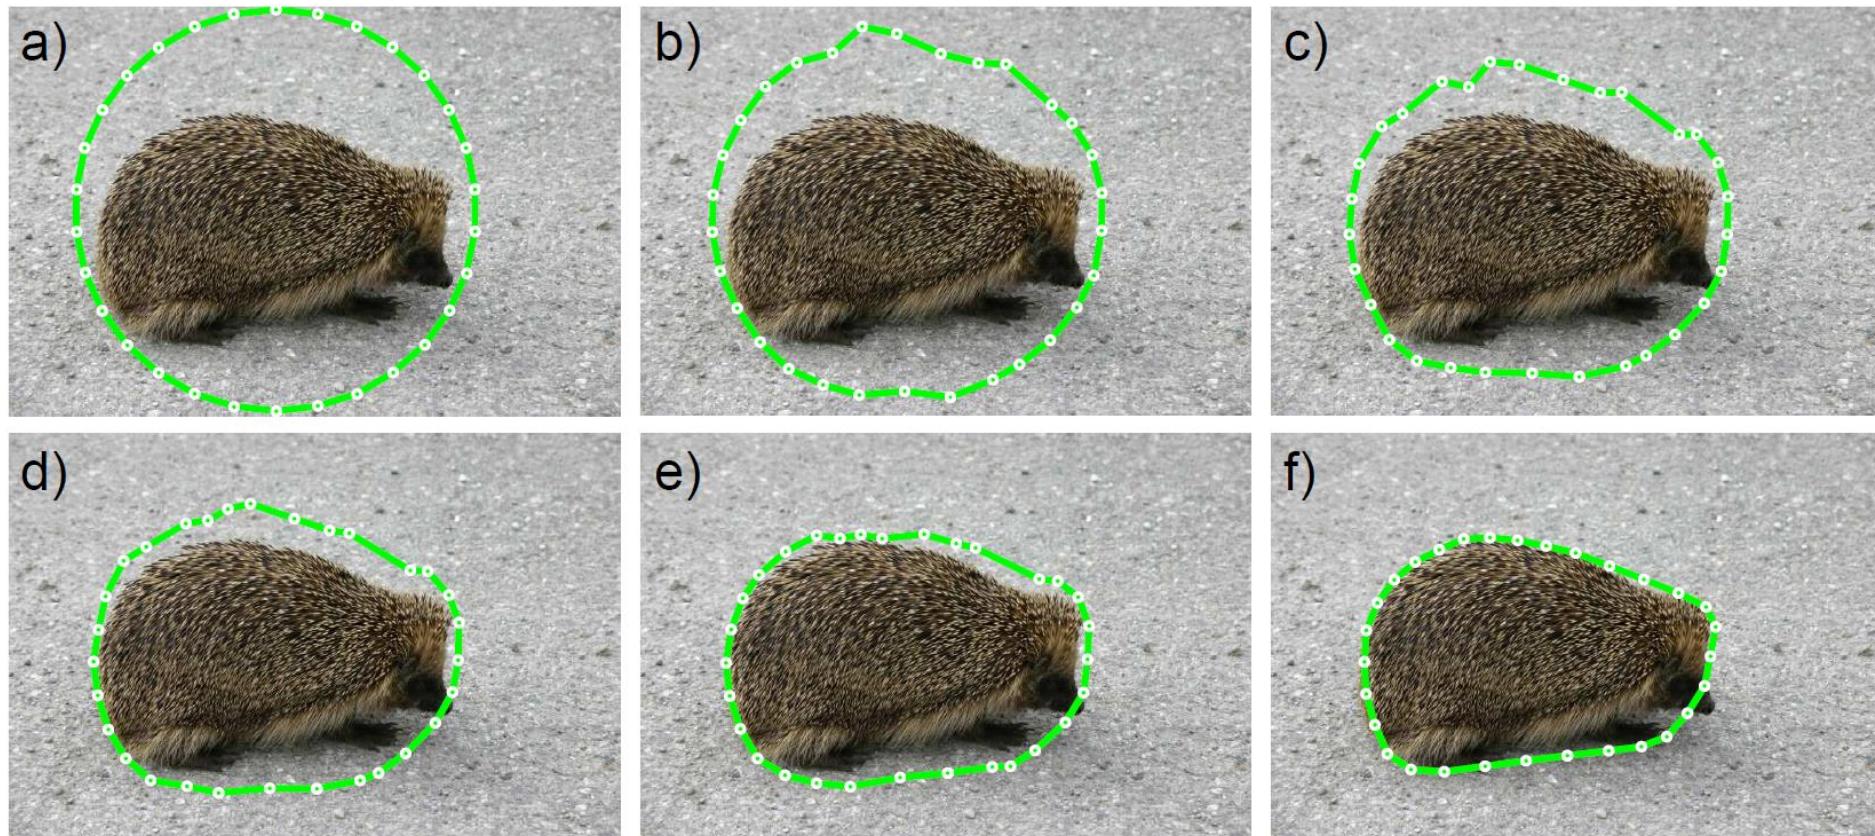
Strong prior: shape

Snakes

What is the name of the animal?

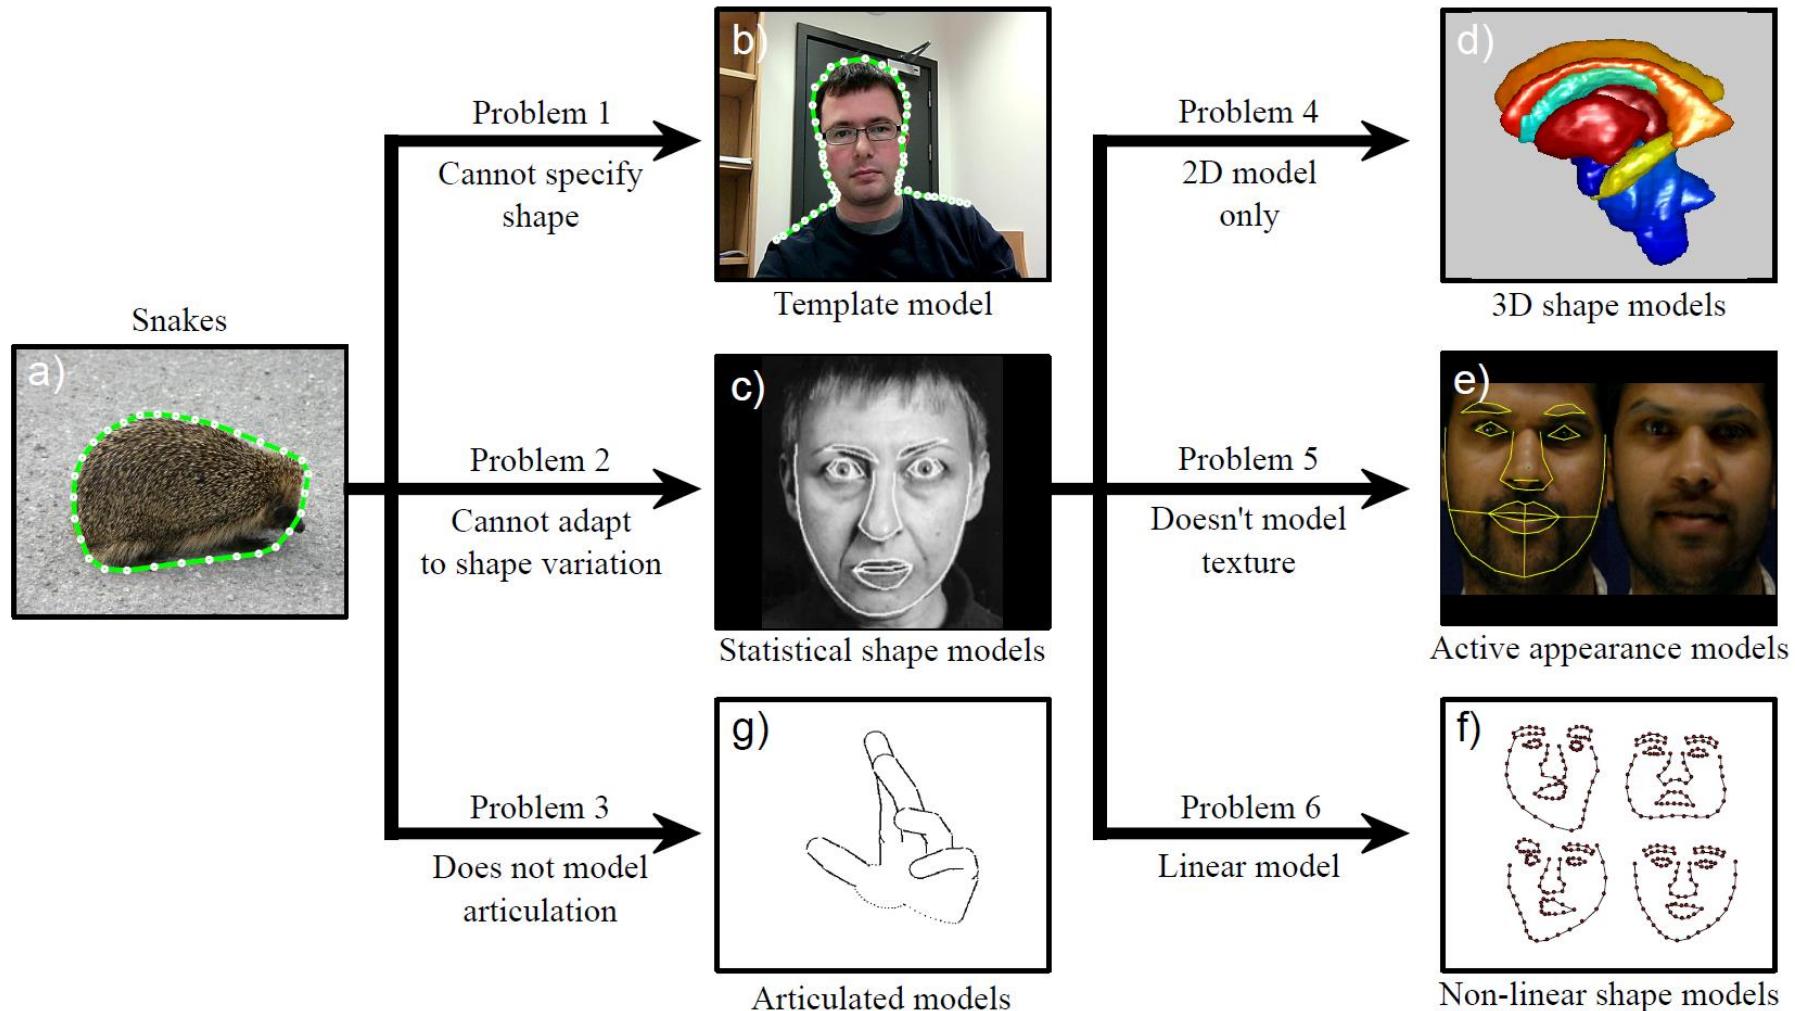


Snakes



Notice failure at nose – falls between points.
Resulting shape does not represent hedgehog

Relationships between models



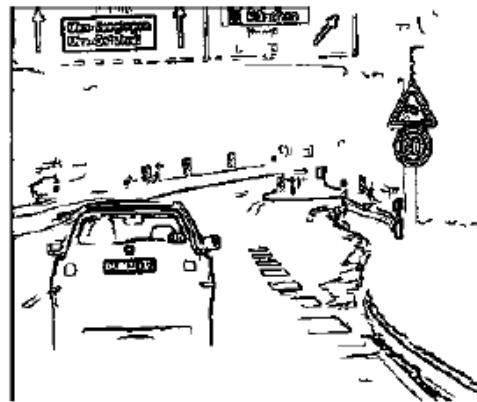
Shape template model

- Shape based on landmark points $\mathbf{W} = \{\mathbf{w}_n\}_{n=1}^N$
- These points are assumed known
- Mapped into the image by transformation $\text{trans}[\mathbf{w}, \Psi]$
- What is left is to find parameters of transformation Ψ
- Likelihood is based on distance transform:

$$Pr(\mathbf{x}|\mathbf{W}, \Psi) \propto \prod_{n=1}^N \exp \left[- (\text{dist} [\mathbf{x}, \text{trans}[\mathbf{w}_n, \Psi]])^2 \right]$$

- No prior on parameters (but could have one)

Recap: Chamfer distance



Edge image



Distance transform image

Inference

- Use maximum likelihood approach

$$\begin{aligned}\hat{\Psi} = \operatorname{argmax}_{\Psi} L &= \operatorname{argmax}_{\Psi} [\log [Pr(\mathbf{x}|\mathbf{W}, \Psi)]] \\ &= \operatorname{argmax}_{\Psi} \left[\sum_{n=1}^N -(\operatorname{dist}[\mathbf{x}, \operatorname{trans}[\mathbf{w}_n, \Psi]])^2 \right]\end{aligned}$$

- No closed form solution
- Must use non-linear optimization or brute-force search
- Use chain rule to compute derivatives

$$\frac{\partial L}{\partial \Psi} = - \sum_{n=1}^N \sum_{j=1}^2 \frac{\partial (\operatorname{dist}[\mathbf{x}, \mathbf{w}'_n])^2}{\partial w'_{jn}} \frac{\partial w'_{jn}}{\partial \Psi} \quad \text{where } \mathbf{w}'_n = \operatorname{trans}[\mathbf{w}_n, \Psi]$$

Derivative

Pre-compute Euclidean distance transform D:



Pre-compute gradient by central difference:

$$\frac{\partial}{\partial w'_{jn}} \frac{1}{2} (\text{dist}[x, w'_n])^2 \approx D(w'_n) G_j(w'_n) = D(w'_n) \frac{1}{2} (D(w'_n + e_j) - D(w'_n - e_j))$$

Interpolate w' from nearby pixels

Linear Transformations

- Transformations Ψ

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a & b \\ c & d \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

- Linear 2D transformations can be represented with a 2x2 matrix and are combinations of ...
 - Scale
 - Rotation,
 - Shear
 - Mirror

Examples

2D Scaling

$$x' = s_x * x$$

$$y' = s_y * y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} s_x & 0 \\ 0 & s_y \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

2D Rotate around (0,0)

$$x' = \cos \Theta * x - \sin \Theta * y$$

$$y' = \sin \Theta * x + \cos \Theta * y$$

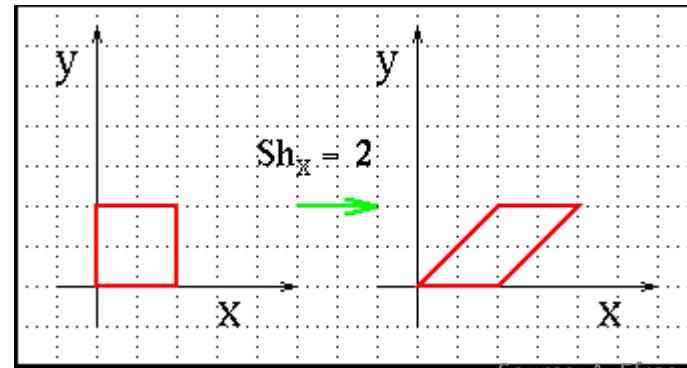
$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta \\ \sin \Theta & \cos \Theta \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

2D Shear

$$x' = x + sh_x * y$$

$$y' = sh_y * x + y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & sh_x \\ sh_y & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$



Source: A. Elros

Examples

2D Mirror about Y axis

$$x' = -x$$

$$y' = y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

2D Mirror over (0,0)

$$x' = -x$$

$$y' = -y$$

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} -1 & 0 \\ 0 & -1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

2D Translation

$$x' = x + t_x$$

$$y' = y + t_y$$

NO!

Homogeneous coordinates

To convert to homogeneous coordinates:

$$(x, y) \Rightarrow \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

homogeneous image
coordinates

Converting *from* homogeneous coordinates

$$\begin{bmatrix} x \\ y \\ w \end{bmatrix} \Rightarrow (x/w, y/w)$$

Homogeneous Coordinates

- Q: How can we represent 2d translation as a 3x3 matrix using homogeneous coordinates?

$$x' = x + t_x$$

$$y' = y + t_y$$

- A: Using the rightmost column:

$$\text{Translation} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix}$$

Basic 2D Transformations

- Basic 2D transformations as 3x3 matrices

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta & 0 \\ \sin \Theta & \cos \Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Rotate

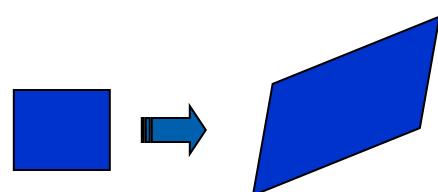
$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Shear

2D Affine Transformations

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- Affine transformations are combinations of
 - Linear transformations, and
 - Translations
- Parallel lines remain parallel



Inference

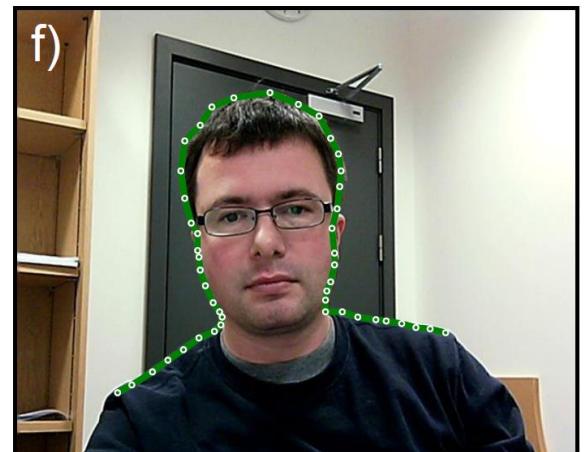
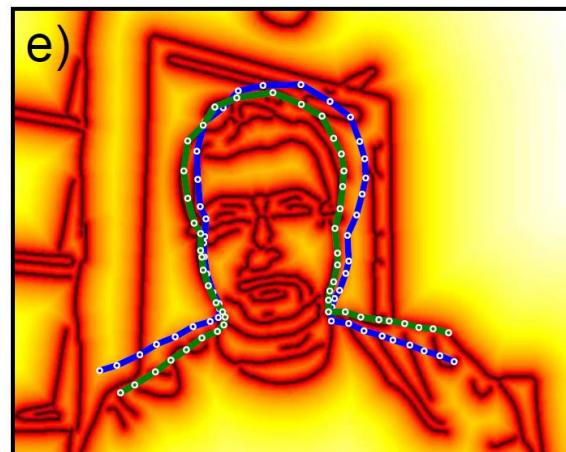
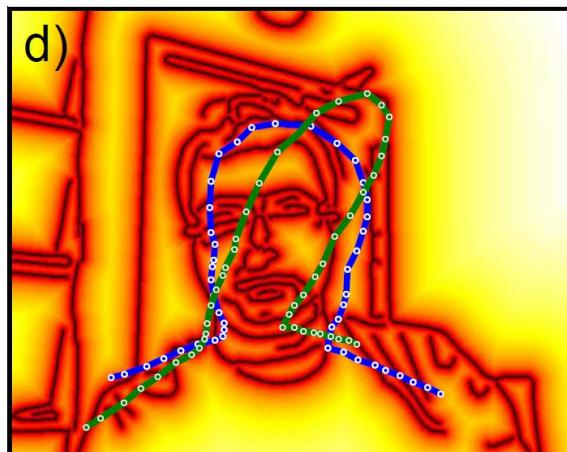
- Use maximum likelihood approach

$$\begin{aligned}\hat{\Psi} = \operatorname{argmax}_{\Psi} L &= \operatorname{argmax}_{\Psi} [\log [Pr(\mathbf{x}|\mathbf{W}, \Psi)]] \\ &= \operatorname{argmax}_{\Psi} \left[\sum_{n=1}^N -(\operatorname{dist}[\mathbf{x}, \operatorname{trans}[\mathbf{w}_n, \Psi]])^2 \right]\end{aligned}$$

- No closed form solution
- Must use non-linear optimization or brute-force search
- Use chain rule to compute derivatives

$$\frac{\partial L}{\partial \Psi} = - \sum_{n=1}^N \sum_{j=1}^2 \frac{\partial (\operatorname{dist}[\mathbf{x}, \mathbf{w}'_n])^2}{\partial w'_{jn}} \frac{\partial w'_{jn}}{\partial \Psi} \quad \text{where } \mathbf{w}'_n = \operatorname{trans}[\mathbf{w}_n, \Psi]$$

Shape template model



Local optimum

Global optimum

Blue initialization, green estimate

Source: S. Prince

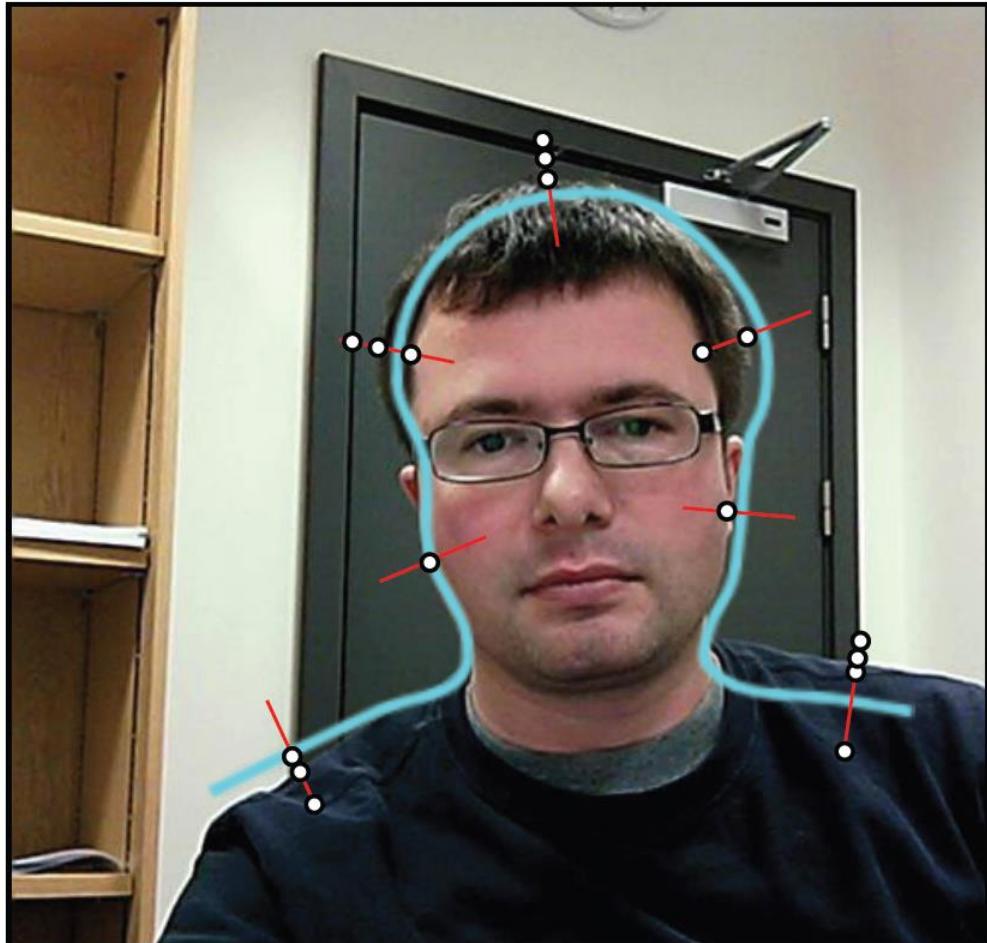
Iterative closest points

- Maximum likelihood approach

$$\frac{\partial L}{\partial \Psi} = - \sum_{n=1}^N \sum_{j=1}^2 \frac{\partial(\text{dist}[\mathbf{x}, \mathbf{w}'_n])^2}{\partial w'_{jn}} \frac{\partial w'_{jn}}{\partial \Psi} \quad \text{where } \mathbf{w}'_n = \text{trans}[\mathbf{w}_n, \Psi]$$

- Iterative closest point:
 - Find closest edge point or line search along normal
 - Estimate Ψ from correspondences (Closed form solution exist)

Iterative closest points



- Find nearest edge point to each landmark point
- Compute transformation in closed form
- Repeat

Correspondences

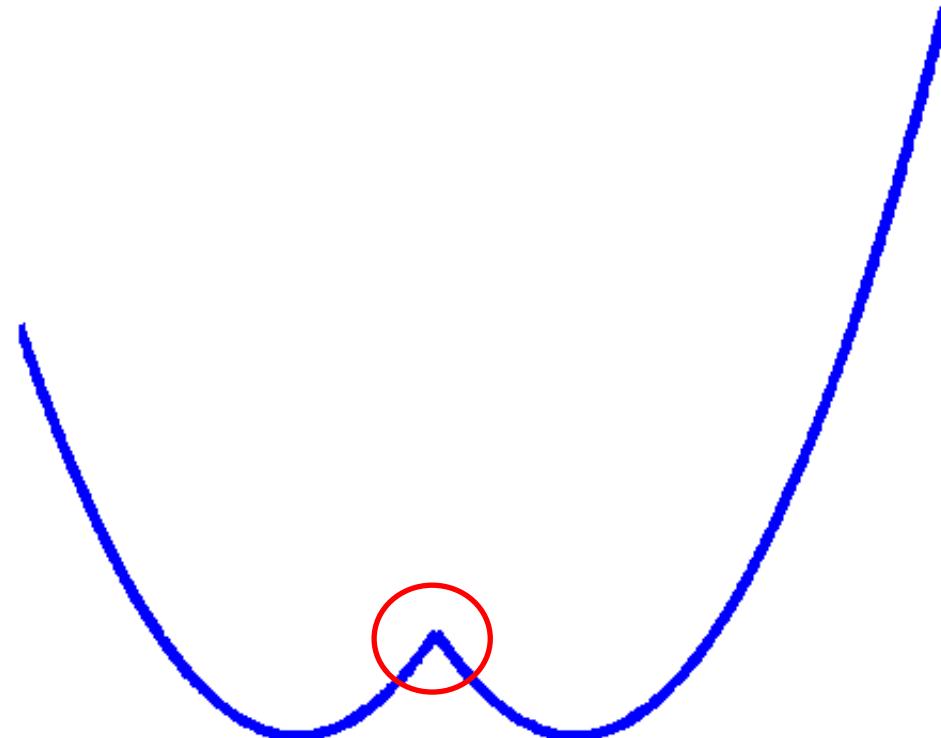
- Line search
- Closest point based on distance transform D:

$$x_n = f(w'_n) = w'_n - \frac{D(w'_n)}{\sqrt{G_x(w'_n)^2 + G_y(w'_n)^2}}(G_x(w'_n)G_y(w'_n))$$

- 1D example:
 - w: 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14
 - E: 0 0 0 0 1 0 0 0 0 0 1 0 0 0 0
 - D: 4 3 2 1 0 1 2 3 2 1 0 1 2 3 4
 - G: -1 -1 -1 -1 0 1 1 0 -1 -1 0 1 1 1 1
 - x: 4 4 4 4 4 4 4 7 10 10 10 10 10 10 10

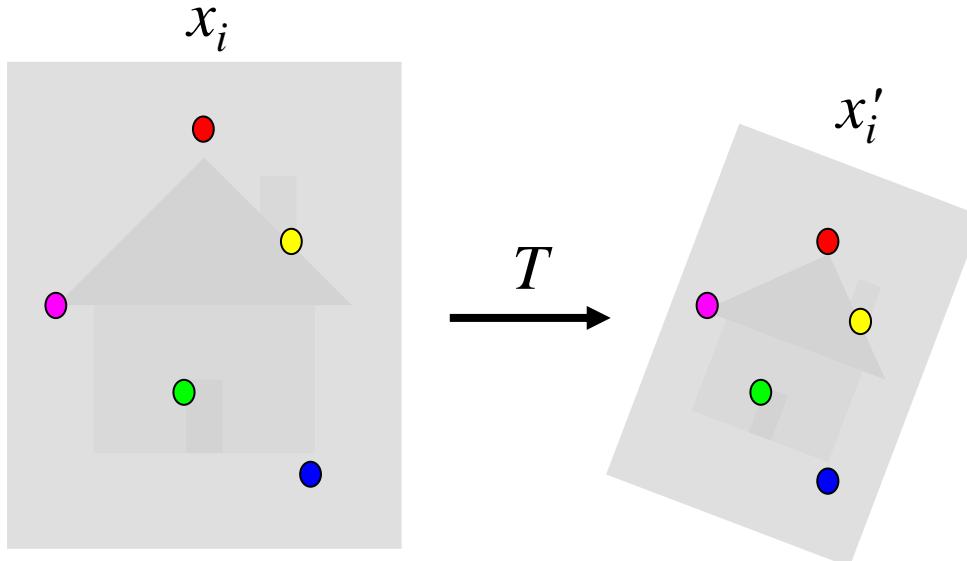
Derivative

Function is not everywhere differentiable

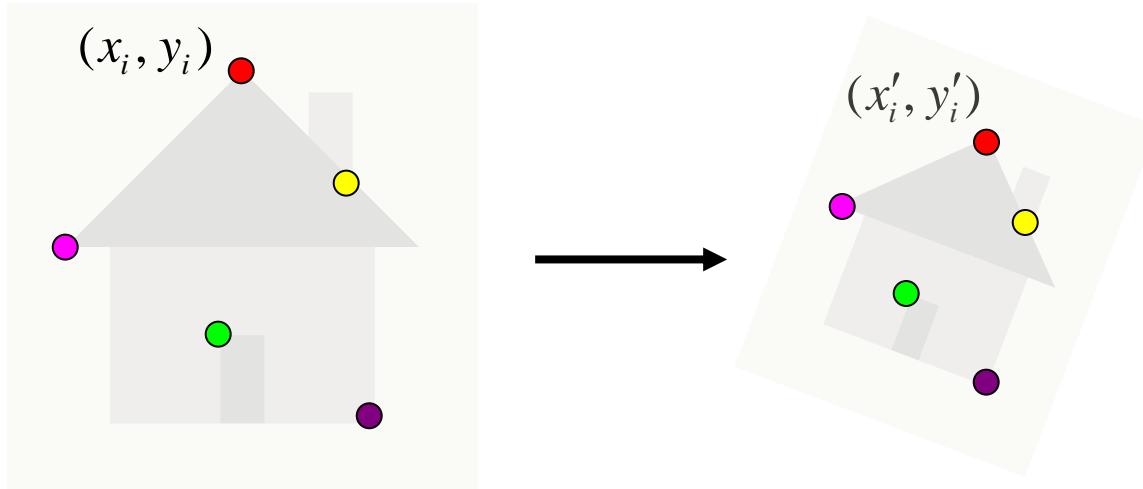


Alignment problem

Fit the parameters of some **transformation** according to a set of matching feature pairs (“correspondences”).



Fitting an affine transformation



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \end{bmatrix}$$

Source: K. Grauman

Least squares problem

A $m \times n$ matrix of rank n :

$$Ax = b$$

Solution with pseudo-inverse:

$$A^T(Ax - b) = 0$$

$$x = (A^T A)^{-1} A^T b$$

Pseudo-inverse \mathbf{A}^+

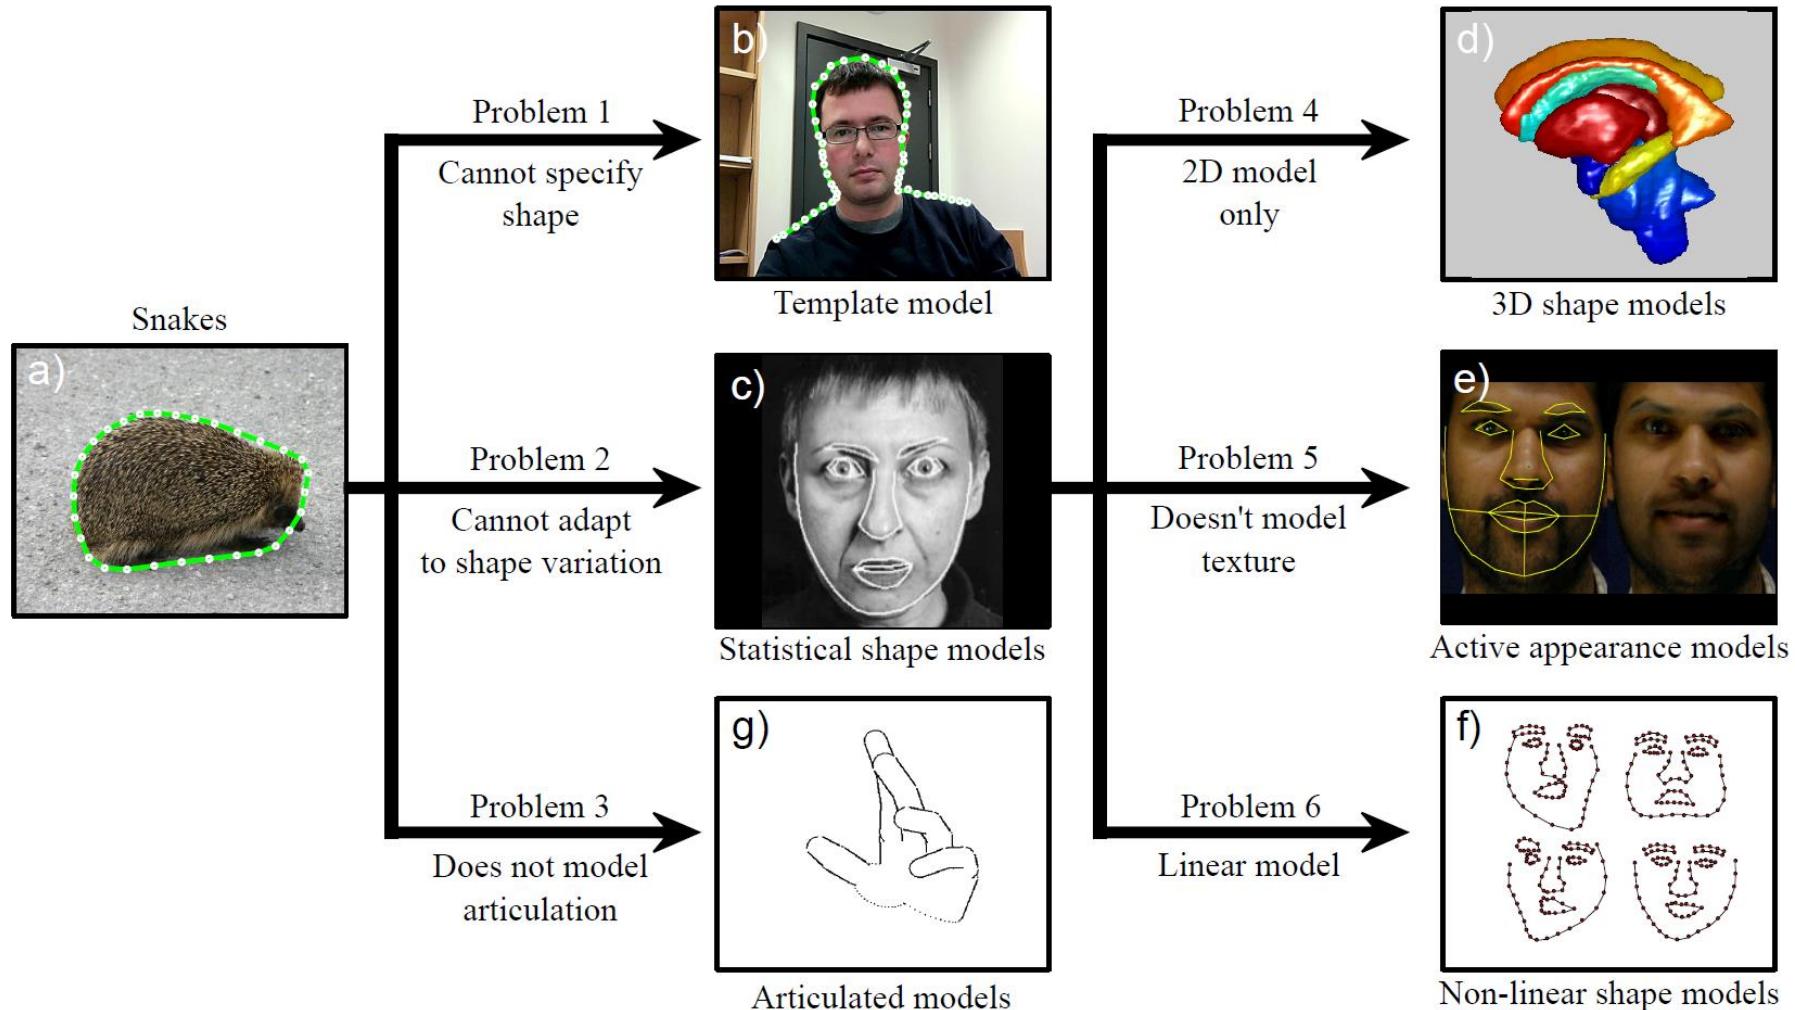
Solution with SVD:

- SVD: $\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{V}^T$; $\mathbf{A}^+ = \mathbf{V}\mathbf{D}^{-1}\mathbf{U}^T$
- $\mathbf{b}' = \mathbf{U}^T\mathbf{b}$
- $y_i = b'_i / D_{ii}$
- $\mathbf{x} = \mathbf{V}\mathbf{y}$

There are different ways to solve the problem. The optimal method depends on size and sparseness of \mathbf{A}

Specific C++ linear algebra libraries exist!

Relationships between models



Statistical shape models

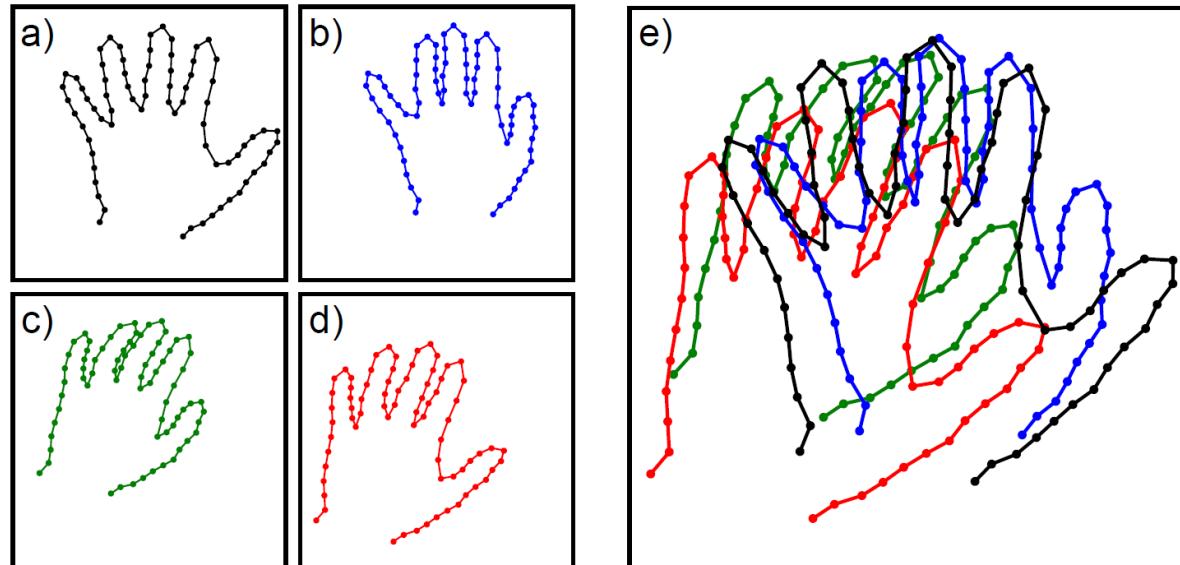
- Also called
 - Point distribution models
 - Active shape models (as they adapt to the image)
- Likelihood:

$$Pr(\mathbf{x}_i | \mathbf{w}_i) \propto \prod_{n=1}^N \exp [-(\text{dist} [\mathbf{x}_i, \text{trans}[\mathbf{w}_{in}, \Psi_i]])^2]$$

- Prior:

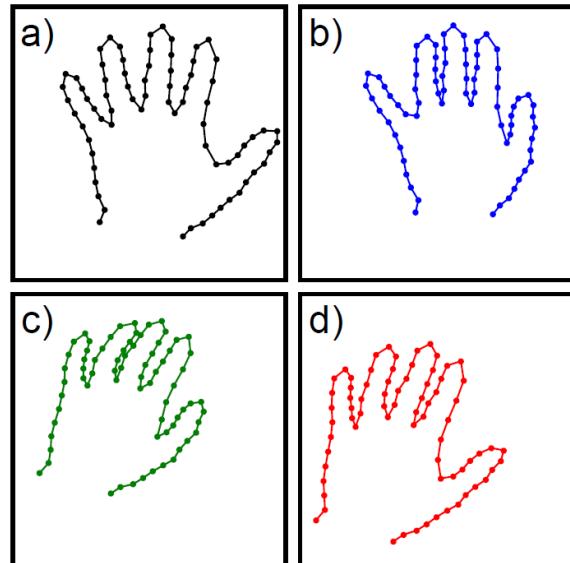
$$Pr(\mathbf{w}_i) = \text{Norm}_{\mathbf{w}_i} [\boldsymbol{\mu}, \boldsymbol{\Sigma}]$$

Learning

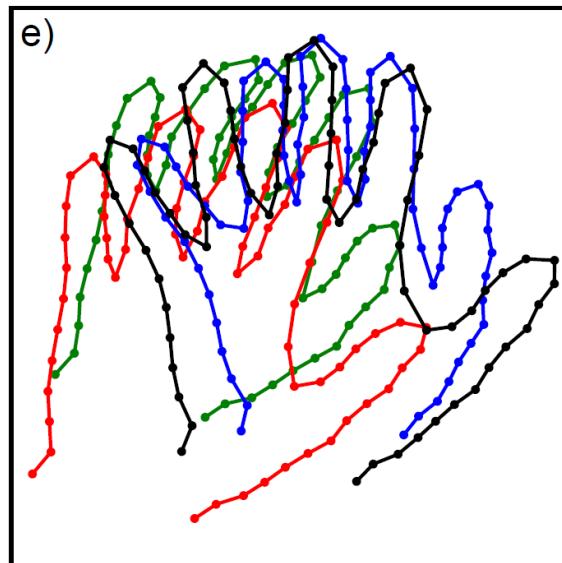


Training data

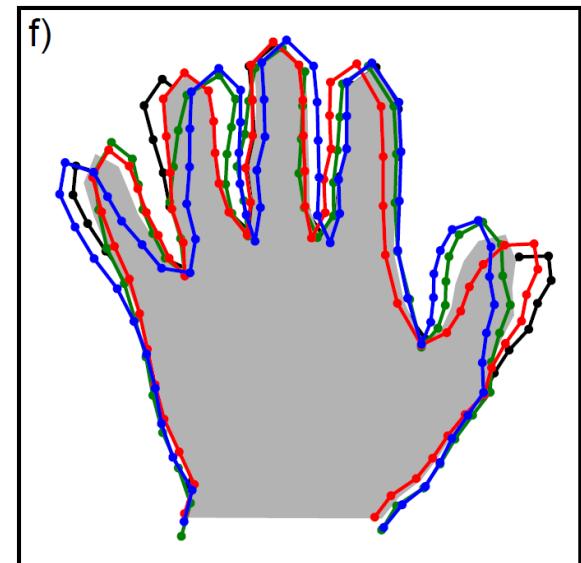
Generalized Procrustes analysis



Training data



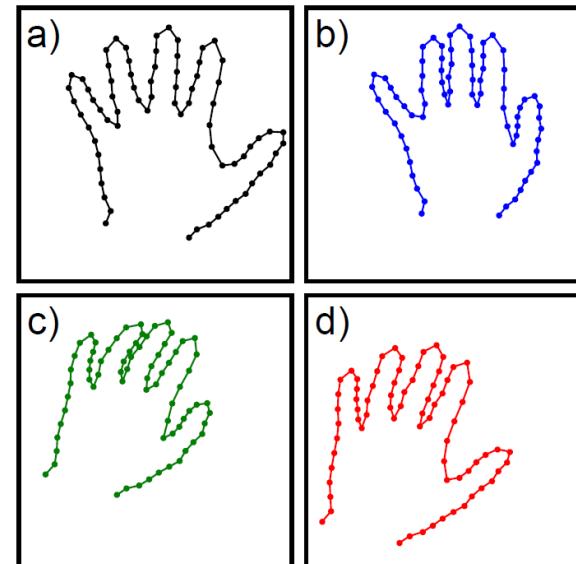
Before alignment



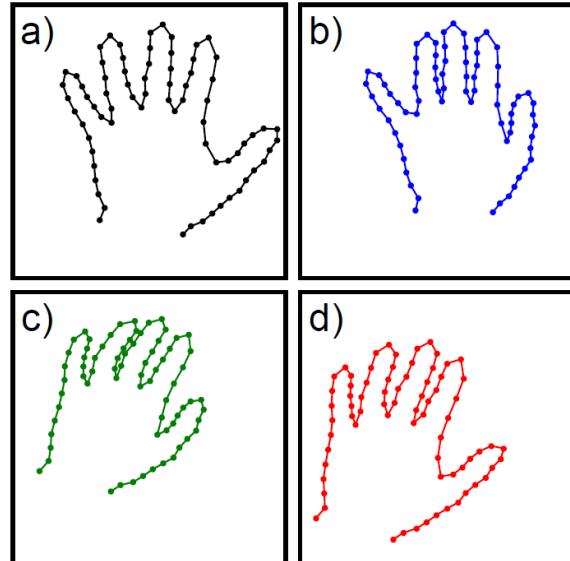
After alignment

Generalized Procrustes analysis

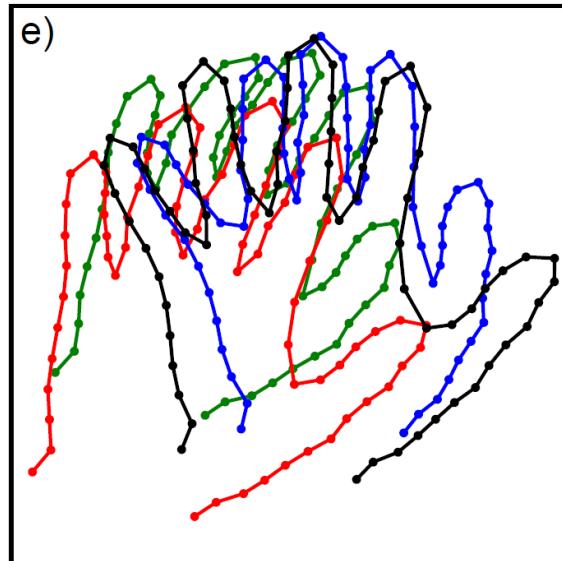
- Given I shapes with N landmarks
- Correspondences between shapes are known (landmarks)
- Take one shape as mean shape
- Iterate:
 - Solve for affine transformation ψ_i for each shape i to the mean $A_i x = b_i$
 - Apply transformation ψ_i to each shape i
 - Compute the mean μ_n of each landmark (mean shape)



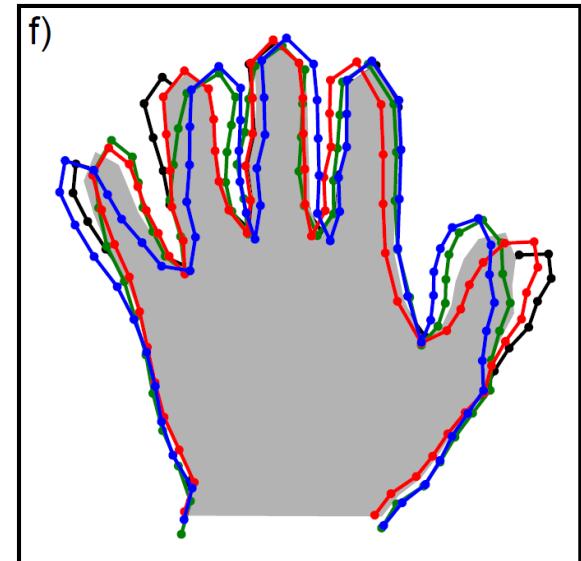
Generalized Procrustes analysis



Training data



Before alignment



After alignment

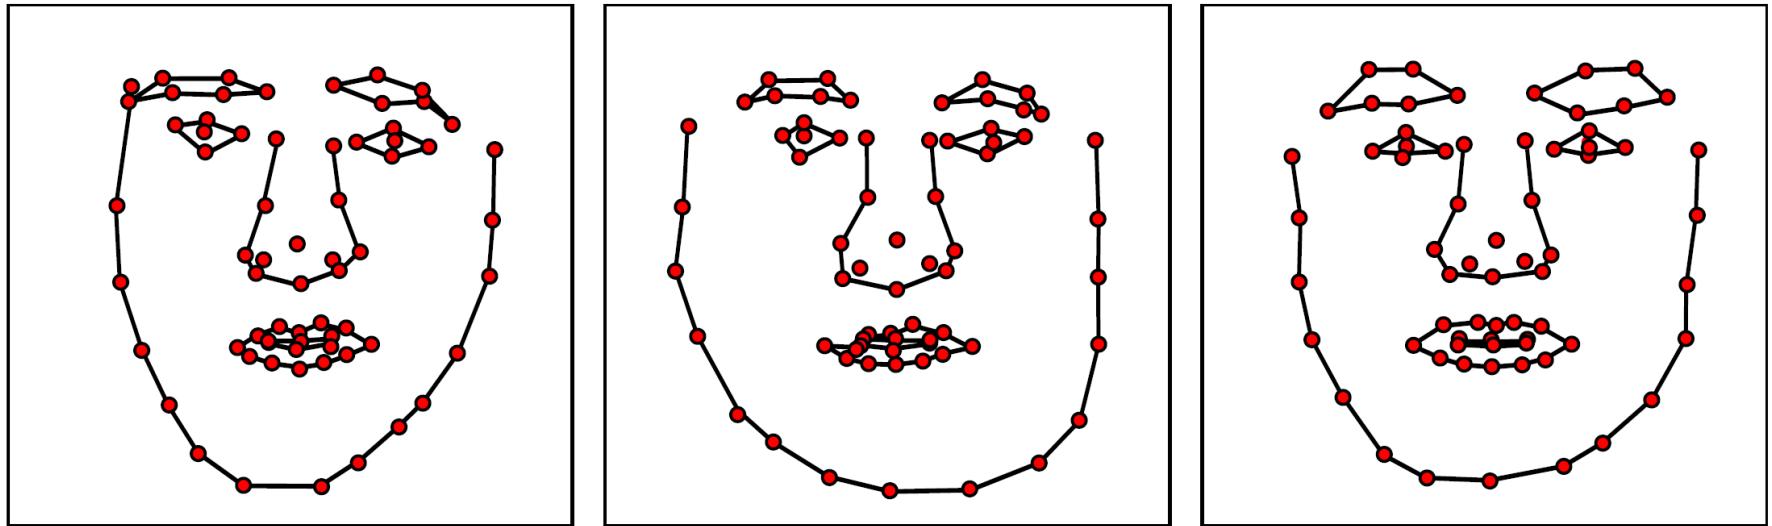
Inference

- Map inference:

$$\hat{\mathbf{w}} = \operatorname{argmax}_{\mathbf{w}} \left[\max_{\Psi} \left[\sum_{n=1}^N -(\operatorname{dist}[\mathbf{x}_i, \operatorname{trans}[\mathbf{w}_n, \Psi]])^2 + \log[\operatorname{Norm}_{\mathbf{w}}[\boldsymbol{\mu}, \boldsymbol{\Sigma}]] \right] \right]$$

- No closed form solution
- Use non-linear optimisation
- Or use ICP approach
- However, many parameters, and not clear they are all needed
- More efficient to use subspace model

Face model



Three samples from learned model for faces

Subspace shape model

- Generate data from model:

$$\mathbf{w}_i = \boldsymbol{\mu} + \Phi \mathbf{h}_i + \epsilon_i$$

- $\boldsymbol{\mu}$ is the mean shape
- the matrix $\Phi = [\phi_1, \phi_2, \dots, \phi_K]$ contains K basis functions in its columns
- ϵ_i is normal noise with covariance $\sigma^2 \mathbf{I}$

- Can alternatively write

$$\mathbf{w}_i = \boldsymbol{\mu} + \sum_{k=1}^K \phi_k h_{ik} + \epsilon_i$$

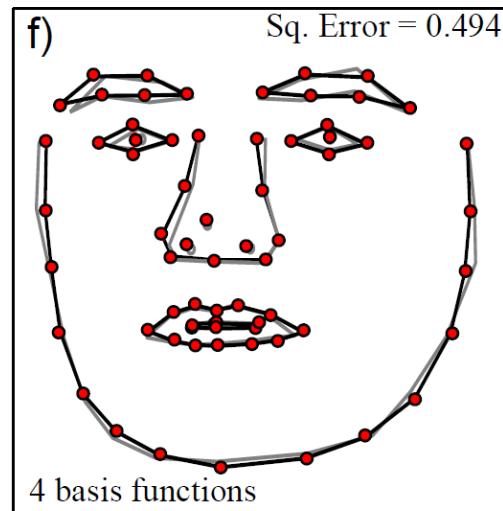
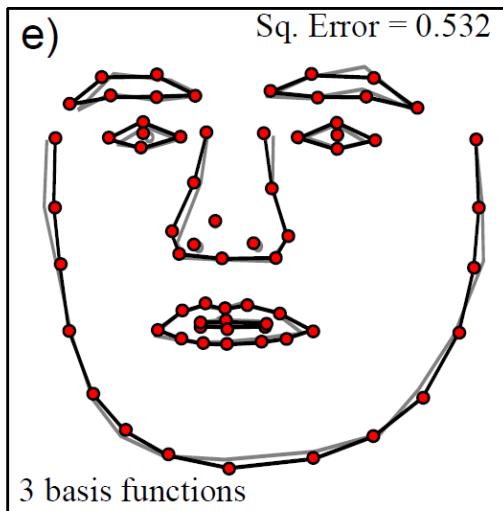
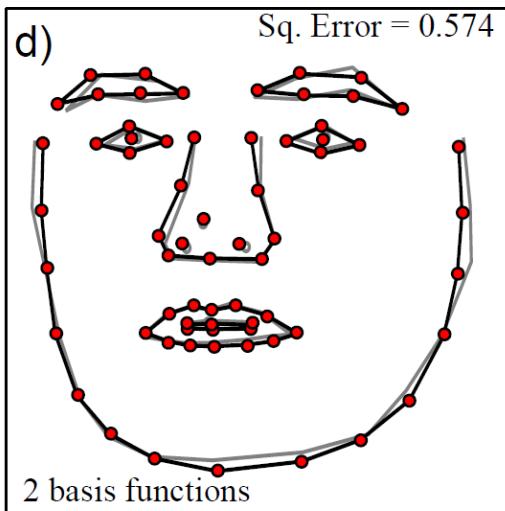
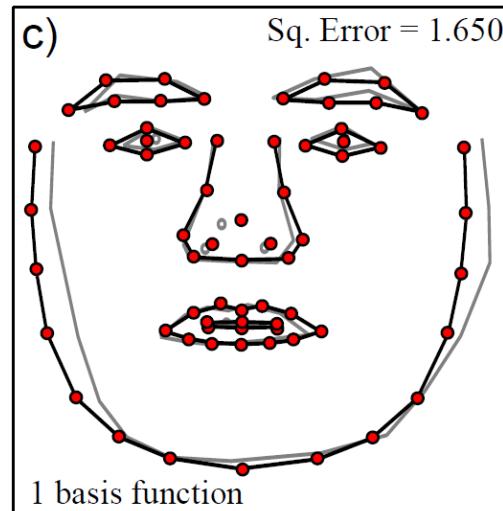
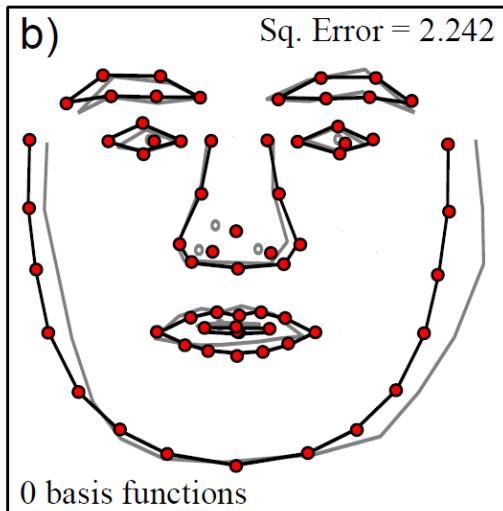
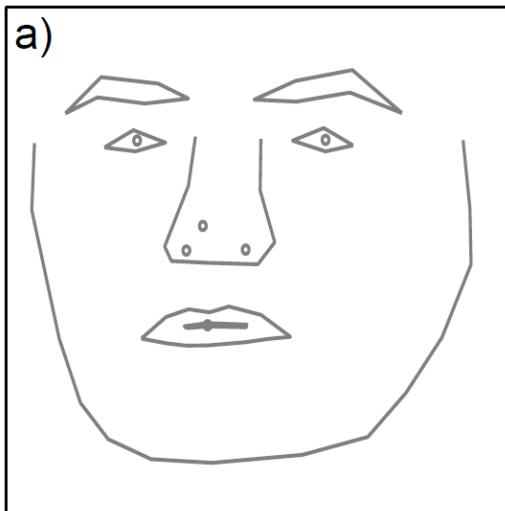
Approximating with subspace

Can approximate a vector \mathbf{w} with a weighted sum of the basis functions

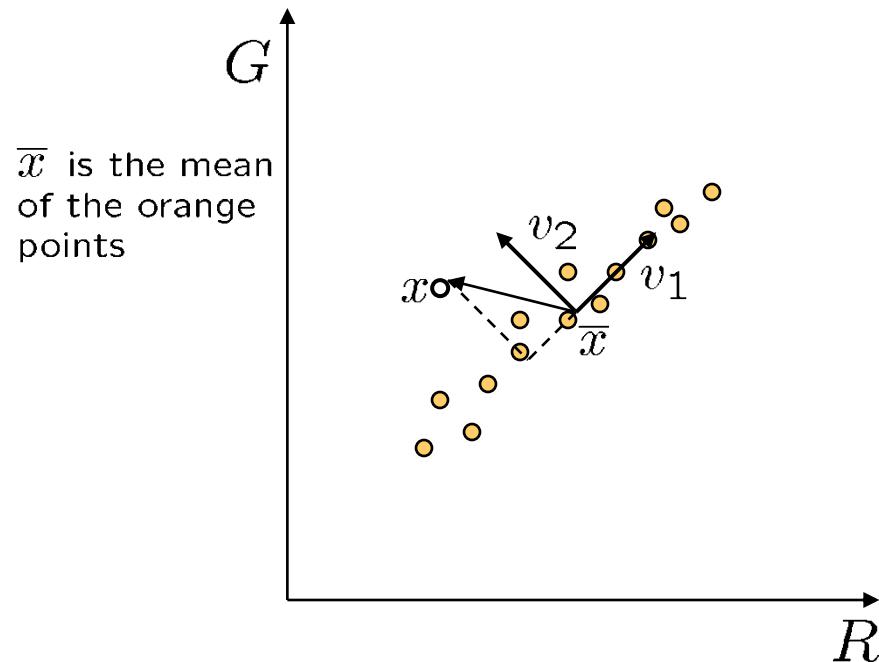
$$\mathbf{w}_i \approx \mu + \sum_{k=1}^K \phi_k h_{ik}$$

Surprising how well this works even with a small number of basis functions

Subspace shape model



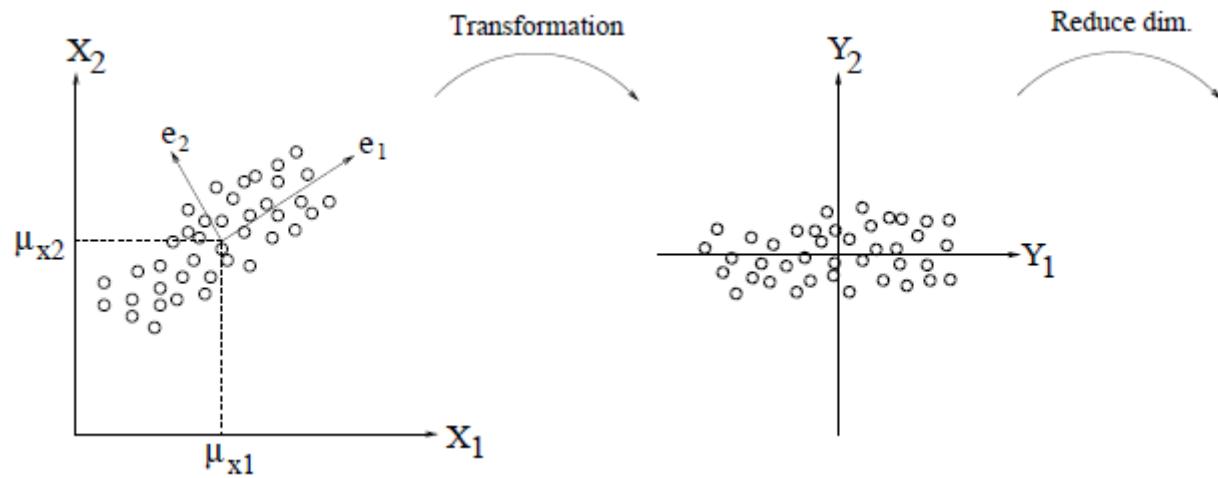
Linear subspaces



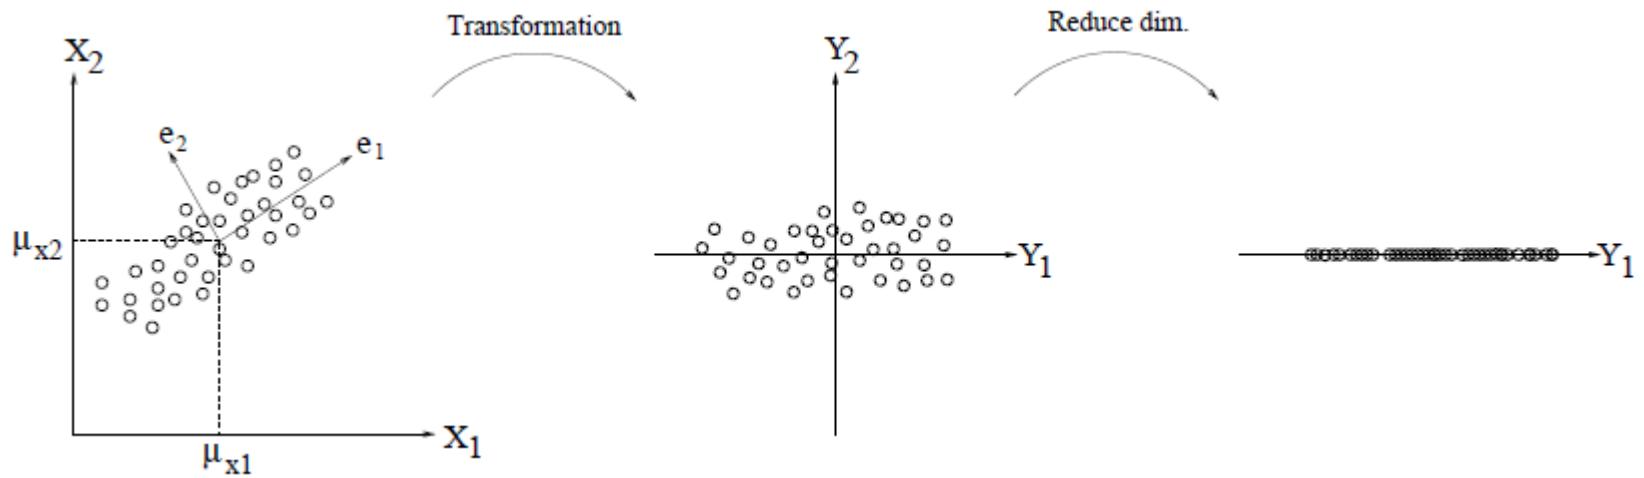
- convert \mathbf{x} into $\mathbf{v}_1, \mathbf{v}_2$ coordinates
 $\mathbf{x} \rightarrow ((\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_1, (\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_2)$
- \mathbf{v}_2 is distance to line
- \mathbf{v}_1 is position along line

- Represent high-dimensional space by linear subspace
- Idea: Estimate new coordinate system from data points

Dimensionality reduction

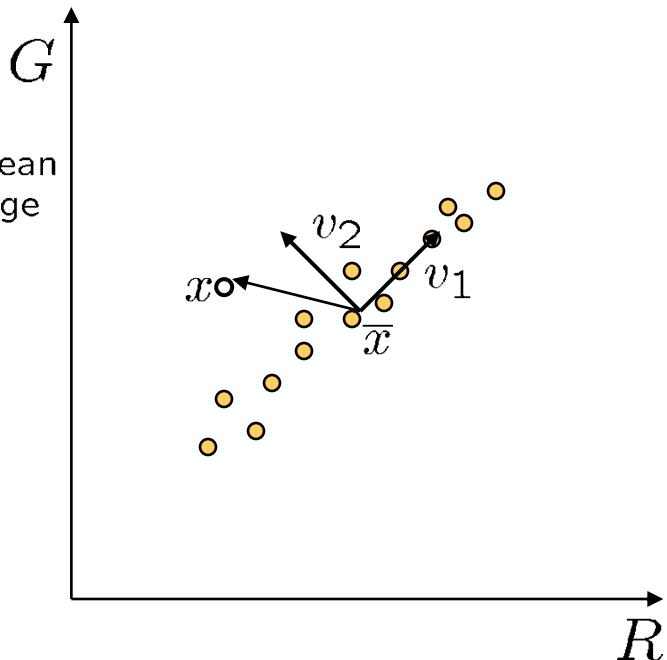


Dimensionality reduction



Linear subspaces

\bar{x} is the mean
of the orange
points



Consider the variation along direction \mathbf{v} among all of the orange points:

$$var(\mathbf{v}) = \sum_{\text{orange point } \mathbf{x}} \|(\mathbf{x} - \bar{\mathbf{x}})^T \cdot \mathbf{v}\|^2$$

$$\begin{aligned} var(\mathbf{v}) &= \sum_{\mathbf{x}} \|(\mathbf{x} - \bar{\mathbf{x}})^T \cdot \mathbf{v}\|^2 \\ &= \sum_{\mathbf{x}} \mathbf{v}^T (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{v} \\ &= \mathbf{v}^T \left[\sum_{\mathbf{x}} (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \right] \mathbf{v} \\ &= \mathbf{v}^T \mathbf{A} \mathbf{v} \quad \text{where } \mathbf{A} = \sum_{\mathbf{x}} (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \end{aligned}$$

Recall: Eigenvalue Problem

- Maximize objective J (Rayleigh Quotient):

$$J(\mathbf{w}) = \frac{\mathbf{w}^T S_B \mathbf{w}}{\mathbf{w}^T S_W \mathbf{w}}$$

- Equivalent to

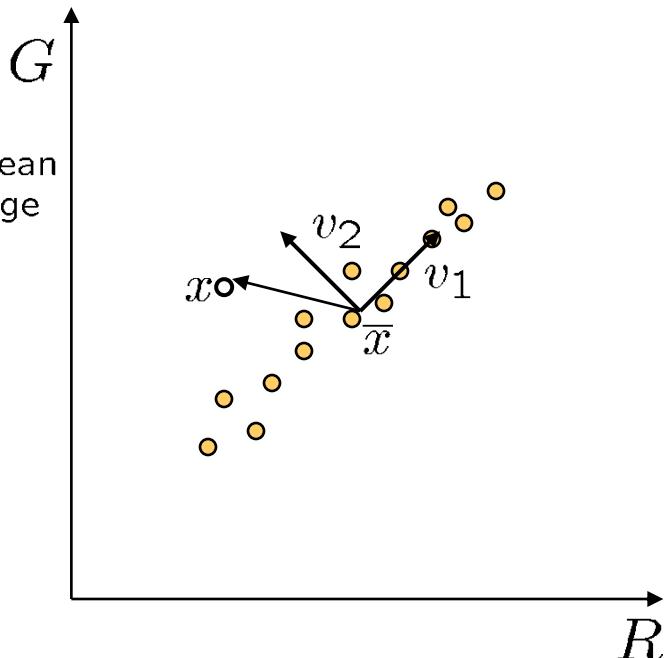
$$\begin{aligned} \min_{\mathbf{w}} \quad & -\frac{1}{2} \mathbf{w}^T S_B \mathbf{w} \\ \text{s.t.} \quad & \mathbf{w}^T S_W \mathbf{w} = 1 \end{aligned}$$

- Generalized eigenvalue problem:

$$S_B \mathbf{w} = \lambda S_W \mathbf{w}$$

Linear subspaces

\bar{x} is the mean of the orange points



Consider the variation along direction \mathbf{v} among all of the orange points:

$$\text{var}(\mathbf{v}) = \sum_{\text{orange point } \mathbf{x}} \|(\mathbf{x} - \bar{\mathbf{x}})^T \cdot \mathbf{v}\|^2$$

$$\begin{aligned} \text{var}(\mathbf{v}) &= \sum_{\mathbf{x}} \|(\mathbf{x} - \bar{\mathbf{x}})^T \cdot \mathbf{v}\|^2 \\ &= \sum_{\mathbf{x}} \mathbf{v}^T (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{v} \\ &= \mathbf{v}^T \left[\sum_{\mathbf{x}} (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \right] \mathbf{v} \\ &= \mathbf{v}^T \mathbf{A} \mathbf{v} \quad \text{where } \mathbf{A} = \sum_{\mathbf{x}} (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \end{aligned}$$

Solution: \mathbf{v}_1 is eigenvector of \mathbf{A} with *largest* eigenvalue

\mathbf{v}_2 is eigenvector of \mathbf{A} with *smallest* eigenvalue

Principal component analysis (PCA)

- Suppose each data point is N-dimensional
- Same procedure applies:

$$\begin{aligned} \text{var}(\mathbf{v}) &= \sum_{\mathbf{x}} \|(\mathbf{x} - \bar{\mathbf{x}})^T \cdot \mathbf{v}\| \\ &= \mathbf{v}^T \mathbf{A} \mathbf{v} \quad \text{where } \mathbf{A} = \sum_{\mathbf{x}} (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \end{aligned}$$

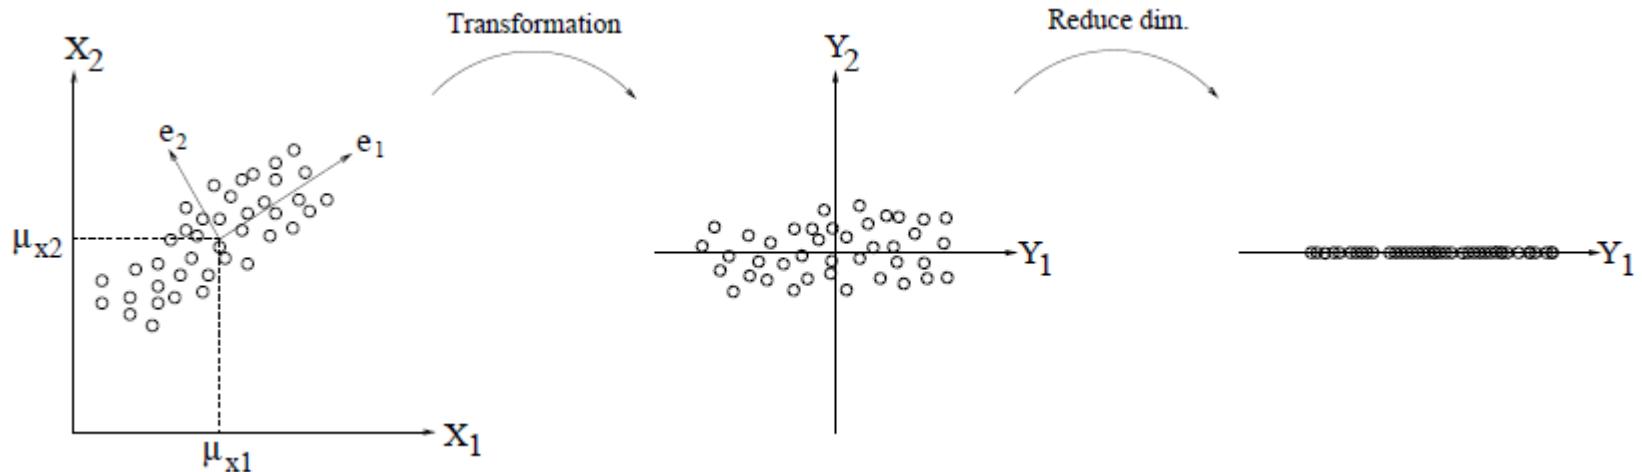
- The eigenvectors of \mathbf{A} define a new coordinate system
 - eigenvector with largest eigenvalue captures the most variation among training vectors \mathbf{x}
 - eigenvector with smallest eigenvalue has least variation
- We can compress the data by only using the top few eigenvectors
 - corresponds to choosing a “linear subspace”
 - represent points on a line, plane, or “hyper-plane”
 - these eigenvectors are known as the ***principal components***

Principal component analysis (PCA)

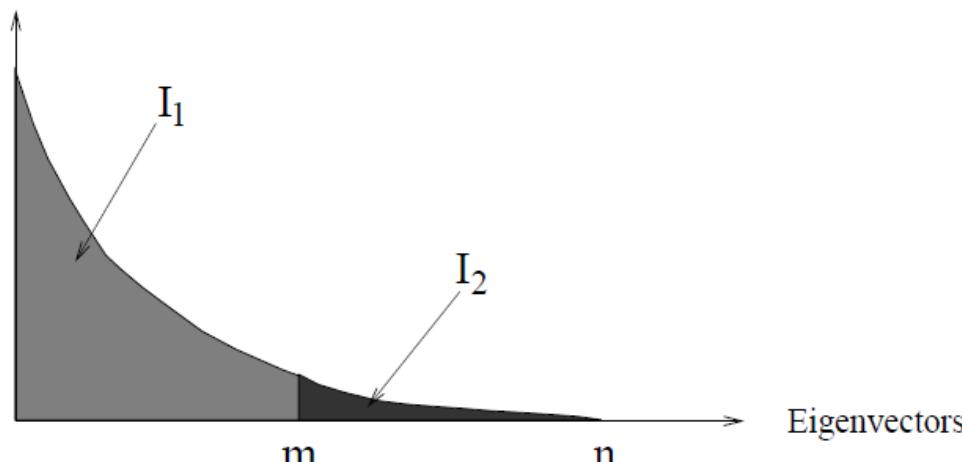
- Given data points w_i
- Compute mean μ
- $W = (w_1 - \mu, \dots, w_n - \mu)$
- Eigenvalue decomposition:
 $WW^T = U\Sigma V^T V\Sigma U^T = UL^2U^T$ (u_i basis vectors)
- SVD: $W = U\Sigma V^T$

Principal component analysis (PCA)

- Dimensionality reduction



Eigenvalues

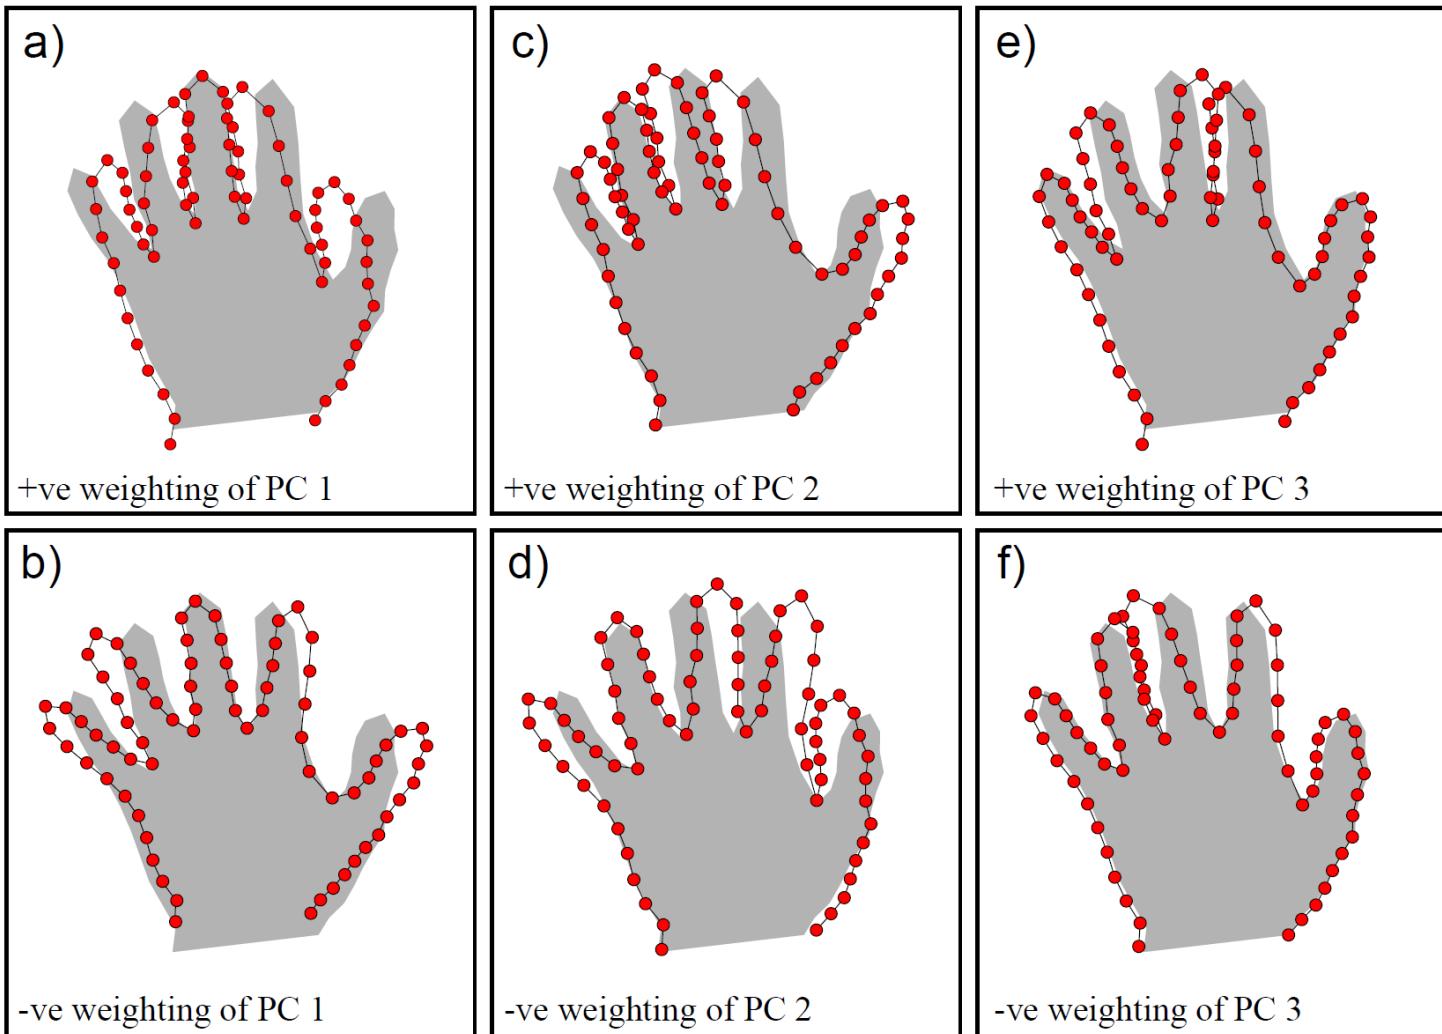


$$\frac{\sum_{i=1}^{i=K} L_{ii}^2}{\sum_{i=1}^{i=D} L_{ii}^2} > 1 - \epsilon$$

Source: T. Moeslund

Learned hand model

$$\mu + \sum_{k=1}^K \phi_k h_{ik}$$



Probabilistic PCA

Generative eq: $\mathbf{w}_i = \boldsymbol{\mu} + \Phi \mathbf{h}_i + \epsilon_i$

Probabilistic version:

$$Pr(\mathbf{w}_i | \mathbf{h}_i, \boldsymbol{\mu}, \Phi, \sigma^2) = \text{Norm}_{\mathbf{w}_i} [\boldsymbol{\mu} + \Phi \mathbf{h}_i, \sigma^2 \mathbf{I}]$$

Add prior: $Pr(\mathbf{h}_i) = \text{Norm}_{\mathbf{h}_i} [\mathbf{0}, \mathbf{I}]$

$$\begin{aligned}\text{Density: } Pr(\mathbf{w}_i) &= \int Pr(\mathbf{w}_i | \mathbf{h}_i) Pr(\mathbf{h}_i) d\mathbf{h}_i \\ &= \int \text{Norm}_{\mathbf{w}_i} [\boldsymbol{\mu} + \Phi \mathbf{h}_i, \sigma^2 \mathbf{I}] \text{Norm}_{\mathbf{h}_i} [\mathbf{0}, \mathbf{I}] d\mathbf{h}_i \\ &= \text{Norm}_{\mathbf{w}_i} [\boldsymbol{\mu}, \Phi \Phi^T + \sigma^2 \mathbf{I}].\end{aligned}$$

Probabilistic PCA

$$\int \text{Norm}_{w_i}[\mu + \Phi h_i, \sigma^2 I] \text{Norm}_{h_i}[0, I] dh_i$$

using $w = \mu + \Phi h_i, \quad \Phi^{-1}(w - \mu) = h_i$

$$= \int \text{Norm}_{w_i}[\mu + \Phi\Phi^{-1}(w - \mu), \sigma^2 I] \text{Norm}_w[\Phi 0 + \mu, \Phi I \Phi^T] dw$$

$$= \int \text{Norm}_{w_i}[w, \sigma^2 I] \text{Norm}_w[\mu, \Phi\Phi^T] dw$$

using $\int \text{Norm}_x[a, A] \text{Norm}_x[b, B] dx = \text{Norm}_a[b, A + B]$

$$= \text{Norm}_{w_i}[\mu, \sigma^2 I + \Phi\Phi^T]$$

Learning PPCA

Learn parameters μ , Φ and σ^2 from data $\{\mathbf{w}_i\}_{i=1}^I$
where $\mathbf{w}_i = [\mathbf{w}_{i1}^T, \mathbf{w}_{i2}^T, \dots, \mathbf{w}_{iN}^T]^T$.

Learn mean: $\mu = \frac{\sum_{i=1}^I \mathbf{w}_i}{I}$

Then set $\mathbf{W} = [\mathbf{w}_1 - \mu, \mathbf{w}_2 - \mu, \dots, \mathbf{w}_I - \mu]$
and compute eigen-decomposition

$$\mathbf{W}\mathbf{W}^T = \mathbf{U}\mathbf{L}^2\mathbf{U}^T$$

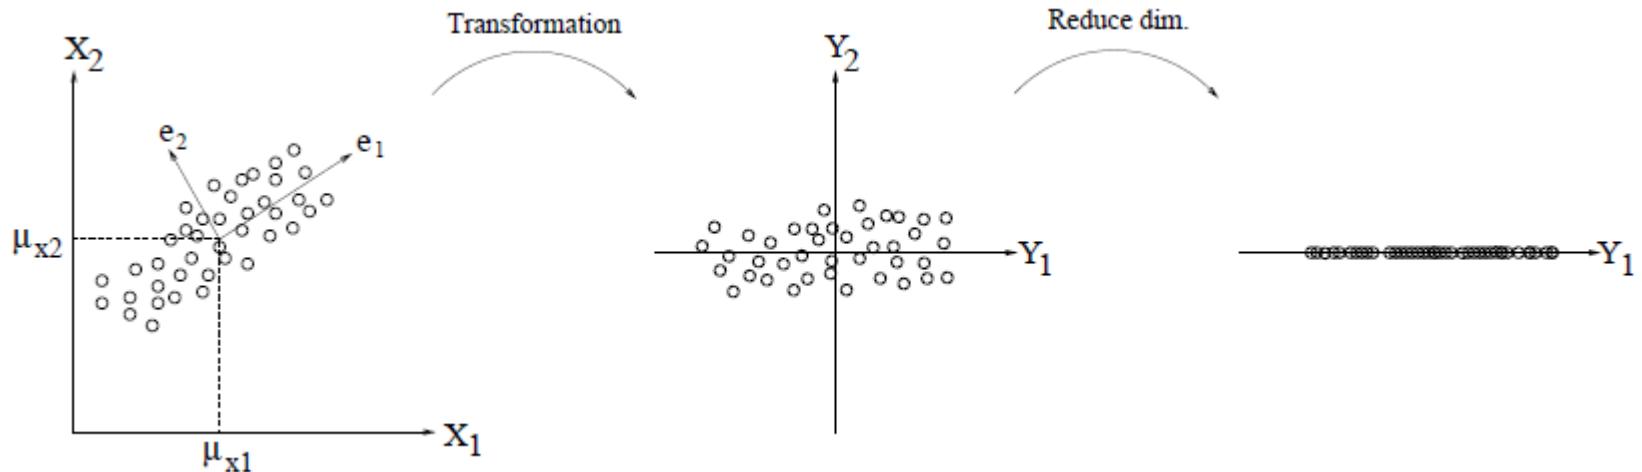
Choose parameters

$$\hat{\sigma}^2 = \frac{1}{D-K} \sum_{j=K+1}^D L_{jj}^2$$

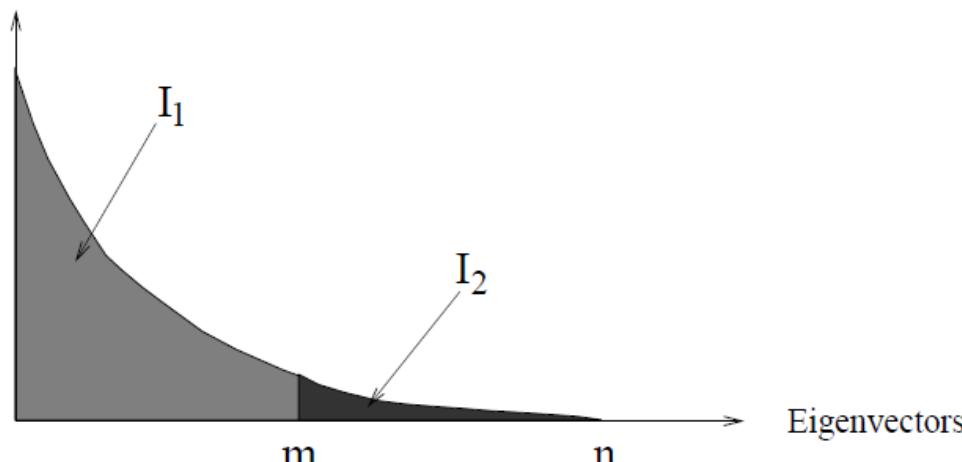
$$\hat{\Phi} = \mathbf{U}_K (\mathbf{L}_K^2 - \hat{\sigma}^2 \mathbf{I})^{1/2}$$

Principal component analysis (PCA)

- Dimensionality reduction



Eigenvalues



$$\frac{\sum_{i=1}^{i=K} L_{ii}^2}{\sum_{i=1}^{i=D} L_{ii}^2} > 1 - \epsilon$$

Source: T. Moeslund

Properties of basis functions

Learning of parameters based on eigen-decomposition:

$$\mathbf{W}\mathbf{W}^T = \mathbf{U}\mathbf{L}^2\mathbf{U}^T$$

Parameters for representing D-dim. Data by K principal components:

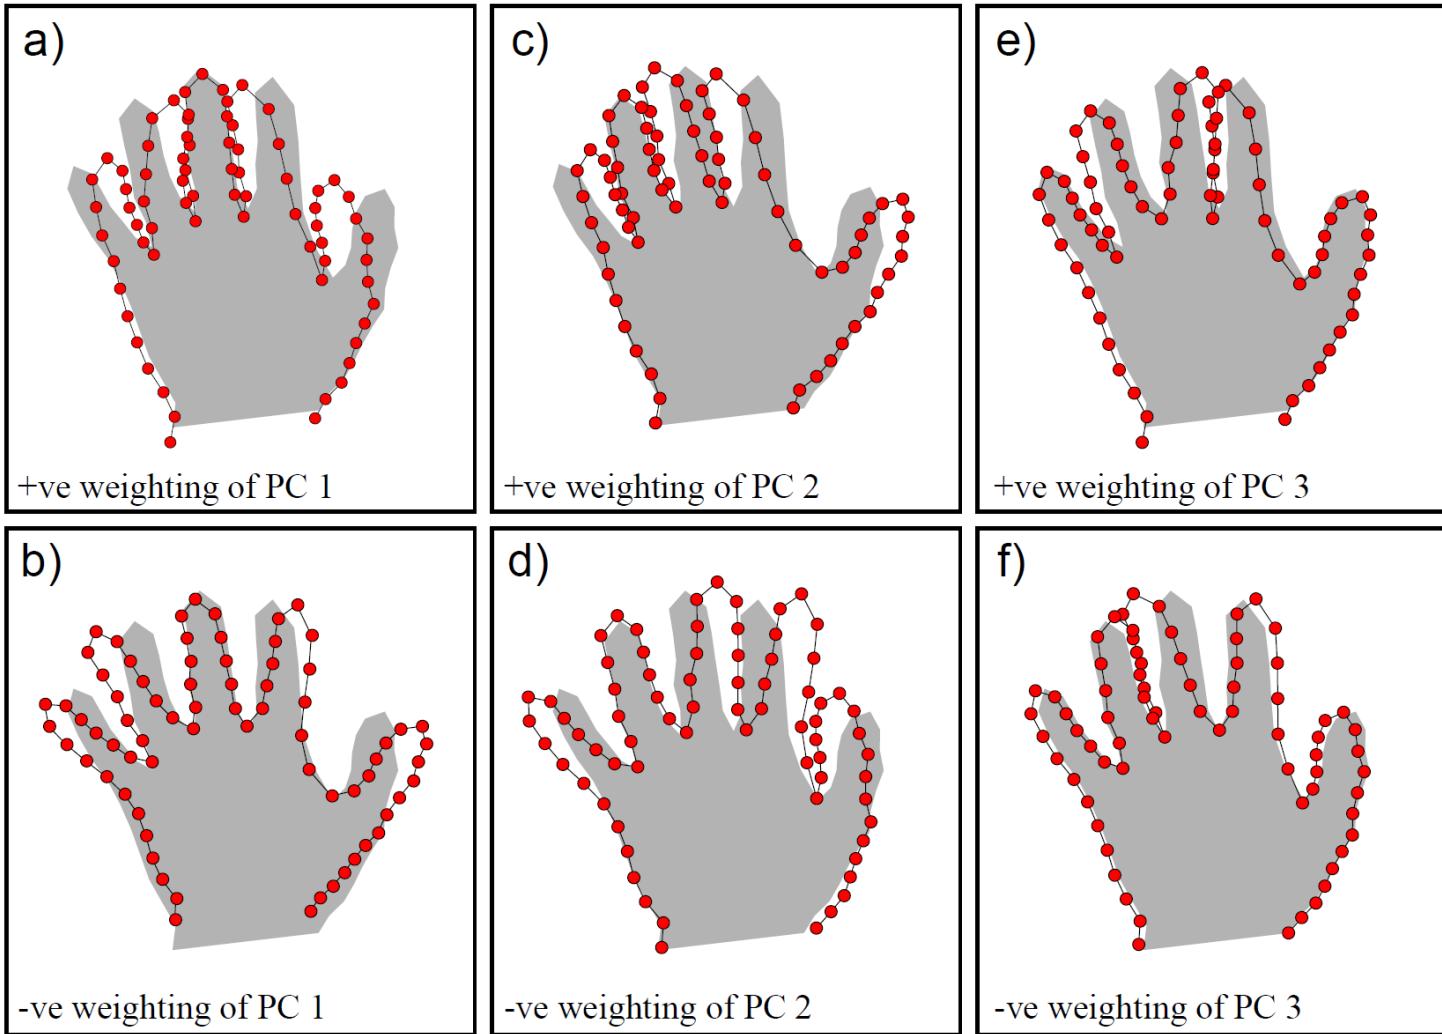
$$\hat{\sigma}^2 = \frac{1}{D - K} \sum_{j=K+1}^D L_{jj}^2 \quad \text{Estimate noise}$$
$$\hat{\Phi} = \mathbf{U}_K (\mathbf{L}_K^2 - \hat{\sigma}^2 \mathbf{I})^{1/2} \quad \text{As PCA but remove noise}$$

Notice that:

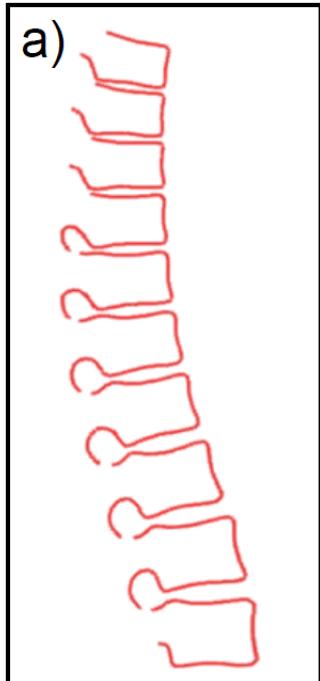
- Basis functions in $\hat{\Phi}$ are orthogonal
- Basis functions in $\hat{\Phi}$ are ordered

Learned hand model

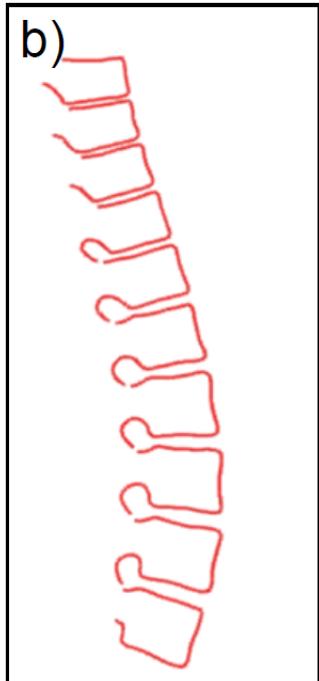
$$\mu + \sum_{k=1}^K \phi_k h_{ik}$$



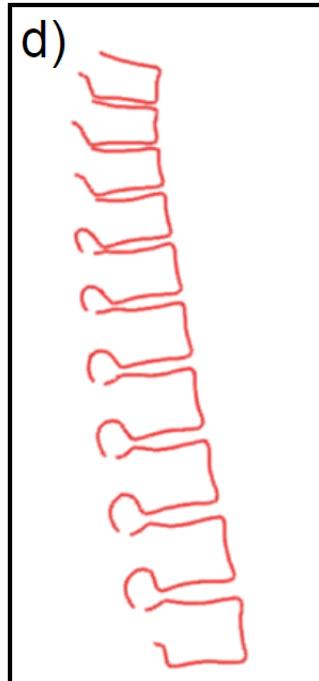
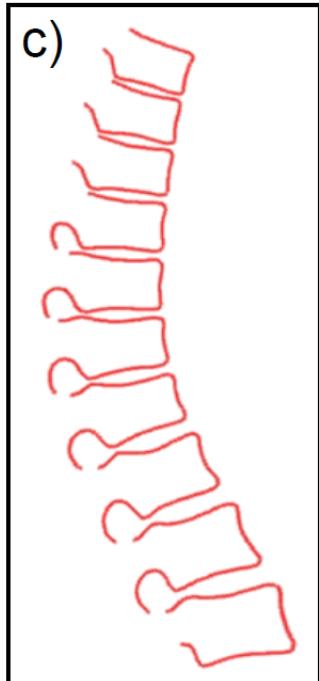
Learned spine model



Mean



Manipulating first
principal component



Manipulating second
principal component

Inference

To fit model to an image:

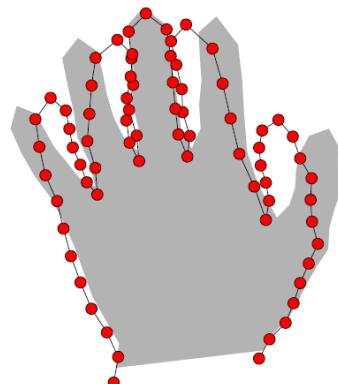
$$\hat{\mathbf{h}} = \operatorname{argmax}_{\mathbf{h}} \left[\max_{\Psi} \left[\sum_{n=1}^N \left(-\frac{(\operatorname{dist} [\mathbf{x}_i, \operatorname{trans}[\boldsymbol{\mu}_n + \Phi_n \mathbf{h}, \Psi]])^2}{\sigma^2} \right) + \log[\operatorname{Norm}_{\mathbf{h}}[\mathbf{0}, \mathbf{I}]] \right] \right]$$

likelihood

prior

ICP Approach:

- Find closest points to current prediction
- Update weightings \mathbf{h}
- Find closest points to current prediction
- Update transformation parameters Ψ



Source: S. Prince

Inference

1. Update weightings \mathbf{h}

$$\begin{aligned}\hat{\mathbf{h}} &= \operatorname{argmax}_{\mathbf{h}} \left[\sum_{n=1}^N \log[Pr(\mathbf{y}_n|\mathbf{h}), \Psi] + \log[Pr(\mathbf{h})] \right] \\ &= \operatorname{argmax}_{\mathbf{h}} \left[\sum_{n=1}^N -(\mathbf{y}_n - \text{trans}[\boldsymbol{\mu}_n + \Phi_n \mathbf{h}, \Psi])^2 / \sigma^2 - \log[\mathbf{h}^T \mathbf{h}] \right]\end{aligned}$$

If transformation parameters can be represented as $\mathbf{Aw} + \mathbf{b}$

$$\hat{\mathbf{h}} = \left(\sigma^2 \mathbf{I} + \sum_{n=1}^N \Phi_n^T \mathbf{A}^T \mathbf{A} \Phi_n \right)^{-1} \sum_{n=1}^N \Phi_n^T \mathbf{A}^T (\mathbf{y}_n - \mathbf{A} \boldsymbol{\mu}_n - \mathbf{b})$$

2. Update transformation parameters ψ

- Using one of the closed form solutions

Simple version with PCA

1. Initialise the shape parameters, \mathbf{b} , to zero (the mean shape).
2. Generate the model point positions using $\mathbf{x} = \bar{\mathbf{x}} + \mathbf{P}\mathbf{b}$
3. Find the pose parameters (X_t, Y_t, s, θ) which best align the model points \mathbf{x} to the current found points \mathbf{Y}

$$T_{X_t, Y_t, s, \theta} \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} X_t \\ Y_t \end{pmatrix} + \begin{pmatrix} s \cos \theta & -s \sin \theta \\ s \sin \theta & s \cos \theta \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix}$$

4. Project \mathbf{Y} into the model co-ordinate frame by inverting the transformation T :

$$\mathbf{y} = T_{X_t, Y_t, s, \theta}^{-1}(\mathbf{Y})$$

5. Project \mathbf{y} into the tangent plane to $\bar{\mathbf{x}}$ by scaling: $\mathbf{y}' = \mathbf{y}/(\mathbf{y} \cdot \bar{\mathbf{x}})$. (optional)

6. Update the model parameters to match to \mathbf{y}'

$$\mathbf{b} = \mathbf{P}^T (\mathbf{y}' - \bar{\mathbf{x}})$$

7. If not converged, return to step 2.

T. Cootes. An Introduction to Active Shape Models.

Fitting model

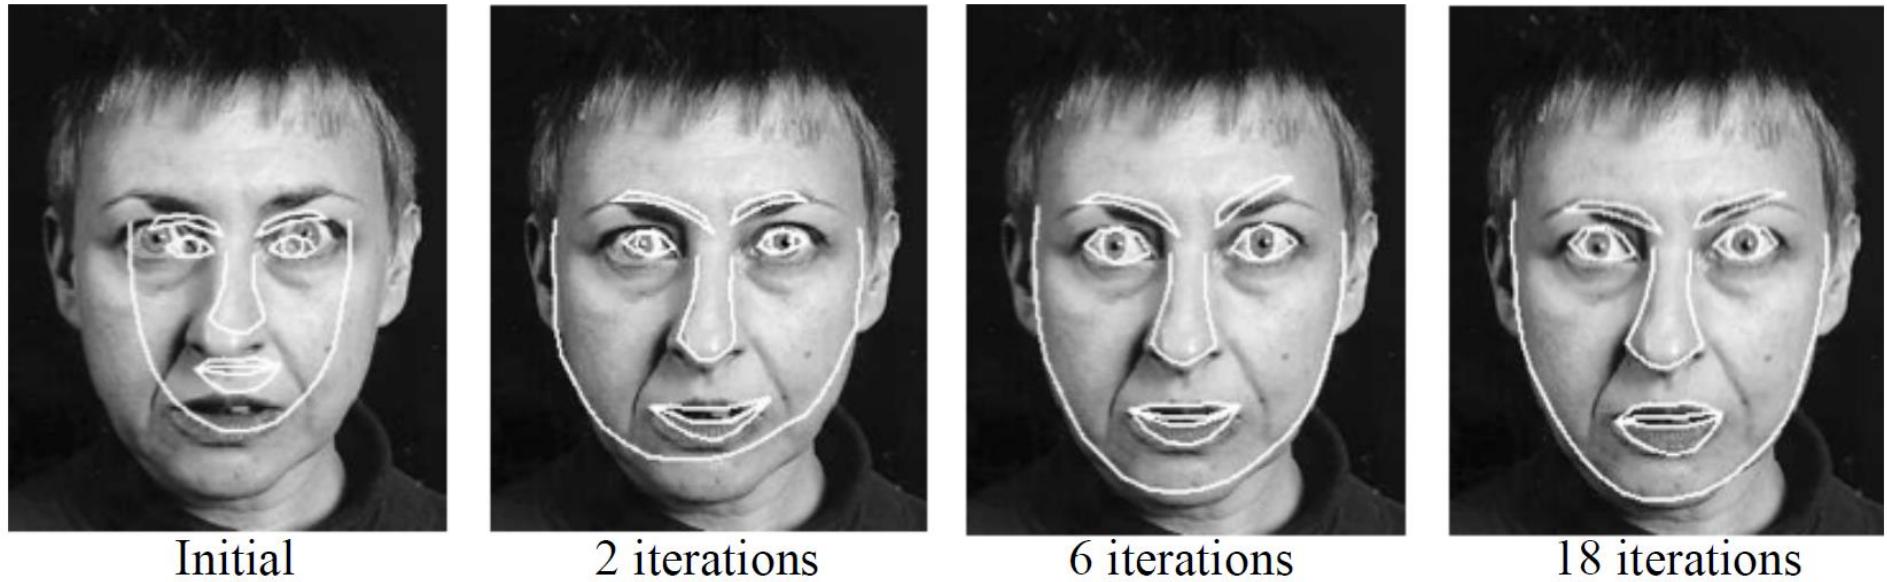
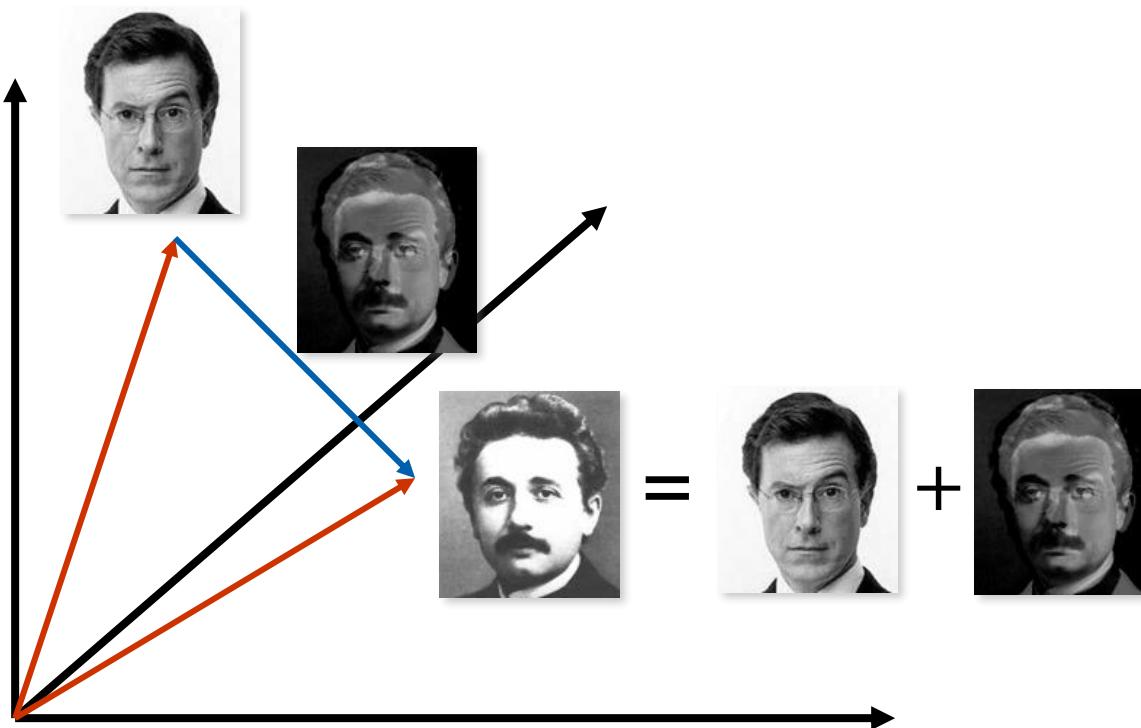


Figure provided by Tim Cootes

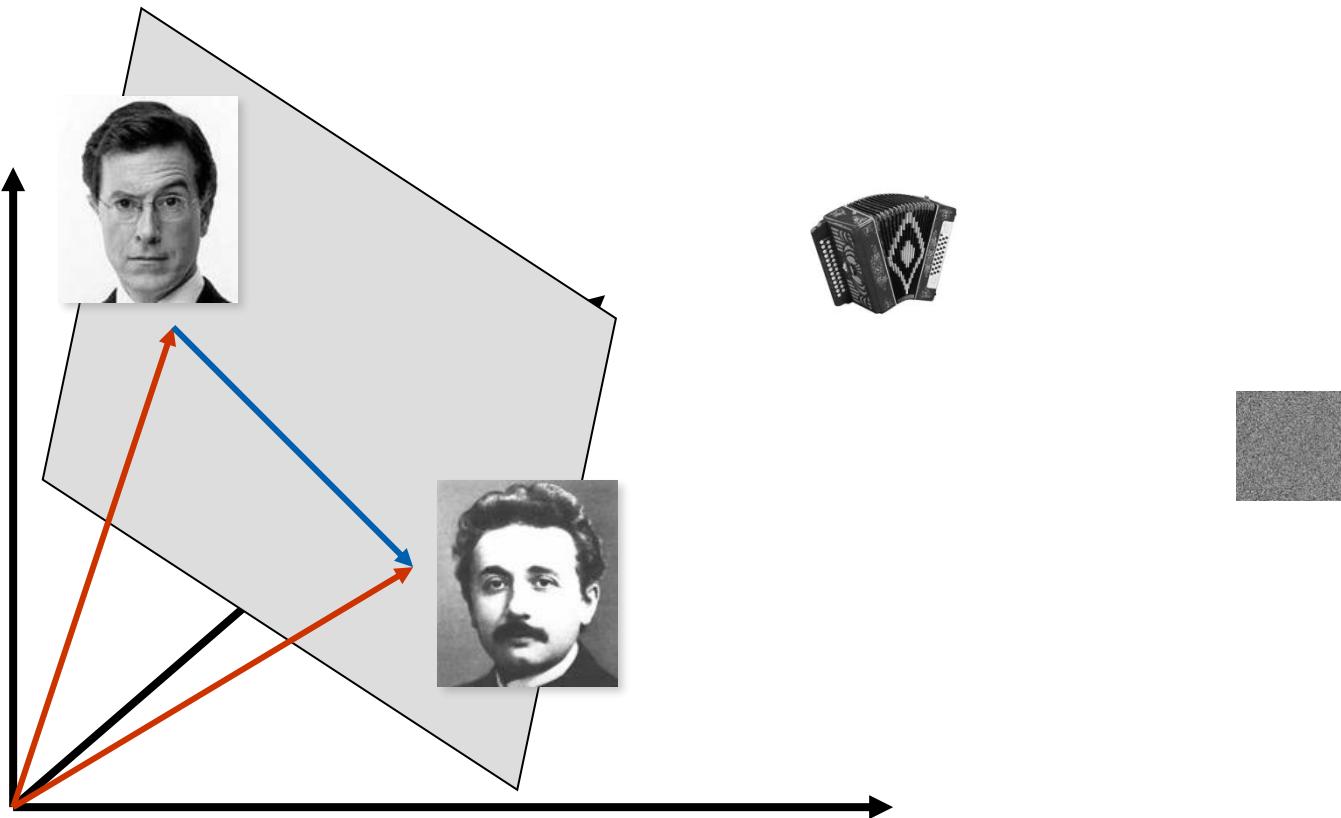
Much better to use statistical classifier instead of just distance from edges

The space of faces



- An image is a point in a high dimensional space
 - An $N \times M$ intensity image is a point in R^{NM}
 - We can define vectors in this space as we did in the 2D case

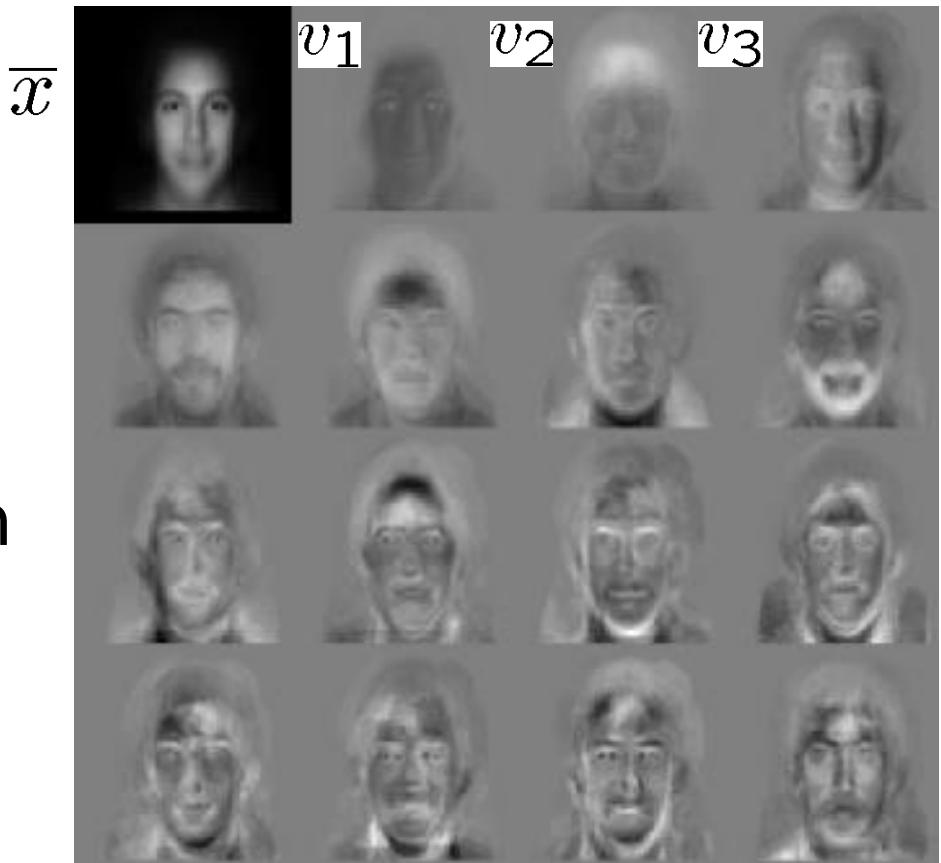
Dimensionality reduction



- The set of faces is a “subspace” of the set of images
 - Suppose it is K dimensional
 - We can find the best subspace using PCA
 - This is like fitting a “hyper-plane” to the set of faces
 - spanned by vectors v_1, v_2, \dots, v_K
 - any face $x \approx \bar{x} + a_1v_1 + a_2v_2 + \dots + a_kv_k$

Eigenfaces

- PCA extracts the eigenvectors of \mathbf{A}
 - Gives a set of vectors $\mathbf{v}_1, \mathbf{v}_2, \mathbf{v}_3, \dots$
 - Each one of these vectors is a direction in face space
 - what do these look like?

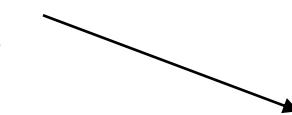


Projecting onto the eigenfaces

- The eigenfaces $\mathbf{v}_1, \dots, \mathbf{v}_K$ span the space of faces
- A face is converted to eigenface coordinates by

$$\mathbf{x} \rightarrow ((\underbrace{(\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_1}_{a_1}, \underbrace{(\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_2}_{a_2}, \dots, \underbrace{(\mathbf{x} - \bar{\mathbf{x}}) \cdot \mathbf{v}_K}_{a_K})$$

$$\mathbf{x} \approx \bar{\mathbf{x}} + a_1 \mathbf{v}_1 + a_2 \mathbf{v}_2 + \dots + a_K \mathbf{v}_K$$

 x 

Detection and recognition with eigenfaces

Algorithm

1. Process the image database (set of images with labels)

- Run PCA—compute eigenfaces
- Calculate the K coefficients for each image

2. Given a new image (to be recognized) \mathbf{x} , calculate K coefficients

$$\mathbf{x} \rightarrow (a_1, a_2, \dots, a_K)$$

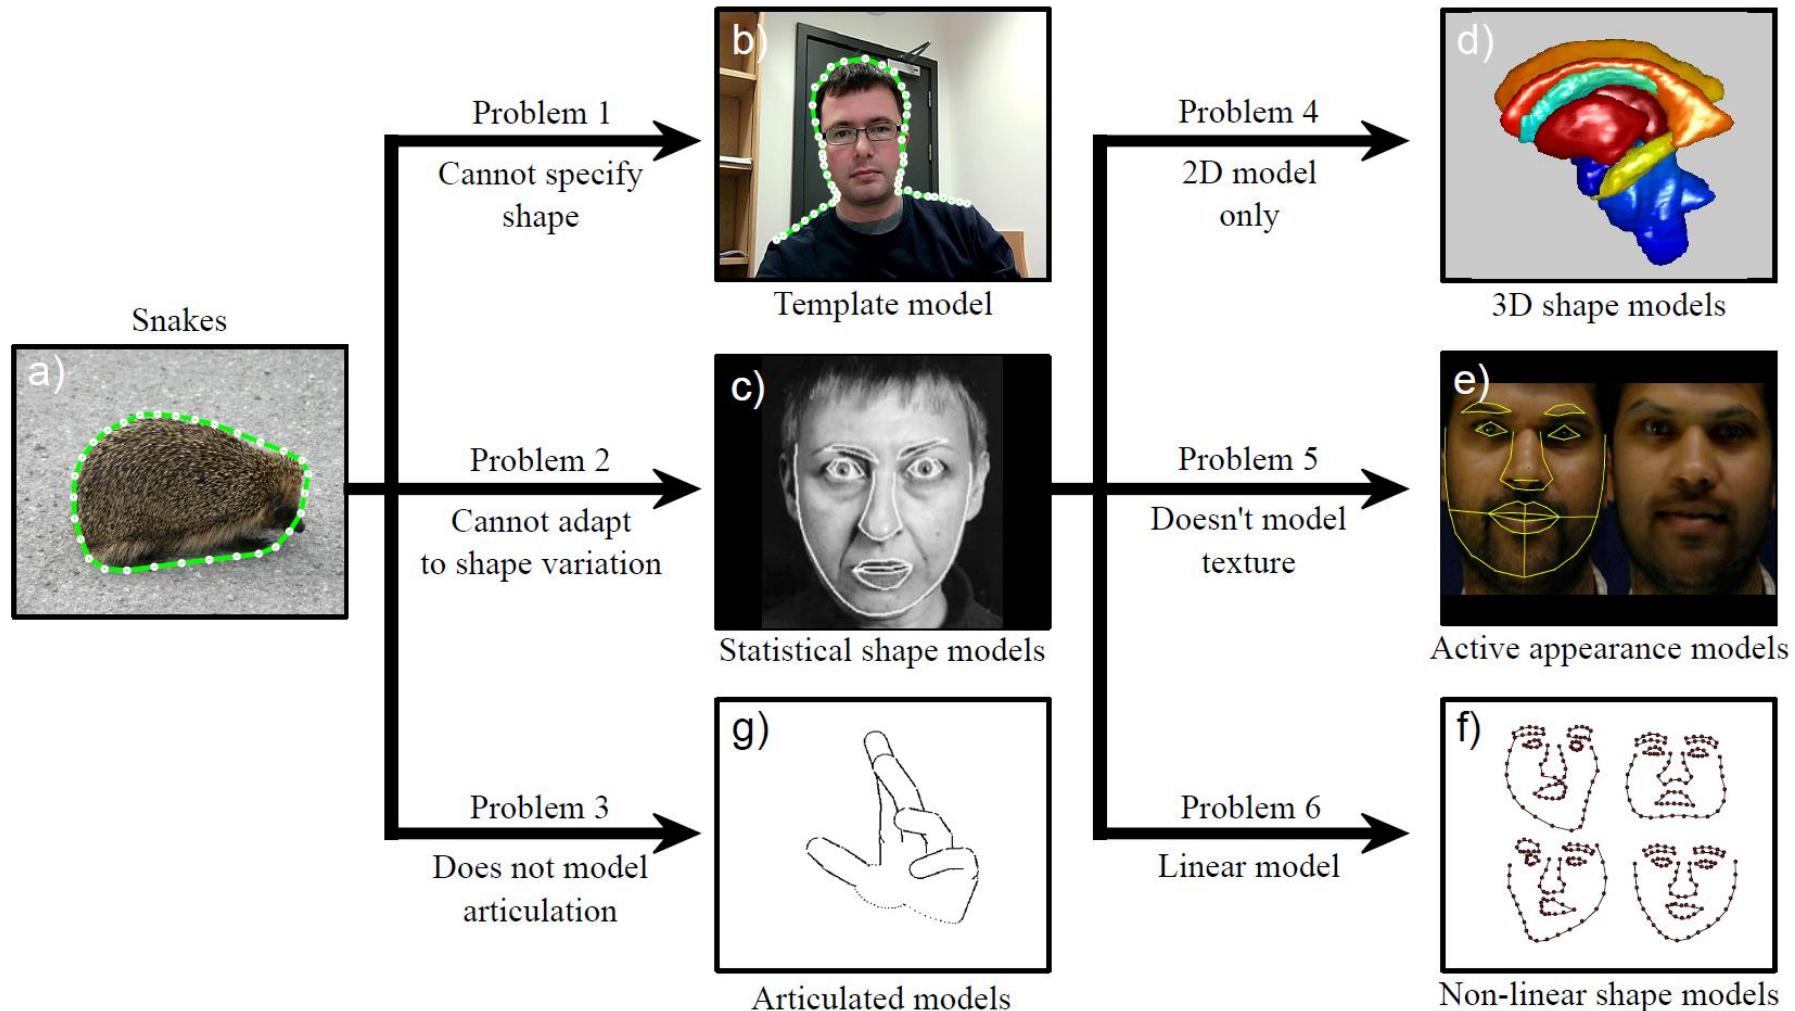
3. Detect if \mathbf{x} is a face

$$\|\mathbf{x} - (\bar{\mathbf{x}} + a_1\mathbf{v}_1 + a_2\mathbf{v}_2 + \dots + a_K\mathbf{v}_K)\| < \text{threshold}$$

4. If it is a face, who is it?

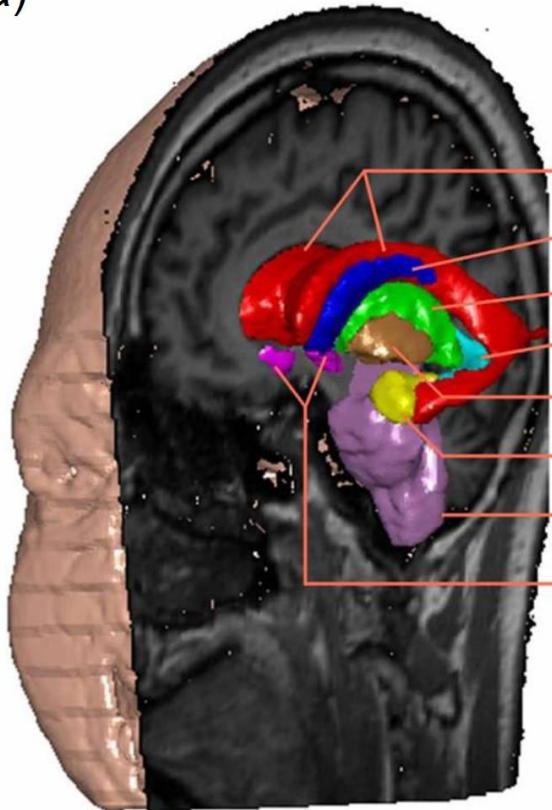
- Find closest labeled face in database (nearest-neighbor in K-dimensional space)

Relationships between models

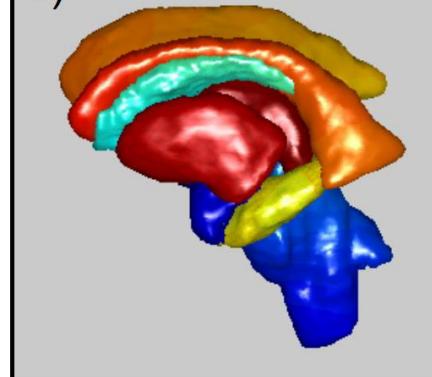


3D shape models

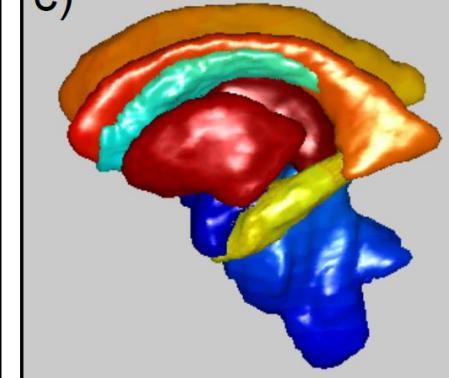
a)



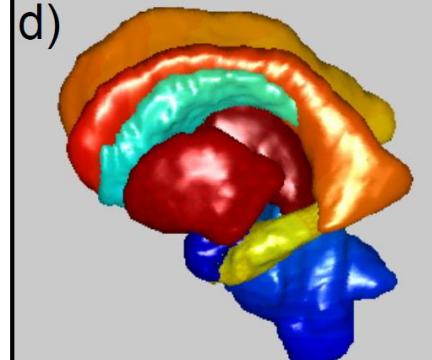
b)



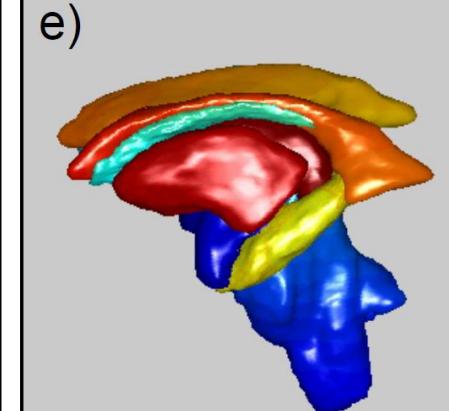
c)



d)



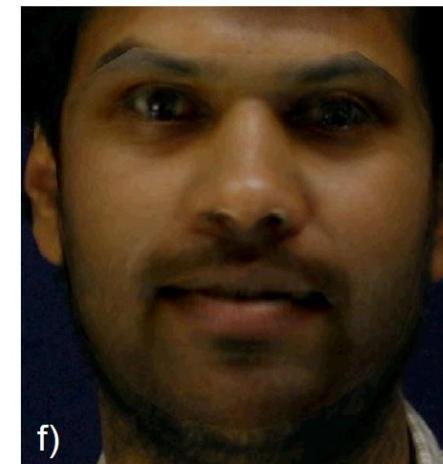
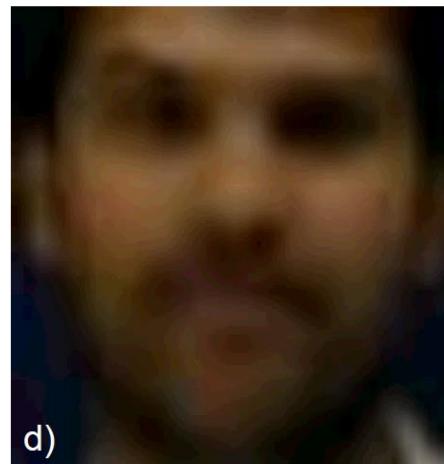
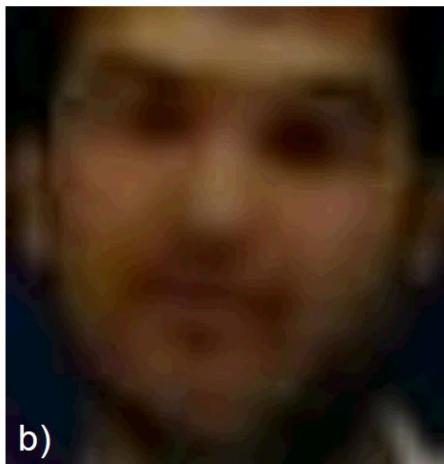
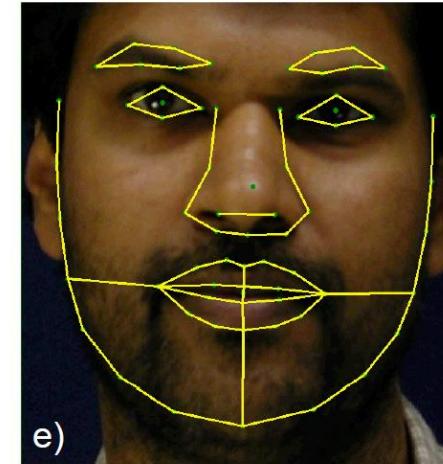
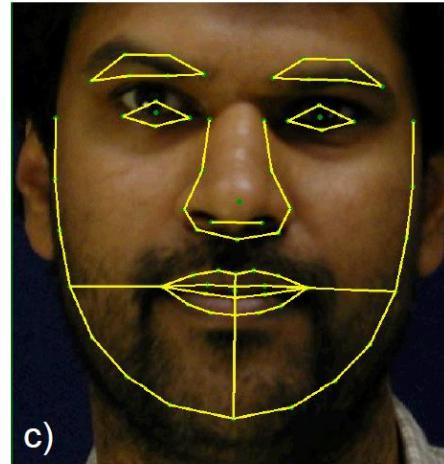
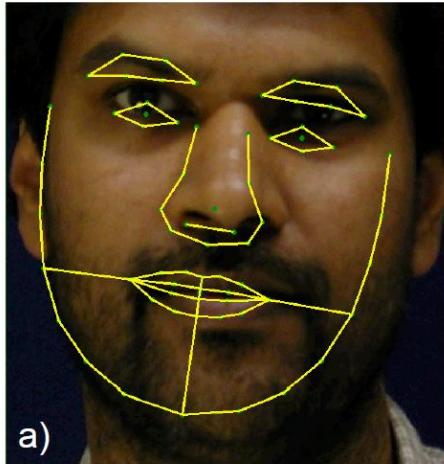
e)



Babalola *et al.* (2008)

Source: S. Prince

Statistical models for shape and appearance



Source: T. Cootes

Statistical models for shape and appearance

$$Pr(\mathbf{h}_i) = \text{Norm}_{\mathbf{h}_i}[\mathbf{0}, \mathbf{I}]$$

$$Pr(\mathbf{w}_i | \mathbf{h}_i) = \text{Norm}_{\mathbf{w}_i}[\boldsymbol{\mu}_w + \boldsymbol{\Phi}_w \mathbf{h}_i, \sigma_w^2 \mathbf{I}]$$

1. We draw a hidden variable \mathbf{h} from a prior
2. We draw landmark points \mathbf{w} from a subspace model

Statistical models for shape and appearance

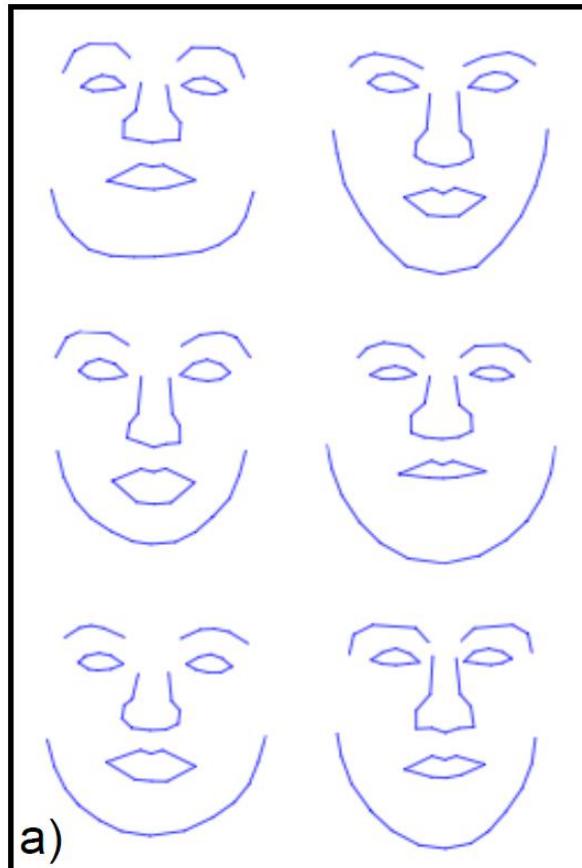
$$Pr(\mathbf{h}_i) = \text{Norm}_{\mathbf{h}_i}[\mathbf{0}, \mathbf{I}]$$

$$Pr(\mathbf{w}_i | \mathbf{h}_i) = \text{Norm}_{\mathbf{w}_i}[\boldsymbol{\mu}_w + \boldsymbol{\Phi}_w \mathbf{h}_i, \sigma_w^2 \mathbf{I}]$$

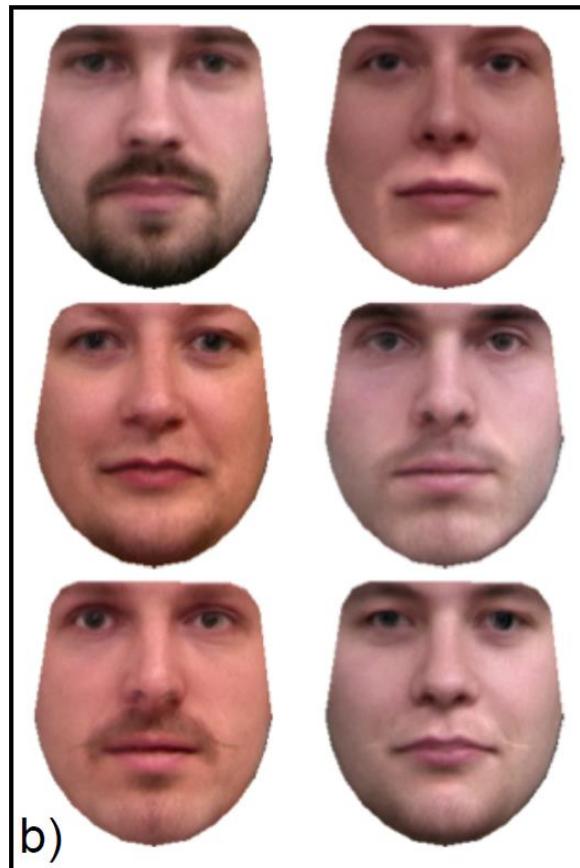
$$Pr(\mathbf{x}_i | \mathbf{w}_i, \mathbf{h}_i) = \text{Norm}_{\mathbf{x}_i}[\text{warp}[\boldsymbol{\mu}_x + \boldsymbol{\Phi}_x \mathbf{h}_i, \mathbf{w}_i, \boldsymbol{\Psi}_i], \sigma_x^2 \mathbf{I}]$$

1. We draw a hidden variable \mathbf{h} from a prior
2. We draw landmark points \mathbf{w} from a subspace model
3. We draw image intensities \mathbf{x} .
 - Generate image intensities in standard template shape
 - Transform the landmark points (parameters ψ)
 - Transform the image to landmark points
 - Add noise

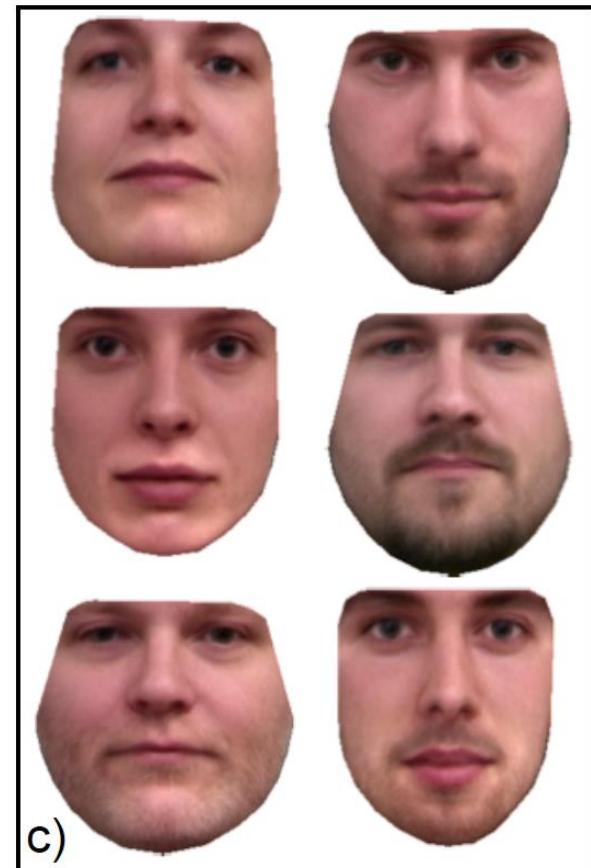
Shape and appearance model



Shape
model

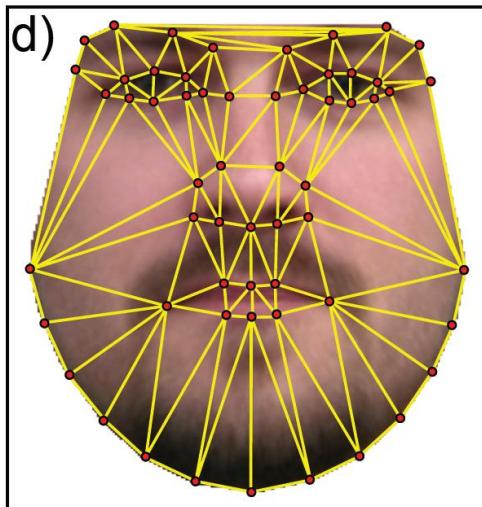
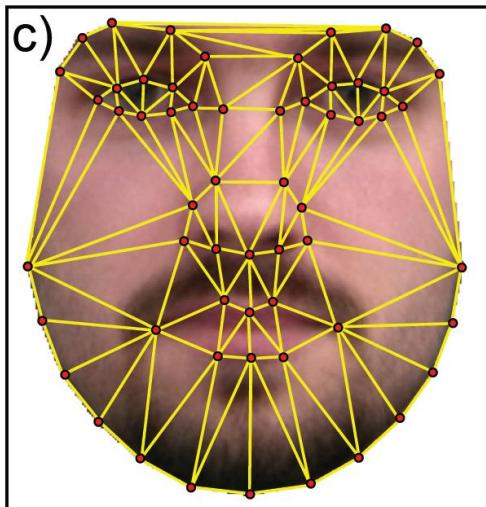
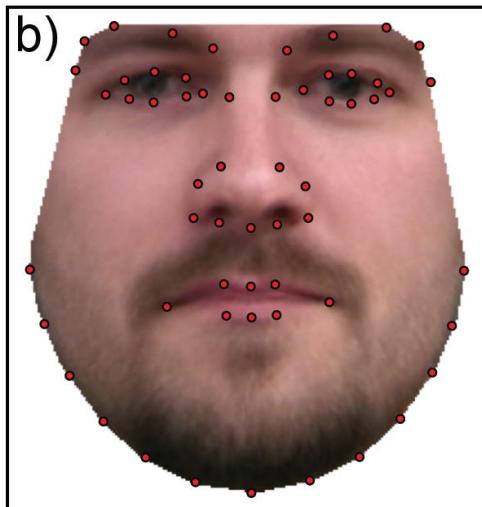
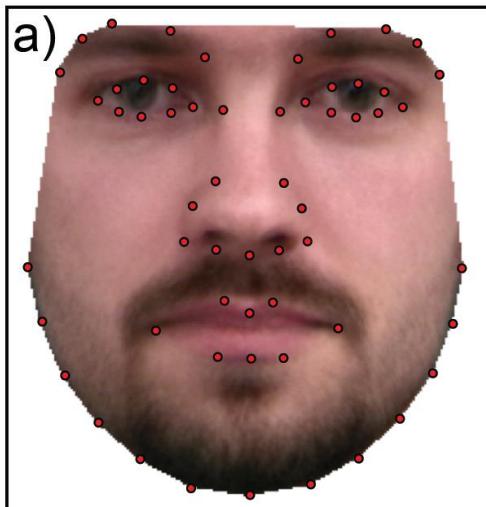


Intensity model



Shape and
intensity

Warping images



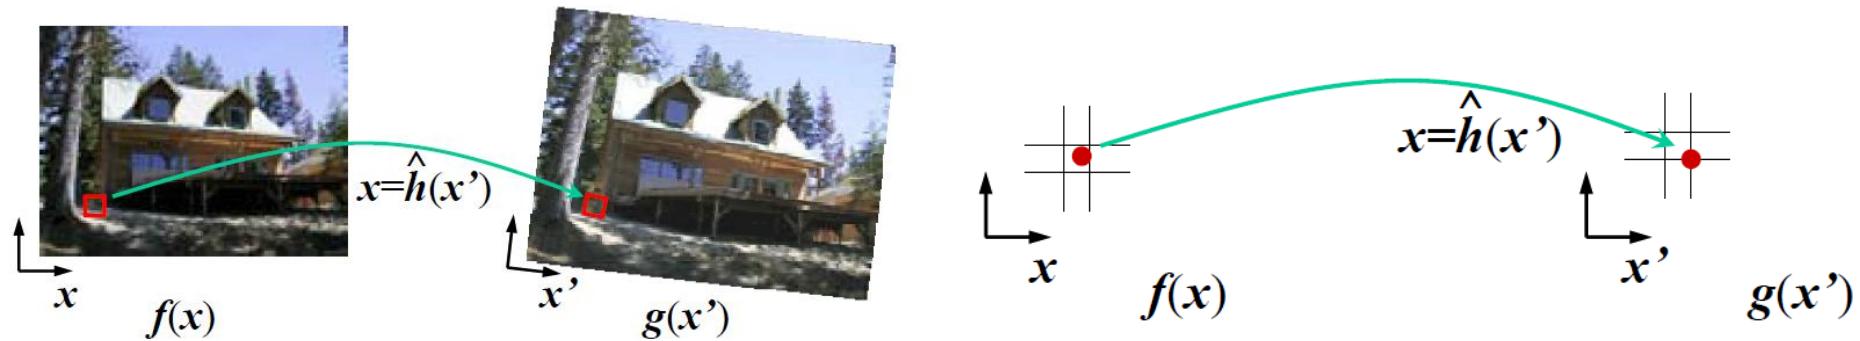
Piecewise affine transformation

Triangulate image points using Delaunay triangulation.

Image in each triangle is warped by an affine transformation.

Warping

- GPU/OpenGL rendering
- Inverse map with interpolation:



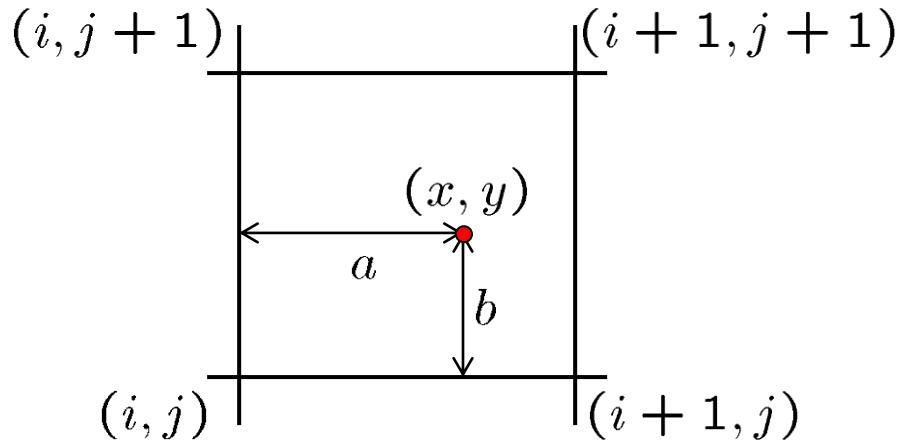
procedure *inverseWarp*(f, h , **out** g):

For every pixel x' in $g(x')$

1. Compute the source location $x = \hat{h}(x')$
2. Resample $f(x)$ at location x and copy to $g(x')$

Bilinear interpolation

Sampling at $f(x, y)$:



$$\begin{aligned} f(x, y) = & (1 - a)(1 - b) f[i, j] \\ & + a(1 - b) f[i + 1, j] \\ & + ab f[i + 1, j + 1] \\ & + (1 - a)b f[i, j + 1] \end{aligned}$$

Learning

$$Pr(\mathbf{h}_i) = \text{Norm}_{\mathbf{h}_i}[\mathbf{0}, \mathbf{I}]$$

$$Pr(\mathbf{w}_i | \mathbf{h}_i) = \text{Norm}_{\mathbf{w}_i}[\boldsymbol{\mu}_w + \boldsymbol{\Phi}_w \mathbf{h}_i, \sigma_w^2 \mathbf{I}]$$

$$Pr(\mathbf{x}_i | \mathbf{w}_i, \mathbf{h}_i) = \text{Norm}_{\mathbf{x}_i}[\text{warp}[\boldsymbol{\mu}_x + \boldsymbol{\Phi}_x \mathbf{h}_i, \mathbf{w}_i, \boldsymbol{\Psi}_i], \sigma_x^2 \mathbf{I}]$$

Goal is to learn parameters :

Problem $\{\boldsymbol{\mu}_w, \boldsymbol{\Phi}_w, \sigma_w^2, \boldsymbol{\mu}_x, \boldsymbol{\Phi}_x, \sigma_x^2\}$

- We are given the transformed landmark points
- We are given the warped and transformed images

Solution

- Use Procrustes analysis to un-transform landmark points
- Warp observed images to template shape

Learning

Now have aligned landmark points \mathbf{w} , and aligned images \mathbf{x} , we can learn the simpler model:

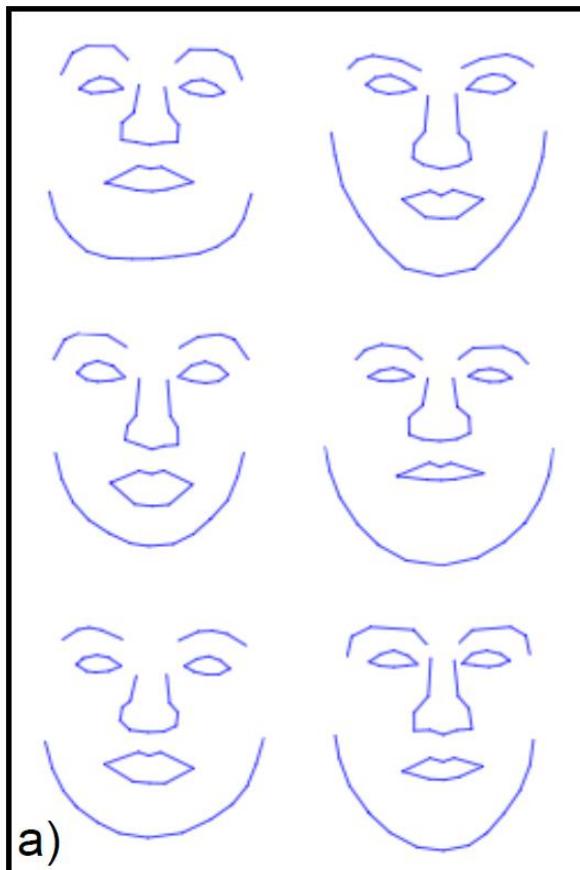
$$\begin{aligned} Pr(\mathbf{h}_i) &= \text{Norm}_{\mathbf{h}_i}[\mathbf{0}, \mathbf{I}] \\ Pr(\mathbf{w}_i|\mathbf{h}_i) &= \text{Norm}_{\mathbf{w}_i}[\boldsymbol{\mu}_w + \boldsymbol{\Phi}_w \mathbf{h}_i, \sigma_w^2 \mathbf{I}] \\ Pr(\mathbf{x}_i|\mathbf{h}_i) &= \text{Norm}_{\mathbf{x}_i}[\boldsymbol{\mu}_x + \boldsymbol{\Phi}_x \mathbf{h}_i, \sigma_x^2 \mathbf{I}]. \end{aligned}$$

Can write generative equation as:

$$\begin{bmatrix} \mathbf{w}_i \\ \mathbf{x}_i \end{bmatrix} = \begin{bmatrix} \boldsymbol{\mu}_w \\ \boldsymbol{\mu}_x \end{bmatrix} + \begin{bmatrix} \boldsymbol{\Phi}_w \\ \boldsymbol{\Phi}_x \end{bmatrix} \mathbf{h}_i + \begin{bmatrix} \boldsymbol{\epsilon}_{wi} \\ \boldsymbol{\epsilon}_{xi} \end{bmatrix}$$

Has the form: $\mathbf{x}' = \boldsymbol{\mu}' + \boldsymbol{\Phi}' \mathbf{h} + \boldsymbol{\epsilon}'$

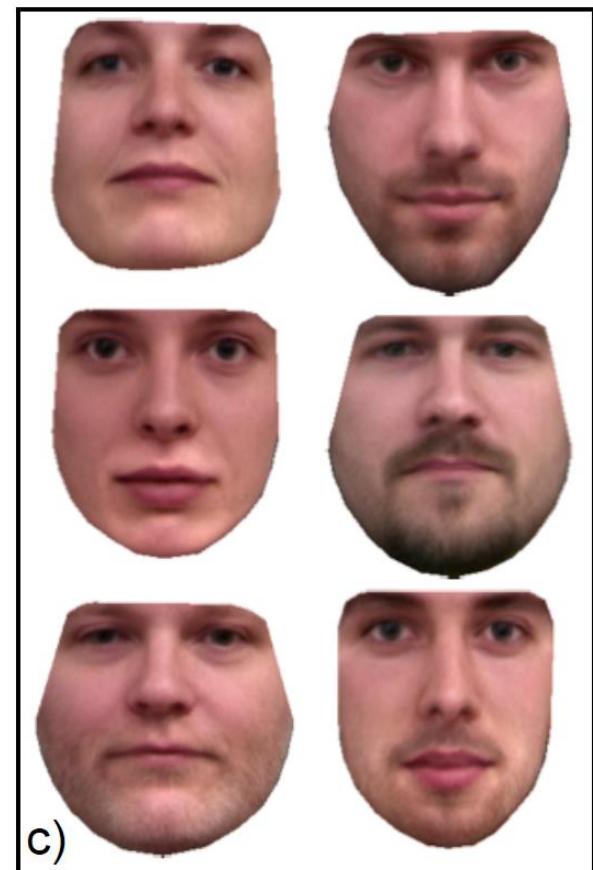
Shape and appearance model



Shape
model



Intensity model



Shape and
intensity

Inference

Likelihood of observed intensities

$$Pr(\mathbf{x}|\mathbf{h}) = \text{Norm}_{\mathbf{x}}[\mathbf{warp}[\boldsymbol{\mu}_x + \boldsymbol{\Phi}_x \mathbf{h}, \boldsymbol{\mu} + \boldsymbol{\Phi} \mathbf{h}, \boldsymbol{\Psi}], \sigma_x^2 \mathbf{I}]$$

To fit the model use maximum likelihood

$$\begin{aligned}\hat{\mathbf{h}}, \hat{\boldsymbol{\Psi}} &= \underset{\mathbf{h}, \boldsymbol{\Psi}}{\text{argmax}} [\log [Pr(\mathbf{x}|\mathbf{h})]] \\ &= \underset{\mathbf{h}, \boldsymbol{\Psi}}{\text{argmin}} [(\mathbf{x} - \mathbf{warp}[\boldsymbol{\mu}_x + \boldsymbol{\Phi}_x \mathbf{h}, \boldsymbol{\mu}_w + \boldsymbol{\Phi}_w \mathbf{h}, \boldsymbol{\Psi}])^T \\ &\quad (\mathbf{x} - \mathbf{warp}[\boldsymbol{\mu}_x + \boldsymbol{\Phi}_x \mathbf{h}, \boldsymbol{\mu}_w + \boldsymbol{\Phi}_w \mathbf{h}, \boldsymbol{\Psi}])]\end{aligned}$$

This has the least squares form

$$f[\boldsymbol{\theta}] = \mathbf{z}[\boldsymbol{\theta}]^T \mathbf{z}[\boldsymbol{\theta}]$$

Inference

This has the least squares form

$$f[\theta] = \mathbf{z}[\theta]^T \mathbf{z}[\theta]$$

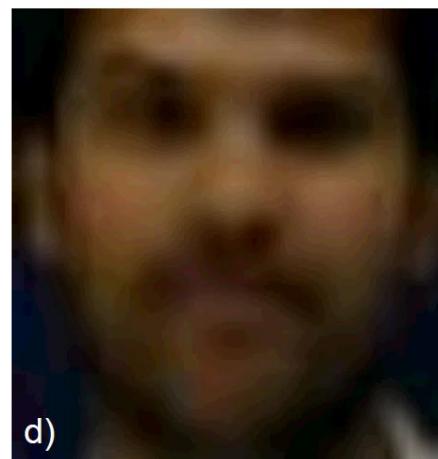
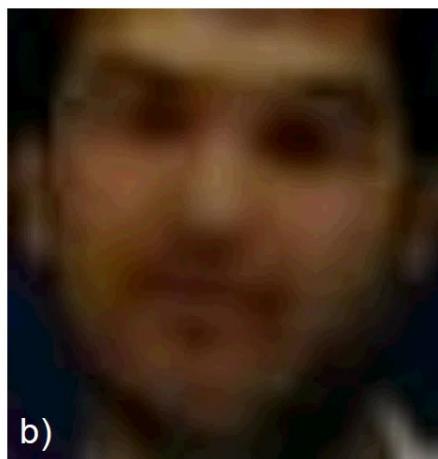
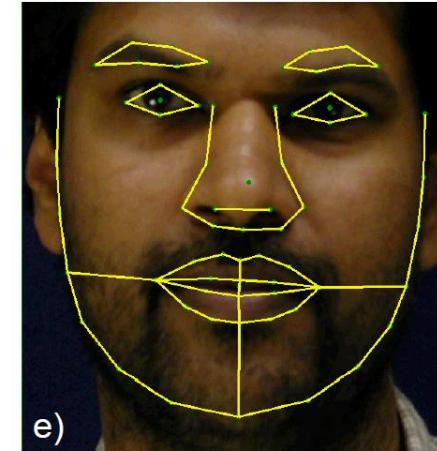
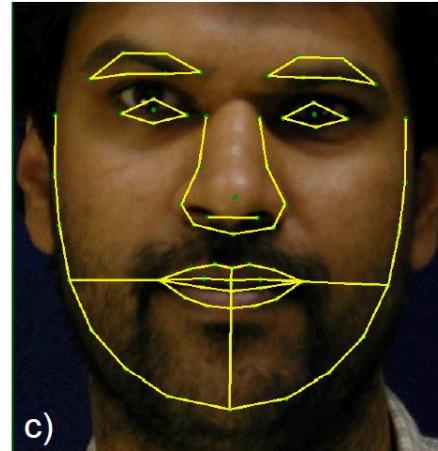
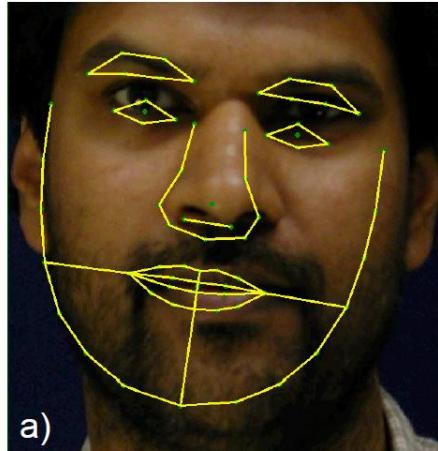
Use Gauss-Newton method or similar

$$\theta^{[t]} = \theta^{[t-1]} + \lambda(\mathbf{J}^T \mathbf{J})^{-1} \frac{\partial f}{\partial \theta}$$

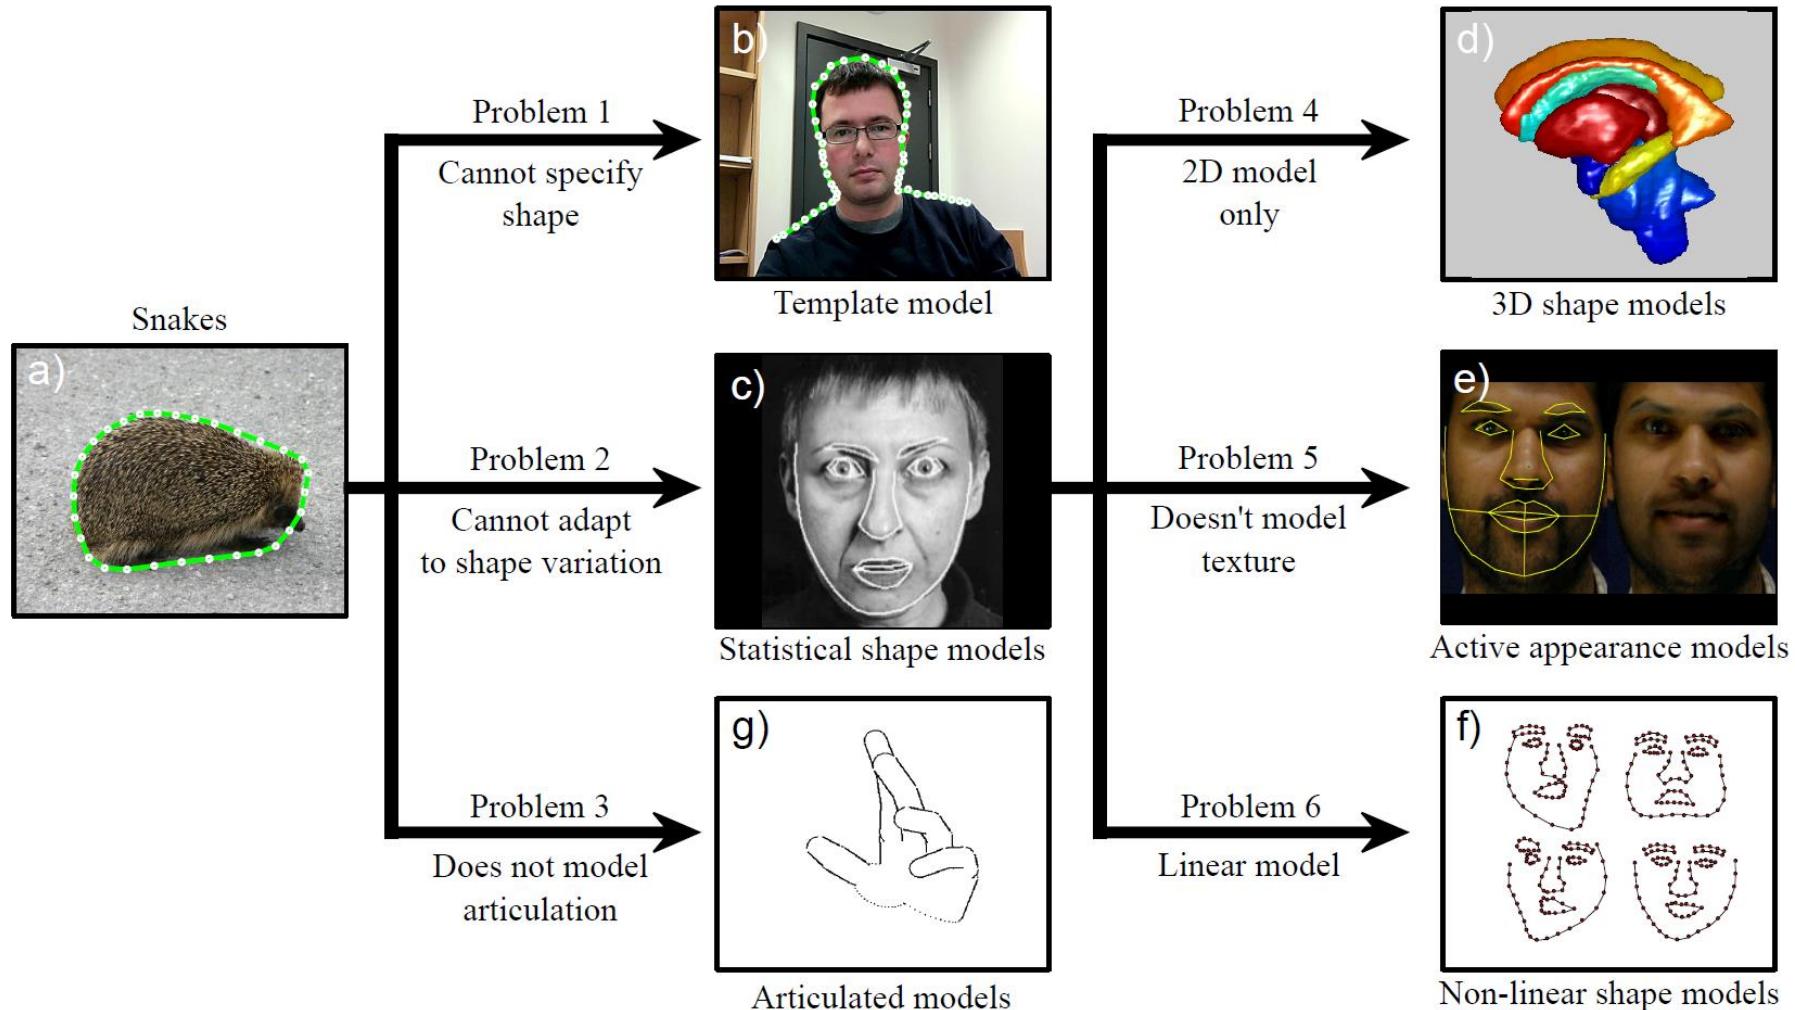
Where the Jacobian \mathbf{J} is a matrix with elements

$$J_{mn} = \frac{\partial z_m}{\partial \theta_n}$$

Statistical models for shape and appearance



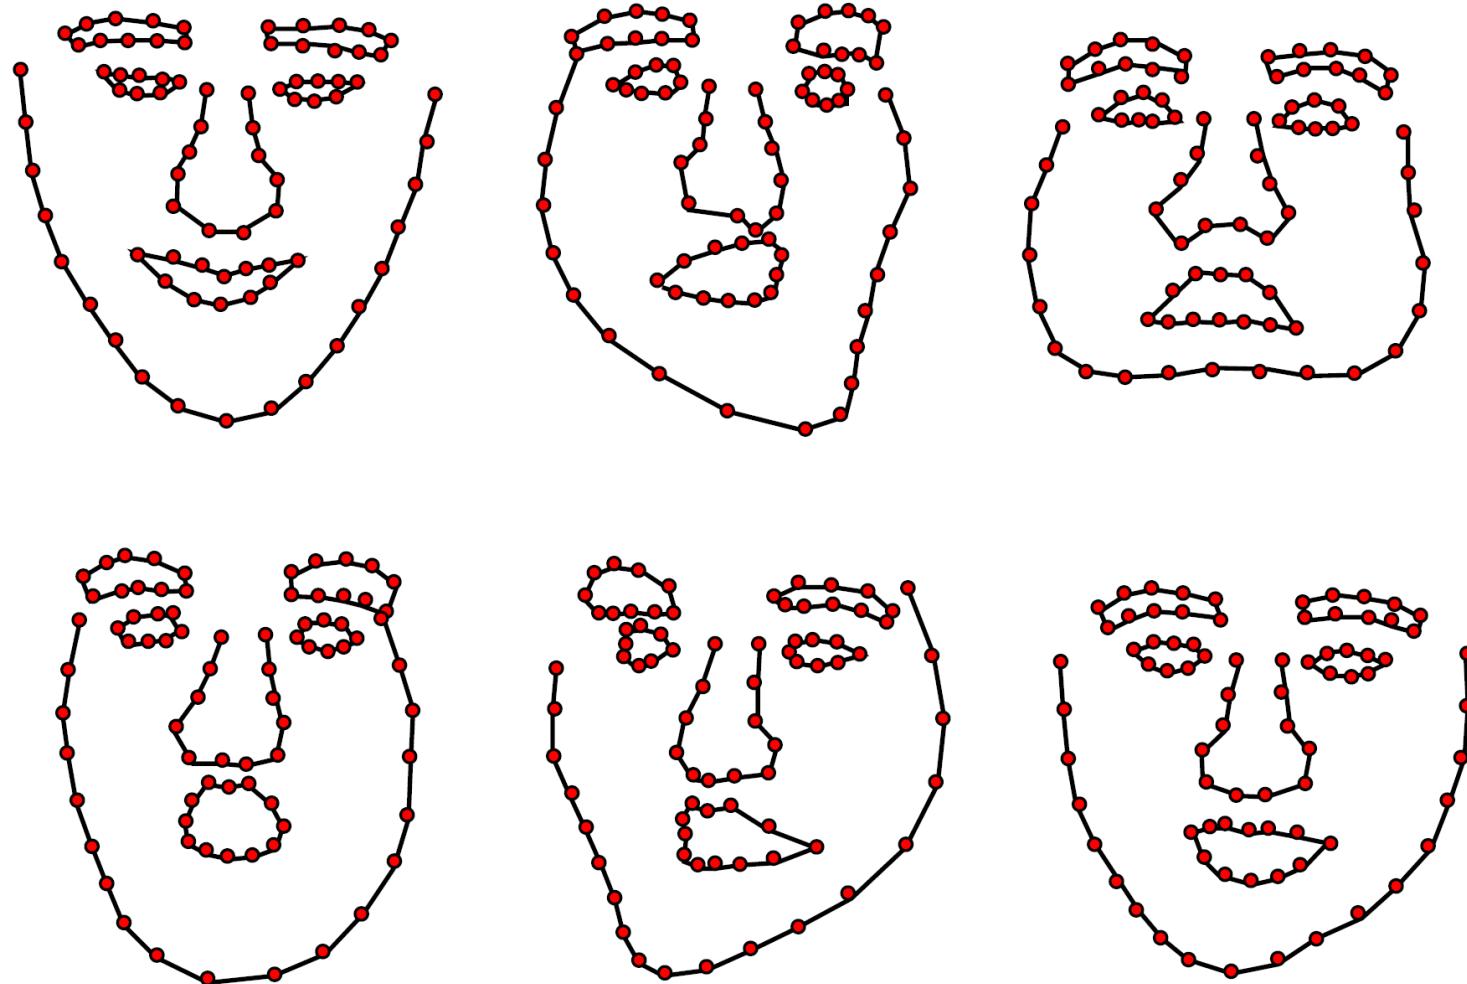
Relationships between models



Non-linear models

- The shape and appearance models that we have studied so far are based on the normal distribution
- But more complex shapes might need more complex distributions
 - Could use mixture of PPCAs or similar
 - Or use a non-linear subspace model
- Example: Gaussian process latent variable model (GPLVM) – Computer Vision II

GPLVM Shape models



3D morphable models

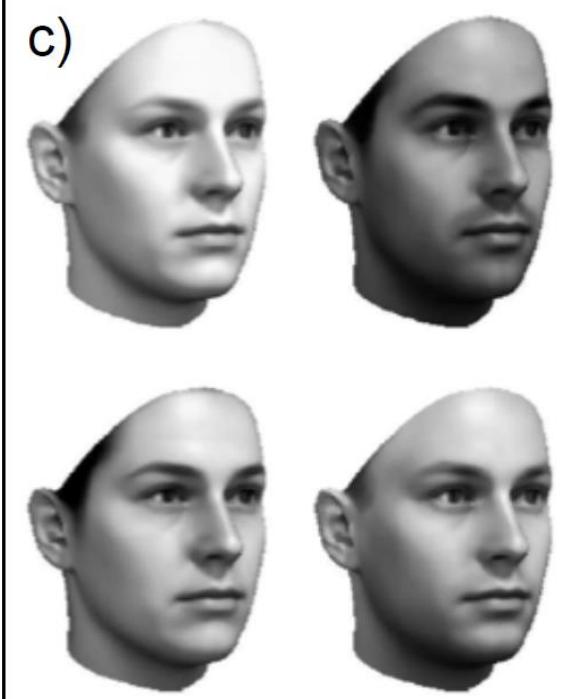
a)



b)

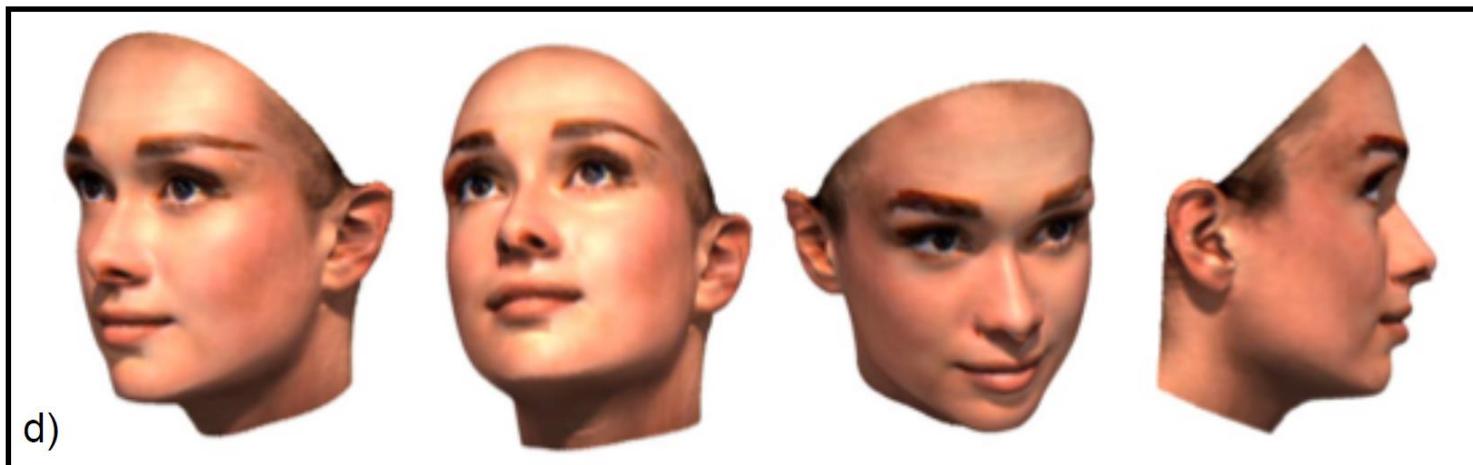
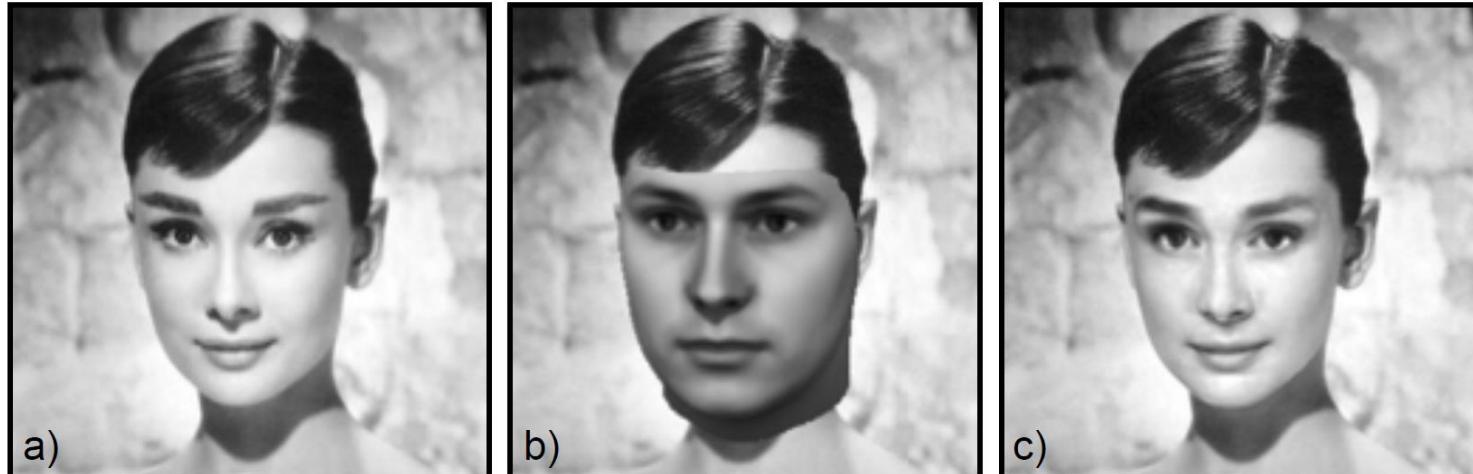


c)



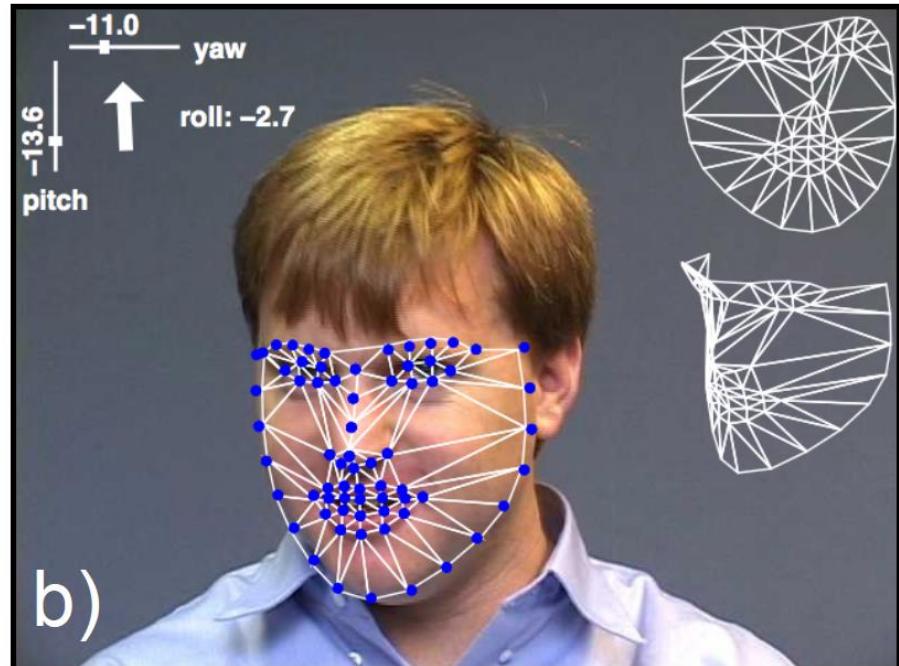
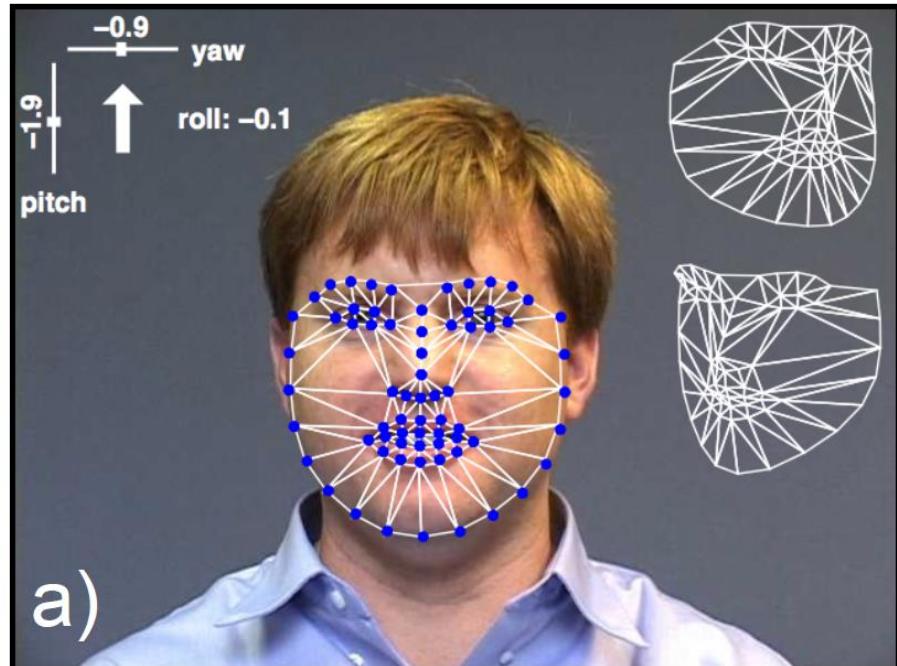
Adapted from Blanz & Vetter (2003). ©2003 IEEE

3D morphable models



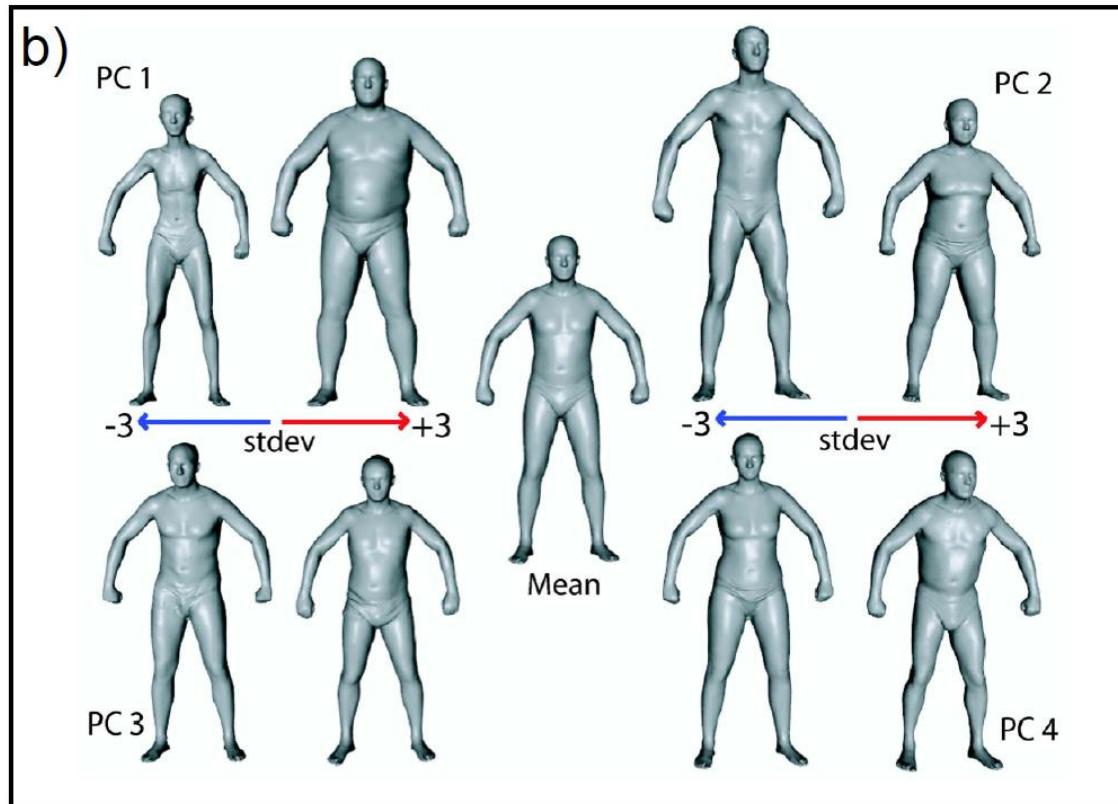
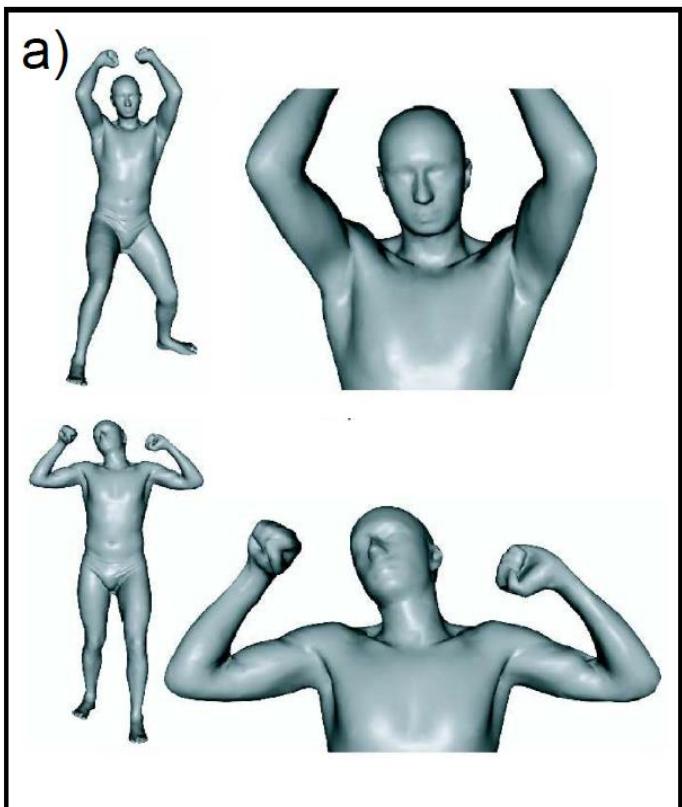
Adapted from Blanz & Vetter (1999).

3D morphable models



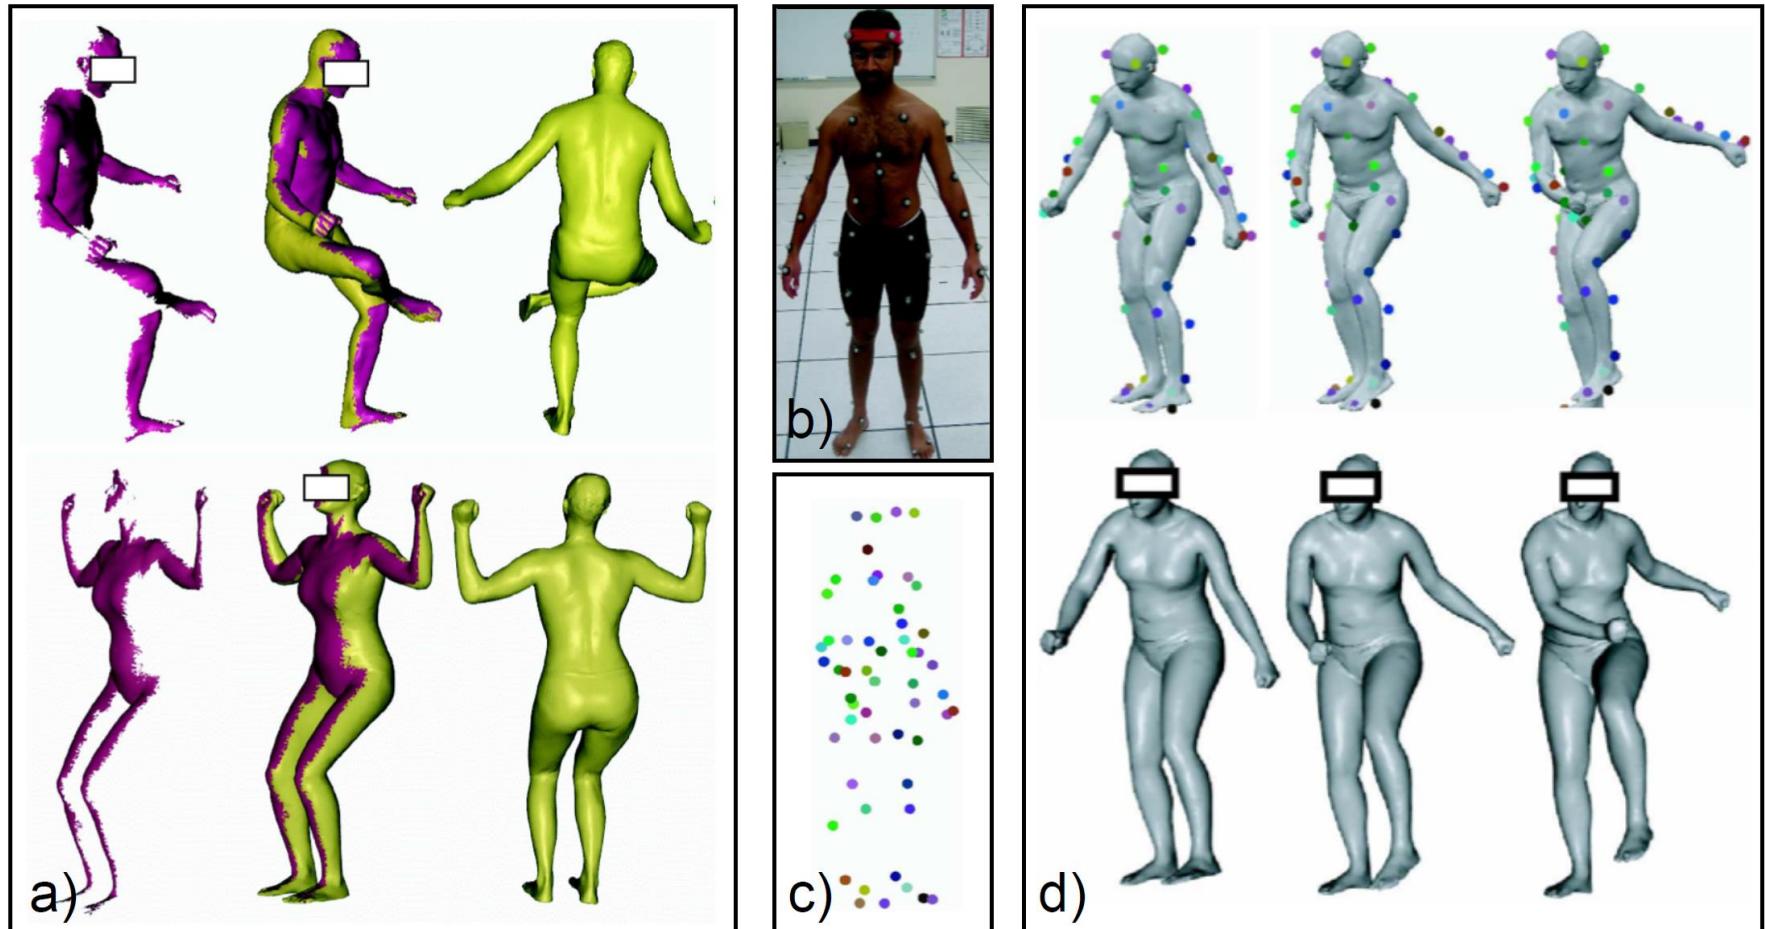
Matthews *et al.* (2007). ©2007 Springer.

3D body model



Adapted from Anguelov *et al.* (2005). ©2005 ACM

3D body model applications

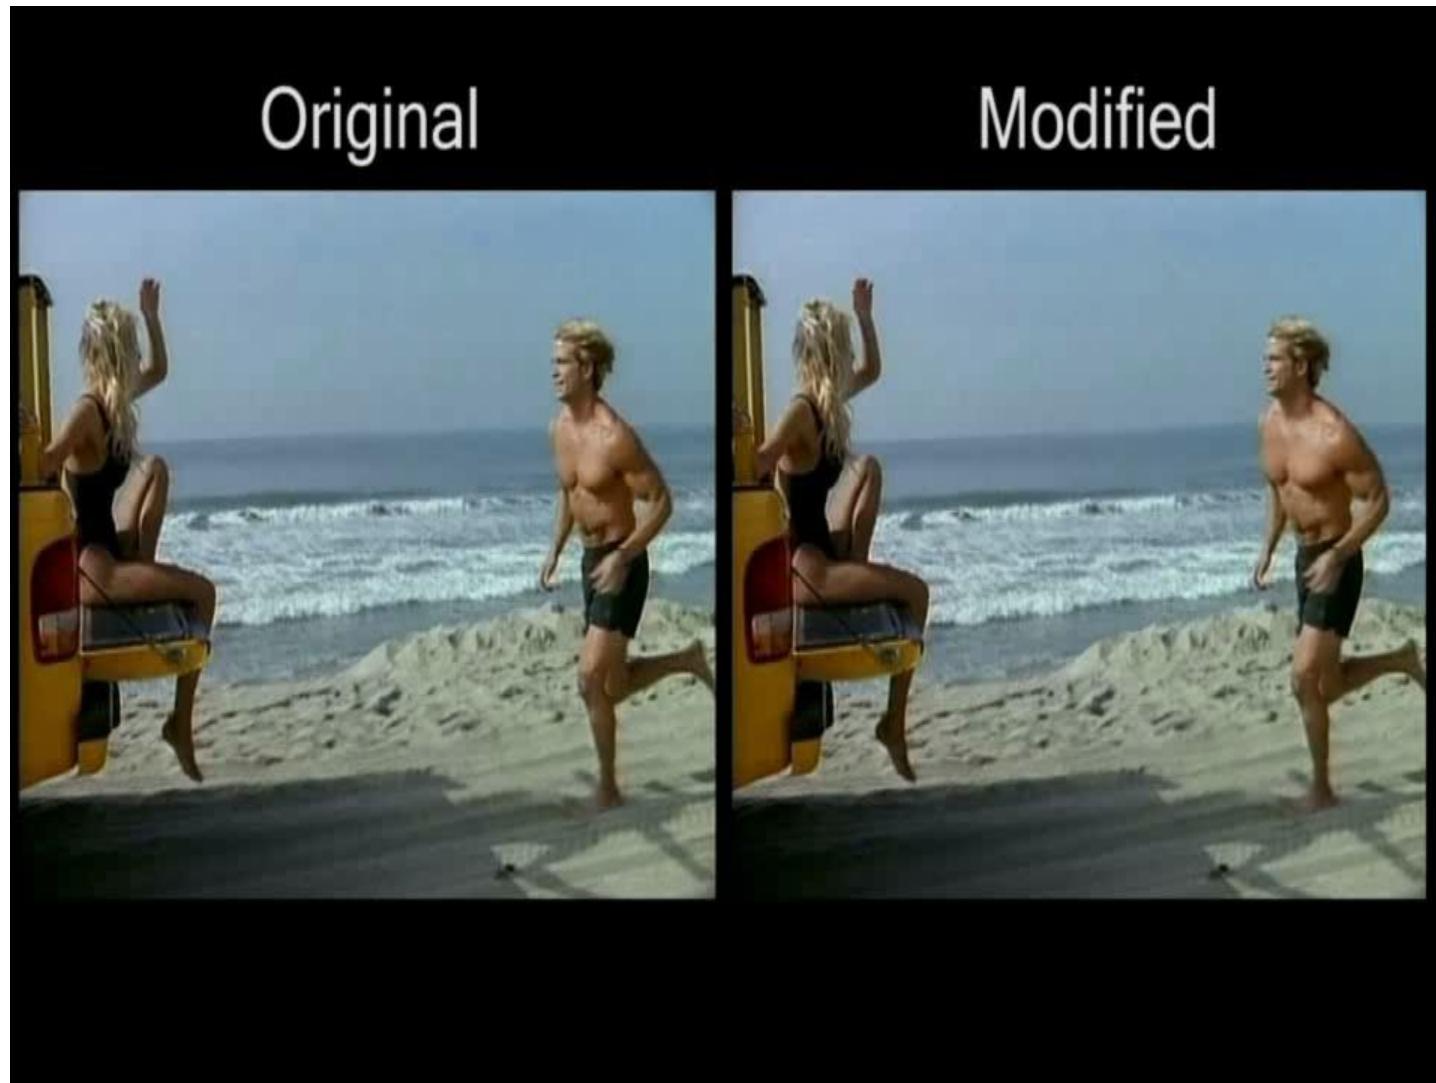


Adapted from Anguelov *et al.* (2005). ©2005 ACM

Shape model



Video editing



Jain et al. MovieReshape: Tracking and Reshaping of Humans in Videos. ToG 2010.

UNIVERSITÄT BONN

