


The logo of the University of Bonn, featuring a blue square with a white curved line and a grey square.

UNIVERSITÄT **BONN**


Juergen Gall

RANSAC and Image Alignment
MA-INF 2201 - Computer Vision
WS24/25

Motivation for feature-based alignment: Image mosaics




Robust feature-based alignment




- Extract features

Robust feature-based alignment




- Extract features
- Compute *putative matches*

Robust feature-based alignment



- Extract features
- Compute *putative matches*
- Loop:
 - *Hypothesize* transformation T (small group of putative matches that are related by T)

Robust feature-based alignment




- Extract features
- Compute *putative matches*
- Loop:
 - *Hypothesize* transformation T (small group of putative matches that are related by T)
 - *Verify* transformation (search for other matches consistent with T)

Robust feature-based alignment



- Extract features
- Compute *putative matches*
- Loop:
 - *Hypothesize* transformation T (small group of putative matches that are related by T)
 - *Verify* transformation (search for other matches consistent with T)

Recall: Fitting an affine transformation



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\begin{bmatrix} x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$

Recall: Least squares problem

A $m \times n$ matrix of rank n :

$$Ax = b$$

Solution with pseudo-inverse:

$$A^T (Ax - b) = 0$$

$$x = (A^T A)^{-1} A^T b$$


Pseudo-inverse \mathbf{A}^+

Solution with SVD:


- SVD: $\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{V}^T$; $\mathbf{A}^+ = \mathbf{V}\mathbf{D}^{-1}\mathbf{U}^T$
 - $\mathbf{b}' = \mathbf{U}^T \mathbf{b}$
 - $y_i = b'_i / D_{ii}$
 - $\mathbf{x} = \mathbf{V}\mathbf{y}$
- There are different ways to solve the problem. The optimal method depends on size and sparseness of \mathbf{A}
- Specific C++ linear algebra libraries exist!

Outliers can hurt the quality of our parameter estimates, e.g.,


- an erroneous pair of matching points from two images
- an edge point that is noise, or doesn't belong to the line we are fitting.



Outliers affect least squares fit



Outliers affect least squares fit




RANdom Sample Consensus (RANSAC)

RANSAC loop:


1. Randomly select a *seed group* of points on which to base transformation estimate (e.g., a group of matches)
2. Compute transformation from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute estimate of transformation on all of the inliers

Keep the transformation with the largest number of inliers

RANSAC for line fitting example




RANSAC for line fitting example




Least-squares fit

RANSAC for line fitting example




1. Randomly select minimal subset of points

RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model

RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model
3. **Compute error function**

RANSAC for line fitting example




1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. **Select points consistent with model**

RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop


RANSAC for line fitting example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop


RANSAC for line fitting example

Uncontaminated sample



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

RANSAC for line fitting example



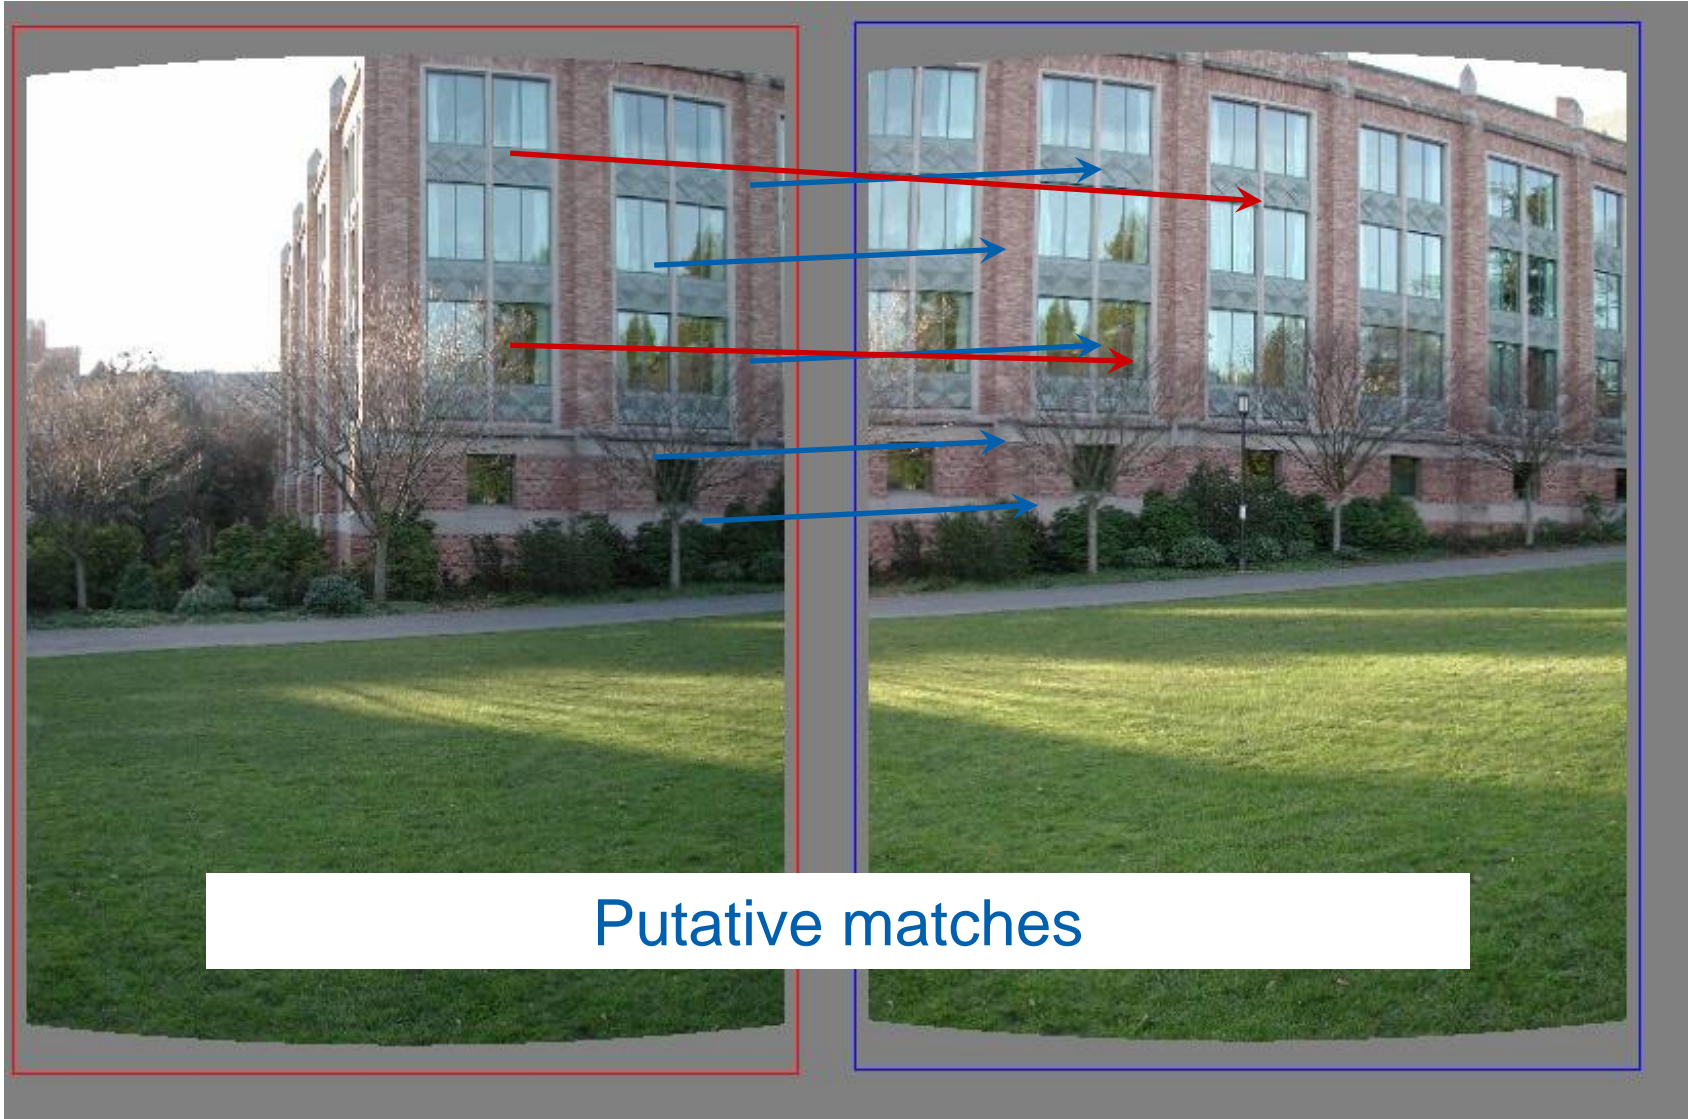
1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. **Repeat**
hypothesize-and-verify loop

RANSAC for line fitting

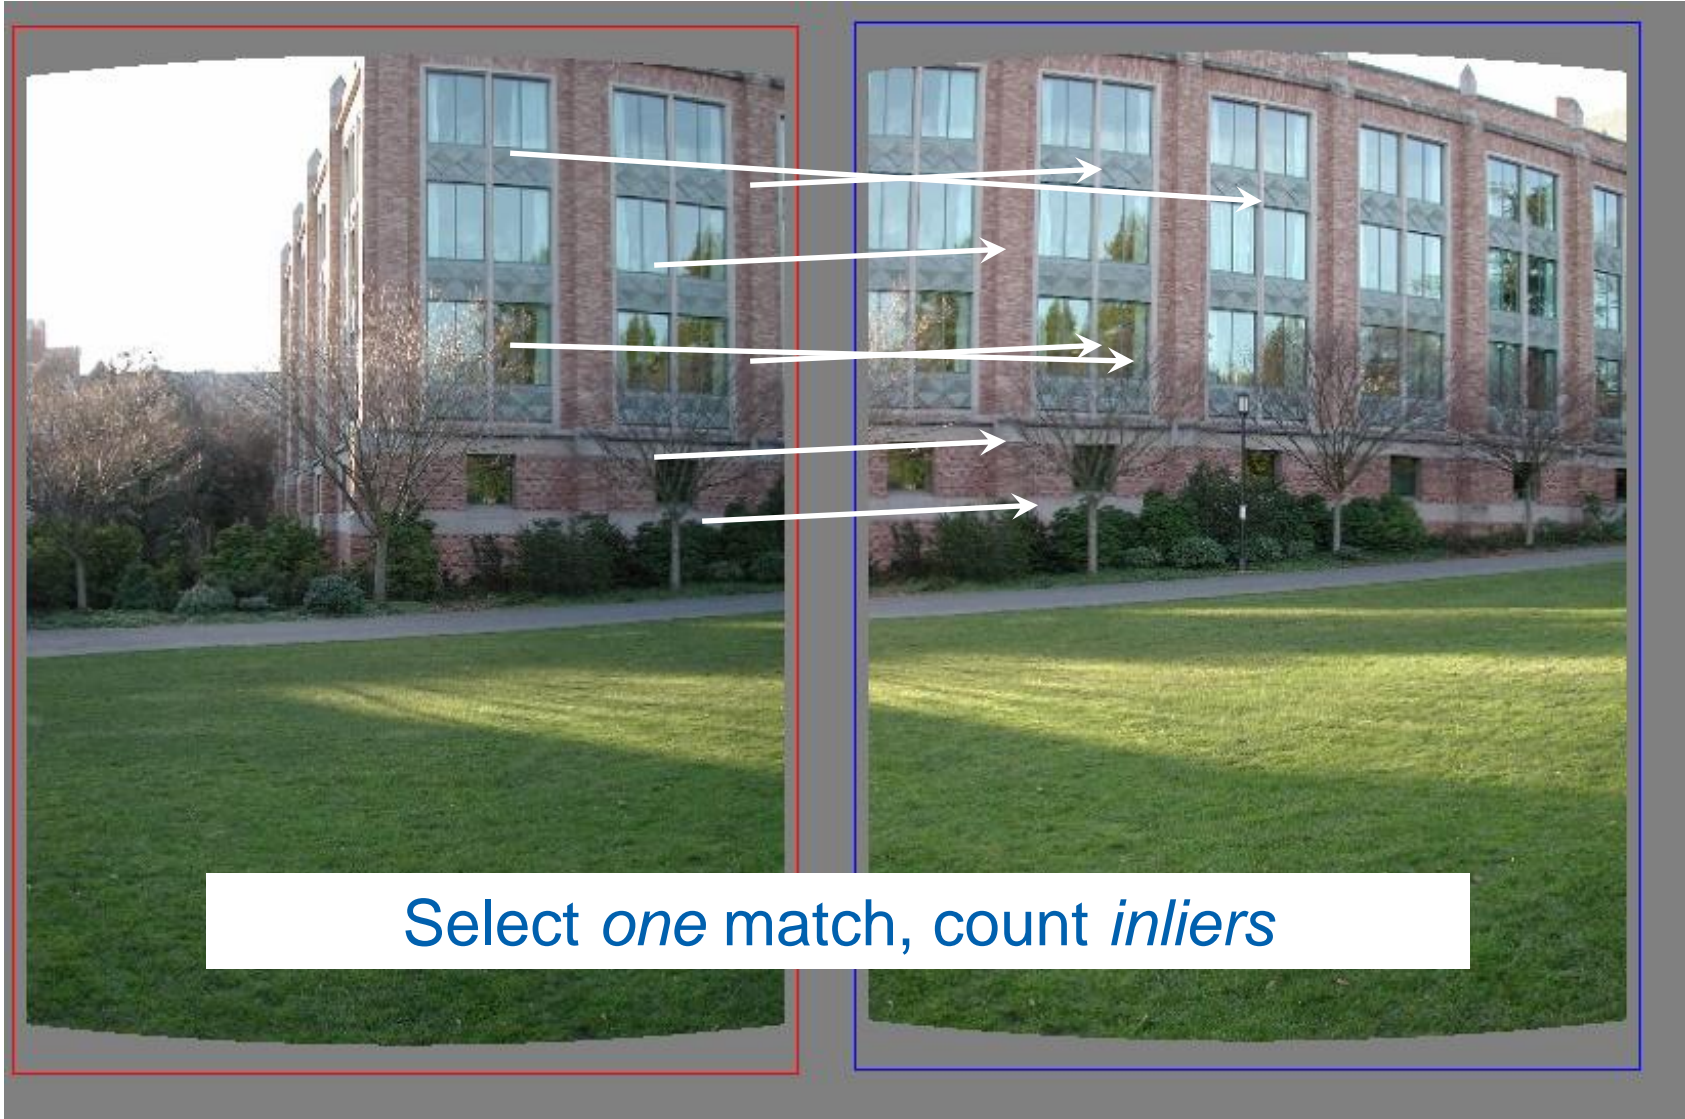
Repeat **N** times:

- Draw **s** points uniformly at random
- Fit line to these **s** points
- Find inliers to this line among the remaining points (i.e., points whose distance from the line is less than **t**)
- If there are **d** or more inliers, accept the line and refit using all inliers

RANSAC example: Translation

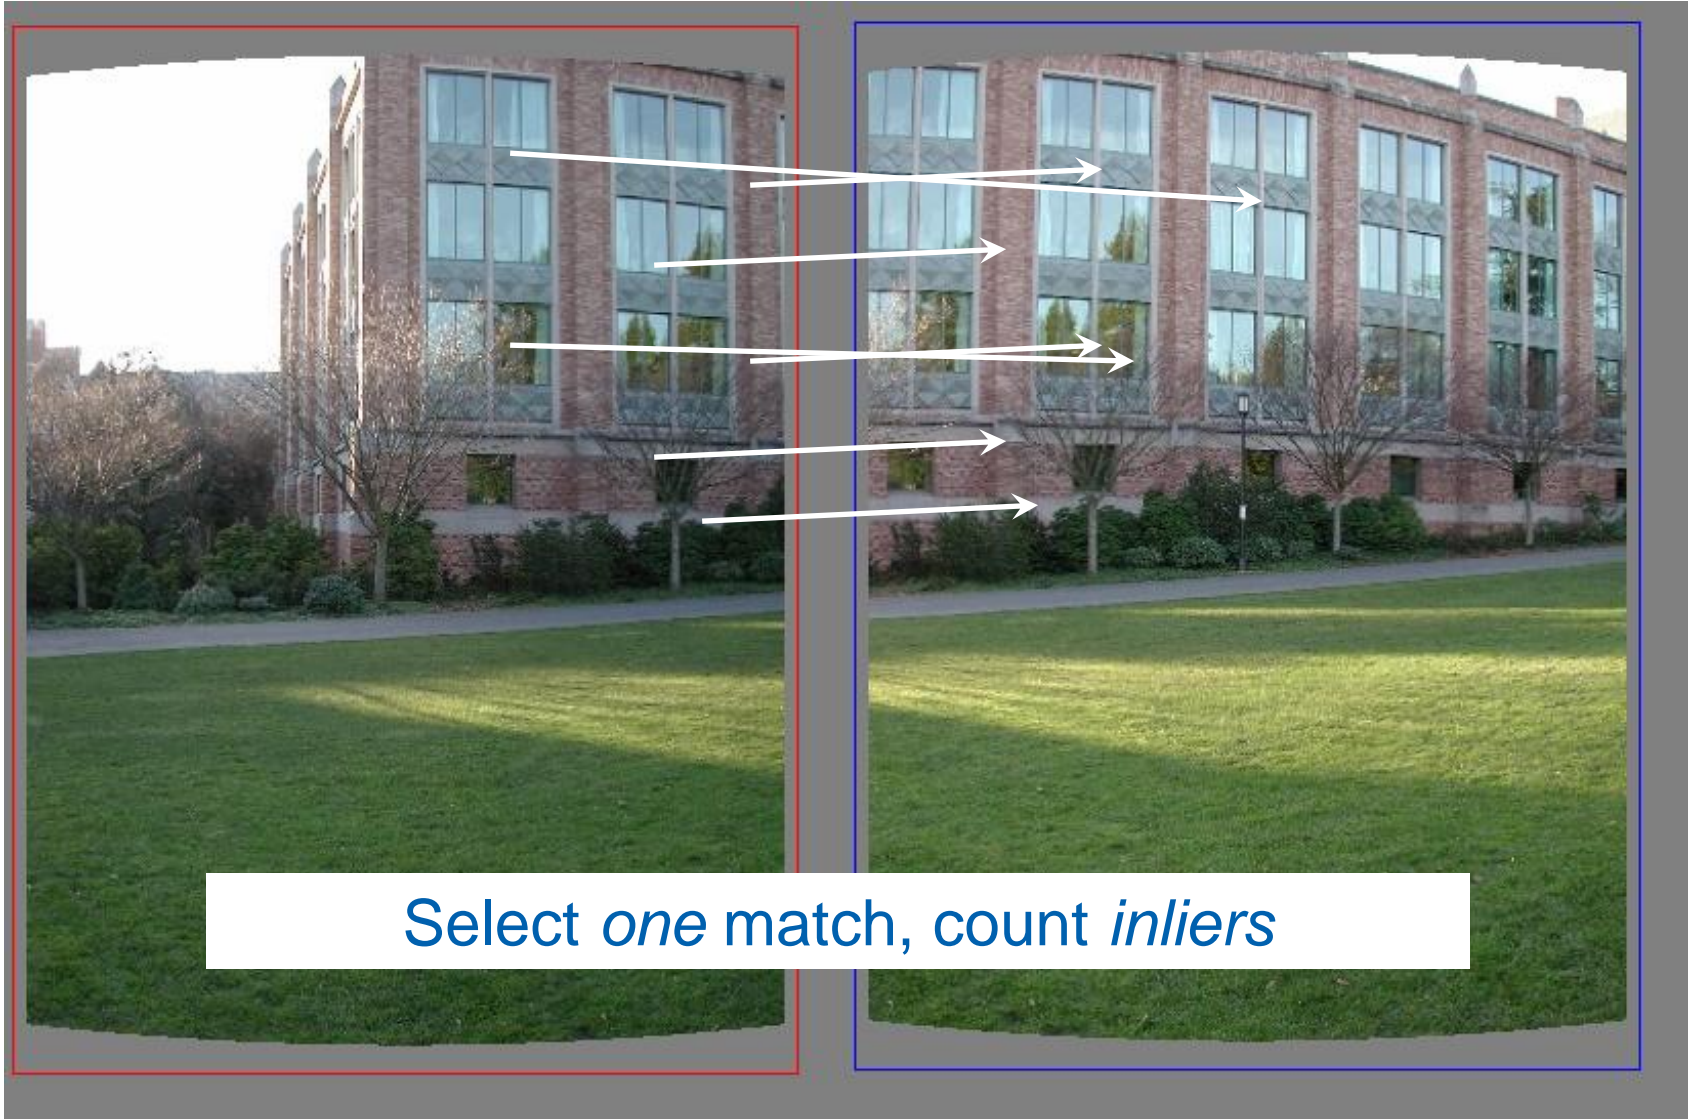


RANSAC example: Translation



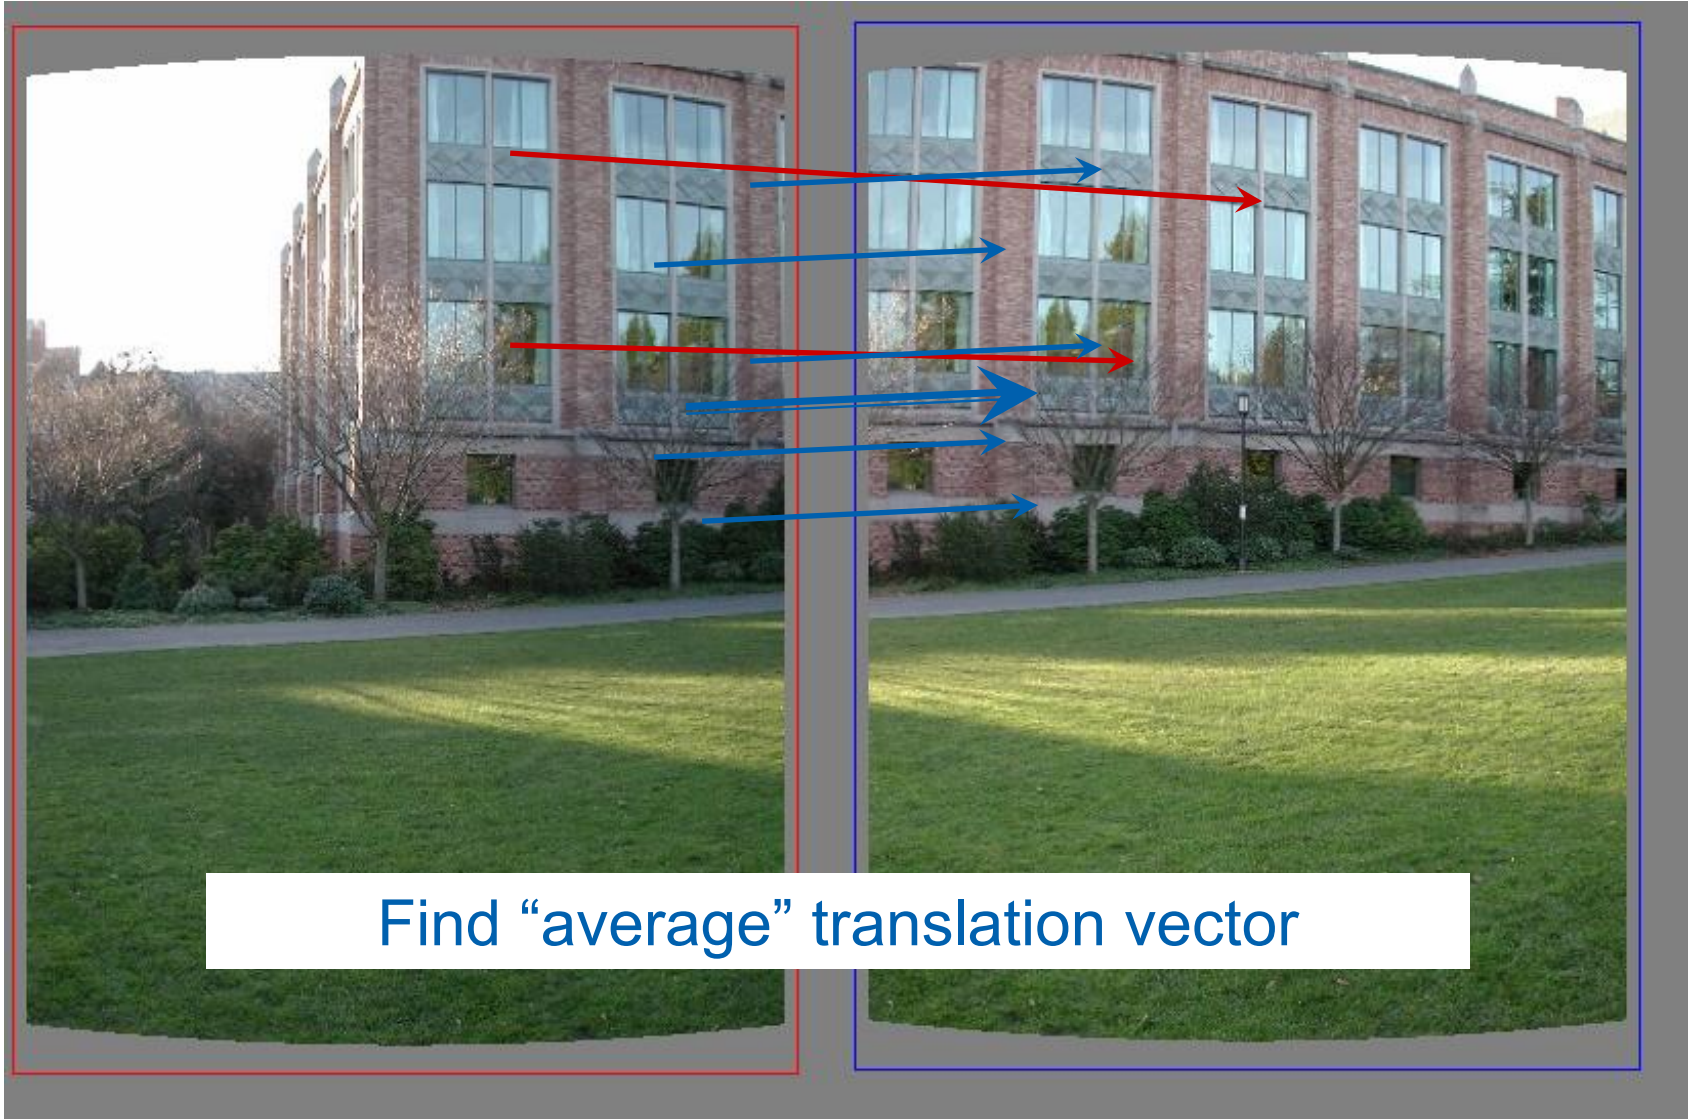
Source: Rick Szeliski

RANSAC example: Translation



Source: Rick Szeliski

RANSAC example: Translation



How Many Samples?

On average

N ... number of point
 I ... number of inliers
 m ... size of the sample

$$P(\text{good}) = \frac{\binom{I}{m}}{\binom{N}{m}} = \prod_{j=0}^{m-1} \frac{I - j}{N - j}$$

mean time before the success

$$E(k) = 1 / P(\text{good})$$

How Many Samples?

With confidence p

How large k ?

... to hit at least one pair of points on the line l with probability larger than p (0.95)

Equivalently

... the probability of not hitting any pair of points on l is $\leq 1 - p$

How Many Samples?

With confidence p

N ... number of point
 I ... number of inliers
 m ... size of the sample

$$P(\text{good}) = \frac{\binom{I}{m}}{\binom{N}{m}} = \prod_{j=0}^{m-1} \frac{I - j}{N - j}$$

$$P(\text{bad}) = 1 - P(\text{good})$$

$$P(\text{bad } k \text{ times}) = (1 - P(\text{good}))^k$$

How Many Samples?

With confidence p

$$P(\text{bad } k \text{ times}) = \left(1 - P(\text{good})\right)^k \leq 1 - p$$

$$k \log \left(1 - P(\text{good})\right) \leq \log(1 - p)$$

$$k \geq \log(1 - p) / \log \left(1 - P(\text{good})\right)$$

$$P(\text{good}) = \frac{\binom{I}{m}}{\binom{N}{m}} = \prod_{j=0}^{m-1} \frac{I - j}{N - j}$$

How Many Samples


// N [%]

Size of the sample m

	15%	20%	30%	40%	50%	70%
2	132	73	32	17	10	4
4	5916	1871	368	116	46	11
7	$1.75 \cdot 10^6$	$2.34 \cdot 10^5$	$1.37 \cdot 10^4$	1827	382	35
8	$1.17 \cdot 10^7$	$1.17 \cdot 10^6$	$4.57 \cdot 10^4$	4570	765	50
12	$2.31 \cdot 10^{10}$	$7.31 \cdot 10^8$	$5.64 \cdot 10^6$	$1.79 \cdot 10^5$	$1.23 \cdot 10^4$	215
18	$2.08 \cdot 10^{15}$	$1.14 \cdot 10^{13}$	$7.73 \cdot 10^9$	$4.36 \cdot 10^7$	$7.85 \cdot 10^5$	1838
30	∞	∞	$1.35 \cdot 10^{16}$	$2.60 \cdot 10^{12}$	$3.22 \cdot 10^9$	$1.33 \cdot 10^5$
40	∞	∞	∞	$2.70 \cdot 10^{16}$	$3.29 \cdot 10^{12}$	$4.71 \cdot 10^6$

RANSAC pros and cons


- Pros
 - Simple and general
 - Applicable to many different problems
 - Often works well in practice
- Cons
 - Doesn't work well for low inlier ratios (too many iterations, or can fail completely)
 - Can't always get a good initialization of the model based on the minimum number of samples



Recall: 2D Affine Transformations

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- Affine transformations are combinations of
 - ...
 - Linear transformations, and
 - Translations
- Parallel lines remain parallel




Projective Transformations


Projective transformations:

- Affine transformations, and
- Projective warps

$$\begin{bmatrix} x' \\ y' \\ w' \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

Parallel lines do not necessarily remain parallel





...



image from S. Seitz

Obtain a wider angle view by combining multiple images.

How to stitch together a panorama (a.k.a. mosaic)?

Basic Procedure


- Take a sequence of images from the same position
 - Rotate the camera about its optical center
- Compute transformation between second image and first
- Transform the second image to overlap with the first
- Blend the two together to create a mosaic
- (If there are more images, repeat)

...but **wait**, why should this work at all?

- What about the 3D geometry of the scene?
- Why aren't we using it?

Pinhole camera

Pinhole camera is a simple model to approximate imaging process, perspective **projection**.



If we treat pinhole as a point, only one ray from any given point can enter the camera.

Mosaics




image from S. Seitz

Obtain a wider angle view by combining multiple images.

Mosaics: generating synthetic views





Image reprojection



The mosaic has a natural interpretation in 3D

- The images are reprojected onto a common plane
- The mosaic is formed on this plane
- Mosaic is a *synthetic wide-angle camera*

Image reprojection

- How to relate two images from the same camera center? How to map a pixel from PP1 to PP2?
- Cast a ray through each pixel in PP1
- Draw the pixel where that ray intersects PP2
- Rather than thinking of this as a 3D reprojection, think of it as a 2D image warp from one image to another.





Image reprojection: Homography

A projective transform is a mapping between any two PPs with the same center of projection


Called **Homography**

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} \begin{bmatrix} x \\ y \\ l \end{bmatrix}$$

$\mathbf{p}' \qquad \mathbf{H} \qquad \mathbf{p}$



Homography



To **compute** the homography given pairs of corresponding points in the images, we need to set up an equation where the parameters of **H** are the unknowns...

Solving for homographies

$$\begin{bmatrix} {}^w x'_i \\ {}^w y'_i \\ {}^w 1 \end{bmatrix} \cong \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

Solving for homographies

$$\begin{array}{c}
 \begin{bmatrix}
 x_1 & y_1 & 1 & 0 & 0 & 0 & -x'_1 x_1 & -x'_1 y_1 & -x'_1 \\
 0 & 0 & 0 & x_1 & y_1 & 1 & -y'_1 x_1 & -y'_1 y_1 & -y'_1 \\
 & & & & & & \vdots & & \\
 x_n & y_n & 1 & 0 & 0 & 0 & -x'_n x_n & -x'_n y_n & -x'_n \\
 0 & 0 & 0 & x_n & y_n & 1 & -y'_n x_n & -y'_n y_n & -y'_n
 \end{bmatrix}
 \begin{bmatrix}
 h_{00} \\
 h_{01} \\
 h_{02} \\
 h_{10} \\
 h_{11} \\
 h_{12} \\
 h_{20} \\
 h_{21} \\
 h_{22}
 \end{bmatrix}
 =
 \begin{bmatrix}
 0 \\
 0 \\
 0 \\
 \vdots \\
 0 \\
 0
 \end{bmatrix}
 \\
 \mathbf{A} \qquad \qquad \mathbf{h} \qquad \qquad \mathbf{0} \\
 2n \times 9 \qquad \qquad 9 \qquad \qquad 2n
 \end{array}$$

Defines a least squares problem: minimize $\|\mathbf{A}\mathbf{h} - \mathbf{0}\|^2$

- Since \mathbf{h} is only defined up to scale, solve for unit vector $\hat{\mathbf{h}}$
- Solution: $\hat{\mathbf{h}}$ = eigenvector of $\mathbf{A}^T \mathbf{A}$ with smallest eigenvalue

Homography

 (x, y)


$$\begin{pmatrix} wx'/w & wy'/w \end{pmatrix} = (x', y')$$

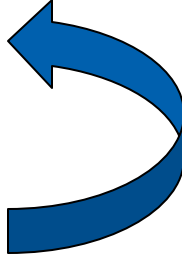
To **apply** a given homography **H**

- Compute $\mathbf{p}' = \mathbf{H}\mathbf{p}$ (regular matrix multiply)
- Convert \mathbf{p}' from homogeneous to image coordinates

$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$


$\mathbf{p}' \qquad \qquad \mathbf{H} \qquad \qquad \mathbf{p}$

RANSAC for estimating homography

- RANSAC loop:
 - 1. Select four feature pairs (at random)
 - 2. Compute homography H (exact)
 - 3. Compute *inliers* where $SSD(p_i', H p_i) < \varepsilon$
 - 4. Keep largest set of inliers
 - 5. Re-compute least-squares H estimate on all of the inliers
- 

Recall: Warping

- GPU/OpenGL rendering
- Inverse map with interpolation:



procedure *inverseWarp*(f, h , **out** g):

For every pixel x' in $g(x')$

1. Compute the source location $x = \hat{h}(x')$
2. Resample $f(x)$ at location x and copy to $g(x')$

Recap: How to stitch together a panorama (a.k.a. mosaic)?

Basic Procedure

- Take a sequence of images from the same position
 - Rotate the camera about its optical center
- Compute transformation (homography) between second image and first using corresponding points.
- Transform the second image to overlap with the first.
- Blend the two together to create a mosaic.
- (If there are more images, repeat)

Image warping with homographies




image plane in front




image plane below

Image rectification





Image warping with homographies






image plane in front




image plane below

black area
where no pixel
maps to

Image warping with homographies






image plane in front



black area
where no pixel
maps to



Analysing patterns and shapes

What is the shape of the b/w floor pattern?



Homography




The floor (enlarged)



Automatically
rectified floor

Analysing patterns and shapes

Automatic rectification



From Martin Kemp *The Science of Art*
(*manual reconstruction*)

Analysing patterns and shapes



Automatically rectified floor

St. Lucy Altarpiece, D. Veneziano

Analysing patterns and shapes


Automatic
rectification



From Martin Kemp, *The Science of Art*
(*manual reconstruction*)


Changing camera center

Does it still work?




Source: Alyosha Efros

Recall: same camera center




Can generate synthetic camera view
as long as it has **the same center of projection.**

...or planar scene




- PP3 is a projection plane of both centers of projection, so we are OK!

...or planar scene (or far away)



- PP3 is a projection plane of both centers of projection, so we are OK!
- This is how big aerial photographs are made



Map

Satellite

Hybrid

©2007 Google Kristen Grauman

The logo of the University of Bonn, featuring a blue square with a white curved line and a grey square.

UNIVERSITÄT **BONN**