

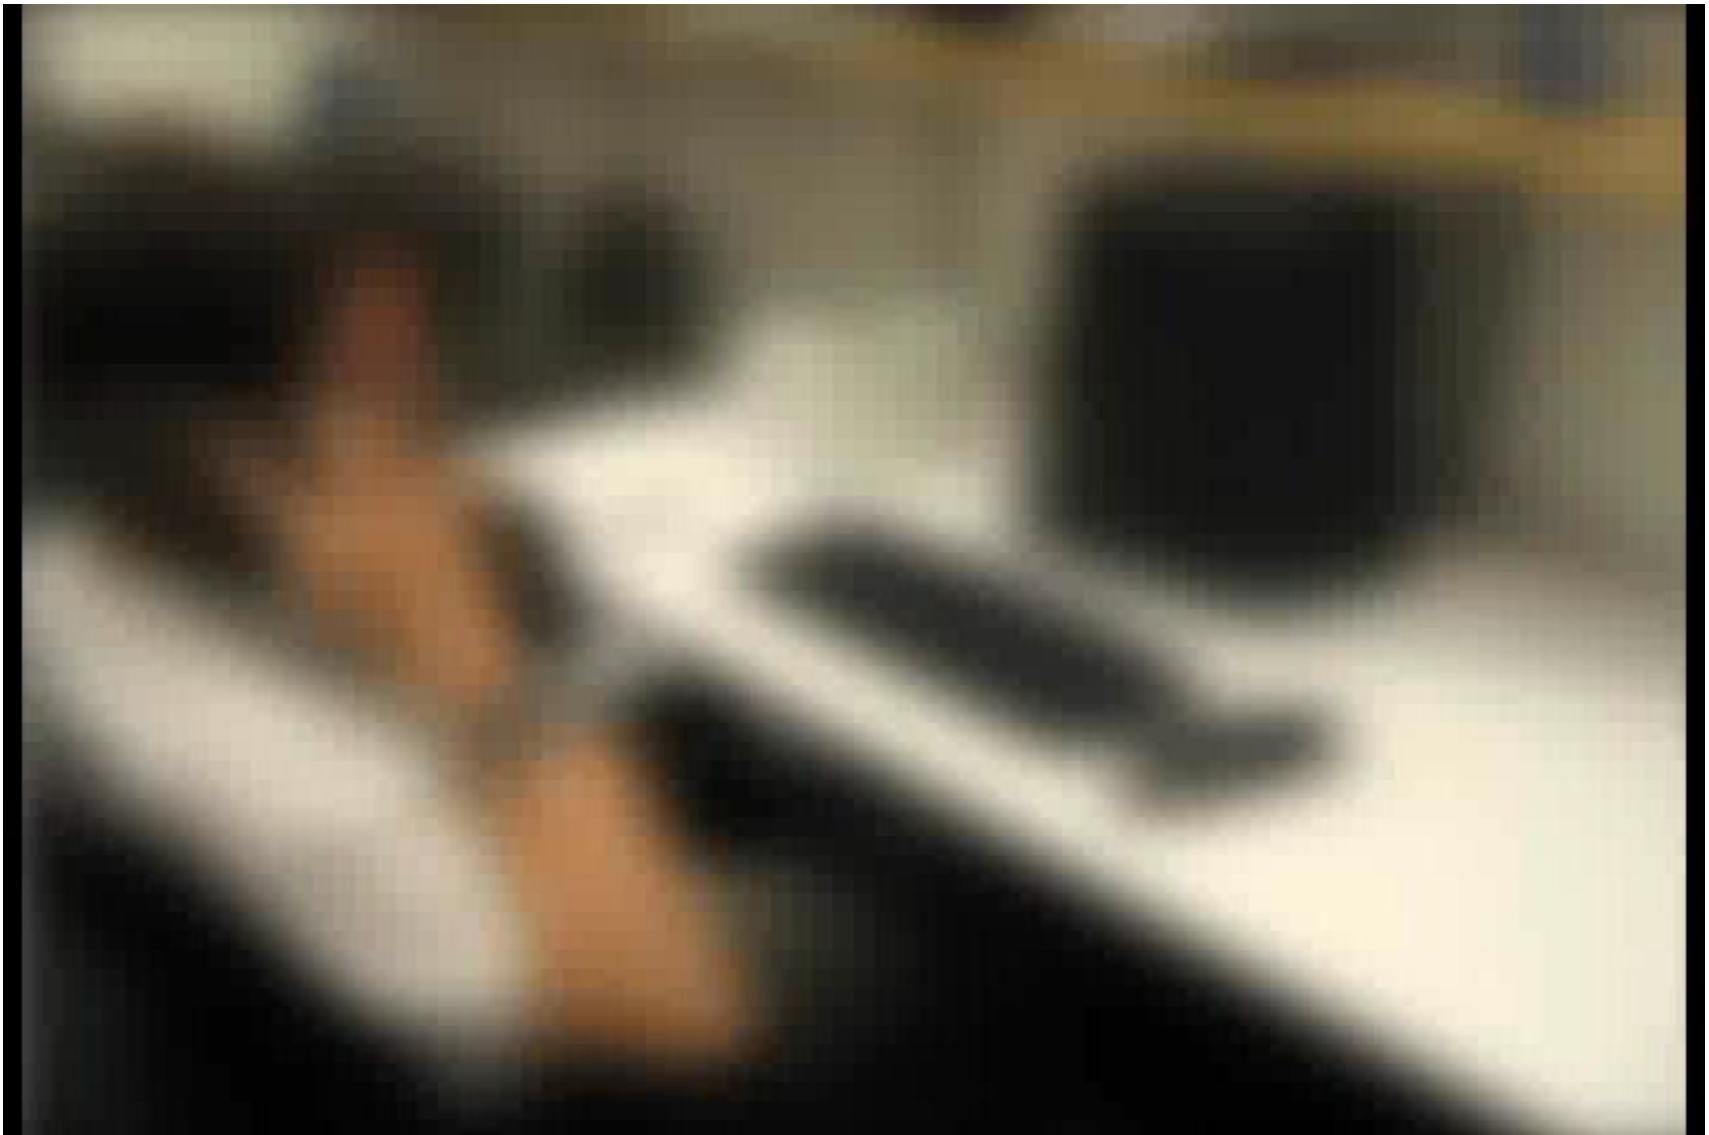


UNIVERSITÄT **BONN**

Juergen Gall

Recapitulation 1  
MA-INF 2201 - Computer Vision  
WS24/25

# Context



Source: A. Torralba

# Humans work with strong priors



Source: A. Torralba

# Conclusion

Human vision relies not **only** on data term (pixel information)

Priors are important

They explain the **most likely** scene if data is ambiguous:

$$\operatorname{argmax}_y P(y|I) = \operatorname{argmax}_y P(I|y)P(y)$$

Bayes formula:

$$P(y|I) = \frac{P(I|y)P(y)}{P(I)}$$

Posterior = Likelihood x Prior

# Image as function



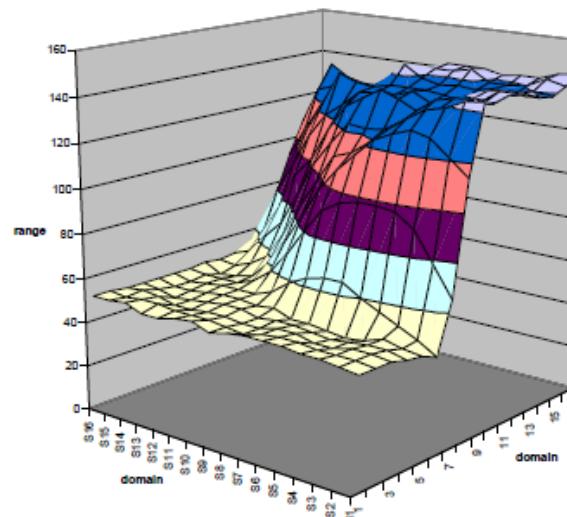
45	60	98	127	132	133	137	133
46	65	98	123	126	128	131	133
47	65	96	115	119	123	135	137
47	63	91	107	113	122	138	134
50	59	80	97	110	123	133	134
49	53	68	83	97	113	128	133
50	50	58	70	84	102	116	126
50	50	52	58	69	86	101	120

Mapping from image domain (pixels)  
to values

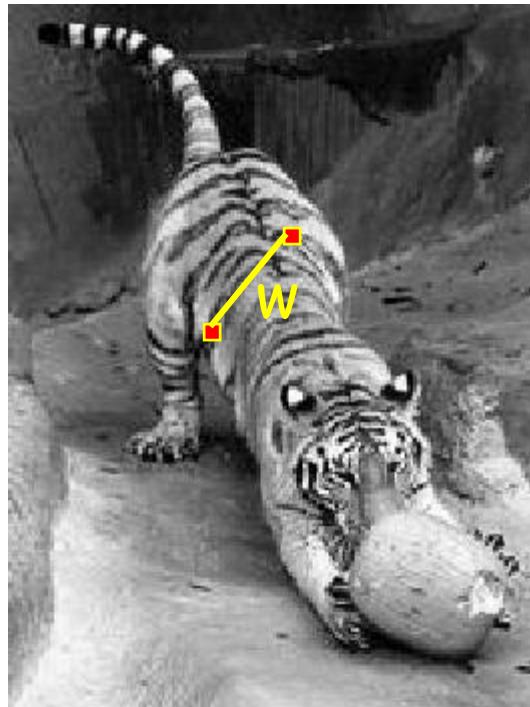
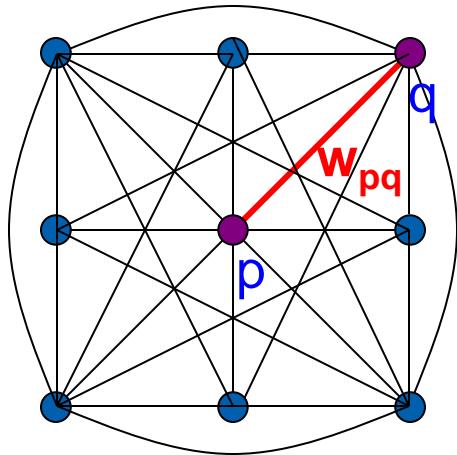
Gray image:  $f : \Omega \rightarrow \mathbb{R}$

Color images:  $f : \Omega \rightarrow \mathbb{R}^3$

Image represented as matrix or  
tensor



# Images as graphs



- *Fully-connected graph*
  - node (vertex) for every pixel
  - link between *every* pair of pixels,  $p, q$
  - affinity weight  $w_{pq}$  for each link (edge)
    - $w_{pq}$  measures *similarity*
      - similarity is *inversely proportional* to difference (in color and position...)

# Histogram equalization

Pixels as distribution:

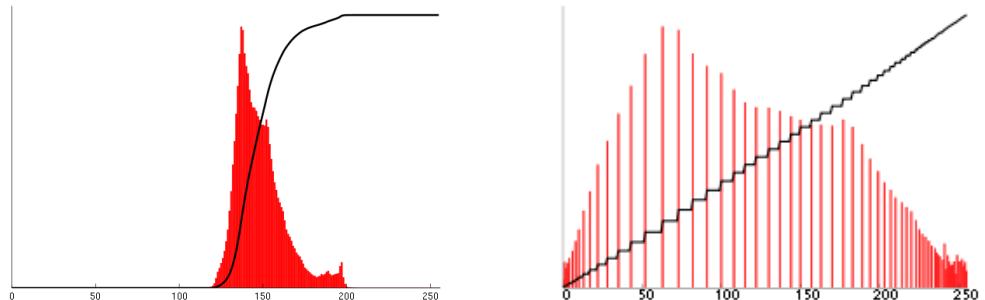
$$p_f(y) = \frac{|\{x \in \Omega : f(x) = y\}|}{|\Omega|}$$

$$F_f(y) = \int_0^y p_f(y') dy' = \sum_{y'=0}^y p_f(y')$$

Cumulative distribution function

Make CDF linear:

$$h(f(x)) = F(f(x)) \cdot \max_{x \in \Omega} f(x)$$



$$p_h(z) = \frac{|\{x \in \Omega : h(f(x)) = z\}|}{|\Omega|} = \frac{1}{\max_{x \in \Omega} f(x)}$$

# Linear filtering



Source: S. Lazebnik

# Defining convolution

Let  $f$  be the image and  $h$  be the kernel. The output of convolving  $f$  with  $h$  is denoted:  $g = f * h$

$$g(i, j) = \sum_{k,l} f(i - k, j - l)h(k, l) = \sum_{k,l} f(k, l)h(i - k, j - l)$$

45	60	98	127	132	133	137	133
46	65	98	123	126	128	131	133
47	65	96	115	119	123	135	137
47	63	91	107	113	122	138	134
50	59	80	97	110	123	133	134
49	53	68	83	97	113	128	133
50	50	58	70	84	102	116	126
50	50	52	58	69	86	101	120

\*

0.1	0.1	0.1
0.1	0.2	0.1
0.1	0.1	0.1

=

69	95	116	125	129	132
68	92	110	120	126	132
66	86	104	114	124	132
62	78	94	108	120	129
57	69	83	98	112	124
53	60	71	85	100	114

 $f(x,y)$ 
 $h(x,y)$ 
 $g(x,y)$

# Key properties

Linearity:  $\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$

Shift invariance: same behavior regardless of pixel location:  $\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$

Theoretical result: any linear shift-invariant operator can be represented as a convolution

# Properties in more detail

Commutative:  $a * b = b * a$

- Conceptually no difference between filter and signal

Associative:  $a * (b * c) = (a * b) * c$

- Often apply several filters one after another:  $((a * b1) * b2) * b3)$
- This is equivalent to applying one filter:  $a * (b1 * b2 * b3)$

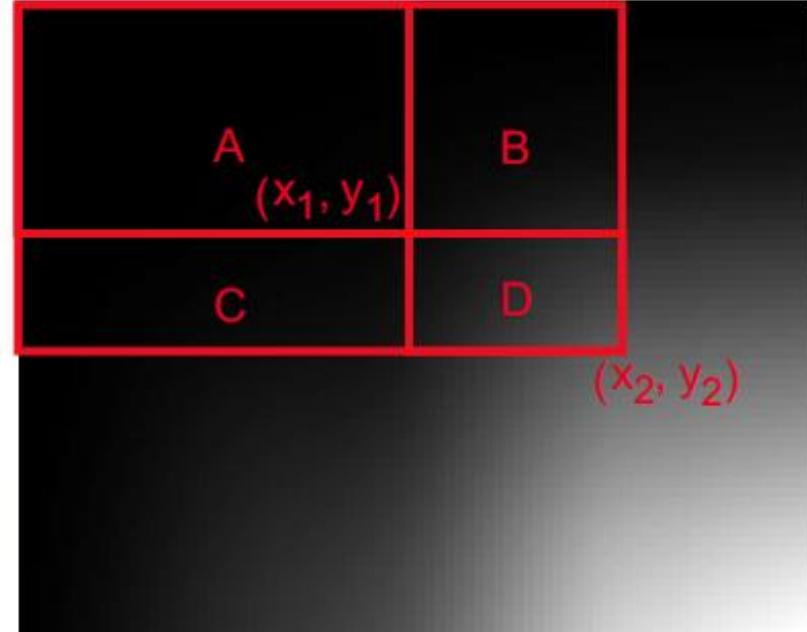
Distributes over addition:  $a * (b + c) = (a * b) + (a * c)$

Scalars factor out:  $ka * b = a * kb = k(a * b)$

Identity: unit impulse  $e = [\dots, 0, 0, 1, 0, 0, \dots]$ ,  
 $a * e = a$

# Integral image

Precompute average



Get average values of any size by reading only 4 values!

# Integral image

3	2	7	2	3
1	5	1	3	4
5	1	3	5	1
4	3	2	1	6
2	4	1	4	8

Original image

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

Integral image

$$s(i, j) = \sum_{k=0}^i \sum_{l=0}^j f(k, l)$$

Recursive:

$$s(i, j) = s(i - 1, j) + s(i, j - 1) - s(i - 1, j - 1) + f(i, j)$$

$$17+19-11+3 = 28$$

# Integral image

3	2	7	2	3
1	5	1	3	4
5	1	3	5	1
4	3	2	1	6
2	4	1	4	8

Original image

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

Integral image

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

Integral image

4 values independent of size:

$$s(i_1, j_1) - s(i_1, j_0 - 1) - s(i_0 - 1, j_1) + s(i_0 - 1, j_0 - 1)$$

$$48 - 13 - 14 + 3 = 24 = 5 + 1 + 3 + 1 + 3 + 5 + 3 + 2 + 1$$

# Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1

?

(Note that filter sums to 1)

# Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$	1	1	1
1	1	1	1
1	1	1	1



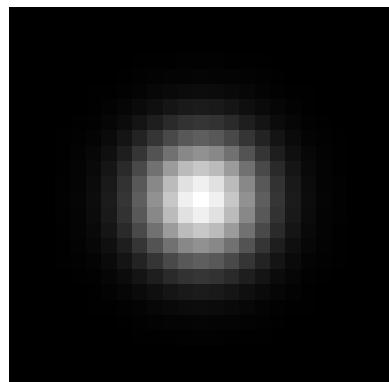
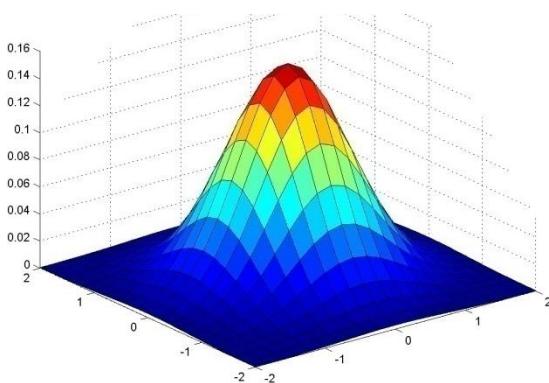
## Sharpening filter

- Accentuates differences with local average

# Gaussian Kernel

Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should renormalize weights to sum to 1 in any case)

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

5 x 5,  $\sigma = 1$

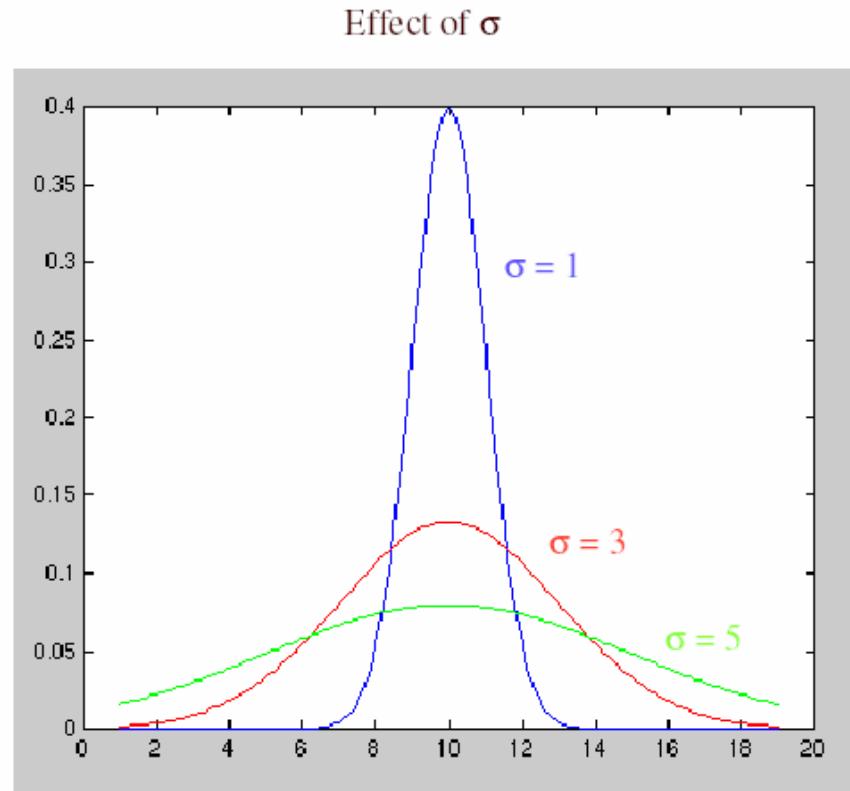
# Choosing kernel width

Rule of thumb: set filter half-width to about  $3\sigma$

$$P(\mu - \sigma < X \leq \mu + \sigma) \approx 68\%$$

$$P(\mu - 2\sigma < X \leq \mu + 2\sigma) \approx 95.5\%$$

$$P(\mu - 3\sigma < X \leq \mu + 3\sigma) \approx 99.7\%$$



# Gaussian filters

Remove “high-frequency” components from the image (low-pass filter)

Convolution with self is another Gaussian

- So can smooth with small- $\sigma$  kernel, repeat, and get same result as larger- $\sigma$  kernel would have
- Convolving two times with Gaussian kernel with std.  $\sigma$  is same as convolving once with kernel with std. dev.  $\sigma\sqrt{2}$

Separable kernel

- Factors into product of two 1D Gaussians

# Why is separability useful?

What is the complexity of filtering an  $n \times n$  image with an  $m \times m$  kernel?

- $O(n^2 m^2)$

What if the kernel is separable?

- $O(n^2 m)$

# Is a kernel separable?

- Kernel as matrix:

$$\mathbf{K} = \sum_i \sigma_i \mathbf{u}_i \mathbf{v}_i^T$$

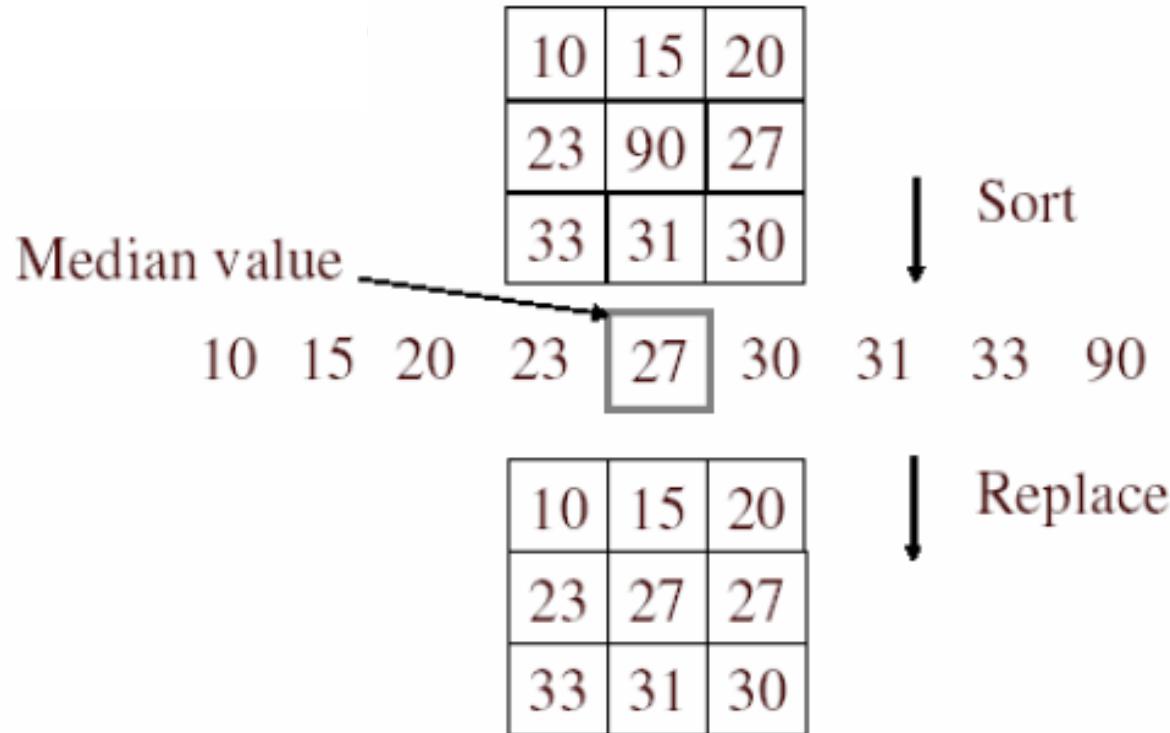
- Singular Value Decomposition (SVD)
- If only the first singular value  $\sigma_0$  is non-zero, kernel is separable
- Vertical and horizontal kernels:

$$\sqrt{\sigma_0} \mathbf{u}_0 \text{ and } \sqrt{\sigma_0} \mathbf{v}_0^T$$

- Approximation:  $\mathbf{K} \approx \sigma_0 \mathbf{u}_0 (\mathbf{v}_0)^T$

# Alternative idea: Median filtering

A median filter operates over a window by selecting the median intensity in the window

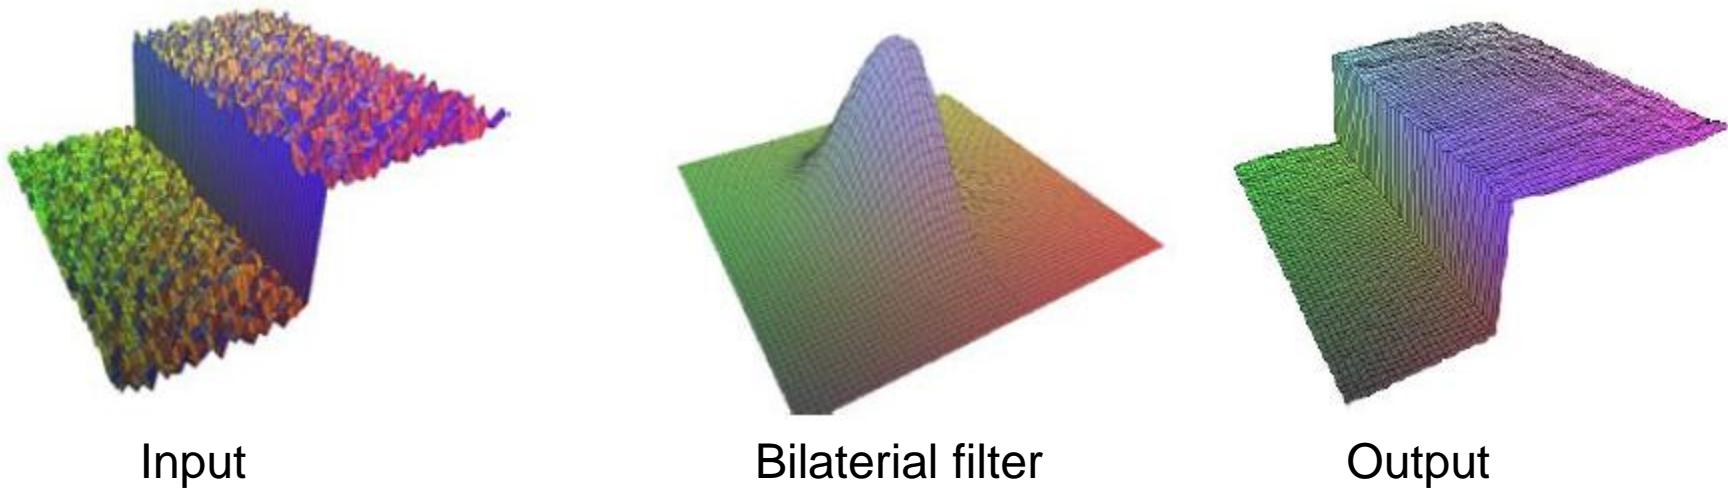


- Is median filtering linear?

# Bilateral filter

Filter is data dependent

$$g(i, j) = \frac{\sum_{k,l} f(k, l)w(i, j, k, l)}{\sum_{k,l} w(i, j, k, l)}$$



Input

Bilateral filter

Output

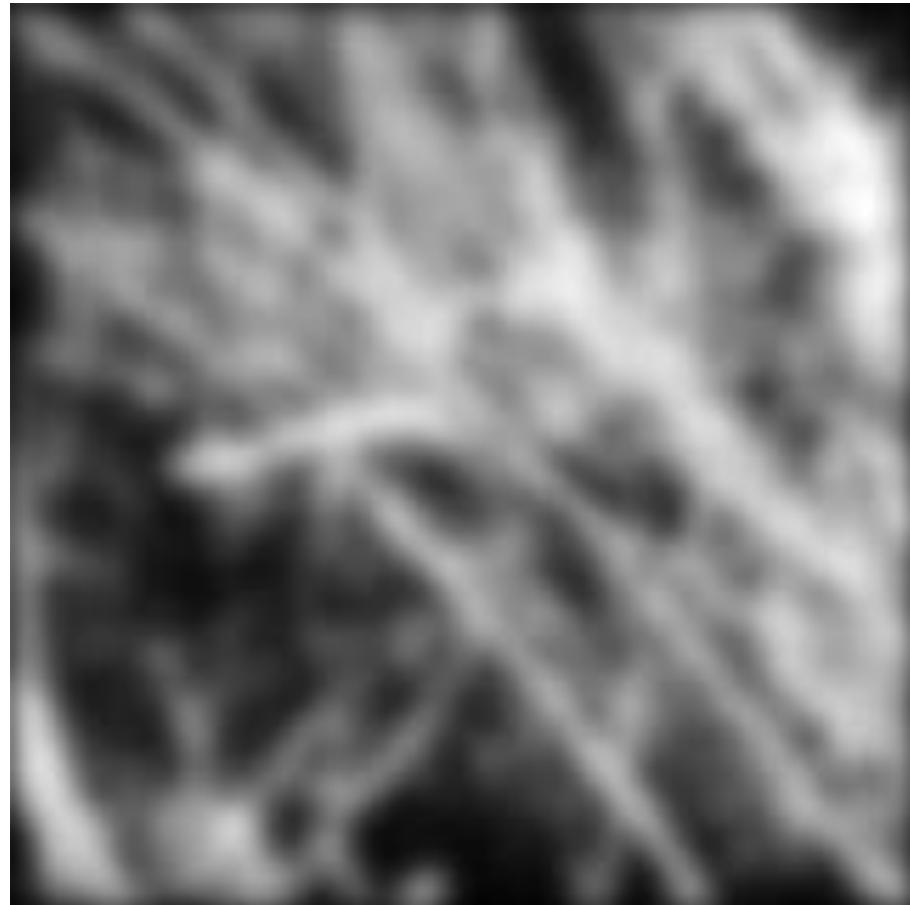
$$w(i, j, k, l) = \exp \left( -\frac{(i - k)^2 + (j - l)^2}{2\sigma_d^2} - \frac{\|f(i, j) - f(k, l)\|^2}{2\sigma_r^2} \right)$$

Slow, but methods for approximation exist

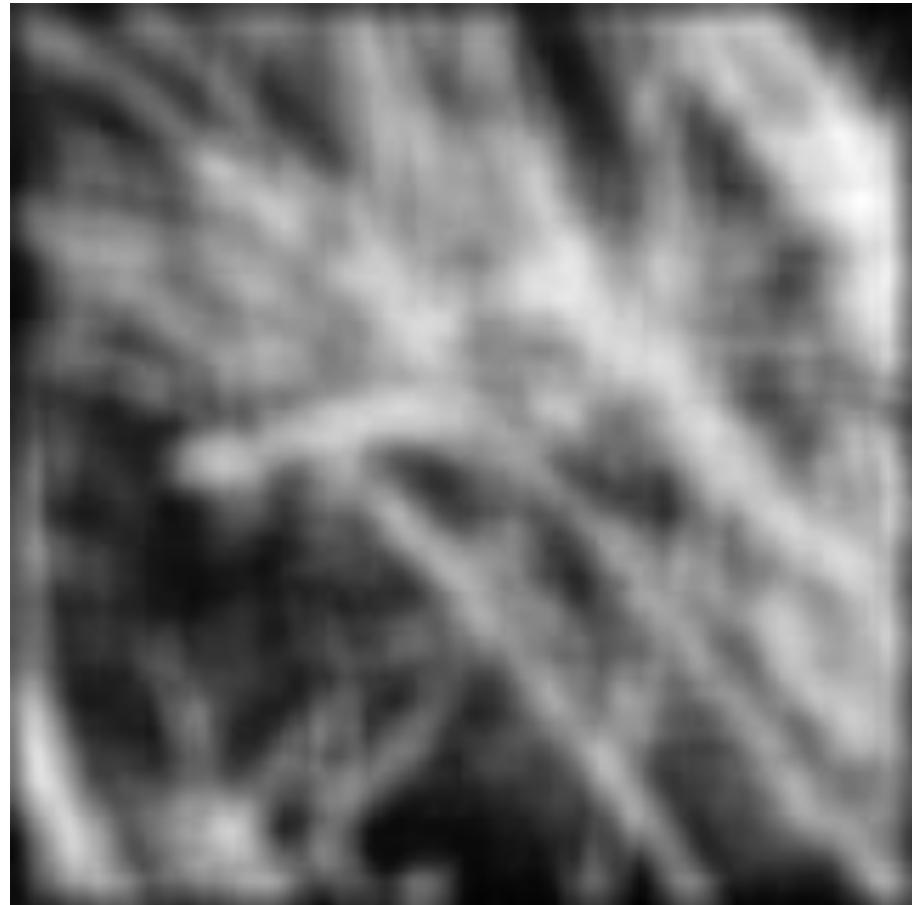
# Linear filtering

## Gauss vs. box filter

Gaussian



Box filter



# A sum of sines

Our building block:

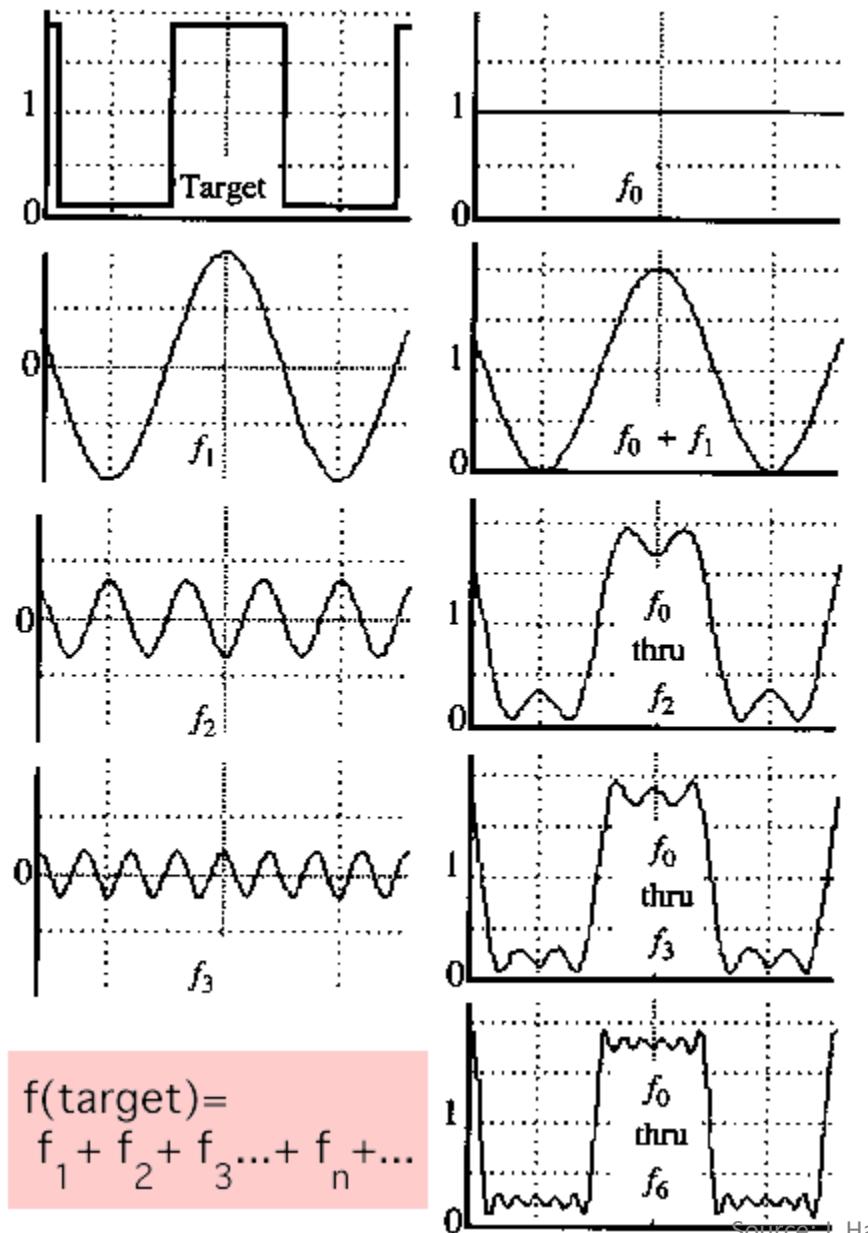
$$A \sin(\omega x + \phi)$$

A: Amplitude

$\omega$ : angular frequency

$\phi$ : phase

Add enough of them to get any signal  $f(x)$  you want!



Source: J. Hays

# Fourier Transform

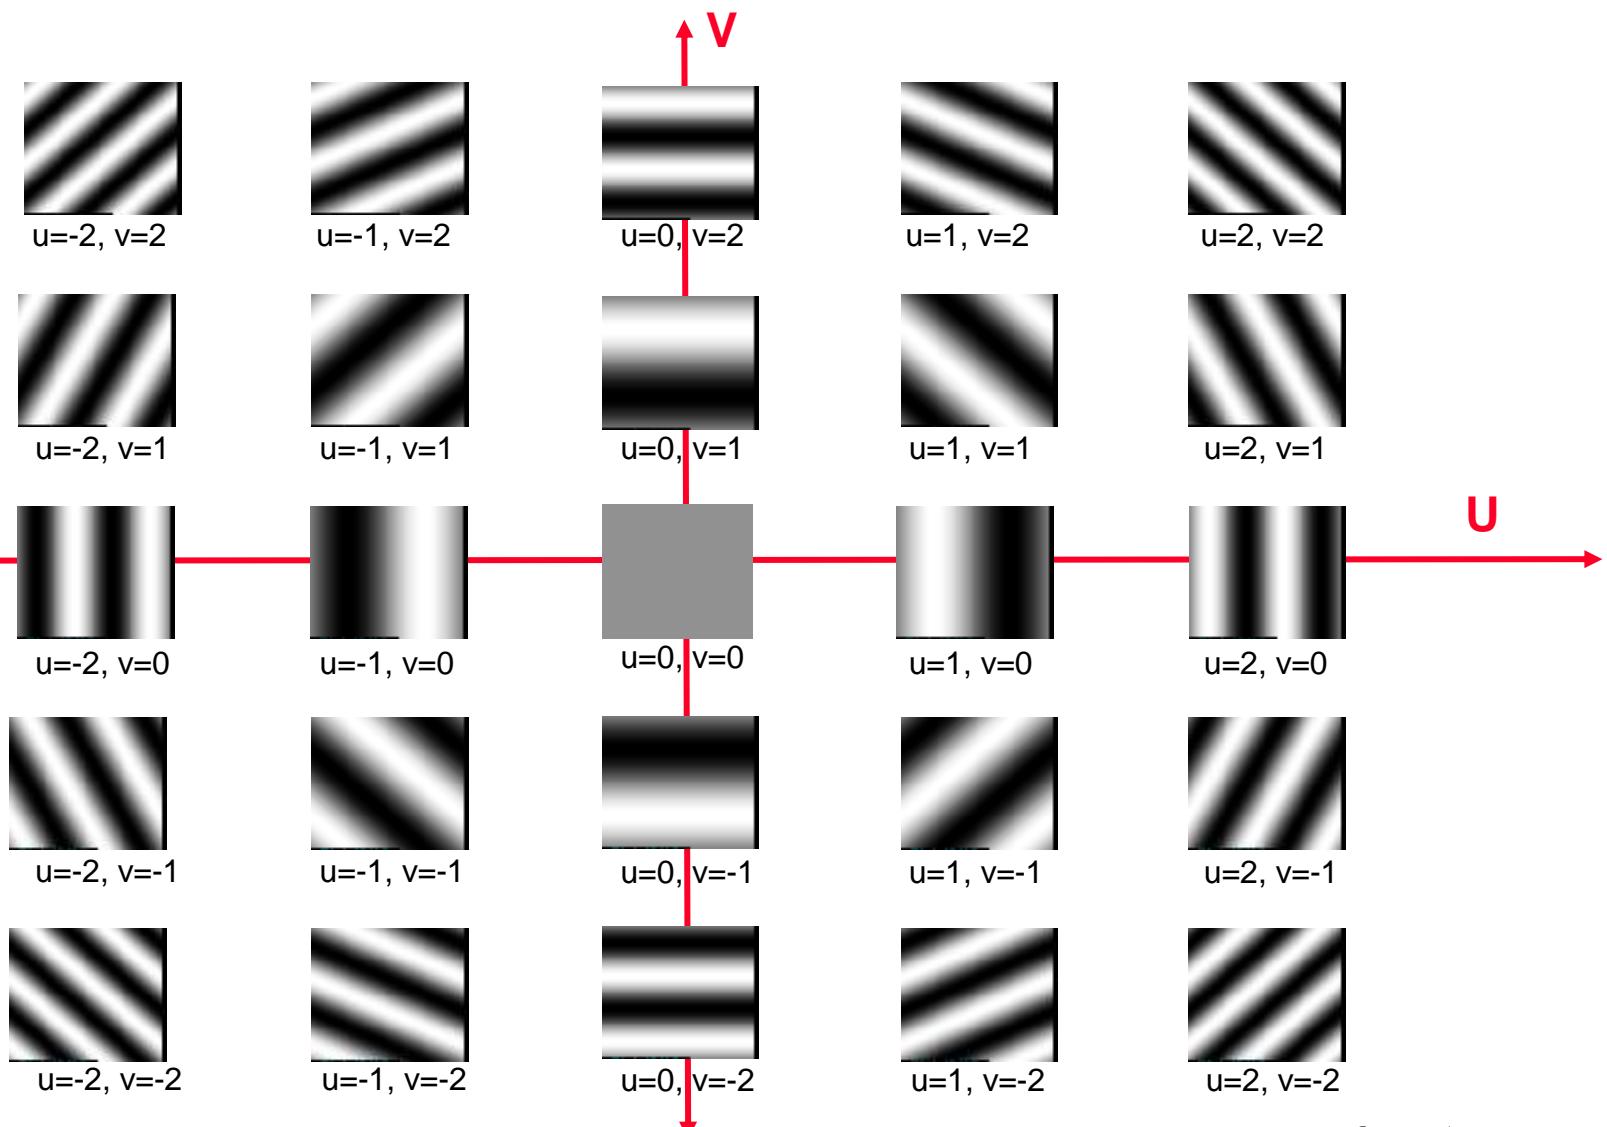
Fourier transform stores the magnitude and phase at each frequency

- Magnitude encodes how much signal there is at a particular frequency
- Phase encodes spatial information (indirectly)
- For mathematical convenience, this is often notated in terms of real and complex numbers

$$\text{Amplitude: } A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

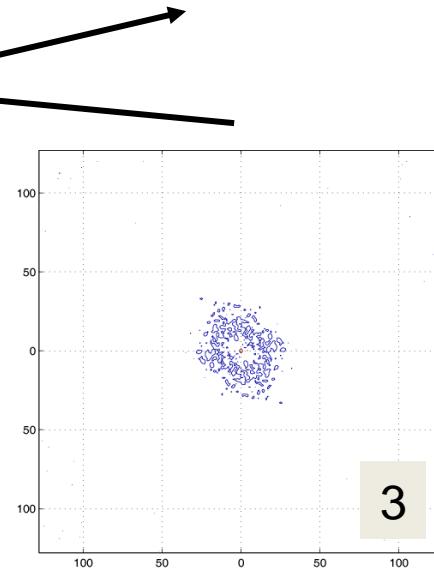
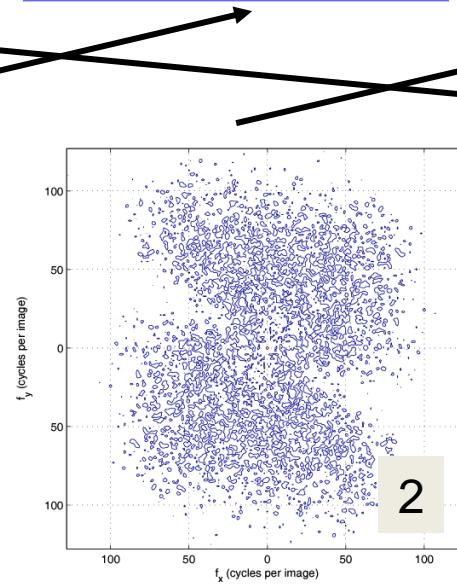
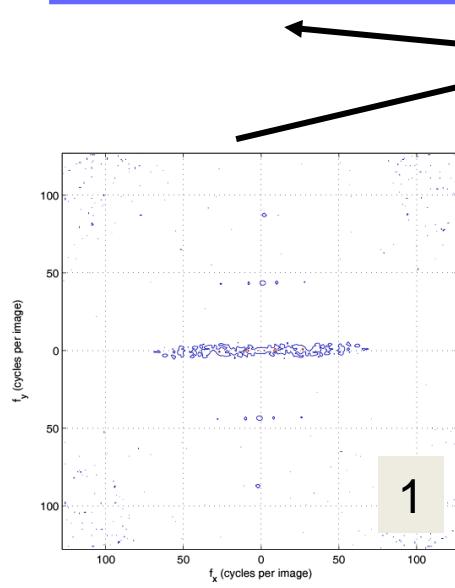
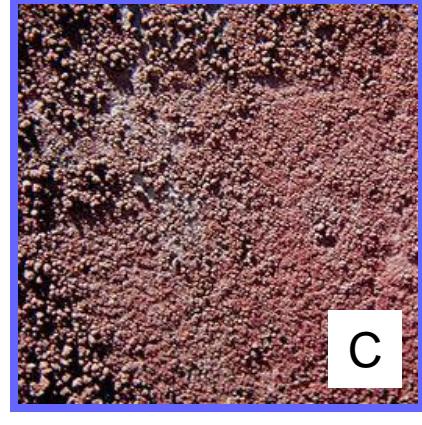
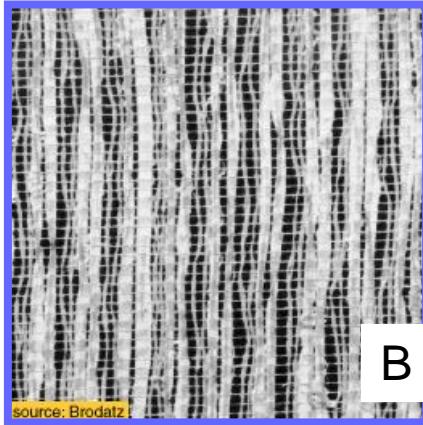
$$\text{Phase: } \phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

# The 2D Basis Functions



$$e^{2\pi i(ux+vy)}$$

# Fourier Amplitude Spectrum



# The Convolution Theorem

- The Fourier transform of the convolution of two functions is the product of their Fourier transforms

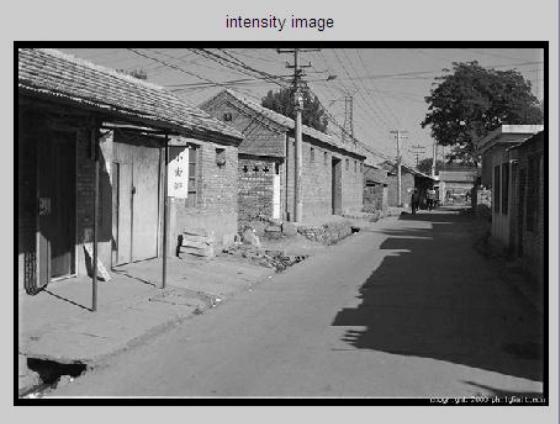
$$\mathcal{F}[g * h] = \mathcal{F}[g]\mathcal{F}[h]$$

- The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms

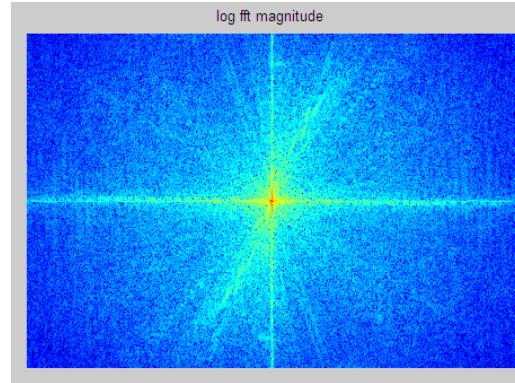
$$\mathcal{f}[gh] = \mathcal{f}[g] * \mathcal{f}[h]$$

- **Convolution** in spatial domain is equivalent to **multiplication** in frequency domain!

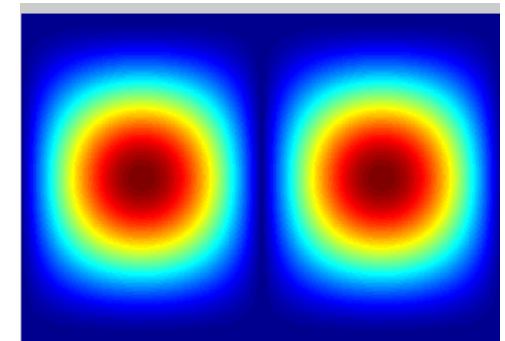
# Filtering in frequency domain



FFT

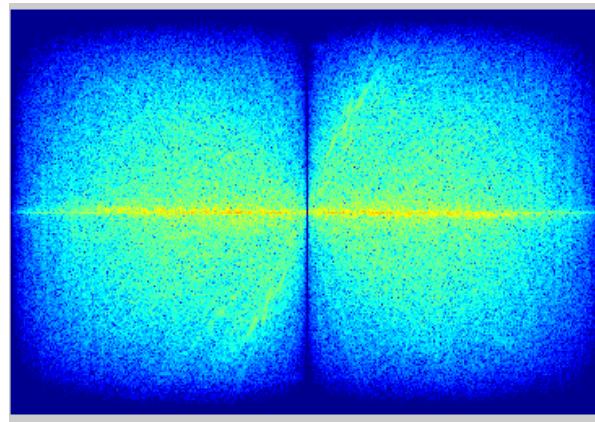


X



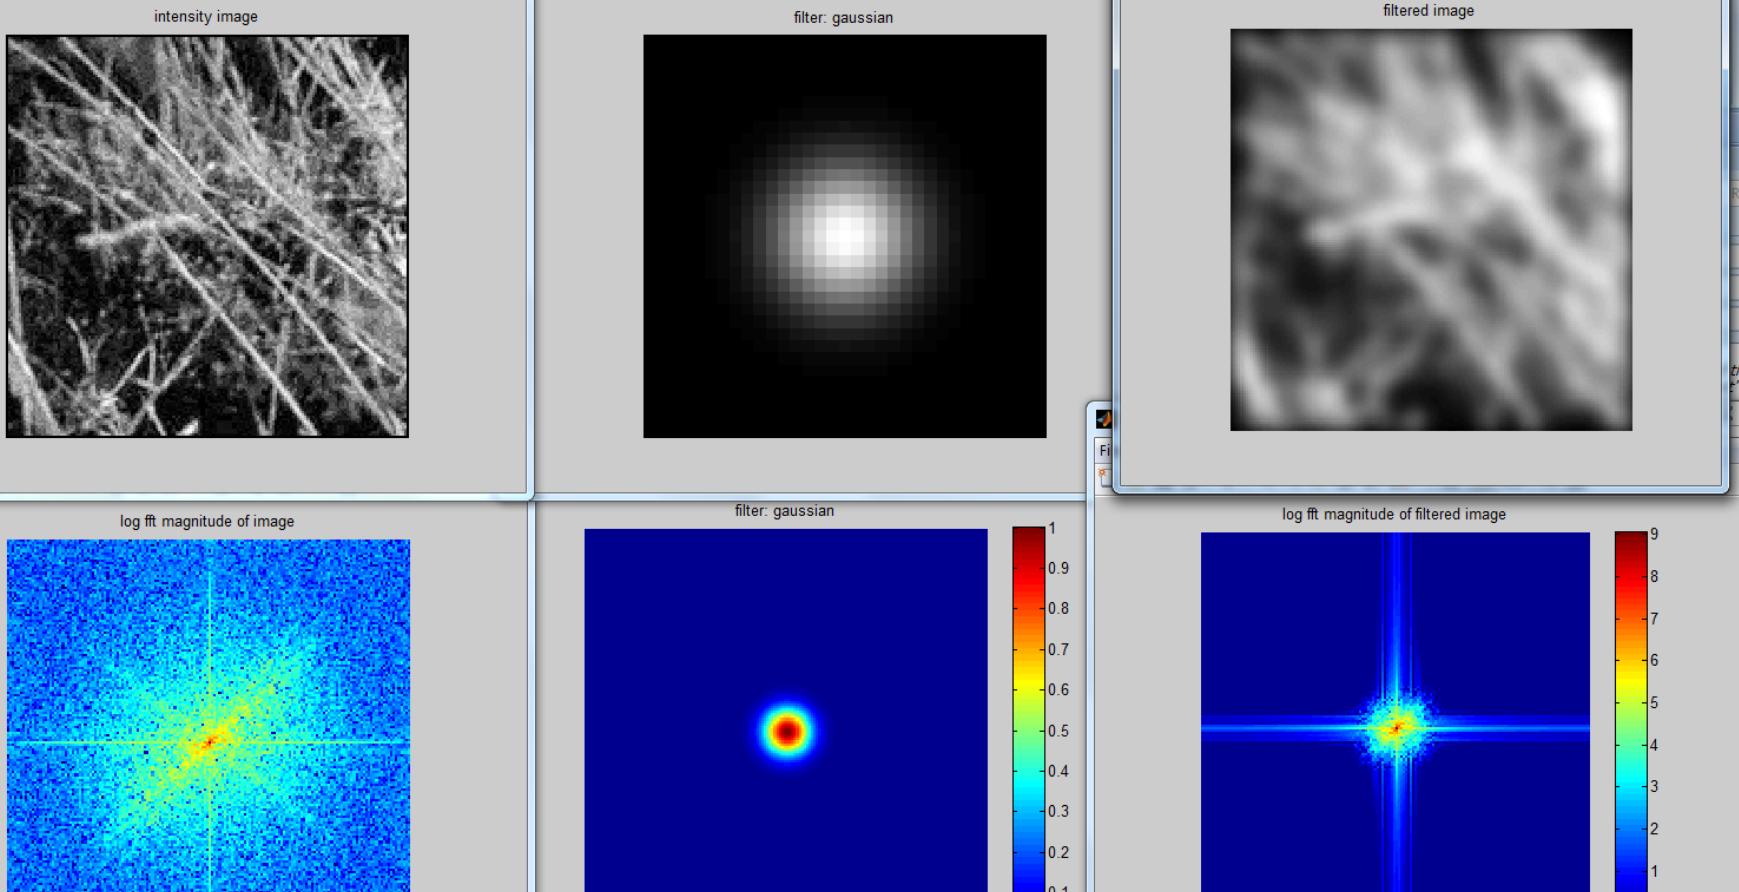
||

Inverse FFT



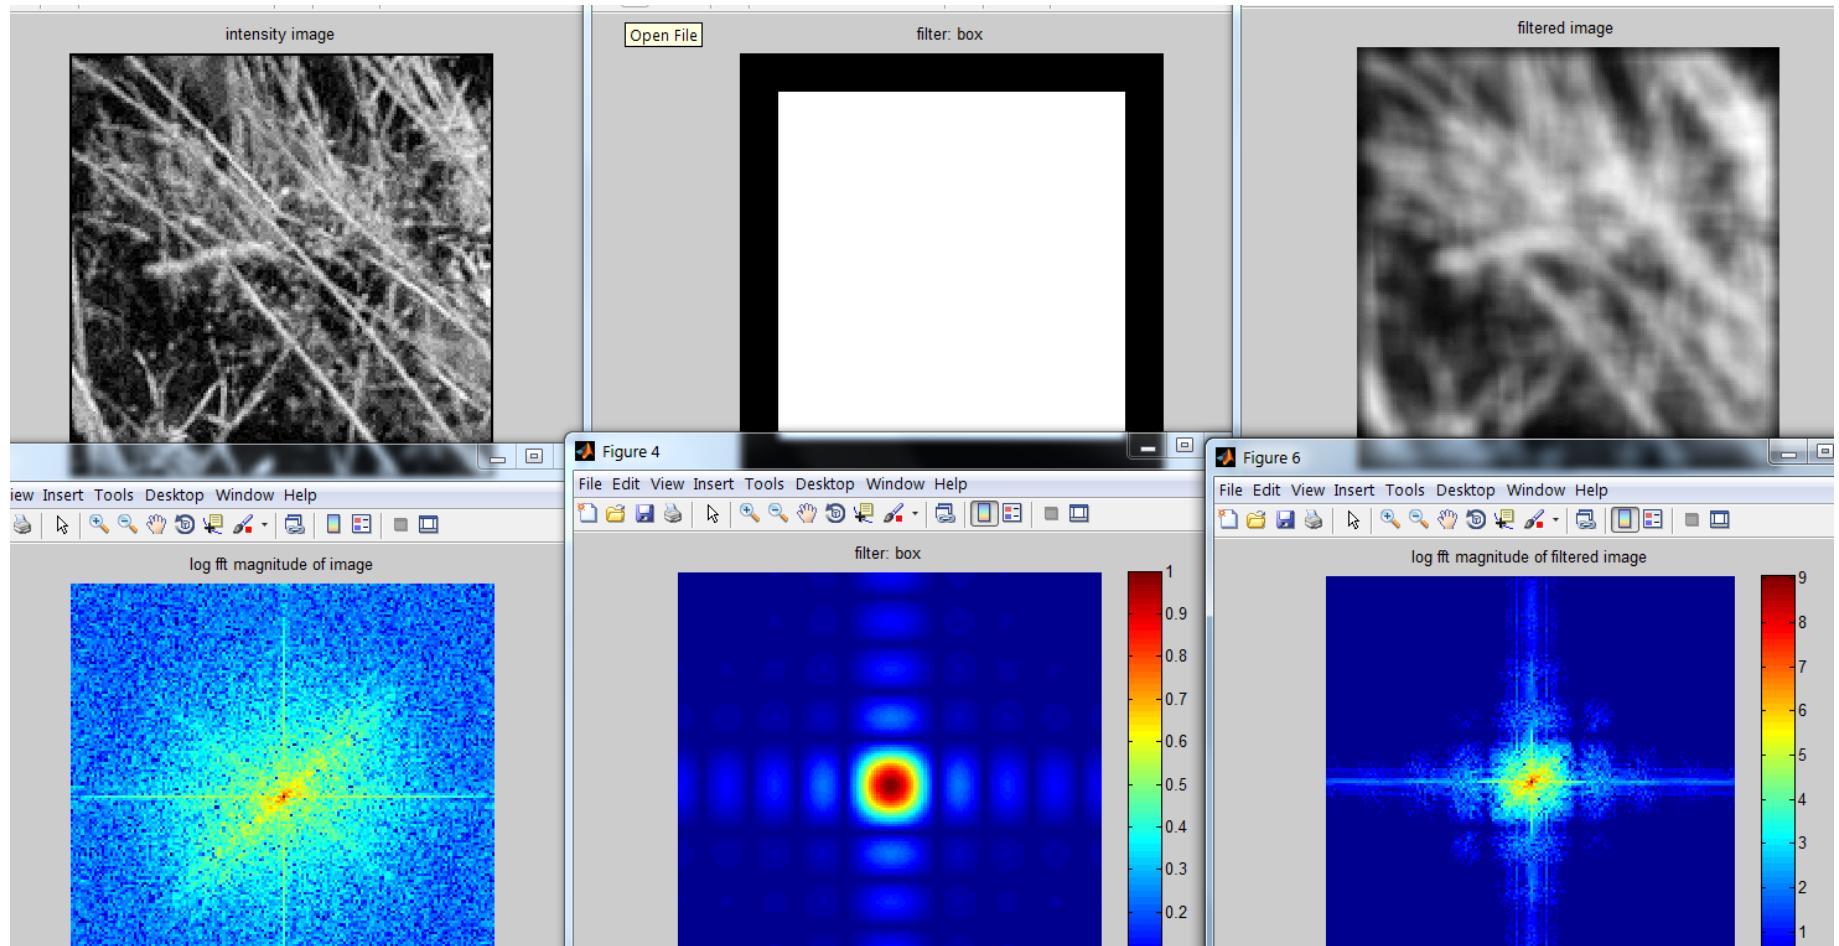
Source: D. Hoiem

# Gaussian



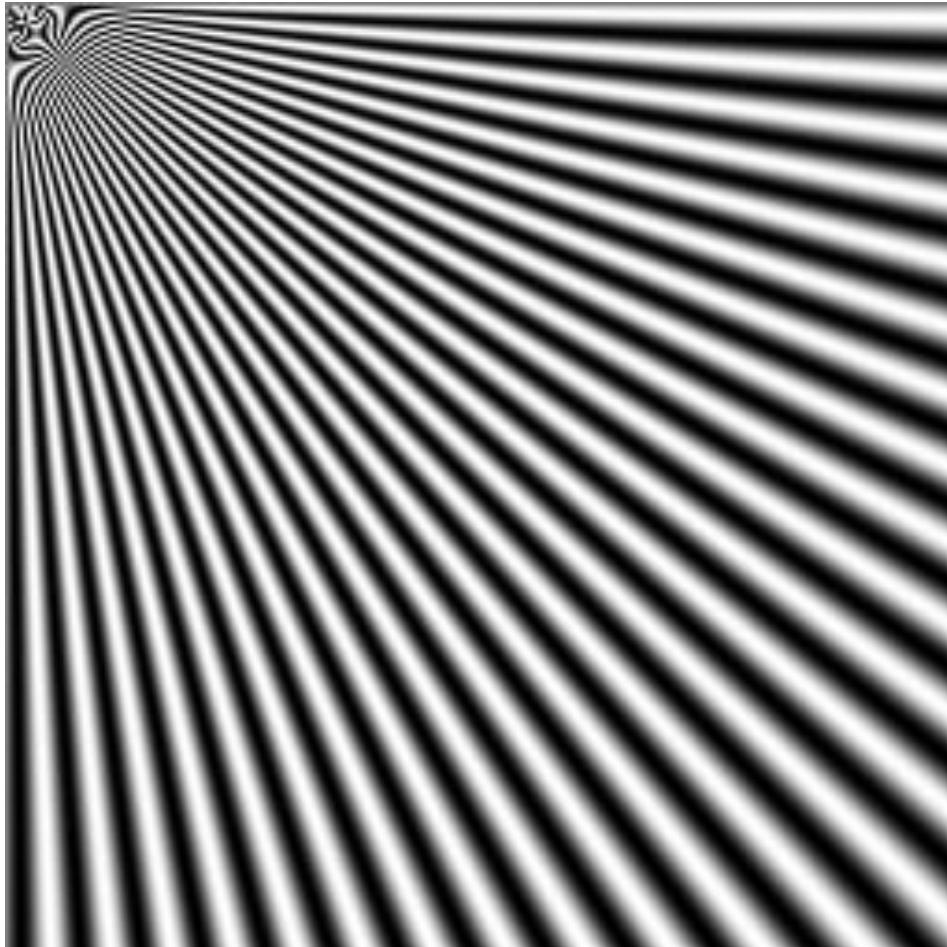
Source: J. Hays

# Box filter



Source: J. Hays

# Aliasing



Source: S. Narasimhan

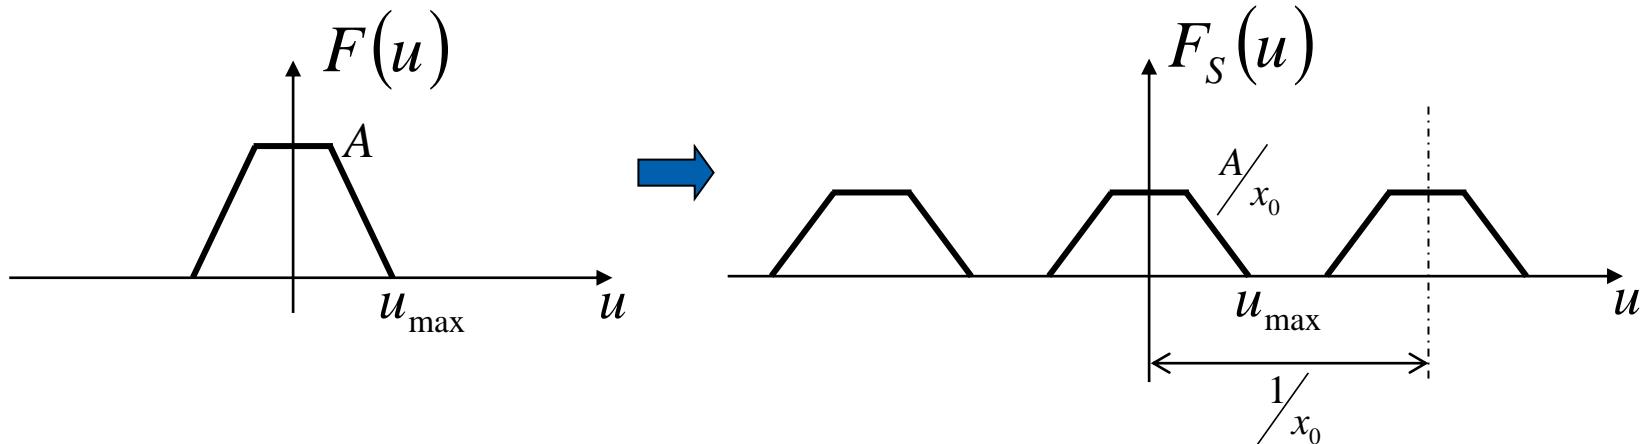
# Sampling Theorem

Sampled function:

$$f_s(x) = f(x)s(x) = f(x) \sum_{n=-\infty}^{\infty} \delta(x - nx_0)$$

Sampling frequency  $\frac{1}{x_0}$

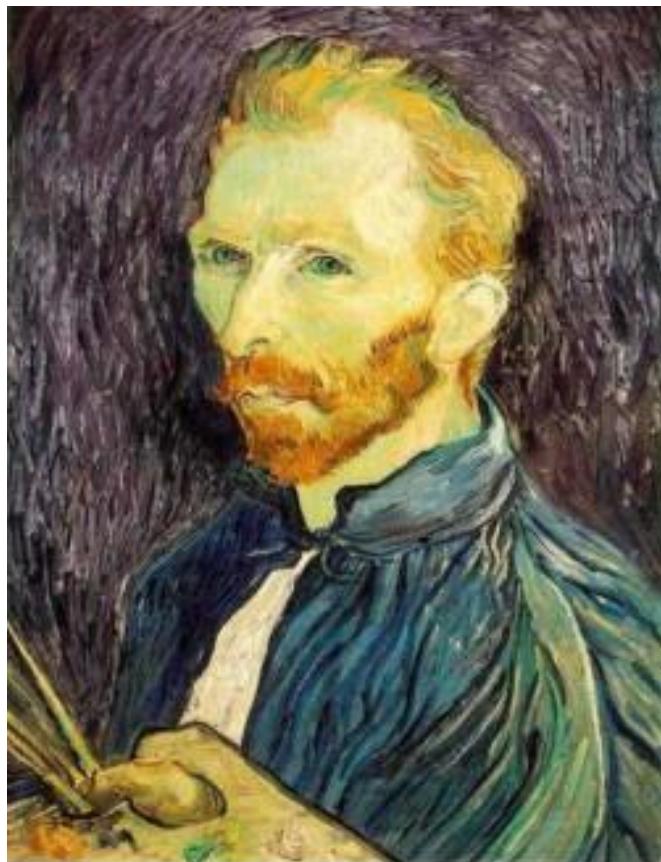
$$F_s(u) = F(u) * S(u) = F(u) * \frac{1}{x_0} \sum_{n=-\infty}^{\infty} \delta\left(u - \frac{n}{x_0}\right)$$



Only if  $u_{\max} \leq \frac{1}{2x_0}$

Source: S. Narasimhan

# Sub-Sampling with Gaussian Pre-Filtering



Gaussian 1/2



G 1/4



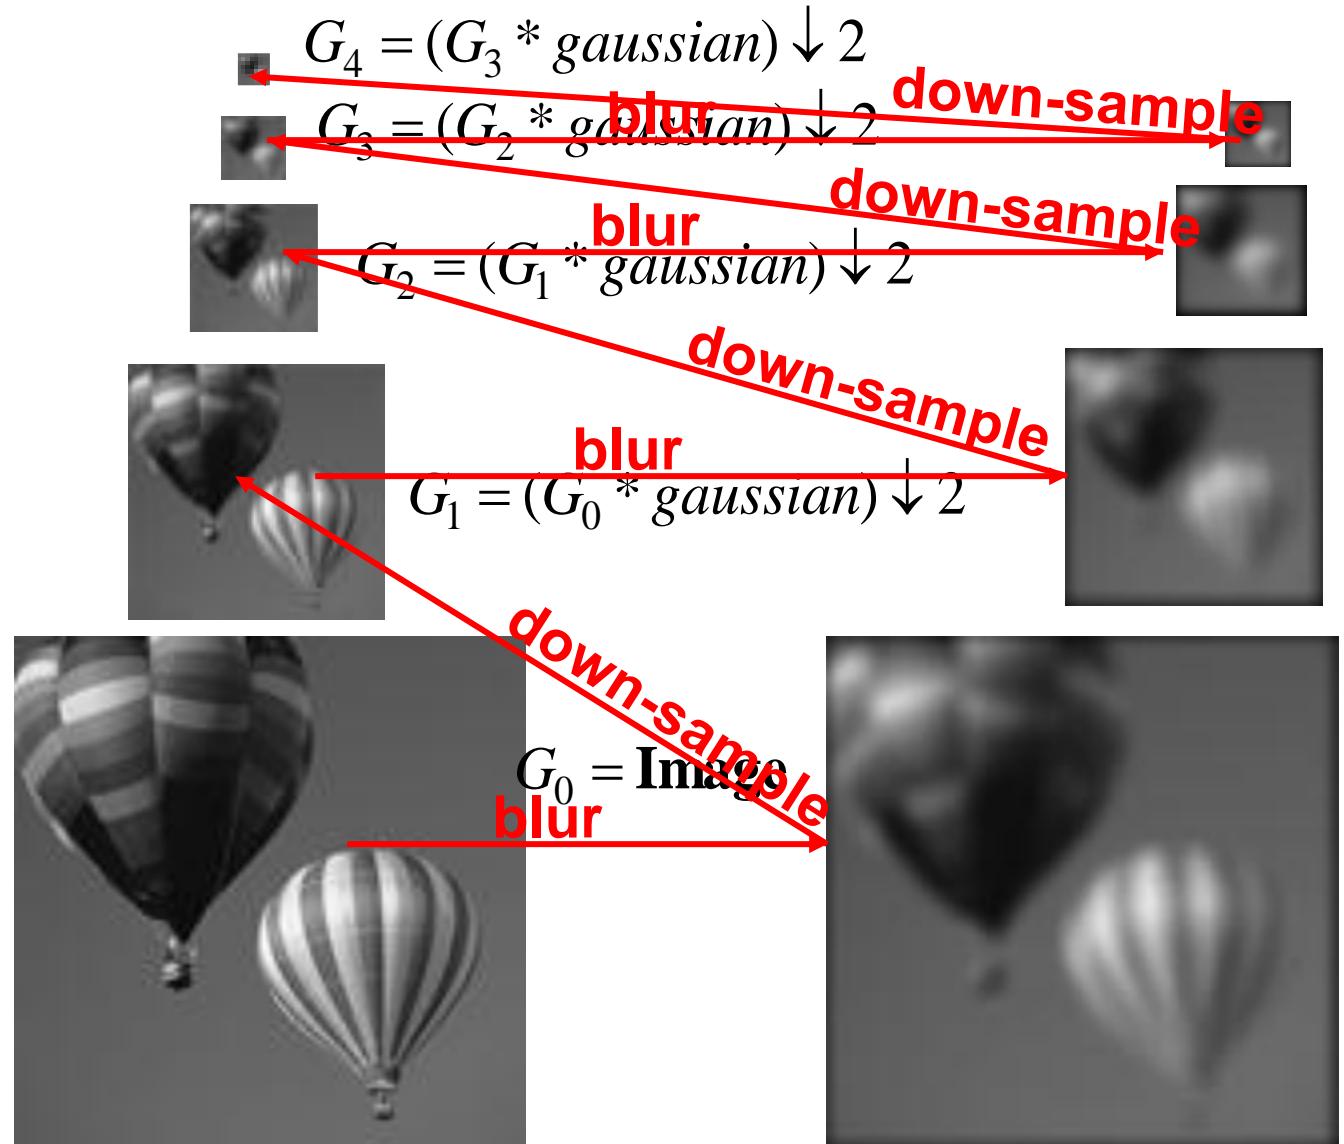
G 1/8

- Solution: filter the image, *then* subsample
  - Filter size should double for each  $\frac{1}{2}$  size reduction.

Source: S. Narasimhan

# Gaussian pyramid

Low resolution

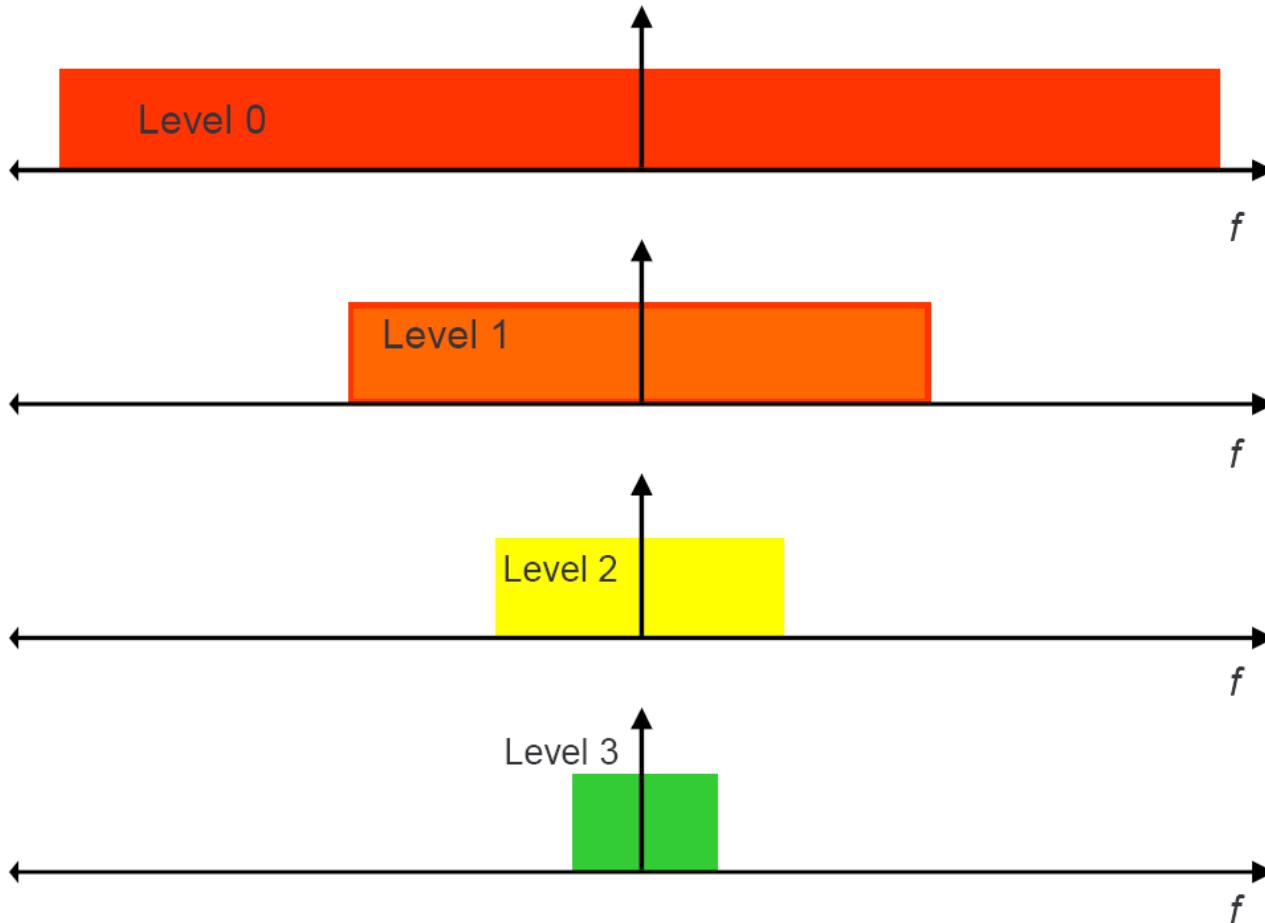


High resolution

Source: S. Narasimhan

# Frequencies

## Gaussian Pyramid Frequency Composition



Source: S. Narasimhan

# Sharpening

What does blurring take away?



-



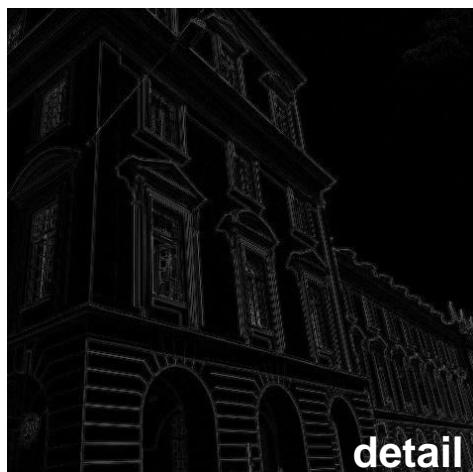
=



Let's add it back:



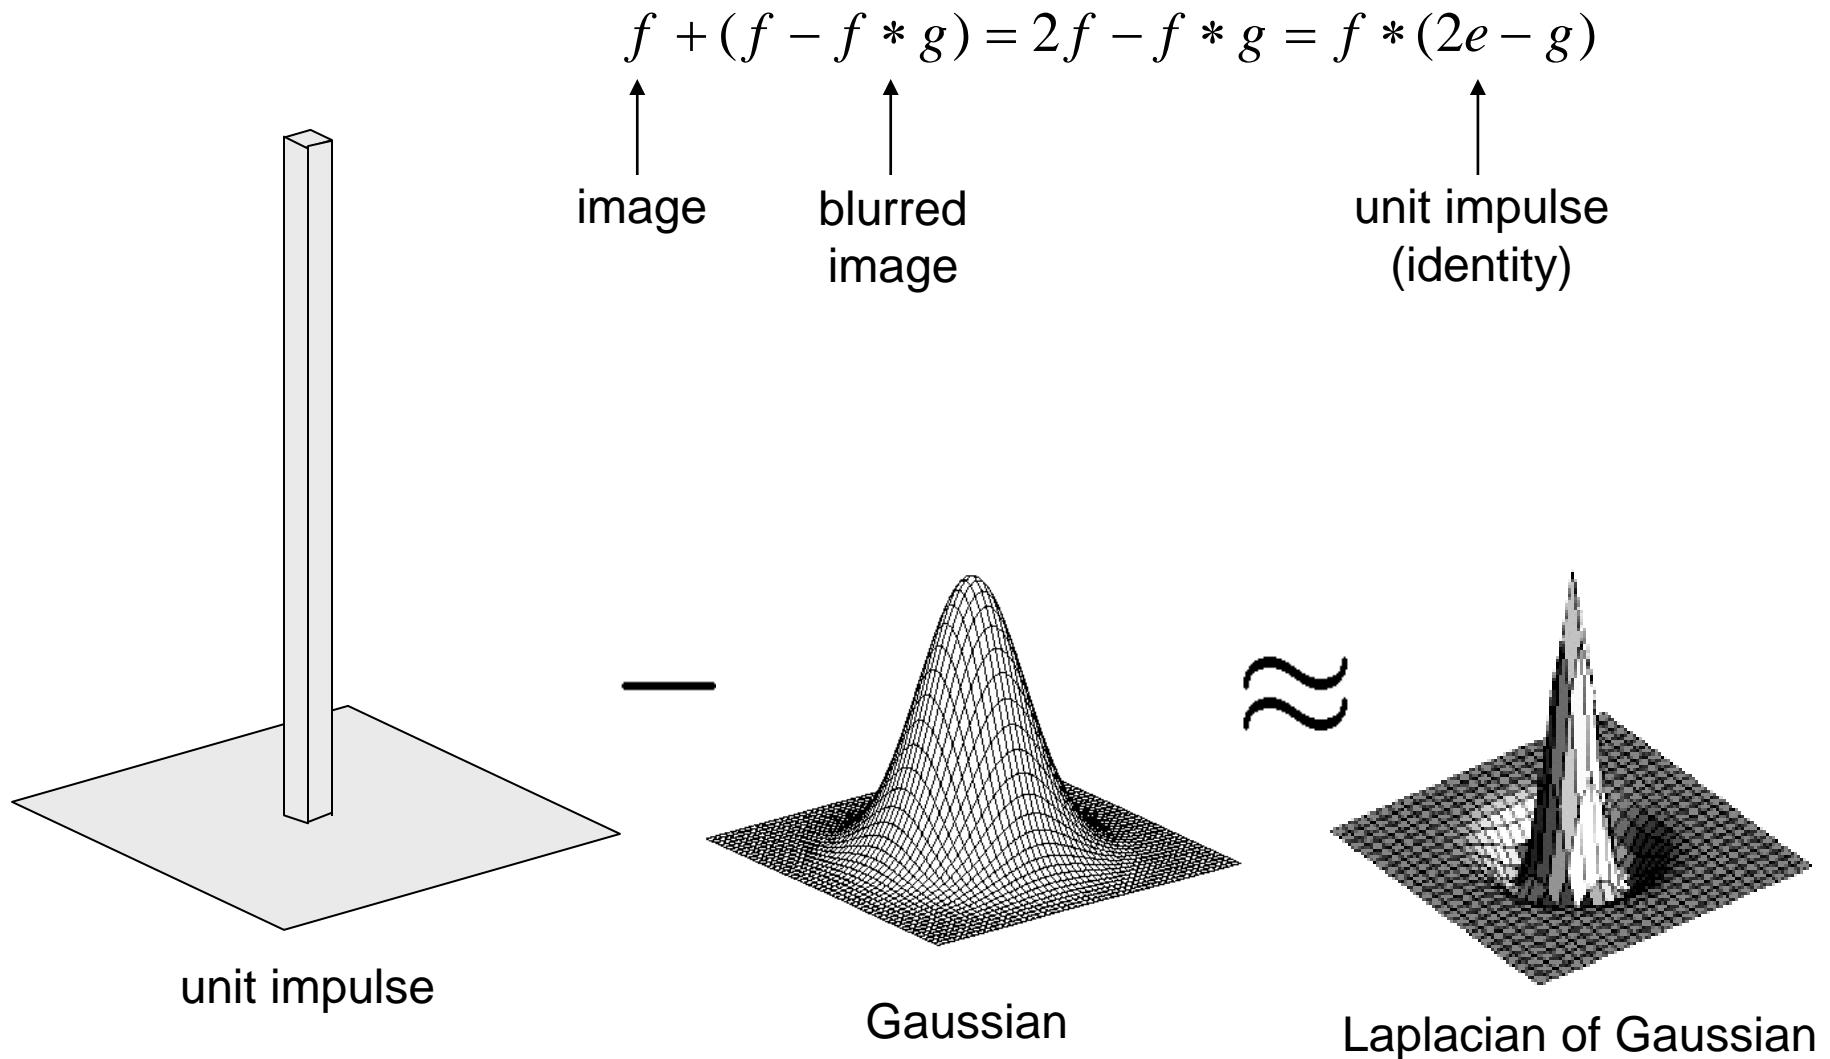
+



=



# Laplacian of Gaussian

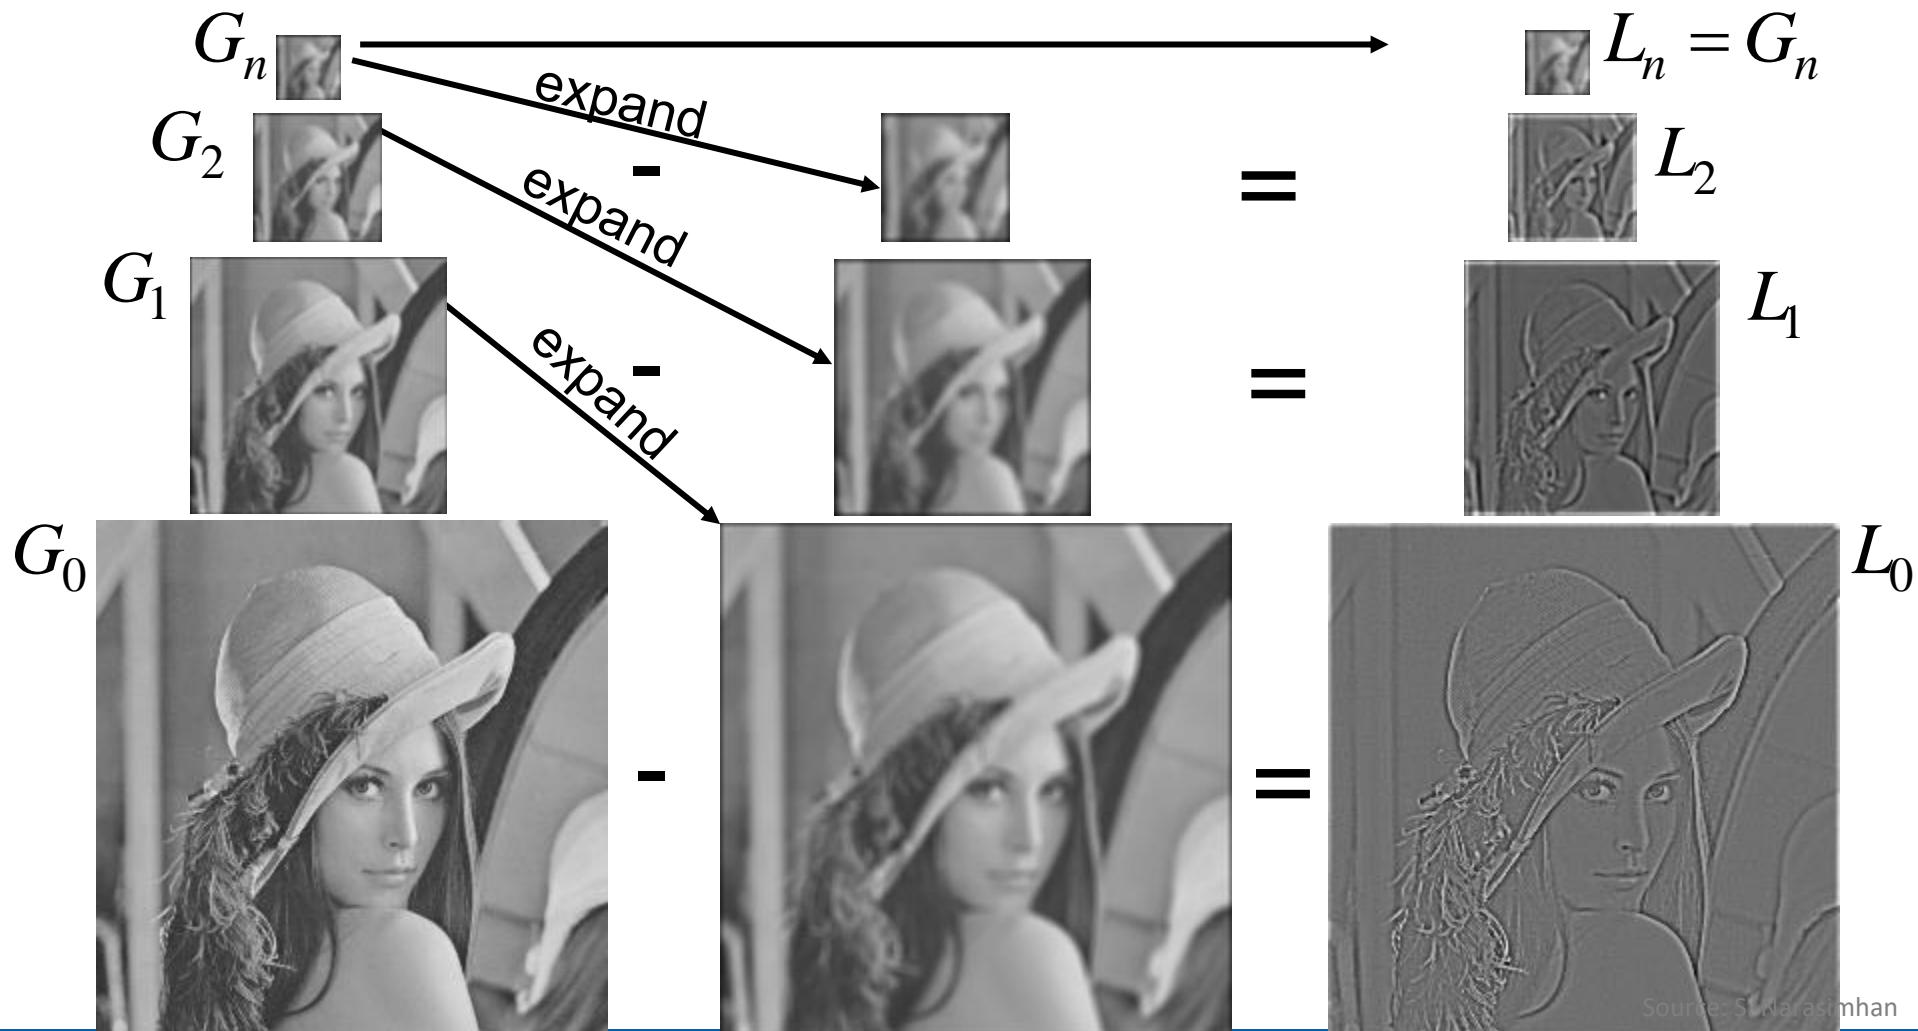


# Laplacian pyramid

Gaussian Pyramid

$$L_i = G_i - \text{expand}(G_{i+1})$$

Laplacian Pyramid



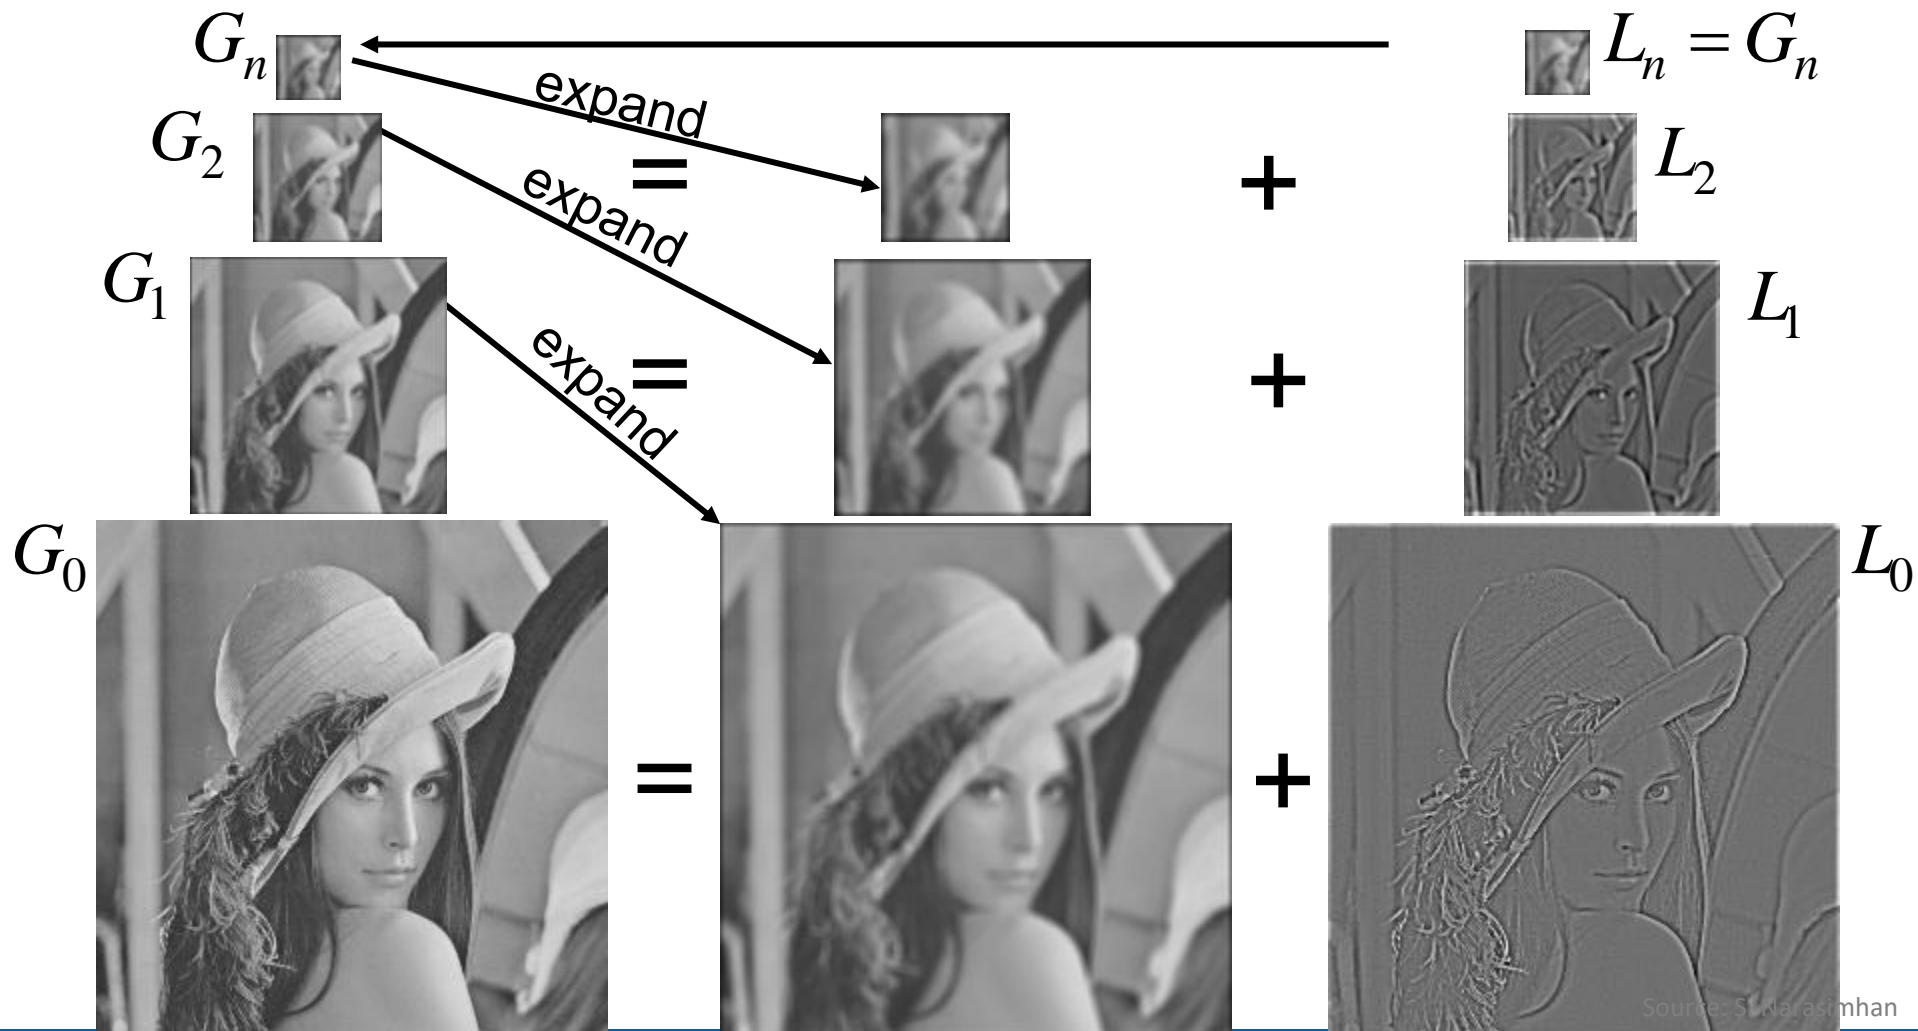
Source: S. Narasimhan

# Laplacian pyramid

Gaussian Pyramid

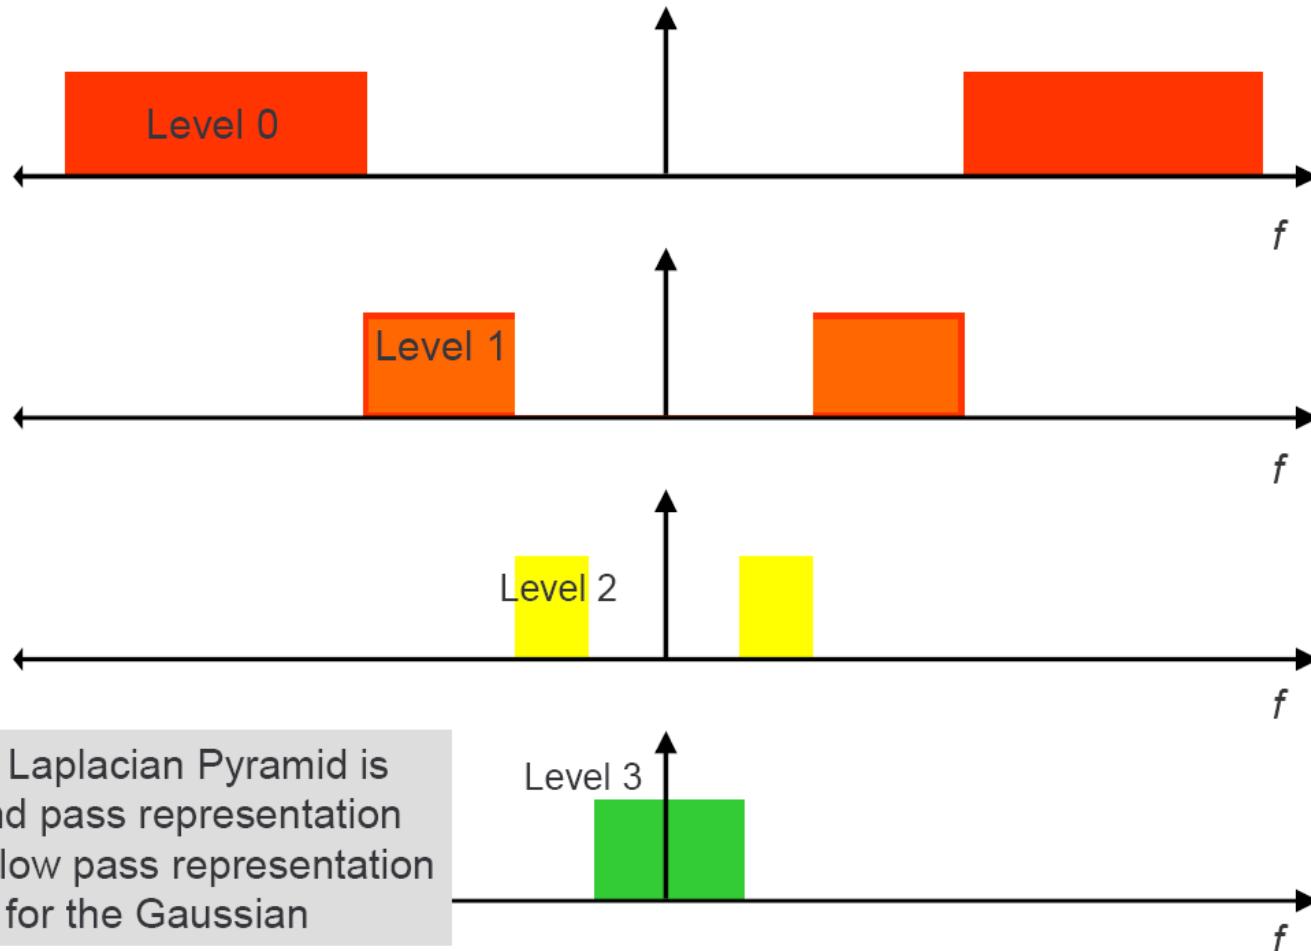
$$G_i = L_i + \text{expand}(G_{i+1})$$

Laplacian Pyramid



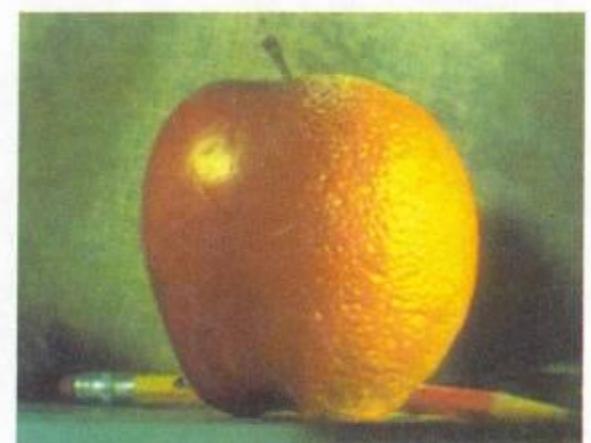
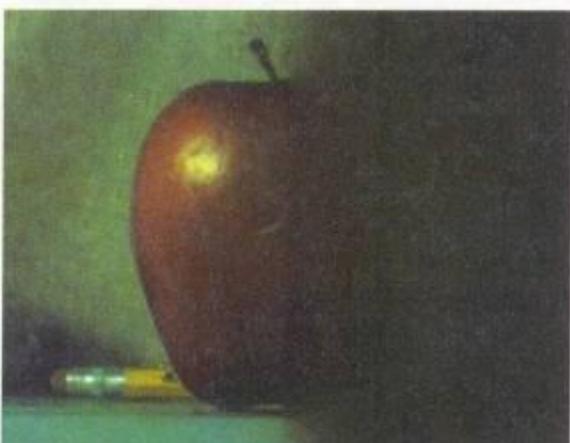
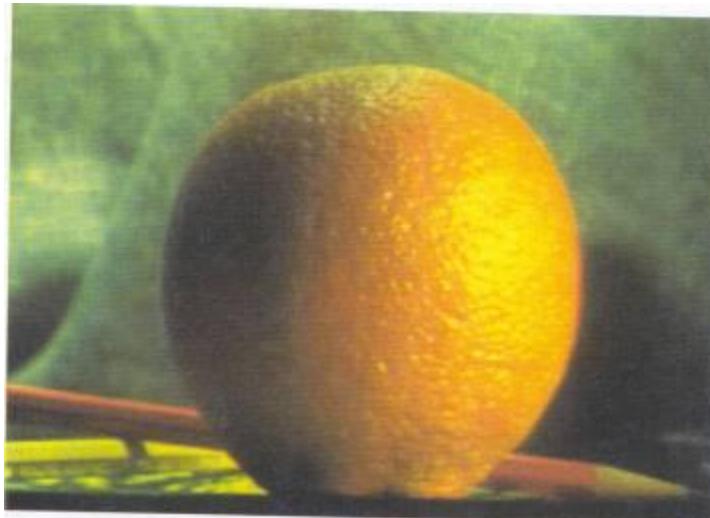
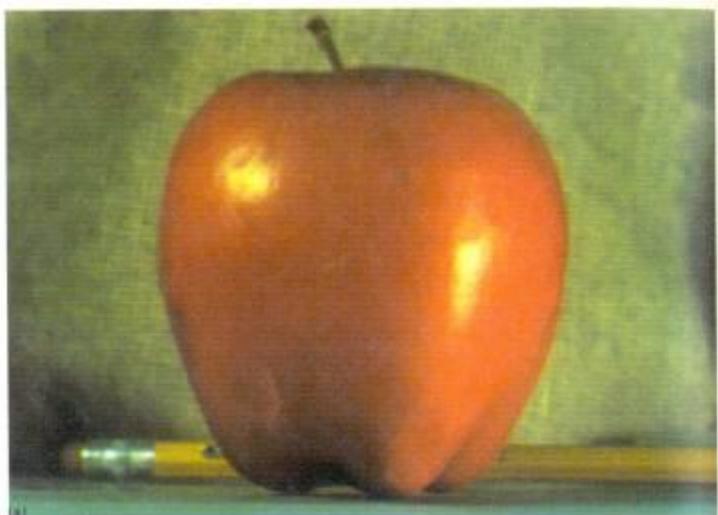
# Frequencies

## Laplacian Pyramid Frequency Composition



Source: S. Narasimhan

# Pyramid Blending



Source: A. Efros

# Finite difference filters

Other approximations of derivative filters exist:

Prewitt:  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

Sobel:  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

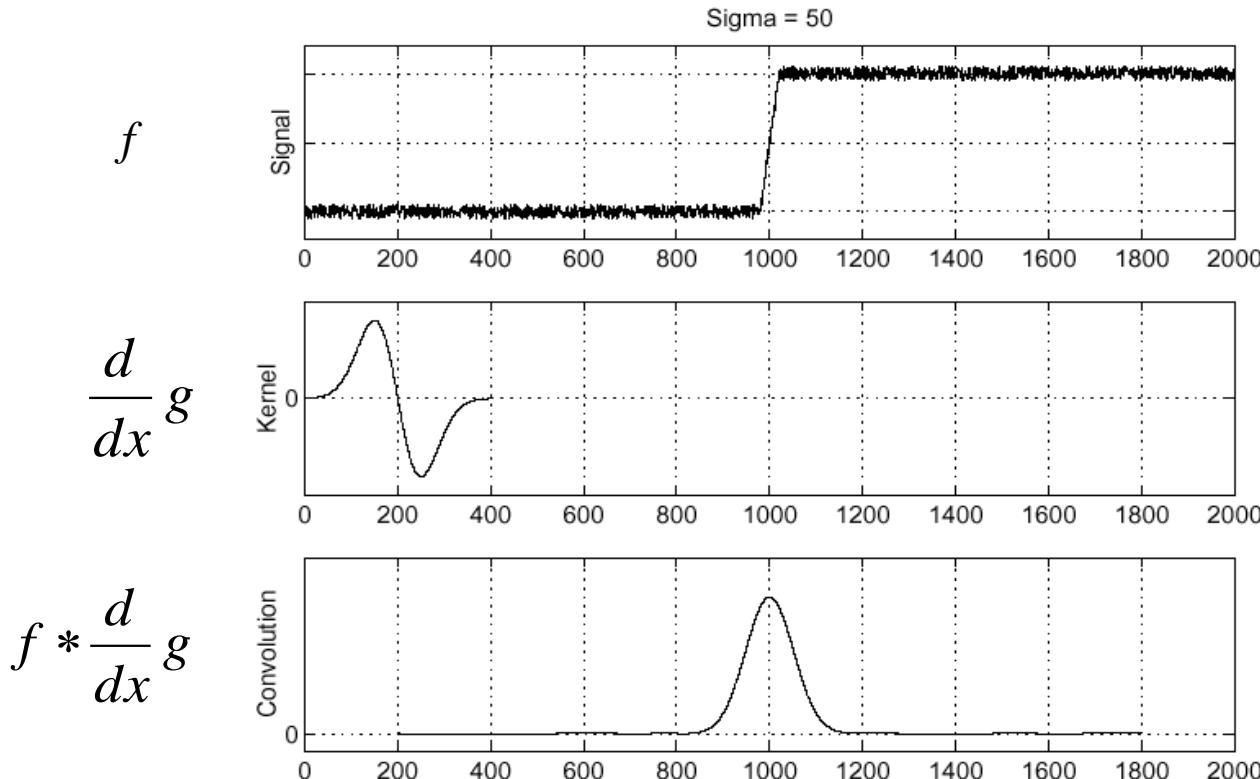
Roberts:  $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

# Derivative theorem of convolution

Differentiation is convolution,  
and convolution is associative:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

This saves us one operation:



Source: S. Seitz

# The Canny edge detector



original image



high threshold  
(strong edges)



low threshold  
(weak edges)



hysteresis threshold

Source: L. Fei-Fei

# Canny edge detector

1. Filter image with derivative of Gaussian
2. Find magnitude and orientation of gradient
3. **Non-maximum suppression:**
  - Thin wide “ridges” down to single pixel width
4. **Linking and thresholding (hysteresis):**
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them

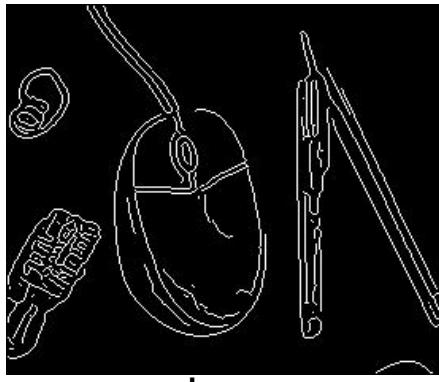
J. Canny, ***A Computational Approach To Edge Detection***, IEEE Trans. Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

Source: S. Lazebnik

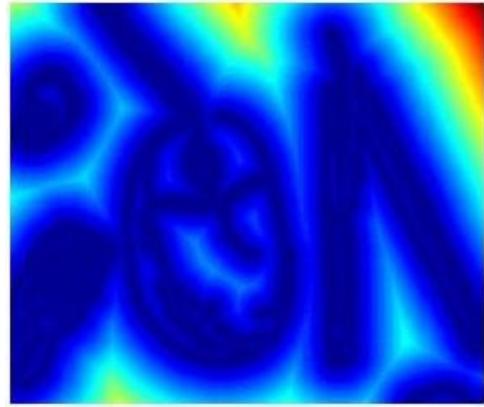
# Distance transform



original



edges



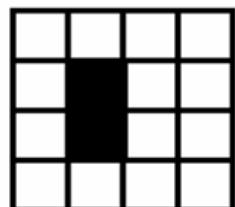
distance transform

Value at  $(x,y)$  tells how far that position is from the nearest edge point (or other binary mage structure)

# Distance Transform (2D)

- 2D case analogous to 1D
  - Initialization
  - Forward and backward pass
    - Fwd pass finds closest above and to left
    - Bwd pass finds closest below and to right

-	1
1	0
0	1
1	-



$\infty$	$\infty$	$\infty$	$\infty$
$\infty$	0	$\infty$	$\infty$
$\infty$	0	$\infty$	$\infty$
$\infty$	$\infty$	$\infty$	$\infty$

$\infty$	$\infty$	$\infty$	$\infty$
$\infty$	0	1	$\infty$
$\infty$	0	$\infty$	$\infty$
$\infty$	$\infty$	$\infty$	$\infty$

$\infty$	$\infty$	$\infty$	$\infty$
$\infty$	0	1	2
$\infty$	0	1	2
$\infty$	1	2	3

2	1	2	3
1	0	1	2
1	0	1	2
2	1	2	3

Source: D. Huttenlocher

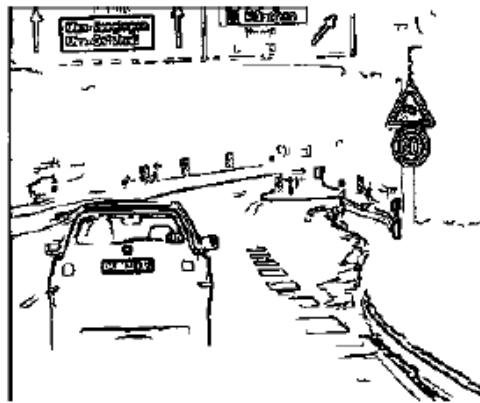
# Chamfer distance

- Average distance to nearest feature

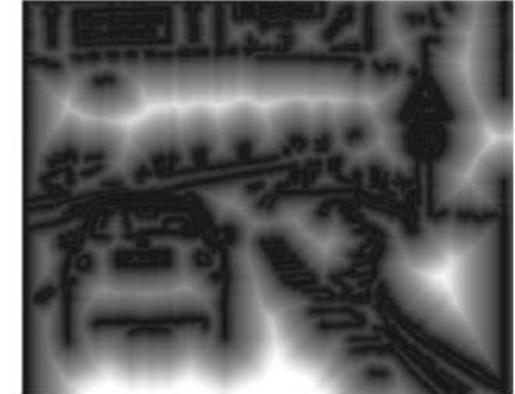
$$D_{chamfer}(T, I) \equiv \frac{1}{|T|} \sum_{t \in T} d_I(t)$$

- $T$ : template shape → a set of points
- $I$ : image to search → a set of points
- $d_I(t)$ : min distance for point  $t$  to some point in  $I$

# Chamfer distance

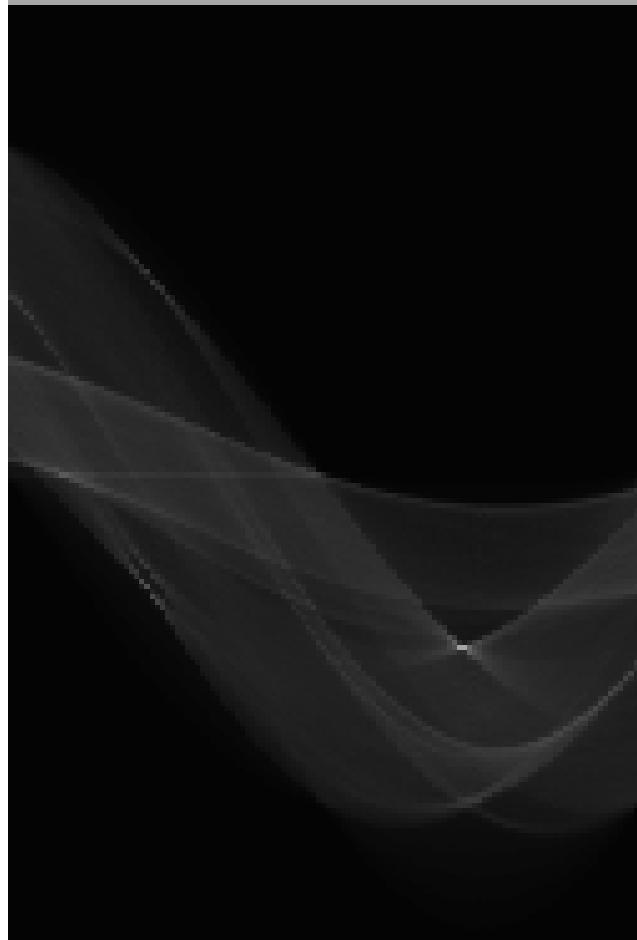
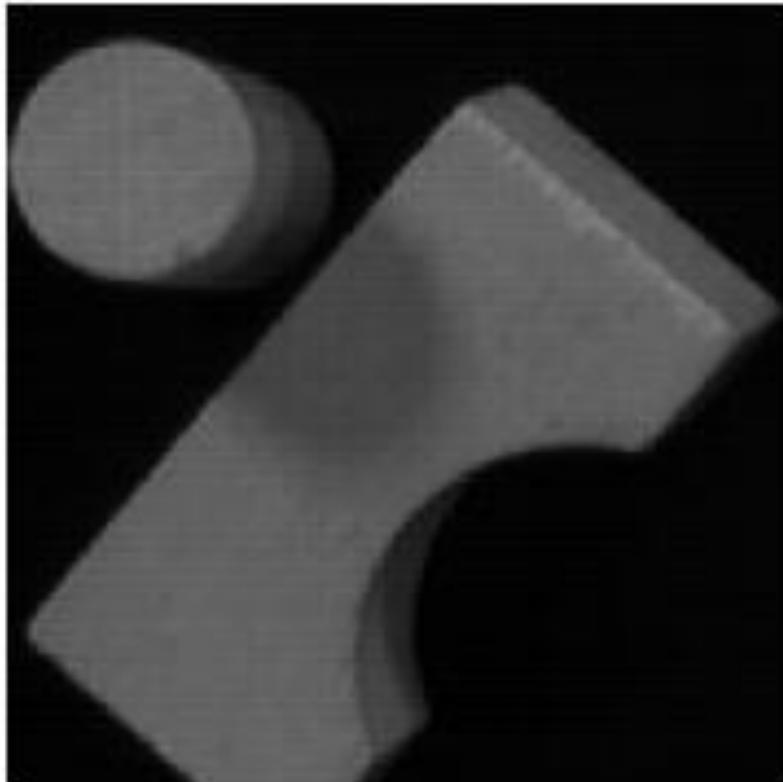


Edge image



Distance transform image

# Hough transform



Source: K. Grauman

# Hough transform algorithm

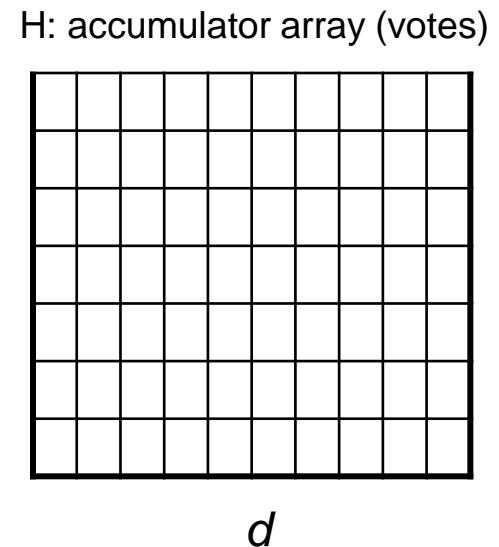
Using the polar parameterization:

$$x \cos \theta - y \sin \theta = d$$

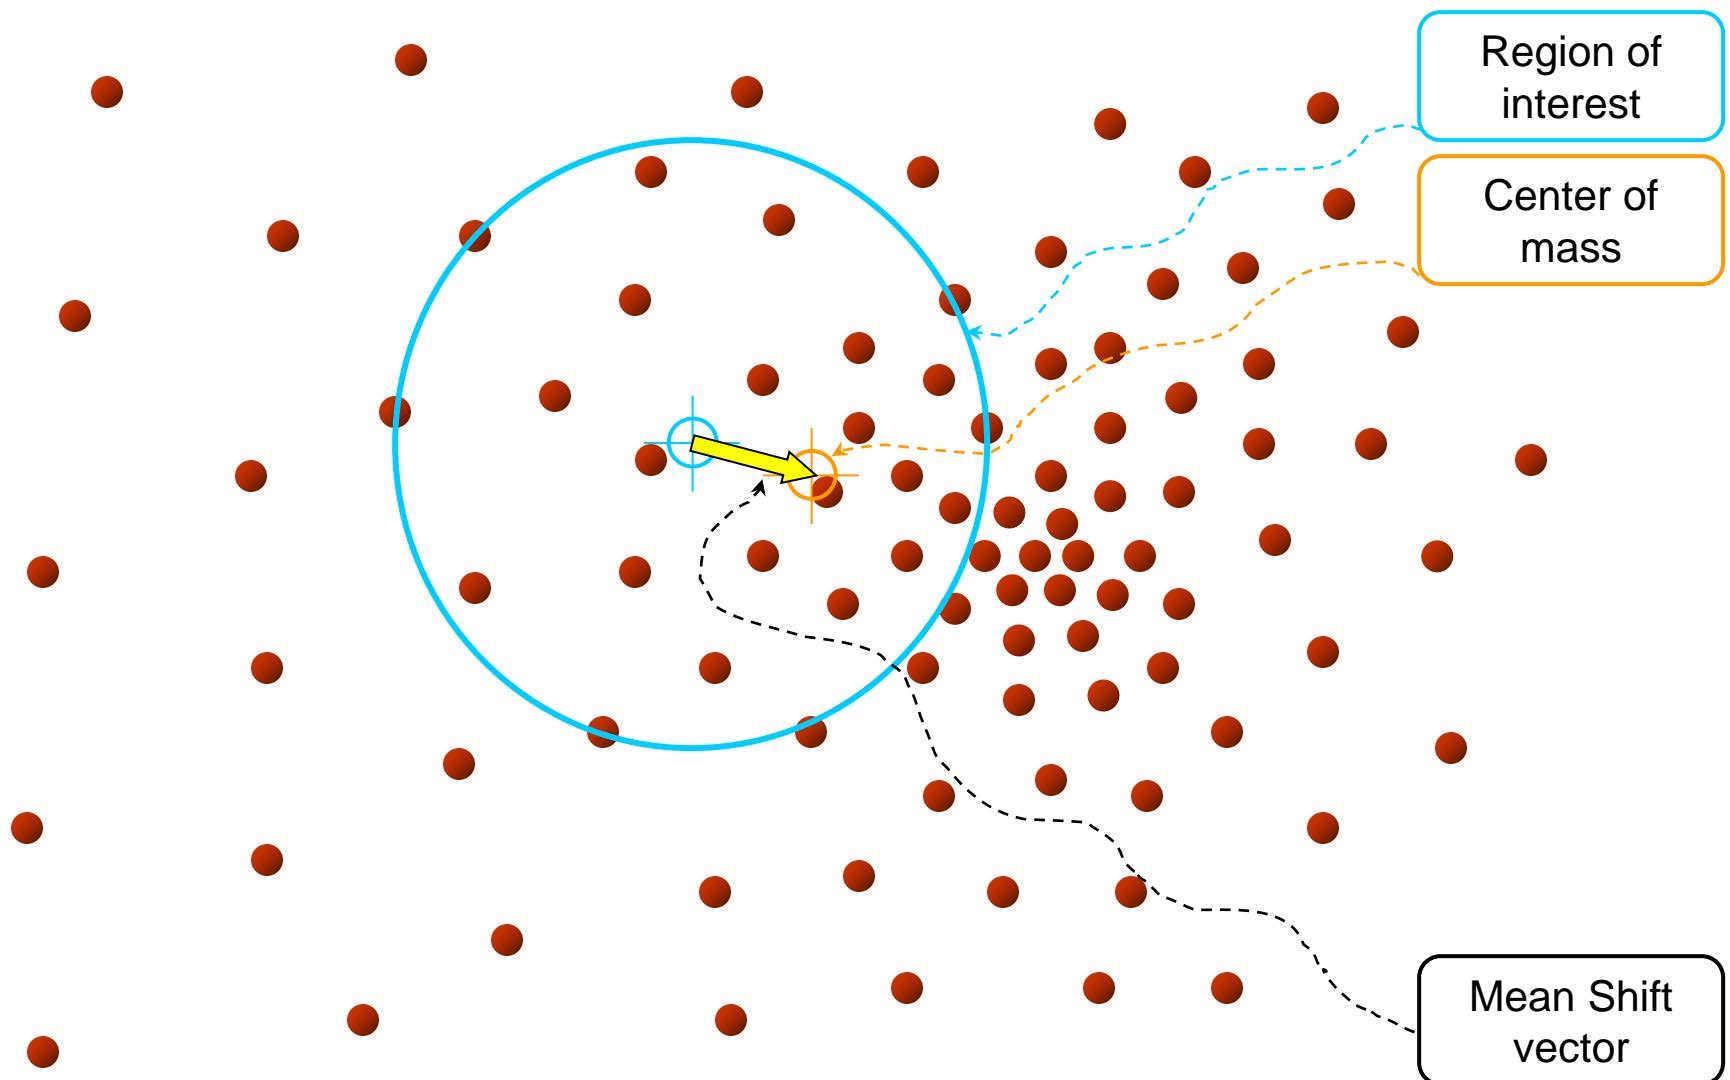
Basic Hough transform algorithm

1. Initialize  $H[d, \theta] = 0$
2. for each edge point  $I[x, y]$  in the image
  - for  $\theta = 0$  to  $180$  // some quantization
 
$$d = x \cos \theta - y \sin \theta$$

$$H[d, \theta] += 1$$
3. Find the value(s) of  $(d, \theta)$  where  $H[d, \theta]$  is maximum
4. The detected line in the image is given by  $d = x \cos \theta - y \sin \theta$



# Mean Shift



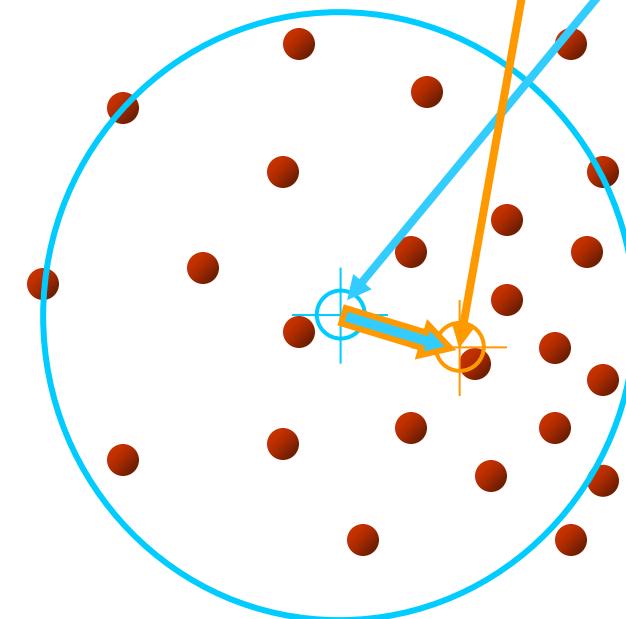
**Objective : Find the densest region**  
Distribution of identical billiard balls

Source: Y. Ukrainitz and B. Sarel

# Mean shift

$$\nabla P(\mathbf{x}) = \frac{c}{n} \sum_{i=1}^n \nabla k_i = \frac{2c}{N} \sum_i g\left(\|\mathbf{x} - \mathbf{x}_i\|^2\right)$$

$$\left( \frac{\sum_i x_i g\left(\|\mathbf{x} - \mathbf{x}_i\|^2\right)}{\sum_i g\left(\|\mathbf{x} - \mathbf{x}_i\|^2\right)} - \mathbf{x} \right)$$



## Simple Mean Shift procedure:

- Compute mean shift vector

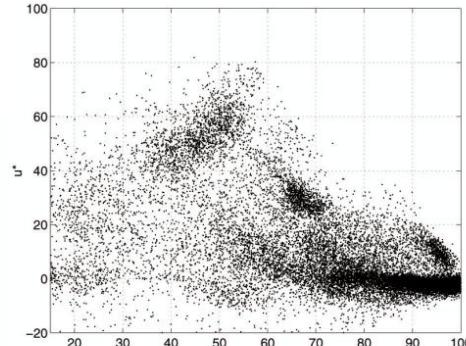
$$\mathbf{m}(\mathbf{x}) = \left[ \frac{\sum_{i=1}^n \mathbf{x}_i g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)}{\sum_{i=1}^n g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)} - \mathbf{x} \right]$$

- Translate the Kernel window by  $\mathbf{m}(\mathbf{x})$

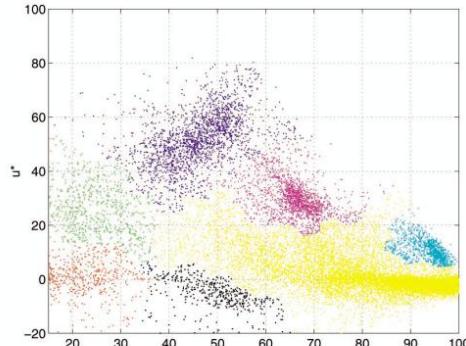
Source: Y. Ukrainitz and B. Sarel

# Mean shift clustering/segmentation

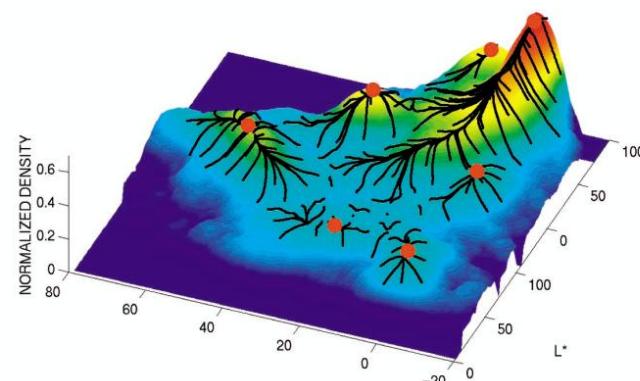
- Find features (color, gradients, texture, etc)
- Initialize windows at individual feature points
- Perform mean shift for each window until convergence
- Merge windows that end up near the same “peak” or mode



(a)



(b)



Source: K. Grauman

# Mean Shift Tracking

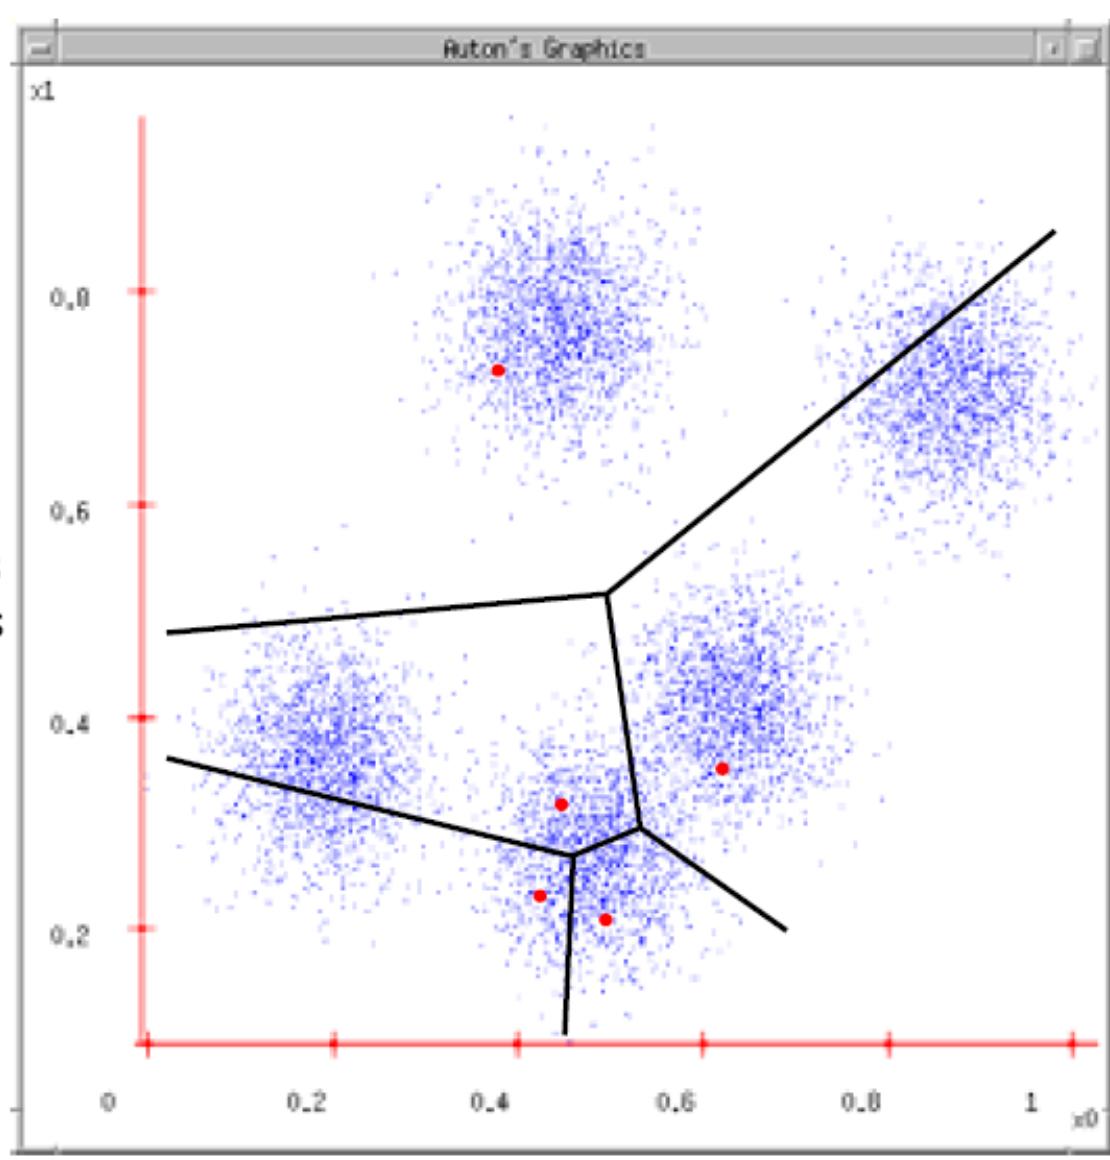
Goal: Mark object  
in first frame and  
locate it in all  
frames



D. Comaniciu. **Real-Time  
Tracking of Non-Rigid  
Objects using Mean Shift.**  
CVPR 2000

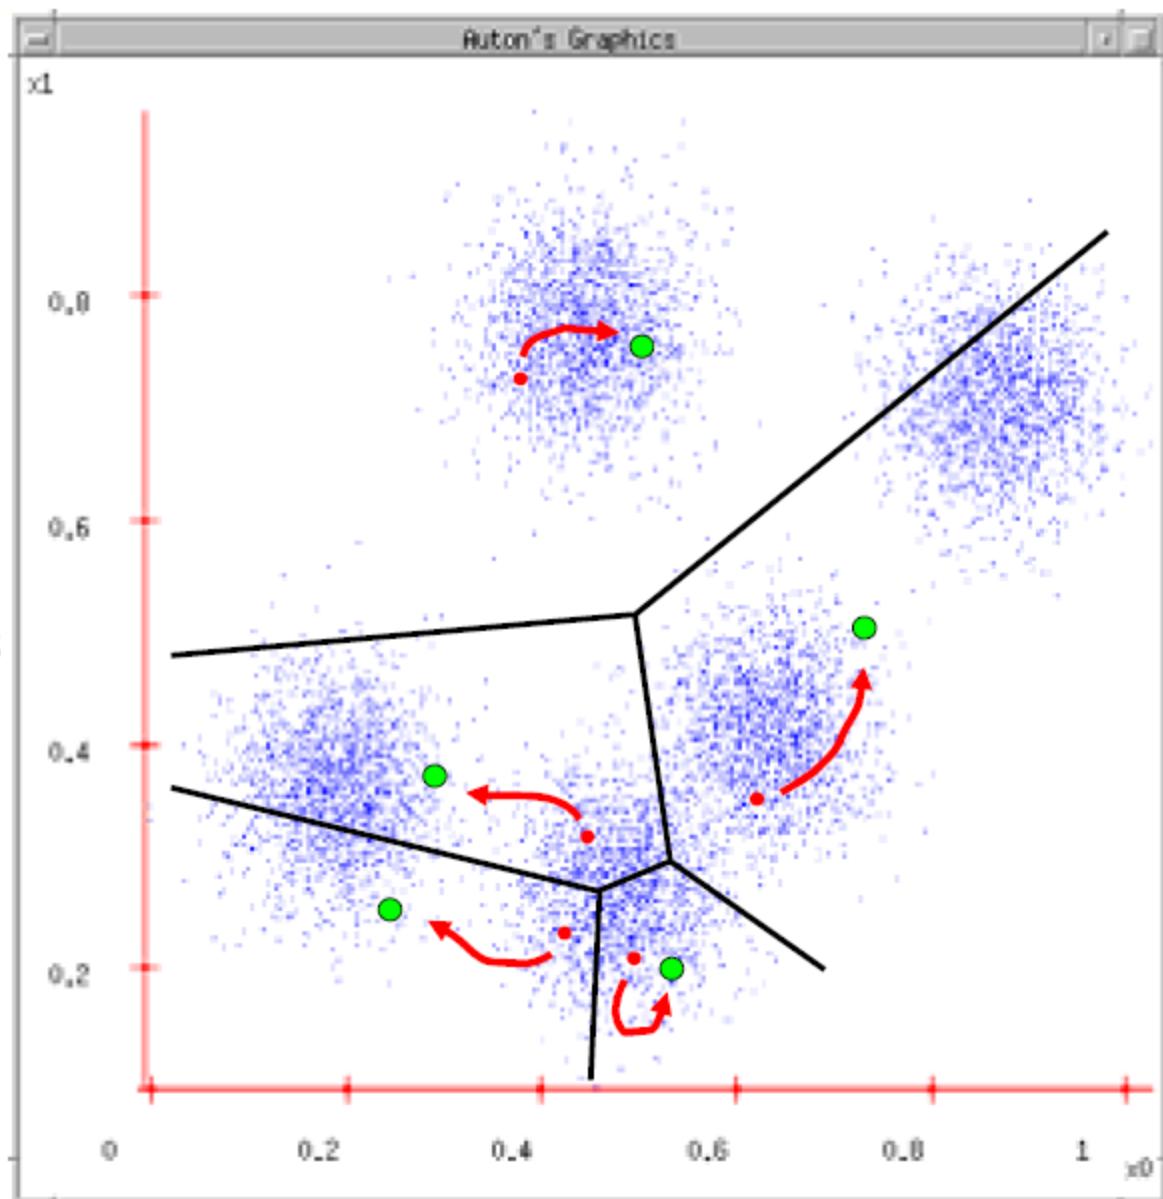
# K-means

1. Ask user how many clusters they'd like.  
*(e.g. k=5)*
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to. (Thus each Center "owns" a set of datapoints)



# K-means

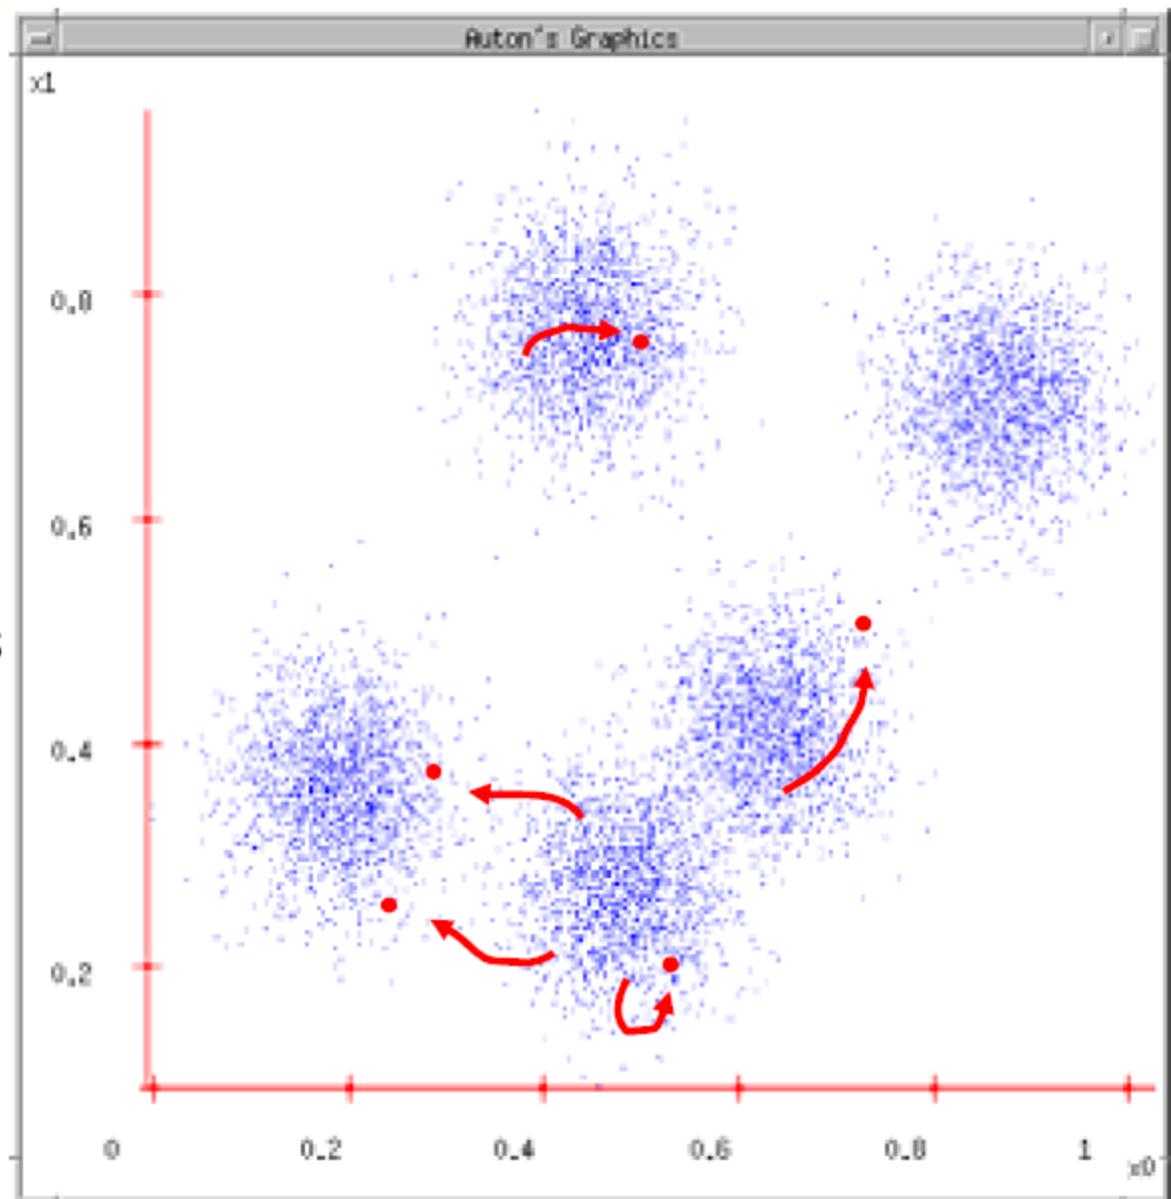
1. Ask user how many clusters they'd like.  
*(e.g. k=5)*
2. Randomly guess k cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns



Source: A. Moore

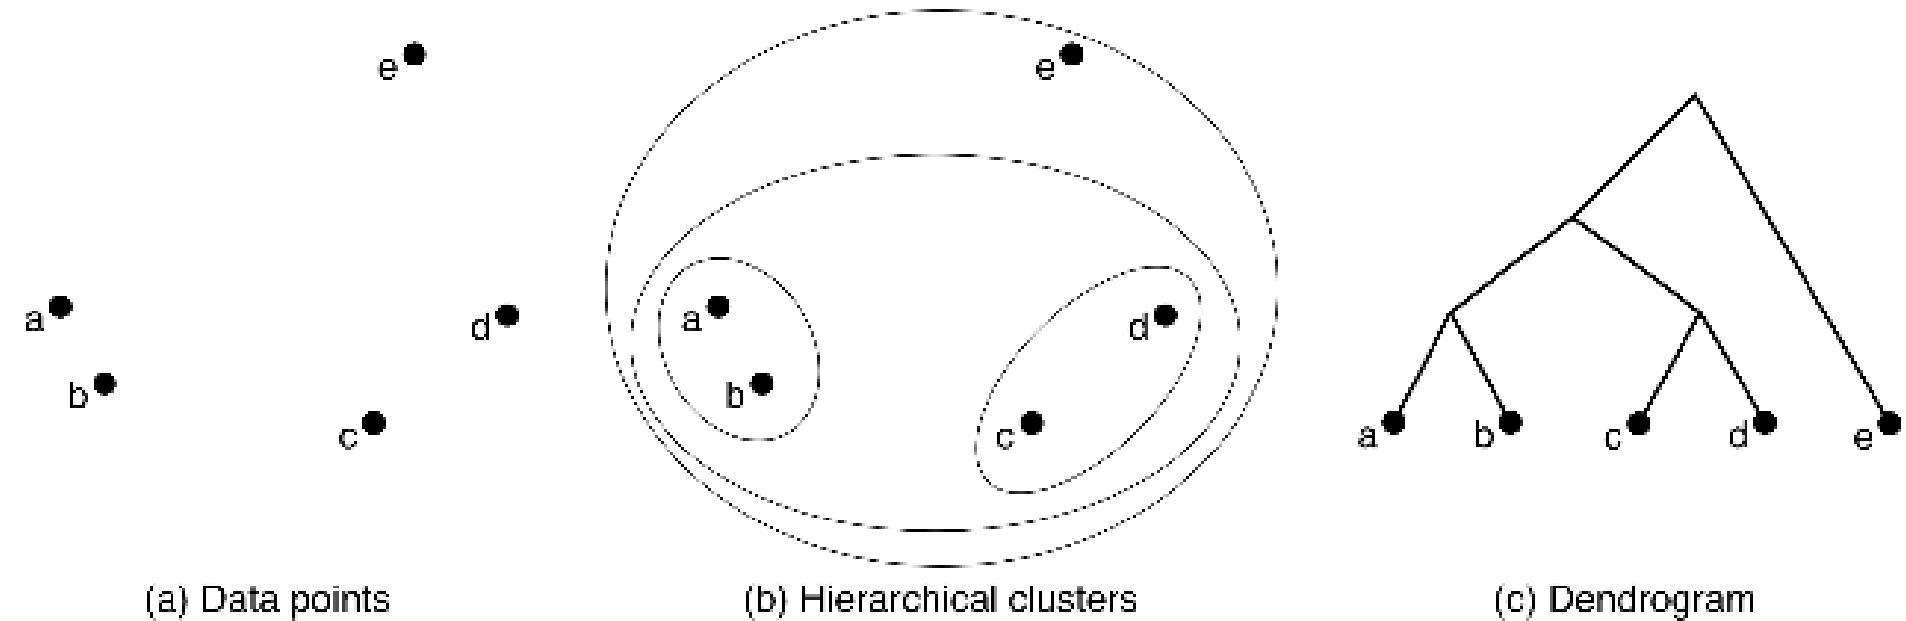
# K-means

1. Ask user how many clusters they'd like.  
*(e.g.  $k=5$ )*
2. Randomly guess  $k$  cluster Center locations
3. Each datapoint finds out which Center it's closest to.
4. Each Center finds the centroid of the points it owns...
5. ...and jumps there
6. ...Repeat until terminated!



Source: A. Moore

# Agglomerative Clustering



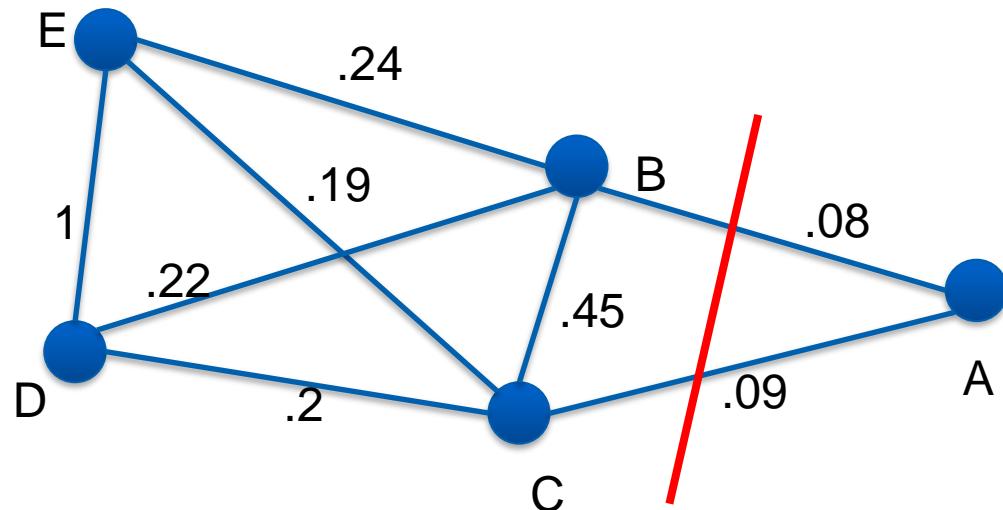
(a) Data points

(b) Hierarchical clusters

(c) Dendrogram

# Spectral Clustering Example

- Minimum Cut

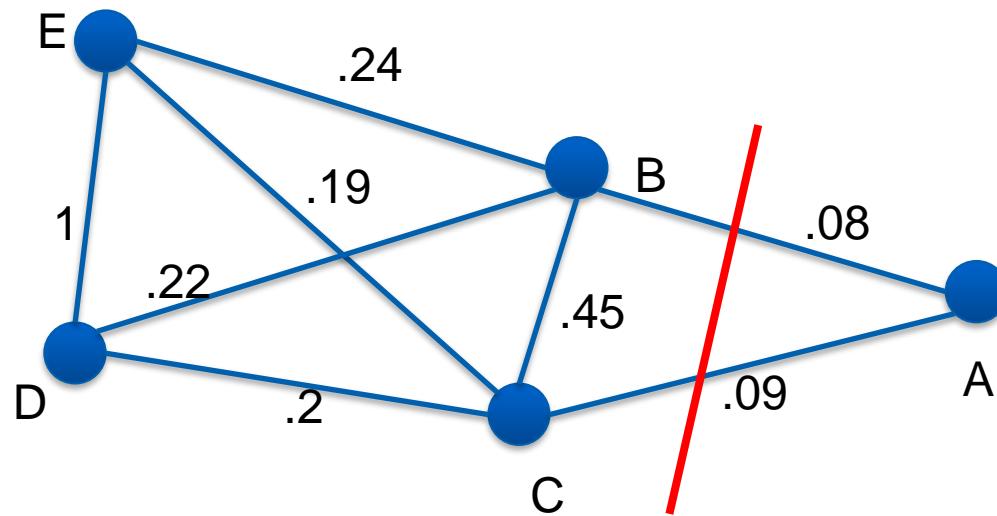


$$\text{Cut}(\text{BCDE}, \text{A}) = 0.08 + 0.09 = 0.17$$

# Spectral Clustering Example

- Normalized Minimum Cut

$$NormCut(C_1, C_2) = \frac{Cut(C_1, C_2)}{Vol(C_1)} + \frac{Cut(C_1, C_2)}{Vol(C_2)}$$

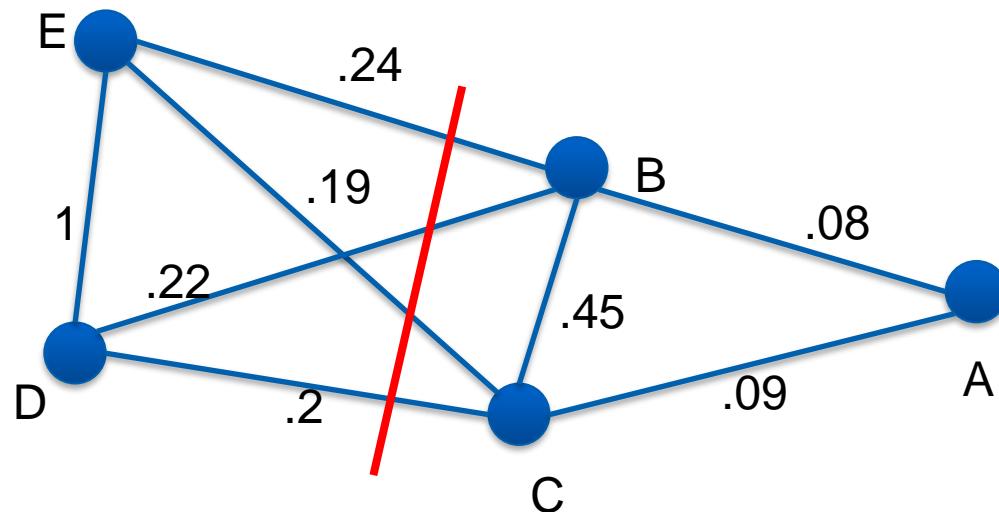


$$Cut(BCDE, A) = 0.17 / (0.99 + 0.93 + 1.42 + 1.43) + 0.17 / 0.17 = 1.036$$

# Spectral Clustering Example

- Normalized Minimum Cut

$$NormCut(C_1, C_2) = \frac{Cut(C_1, C_2)}{Vol(C_1)} + \frac{Cut(C_1, C_2)}{Vol(C_2)}$$

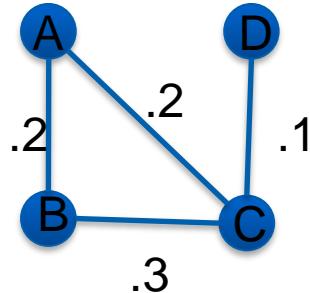


$$Cut(BCDE, A) = 0.17 / (0.99 + 0.93 + 1.42 + 1.43) + 0.17 / 0.17 = 1.036$$

$$Cut(DE, ABC) = 0.85 / (1.42 + 1.43) + 0.85 / (0.17 + 0.99 + 0.93) = 0.705$$

# Spectral Clustering

- Construct an **affinity matrix**



W

	A	B	C	D
A	0	.2	.2	0
B	.2	0	.3	0
C	.2	.3	0	.1
D	0	0	.1	0

- Graph Laplacian

$$L = D - W$$

	A	B	C	D
A	.4	-.2	-.2	-0
B	-.2	.5	-.3	-0
C	-.2	-.3	.6	-.1
D	0	0	-.1	.1

# Spectral Clustering

Normalized Minimum Cut (Relaxed):

$$\min_x Ncut(\mathbf{x}) = \min_y \frac{\mathbf{y}^T (\mathbf{D} - \mathbf{W}) \mathbf{y}}{\mathbf{y}^T \mathbf{D} \mathbf{y}}$$

Generalized eigenvalue problem:

$$(\mathbf{D} - \mathbf{W})\mathbf{y} = \lambda \mathbf{D}\mathbf{y}$$

Equal eigenvalue problem:

$$\mathbf{D}^{-\frac{1}{2}}(\mathbf{D} - \mathbf{W})\mathbf{D}^{-\frac{1}{2}}\mathbf{z} = \lambda \mathbf{z} \quad \mathbf{z} = \mathbf{D}^{\frac{1}{2}}\mathbf{y}$$

Eigenvector corresponding to the second smallest eigenvalue

# More than two clusters, Laplacian eigenmaps

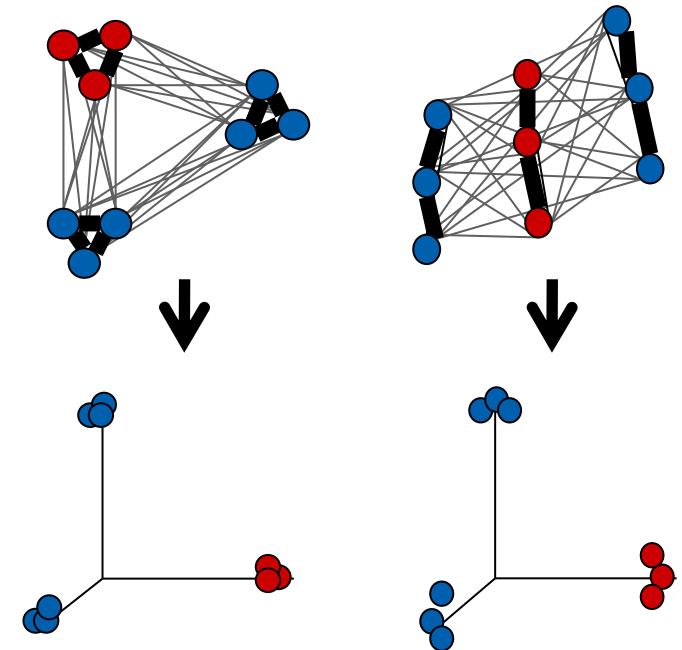
The eigenvector to the third smallest eigenvalue indicates an alternative partitioning of the graph

In the general case we compute the  $k$  smallest eigenvalues

The corresponding eigenvectors span a subspace

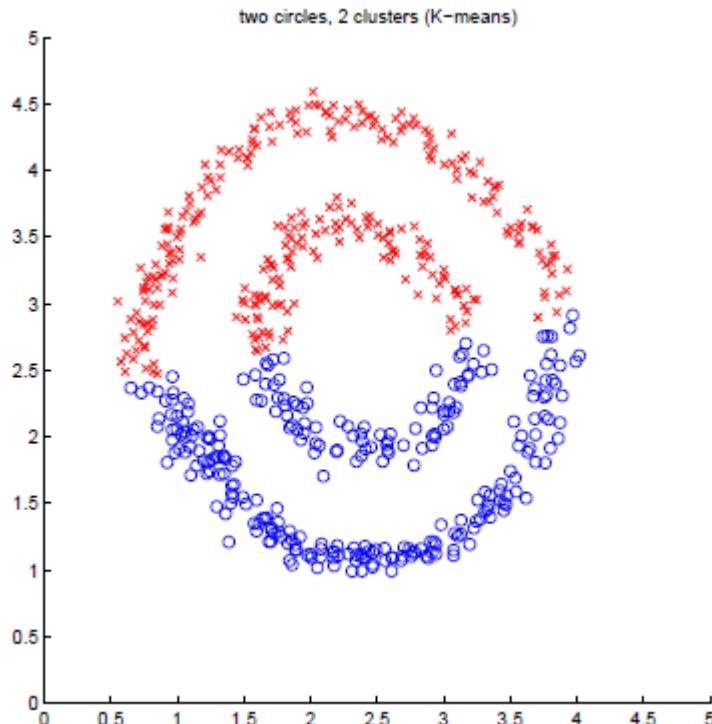
- $Y = (y_1, y_2, \dots, y_k) \in \mathbb{R}^{N \times k}$  matrix
- Each data point  $1, 2, \dots, N$  maps to a  $k$ -dimensional vector (i-th row of  $Y$ )
- Mapping is called **Laplacian eigenmap**

Standard clustering techniques (k-means) to convert the real-valued vectors into integer labels  
→ **spectral clustering**

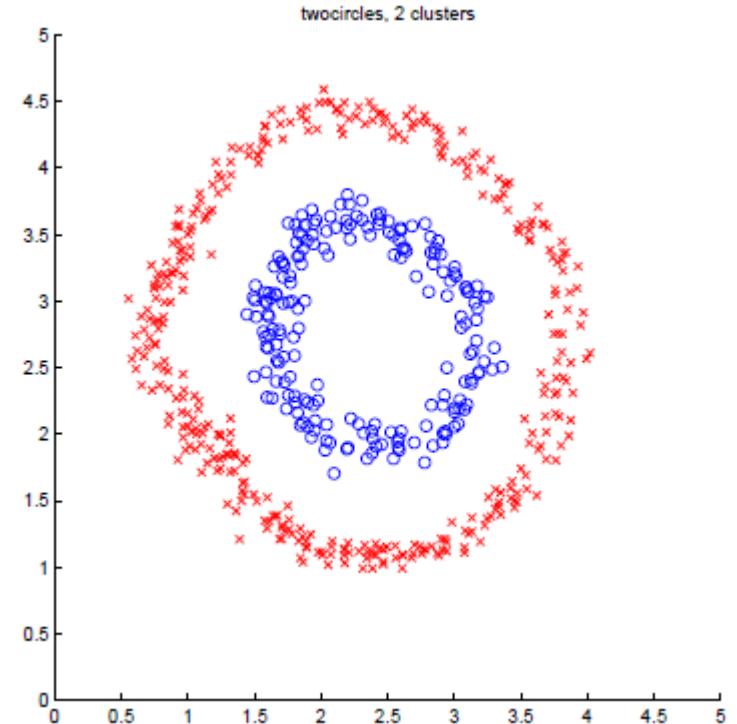


Mappings from a graph to its eigenspace, clustering in the eigenspace

# K means vs. spectral clustering



K means



Spectral clustering

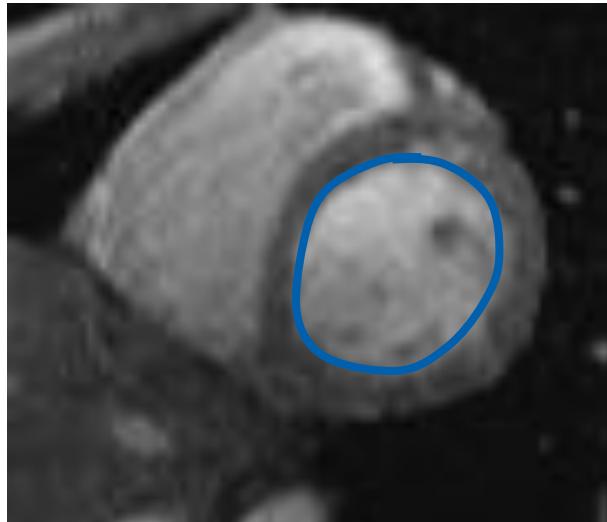
A. Ng et al. On spectral clustering: Analysis and an algorithm. NIPS 2001

# Deformable contours

a.k.a. active contours, snakes

**Given:** initial contour (model) near desired object

**Goal:** evolve the contour to fit exact object boundary



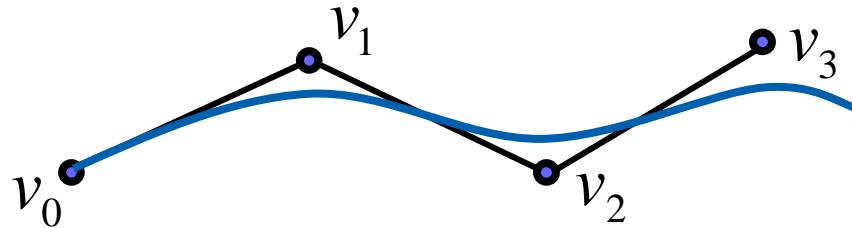
**Main idea:** elastic band is iteratively adjusted so as to

- be near image positions with high gradients, **and**
- satisfy shape “preferences” or contour priors

[ Snakes: Active contour models, Kass, Witkin, & Terzopoulos, ICCV1987 ]

# B-spline snakes

Explicit representation of contour:  $C(s) = \sum_{i=0}^{N-1} v_i \cdot B_i(s)$



Smooth (differentiable) curve passes close to the poly-line  
joining the knot (control) points

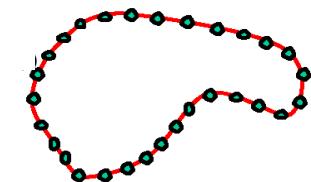
*B-spline snakes* - parametric representation of continuous  
contours via finite number of parameters:

$$\{v_0, v_1, v_2, \dots, v_{N-1}\}$$

# Energy function

The total energy (cost) of the current snake is defined as:

$$E_{total} = E_{internal} + E_{external}$$



**Internal** energy: encourage *prior* shape preferences: e.g., smoothness, elasticity, particular known shape.

**External** energy (“image” energy): encourage contour to fit on places where image structures exist, e.g., edges.

A good fit between the current deformable contour and the target shape in the image will yield a **low** value for this cost function.

# Total energy: function of the weights

$$E_{total} = E_{internal} + \gamma E_{external}$$

$$E_{external} = - \sum_{i=0}^{n-1} |G_x(x_i, y_i)|^2 + |G_y(x_i, y_i)|^2$$

$$E_{internal} = \sum_{i=0}^{n-1} \alpha (\bar{d} - \|v_{i+1} - v_i\|)^2 + \beta \|v_{i+1} - 2v_i + v_{i-1}\|^2$$

# The Snakes model (continuous)

Kass, Witkin, and Terzopoulos proposed the following energy functional:

$$E(C) = - \int_0^1 |\nabla I(C(s))|^2 ds + \alpha \int_0^1 |C_s(s)|^2 ds + \beta \int_0^1 |C_{ss}(s)|^2 ds$$

External energy (data)    Internal energy (prior)

$C : [0, 1] \rightarrow \Omega$  is a parametric contour.  $C_s$  and  $C_{ss}$  denote the first and second derivative of this contour.

The first term is called **external energy**, since it depends on the (external) input image. The second two terms are called **internal energy**, since they are inherent to the model and independent of the data.

Minimizing the external energy drives the contour to go along maxima of the gradient.

[ M. Kass, A. Witkin, D. Terzopoulos. Snakes: Active contour models. IJCV 1987 ]

Source: T. Brox

# Energy minimization

- Several algorithms have been proposed to fit deformable contours.
- We'll look at two:
  - Greedy search
  - Dynamic programming (for 2d snakes)
  - Variational optimization

# Dynamic programming

- General minimization problem with unary  $U_n$  and pairwise terms  $P_n$ :

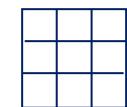
$$\operatorname{argmin}_{w_1 \dots N} \left[ \sum_{n=1}^N U_n(w_n) + \sum_{n=1}^{N-1} P_n(w_n, w_{n+1}) \right]$$

- Snakes with control points  $w_n$ :

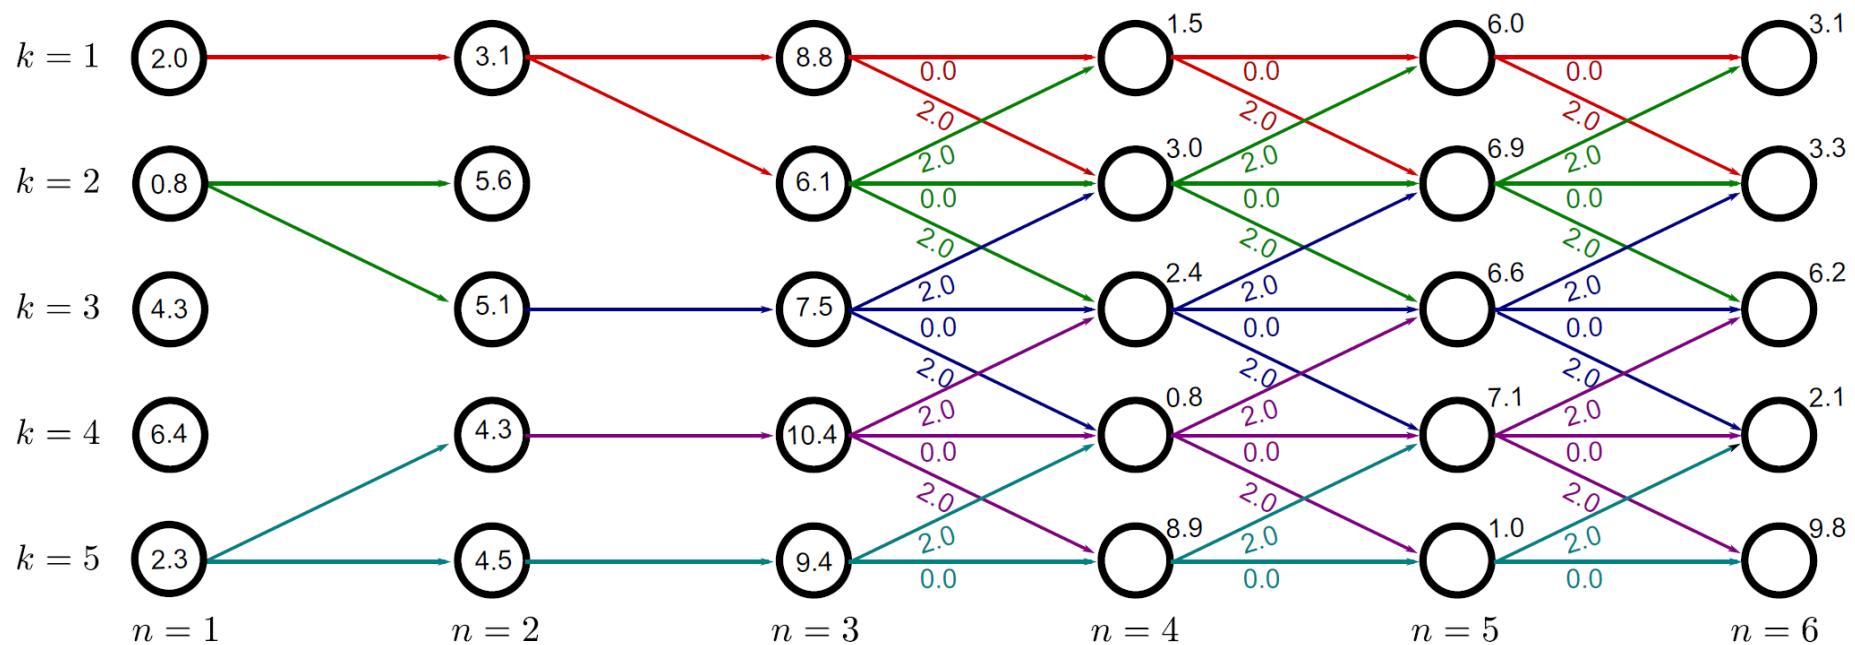
$$U_n(w_n) = -\|G(w_n)\|^2$$

$$P_n(w_n, w_{n+1}) = \alpha \|w_{n+1} - w_n\|^2$$

- States of  $w_n=k$  are pixels in neighborhood:

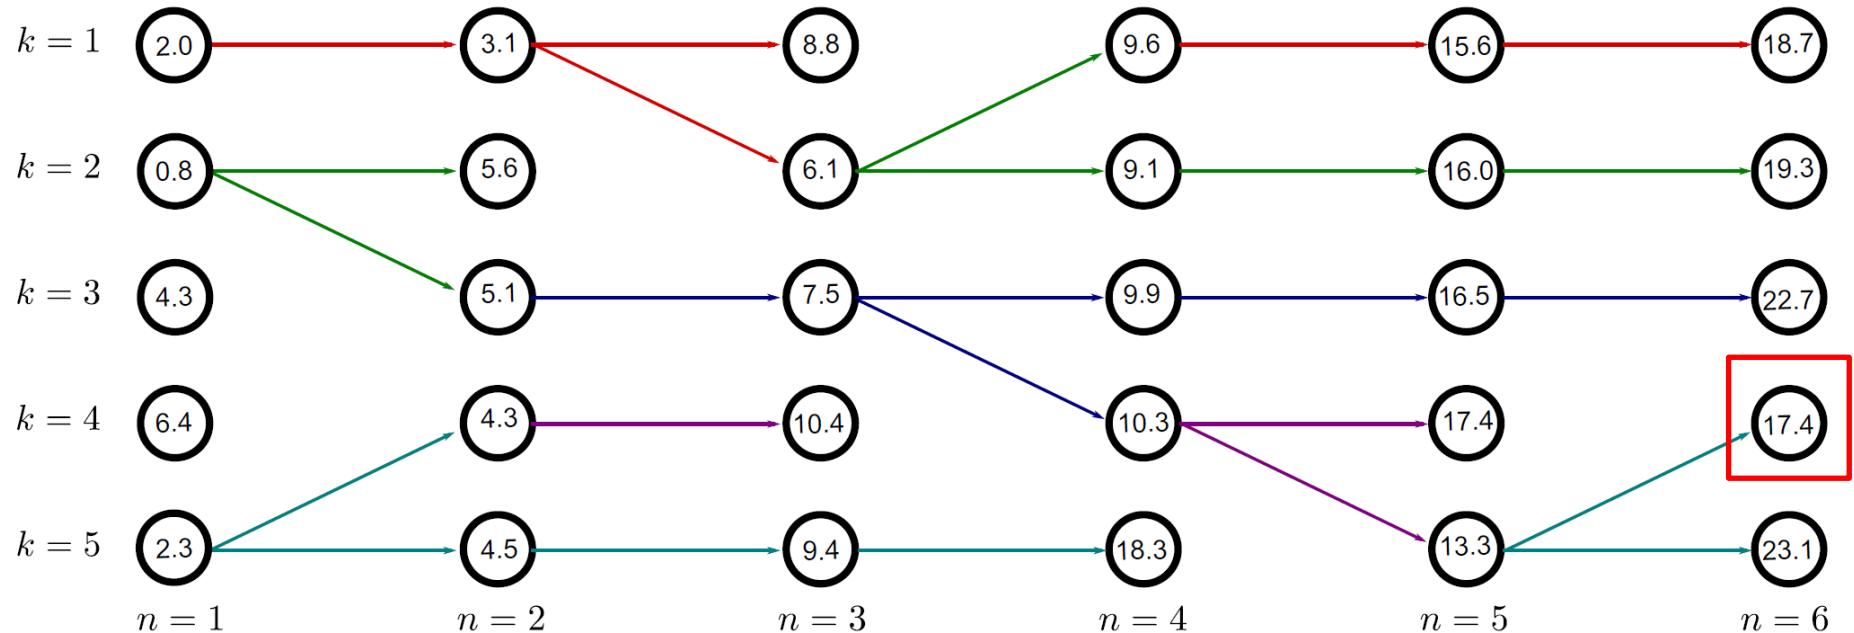


# Worked example



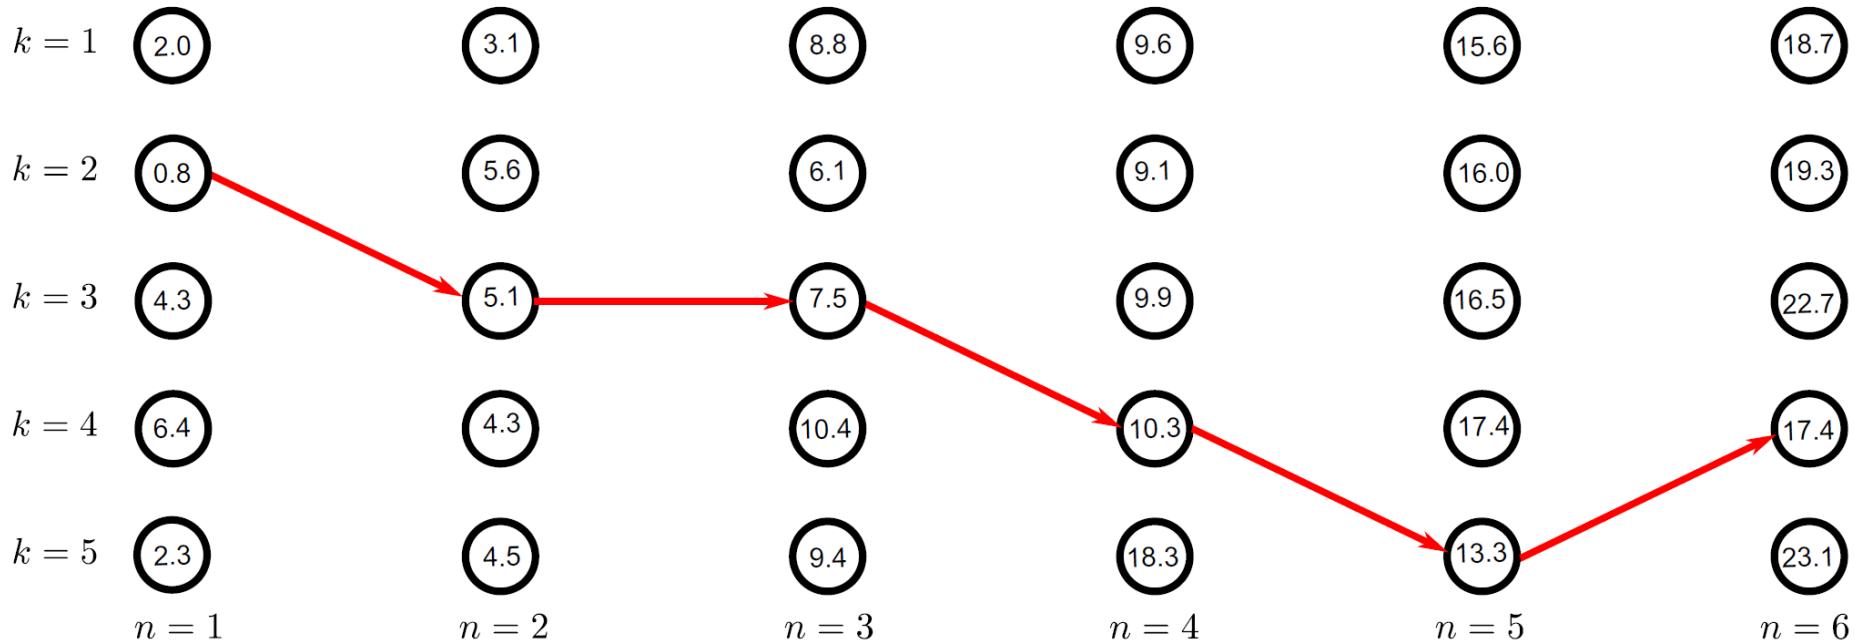
Work through the graph, computing the minimum cost to reach each node

# Worked example



Find the minimum possible cost to reach the final column

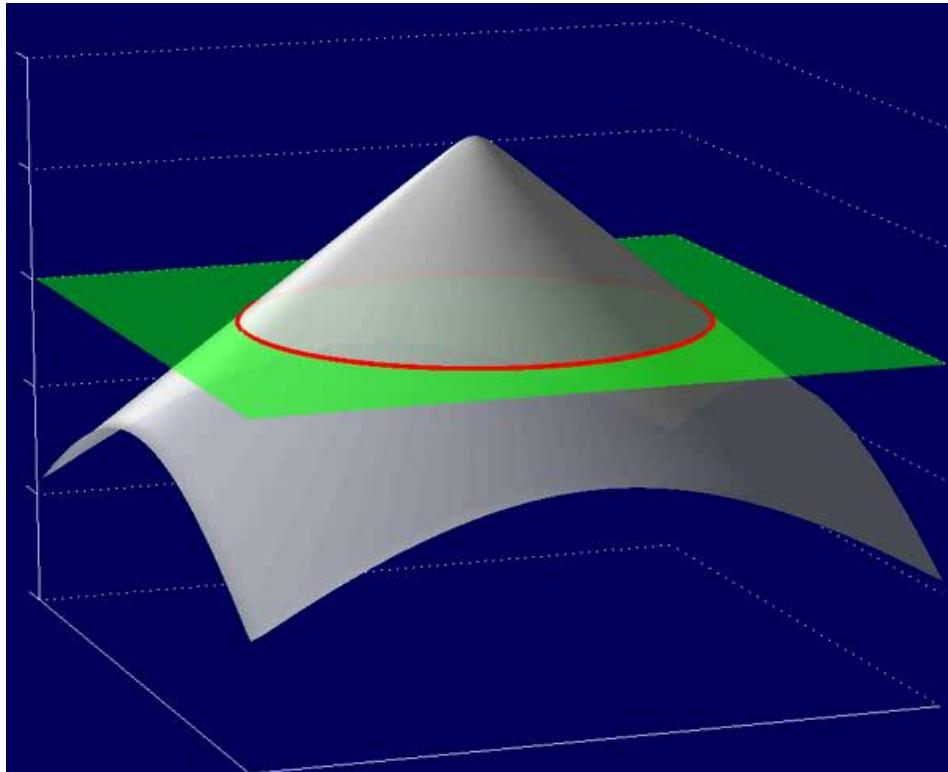
# Worked example



Trace back the route that we arrived here by – this is the minimum configuration

# Level set representation of contours

- Introduce an **embedding function**  $\phi : \Omega \rightarrow \mathbb{R}$
- The zero-level line represents the contour (can be any other level)  
$$C = \{\mathbf{x} \in \Omega \mid \phi(\mathbf{x}) = 0\}$$
- For evolving  $C$  evolve  $\phi$
- No regridding necessary
- Allows for topological changes
- Can be applied in any dimension
- Represents the contour *and* the enclosed region



# Geodesic active contours

Caselles et al. **Geodesic Active Contours**. IJCV 1997 introduced two improvements of the snake model.

1. A level set representation of the contour

2. A variation of the model, called the **geodesic active contour**:

$$E(C) = \int_0^1 w(|\nabla I(C(s))|^2) |C_s(s)| ds = \int_C w(|\nabla I|^2) ds$$

It consists of a single term that penalizes the **geodesic length** of the contour using a metric induced by the image gradient.

For comparison: measuring the Euclidean length of a contour reads:

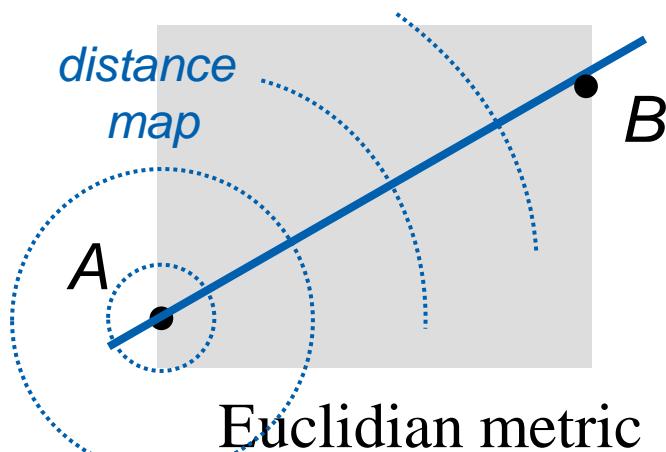
$$E(C) = \int_0^1 |C_s(s)| ds = \int_C ds$$

→ Combines the internal and external energy of the Snake model

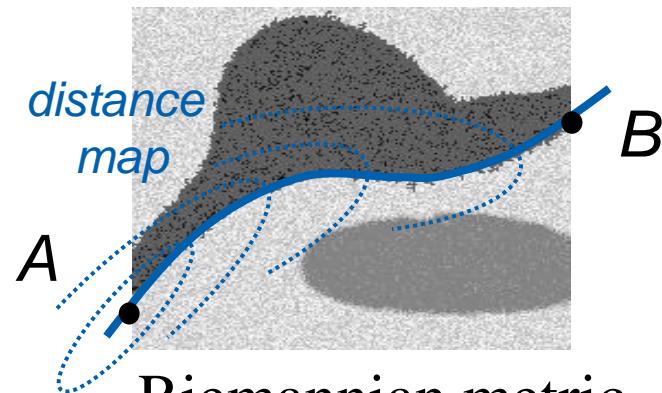
Source: T. Brox

# Geodesic active contours

- Energy  $E(C) = \int g(|\nabla I_{C(s)}|)ds$  = “weighted length” of  $C$
- *Geodesic*: Shortest curve between two points (wrt. E).

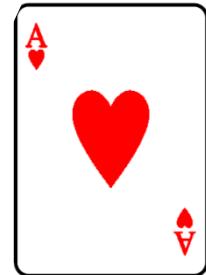


$E(C)$  = Euclidean length of  $C$

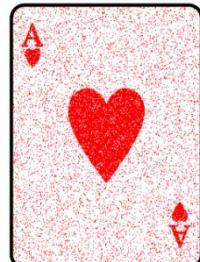


$E(C)$  = image-weighted length of  $C$

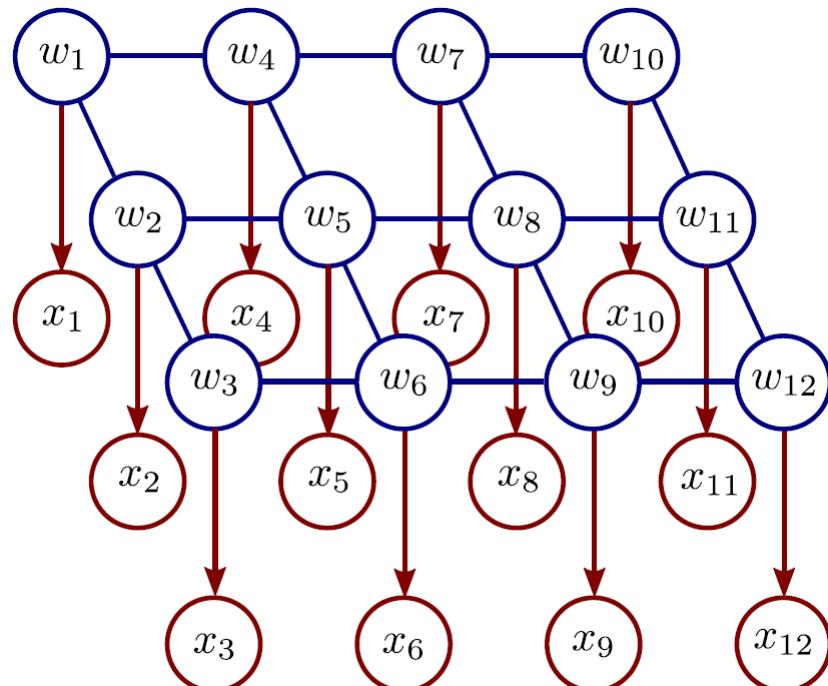
# Denoising with MRFs



Original image,  $\mathbf{w}$



Observed image,  $\mathbf{x}$



Inference :

$$Pr(w_1 \dots w_N | x_1 \dots x_N) = \frac{\prod_{n=1}^N Pr(x_n | w_n) Pr(w_1 \dots w_N)}{Pr(x_1 \dots x_N)}$$

MRF Prior (pairwise cliques)

$$Pr(w_1 \dots w_N) = \frac{1}{Z} \exp \left[ - \sum_{(m,n) \in \mathcal{C}} \psi[w_m, w_n, \theta] \right]$$

Likelihoods

$$Pr(x_n | w_n = 0) = \text{Bern}_{x_n}[\rho]$$

$$Pr(x_n | w_n = 1) = \text{Bern}_{x_n}[1 - \rho]$$

# Graph Cuts Overview

Graph cuts used to optimise this cost function:

$$\operatorname{argmin}_{w_1 \dots w_N} \left[ \sum_{n=1}^N U_n(w_n) + \sum_{(m,n) \in \mathcal{C}} P_{mn}(w_m, w_n) \right]$$

**Unary terms**

(compatability of data with label y)

**Pairwise terms**

(compatability of neighboring labels)

Approach:

Convert minimization into the form of a standard CS problem,

MAXIMUM FLOW or MINIMUM CUT ON A GRAPH

Polynomial-time methods for solving this problem are known

# General Pairwise costs

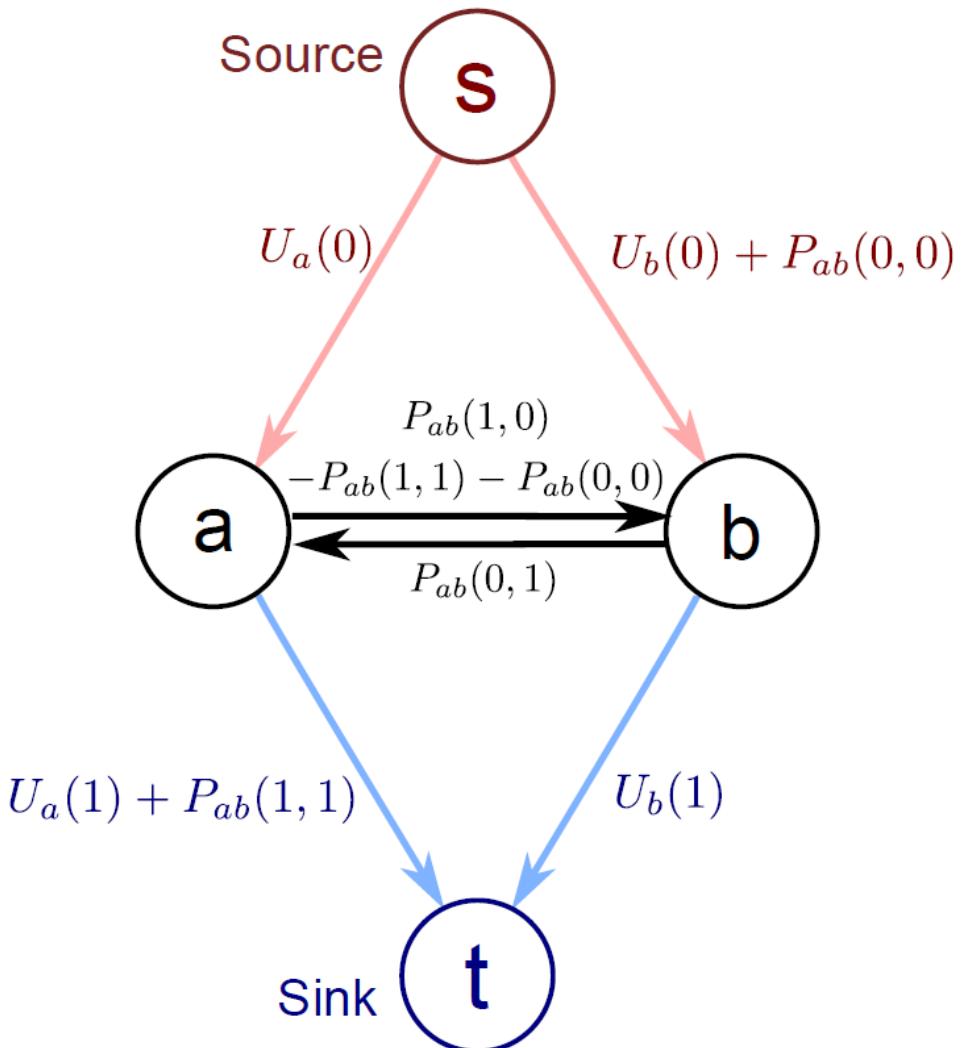
$$P_{m,n}(0,0) = \theta_{00} \quad P_{m,n}(1,0) = \theta_{10}$$

$$P_{m,n}(0,1) = \theta_{01} \quad P_{m,n}(1,1) = \theta_{11}$$

Modify graph to

- Add  $P(0,0)$  to edge s-b
  - Implies that solutions 0,0 and 1,0 also pay this cost
- Subtract  $P(0,0)$  from edge b-a
  - Solution 1,0 has this cost removed again

Similar approach for  $P(1,1)$



# Submodularity

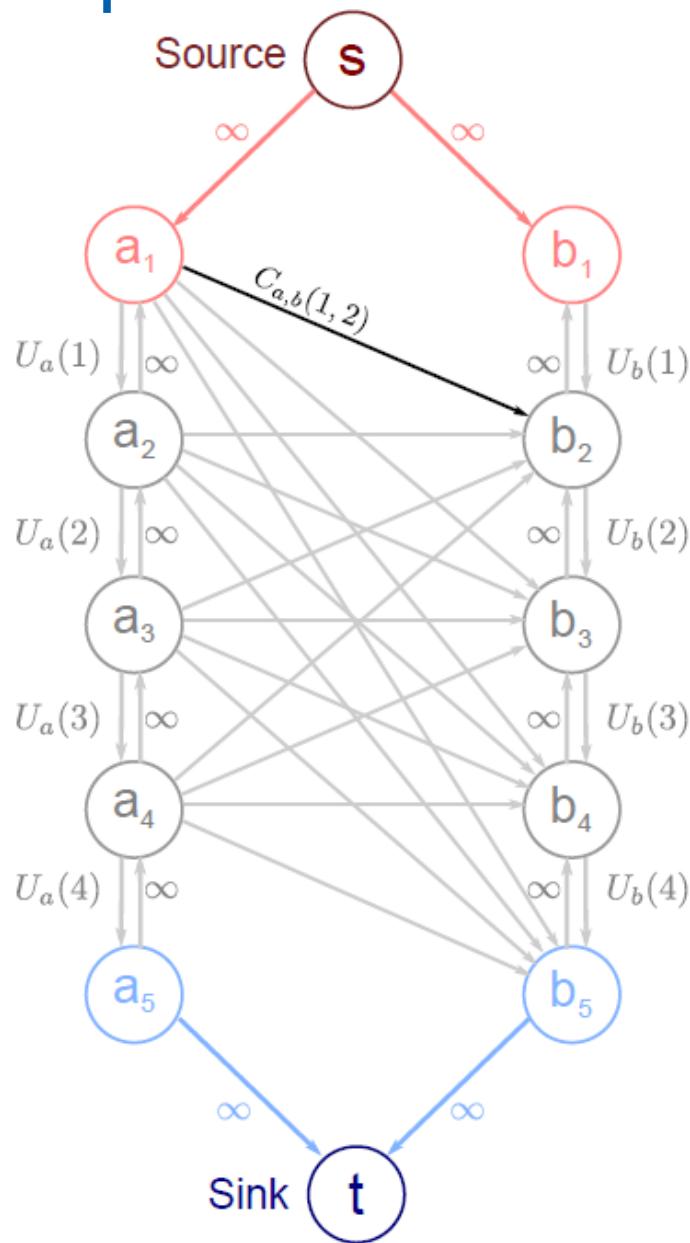
$$\theta_{10} + \theta_{01} - \theta_{11} - \theta_{00} \geq 0$$

If this condition is obeyed, it is said that the problem is “submodular” and it can be solved in polynomial time.

If it is not obeyed then the problem is NP hard.

Usually it is not a problem as we tend to favour smooth solutions.

# Multiple Labels



Inter-pixel edges have costs defined as:

$$C_{ab}(i, j) =$$

$$P_{ab}(i, j - 1) + P_{ab}(i - 1, j)$$

$$- P_{a,b}(i, j) - P_{ab}(i - 1, j - 1)$$

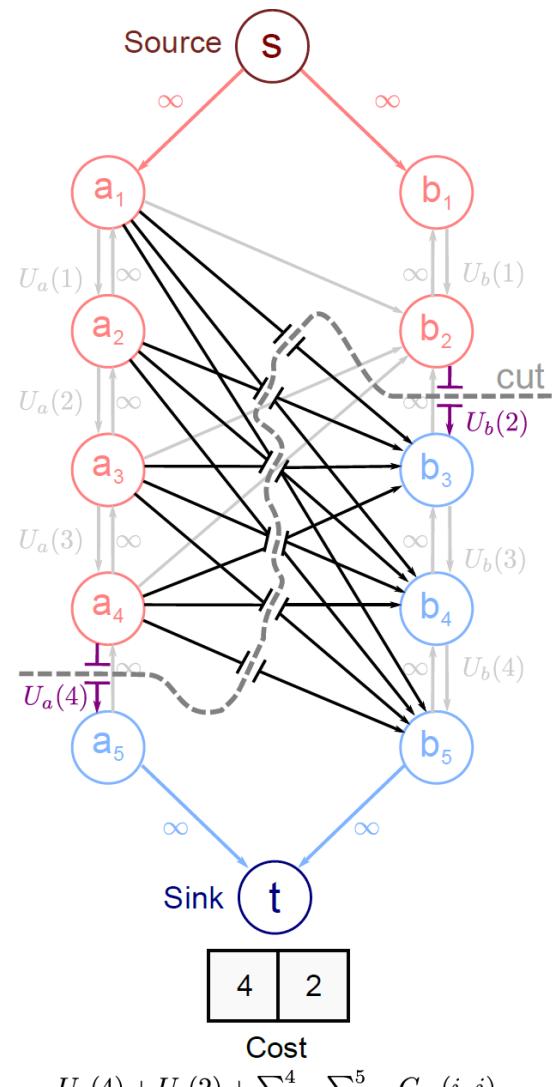
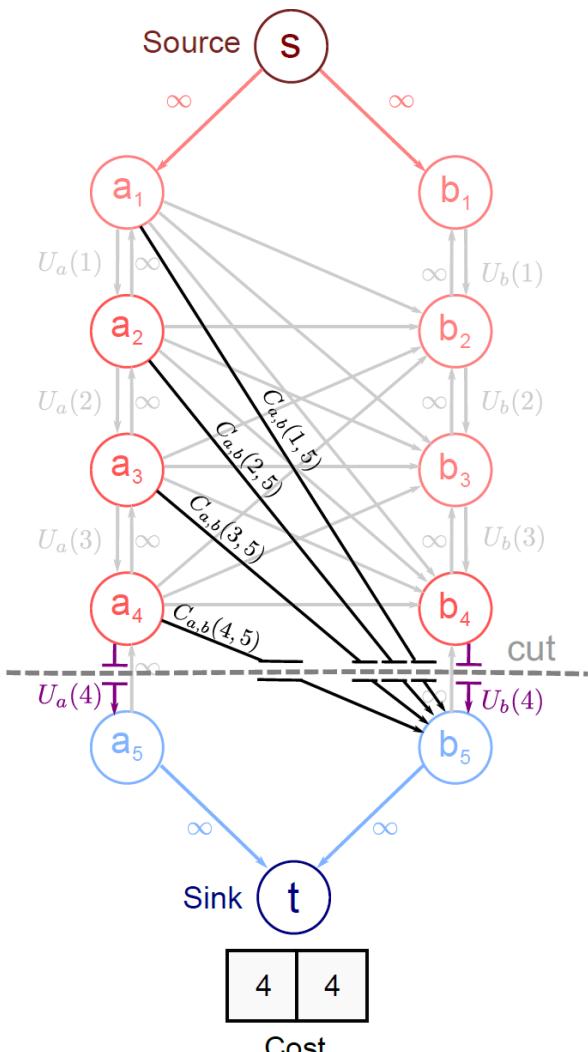
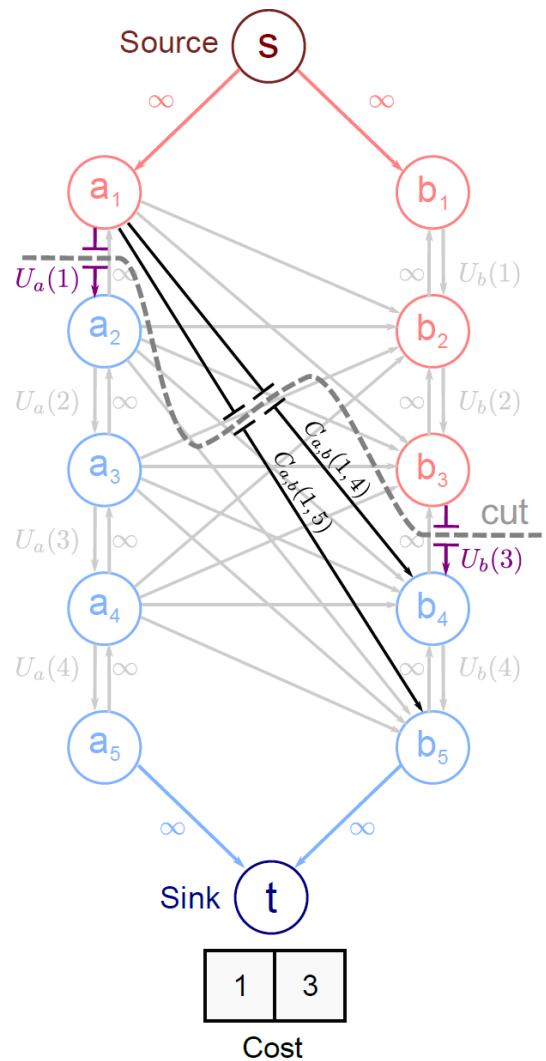
Superfluous terms :

$$P_{ab}(i, 0) = 0 \quad P_{ab}(i, K + 1) = 0$$

$$P_{ab}(0, j) = 0 \quad P_{ab}(K + 1, j) = 0$$

For all i,j where K is number of labels

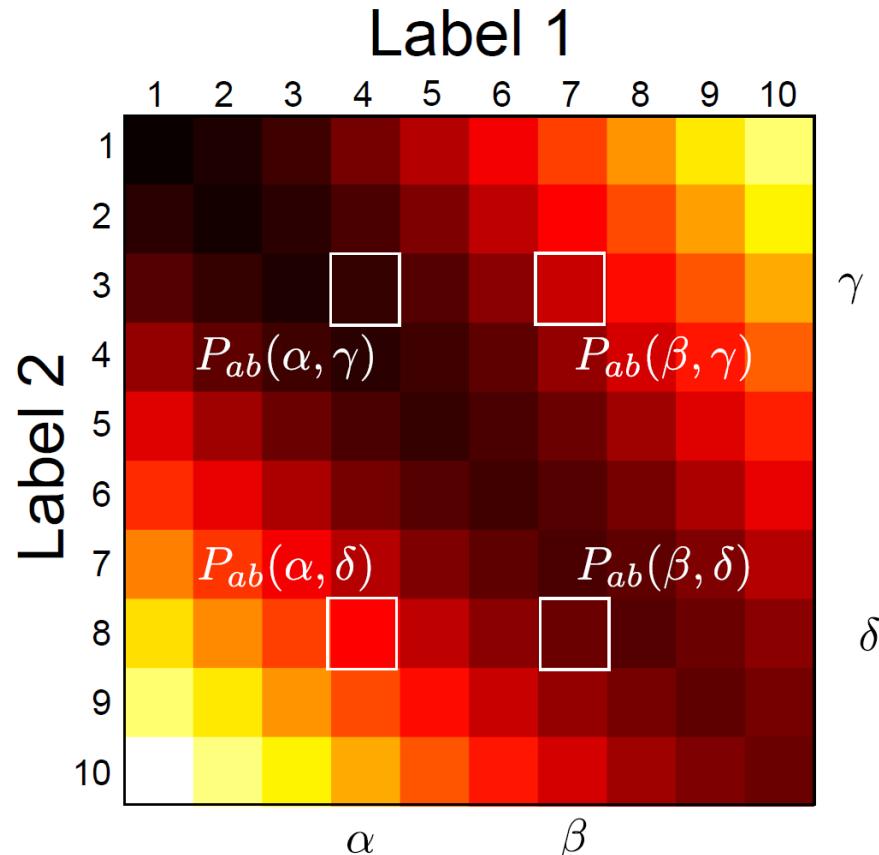
# Example Cuts



Must cut links from before cut on pixel a to after cut on pixel b.

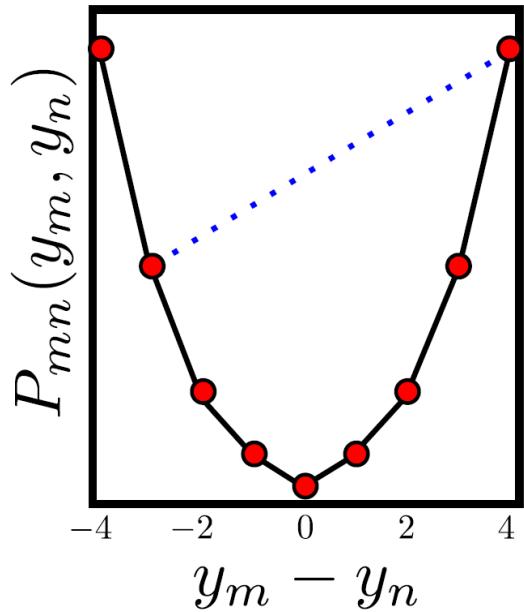
# Submodularity

$$P_{ab}(\beta, \gamma) + P_{ab}(\alpha, \delta) - P_{ab}(\beta, \delta) - P_{ab}(\alpha, \gamma) \geq 0 \quad \beta > \alpha \text{ and } \delta > \gamma$$



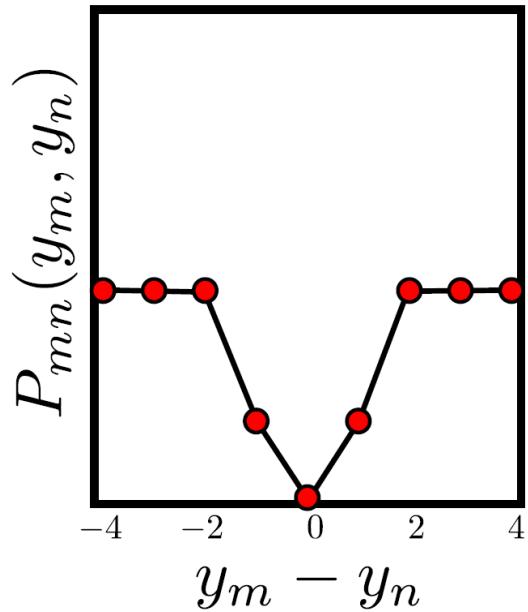
If not submodular, then the problem is NP hard.

# Convex vs. non-convex costs



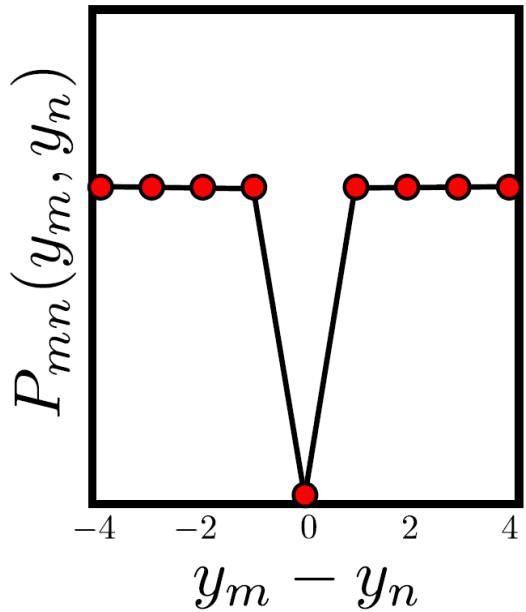
- Quadratic
- Convex
  - Submodular

$$\kappa(w_m - w_n)^2$$



- Truncated Quadratic
- Not Convex
  - Not Submodular

$$\min(\kappa_1, \kappa_2(w_m - w_n)^2)$$

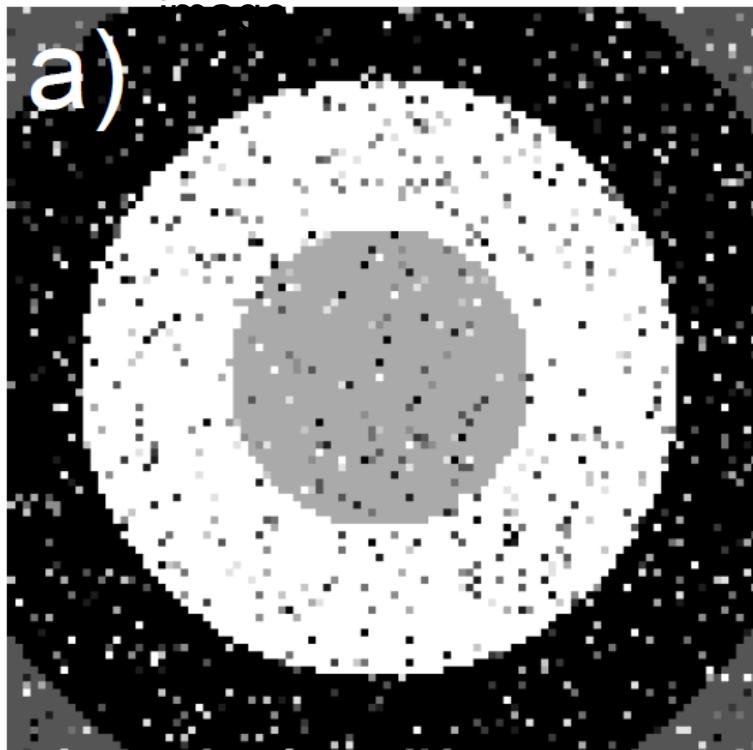


- Potts Model
- Not Convex
  - Not Submodular

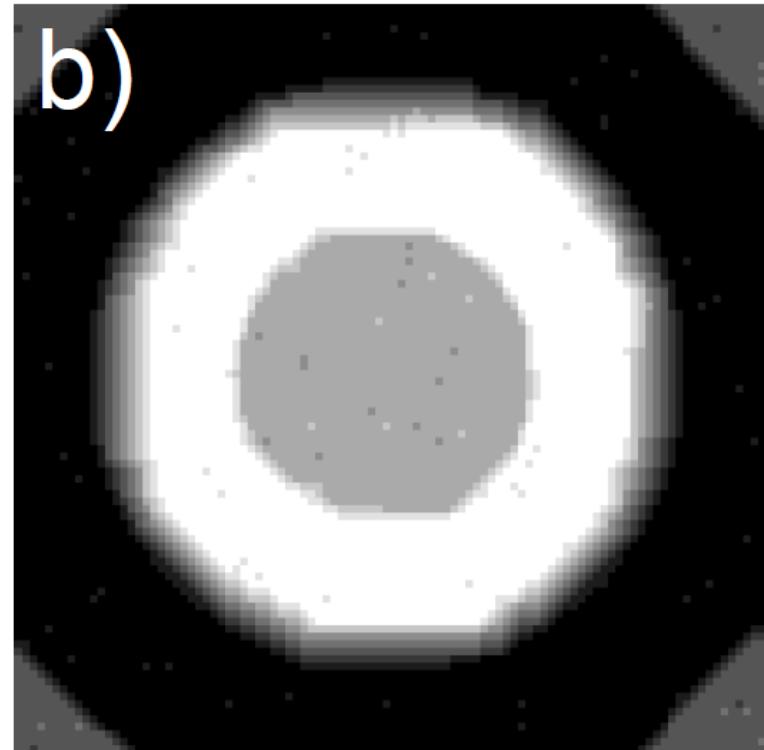
$$\kappa(1 - \delta(w_m - w_n))$$

# What is wrong with convex costs?

Observed noisy



Denoised result

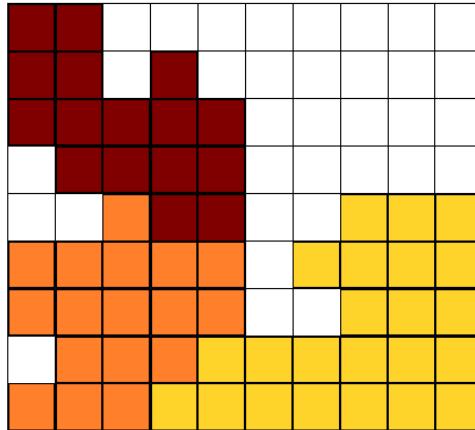


- Pay lower price for many small changes than one large one
- Result: blurring at large changes in intensity

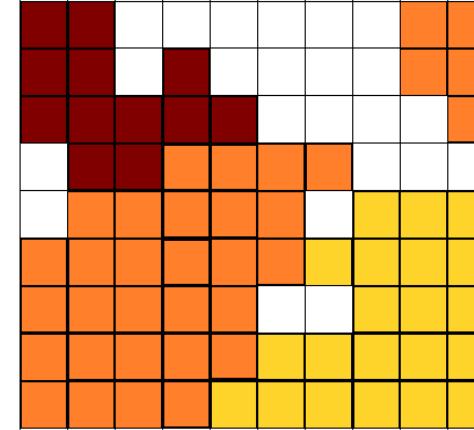
# Alpha Expansion Algorithm

- break multilabel problem into a series of binary problems
- at each iteration, pick label  $\alpha$  and expand (retain original or change to  $\alpha$ )

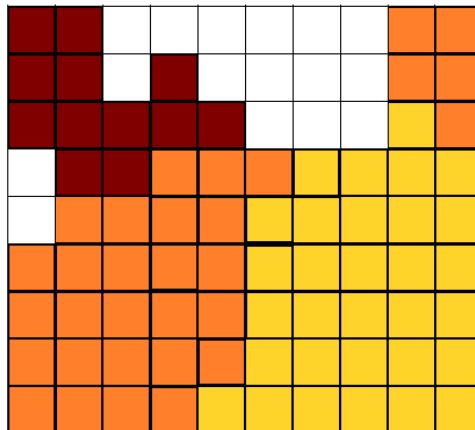
Initial labelling



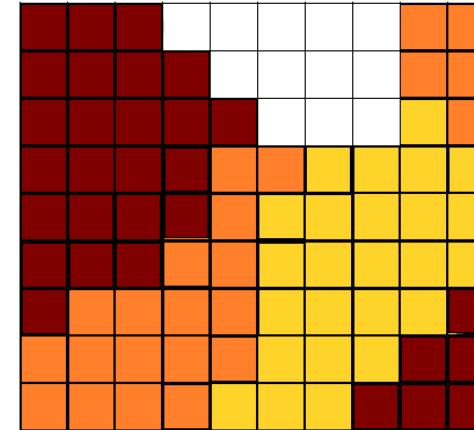
Iteration 1  
(orange)



Iteration 2  
(yellow)



Iteration 3  
(red)



# Alpha Expansion Ideas

- For every iteration
  - For every label
  - Expand label using optimal graph cut solution

Co-ordinate descent in label space.

Each step optimal, but overall global maximum not guaranteed

Proved to be within a factor of 2 of global optimum.

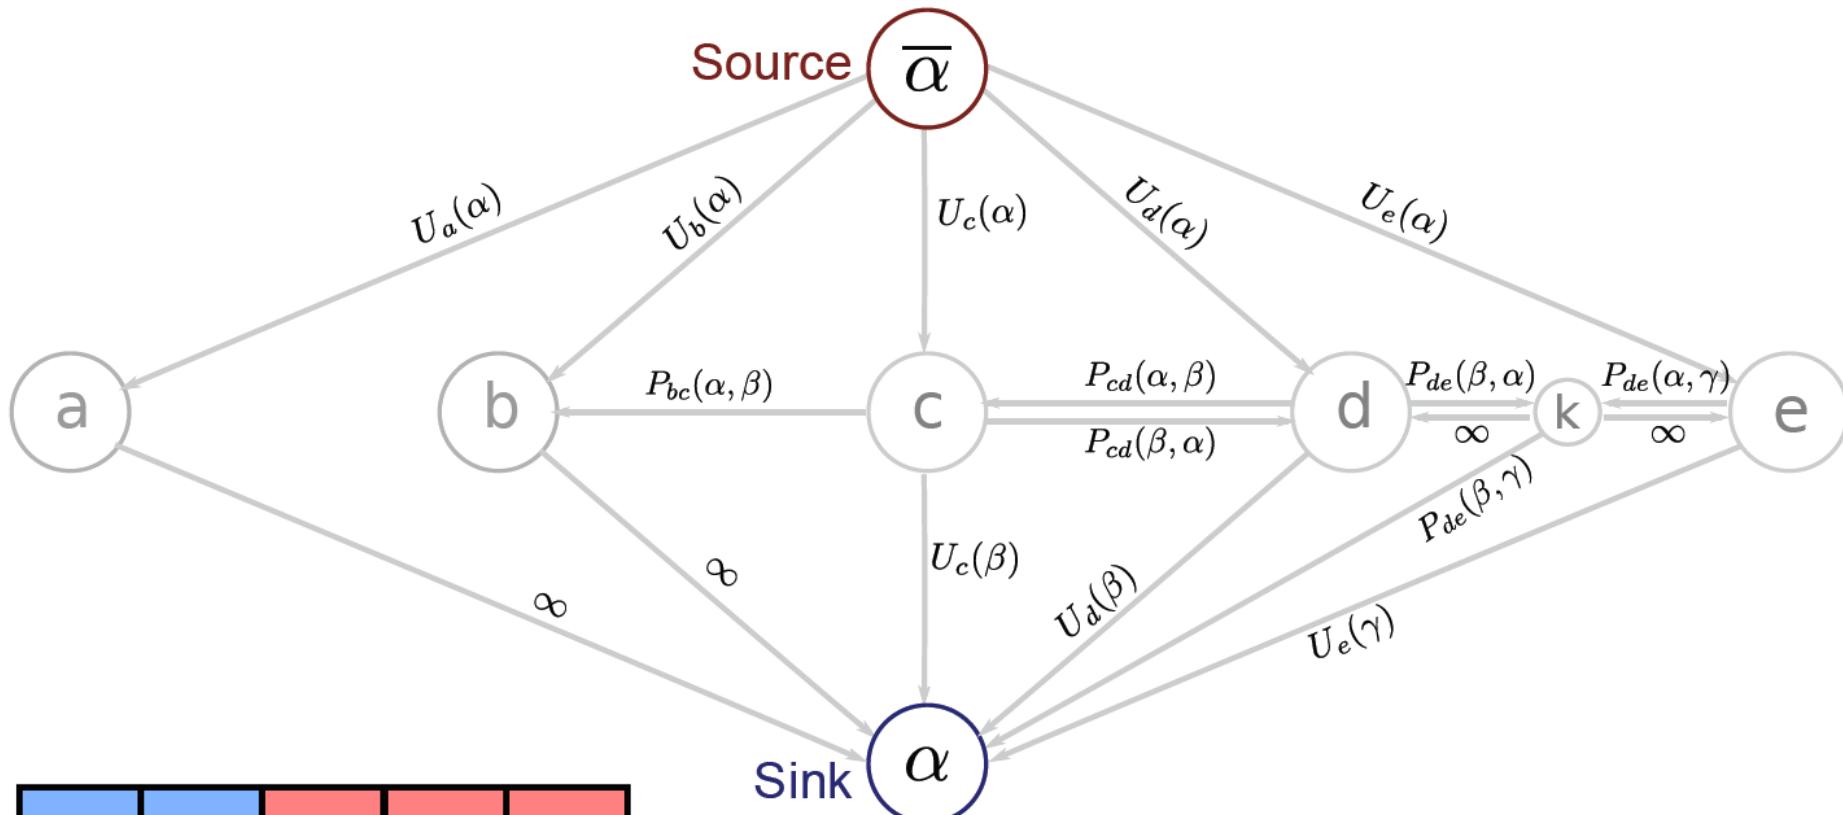
Requires that pairwise costs form a metric:

$$P(\alpha, \beta) = 0 \Leftrightarrow \alpha = \beta$$

$$P(\alpha, \beta) = P(\beta, \alpha) \geq 0$$

$$P(\alpha, \beta) \leq P(\alpha, \gamma) + P(\gamma, \beta)$$

# Alpha Expansion Construction



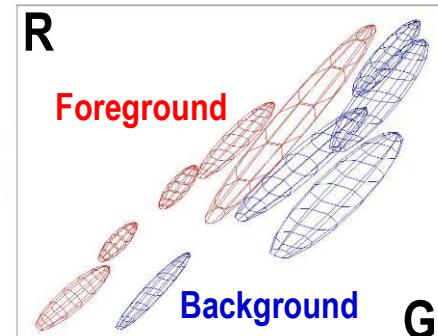
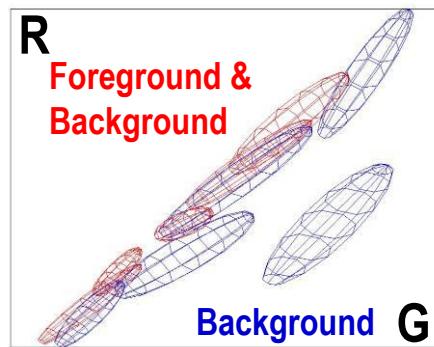
$\alpha$	$\alpha$	$\beta$	$\beta$	$\gamma$
States before				

Binary graph cut – either cut link to source (assigned to  $\alpha$ ) or to sink (retain current label)

Unary costs attached to links between source, sink and pixel nodes appropriately.

# Grab Cut

- Loosely specify foreground region
- Iterated graph cut



Gaussian Mixture Model (typically 5-8 components)

Source: K. Grauman

# Expectation Maximization

An algorithm specialized to fitting pdfs which are the marginalization of a joint distribution

$$\hat{\boldsymbol{\theta}} = \operatorname{argmax}_{\boldsymbol{\theta}} \left[ \sum_{i=1}^I \log \left[ \int Pr(\mathbf{x}_i, \mathbf{h}_i | \boldsymbol{\theta}) d\mathbf{h}_i \right] \right]$$

Defines a lower bound on log likelihood and increases bound iteratively

$$\begin{aligned} \mathcal{B} [\{q_i(\mathbf{h}_i)\}, \boldsymbol{\theta}] &= \sum_{i=1}^I \int q_i(\mathbf{h}_i) \log \left[ \frac{Pr(\mathbf{x}, \mathbf{h}_i | \boldsymbol{\theta})}{q_i(\mathbf{h}_i)} \right] d\mathbf{h}_{1\dots I} \\ &\leq \sum_{i=1}^I \log \left[ \int Pr(\mathbf{x}, \mathbf{h} | \boldsymbol{\theta}) d\mathbf{h}_{1\dots I} \right]. \end{aligned}$$

# E-Step & M-Step

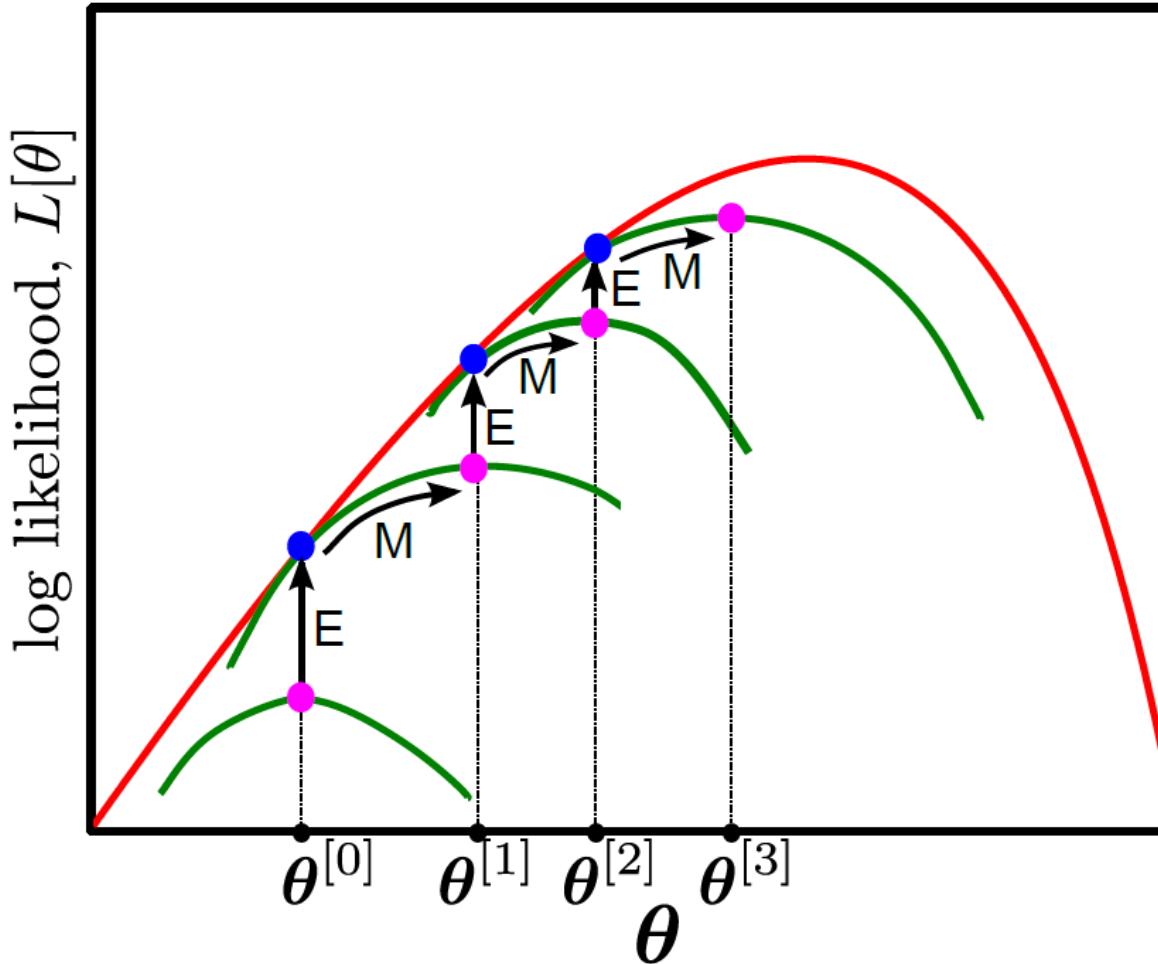
E-Step – Maximize bound w.r.t. distributions  $q_i(\mathbf{h}_i)$

$$\hat{q}_i(\mathbf{h}_i) = Pr(\mathbf{h}_i | \mathbf{x}_i, \boldsymbol{\theta}^{[t]}) = \frac{Pr(\mathbf{x}_i | \mathbf{h}_i, \boldsymbol{\theta}^{[t]}) Pr(\mathbf{h}_i | \boldsymbol{\theta}^{[t]})}{Pr(\mathbf{x}_i)}$$

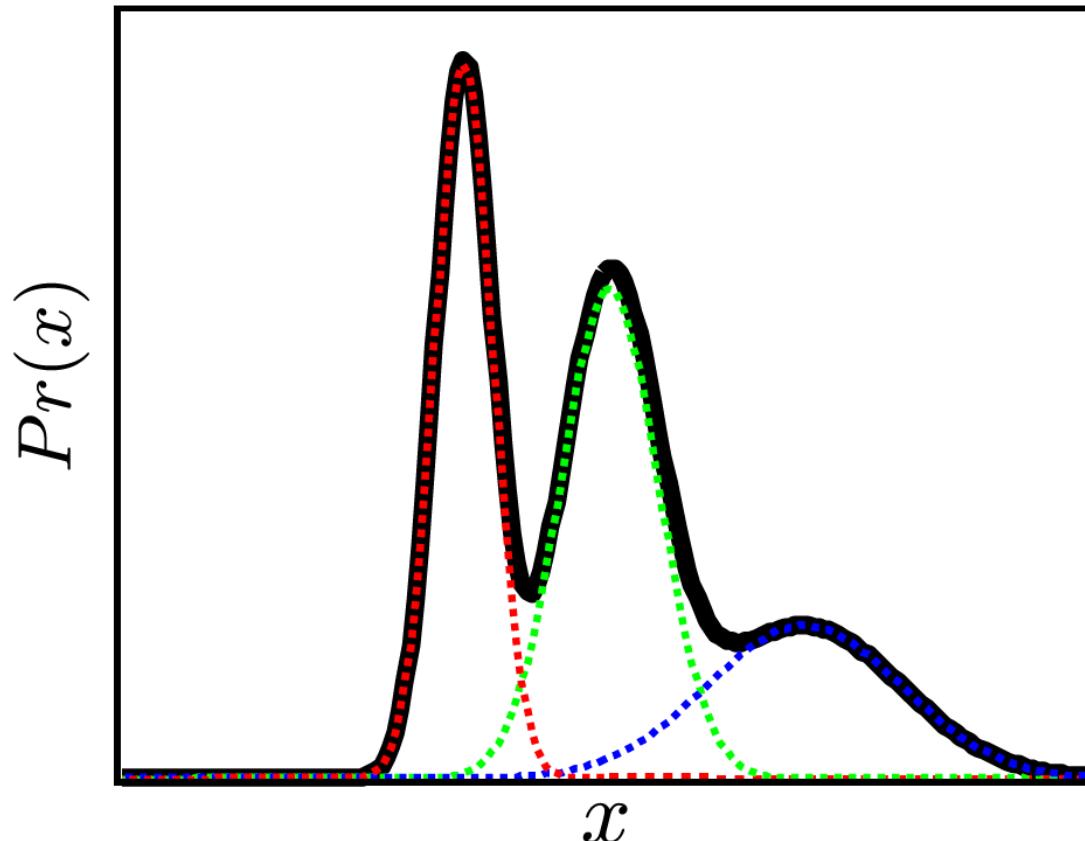
M-Step – Maximize bound w.r.t. parameters  $\boldsymbol{\theta}$

$$\hat{\boldsymbol{\theta}}^{[t+1]} = \operatorname{argmax}_{\boldsymbol{\theta}} \left[ \sum_{i=1}^I \int \hat{q}_i(\mathbf{h}_i) \log [Pr(\mathbf{x}_i, \mathbf{h}_i | \boldsymbol{\theta})] d\mathbf{h}_i \right]$$

# E-Step & M-Step



# Mixture of Gaussians (MoG)

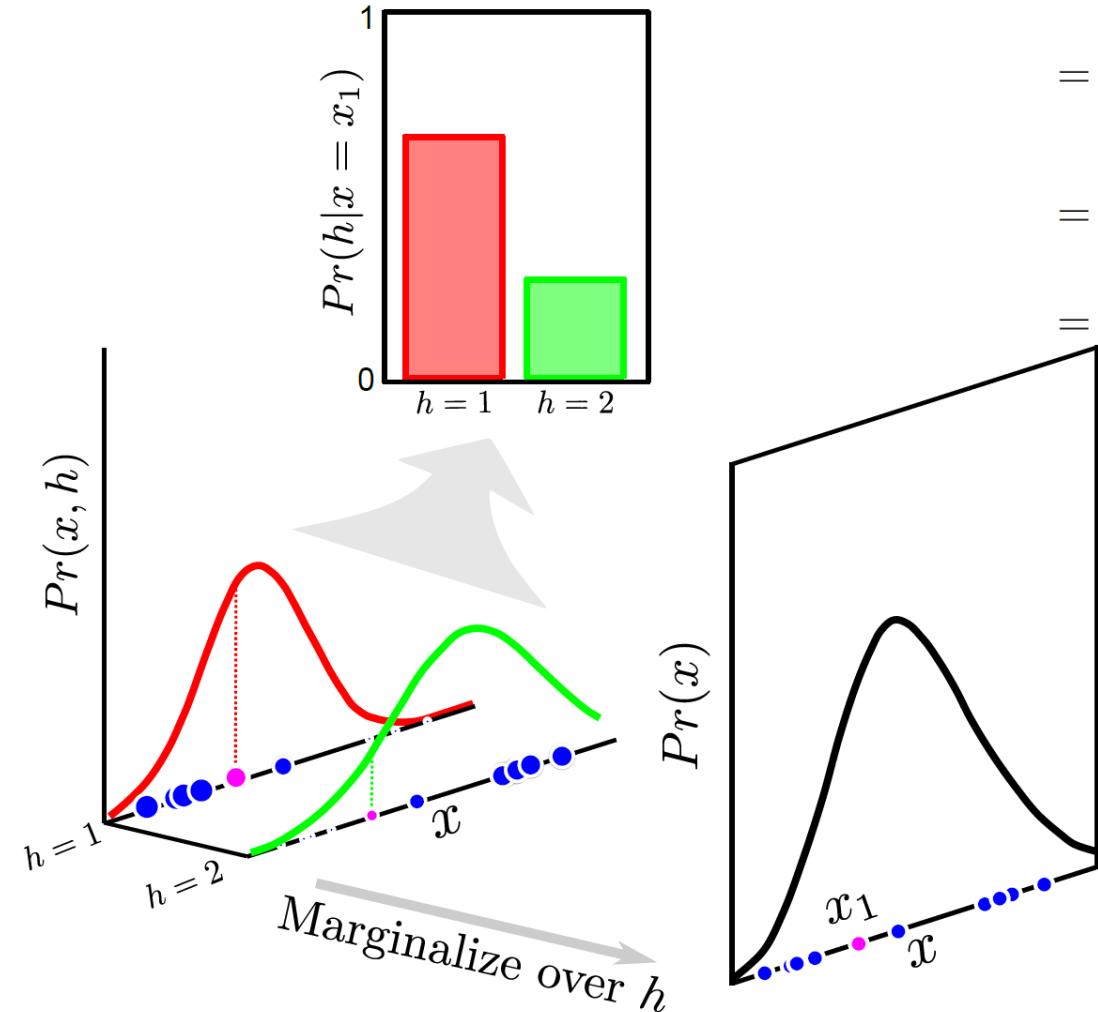


$$Pr(\mathbf{x}|\boldsymbol{\theta}) = \sum_{k=1}^K \lambda_k \text{Norm}_{\mathbf{x}}[\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k]$$

# E-Step

$$Pr(h_i = k | \mathbf{x}_i, \boldsymbol{\theta}^{[t]})$$

$$\begin{aligned}
&= \frac{Pr(\mathbf{x}_i | h_i = k, \boldsymbol{\theta}^{[t]}) Pr(h_i = k, \boldsymbol{\theta}^{[t]})}{\sum_{j=1}^K Pr(\mathbf{x}_i | h_i = j, \boldsymbol{\theta}^{[t]}) Pr(h_i = j, \boldsymbol{\theta}^{[t]})} \\
&= \frac{\lambda_k \text{Norm}_{\mathbf{x}_i} [\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k]}{\sum_{j=1}^K \lambda_j \text{Norm}_{\mathbf{x}_i} [\boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j]} \\
&= r_{ik}
\end{aligned}$$



We'll call this the responsibility of the  $k^{\text{th}}$  Gaussian for the  $i^{\text{th}}$  data point

Repeat this procedure for every datapoint!

# M-Step

$$\begin{aligned}
\hat{\boldsymbol{\theta}}^{[t+1]} &= \operatorname{argmax}_{\boldsymbol{\theta}} \left[ \sum_{i=1}^I \sum_{k=1}^K \hat{q}_i(h_i = k) \log [Pr(\mathbf{x}_i, h_i = k | \boldsymbol{\theta})] \right] \\
&= \operatorname{argmax}_{\boldsymbol{\theta}} \left[ \sum_{i=1}^I \sum_{k=1}^K r_{ik} \log [\lambda_k \operatorname{Norm}_{\mathbf{x}_i}[\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k]] \right].
\end{aligned}$$

Take derivative, equate to zero and solve (Lagrange multipliers for  $\lambda$ )

$$\begin{aligned}
\lambda_k^{[t+1]} &= \frac{\sum_{i=1}^I r_{ik}}{\sum_{j=1}^K \sum_{i=1}^I r_{ij}} \\
\boldsymbol{\mu}_k^{[t+1]} &= \frac{\sum_{i=1}^I r_{ik} \mathbf{x}_i}{\sum_{i=1}^I r_{ik}} \\
\boldsymbol{\Sigma}_k^{[t+1]} &= \frac{\sum_{i=1}^I r_{ik} (\mathbf{x}_i - \boldsymbol{\mu}_k^{[t+1]})(\mathbf{x}_i - \boldsymbol{\mu}_k^{[t+1]})^T}{\sum_{i=1}^I r_{ik}}
\end{aligned}$$

# Background subtraction

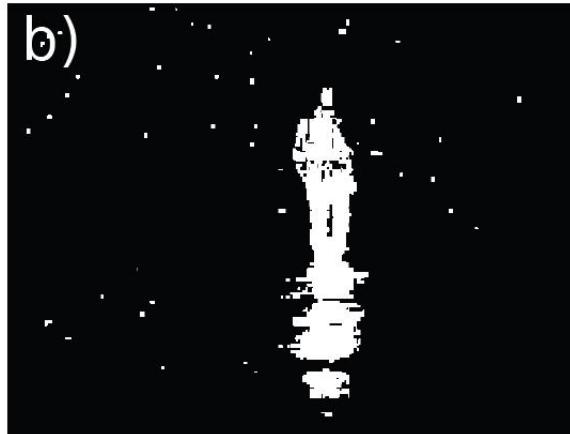


Model background distribution of each pixel by GMM:

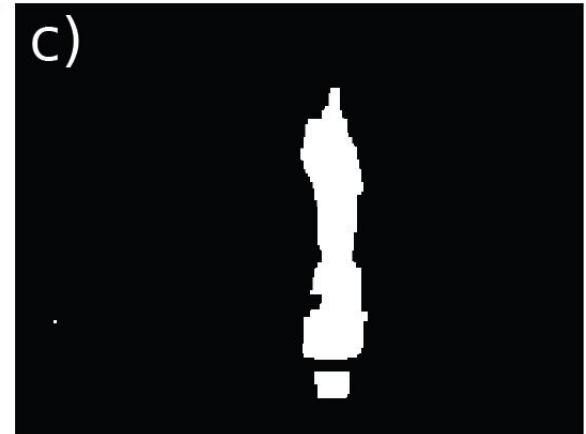
$$\Pr(x_n|w = 0) = \sum_i \lambda_{n,i} \text{Norm}_{x_n} [\mu_{n,i}, \Sigma_{n,i}]$$
$$\Pr(\mathbf{x}_n|w = 1) = \kappa,$$



Test image

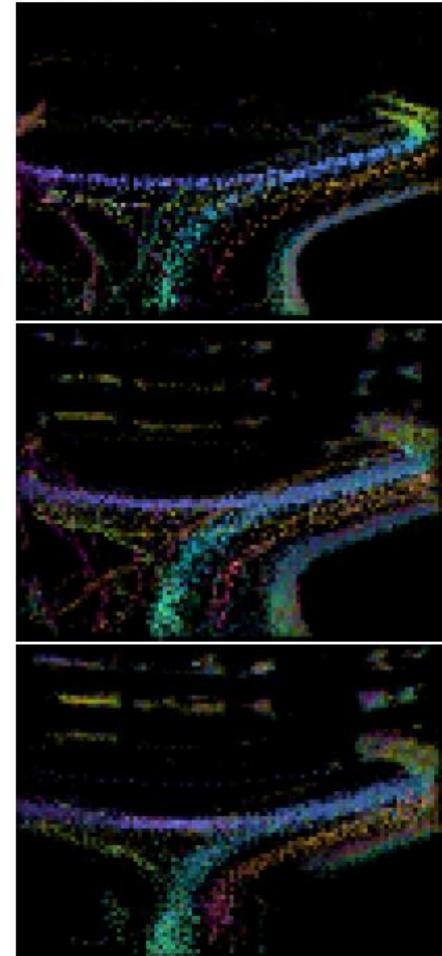
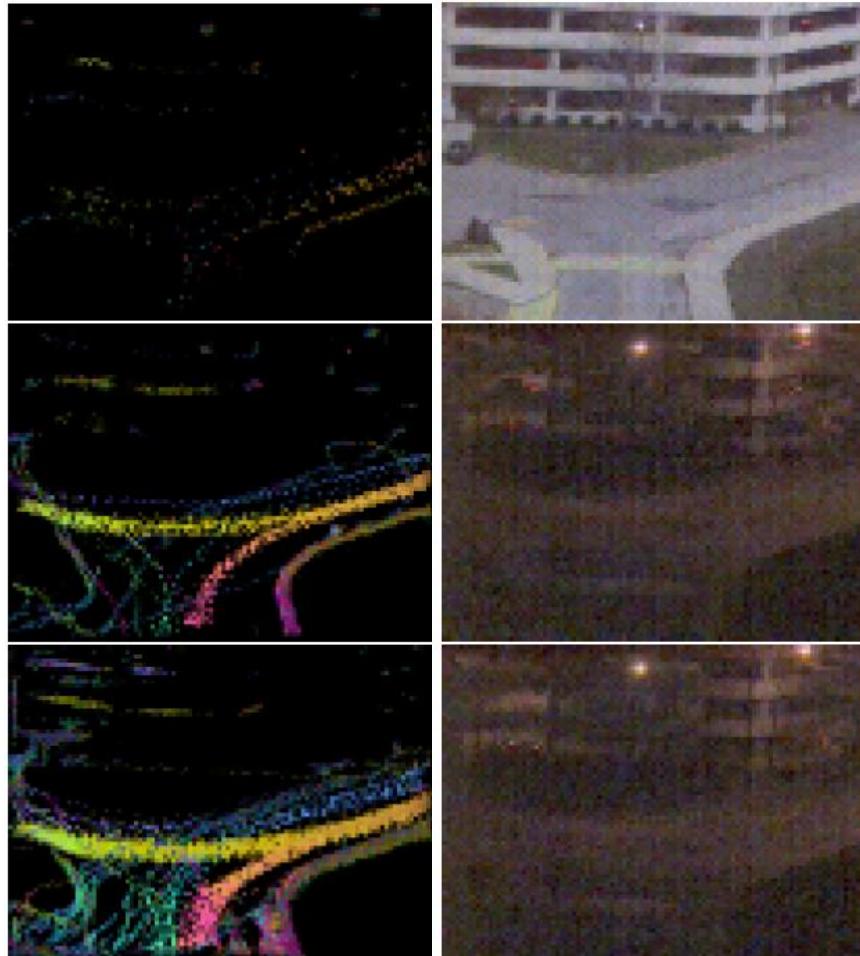


Pixel-wise classification



MRF

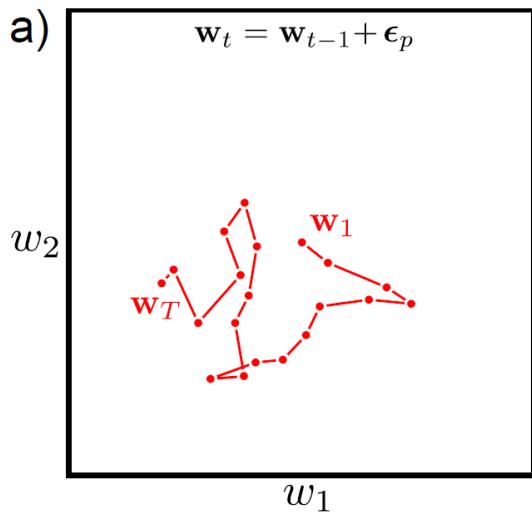
# Cars and pedestrians



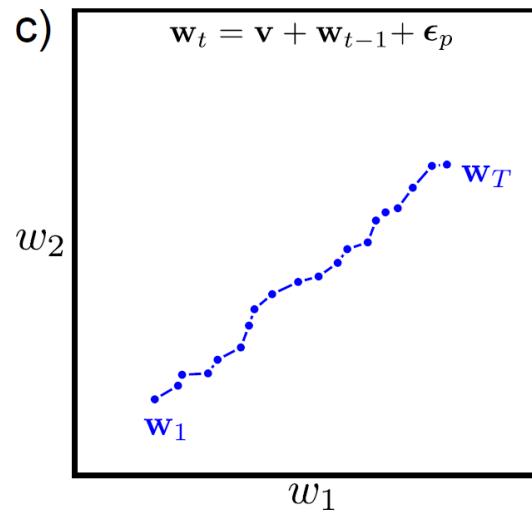
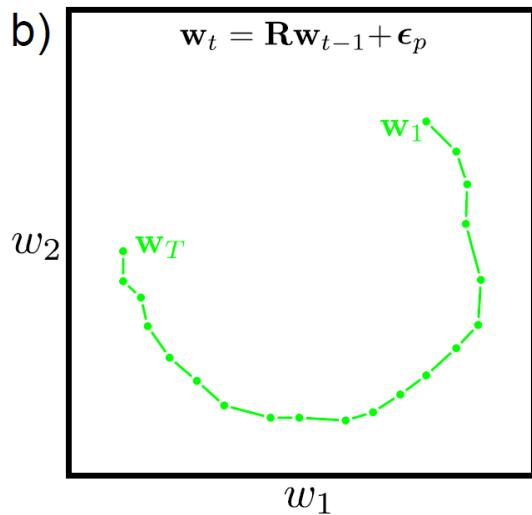
1 day

# Temporal Models

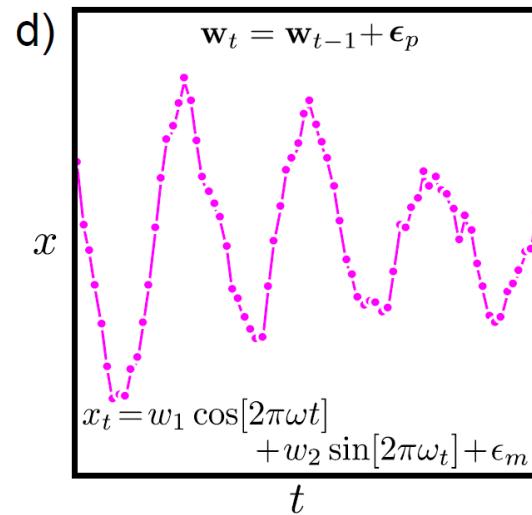
Brownian motion



Constant rotation



Constant velocity



Brownian motion with oscillatory measurements

# Assumptions

- Only the immediate past matters (Markov)
  - the probability of the state at time t is conditionally independent of states at times 1...t-2 given the state at time t-1, i.e.,  $P(w_t|w_{t-1}, \dots, w_1) = P(w_t|w_{t-1})$
- Measurements depend on only the current state
  - the likelihood of the measurements at time t is conditionally independent of all of the other measurements and the states at times 1...t-1 given the state at time t,  $P(x_t|w_t, \dots, w_1, x_{t-1}, \dots, x_1) = P(x_t|w_t)$

# Recursive Estimation

Time t

$$\begin{aligned} Pr(w_t|x_t, x_{1\dots t-1}) &= \frac{Pr(x_t|w_t, x_{1\dots t-1})Pr(w_t|x_{1\dots t-1})}{Pr(x_t|x_{1\dots t-1})} \\ &= \frac{Pr(x_t|w_t)Pr(w_t|x_{1\dots t-1})}{\int Pr(x_t|w_t, x_{1\dots t-1})Pr(w_t|x_{1\dots t-1})dw_t} \\ &= \frac{Pr(x_t|w_t)Pr(w_t|x_{1\dots t-1})}{\int Pr(x_t|w_t)Pr(w_t|x_{1\dots t-1})dw_t} \end{aligned}$$

Measurement model                          Prediction from temporal model

# Summary

Alternate between:

Temporal Evolution

Temporal model

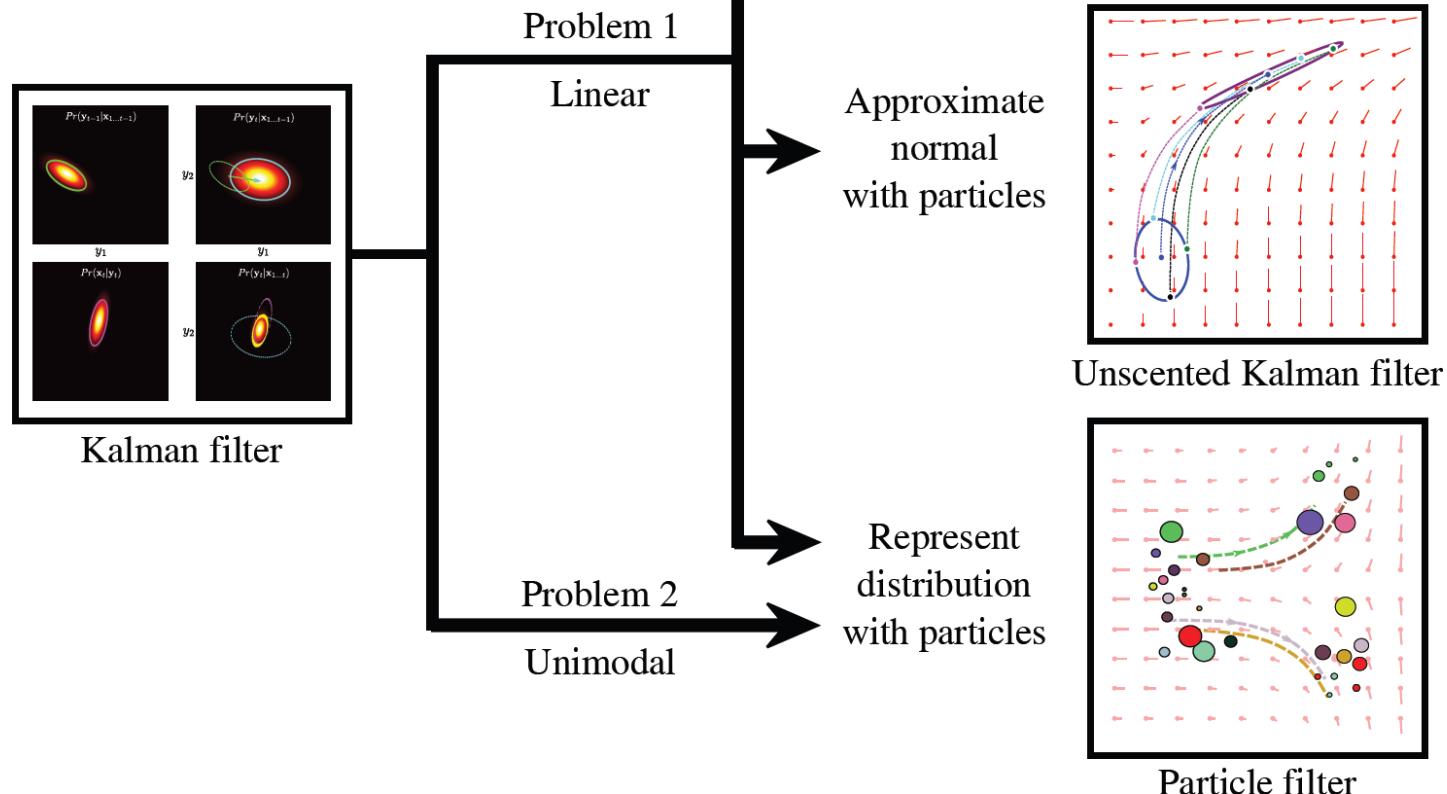
$$Pr(\mathbf{w}_t | \mathbf{x}_{1\dots t-1}) = \int Pr(\mathbf{w}_t | \mathbf{w}_{t-1}) Pr(\mathbf{w}_{t-1} | \mathbf{x}_{1\dots t-1}) d\mathbf{w}_{t-1}$$

Measurement Update

Measurement model

$$Pr(\mathbf{w}_t | \mathbf{x}_{1\dots t}) = \frac{Pr(\mathbf{x}_t | \mathbf{w}_t) Pr(\mathbf{w}_t | \mathbf{x}_{1\dots t-1})}{\int Pr(\mathbf{x}_t | \mathbf{w}_t) Pr(\mathbf{w}_t | \mathbf{x}_{1\dots t-1}) d\mathbf{w}_t}$$

# Roadmap



# Kalman Filter Definition

## Time evolution equation

$$Pr(\mathbf{w}_t | \mathbf{w}_{t-1}) = \text{Norm}_{\mathbf{w}_t} [\boldsymbol{\mu}_p + \boldsymbol{\Psi} \mathbf{w}_{t-1}, \boldsymbol{\Sigma}_p]$$

State transition matrix

Additive Gaussian noise

## Measurement equation

$$Pr(\mathbf{x}_t | \mathbf{w}_t) = \text{Norm}_{\mathbf{x}_t} [\boldsymbol{\mu}_m + \boldsymbol{\Phi} \mathbf{w}_t, \boldsymbol{\Sigma}_m]$$

Relates state and measurement

Additive Gaussian noise

# Kalman Filter Summary

## Time evolution equation

$$Pr(\mathbf{w}_t | \mathbf{w}_{t-1}) = \text{Norm}_{\mathbf{w}_t} [\boldsymbol{\mu}_p + \boldsymbol{\Psi} \mathbf{w}_{t-1}, \boldsymbol{\Sigma}_p]$$

## Measurement equation

$$Pr(\mathbf{x}_t | \mathbf{w}_t) = \text{Norm}_{\mathbf{x}_t} [\boldsymbol{\mu}_m + \boldsymbol{\Phi} \mathbf{w}_t, \boldsymbol{\Sigma}_m]$$

## Inference

State Prediction:

$$\boldsymbol{\mu}_+ = \boldsymbol{\mu}_p + \boldsymbol{\Psi} \boldsymbol{\mu}_{t-1}$$

Covariance Prediction:

$$\boldsymbol{\Sigma}_+ = \boldsymbol{\Sigma}_p + \boldsymbol{\Psi} \boldsymbol{\Sigma}_{t-1} \boldsymbol{\Psi}^T$$

State Update:

$$\boldsymbol{\mu}_t = \boldsymbol{\mu}_+ + \mathbf{K} (\mathbf{x}_t - \boldsymbol{\mu}_m - \boldsymbol{\Phi} \boldsymbol{\mu}_+)$$

Covariance Update:

$$\boldsymbol{\Sigma}_t = (\mathbf{I} - \mathbf{K} \boldsymbol{\Phi}) \boldsymbol{\Sigma}_+,$$

$$\mathbf{K} = \boldsymbol{\Sigma}_+ \boldsymbol{\Phi}^T (\boldsymbol{\Sigma}_m + \boldsymbol{\Phi} \boldsymbol{\Sigma}_+ \boldsymbol{\Phi}^T)^{-1}$$

# Fixed lag smoother

## State evolution equation

$$\begin{bmatrix} \mathbf{w}_t \\ \mathbf{w}_t^{[1]} \\ \mathbf{w}_t^{[2]} \\ \vdots \\ \mathbf{w}_t^{[\tau]} \end{bmatrix} = 
 \begin{bmatrix} \Psi & 0 & \dots & 0 & 0 \\ I & 0 & \dots & 0 & 0 \\ 0 & I & \dots & 0 & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & I & 0 \end{bmatrix} 
 \begin{bmatrix} \mathbf{w}_{t-1} \\ \mathbf{w}_{t-1}^{[1]} \\ \mathbf{w}_{t-1}^{[2]} \\ \vdots \\ \mathbf{w}_{t-1}^{[\tau]} \end{bmatrix} + 
 \begin{bmatrix} \epsilon_p \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

Estimate delayed by  $\tau$  

## Measurement equation

$$\mathbf{x}_t = [\Phi \quad 0 \quad 0 \quad \dots \quad 0] \begin{bmatrix} \mathbf{w}_t \\ \mathbf{w}_t^{[1]} \\ \mathbf{w}_t^{[2]} \\ \vdots \\ \mathbf{w}_t^{[\tau]} \end{bmatrix} + \epsilon_m$$

# Fixed interval smoothing

Having  $\mu_t$ ,  $\Sigma_t$ ,  $\mu_{+|t}$ , and  $\Sigma_{+|t}$  from forward pass where  
 $\Pr(w_{t+1}|x_1 \dots t) = \text{Norm}(\mu_{+|t}, \Sigma_{+|t})$

Backward set of recursions

$$\begin{aligned}\boldsymbol{\mu}_{t|T} &= \boldsymbol{\mu}_t + \mathbf{C}_t (\boldsymbol{\mu}_{t+1|T} - \boldsymbol{\mu}_{+|t}) \\ \boldsymbol{\Sigma}_{t|T} &= \boldsymbol{\Sigma}_t + \mathbf{C}_t (\boldsymbol{\Sigma}_{t+1|T} - \boldsymbol{\Sigma}_{+|t}) \mathbf{C}_t^T\end{aligned}$$

where

$$\mathbf{C}_t = \boldsymbol{\Sigma}_t \boldsymbol{\Psi}^T \boldsymbol{\Sigma}_{+|t}^{-1}$$

# Learning parameters

Parameters of equations:

$$w_t = \Psi w_{t-1} + \epsilon_p$$

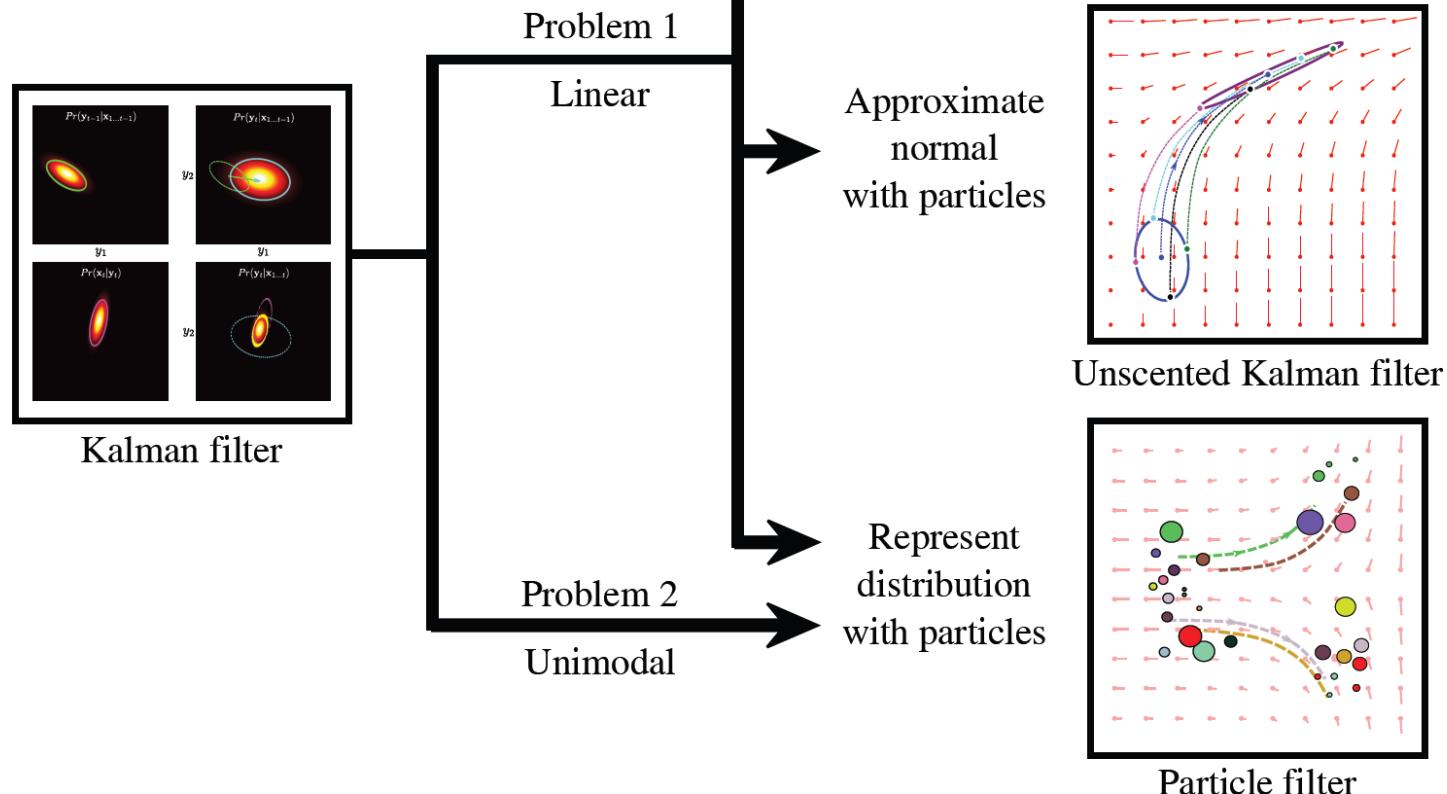
$$\theta = \{\Psi, \Sigma_p, \Phi, \Sigma_m, \mu_0, \Sigma_0\}$$

$$x_t = \Phi w_t + \epsilon_m$$

Expectation-Maximation

- E step:
  - Fixed interval smoothing to get ( $w_t$ )
- M step:
  - Maximize log-likelihood to get  $\theta$

# Roadmap



# Extended Kalman Filter

Allows non-linear measurement and temporal equations

$$\mathbf{w}_t = \mathbf{f}[\mathbf{w}_{t-1}, \boldsymbol{\epsilon}_p]$$

$$\mathbf{x}_t = \mathbf{g}[\mathbf{w}_t, \boldsymbol{\epsilon}_m]$$

Key idea: take Taylor expansion and treat as locally linear

# Unscented Kalman Filter

Approximate with particles where  $D_w$  is dimensionality of state:

$$\begin{aligned}
 Pr(\mathbf{w}_{t-1} | \mathbf{x}_{1\dots t-1}) &= \text{Norm}_{\mathbf{w}_{t-1}}[\boldsymbol{\mu}_{t-1}, \boldsymbol{\Sigma}_{t-1}] \\
 &\approx \sum_{j=0}^{2D_w} a_j \delta[\mathbf{w}_{t-1} - \hat{\mathbf{w}}^{[j]}],
 \end{aligned}$$

weights
particles or sigma points

Choose so that

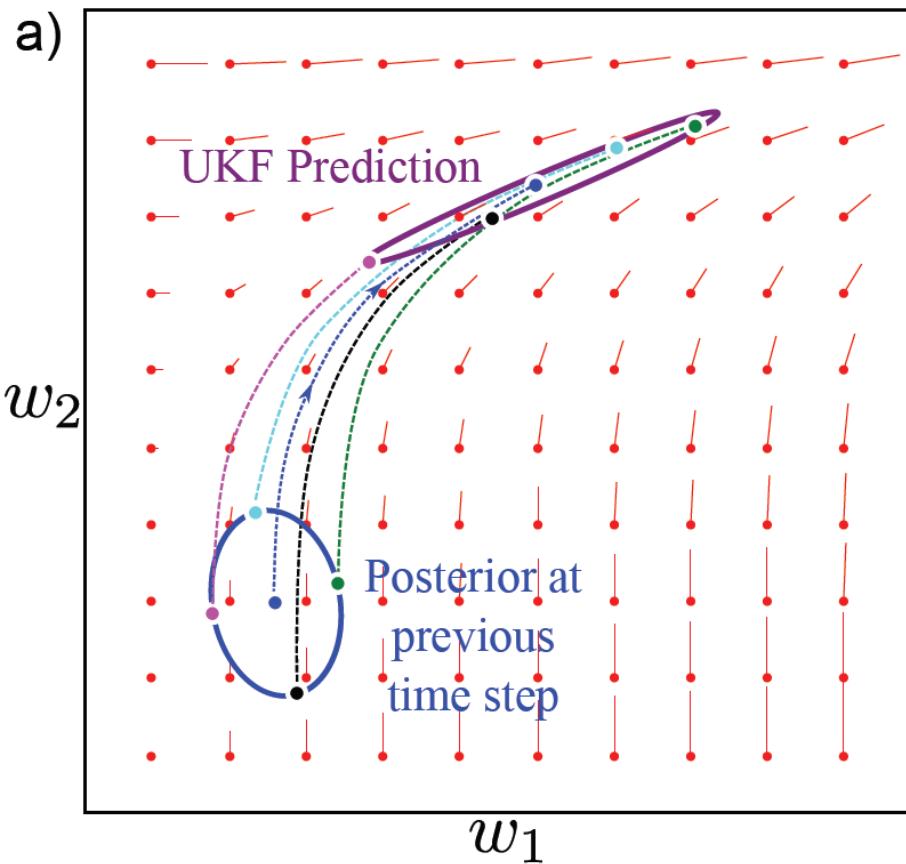
$$\boldsymbol{\mu}_{t-1} = \sum_{j=0}^{2D_w} a_j \hat{\mathbf{w}}^{[j]}$$

$$\boldsymbol{\Sigma}_{t-1} = \sum_{j=0}^{2D_w} a_j (\hat{\mathbf{w}}^{[j]} - \boldsymbol{\mu}_{t-1})(\hat{\mathbf{w}}^{[j]} - \boldsymbol{\mu}_{t-1})^T$$

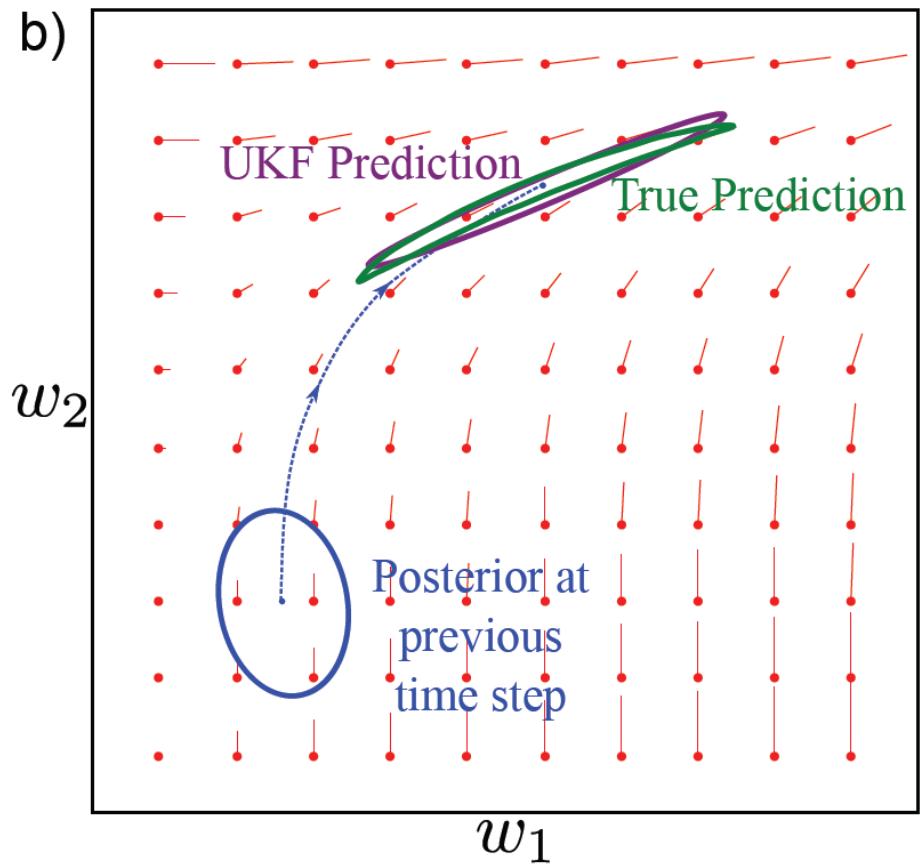
$$\sum_{j=0}^{2D_w} a_j = 1 \quad \delta(x) = \begin{cases} 1 & x = 0 \\ 0 & \text{otherwise} \end{cases}$$

# Unscented Kalman Filter

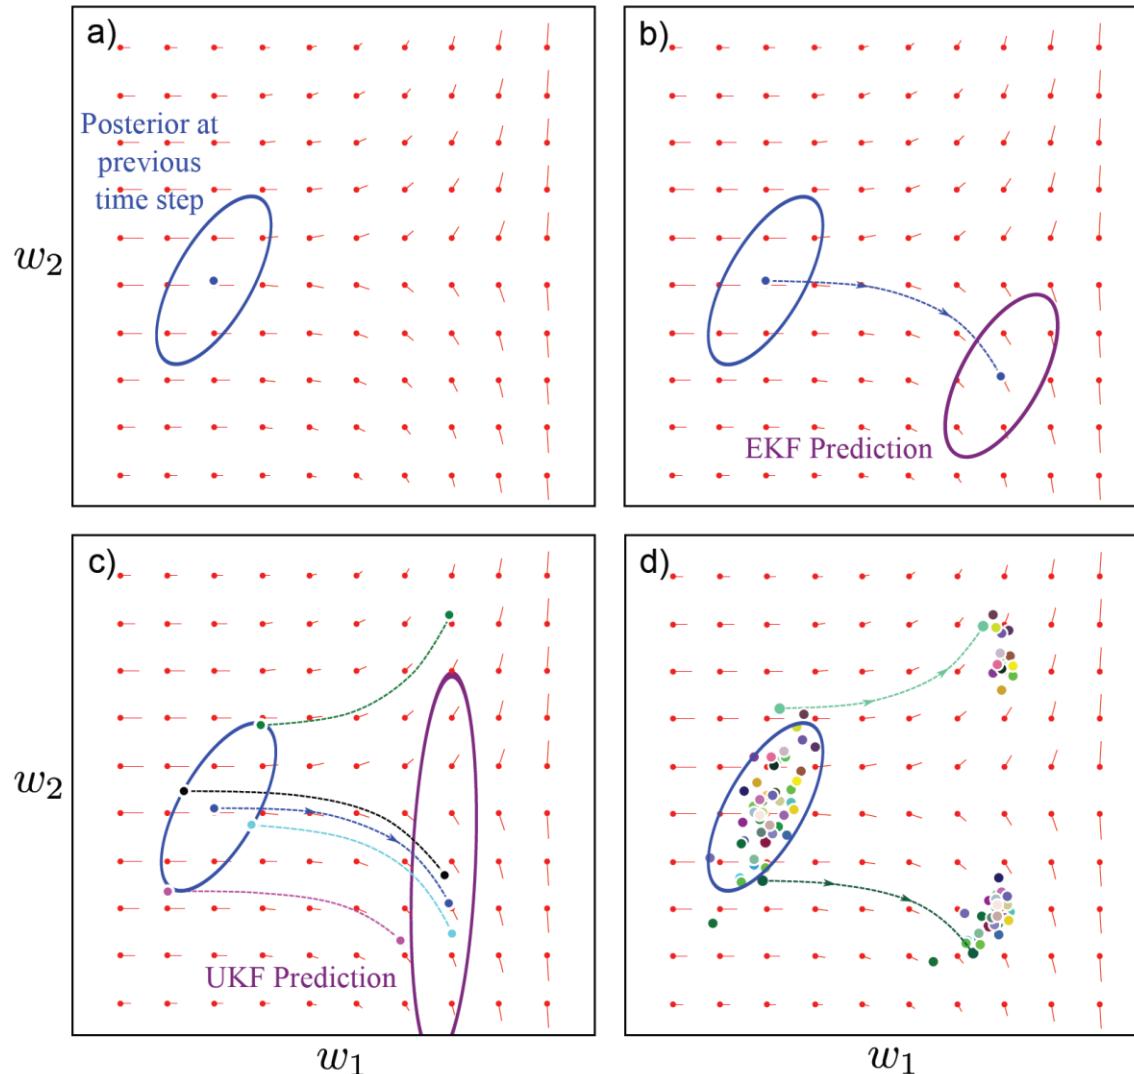
a)



b)



# Problems with UKF



# Particle filters

Key idea:

- Represent probability distribution as a set of weighted particles

$$Pr(\mathbf{w}_{t-1} | \mathbf{x}_{1\dots t-1}) = \sum_{j=1}^J a_j \delta[\mathbf{w}_{t-1} - \hat{\mathbf{w}}_{t-1}^{[j]}]$$

Advantages and disadvantages:

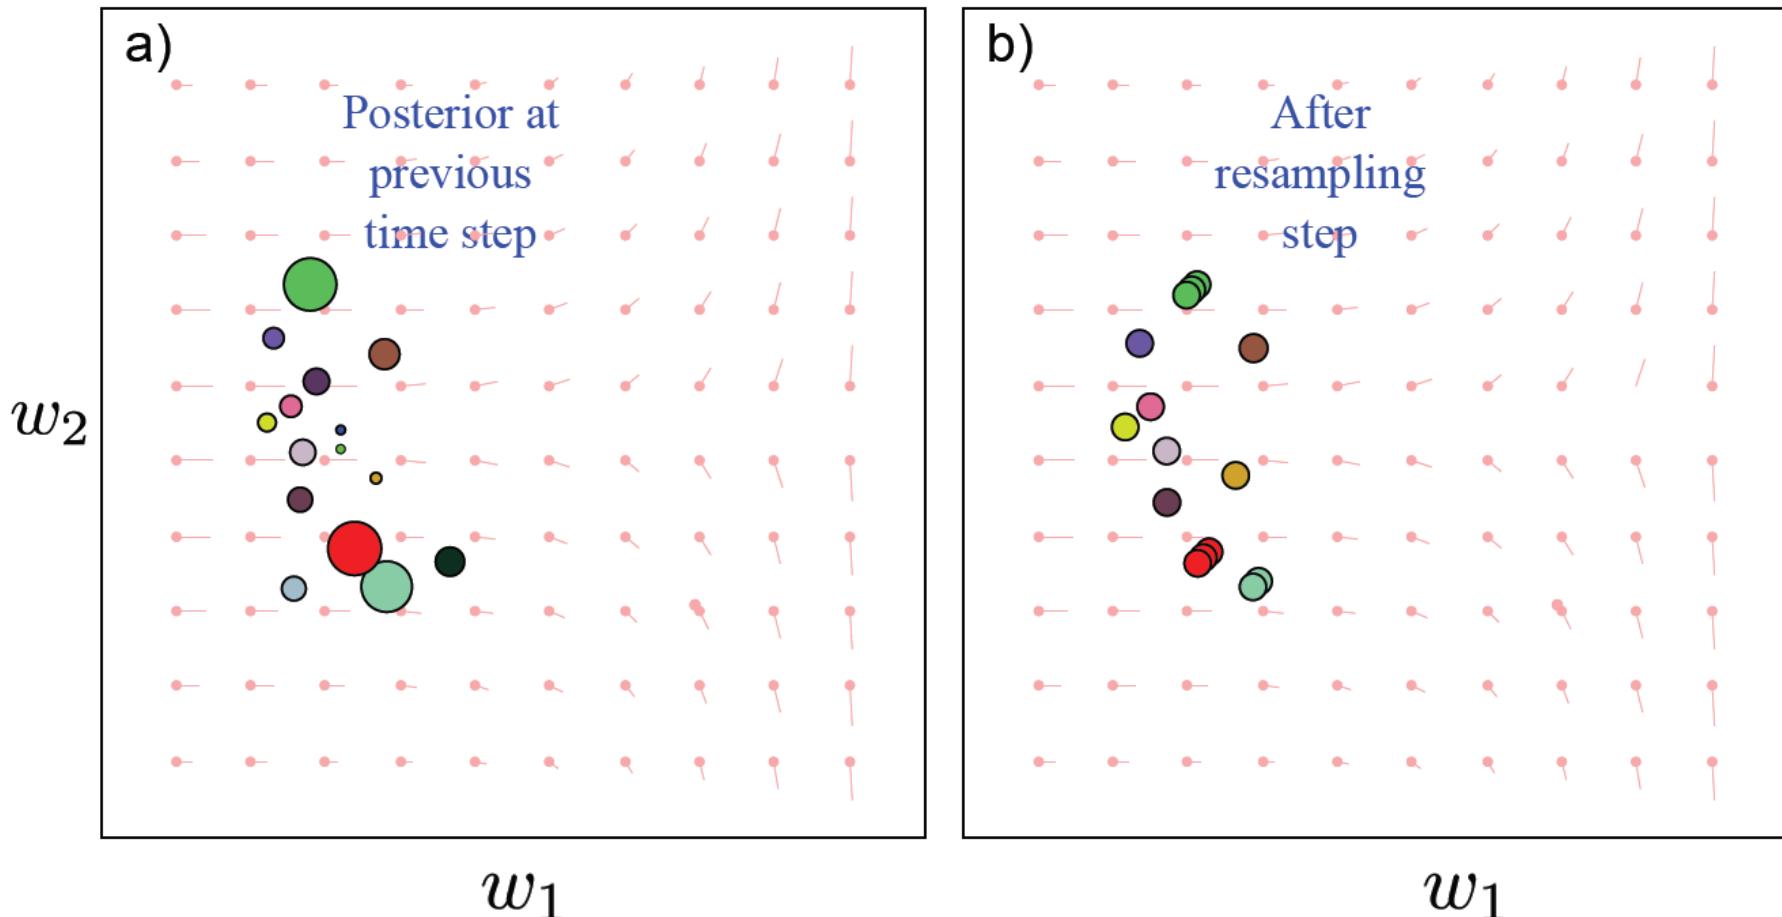
- + Can represent non-Gaussian multimodal densities
- + No need for data association
- Expensive

N. Gordon et al. **Novel approach to nonlinear/non-Gaussian Bayesian state estimation.** 1993  
M. Isard and A. Blake. **Contour tracking by stochastic propagation of conditional density.** ECCV 1996

Source: S. Prince

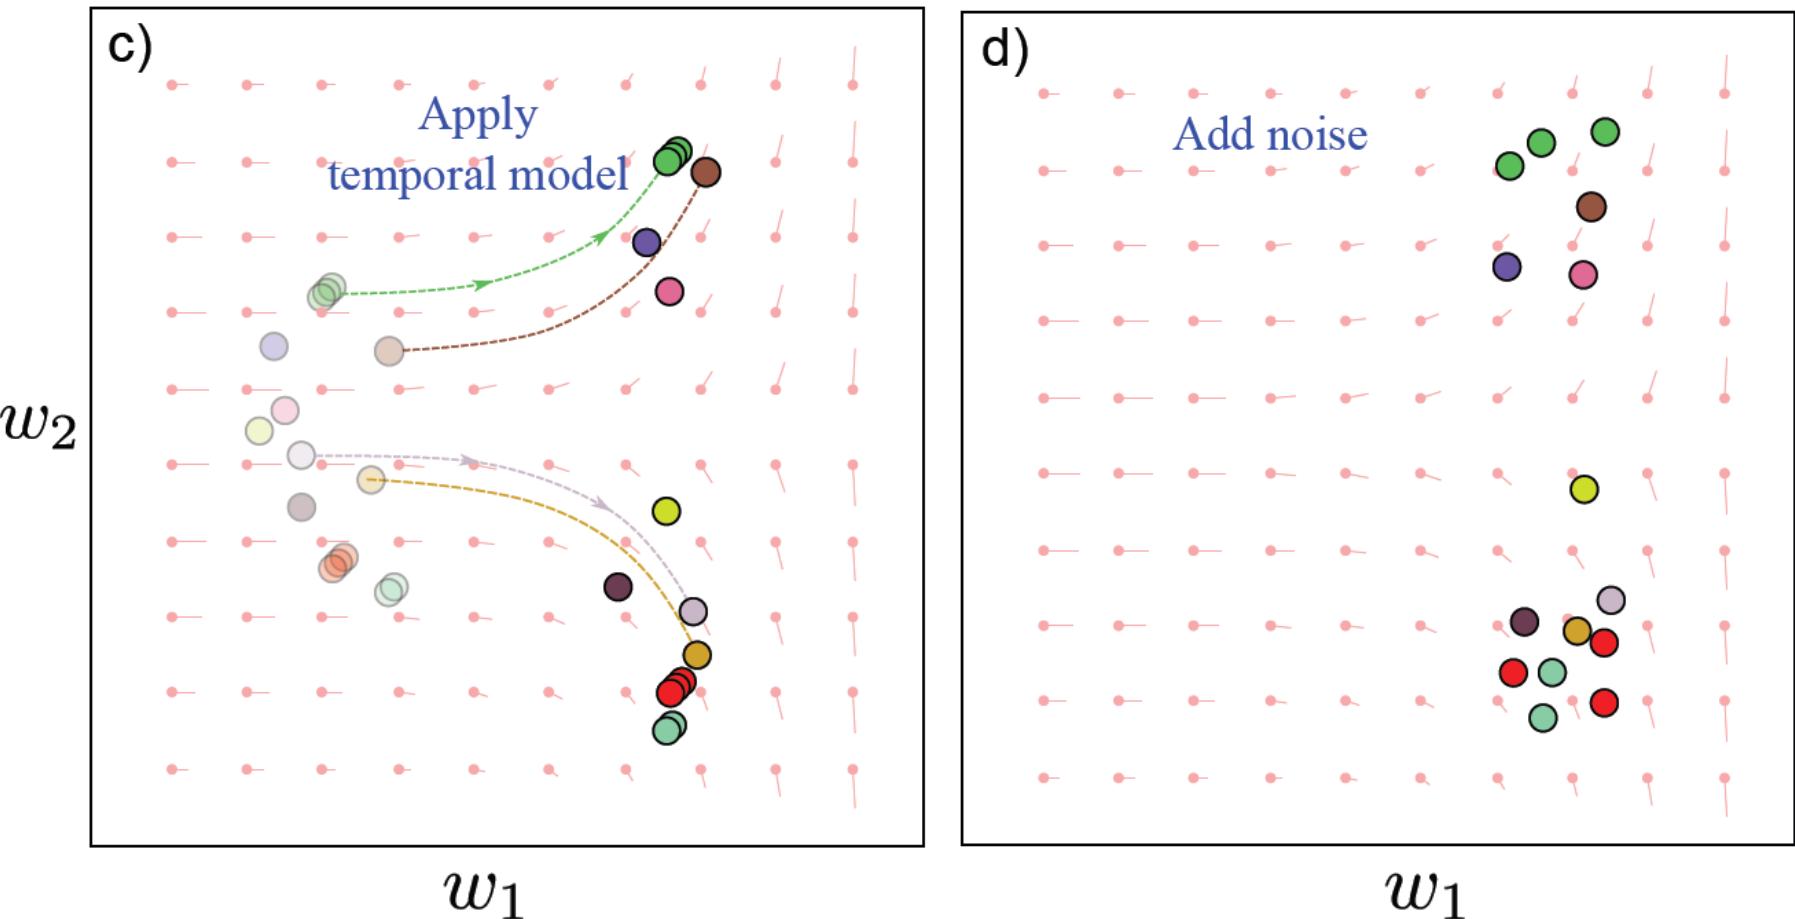
# Particle filter

Stage 1: Resample from weighted particles according to their weight to get unweighted particles



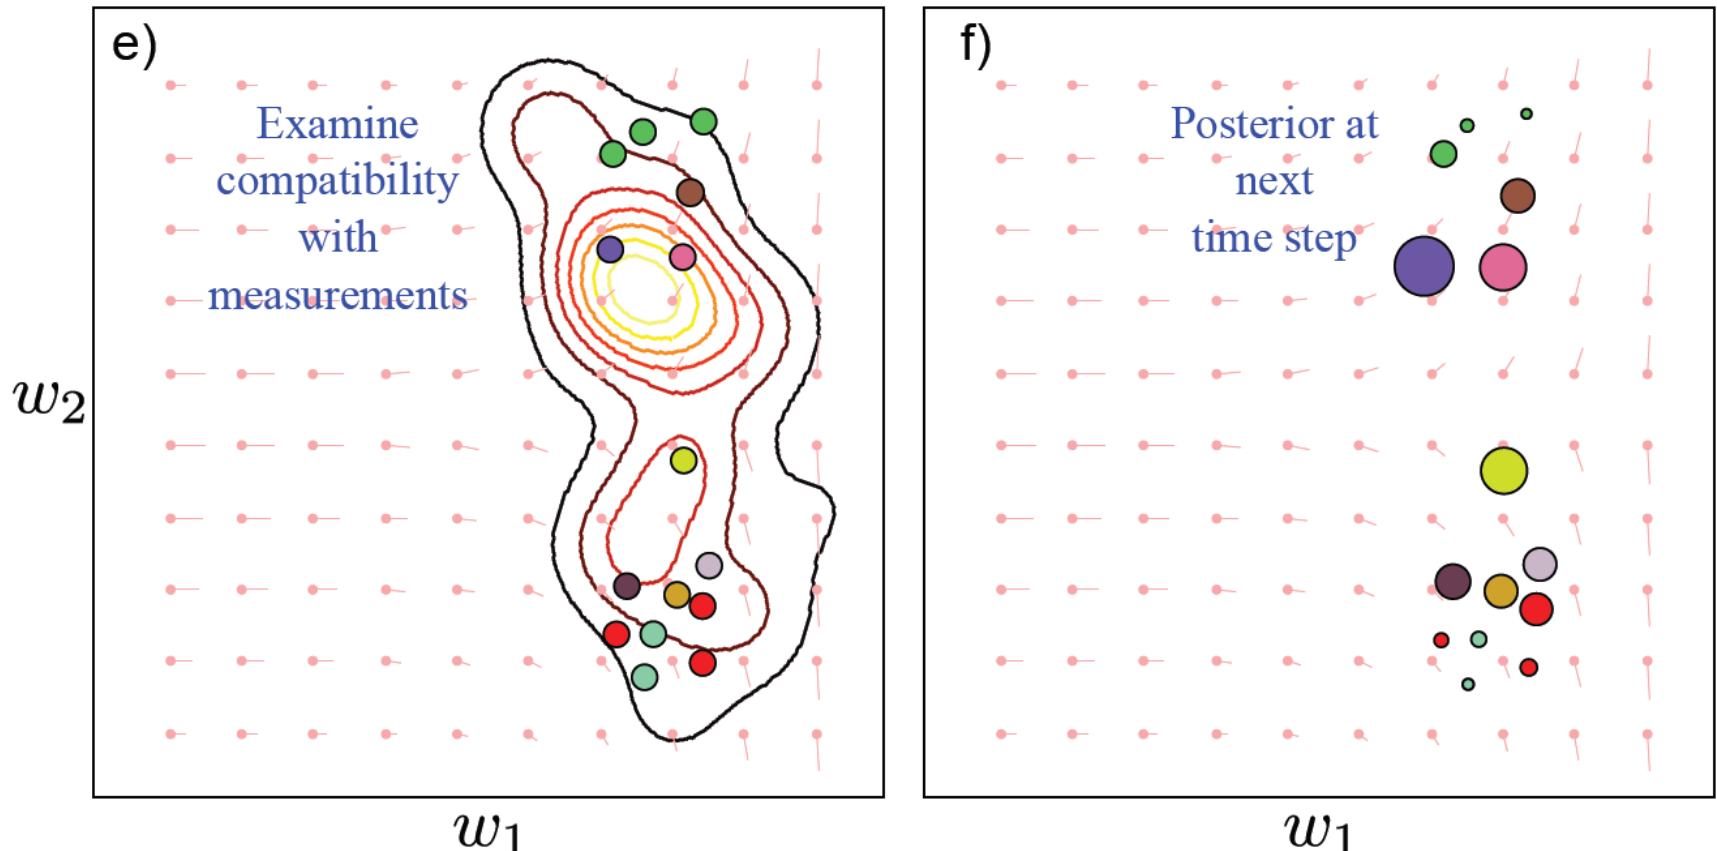
# Particle filter

Stage 2: Pass unweighted samples through temporal model and add noise



# Particle filter

## Stage 3: Weight samples by measurement density



# Generic particle filter

## 1. Initialisation

- $t \leftarrow 0$
- For  $i = 1, \dots, n$ , sample  $x_{0,0}^{(i)}$  from  $\eta_0$

## 2. Prediction

Temporal model

- For  $i = 1, \dots, n$ , sample  $\bar{x}_{t+1,0}^{(i)}$  from  $K_t(x_{t,0}^{(i)}, \cdot)$

## 3. Updating

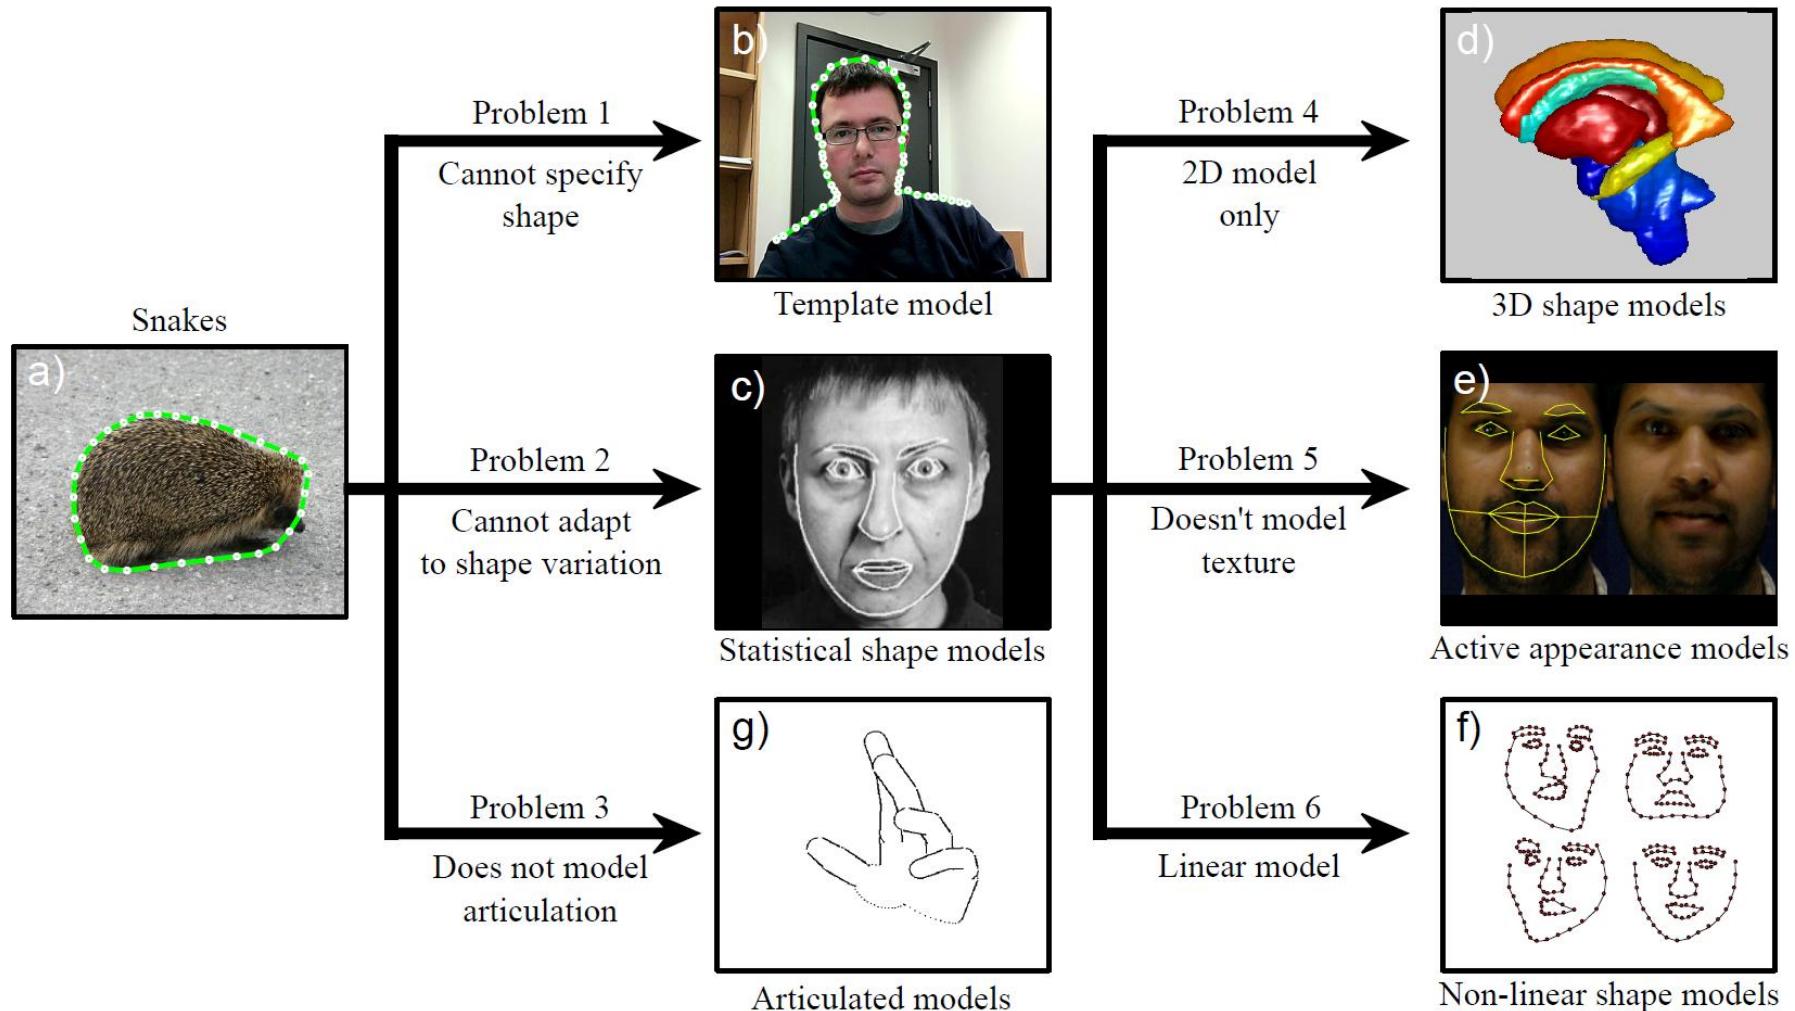
Likelihood

- For  $i = 1, \dots, n$ , set  $\pi_{t+1,0}^{(i)} \leftarrow g_{t+1}(y_{t+1} - h_{t+1}(\bar{x}_{t+1,0}^{(i)}))$
- For  $i = 1, \dots, n$ , set  $\pi_{t+1,0}^{(i)} \leftarrow \frac{\pi_{t+1,0}^{(i)}}{\sum_{j=1}^n \pi_{t+1,0}^{(j)}}$

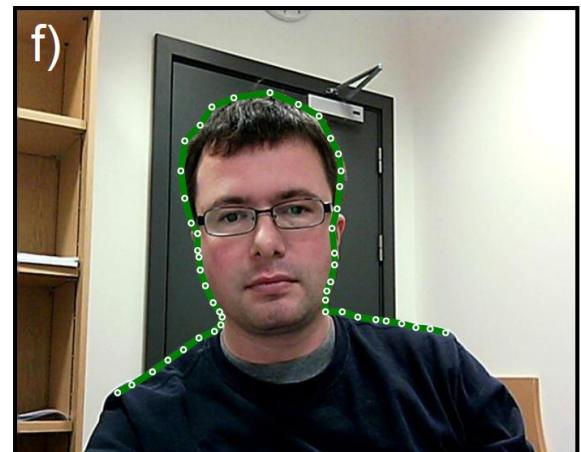
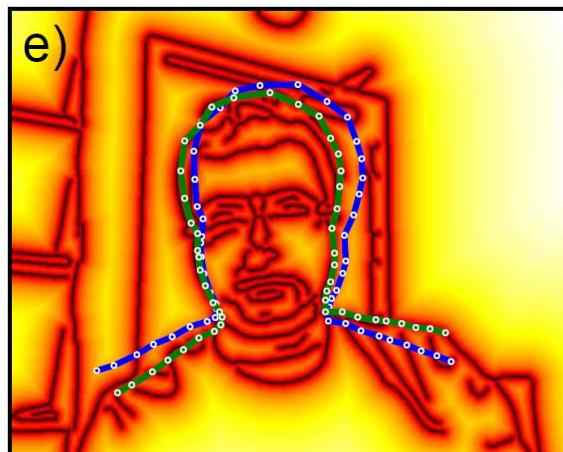
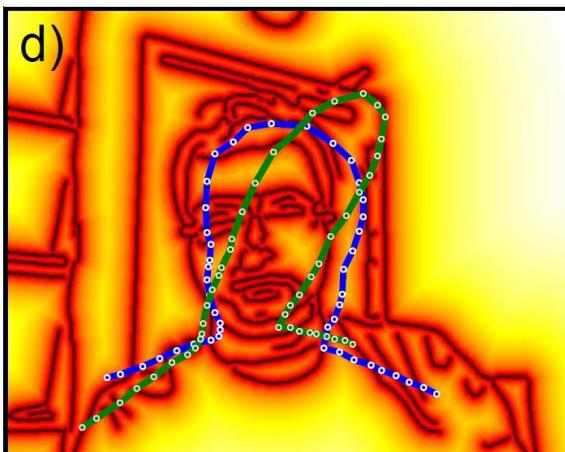
## 4. Resampling

- For  $i = 1, \dots, n$ , set  $x_{t+1,0}^{(i)} \leftarrow \bar{x}_{t+1,0}^{(j)}$  with probability  $\pi_{t+1,0}^{(j)}$
- $t \leftarrow t + 1$  and go to step 2

# Relationships between models



# Shape template model



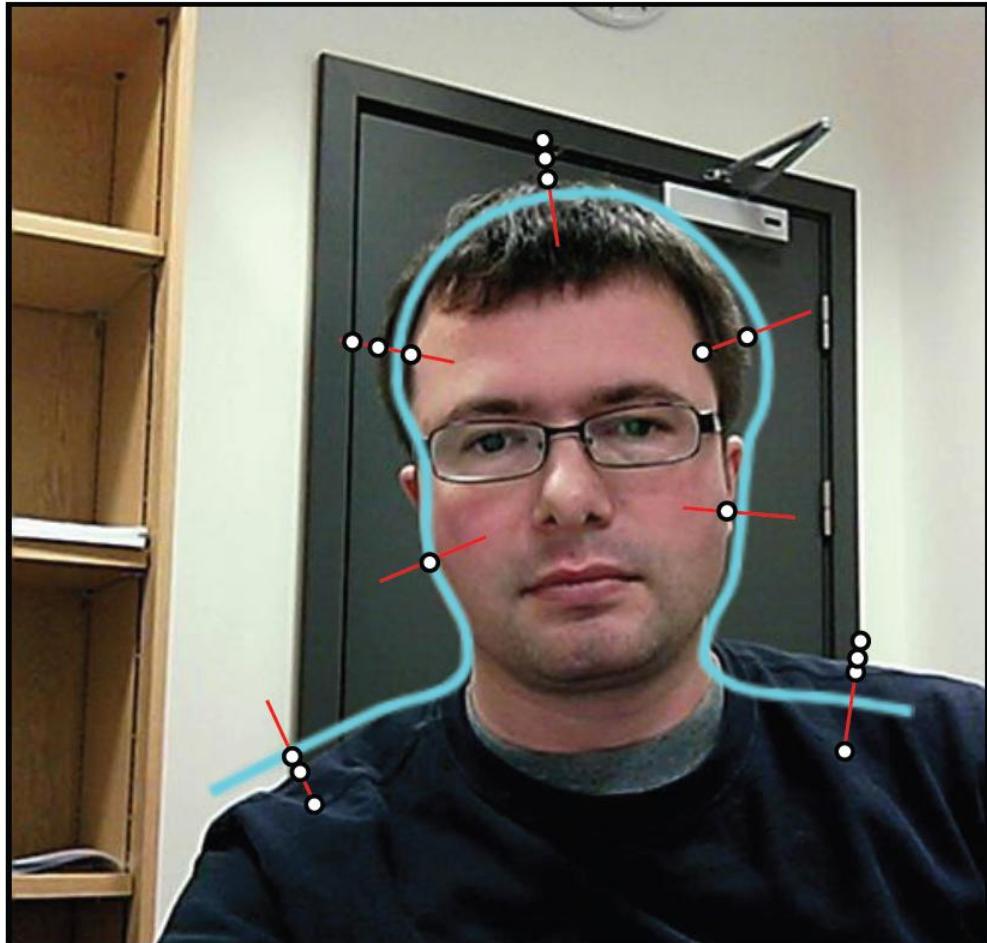
Local optimum

Global optimum

Blue initialization, green estimate

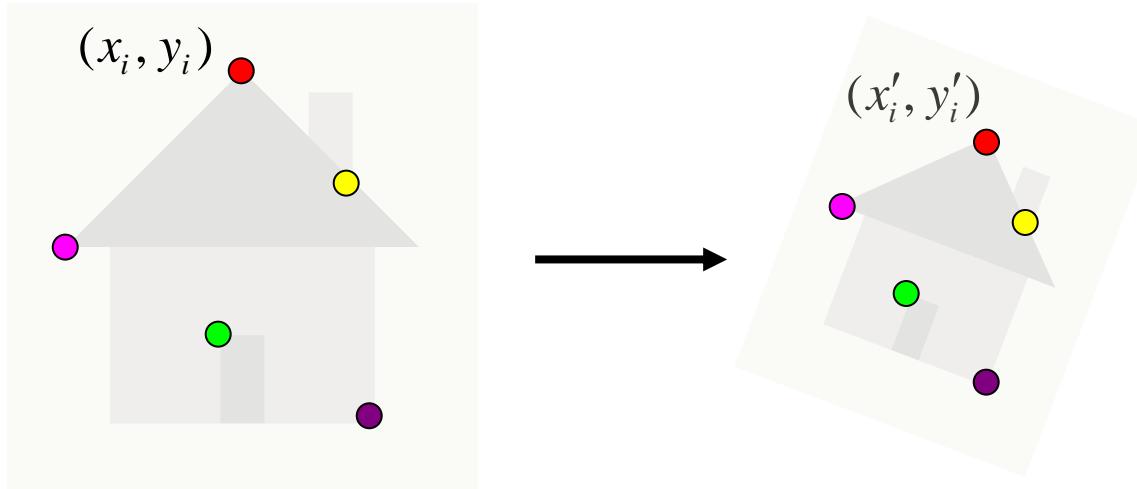
Source: S. Prince

# Iterative closest points



- Find nearest edge point to each landmark point
- Compute transformation in closed form
- Repeat

# Fitting an affine transformation



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\begin{bmatrix} x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} x'_i \\ y'_i \\ \dots \end{bmatrix}$$

Source: K. Grauman

# Least squares problem

A  $m \times n$  matrix of rank  $n$ :

$$Ax = b$$

Solution with pseudo-inverse:

$$A^T(Ax - b) = 0$$

$$x = (A^T A)^{-1} A^T b$$

Pseudo-inverse  $\mathbf{A}^+$

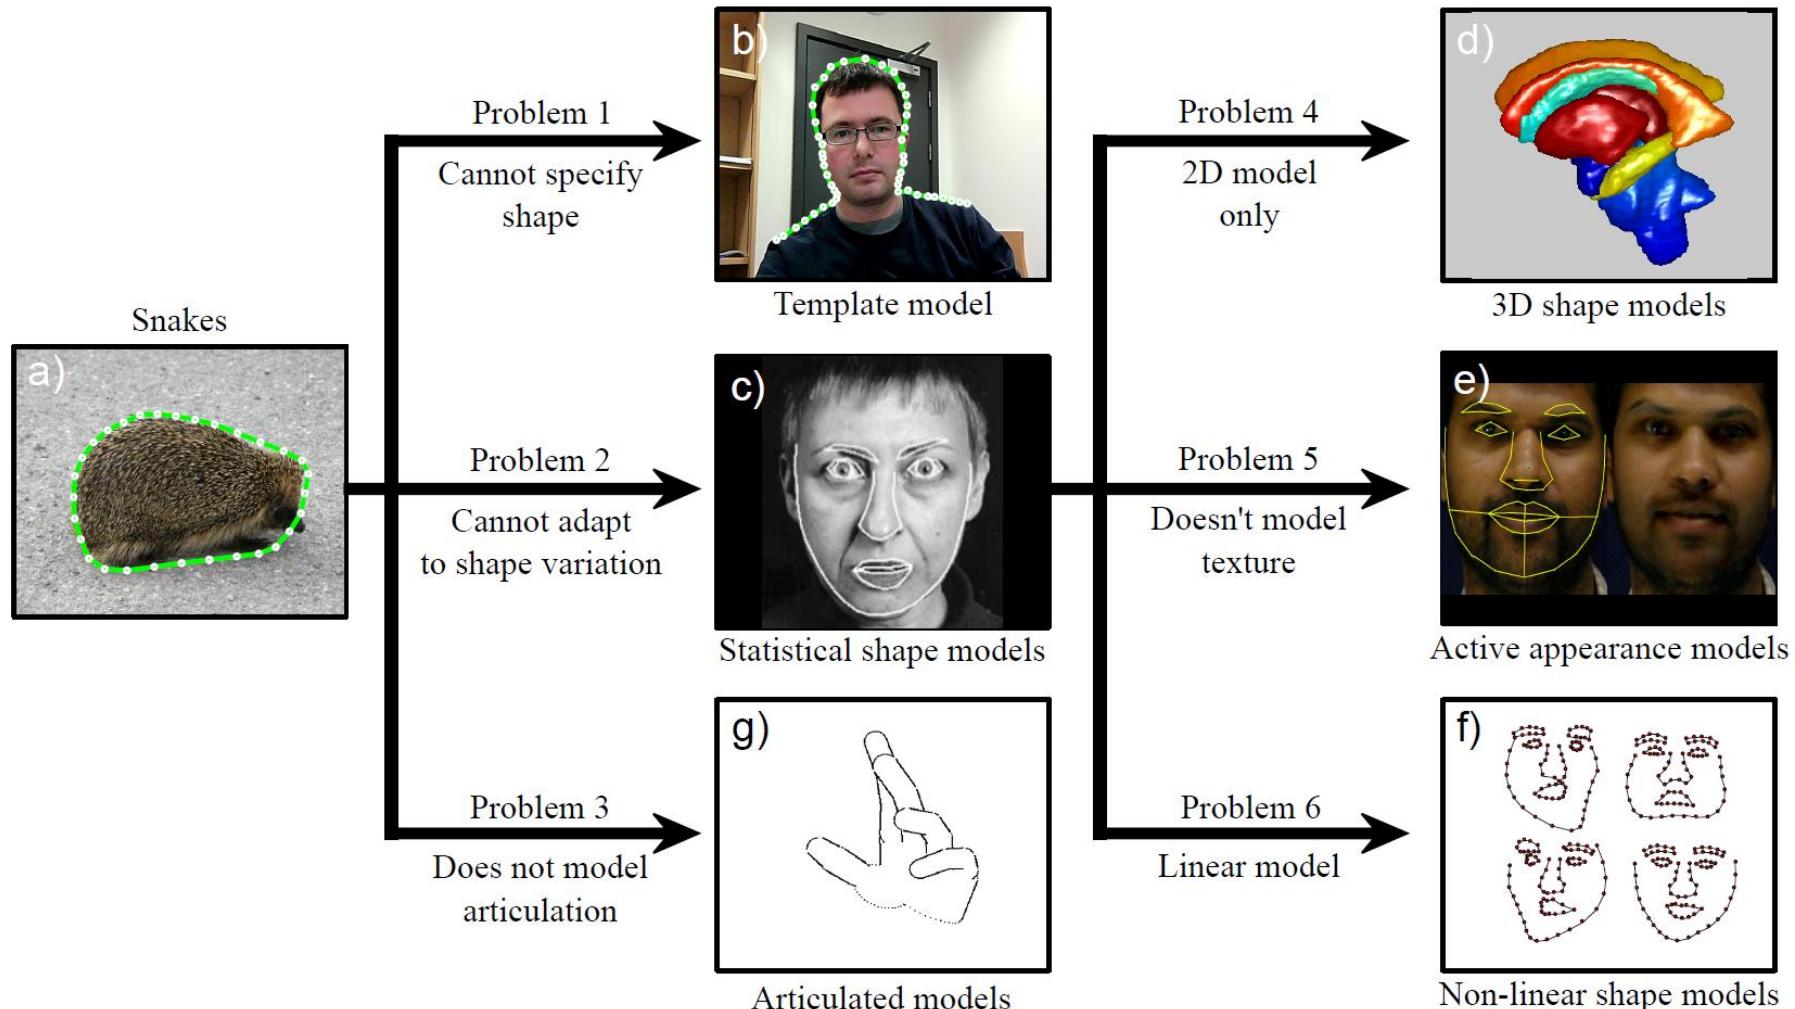
Solution with SVD:

- SVD:  $\mathbf{A} = \mathbf{U}\mathbf{D}\mathbf{V}^T$  ;  $\mathbf{A}^+ = \mathbf{V}\mathbf{D}^{-1}\mathbf{U}^T$
- $\mathbf{b}' = \mathbf{U}^T\mathbf{b}$
- $y_i = b'_i / D_{ii}$
- $\mathbf{x} = \mathbf{V}\mathbf{y}$

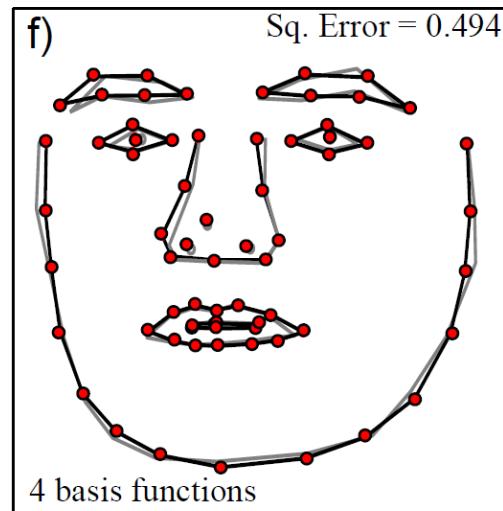
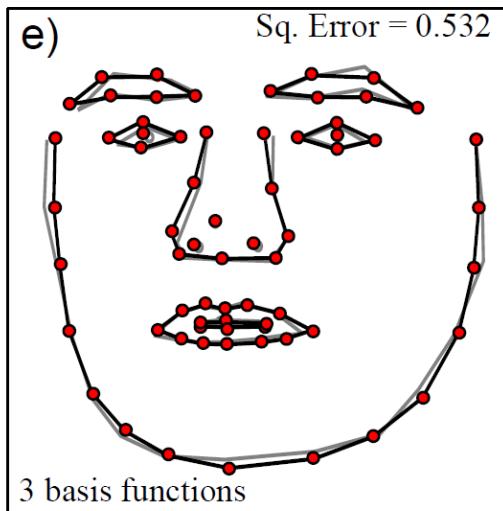
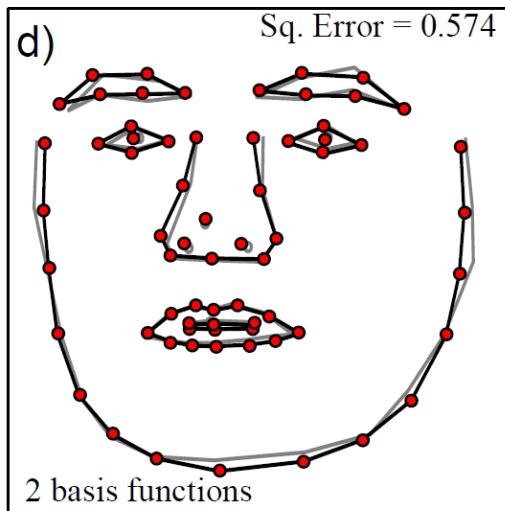
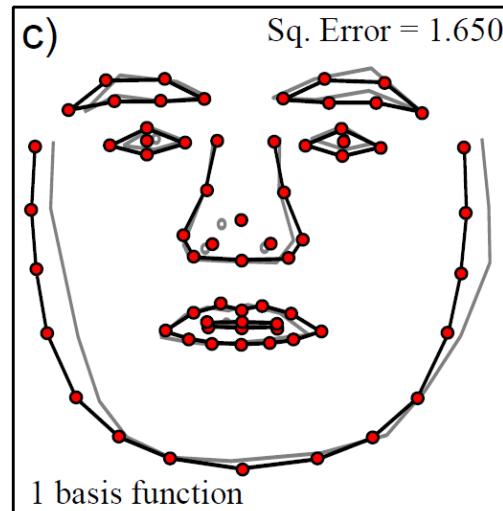
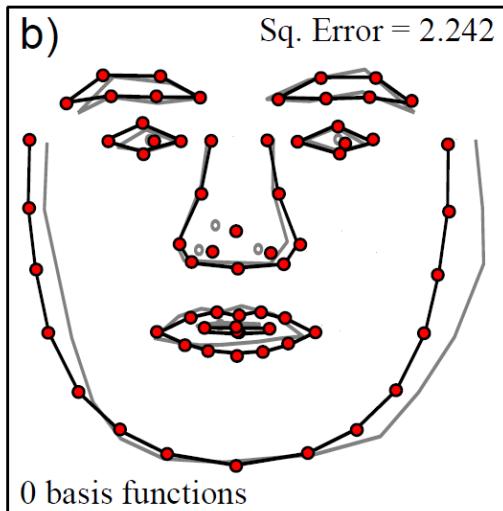
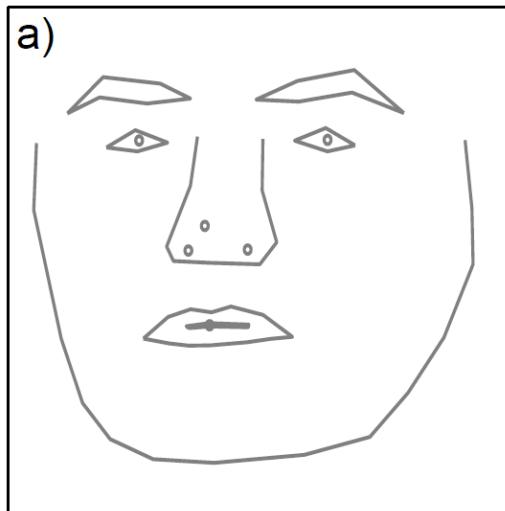
There are different ways to solve the problem. The optimal method depends on size and sparseness of  $\mathbf{A}$

Specific C++ linear algebra libraries exist!

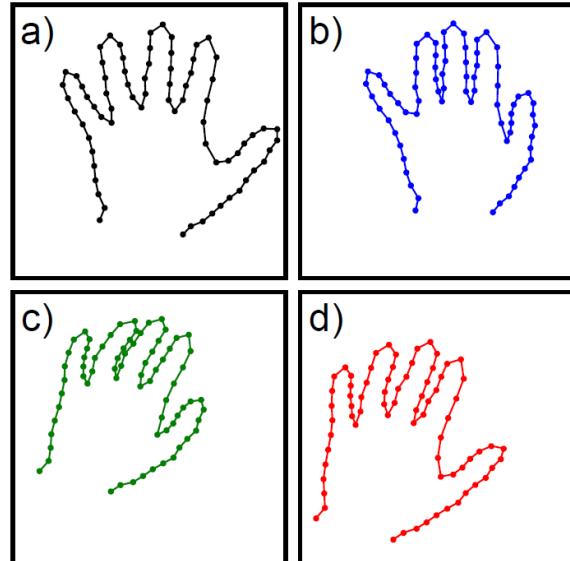
# Relationships between models



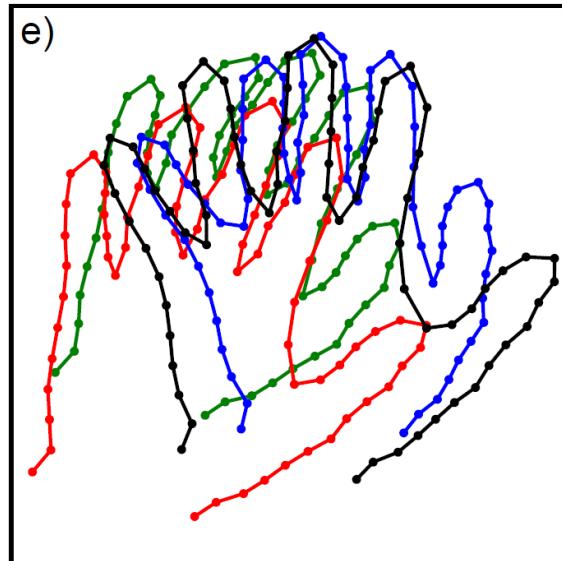
# Subspace shape model



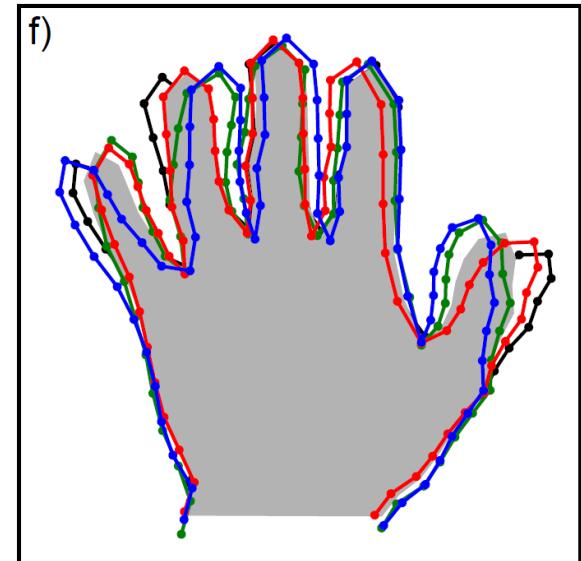
# Generalized Procrustes analysis



Training data



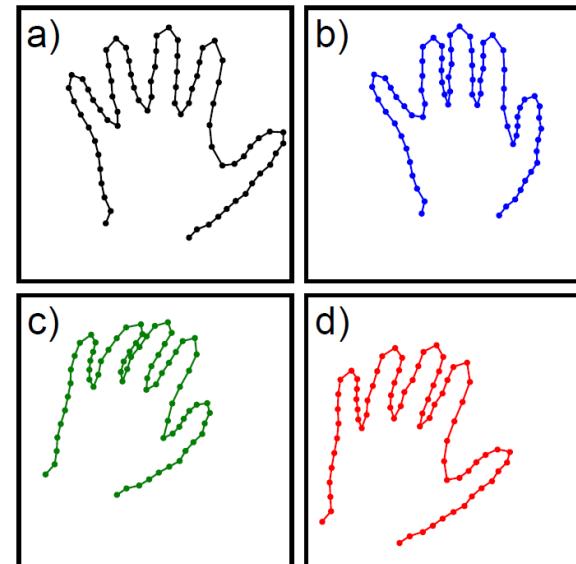
Before alignment



After alignment

# Generalized Procrustes analysis

- Given  $I$  shapes with  $N$  landmarks
- Correspondences between shapes are known (landmarks)
- Take one shape as mean shape
- Iterate:
  - Solve for affine transformation  $\Psi_i$  for each shape  $i$  to the mean  $A_i x = b_i$
  - Apply transformation  $\Psi_i$  to each shape  $i$
  - Compute the mean  $\mu_n$  of each landmark (mean shape)



# Subspace shape model

- Generate data from model:

$$\mathbf{w}_i = \boldsymbol{\mu} + \Phi \mathbf{h}_i + \epsilon_i$$

- $\boldsymbol{\mu}$  is the mean shape
- the matrix  $\Phi = [\phi_1, \phi_2, \dots, \phi_K]$  contains K basis functions in its columns
- $\epsilon_i$  is normal noise with covariance  $\sigma^2 \mathbf{I}$

- Can alternatively write

$$\mathbf{w}_i = \boldsymbol{\mu} + \sum_{k=1}^K \phi_k h_{ik} + \epsilon_i$$

# Approximating with subspace

## Subspace model

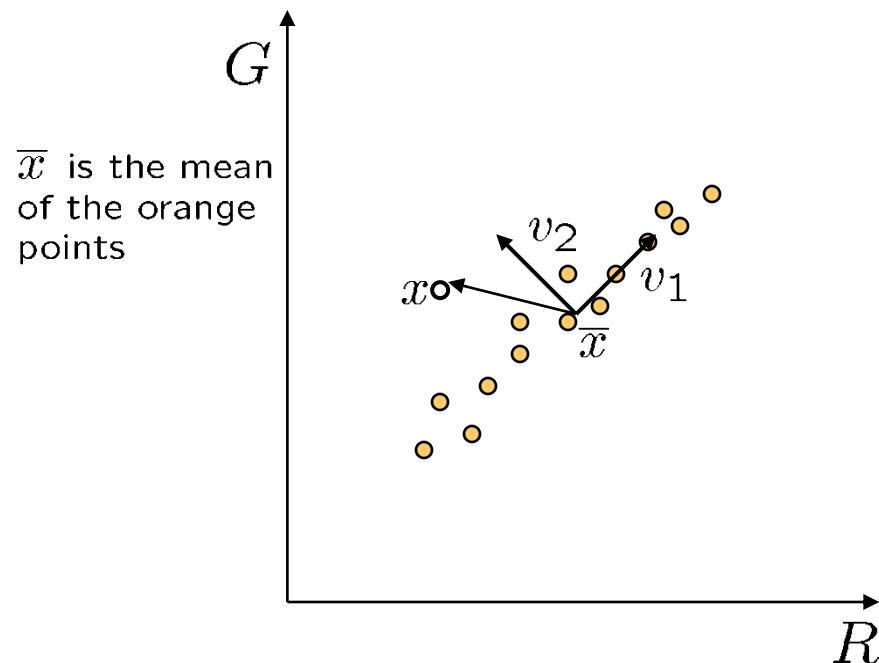
$$\mathbf{w}_i = \mu + \sum_{k=1}^K \phi_k h_{ik} + \epsilon_i$$

Can approximate a vector  $\mathbf{w}$  with a weighted sum of the basis functions

$$\mathbf{w}_i \approx \mu + \sum_{k=1}^K \phi_k h_{ik}$$

Surprising how well this works even with a small number of basis functions

# Linear subspaces



Consider the variation along direction  $\mathbf{v}$  among all of the orange points:

$$var(\mathbf{v}) = \sum_{\text{orange point } \mathbf{x}} \|(\mathbf{x} - \bar{\mathbf{x}})^T \cdot \mathbf{v}\|^2$$

What unit vector  $\mathbf{v}$  minimizes  $var$ ?

$$\mathbf{v}_2 = \min_{\mathbf{v}} \{var(\mathbf{v})\}$$

What unit vector  $\mathbf{v}$  maximizes  $var$ ?

$$\mathbf{v}_1 = \max_{\mathbf{v}} \{var(\mathbf{v})\}$$

$$\begin{aligned} var(\mathbf{v}) &= \sum_{\mathbf{x}} \|(\mathbf{x} - \bar{\mathbf{x}})^T \cdot \mathbf{v}\|^2 \\ &= \sum_{\mathbf{x}} \mathbf{v}^T (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \mathbf{v} \\ &= \mathbf{v}^T \left[ \sum_{\mathbf{x}} (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \right] \mathbf{v} \\ &= \mathbf{v}^T \mathbf{A} \mathbf{v} \quad \text{where } \mathbf{A} = \sum_{\mathbf{x}} (\mathbf{x} - \bar{\mathbf{x}})(\mathbf{x} - \bar{\mathbf{x}})^T \end{aligned}$$

Solution:  $\mathbf{v}_1$  is eigenvector of  $\mathbf{A}$  with *largest* eigenvalue

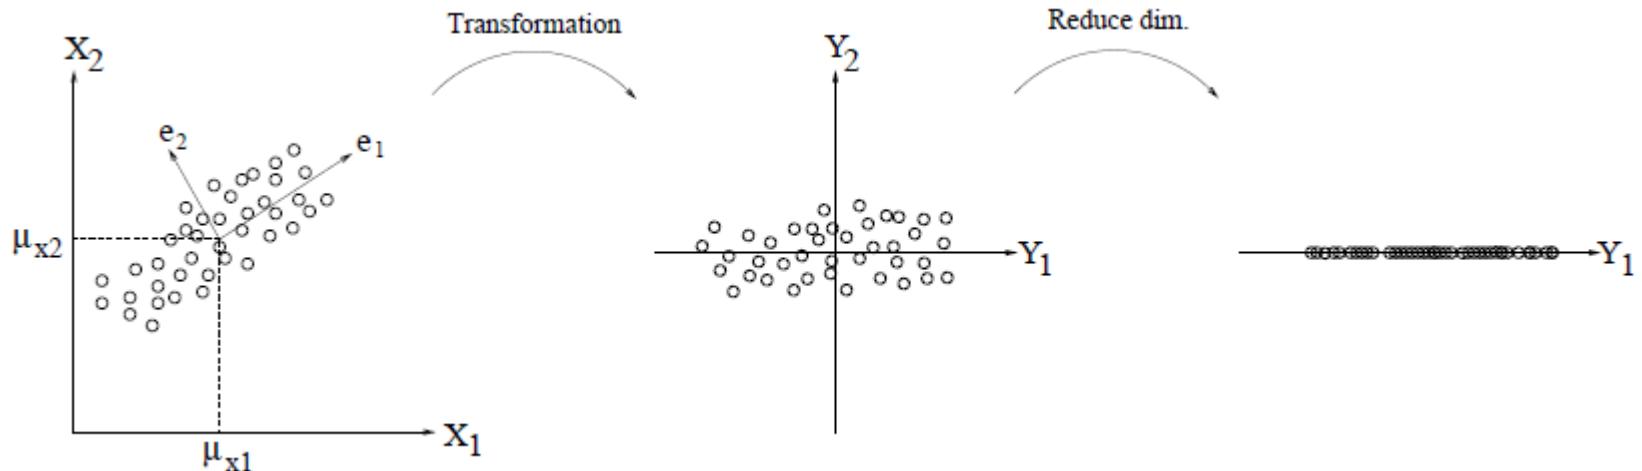
$\mathbf{v}_2$  is eigenvector of  $\mathbf{A}$  with *smallest* eigenvalue

# Principal component analysis (PCA)

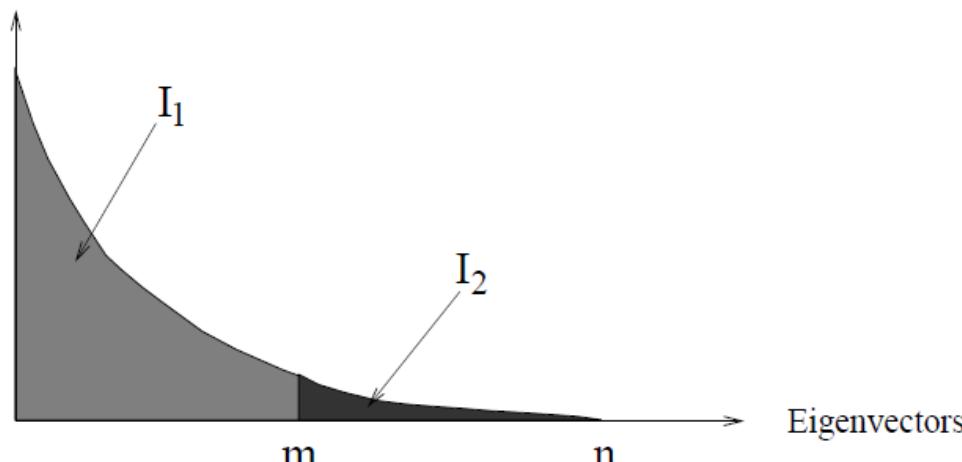
- Given data points  $w_i$
- Compute mean  $\mu$
- $W = (w_1 - \mu, \dots, w_n - \mu)$
- Eigenvalue decomposition:  
 $WW^T = U\Sigma V^T V\Sigma U^T = UL^2U^T$  ( $u_i$  basis vectors)
- SVD:  $W = U\Sigma V^T$

# Principal component analysis (PCA)

- Dimensionality reduction



Eigenvalues



$$\frac{\sum_{i=1}^{i=K} L_{ii}^2}{\sum_{i=1}^{i=D} L_{ii}^2} > 1 - \epsilon$$

Source: T. Moeslund

# Learning PPCA

Learn parameters  $\mu$ ,  $\Phi$  and  $\sigma^2$  from data  $\{\mathbf{w}_i\}_{i=1}^I$   
where  $\mathbf{w}_i = [\mathbf{w}_{i1}^T, \mathbf{w}_{i2}^T, \dots, \mathbf{w}_{iN}^T]^T$ .

Learn mean:  $\mu = \frac{\sum_{i=1}^I \mathbf{w}_i}{I}$

Then set  $\mathbf{W} = [\mathbf{w}_1 - \mu, \mathbf{w}_2 - \mu, \dots, \mathbf{w}_I - \mu]$   
and compute eigen-decomposition

$$\mathbf{W}\mathbf{W}^T = \mathbf{U}\mathbf{L}^2\mathbf{U}^T$$

Choose parameters

$$\hat{\sigma}^2 = \frac{1}{D-K} \sum_{j=K+1}^D L_{jj}^2$$

$$\hat{\Phi} = \mathbf{U}_K (\mathbf{L}_K^2 - \hat{\sigma}^2 \mathbf{I})^{1/2}$$

# Simple version with PCA

1. Initialise the shape parameters,  $\mathbf{b}$ , to zero (the mean shape).
2. Generate the model point positions using  $\mathbf{x} = \bar{\mathbf{x}} + \mathbf{P}\mathbf{b}$
3. Find the pose parameters  $(X_t, Y_t, s, \theta)$  which best align the model points  $\mathbf{x}$  to the current found points  $\mathbf{Y}$

$$T_{X_t, Y_t, s, \theta} \begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} X_t \\ Y_t \end{pmatrix} + \begin{pmatrix} s \cos \theta & -s \sin \theta \\ s \sin \theta & s \cos \theta \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix}$$

4. Project  $\mathbf{Y}$  into the model co-ordinate frame by inverting the transformation  $T$ :

$$\mathbf{y} = T_{X_t, Y_t, s, \theta}^{-1}(\mathbf{Y})$$

5. Project  $\mathbf{y}$  into the tangent plane to  $\bar{\mathbf{x}}$  by scaling:  $\mathbf{y}' = \mathbf{y}/(\mathbf{y} \cdot \bar{\mathbf{x}})$ . (optional)

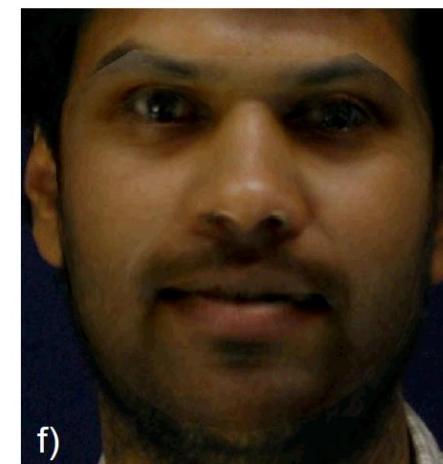
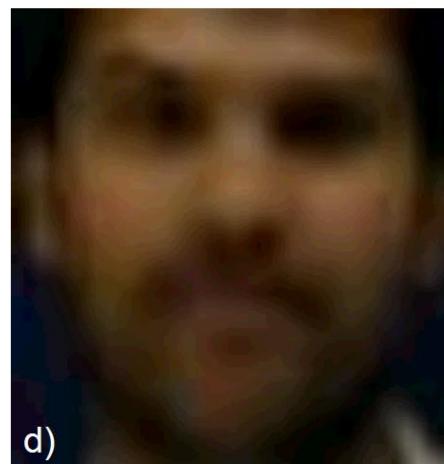
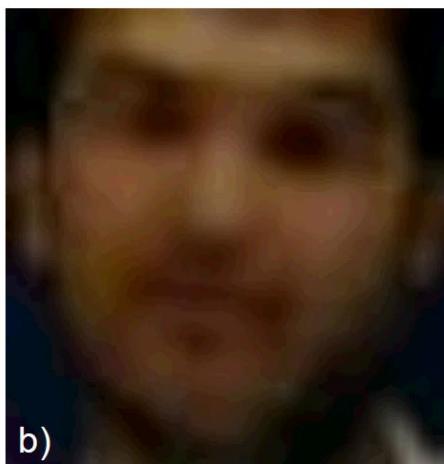
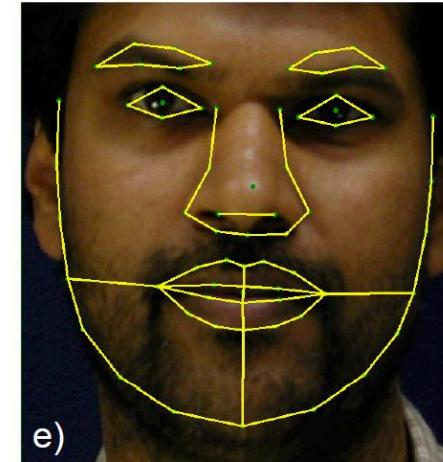
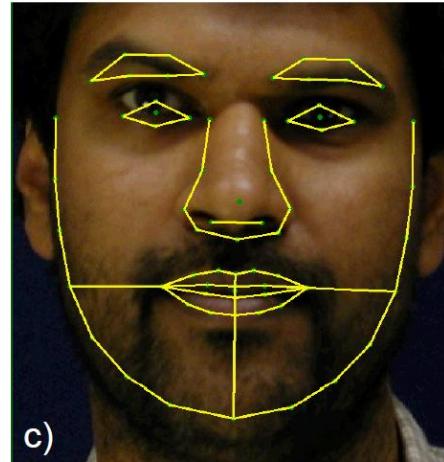
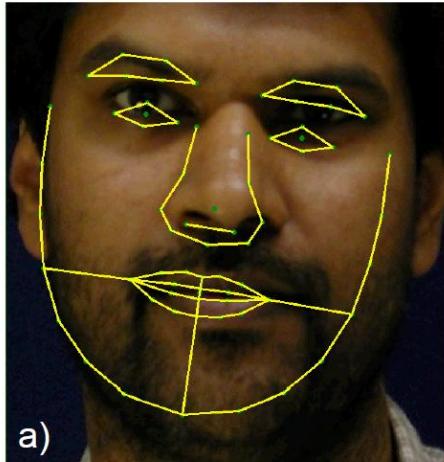
6. Update the model parameters to match to  $\mathbf{y}'$

$$\mathbf{b} = \mathbf{P}^T (\mathbf{y}' - \bar{\mathbf{x}})$$

7. If not converged, return to step 2.

T. Cootes. An Introduction to Active Shape Models.

# Statistical models for shape and appearance



Source: T. Cootes

# Motivation: mosaics

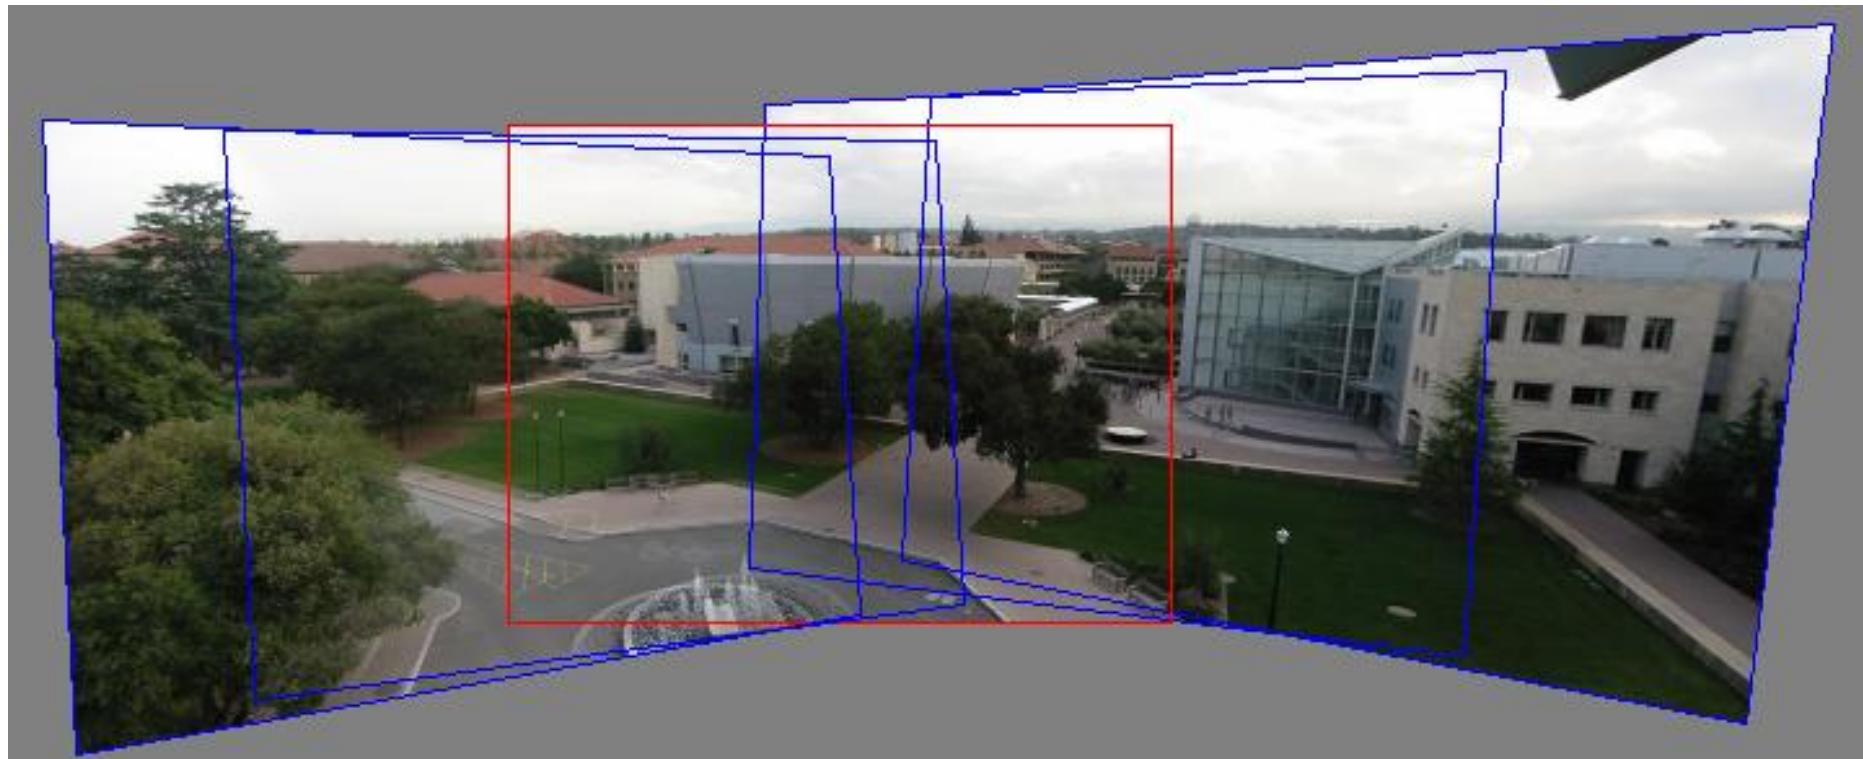
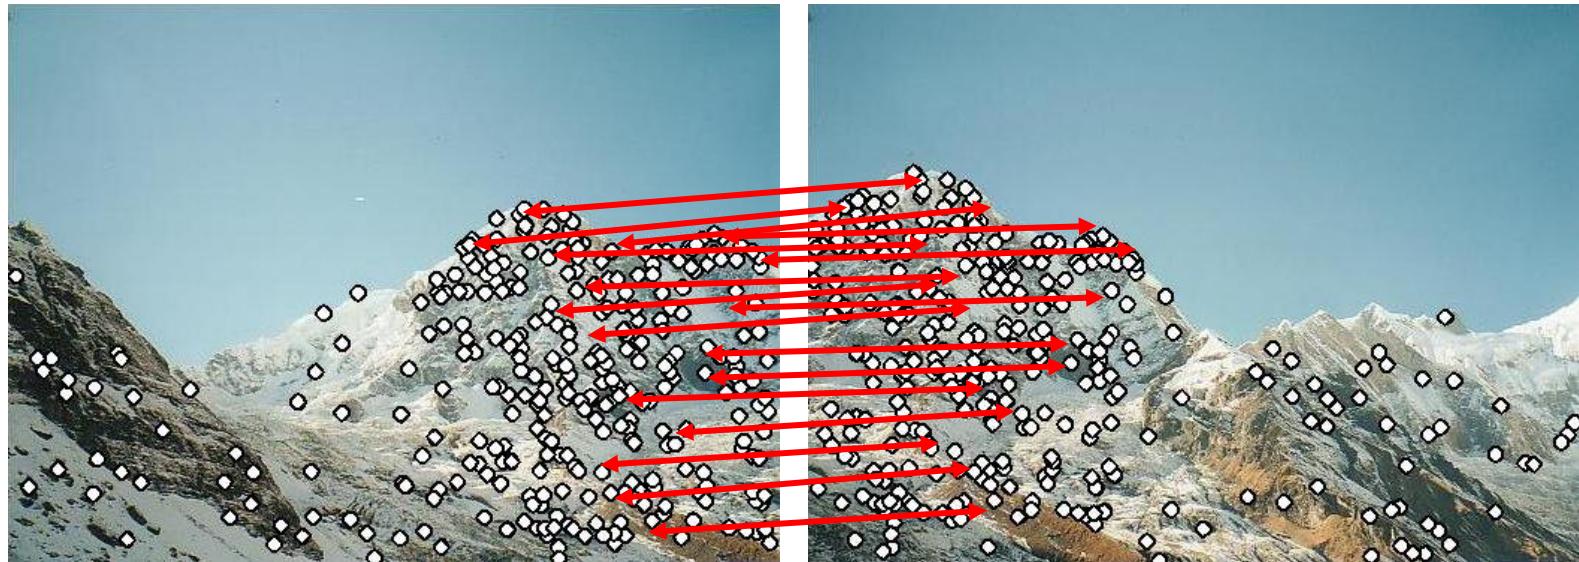


Image from [http://graphics.cs.cmu.edu/courses/15-463/2010\\_fall/](http://graphics.cs.cmu.edu/courses/15-463/2010_fall/)

# Robust feature-based alignment



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$  (small group of putative matches that are related by  $T$ )
  - *Verify* transformation (search for other matches consistent with  $T$ )

# Robust feature-based alignment



- Extract features
- Compute *putative matches*
- Loop:
  - *Hypothesize* transformation  $T$  (small group of putative matches that are related by  $T$ )
  - *Verify* transformation (search for other matches consistent with  $T$ )

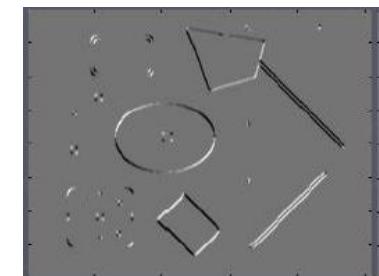
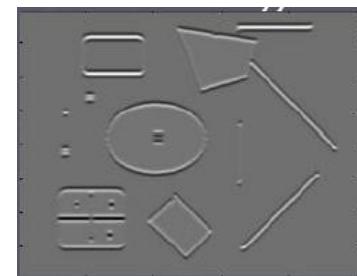
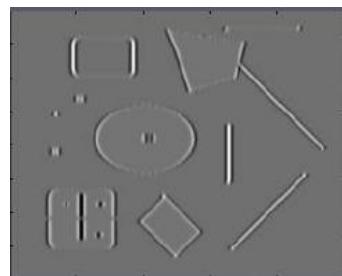
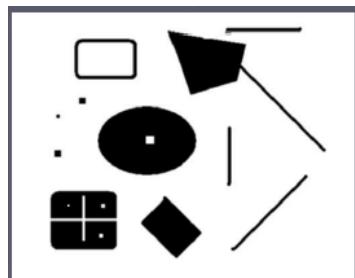
# Local features: desired properties

- Repeatability
  - The same feature can be found in several images despite geometric and photometric transformations
- Saliency
  - Each feature has a distinctive description
- Locality
  - A feature occupies a relatively small area of the image; robust to clutter and occlusion
- Sometimes: Compactness and efficiency
  - Low dimensionality

# Corners as distinctive interest points

$$M = \sum w(x, y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

2 x 2 matrix of image derivatives (averaged in neighborhood of a point).



Notation:

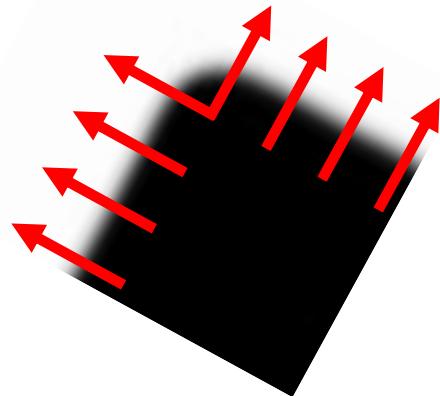
$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$

# What does this matrix reveal?

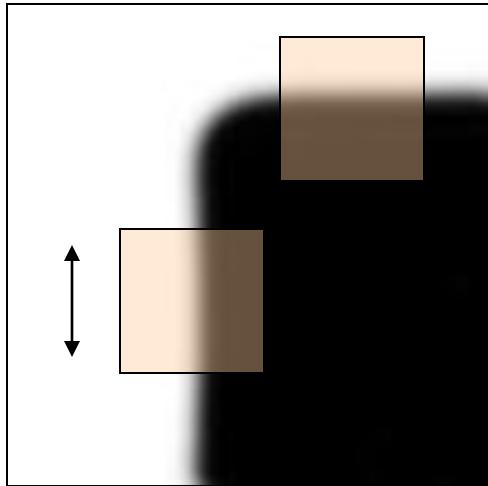
Since  $M$  is symmetric, we have  $M = X \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} X^T$



$$Mx_i = \lambda_i x_i$$

The *eigenvalues* of  $M$  reveal the amount of intensity change in the two principal orthogonal gradient directions in the window.

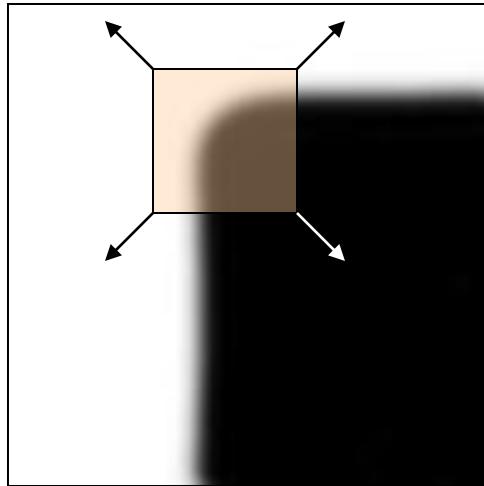
# Corner response function



“edge”:

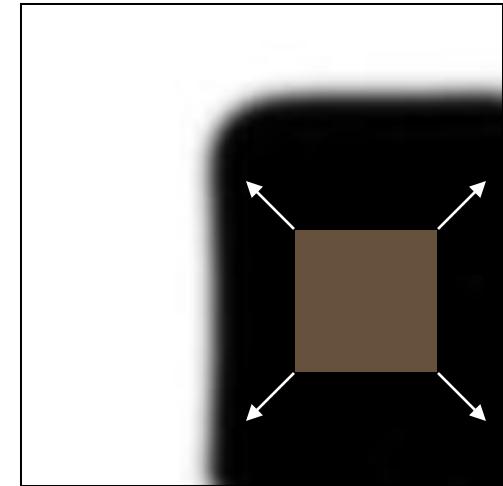
$$\lambda_1 \gg \lambda_2$$

$$\lambda_2 \gg \lambda_1$$



“corner”:

$\lambda_1$  and  $\lambda_2$  are large,  
 $\lambda_1 \sim \lambda_2$ ;



“flat” region

$\lambda_1$  and  $\lambda_2$  are  
small;

$$f = |M| - k (\text{Tr}(M))^2 = \lambda_1 \lambda_2 - k (\lambda_1 + \lambda_2)^2$$

# Harris corner detector

- 1) Compute  $M$  matrix for each image window to get their *cornerness* scores.
- 2) Find points whose surrounding window gave large corner response ( $f > \text{threshold}$ )
- 3) Take the points of local maxima, i.e., perform non-maximum suppression

# Example of Harris application



# Automatic scale selection

## Intuition:

- Find scale that gives local maxima of some function  $f$  in both position and scale.

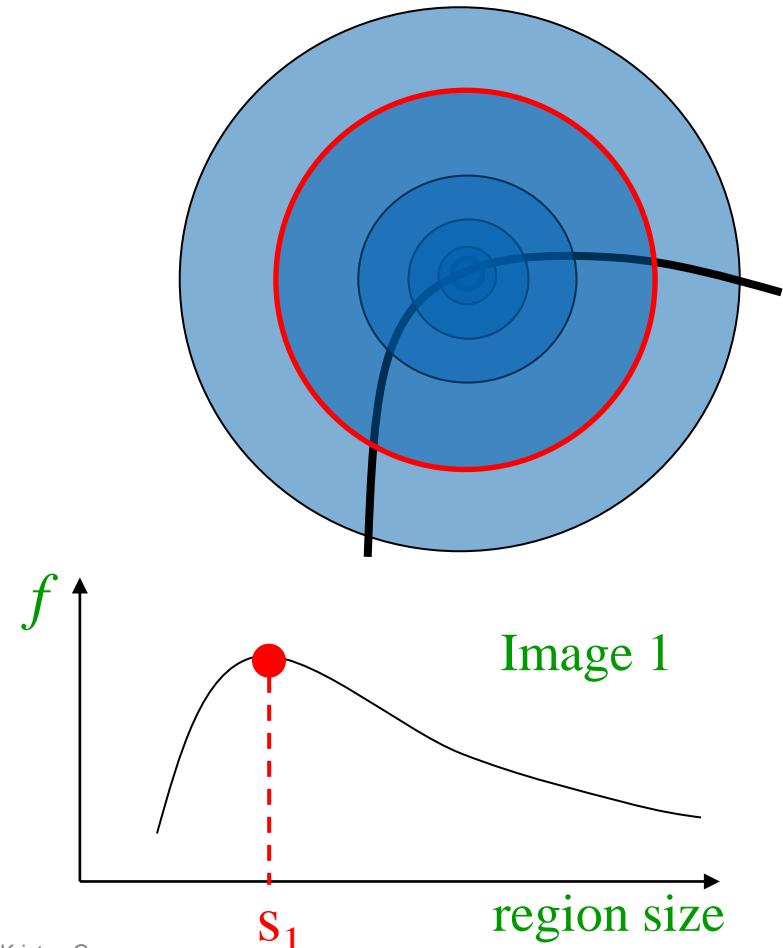


Image 1

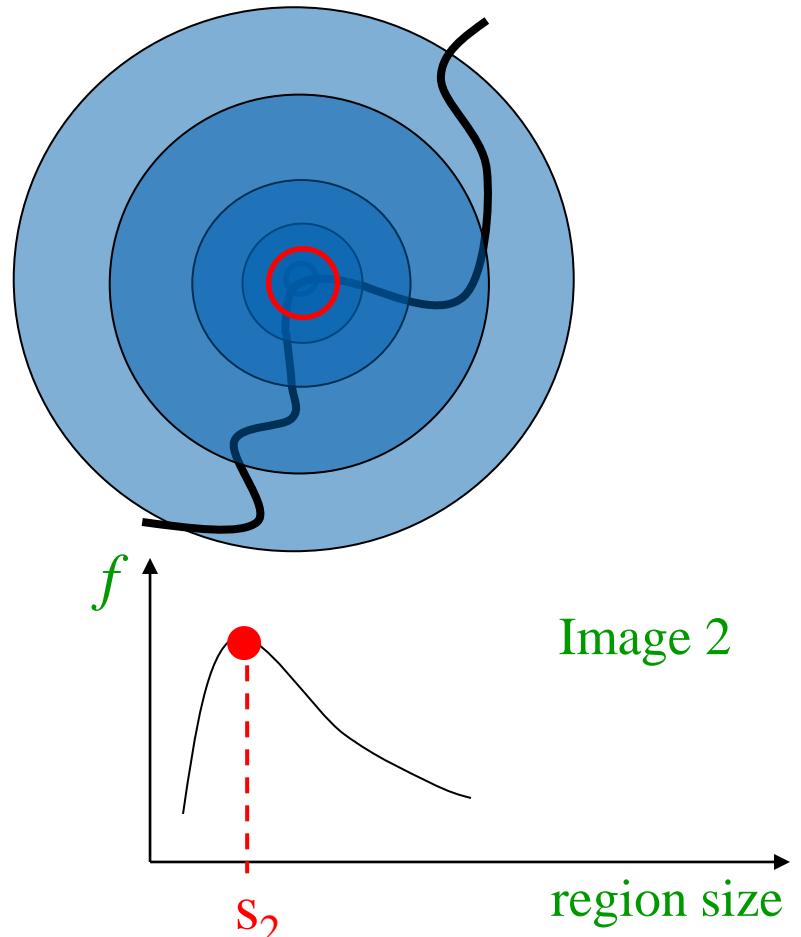
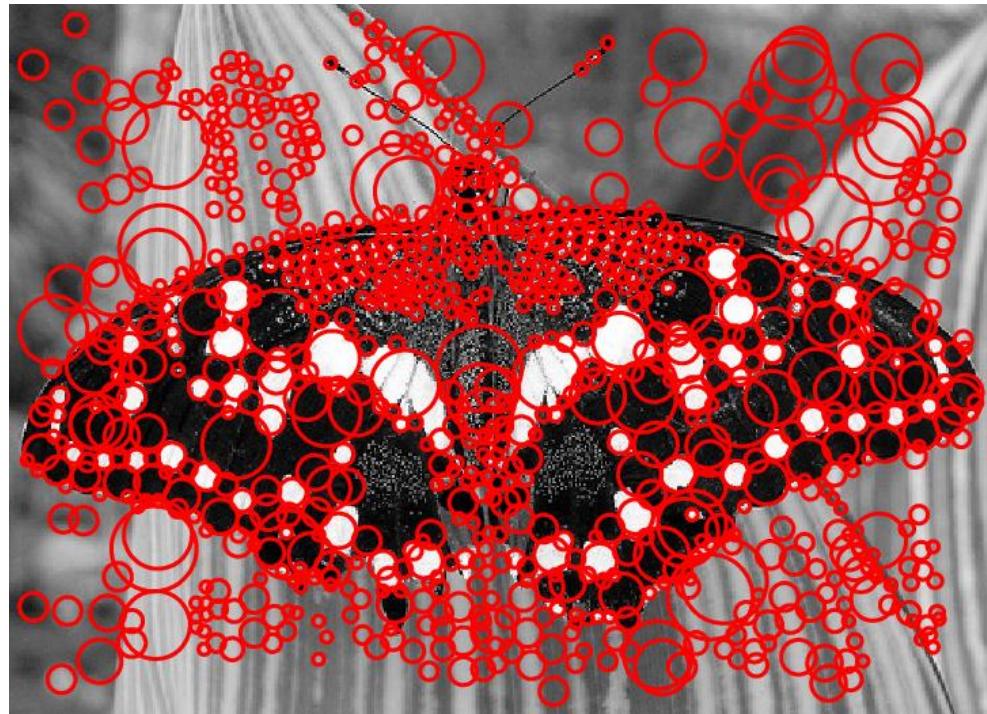
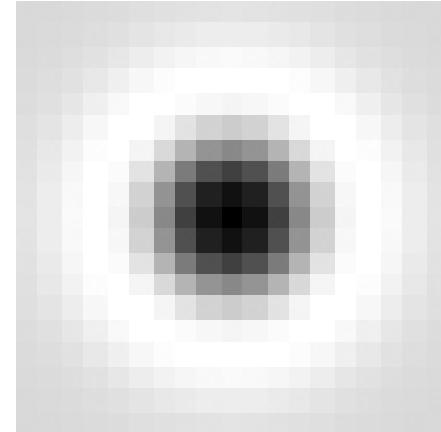
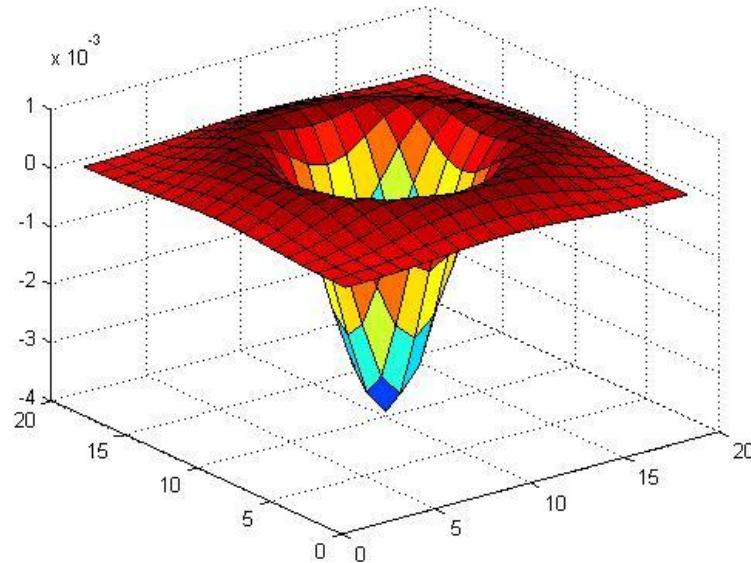


Image 2

# Scale invariant blob detection



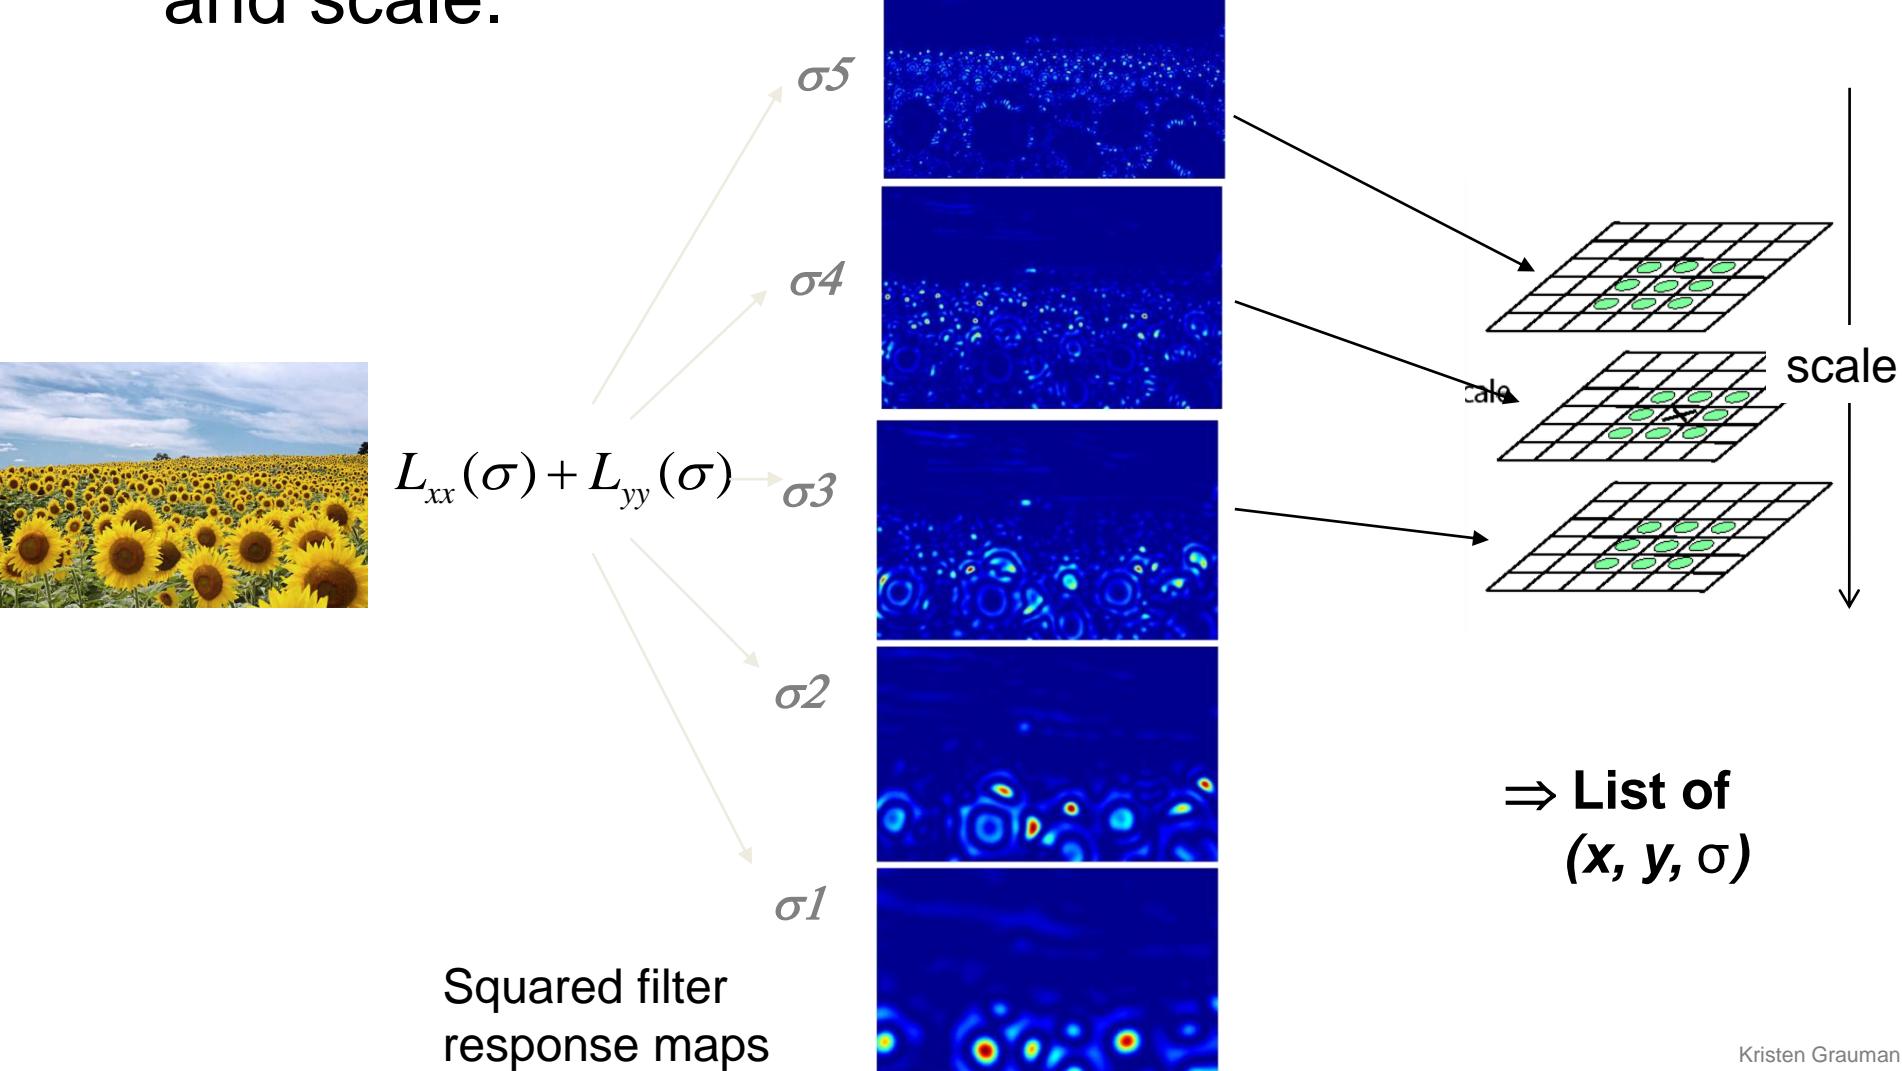
# Blob detection in 2D



- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D

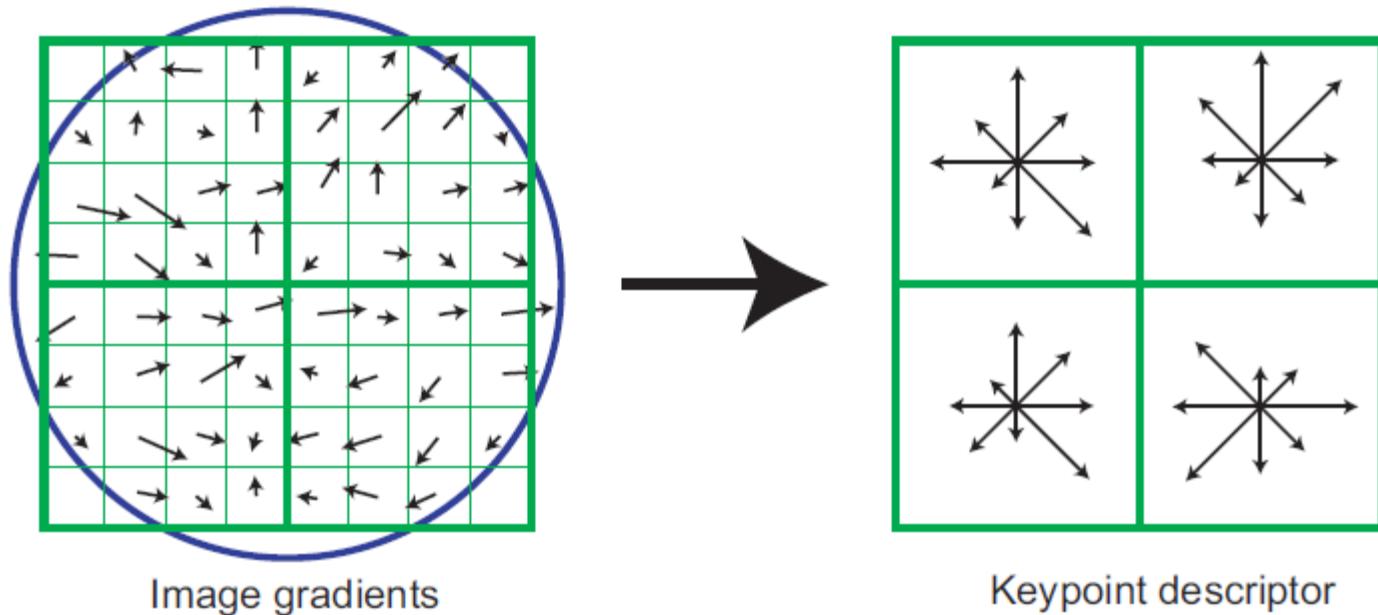
# Scale invariant interest points

Interest points are local maxima in both position and scale.



# SIFT descriptor [Lowe 2004]

- 128 D vector



- Normalize to 1 (robust to illumination changes)
- Entries above 0.2 are set to 0.2 (reduce impact of very strong gradients e.g. shadow)
- Normalize to 1

# Matching local features



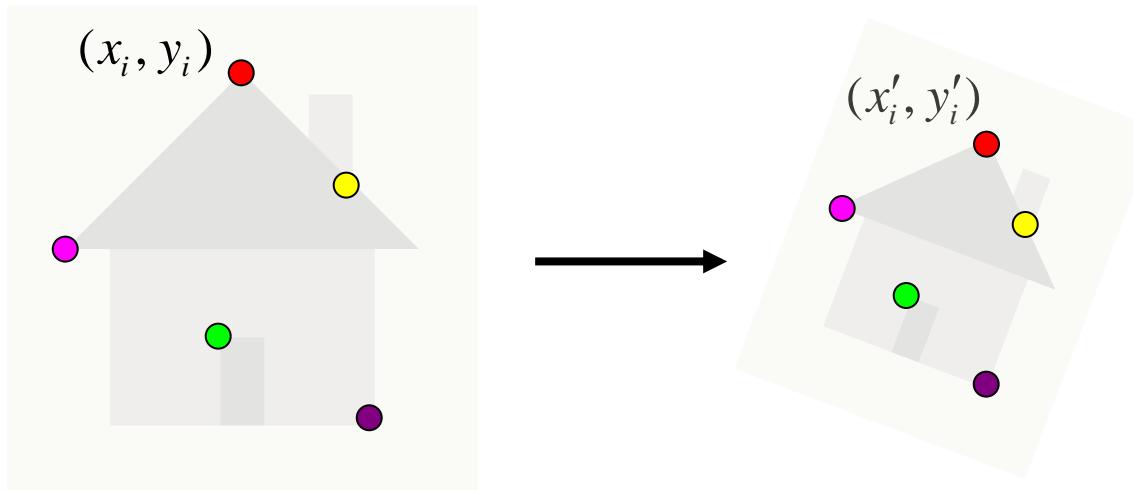
Image 1



Image 2

To add robustness to matching, can consider **ratio** :  
distance to best match / distance to second best match  
If low, first match looks good.  
If high and close to one, could be ambiguous match.

# Fitting an affine transformation

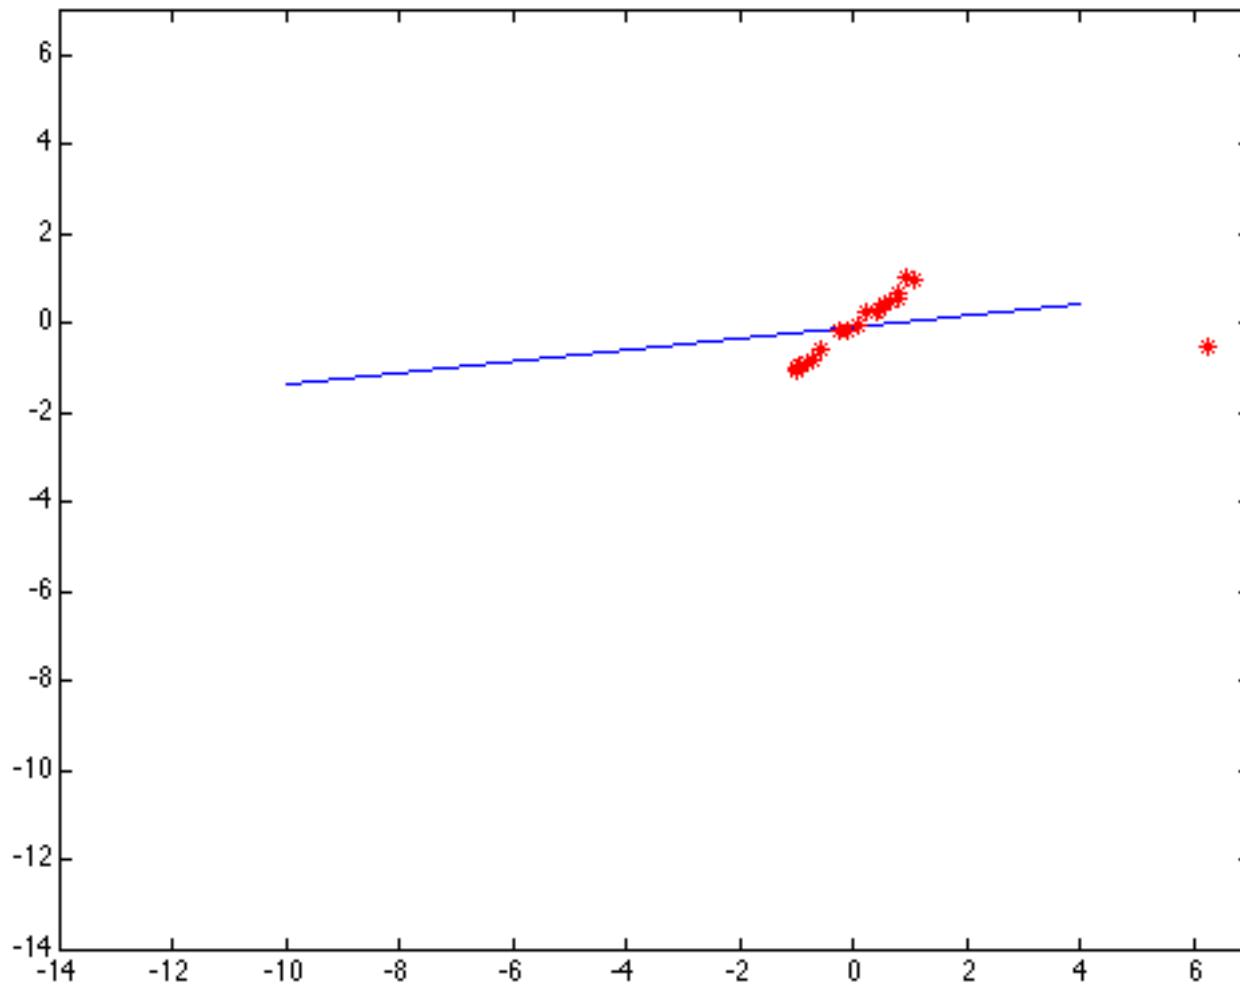


$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

$$\begin{bmatrix} x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} x'_i \\ y'_i \\ \dots \end{bmatrix}$$

Source: K. Grauman

# Outliers affect least squares fit



# RANSAC: General form

## RANSAC loop:

1. Randomly select a *seed group* of points on which to base transformation estimate (e.g., a group of matches)
2. Compute transformation from seed group
3. Find *inliers* to this transformation
4. If the number of inliers is sufficiently large, re-compute estimate of transformation on all of the inliers

Keep the transformation with the largest number of inliers

# How Many Samples?

With confidence  $p$

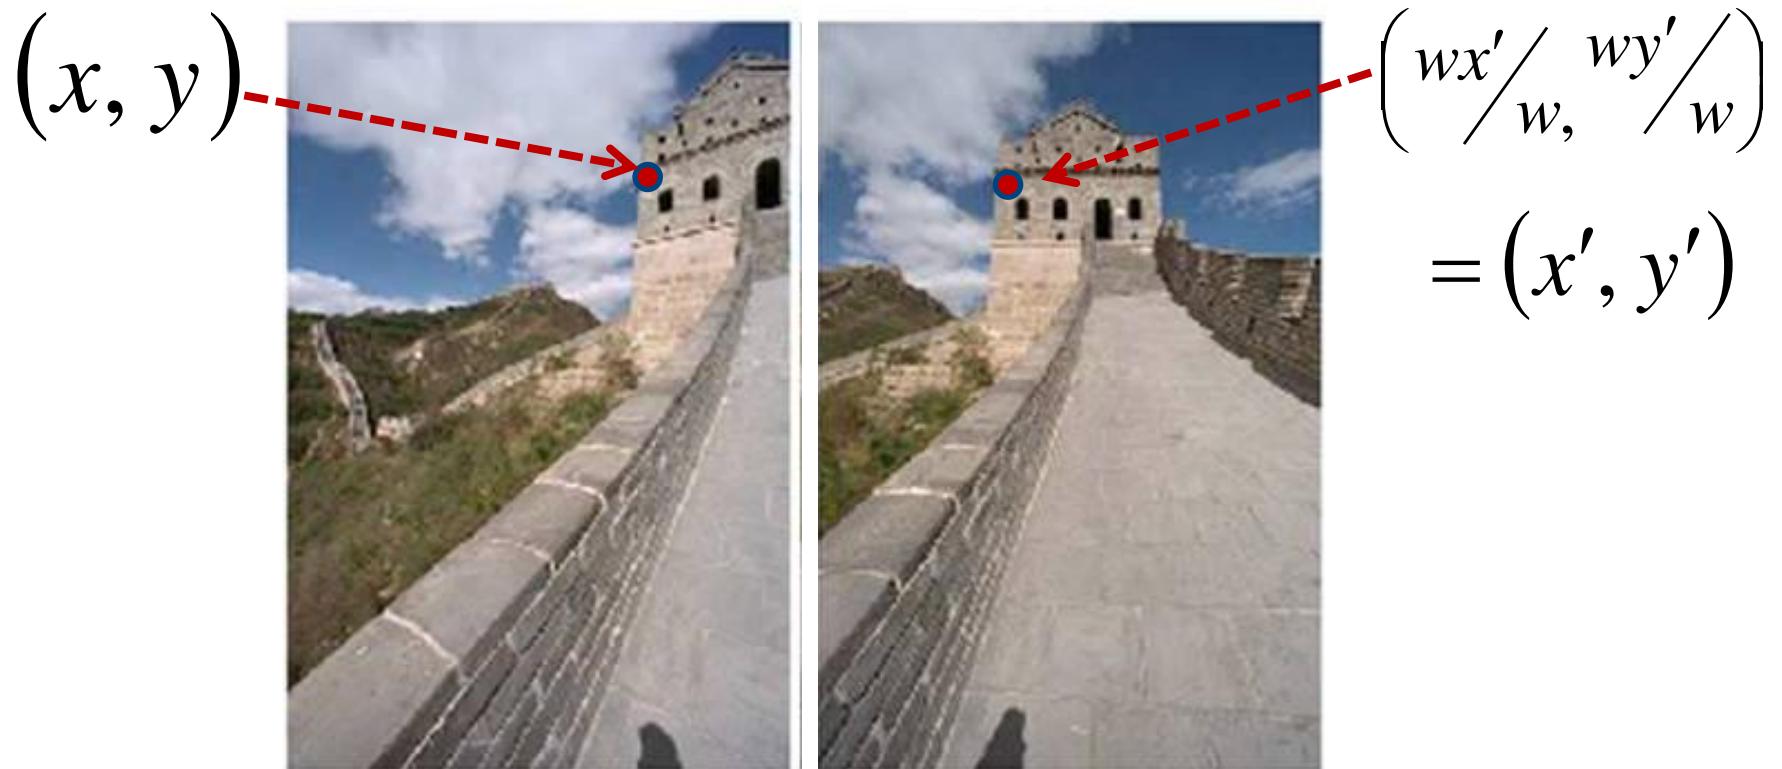
$$P(\text{bad } k \text{ times}) = (1 - P(\text{good}))^k \leq 1 - p$$

$$k \log(1 - P(\text{good})) \leq \log(1 - p)$$

$$k \geq \log(1 - p) / \log(1 - P(\text{good}))$$

$$P(\text{good}) = \frac{\binom{I}{m}}{\binom{N}{m}} = \prod_{j=0}^{m-1} \frac{I-j}{N-j}$$

# Homography



$$\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix} = \begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$\mathbf{p}' \qquad \mathbf{H} \qquad \mathbf{p}$

# Solving for homographies

$$\begin{bmatrix} {}^w x'_i \\ {}^w y'_i \\ w \end{bmatrix} \cong \begin{bmatrix} h_{00} & h_{01} & h_{02} \\ h_{10} & h_{11} & h_{12} \\ h_{20} & h_{21} & h_{22} \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

$$x'_i = \frac{h_{00}x_i + h_{01}y_i + h_{02}}{h_{20}x_i + h_{21}y_i + h_{22}}$$

$$y'_i = \frac{h_{10}x_i + h_{11}y_i + h_{12}}{h_{20}x_i + h_{21}y_i + h_{22}}$$

$$x'_i(h_{20}x_i + h_{21}y_i + h_{22}) = h_{00}x_i + h_{01}y_i + h_{02}$$

$$y'_i(h_{20}x_i + h_{21}y_i + h_{22}) = h_{10}x_i + h_{11}y_i + h_{12}$$

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 & -x'_i x_i & -x'_i y_i & -x'_i \\ 0 & 0 & 0 & x_i & y_i & 1 & -y'_i x_i & -y'_i y_i & -y'_i \end{bmatrix} \begin{bmatrix} h_{00} \\ h_{01} \\ h_{02} \\ h_{10} \\ h_{11} \\ h_{12} \\ h_{20} \\ h_{21} \\ h_{22} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

# Solving for homographies

$$\begin{bmatrix}
 x_1 & y_1 & 1 & 0 & 0 & 0 & -x'_1 x_1 & -x'_1 y_1 & -x'_1 \\
 0 & 0 & 0 & x_1 & y_1 & 1 & -y'_1 x_1 & -y'_1 y_1 & -y'_1 \\
 & & & & & \vdots & & & \\
 x_n & y_n & 1 & 0 & 0 & 0 & -x'_n x_n & -x'_n y_n & -x'_n \\
 0 & 0 & 0 & x_n & y_n & 1 & -y'_n x_n & -y'_n y_n & -y'_n
 \end{bmatrix} \begin{bmatrix} h_{00} \\ h_{01} \\ h_{02} \\ h_{10} \\ h_{11} \\ h_{12} \\ h_{20} \\ h_{21} \\ h_{22} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 0 \end{bmatrix}$$

A  
 $2n \times 9$ 
h  
 $9$ 0  
 $2n$

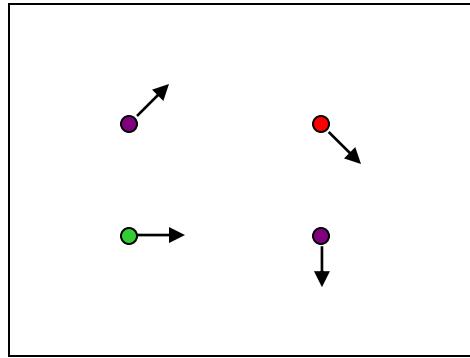
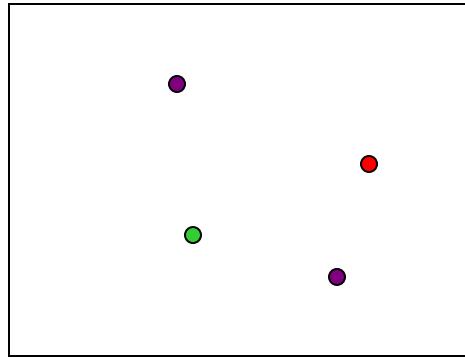
Defines a least squares problem: minimize  $\|Ah - 0\|^2$

- Since  $h$  is only defined up to scale, solve for unit vector  $\hat{h}$
- Solution:  $\hat{h} = \text{eigenvector of } A^T A \text{ with smallest eigenvalue}$
- Works with 4 or more points

# Visual motion

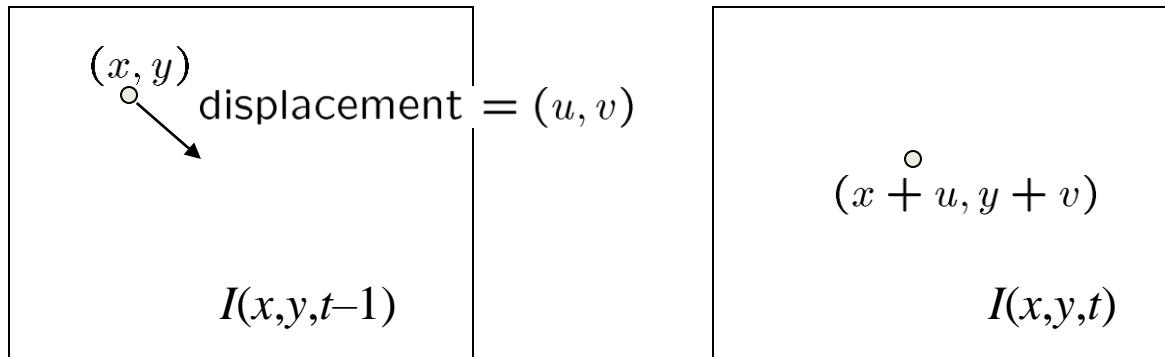


# Estimating optical flow

 $I(x,y,t-1)$  $I(x,y,t)$ 

- Given two subsequent frames, estimate the apparent motion field  $u(x,y)$  and  $v(x,y)$  between them
- Key assumptions
  - Brightness constancy: projection of the same point looks the same in every frame
  - Small motion: points do not move very far
  - Spatial coherence: points move like their neighbors

# The brightness constancy constraint



- Brightness Constancy Equation:

$$I(x, y, t-1) = I(x + u(x, y), y + v(x, y), t)$$

Linearizing the right side using Taylor expansion:

$$I(x, y, t-1) \approx I(x, y, t) + I_x \cdot u(x, y) + I_y \cdot v(x, y)$$

Hence,

$$I_x \cdot u + I_y \cdot v + I_t \approx 0$$

# Lucas-Kanade flow

- Linear least squares problem

$$\begin{bmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_{25}) & I_y(p_{25}) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_{25}) \end{bmatrix}$$

$A \quad d = b$   
 $25 \times 2 \quad 2 \times 1 \quad 25 \times 1$

Solution given by  $(A^T A) d = A^T b$

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$

$A^T A \qquad \qquad \qquad A^T b$

The summations are over all pixels in the window

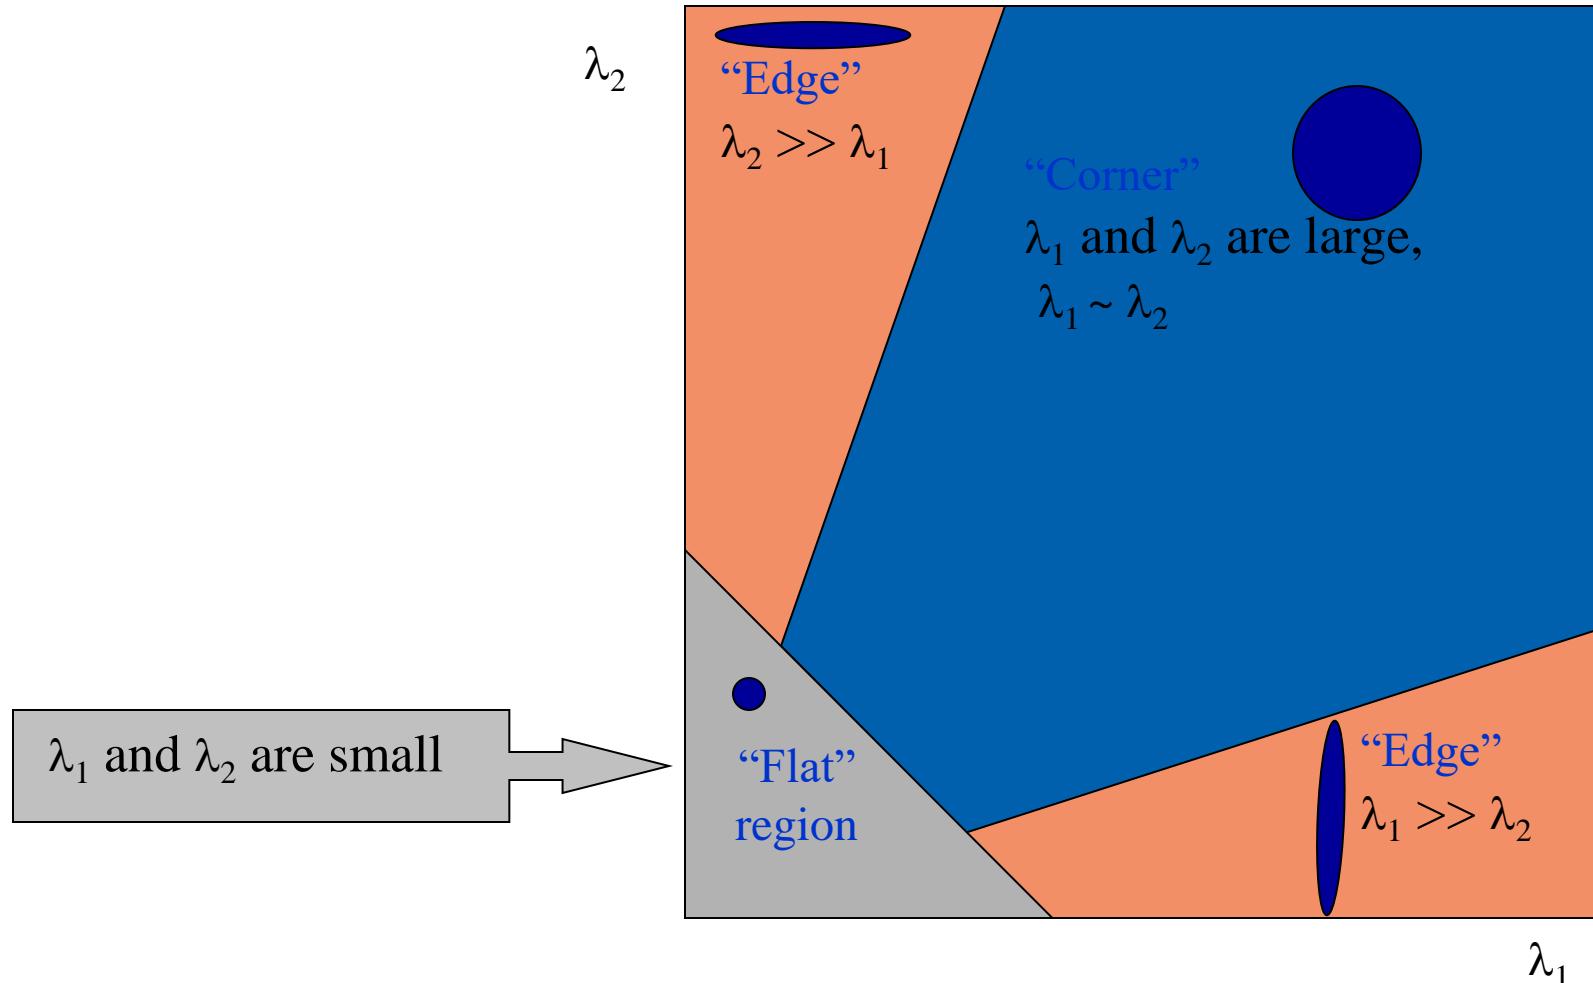
# Lucas-Kanade flow

$$\begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \begin{bmatrix} \sum I_x I_t \\ \sum I_y I_t \end{bmatrix}$$
$$A^T A \quad A^T b$$

- Recall the Harris corner detector:  $M = A^T A$  is the *second moment matrix*
- We can figure out whether the system is solvable by looking at the eigenvalues of the second moment matrix
  - The eigenvectors and eigenvalues of  $M$  relate to edge direction and magnitude
  - The eigenvector associated with the larger eigenvalue points in the direction of fastest intensity change, and the other eigenvector is orthogonal to it

# Interpreting the eigenvalues

Classification of image points using eigenvalues of the second moment matrix:



# Errors in Lucas-Kanade

- The motion is large (larger than a pixel)
  - Iterative refinement
  - Coarse-to-fine estimation
  - Exhaustive neighborhood search (feature matching)
- A point does not move like its neighbors
  - Motion segmentation
- Brightness constancy does not hold
  - Exhaustive neighborhood search with normalized correlation

# Shi-Tomasi feature tracker

- Find good features using eigenvalues of second-moment matrix
  - Key idea: “good” features to track are the ones whose motion can be estimated reliably
- From frame to frame, track with Lucas-Kanade
  - This amounts to assuming a translation model for frame-to-frame feature movement
- Check consistency of tracks by *affine* registration to the first observed instance of the feature
  - Affine model is more accurate for larger displacements
  - Comparing to the first frame helps to minimize drift

# Tracking example

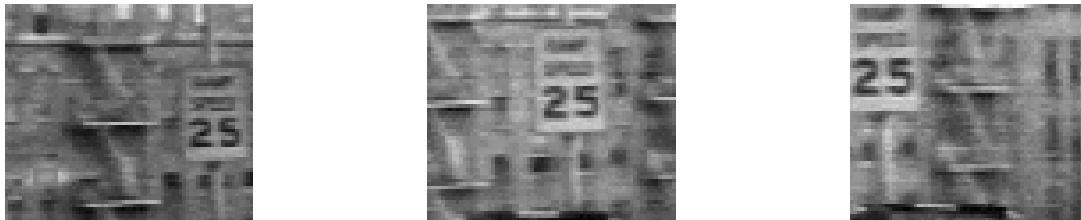


Figure 1: Three frame details from Woody Allen's *Manhattan*. The details are from the 1st, 11th, and 21st frames of a subsequence from the movie.

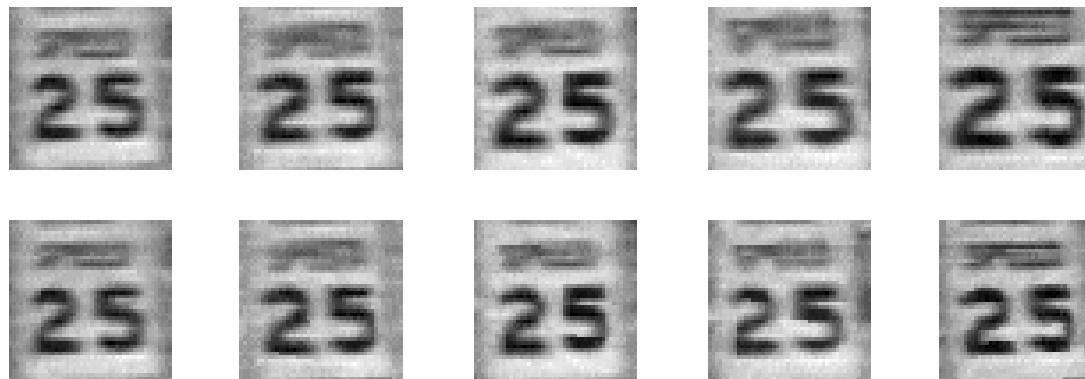


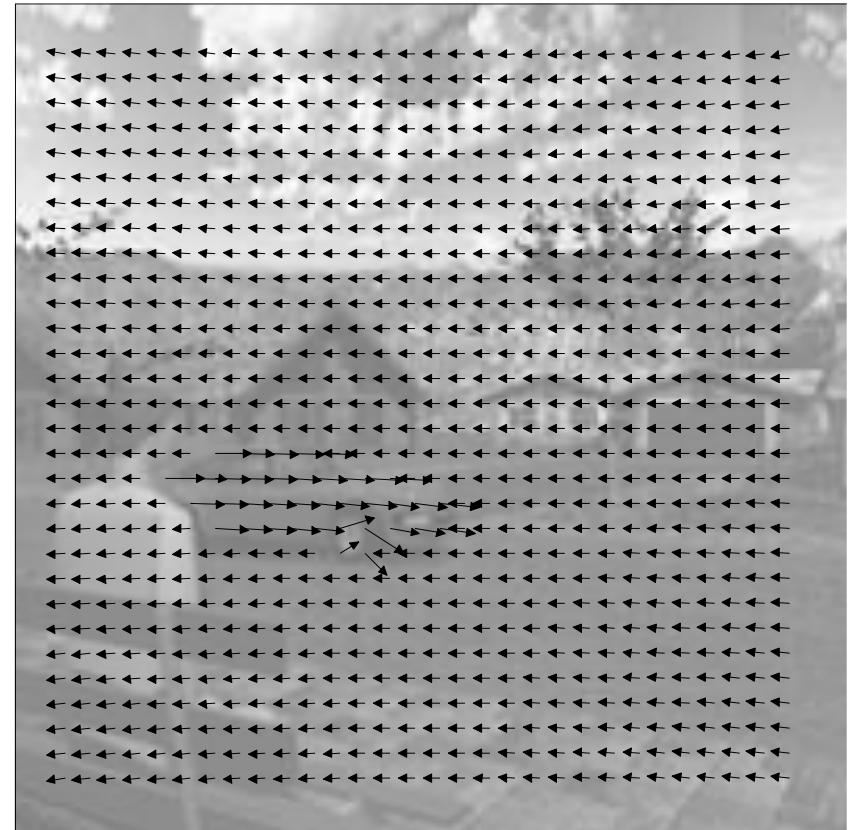
Figure 2: The traffic sign windows from frames 1,6,11,16,21 as tracked (top), and warped by the computed deformation matrices (bottom).

# Dense motion estimation

- For each pixel in an image, optical flow determines the corresponding pixel in the next image
- Can be cast as a variational problem



Image sequence  $I(x, y, t)$



Optical flow field  $(u, v)(x, y, t)$ .

Thomas Brox

# Horn-Schunck model

- Most basic model → good to start with
- Assumption I:  
Gray value constancy (optic flow constraint)

$$I_x u + I_y v + I_t = 0 \quad \text{reminder: subscripts denote partial derivatives}$$

Linearized version of

$$I(x + u, y + v, t + 1) - I(x, y, t) = 0$$

- Assumption II:  
Smoothness of optic flow field

$$|\nabla u|^2 + |\nabla v|^2 \rightarrow \min$$

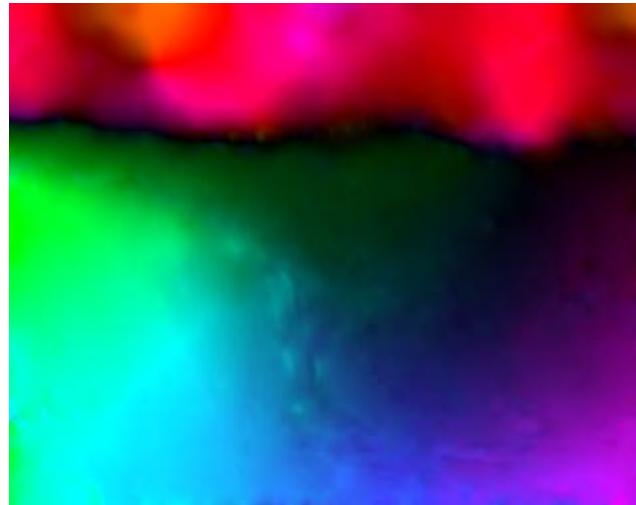
- Combined in an energy functional:

$$E(u, v) = \int_{\Omega} (I_x u + I_y v + I_t)^2 + \alpha(|\nabla u|^2 + |\nabla v|^2) dx dy$$

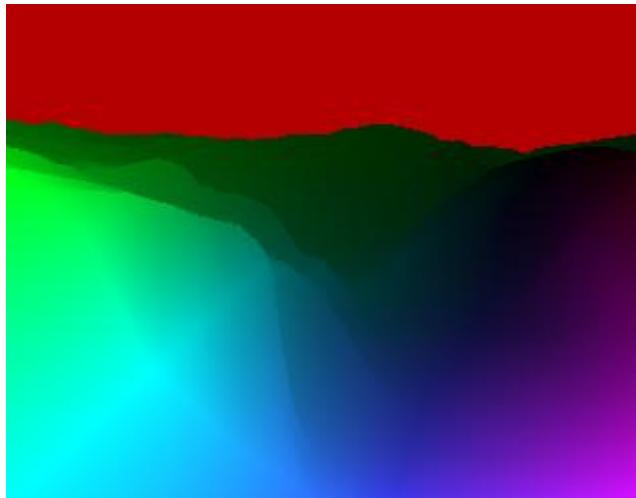
# Results with Horn-Schunck method



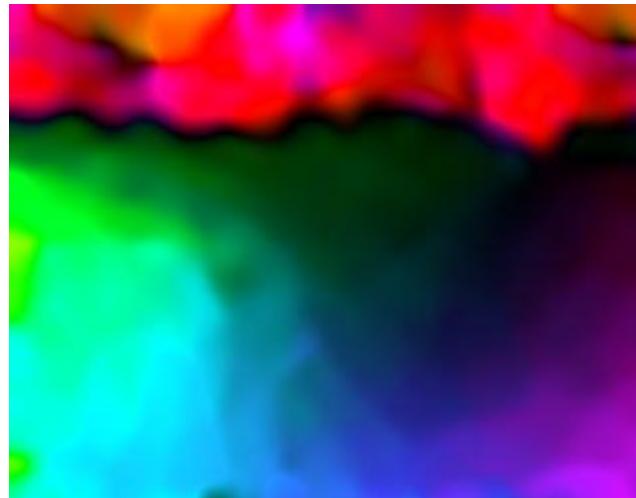
Input sequence



Horn-Schunck ( $7.17^\circ$ )



Ground truth



Lucas-Kanade ( $8.78^\circ$ )



## Motion discontinuities



## Occlusions



## Illumination changes



## Large displacements



## Really large displacements

# Incorporating correspondences from descriptor matching

$$\begin{aligned} E(\mathbf{w}(\mathbf{x})) = & \int \Psi(|I_2(\mathbf{x} + \mathbf{w}(\mathbf{x})) - I_1(\mathbf{x})|^2) d\mathbf{x} \\ & + \gamma \int \Psi(|\nabla I_2(\mathbf{x} + \mathbf{w}(\mathbf{x})) - \nabla I_1(\mathbf{x})|^2) d\mathbf{x} \\ & + \alpha \int \Psi(|\nabla u(\mathbf{x})|^2 + |\nabla v(\mathbf{x})|^2) d\mathbf{x} \end{aligned}$$

point correspondences from HOG descriptors

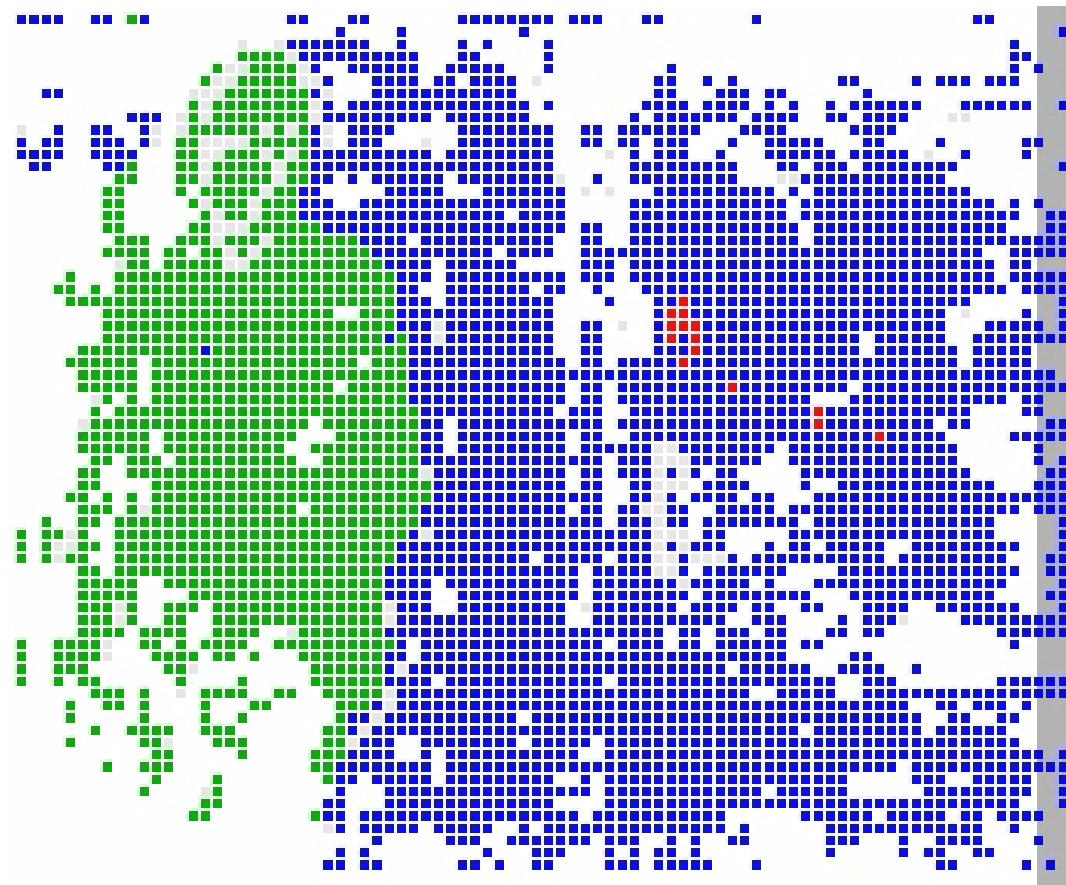
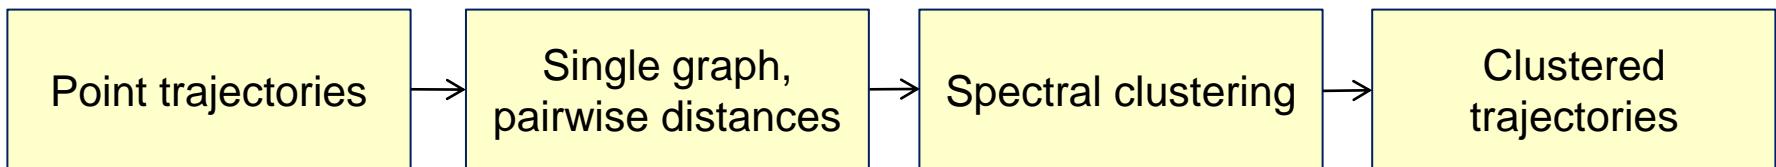
$$+ \beta \int \rho(\mathbf{x}) \Psi((u(\mathbf{x}) - u'(\mathbf{x}))^2 + (v(\mathbf{x}) - v'(\mathbf{x}))^2) d\mathbf{x}$$

Matching score

Robust function

Flow = correspondence vector

# Clustering of point trajectories (Brox-Malik 2010)



Thomas Brox

# UNIVERSITÄT BONN

