

Perspective, Mosaics and Panoramas

Mosaics and Panoramas

- Basic idea
- Registration
- Resampling
- Blending

Why Mosaic?

- Are you getting the whole picture?
 - Compact Camera FOV = $50 \times 35^\circ$



Slide from Brown & Lowe

Why Mosaic?


- Are you getting the whole picture?
 - Compact Camera FOV = $50 \times 35^\circ$
 - Human FOV = $200 \times 135^\circ$



Slide from Brown & Lowe

Why Mosaic?

- Are you getting the whole picture?
 - Compact Camera FOV = $50 \times 35^\circ$
 - Human FOV = $200 \times 135^\circ$
 - Panoramic Mosaic = $360 \times 180^\circ$



Slide from Brown & Lowe

Single vs. Multiple Viewpoint

- Single-viewpoint
 - Necessary for creating pure perspective images.
 - Many vision algorithms assume pinhole cameras.
 - Images that aren't perspective images look distorted.
- Multi-viewpoint
 - Cross-slit panoramas, etc.
 - necessary for scenes which cannot be captured from a single viewpoint

Image Mosaicing

- Register multiple images
- Blend



323

1021

9186

9219

359

1020

9185

9220

322

360

1019

9184

9221

321

361

1018

9183

9222

320

362

1017

9182

319

363


Single Center of Projection

- Take a sequence of images from the same position
 - Rotate the camera about its optical center
 - Where is optical center of thick lens?
- Compute transformation between second image and first
- Transform the second image to overlap with the first
- Blend the two together to create a mosaic
- If there are more images, repeat


- ...why don't we need the 3D geometry?

Image Reprojection

- The images are reprojected onto a common plane
- The mosaic is formed on this plane
- Mosaic is a *synthetic wide-angle camera*




A pencil of rays contains all views



Can generate any synthetic camera view
as long as it has **the same center of projection!**

Image reprojection

- How to relate two images from the same camera center?
- Images contain the same information along the same ray.
- Use 2D image warp instead of ray tracing.




Taxonomy of Projective Transformations

- Homogeneous coordinates: Specify 2D image points using three dimensions.

$$\begin{pmatrix} x \\ y \end{pmatrix} \rightarrow \begin{pmatrix} x \\ y \\ w \end{pmatrix}$$



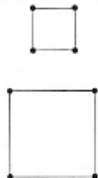

Points that differ only by scalar factor are equivalent.

Projective transformation:
$$\begin{pmatrix} x' \\ y' \\ w' \end{pmatrix} = \begin{pmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{pmatrix} \begin{pmatrix} x \\ y \\ w \end{pmatrix}$$




Taxonomy of Projective Transformations

$$\begin{pmatrix} x' \\ y' \\ w' \end{pmatrix} = \begin{pmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{pmatrix} \begin{pmatrix} x \\ y \\ w \end{pmatrix}$$

Group	Matrix	Distortion	Invariant properties
Projective 8 dof	$\begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix}$		Concurrency, collinearity, order of contact : intersection (1 pt contact); tangency (2 pt contact); inflections (3 pt contact with line); tangent discontinuities and cusps. cross ratio (ratio of ratio of lengths).
Affine 6 dof	$\begin{bmatrix} a_{11} & a_{12} & t_x \\ a_{21} & a_{22} & t_y \\ 0 & 0 & 1 \end{bmatrix}$		Parallelism, ratio of areas, ratio of lengths on collinear or parallel lines (e.g. midpoints), linear combinations of vectors (e.g. centroids). The line at infinity, l_∞ .
Similarity 4 dof	$\begin{bmatrix} sr_{11} & sr_{12} & t_x \\ sr_{21} & sr_{22} & t_y \\ 0 & 0 & 1 \end{bmatrix}$		Ratio of lengths, angle.
Euclidean 3 dof	$\begin{bmatrix} r_{11} & r_{12} & t_x \\ r_{21} & r_{22} & t_y \\ 0 & 0 & 1 \end{bmatrix}$		Length, area

Distortions under Central Projection



- **Similarity:** circle remains circle, square remains square
⇒ line orientation is preserved
- **Affine:** circle becomes ellipse, square becomes rhombus
⇒ parallel lines remain parallel
- **Projective:** imaged object size depends on distance from camera
⇒ parallel lines converge

Homography


- Projective mapping between any two PPs with the same center of projection
 - rectangle should map to arbitrary quadrilateral
 - parallel lines no longer parallel
 - but straight lines preserved
 - same as: project, rotate, reproject
- called Homography

$$\begin{pmatrix} x' \\ y' \\ w \end{pmatrix} = \begin{pmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{pmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

$$\mathbf{p}' \quad \mathbf{H} \quad \mathbf{p}$$

To apply a homography \mathbf{H}


- Compute $\mathbf{p}' = \mathbf{H}\mathbf{p}$ (regular matrix multiply)
- Convert \mathbf{p}' from homogeneous to image coordinates



Removing Projective Distortion



a



b

$$\begin{pmatrix} x' \\ y' \\ w' \end{pmatrix} = \begin{pmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{pmatrix} \begin{pmatrix} x \\ y \\ \mathbf{1} \end{pmatrix}$$

Projective transformation in inhomogeneous form

$$x'^{(2D)} = \frac{x'}{w'} = \frac{h_{11}x + h_{12}y + h_{13}}{h_{31}x + h_{32}y + h_{33}} \quad y'^{(2D)} = \frac{y'}{w'} = \frac{h_{21}x + h_{22}y + h_{23}}{h_{31}x + h_{32}y + h_{33}}$$

4 general point correspondences $(x, y \rightarrow x'^{(2D)}, y'^{(2D)})$ on the planar facade lead to eight linear equations of the type

$$\begin{aligned} x'^{(2D)}(h_{31}x + h_{32}y + h_{33}) &= h_{11}x + h_{12}y + h_{13} \\ y'^{(2D)}(h_{31}x + h_{32}y + h_{33}) &= h_{21}x + h_{22}y + h_{23} \end{aligned}$$

Sufficient to solve for H up to multiplicative factor

Getting it numerically stable is not trivial! **Best use proven libraries like OpenCV**

Panoramic Mosaicing

Rotation about camera center: homography

- choose one image as reference
 - compute homography to map neighboring image to reference image plane
 - projectively warp image, add to reference plane
 - repeat for all images
- ⇒ bow tie shape





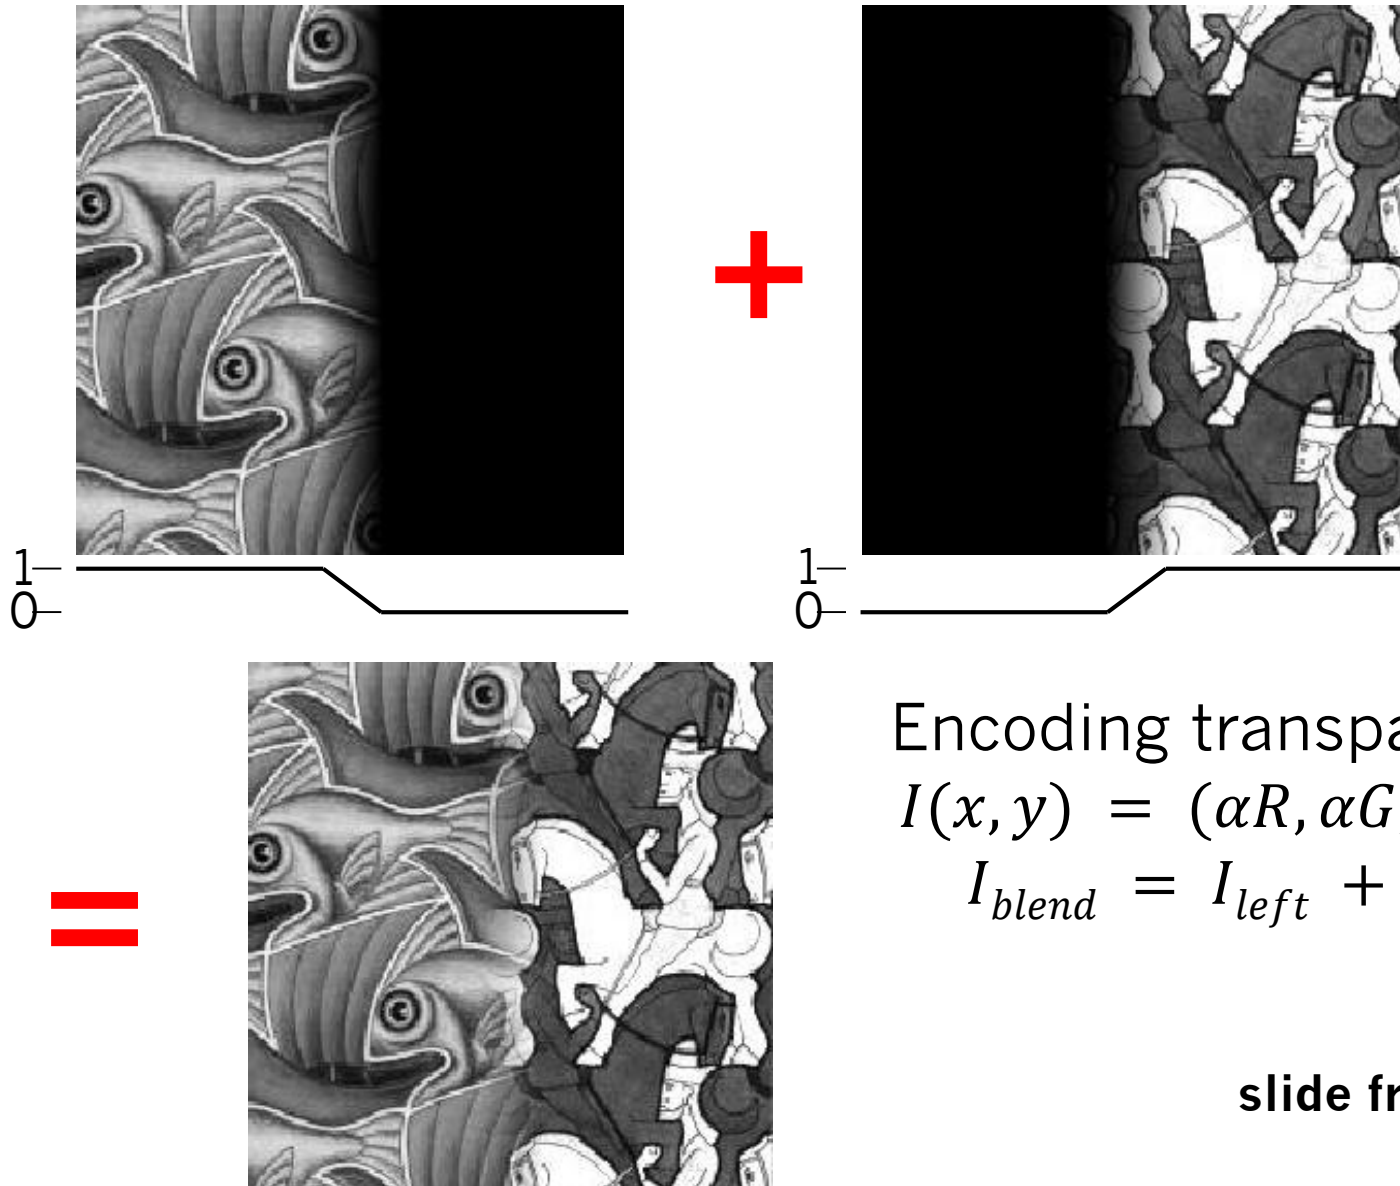


Image Blending



slide from Efros




Feathering



Encoding transparency
 $I(x, y) = (\alpha R, \alpha G, \alpha B, \alpha)$
 $I_{blend} = I_{left} + I_{right}$




slide from Efros

Effect of Window Size




slide from Efras

Effect of Window Size



slide from Efros

Good Window Size




“Optimal” Window: smooth but not ghosted

[Slide: Efros]

What is the Optimal Window?

- To avoid seams
 - $\text{window} \geq \text{size of largest prominent feature}$
- To avoid ghosting
 - $\text{window} \leq 2 * \text{size of smallest prominent feature}$
- Natural to cast this in the *Fourier domain*
 - $\text{largest frequency} \leq 2 * \text{size of smallest frequency}$
 - do blending in different frequency bands

What if the Frequency Spread is Wide



- Idea (Burt and Adelson)
 - Compute $F_{\text{left}} = \text{FFT}(I_{\text{left}})$, $F_{\text{right}} = \text{FFT}(I_{\text{right}})$
 - Decompose Fourier image into octaves (bands)
 - $F_{\text{left}} = F_{\text{left}}^1 + F_{\text{left}}^2 + \dots$
 - Feather corresponding octaves F_{left}^i with F_{right}^i
 - Can compute inverse FFT and feather in spatial domain
 - Sum feathered octave images in frequency domain
- Better implemented in *spatial domain*

What does blurring take away?



original

What does blurring take away?



smoothed (5x5 Gaussian)


High-Pass Filter



smoothed – original


Image Pyramids

Idea: Represent $N \times N$ image as a “pyramid” of $1 \times 1, 2 \times 2, 4 \times 4, \dots, 2^k \times 2^k$ images (assuming $N = 2^k$)



Mipmap or precursor of wavelets

Image Sub-sampling




1/4



1/8

Throw away every other row and column to create a 1/2 size image

Image Sub-sampling



1/2



1/4 (2x zoom)




1/8 (4x zoom)

Why does this look so bad?

Gaussian Pyramid Construction

- Repeat
 - Filter
 - Subsample




filter kernel,
e.g., discrete
Gaussian
with sigma=1

- Until minimum resolution reached

- Good filter kernel: $\frac{1}{256} \begin{bmatrix} 1 \\ 4 \\ 6 \\ 4 \\ 1 \end{bmatrix} [1 \ 4 \ 6 \ 4 \ 1] = \frac{1}{256} \begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$

- Whole pyramid is only 4/3 the size of the original image!

Gaussian pre-filtering



Gaussian $1/2$



G $1/4$



G $1/8$

Solution: filter the image, *then* subsample

- Filter size should double for each $1/2$ size reduction.

Subsampling with Gaussian pre-filtering



Gaussian 1/2



G 1/4




G 1/8

Solution: filter the image, *then* subsample

- Filter size should double for each $\frac{1}{2}$ size reduction.

Compare with...



$1/2$




$1/4$ (2x zoom)



$1/8$ (4x zoom)

Band-pass filtering

Gaussian Pyramid (low-pass images)




- Laplacian Pyramid (subband images)
- Created from Gaussian pyramid by subtraction

Laplacian Pyramid


- How can we reconstruct (collapse) this pyramid into the original image?

Original
image

Need this!




Pyramid Blending





Left pyramid

blend



Right pyramid

Pyramid Blending



(d)




(h)




(l)

Laplacian
level
4



(c)




(g)




(k)


Laplacian
level
2



(b)




(f)




(j)


Laplacian
level
0



(a)



(e)



(i)

left pyramid

right pyramid

blended pyramid

Simplification: Two-band Blending

- Brown & Lowe, 2003
 - Only use two bands: high freq. and low freq.
 - Blends low freq. smoothly
 - Blend high freq. with no smoothing: use binary mask



Still Some Artifacts Left...

- Ghosting—objects move in the scene.
- Differing exposures between images.
 - Pyramid blending does not solve this.

De-Ghosting

- In regions with differences don't blend – choose one.



(A)



[Uyttendaele et al. 2001]

Gradient domain blending

- In pyramid blending, we decomposed our image into 2nd derivatives (Laplacian) and a low-res image
- Let us now look at 1st derivatives (gradients):
- No need for low-res image
 - captures everything (up to a constant)
 - easy to deal with low-frequency differences
- Idea:
 - Differentiate
 - Blend
 - Reintegrate

Gradient domain blending (2D)

- Trickier in 2D:
 - Take partial derivatives dx and dy (gradient field)
 - Modify them (smooth, blend, feather, etc)
- Reintegrate
 - But now $\text{integral}(dx)$ might not equal $\text{integral}(dy)$
- Find the most agreeable solution
 - Equivalent to solving Poisson equation
 - Can use FFT, deconvolution, multigrid solvers, etc.

1. Given:

- a. intended gradient field: $\nabla_x I(x, y), \nabla_y I(x, y)$
- b. boundary conditions: $I_{\text{fix}}(x, y)_j$
fixed color at a couple of pixels $(x, y)_j$

2. Assemble gradient vector v and corresponding matrix A . Goal image g (as vector) $\rightarrow Ag = v$

- a. constraint vector

$$v = \begin{pmatrix} \nabla_x I \\ \nabla_y I \\ I_{\text{fix}} \end{pmatrix} \quad \begin{array}{l} \text{(given gradients)} \\ \text{(given colors)} \end{array}$$

Gradient domain fusion

- 2. Assemble gradient vector v and corresponding matrix A . Goal image g (as vector) $\rightarrow Ag = v$
 - b. assemble matrix A according to

$$g(x+1, y) - g(x, y) = \nabla_x I(x, y)$$

$$g(x, y+1) - g(x, y) = \nabla_y I(x, y)$$

structure similar to

$$A \approx \begin{pmatrix} -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ 0 & -1 & 1 & 0 \\ 0 & -1 & 0 & 1 \end{pmatrix} \quad \text{always one „+1“ and „-1“ per row}$$

Gradient domain fusion

- 2. Assemble gradient vector v and corresponding matrix
A. Goal image g (as vector) $Ag = v$
c. Extend matrix A with boundary conditions

$$g(x, y) = I_{\text{fix}}(x, y)$$

structure
similar to

$$A \approx \begin{pmatrix} -1 & 1 & 0 & 0 \\ -1 & 0 & 1 & 0 \\ 0 & -1 & 1 & 0 \\ 0 & -1 & 0 & 1 \\ \hline 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

always one „+1“ per row

Gradient domain fusion

- 3. Solve for g : $Ag = v$
 - As gradients might be contradicting, use least-squares solution

$$A^T Ag = A^T v \quad (\text{normal equation})$$

$$\Rightarrow g = (A^T A)^{-1} A^T v$$

- Implicitly done by QR decomposition, SVD, etc.
 - Alternatively, use conjugate gradient methods

e.g. <http://www.eecs.berkeley.edu/~demmel/cs267/lecture24/lecture24.html>


Comparisons [Levin et al 2004]



Pyramid blending



Pyramid blending on the gradients



Feathering

More common: Poisson editing

- Use **Laplacian operator** $\Delta = \nabla^2$ instead of $\nabla_{x,y}$

- Discrete Laplacian: convolution with 2D kernel $\begin{bmatrix} -1 & -1 & -1 \\ -1 & +4 & -1 \\ -1 & -1 & -1 \end{bmatrix}$

- This leads to convolution matrix:

$$M_{\Delta} = \begin{bmatrix} \begin{bmatrix} +4 & -1 & -1 \\ -1 & \ddots & -1 \\ & -1 & +4 \end{bmatrix} & \begin{bmatrix} -1 & & \\ & \ddots & \\ & & -1 \end{bmatrix} & & \\ \begin{bmatrix} -1 & & \\ & \ddots & \\ & & -1 \end{bmatrix} & & \begin{bmatrix} -1 & & \\ & \ddots & \\ & & -1 \end{bmatrix} & \\ & \begin{bmatrix} -1 & & \\ & \ddots & \\ & & -1 \end{bmatrix} & \begin{bmatrix} +4 & -1 & -1 \\ -1 & \ddots & -1 \\ & -1 & +4 \end{bmatrix} \end{bmatrix}$$

read, “output pixel $(i,j) = 4 \cdot \text{input pixel } (i,j) - \sum \text{input pixels } (i \pm 1, j \pm 1)$ ”

For a $w \times h$ image, M has $w \times w$ blocks of $h \times h$ elements each

e.g., 307200×307200 matrix for a 640×480 image.

Use sparse matrices, or die of memory exhaustion

Poisson editing continued

- Construct M_Δ
- Compute $v_{1,2} = M_\Delta I_{1,2}$
- Blend $v_{1,2}$ using method of choice to obtain combined v
- Reconstruct I by solving $M_\Delta I = v$

Add constraints as before:

- Constraint matrix M_c with 1 entry per row
- Constraint vector v_c with prescribed pixel values
- Solve $\begin{bmatrix} M_\Delta \\ M_c \end{bmatrix} I = \begin{bmatrix} v \\ v_c \end{bmatrix}$ in least-squares sense

Panoramas and Image blending

- A multiresolution spline with application to image mosaics
P. J. Burt, E. H. Adelson.
ACM Transactions on Graphics. 2(4), pp. 217-236, 1983.
- Recognising Panoramas. M. Brown and D. G. Lowe. In Proceedings of the *9th International Conference on Computer Vision (ICCV2003)*, pages 1218-1225, Nice, France, 2003.
- Multiple View Geometry in Computer Vision (First Edition). Richard Hartley and Andrew Zisserman, Cambridge University Press, March 2001
- Seamless Image Stitching in the Gradient Domain. A. Levin, A. Zomet, S. Peleg and Y. Weiss, In Proc. ECCV 2004.
- Interactive Digital Photomontage. A. Agarwala, M. Dontcheva, M. Agrawala, S. Drucker, A. Colburn, B. Curless, D. Salesin, M. Cohen, In *SIGGRAPH 2004*.